

**Lecture 1:**

# **Course Introduction + Review of Throughput Hardware Concepts**

---

**Visual Computing Systems  
Stanford CS348V : Winter 2018**


**Hi!**

**Me:**



**Prof. Kayvon**

**Your TA!**




**Raj Setaluri**

# Visual computing applications




## 2D/3D graphics



## Computational photography and image processing



## Understanding the contents of images and videos



# **Visual Computing Systems**

## **— Some History**

**(why I get so excited about this topic)**





Ivan Sutherland's Sketchpad on MIT TX-2 (1962)

# The frame buffer


Shoup's SuperPaint (PARC 1972-73)

16 2K shift registers (640 x 486 x 8 bits)




# The frame buffer


Shoup's SuperPaint (PARC 1972-73)





16 2K shift registers (640 x 486 x 8 bits)



# Xerox Alto (1973)



Bravo (WYSIWYG)



TI 74181 ALU

# Goal: render everything you've ever seen


“Road to Pt. Reyes”  
LucasFilm (1983)



# Pixar's Toy Story (1995)




**“We take an average of three hours to draw a single frame on the fastest computer money can buy.”**  
- Steve Jobs



UNC Pixel Planes (1981), computation-enhanced frame buffer

# Ed Clark's Geometry Engine (1982)

ASIC for geometric transforms  
used in real-time graphics.



**Real-time (30 fps) on a NVIDIA Titan X**




**Unreal Engine Kite Demo (Epic Games 2015)**



**NVIDIA Titan X Pascal GPU (2017)**  
**(~ 12 TFLOPs fp32)**

**~ ASCI Q (top US supercomputer circa 2002)**

# Modern GPU: heterogeneous multi-core



Multi-threaded, SIMD cores


Custom circuits for key graphics arithmetic

Custom circuits for HW-assisted graphics-specific DRAM compression

HW logic for scheduling work onto these resources

# Domain-specific languages for heterogeneous computing

## OpenGL Graphics Pipeline (circa 2007)



The OpenGL<sup>TM</sup> Graphics System:  
A Specification  
(Version 1.0)


Mark Segal  
Kurt Akeley

*Editor:*  
Chris Frazier

Version 1.0 - 1 July 1994

# Domain-specific languages for heterogeneous computing

## OpenGL Graphics Pipeline (circa 2007)



```
uniform sampler2D myTexture; ] read-only  
uniform float3 lightDir; ] global variables  
varying vec3 norm; ]  
varying vec2 uv; ] "per-element" inputs
```


```
void myFragmentShader()  
{  
    vec3 kd = texture2D(myTexture, uv);  
    kd *= clamp(dot(lightDir, norm), 0.0,  
1.0);  
    return vec4(kd, 1.0);  
}
```

“fragment shader”  
(a.k.a kernel function mapped onto  
input fragment stream)

per-element output:  
RGBA surface color at pixel

# Emerging state-of-the-art visual computing systems today...

- Intelligent cameras in smartphones
- Cloud servers (“infinite” computing and storage at your disposal as a service)
- Proliferation of specialized compute accelerators
  - For image processing, machine learning
- Proliferation of high-resolution image sensors...



# Capturing pixels to communicate

Ingesting/serving  
the world's photos



2B photo uploads and shares  
per day across Facebook sites  
(incl. Instagram+WhatsApp)  
[FB2015]


Ingesting/streaming  
world's video



PSY - GANGNAM STYLE (강남스타일) M/V




Youtube 2015: 300 hours  
uploaded per minute [Youtube]

Cisco VNI projection:  
80-90% of 2019 internet  
traffic will be video.  
(64% in 2014)




# Richer content: beyond a single image

- Example: Apple's "Live Photos"
- Each photo is not only a single frame, but a few seconds of video before and after the shutter is clicked




# Facebook Live



# VR output





**Example: Google's JumpVR video**  
**Input stream: 16 4K GoPro cameras**

**Register + 3D align video stream (on edge device)**  
**Broadcast encoded video stream across**  
**the country to millions of viewers**



# High resolution, multi-camera Facebook Surround 360 VR video





2048 x 2048 PointGrey Camera @ 30 FPS

14 cameras

8K x 8K stereo panorama output

# VR: high resolution requirements



**Future “retina” VR display:**  
**57 ppd covering 180°**  
**= 10K x 10K display per eye**  
**= 200 MPixel**

**RAW data rate @ 120Hz  $\approx$  72 GB/sec**

**iPhone 6: 4.7 in “retina” display:**

**1.3 MPixel**


**326 ppi  $\rightarrow$  57 ppd**

# Enhancing communication: understanding images to improve acquired content


AutoEnhance:




Photo “fix up” [Hayes 2007]



Portrait Mode:



My bad vacation photo




Part to fix




Similar photos others  
have taken

Fixed!

# On every vehicle: analyzing images for robot navigation




# High-resolution video (moving camera)





[Image Credit: Kundu et al. 2016]

# NVIDIA Drive PX



**Tegra X1 (1 TFlop fp16 at 1GHz)**

# On every corner: analyzing images for urban efficiency



**“Managing urban areas has become one of the most important development challenges of the 21st century. Our success or failure in building sustainable cities will be a major factor in the success of the post-2015 UN development agenda.”**



**- UN Dept. of Economic and Social Affairs**

# High resolution (static camera)




# Sensing human social interactions

[Joo 2015]




**CMU Panoptic Studio**  
**480 video cameras (640 x 480 @ 25fps)**  
**116 GPixel video sensor**  
**(2.9 TPixel /sec)**

# Capturing social interactions



[Courtesy Yaser Sheikh, Tomas Simon, Hanbyul Joo]

# Capturing social interactions



[Courtesy Yaser Sheikh, Tomas Simon, Hanbyul Joo]

# On every human: analyzing egocentric images to augment humans




Gwangjang Market (Seoul)

# AR requires low-latency localization and scene object recognition






# Smart headlight system



# Seeing clearly through precipitation



Idea: Stream Light Between Snowflakes

Goal: High Light Throughput and Accuracy

Future challenge: recording and analyzing the world's visual information, so **computers** can understand and reason about it

# Capturing everything about the visual world

To understand people

To understand the world around vehicles/drones

To understand cities

Mobile

Continuous (always on)

Exceptionally high resolution


**Capture for computers to analyze, not humans to watch**

# What is this course about?

1. Understanding the characteristics of important visual computing workloads
2. Understanding techniques used to achieve efficient system implementations

## VISUAL COMPUTING WORKLOADS

Algorithms for 3D graphics, image processing, compression, etc.



Parallelism  
Exploiting locality  
Minimizing communication

## MACHINE ORGANIZATION



High-throughput hardware designs:  
Parallel, heterogeneous, specialized

DESIGN OF PROGRAMMING ABSTRACTIONS  
FOR VISUAL COMPUTING  
choice of programming primitives  
level of abstraction



# In other words

It is about understanding the **fundamental structure** of problems in the visual computing domain, and then leveraging that understanding to...

To design more efficient algorithms

To build the most efficient hardware to run these applications

To design the right programming systems to make developing new applications simpler, more productive, and highly performant

# **Course Logistics**

# Logistics

- **Course web site:**
  - <http://graphics.stanford.edu/courses/cs348v-18-winter>
- **All announcements will go out via Piazza**
  - <https://piazza.com/class/jc1f626cfne6r6>
- **Kayvon's office hours: Tuesday after class, or by appt.**

# Expectations of you

## ■ 20% participation

- There will be ~1 assigned paper reading per class
- Everyone is expected to come to class and participate in discussions based on readings
- You are encouraged discuss papers and or my lectures on the course discussion board.
- If you form a weekly course reading/study group, I will buy Pizza for said group.

## ■ 30% mini-assignments (3 short programming assignments)

- Assignment 1: analyze parallel program performance on a multi-core CPU
- Assignment 2: implement and optimize a basic RAW image processing pipeline
- Assignment 3: optimize performance of a modern DNN module

## ■ 20% 1 take-home “exam”

## ■ 30% self-selected final project

- I suggest you start thinking about projects now (can be teams of up to two)

# **Major course themes/topics**

## **Part 1: High Efficiency Image and Video Processing**

**Overview of a Modern Digital Camera Processing Pipeline**

**Image Processing Algorithms You Should Know**

**Efficiently Scheduling Image Processing Algorithms on Parallel Hardware**

**Specialized Hardware for Image Processing**

**Lossy Image (JPEG) and Video (H.264) Video Compression**

**Video Processing/Synthesis for Virtual Reality Display**

## **Part 2: Accelerating Deep Learning for Computer Vision (from a systems perspective)**

**Workload Characteristics of DNN Inference for Image Analysis**

**Scheduling and Algorithms for Parallel DNN Training at Scale**

**A Case Study of Algorithmic Optimizations for Object Detection**

**Leveraging Task-Specific DNN Structure for Improving Performance and Accuracy**

**Hardware Accelerators for DNN Inference**

**Design Space of Dataflow Programming Abstractions for Deep Learning**

**Enhancing Efficiency Through Model Specialization (in particular for video)**

**Efficient Inference at Datacenter Scale**

# Algorithmic innovation in image classification

Improving *accuracy-per-unit cost* using better DNN designs?

2014 → 2017 ~ 25x improvement in cost at similar accuracy

	ImageNet Top-1 Accuracy	Num Params	Cost/image (MADDs)	
<b>VGG-16</b>	<b>71.5%</b>	<b>138M</b>	<b>15B</b>	<b>[2014]</b>
<b>GoogleNet</b>	<b>70%</b>	<b>6.8M</b>	<b>1.5B</b>	<b>[2015]</b>
<b>ResNet-18</b>	<b>73% *</b>	<b>11.7M</b>	<b>1.8B</b>	<b>[2016]</b>
<b>MobileNet-224</b>	<b>70.5%</b>	<b>4.2M</b>	<b>0.6B</b>	<b>[2017]</b>

\* 10-crop results (ResNet 1-crop results are similar to other DNNs in this table)

# **Major course themes/topics**

## **Part 3: The GPU Accelerated 3D Graphics Pipeline**

**Real-Time 3D Graphics Pipeline Architecture**

**Hardware Acceleration of Z-Buffering and Texturing**

**Scheduling the Graphics Pipeline onto a GPU**

**Domain Specific Languages for Shading**


# **Review: key principles of modern throughput computing hardware**

# Review concepts

- **What are these design concepts, and what problem/goals do they address?**
  - **Muti-core processing**
  - **SIMD processing**
  - **Hardware multi-threading**
- **What is the motivation for specialization via:**
  - **Multiple types of processors (e.g., CPUs, GPUs)**
  - **Custom hardware units (ASIC)**
- **What is memory bandwidth a major constraint when mapping applications to modern computer systems?**

# Let's crack open a modern smartphone

Samsung Galaxy S7 phone with  
Qualcomm Snapdragon 820 processor



# **Multi-core processing**

# Review: what does a processor do?


**It runs programs!**

**Processor executes instruction  
referenced by the program counter  
(PC)**

**(executing the instruction will modify machine  
state: contents of registers, memory, CPU  
state, etc.)**

**Move to next instruction ...**


**Then execute it...**

**PC** 

**And so on...**


```
_main:  
100000f10: pushq   %rbp  
100000f11: movq %rsp, %rbp  
100000f14: subq $32, %rsp  
100000f18: movl $0, -4(%rbp)  
100000f1f: movl %edi, -8(%rbp)  
100000f22: movq %rsi, -16(%rbp)  
100000f26: movl $1, -20(%rbp)  
100000f2d: movl $0, -24(%rbp)  
100000f34: cmpl $10, -24(%rbp)  
100000f38: jge 23 <_main+0x45>  
100000f3e: movl -20(%rbp), %eax  
100000f41: addl -20(%rbp), %eax  
100000f44: movl %eax, -20(%rbp)  
100000f47: movl -24(%rbp), %eax  
100000f4a: addl $1, %eax  
100000f4d: movl %eax, -24(%rbp)  
100000f50: jmp -33 <_main+0x24>  
100000f55: leaq 58(%rip), %rdi  
100000f5c: movl -20(%rbp), %esi  
100000f5f: movb $0, %al  
100000f61: callq 14  
100000f66: xorl %esi, %esi  
100000f68: movl %eax, -28(%rbp)  
100000f6b: movl %esi, %eax  
100000f6d: addq $32, %rsp  
100000f71: popq %rbp  
100000f72: retq
```

# Executing an instruction stream




# Executing an instruction stream

My very simple processor: executes one instruction per clock




# Executing an instruction stream

# My very simple processor: executes one instruction per clock



# Executing an instruction stream

My very simple processor: executes one instruction per clock



# **Quick aside: Instruction-level parallelism and superscalar execution**

# Instruction level parallelism (ILP) example

$$a = x*x + y*y + z*z$$

Consider the following program:

```
// assume r0=x, r1=y, r2=z


mul r0, r0, r0
mul r1, r1, r1
mul r2, r2, r2
add r0, r0, r1
add r3, r0, r2

// now r3 stores value of program variable 'a'
```

This program has five instructions, so it will take five clocks to execute, correct?  
Can we do better?

# ILP example

$$a = x*x + y*y + z*z$$



# Superscalar execution

$$a = x*x + y*y + z*z$$

// assume r0=x, r1=y, r2=z

1. mul r0, r0, r0
2. mul r1, r1, r1
3. mul r2, r2, r2
4. add r0, r0, r1
5. add r3, r0, r2

// r3 stores value of variable 'a'

**Superscalar execution:** processor automatically finds **independent instructions** in an instruction sequence and executes them in **parallel** on multiple execution units!

In this example: instructions 1, 2, and 3 **can be** executed in parallel  
(on a superscalar processor that determines that the lack of dependencies exists)

But instruction 4 must come after instructions 1 and 2

And instruction 5 must come after instruction 4

# Superscalar execution


Program: computes sin of input  $x$  via Taylor expansion

```
void sinx(int N, int terms, float x)
{
    float value = x;
    float numer = x * x * x;
    int denom = 6; // 3!
    int sign = -1;

    for (int j=1; j<=terms; j++)
    {
        value += sign * numer / denom;
        numer *= x * x;
        denom *= (2*j+2) * (2*j+3);
        sign *= -1;
    }

    return value;
}
```

My single core, superscalar processor:  
executes up to two instructions per clock  
from a single instruction stream.



Independent operations in  
instruction stream

(They are detected by the processor  
at run-time and may be executed in  
parallel on execution units 1 and 2)

**Now consider a program that computes  
the sine of many numbers...**

# Example program



**Compute  $\sin(x)$  using Taylor expansion:**  $\sin(x) = x - x^3/3! + x^5/5! - x^7/7! + \dots$   
**for each element of an array of N floating-point numbers**

```
void sinx(int N, int terms, float* x, float* result)
{
    for (int i=0; i<N; i++)
    {
        float value = x[i];
        float numer = x[i] * x[i] * x[i];
        int denom = 6; // 3!
        int sign = -1;

        for (int j=1; j<=terms; j++)
        {
            value += sign * numer / denom;
            numer *= x[i] * x[i];
            denom *= (2*j+2) * (2*j+3);
            sign *= -1;
        }


        result[i] = value;
    }
}
```

# Multi-core: process multiple instruction streams in parallel




**Sixteen cores, sixteen simultaneous instruction streams**

# Multi-core examples




**Intel "Skylake" Core i7 quad-core CPU  
(2015)**




**NVIDIA GP104 (GTX 1080) GPU  
20 replicated ("SM") cores  
(2016)**

# More multi-core examples




**Intel Xeon Phi "Knights Landing" 76-core CPU  
(2015)**



**Apple A9 dual-core CPU  
(2015)**

# **SIMD processing**

# Add ALUs to increase compute capability



**Idea #2:**

**Amortize cost/complexity of managing an instruction stream across many ALUs**

**SIMD processing**

**Single instruction, multiple data**

**Same instruction broadcast to all ALUs**

**Executed in parallel on all ALUs**

# Scalar program

```
void sinx(int N, int terms, float* x, float* result)
{
    for (int i=0; i<N; i++)
    {
        float value = x[i];
        float numer = x[i] * x[i] * x[i];
        int denom = 6; // 3!
        int sign = -1;

        for (int j=1; j<=terms; j++)
        {
            value += sign * numer / denom;
            numer *= x[i] * x[i];
            denom *= (2*j+2) * (2*j+3);
            sign *= -1;
        }

        result[i] = value;
    }
}
```

**Original compiled program:**

**Processes one array element using scalar instructions on scalar registers (e.g., 32-bit floats)**

ld r0, addr[r1]
mul r1, r0, r0
mul r1, r1, r0
...
...
...
...
...
...
...
...
st addr[r2], r0

# Vector program (using AVX intrinsics)

```
#include <immintrin.h>
void sinx(int N, int terms, float* x, float* sinx)
{
    float three_fact = 6; // 3!
    for (int i=0; i<N; i+=8)
    {
        __m256 origx = _mm256_load_ps(&x[i]);
        __m256 value = origx;
        __m256 numer = _mm256_mul_ps(origx, _mm256_mul_ps(origx, origx));
        __m256 denom = _mm256_broadcast_ss(&three_fact);
        int sign = -1;

        for (int j=1; j<=terms; j++)
        {
            // value += sign * numer / denom
            __m256 tmp =
                _mm256_div_ps(_mm256_mul_ps(_mm256_broadcast_ss(sign), numer), denom);
            value = _mm256_add_ps(value, tmp);


            numer = _mm256_mul_ps(numer, _mm256_mul_ps(origx, origx));
            denom = _mm256_mul_ps(denom, _mm256_broadcast_ss((2*j+2) * (2*j+3)));
            sign *= -1;
        }
        _mm256_store_ps(&sinx[i], value);
    }
}
```

<b>vloadps</b>	<b>xmm0, addr[r1]</b>
<b>vmulps</b>	<b>xmm1, xmm0, xmm0</b>
<b>vmulps</b>	<b>xmm1, xmm1, xmm0</b>
...	
...	
...	
...	
...	
...	
<b>vstoreps</b>	<b>addr[xmm2], xmm0</b>

## Compiled program:

**Processes eight array elements simultaneously using vector instructions on 256-bit vector registers**

# 16 SIMD cores: 128 elements in parallel



16 cores, 128 ALUs, 16 simultaneous instruction streams

# Data-parallel expression

(in Kayvon's fictitious data-parallel language)

```
void sinx(int N, int terms, float* x, float* result)
{
    // declare independent loop iterations
    forall (int i from 0 to N-1)
    {
        float value = x[i];
        float numer = x[i] * x[i] * x[i];
        int denom = 6; // 3!
        int sign = -1;


        for (int j=1; j<=terms; j++)
        {
            value += sign * numer / denom
            numer *= x[i] * x[i];
            denom *= (2*j+2) * (2*j+3);
            sign *= -1;
        }

        result[i] = value;
    }
}
```

**Compiler understands loop iterations are independent, and that same loop body will be executed on a large number of data elements.**

**Abstraction facilitates automatic generation of both multi-core parallel code, and vector instructions to make use of SIMD processing capabilities within a core.**

# What about conditional execution?



(assume logic below is to be executed for each element in input array 'A', producing output into the array 'result')


<unconditional code>

```
float x = A[i];  
  
if (x > 0) {  
    float tmp = exp(x,5.f);  
    tmp *= kMyConst1;  
    x = tmp + kMyConst2;  
} else {  
    float tmp = kMyConst1;  
    x = 2.f * tmp;  
}
```

<resume unconditional code>

```
result[i] = x;
```

# What about conditional execution?



(assume logic below is to be executed for each element in input array 'A', producing output into the array 'result')


<unconditional code>

```
float x = A[i];  
  
if (x > 0) {  
    float tmp = exp(x,5.f);  
    tmp *= kMyConst1;  
    x = tmp + kMyConst2;  
} else {  
    float tmp = kMyConst1;  
    x = 2.f * tmp;  
}
```

<resume unconditional code>

```
result[i] = x;
```

# Mask (discard) output of ALU



(assume logic below is to be executed for each element in input array 'A', producing output into the array 'result')


<unconditional code>

```
float x = A[i];  
  
if (x > 0) {  
    float tmp = exp(x, 5.f);  
    tmp *= kMyConst1;  
    x = tmp + kMyConst2;  
} else {  
    float tmp = kMyConst1;  
    x = 2.f * tmp;  
}
```

<resume unconditional code>

```
result[i] = x;
```

# After branch: continue at full performance



(assume logic below is to be executed for each element in input array 'A', producing output into the array 'result')


<unconditional code>

```
float x = A[i];  
  
if (x > 0) {  
    float tmp = exp(x, 5.f);  
    tmp *= kMyConst1;  
    x = tmp + kMyConst2;  
} else {  
    float tmp = kMyConst1;  
    x = 2.f * tmp;  
}
```

<resume unconditional code>

```
result[i] = x;
```

# Example: eight-core Intel Xeon E5-1660 v4




**8 cores**  
**8 SIMD ALUs per core**  
**(AVX2 instructions)**

**490 GFLOPs (@3.2 GHz)**  
**(140 Watts)**

\* Showing only AVX math units, and fetch/decode unit for AVX (additional capability for integer math)


# Example: NVIDIA GTX 1080 GPU



**20 cores ("SMs")**

**128 SIMD ALUs per core (@1.6 GHz) = 8.1 TFLOPs (180 Watts)**

# Part 2: accessing memory



# **Hardware multi-threading**

# Terminology

## ■ Memory latency

- The amount of time for a memory request (e.g., load, store) from a processor to be serviced by the memory system
- Example: 100 cycles, 100 nsec

## ■ Memory bandwidth

- The rate at which the memory system can provide data to a processor
- Example: 20 GB/s

# Stalls


- A processor “stalls” when it cannot run the next instruction in an instruction stream because of a dependency on a previous instruction.
- Accessing memory is a major source of stalls

```
ld r0 mem[r2] ←  
ld r1 mem[r3] ←  
add r0, r0, r1
```

Dependency: cannot execute 'add' instruction until data at mem[r2] and mem[r3] have been loaded from memory

- Memory access times ~ 100's of cycles
  - Memory “access time” is a measure of latency


# Review: why do modern processors have caches?



# Caches reduce length of stalls (reduce latency)

Processors run efficiently when data is resident in caches

Caches reduce memory access latency \*




\* Caches also provide high bandwidth data transfer to CPU

Stanford CS348V, Winter 2018

# Prefetching reduces stalls (hides latency)

- All modern CPUs have logic for prefetching data into caches
  - Dynamically analyze program's access patterns, predict what it will access soon
- Reduces stalls since data is resident in cache when accessed



**Note: Prefetching can also reduce performance if the guess is wrong (hogs bandwidth, pollutes caches)**

(more detail later in course)


# Multi-threading reduces stalls

- Idea: interleave processing of multiple threads on the same core to hide stalls
- Like prefetching, multi-threading is a latency hiding, not a latency reducing technique


# Hiding stalls with multi-threading




# Hiding stalls with multi-threading




# Hiding stalls with multi-threading



# Hiding stalls with multi-threading



# Throughput computing trade-off




**Key idea of throughput-oriented systems:  
Potentially increase time to complete work by any  
one any one thread, in order to increase overall  
system throughput when running multiple threads.**


**During this time, this thread is runnable, but it is not being executed  
by the processor. (The core is running some other thread.)**

# Kayvon's fictitious multi-core chip


**16 cores**




**8 SIMD ALUs per core  
(128 total)**




**4 threads per core**




**16 simultaneous instruction  
streams**



**64 total concurrent instruction  
streams**





**512 independent pieces of work  
are needed to run chip with  
maximal latency hiding ability**



# GPUs: extreme throughput-oriented processors


## NVIDIA GTX 1080 core ("SM")



 = SIMD function unit,  
control shared across 32 units  
(1 MUL-ADD per clock)

- Instructions operate on 32 pieces of data at a time (instruction streams called "warps").
- Think: warp = thread issuing 32-wide vector instructions
- Different instructions from up to four warps can be executed simultaneously (simultaneous multi-threading)
- Up to 64 warps are interleaved on the SM (interleaved multi-threading)
- Over 2,048 elements can be processed concurrently by a core

# NVIDIA GTX 1080



**There are 20 SM cores on the GTX 1080:**

**That's 40,960 pieces of data being processed concurrently to get maximal latency hiding!**

**Another example:  
for review and to check your understanding  
(if you understand the following sequence you understand this lecture)**

# Running code on a simple processor


**My very simple program:  
compute  $\sin(x)$  using Taylor expansion**

```
void sinx(int N, int terms, float* x, float* result)
{
    for (int i=0; i<N; i++)
    {
        float value = x[i];
        float numer = x[i] * x[i] * x[i];
        int denom = 6; // 3!
        int sign = -1;

        for (int j=1; j<=terms; j++)
        {
            value += sign * numer / denom;
            numer *= x[i] * x[i];
            denom *= (2*j+2) * (2*j+3);
            sign *= -1;
        }

        result[i] = value;
    }
}
```

**My very simple processor:  
completes one instruction per clock**



# Review: superscalar execution

## Unmodified program

```
void sinx(int N, int terms, float* x, float* result)
{
    for (int i=0; i<N; i++)
    {
        float value = x[i];
        float numer = x[i] * x[i] * x[i];
        int denom = 6; // 3!
        int sign = -1;


        for (int j=1; j<=terms; j++)
        {
            value += sign * numer / denom;
            numer *= x[i] * x[i];
            denom *= (2*j+2) * (2*j+3);
            sign *= -1;
        }

        result[i] = value;
    }
}
```

Independent operations in instruction stream

(They are detected by the processor at run-time and may be executed in parallel on execution units 1 and 2)

My single core, superscalar processor:  
executes up to two instructions per clock  
from a single instruction stream.



# Review: multi-core execution (two cores)

Modify program to create two threads of control (two instruction streams)

```
typedef struct {
    int N;
    int terms;
    float* x;
    float* result;
} my_args;


void parallel_sinx(int N, int terms, float* x, float* result) {
    pthread_t thread_id;
    my_args args;

    args.N = N/2;
    args.terms = terms;
    args.x = x;
    args.result = result;

    // launch thread
    pthread_create(&thread_id, NULL, my_thread_start, &args);
    sinx(N - args.N, terms, x + args.N, result + args.N); // do work
    pthread_join(thread_id, NULL);
}

void my_thread_start(void* thread_arg) {
    my_args* thread_args = (my_args*)thread_arg;
    sinx(args->N, args->terms, args->x, args->result); // do work
}
```

**My dual-core processor:  
executes one instruction per clock  
from an instruction stream on each core.**



# Review: multi-core + superscalar execution

Modify program to create two threads of control (two instruction streams)

```
typedef struct {
    int N;
    int terms;
    float* x;
    float* result;
} my_args;


void parallel_sinx(int N, int terms, float* x, float* result) {
    pthread_t thread_id;
    my_args args;

    args.N = N/2;
    args.terms = terms;
    args.x = x;
    args.result = result;

    // launch thread
    pthread_create(&thread_id, NULL, my_thread_start, &args);
    sinx(N - args.N, terms, x + args.N, result + args.N); // do work
    pthread_join(thread_id, NULL);
}

void my_thread_start(void* thread_arg) {
    my_args* thread_args = (my_args*)thread_arg;
    sinx(args->N, args->terms, args->x, args->result); // do work
}
```

My superscalar dual-core processor:  
executes up to two instructions per clock  
from an instruction stream on each core.



# Review: multi-core (four cores)


Modify program to create many threads of control:  
(code written in Kayvon's fictitious data-parallel language)

```
void sinx(int N, int terms, float* x, float* result)
{
    // declare independent loop iterations
    forall (int i from 0 to N-1)
    {
        float value = x[i];
        float numer = x[i] * x[i] * x[i];
        int denom = 6;    // 3!
        int sign = -1;

        for (int j=1; j<=terms; j++)
        {
            value += sign * numer / denom
            numer *= x[i] * x[i];
            denom *= (2*j+2) * (2*j+3);
            sign *= -1;
        }

        result[i] = value;
    }
}
```

My quad-core processor:  
executes one instruction per clock  
from an instruction stream on each core.



# Review: four, 8-wide SIMD cores

**Observation:** program must execute many iterations of the same loop body.

**Optimization:** share instruction stream across execution of multiple iterations (single instruction multiple data = SIMD)


```
void sinx(int N, int terms, float* x, float* result)
{
    // declare independent loop iterations
    forall (int i from 0 to N-1)
    {
        float value = x[i];
        float numer = x[i] * x[i] * x[i];
        int denom = 6; // 3!
        int sign = -1;

        for (int j=1; j<=terms; j++)
        {
            value += sign * numer / denom
            numer *= x[i] * x[i];
            denom *= (2*j+2) * (2*j+3);
            sign *= -1;
        }

        result[i] = value;
    }
}
```

My SIMD quad-core processor:

executes one 8-wide SIMD instruction per clock  
from an instruction stream on each core.



# Review: four SIMD, multi-threaded cores


**Observation: memory operations have very long latency**

**Solution: hide latency of loading data for one iteration by executing arithmetic instructions from other iterations**

```
void sinx(int N, int terms, float* x, float* result)
{
    // declare independent loop iterations
    forall (int i from 0 to N-1)
    {
        float value = x[i]; Memory load
        float numer = x[i] * x[i] * x[i];
        int denom = 6; // 3!
        int sign = -1;

        for (int j=1; j<=terms; j++)
        {
            value += sign * numer / denom
            numer *= x[i] * x[i];
            denom *= (2*j+2) * (2*j+3);
            sign *= -1;
        }
        result[i] = value; Memory store
    }
}
```





My multi-threaded, SIMD quad-core processor:  
executes one SIMD instruction per clock  
from one instruction stream on each core. But  
can switch to processing the other instruction  
stream when faced with a stall.



# Summary: four superscalar, SIMD, multi-threaded cores

My multi-threaded, superscalar, SIMD quad-core processor:  
executes up to two instructions per clock from one instruction stream on each core  
(in this example: one SIMD instruction + one scalar instruction).


Processor can switch to execute the other instruction stream when faced with stall.



# Connecting it all together

Kayvon's simple quad-core processor:

Four cores, two-way multi-threading per core (max eight threads active on chip at once), up to two instructions per clock per core (one of those instructions is 8-wide SIMD)




# Thought experiment

- You write a C application that spawns two pthreads
- The application runs on the processor shown below
  - Two cores, two-execution contexts per core, up to 8 instructions per clock, one instruction is an 8-wide SIMD instruction.
- Question: “who” is responsible for mapping your pthreads to the processor’s thread execution contexts?

Answer: the operating system

- Question: If you were the OS, how would you assign the two threads to the four available execution contexts?
- Another question: How would you assign threads to execution contexts if your C program spawned five pthreads?




# Another thought experiment

**Task: element-wise multiplication of two vectors A and B**

**Assume vectors contain millions of elements**

- **Load input A[i]**
- **Load input B[i]**
- **Compute  $A[i] \times B[i]$**
- **Store result into C[i]**



**Three memory operations (12 bytes) for every MUL**

**NVIDIA GTX 1080 GPU can do 2560 MULs per clock (@ 1.6 GHz)**

**Need ~50 TB/sec of bandwidth to keep functional units busy (only have 320 GB/sec)**

**<1% GPU efficiency... but 4.2x faster than eight-core CPU!**

**(3.2 GHz Xeon E5v4 eight-core CPU connected to 76 GB/sec memory bus will exhibit ~3% efficiency on this computation)**

# Bandwidth limited!

# Bandwidth limited!

**If processors request data at too high a rate, the memory system cannot keep up.**

**No amount of latency hiding helps this.**


**Bandwidth is a critical resource**

**Overcoming bandwidth limits are a common challenge for application developers on throughput-optimized systems.**

# **Hardware specialization**


# Why does energy efficiency matter?

- General mobile processing rule: the longer a task runs the less power it can use
  - Processor's power consumption is limited by heat generated (efficiency is required for more than just maximizing battery life)



# Efficiency benefits of compute specialization



- Rules of thumb: compared to high-quality C code on CPU...
- Throughput-maximized processor architectures: e.g., GPU cores
  - Approximately 10x improvement in perf / watt
  - Assuming code maps well to wide data-parallel execution and is compute bound
- Fixed-function ASIC (“application-specific integrated circuit”)
  - Can approach 100-1000x or greater improvement in perf/watt
  - Assuming code is compute bound and and is not floating-point math



*Efficient Embedded Computing [Dally et al. 08]*

[Figure credit Eric Chung]

# Hardware specialization increases efficiency




# Modern systems use ASICs for...

- Image/video encode/decode (e.g., H.264, JPG)
- Audio recording/playback
- Voice “wake up” (e.g., Ok Google)
- Camera “RAW” processing: processing data acquired by image sensor into images that are pleasing to humans
- Many 3D graphics tasks (rasterization, texture mapping, occlusion using the Z-buffer)
- Significant modern interest in ASICs for deep network evaluation (e.g., Google’s Tensor Processing Unit)


# Qualcomm Hexagon DSP



- Originally used for audio/LTE support on Qualcomm SoC's
- Multi-threaded, VLIW DSP
- Third major programmable unit on modern Qualcomm SoCs
  - Multi-core CPU
  - Multi-core GPU (Adreno)
  - Hexagon DSP



# Summary: choosing the right tool for the job



# Data movement has high energy cost

- Rule of thumb in mobile system design: always seek to reduce amount of data transferred from memory
  - Earlier in class we discussed minimizing communication to reduce stalls (poor performance).  
Now, we wish to reduce communication to reduce energy consumption
- “Ballpark” numbers [Sources: Bill Dally (NVIDIA), Tom Olson (ARM)]
  - Integer op: ~ 1 pJ \*
  - Floating point op: ~20 pJ \*
  - Reading 64 bits from small local SRAM (1mm away on chip): ~ 26 pJ
  - Reading 64 bits from low power mobile DRAM (LPDDR): ~1200 pJ ← Suggests that recomputing values, rather than storing and reloading them, is a better answer when optimizing code for energy efficiency!
- Implications
  - Reading 10 GB/sec from memory: ~1.6 watts
  - Entire power budget for mobile GPU: ~1 watt  
(remember phone is also running CPU, display, radios, etc.)
  - iPhone 6 battery: ~7 watt-hours (note: my Macbook Pro laptop: 99 watt-hour battery)
  - Exploiting locality matters!!!

\* Cost to just perform the logical operation, not counting overhead of instruction decode, load data from registers, etc. Stanford CS348V, Winter 2018

# Welcome to cs348v!

- Make sure you are signed up on Piazza so you get announcements
- Tonight's reading:
  - **“The Compute Architecture of Intel Processor Graphics Gen9” - Intel Technical Report, 2015**
  - **“The Rise of Mobile Visual Computing Systems”, Fatahalian, IEEE Mobile Computing 2016**

# More review

# For the rest of this class, know these terms

- **Multi-core processor**
- **SIMD execution**
- **Coherent control flow**
- **Hardware multi-threading**
  - **Interleaved multi-threading**
  - **Simultaneous multi-threading**
- **Memory latency**
- **Memory bandwidth**
- **Bandwidth bound application**
- **Arithmetic intensity**

# Which program performs better?

## Program 1

```
void add(int n, float* A, float* B, float* C) {
    for (int i=0; i<n; i++)
        C[i] = A[i] + B[i];
}

void mul(int n, float* A, float* B, float* C) {
    for (int i=0; i<n; i++)
        C[i] = A[i] * B[i];
}

float* A, *B, *C, *D, *E, *tmp1, *tmp2;

// assume arrays are allocated here

// compute E = D + ((A + B) * C)
add(n, A, B, tmp1);
mul(n, tmp1, C, tmp2);
add(n, tmp2, D, E);
```

**(Note: an answer probably needs to state its assumptions.)**

## Program 2

```
void fused(int n, float* A, float* B, float* C, float* D, float* E) {
    for (int i=0; i<n; i++)
        E[i] = D[i] + (A[i] + B[i]) * C[i];
}

// compute E = D + (A + B) * C
fused(n, A, B, C, D, E);
```

# More thought questions

## Program 1

```
void add(int n, float* A, float* B, float* C) {
    for (int i=0; i<n; i++)
        C[i] = A[i] + B[i];
}

void mul(int n, float* A, float* B, float* C) {
    for (int i=0; i<n; i++)
        C[i] = A[i] * B[i];
}

float* A, *B, *C, *D, *E, *tmp1, *tmp2;
// assume arrays are allocated here
// compute E = D + ((A + B) * C)
add(n, A, B, tmp1);
mul(n, tmp1, C, tmp2);
add(n, tmp2, D, E);
```

**Which code structuring style would you rather write?**

**Consider running either of these programs: would CPU support for hardware-multi-threading help performance?**

## Program 2

```
void fused(int n, float* A, float* B, float* C, float* D, float* E) {
    for (int i=0; i<n; i++)
        E[i] = D[i] + (A[i] + B[i]) * C[i];
}


// compute E = D + (A + B) * C
fused(n, A, B, C, D, E);
```

# **Visualizing interleaved and simultaneous multi-threading (and combinations thereof)**

# Interleaved multi-threading

Consider a processor with:

- Two execution contexts
- One fetch and decode unit (one instruction per clock)
- One ALU (to execute the instruction)




In an interleaved multi-threading scenario, the threads share the processor.

(This is a visualization of when threads are having their instructions executed by the ALU.)

# Interleaved multi-threading

Consider a processor with:

- Two execution contexts
- One fetch and decode unit (one instruction per clock)
- One ALU (to execute the instruction)




Same as previous slide, but now just a different scheduling order of the threads  
(fine-grained interleaving)

# Simultaneous multi-threading

Consider a processor with:

- Two execution contexts
- Two fetch and decode units (two instructions per clock)
- Two ALUs (to execute the two instructions)




In an simultaneous multi-threading scenario, the threads execute simultaneously on the two ALUs. (note, no ILP in a thread since each thread is run sequentially on one ALU)

# Combining simultaneous and interleaved multi-threading

Consider a processor with:

- **Four execution contexts**
- **Two fetch and decode units (two instructions per clock, choose two of four threads)**
- **Two ALUs (to execute the two instructions)**

time (clocks)




# Another way to visualize execution (ALU-centric view)

Consider a processor with:

- **Four execution contexts**
- **Two fetch and decode units (two instructions per clock, choose two of four threads)**
- **Two ALUs (to execute the two instructions)**

**Now the graph is visualizing what each ALU is doing each clock:**

time (clocks)



# Instructions can be drawn from same thread (ILP)

Consider a processor with:

- Four execution contexts
- Two fetch and decode units (two instructions per clock, choose any two independent instructions from the four threads)
- Two ALUs (to execute the two instructions)

