

Lecture 2:

The Camera Image Processing Pipeline

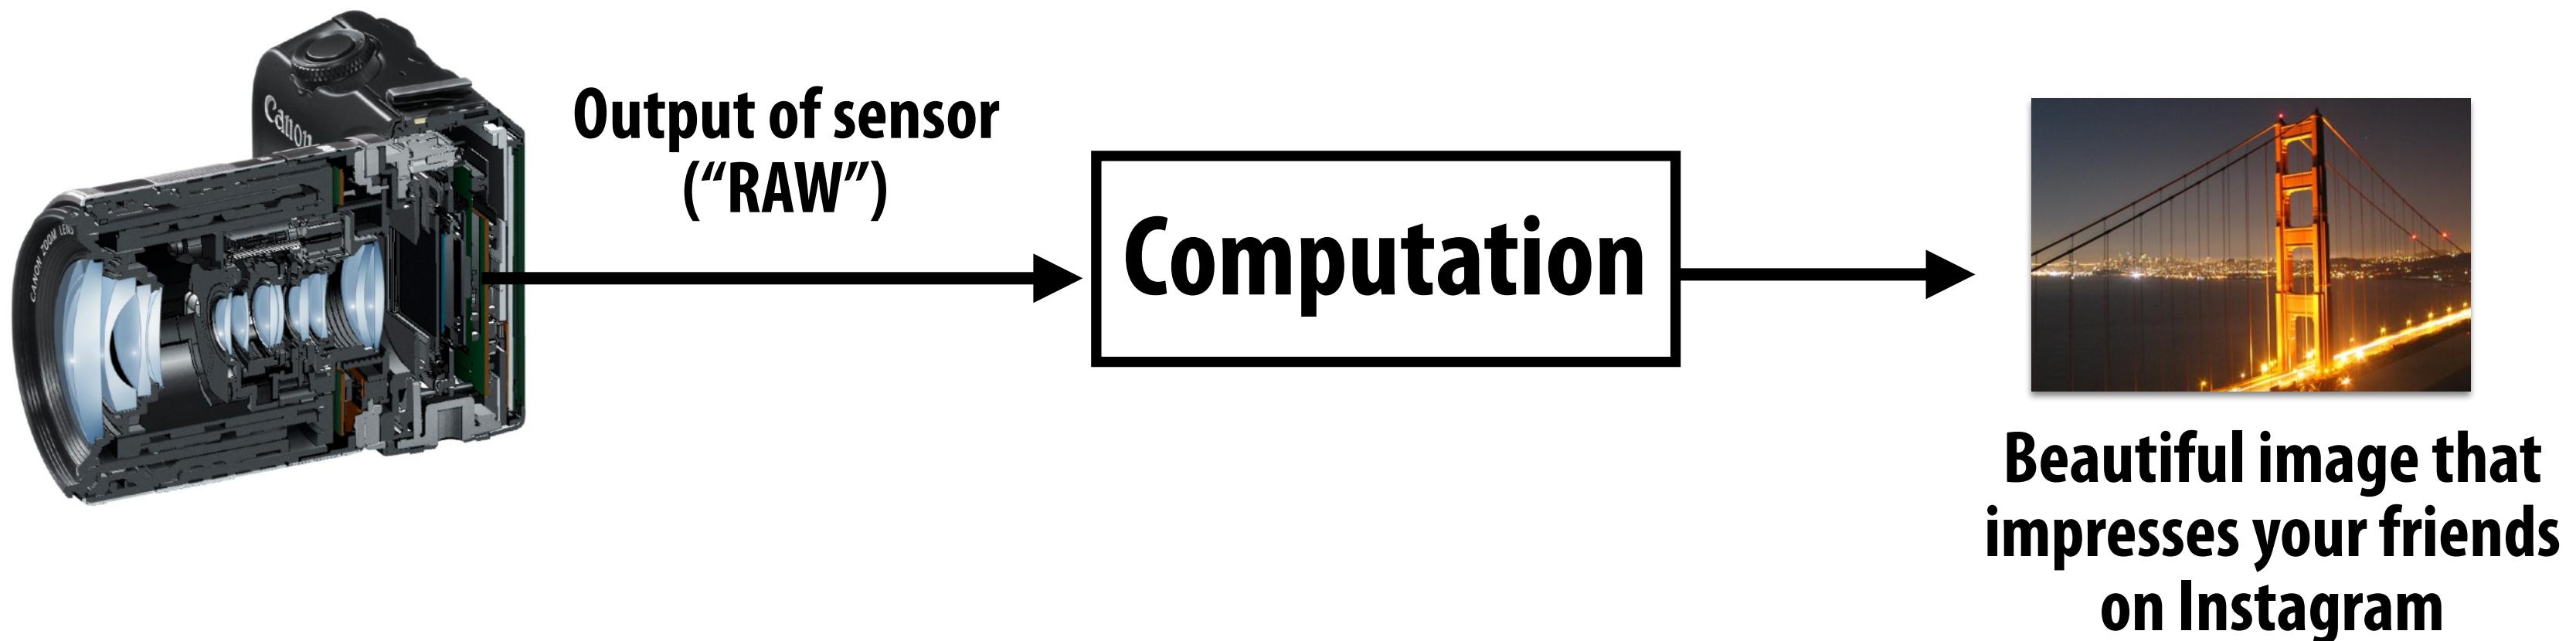
**Visual Computing Systems
Stanford CS348V, Winter 2018**

Acknowledgement to Ren Ng (Berkeley), Marc Levoy (Stanford/Google) for various slides used in this lecture.

Today's lecture

The values of pixels in a picture you see on screen are quite different than the values output by the sensor in a modern digital camera.

Computation is now a fundamental aspect of producing high-quality pictures.



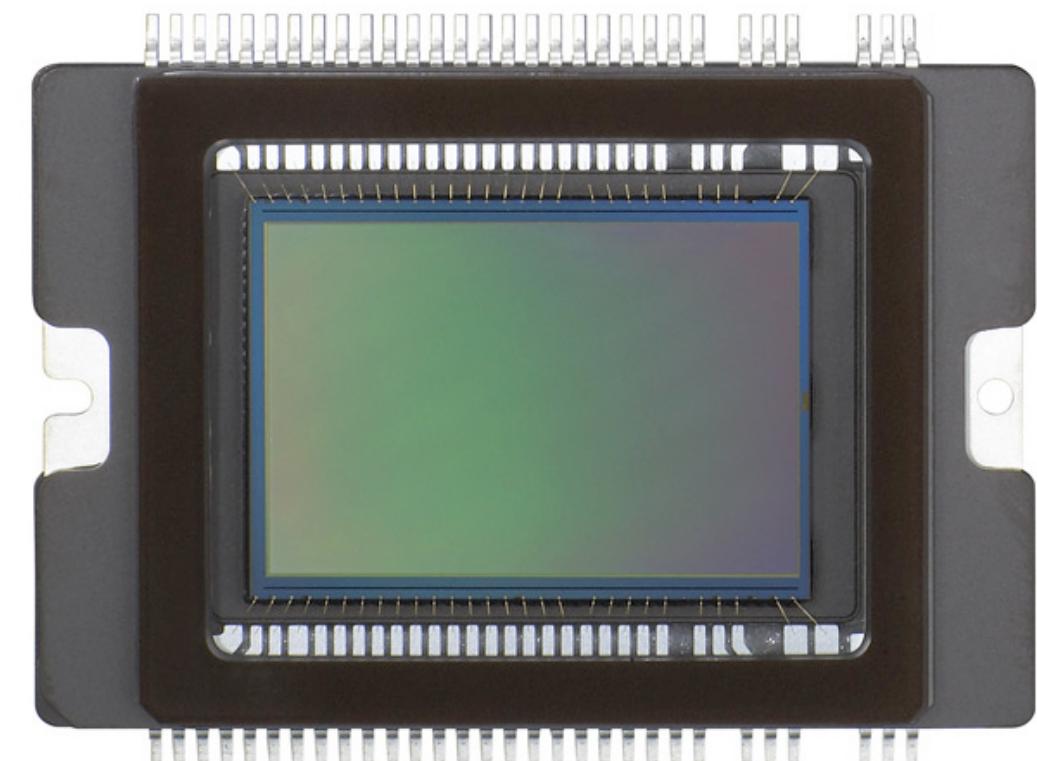
Where we are headed

- I'm about to describe a sequence of operations that take raw image pixels from a sensor (measurements of light) to high-quality pictures
 - Correct for sensor bias (using measurements of optically black pixels)
 - Correct pixel defects
 - Vignetting compensation
 - Dark frame subtract (optional)
 - White balance
 - Demosaic
 - Denoise / sharpen, etc.
 - Color Space Conversion
 - Gamma Correction
 - Color Space Conversion (Y'CbCr)
 - ...
- Today's pipelines are sophisticated, but they only scratch the surface of what future image processing pipelines might do
 - Person identification, action recognition, scene understanding (to automatically compose shot or automatically pick best picture) etc.

Camera cross section



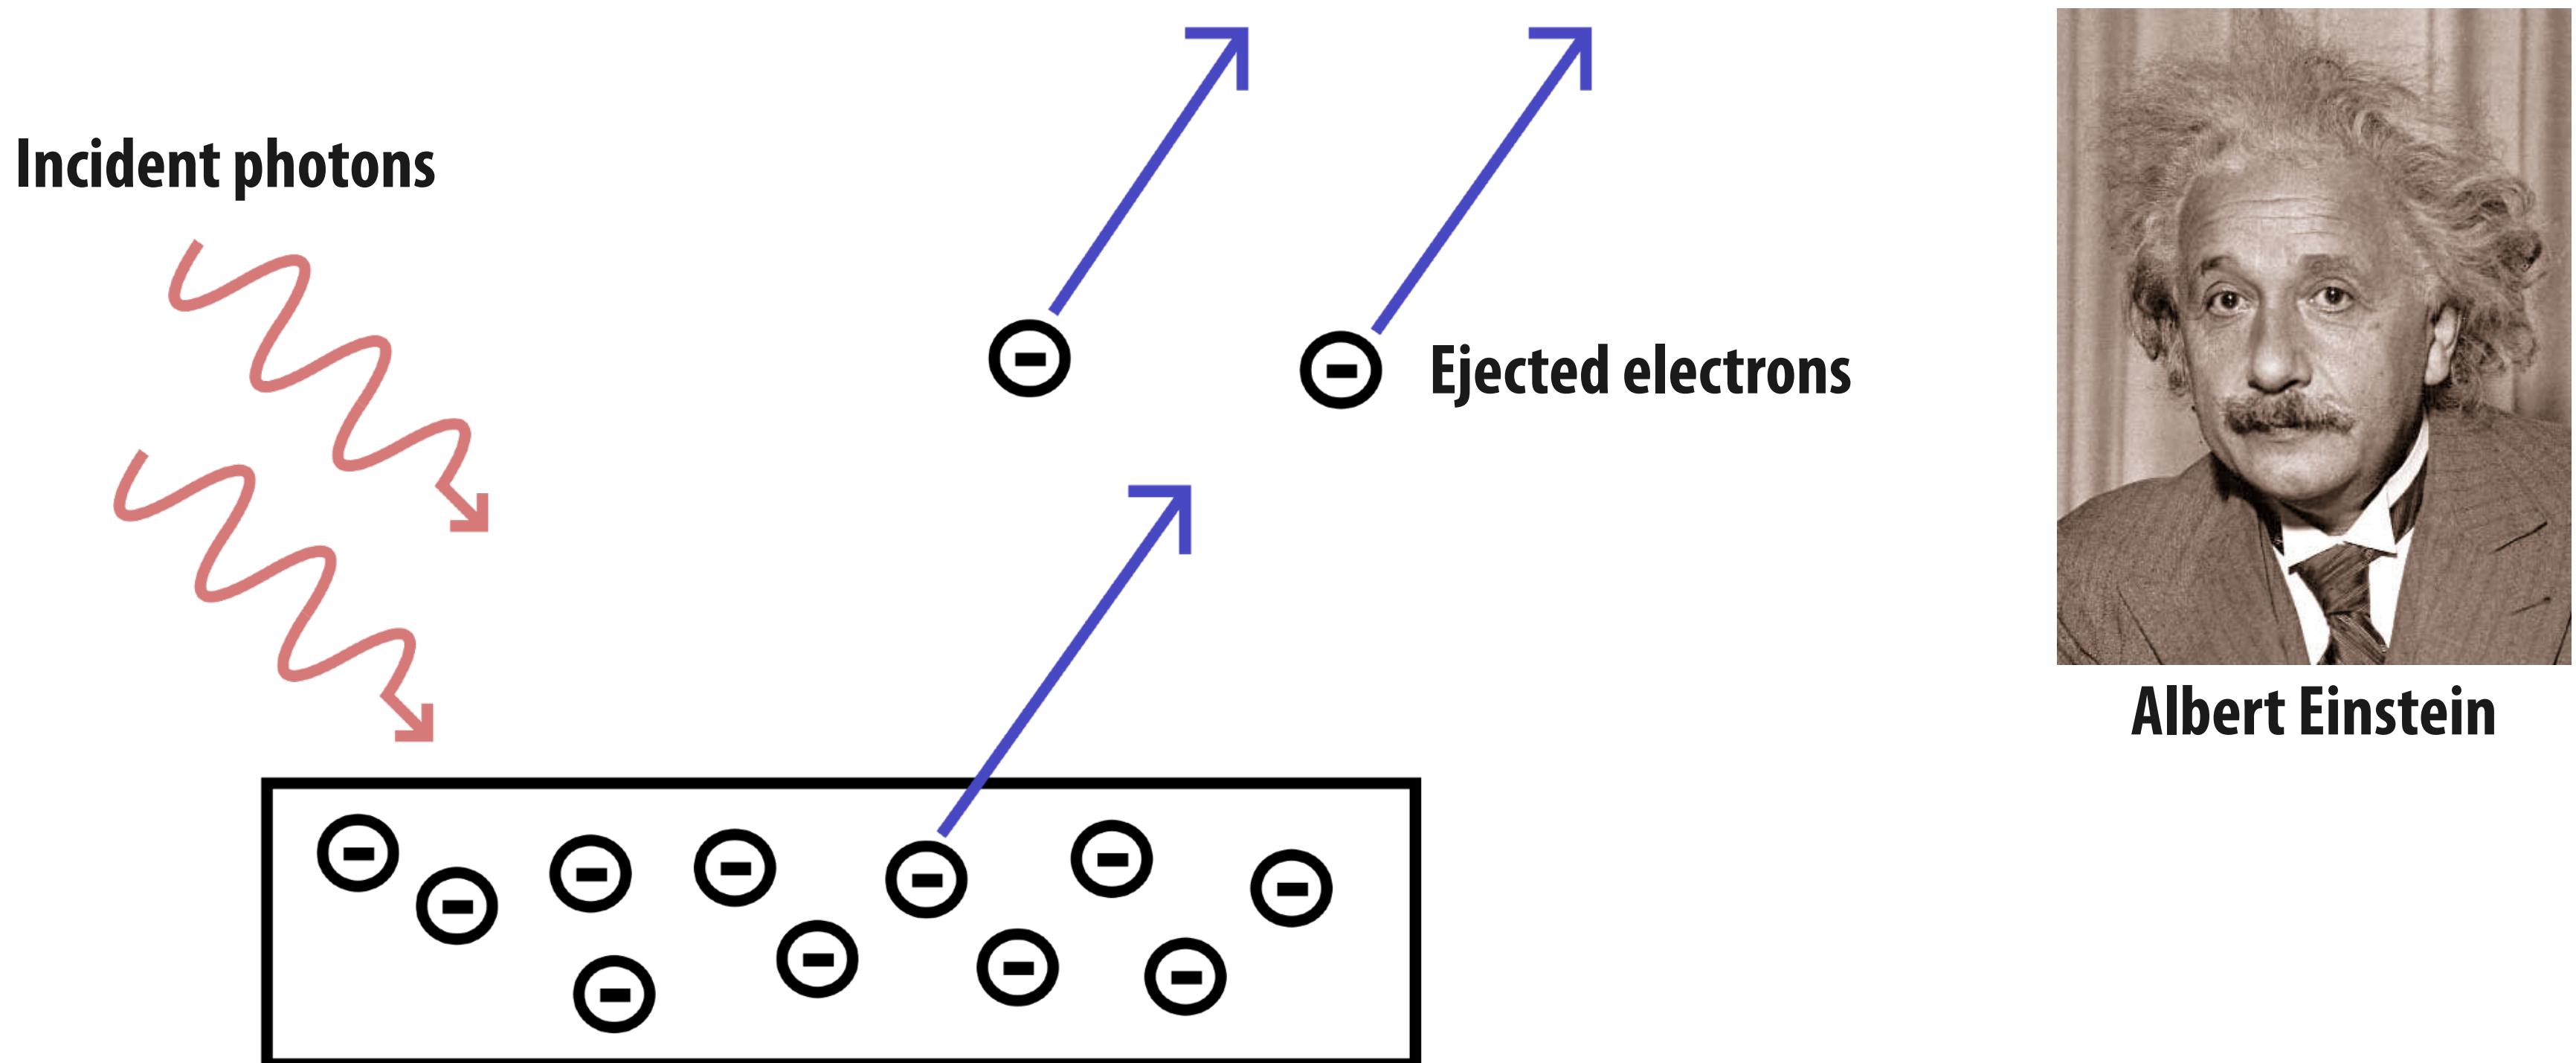
Sensor



**Canon 14 MP CMOS Sensor
(14 bits per pixel)**

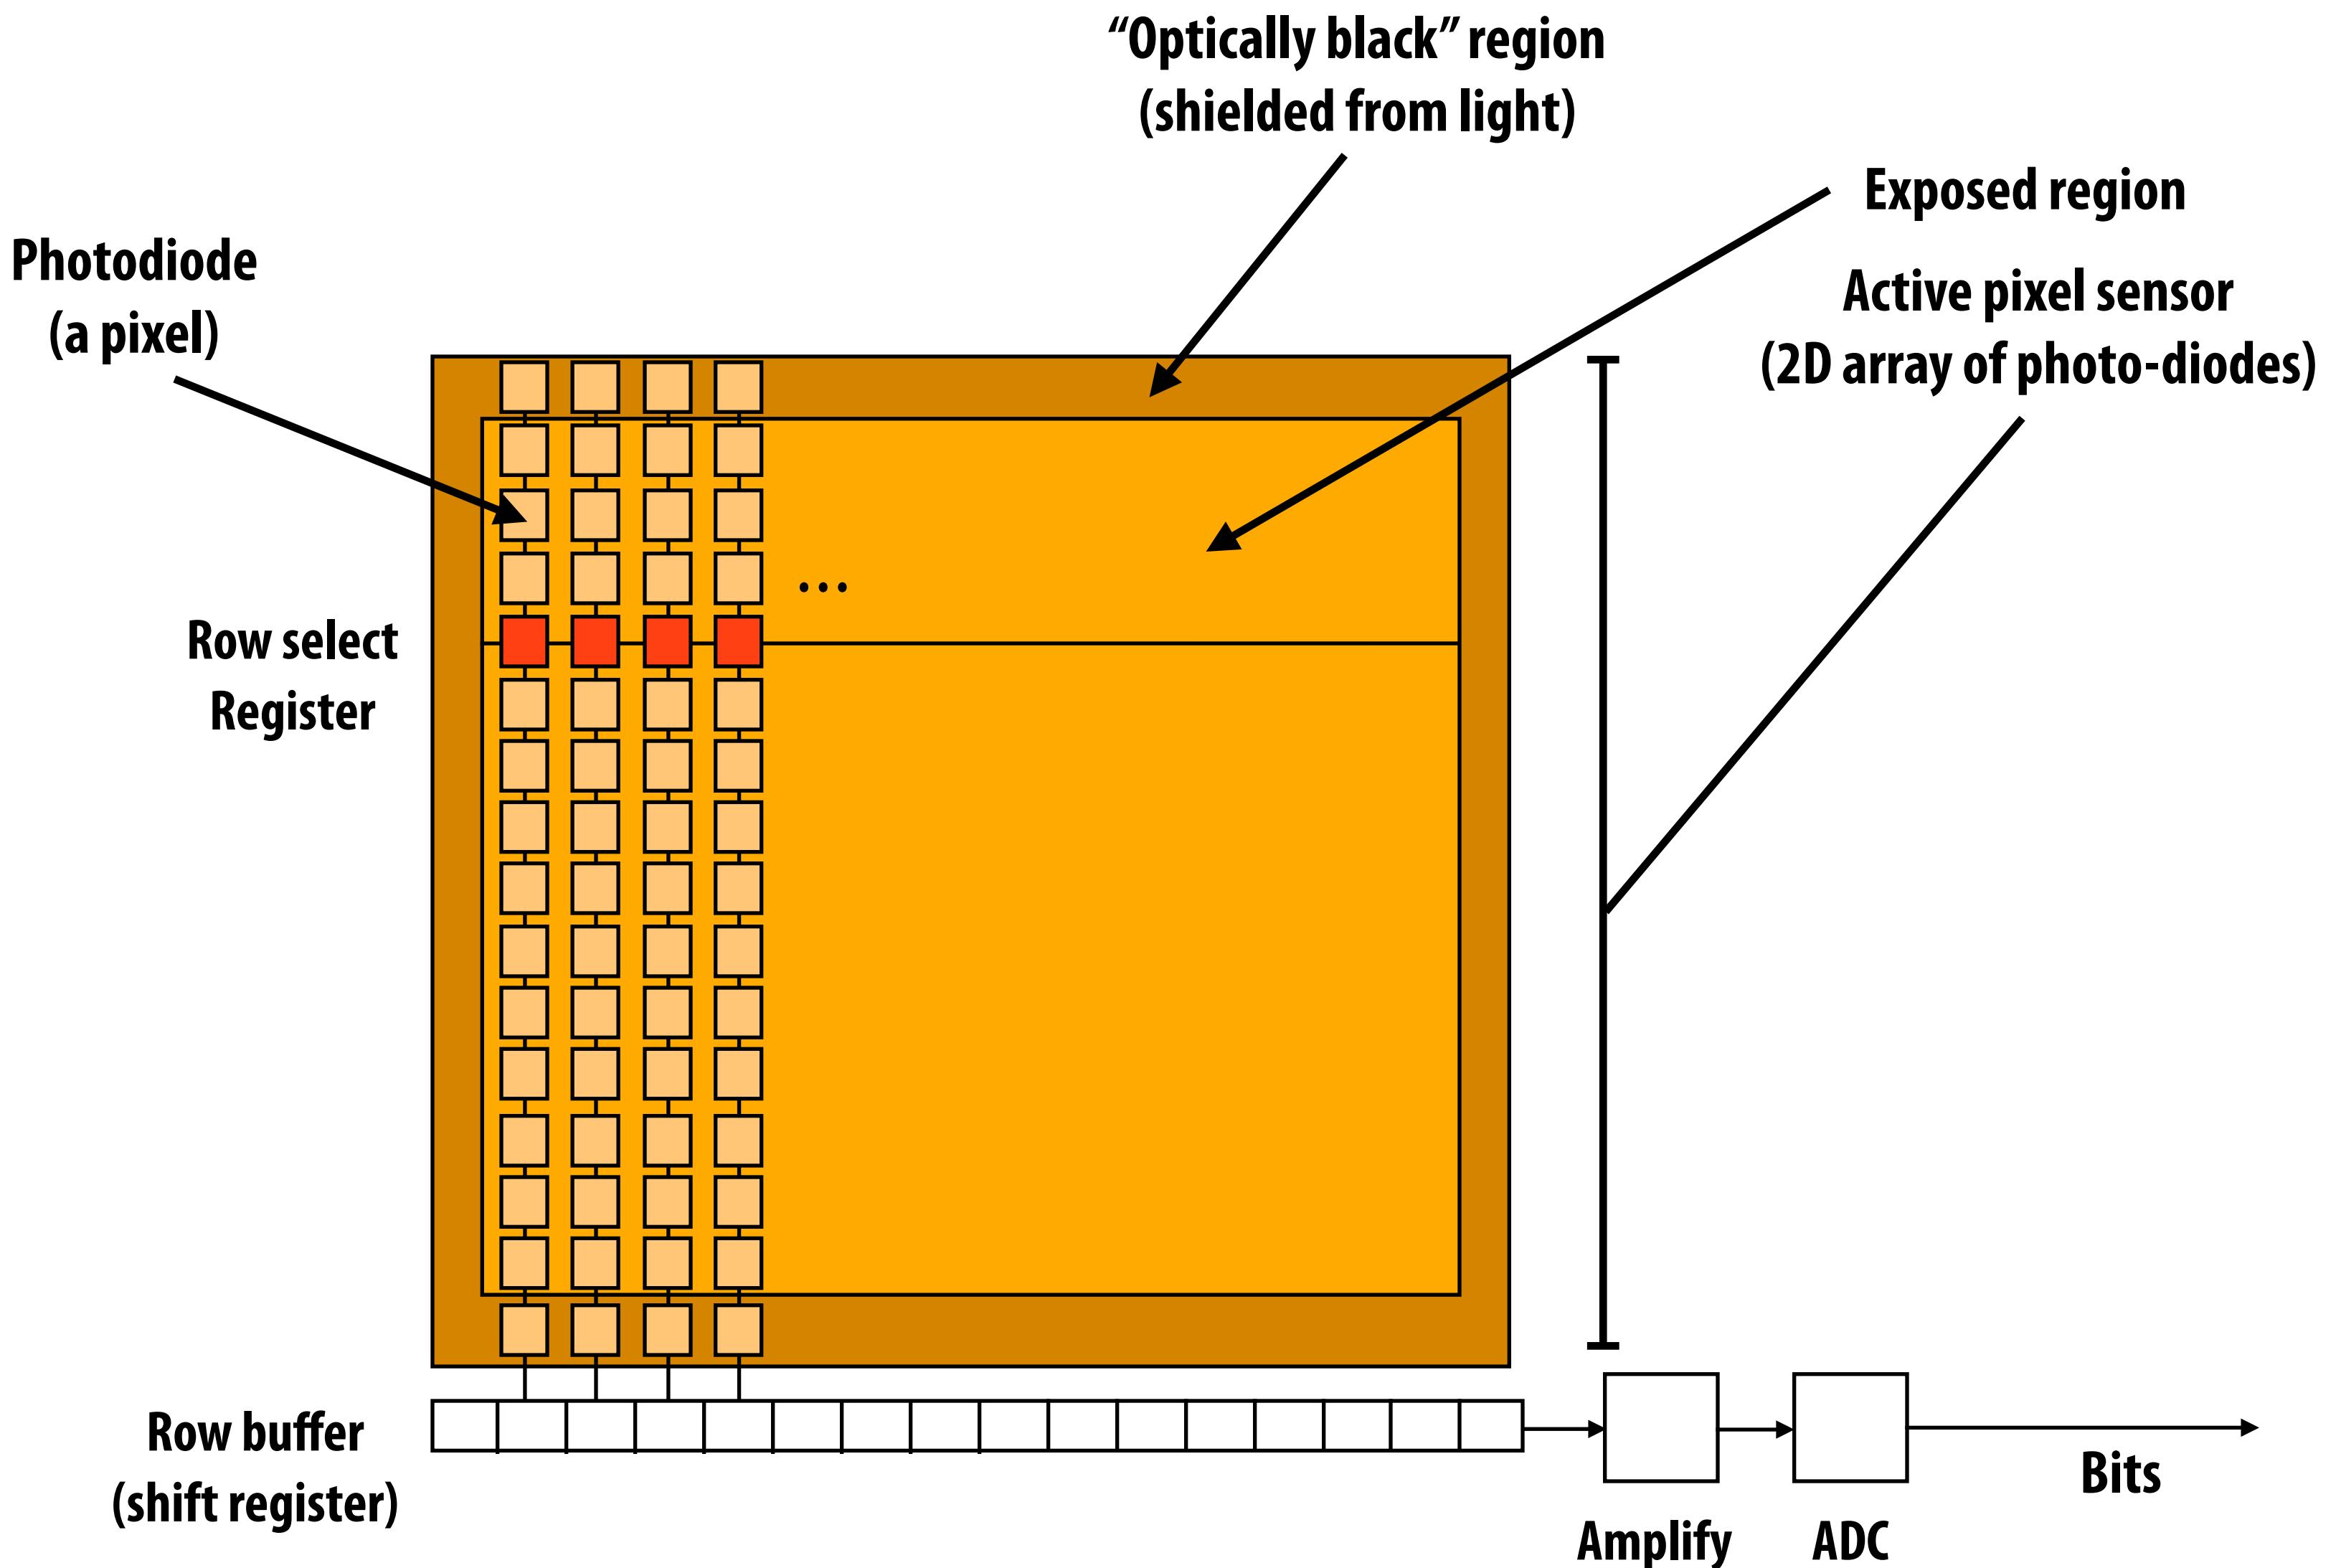
The Sensor

Photoelectric effect

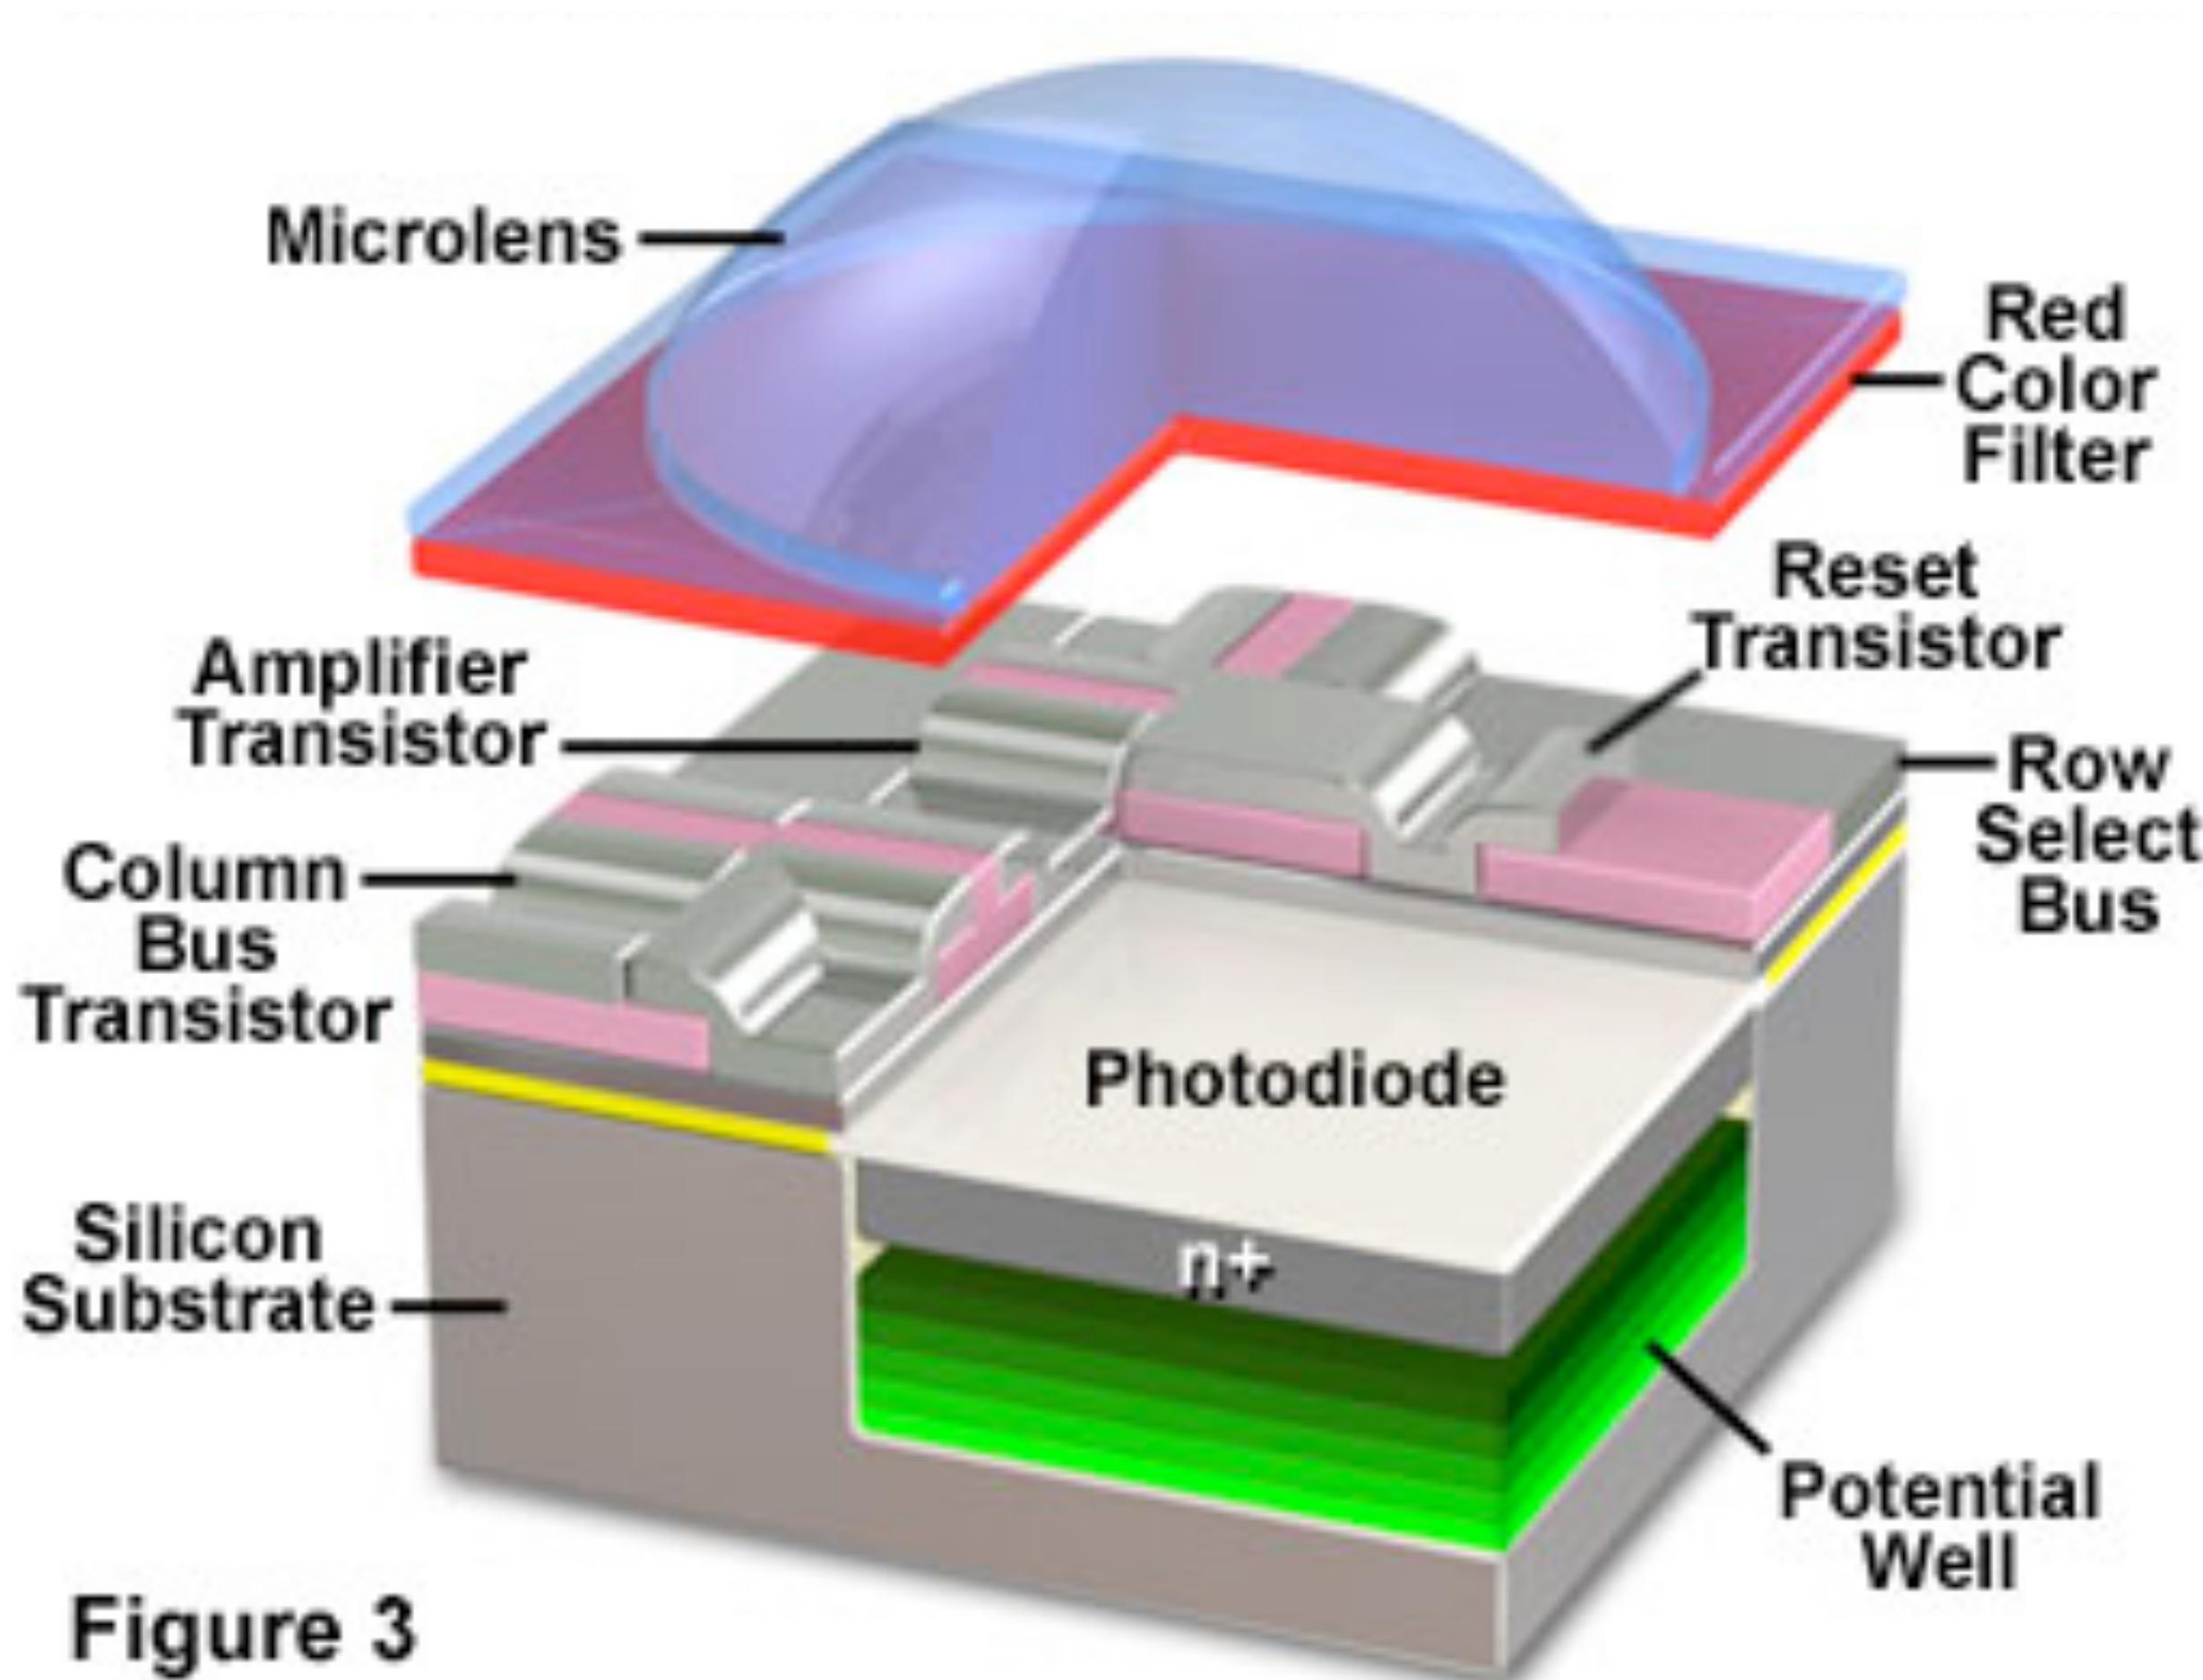


Einstein's Nobel Prize in 1921 "for his services to Theoretical Physics, and especially for his discovery of the law of the photoelectric effect"

CMOS sensor



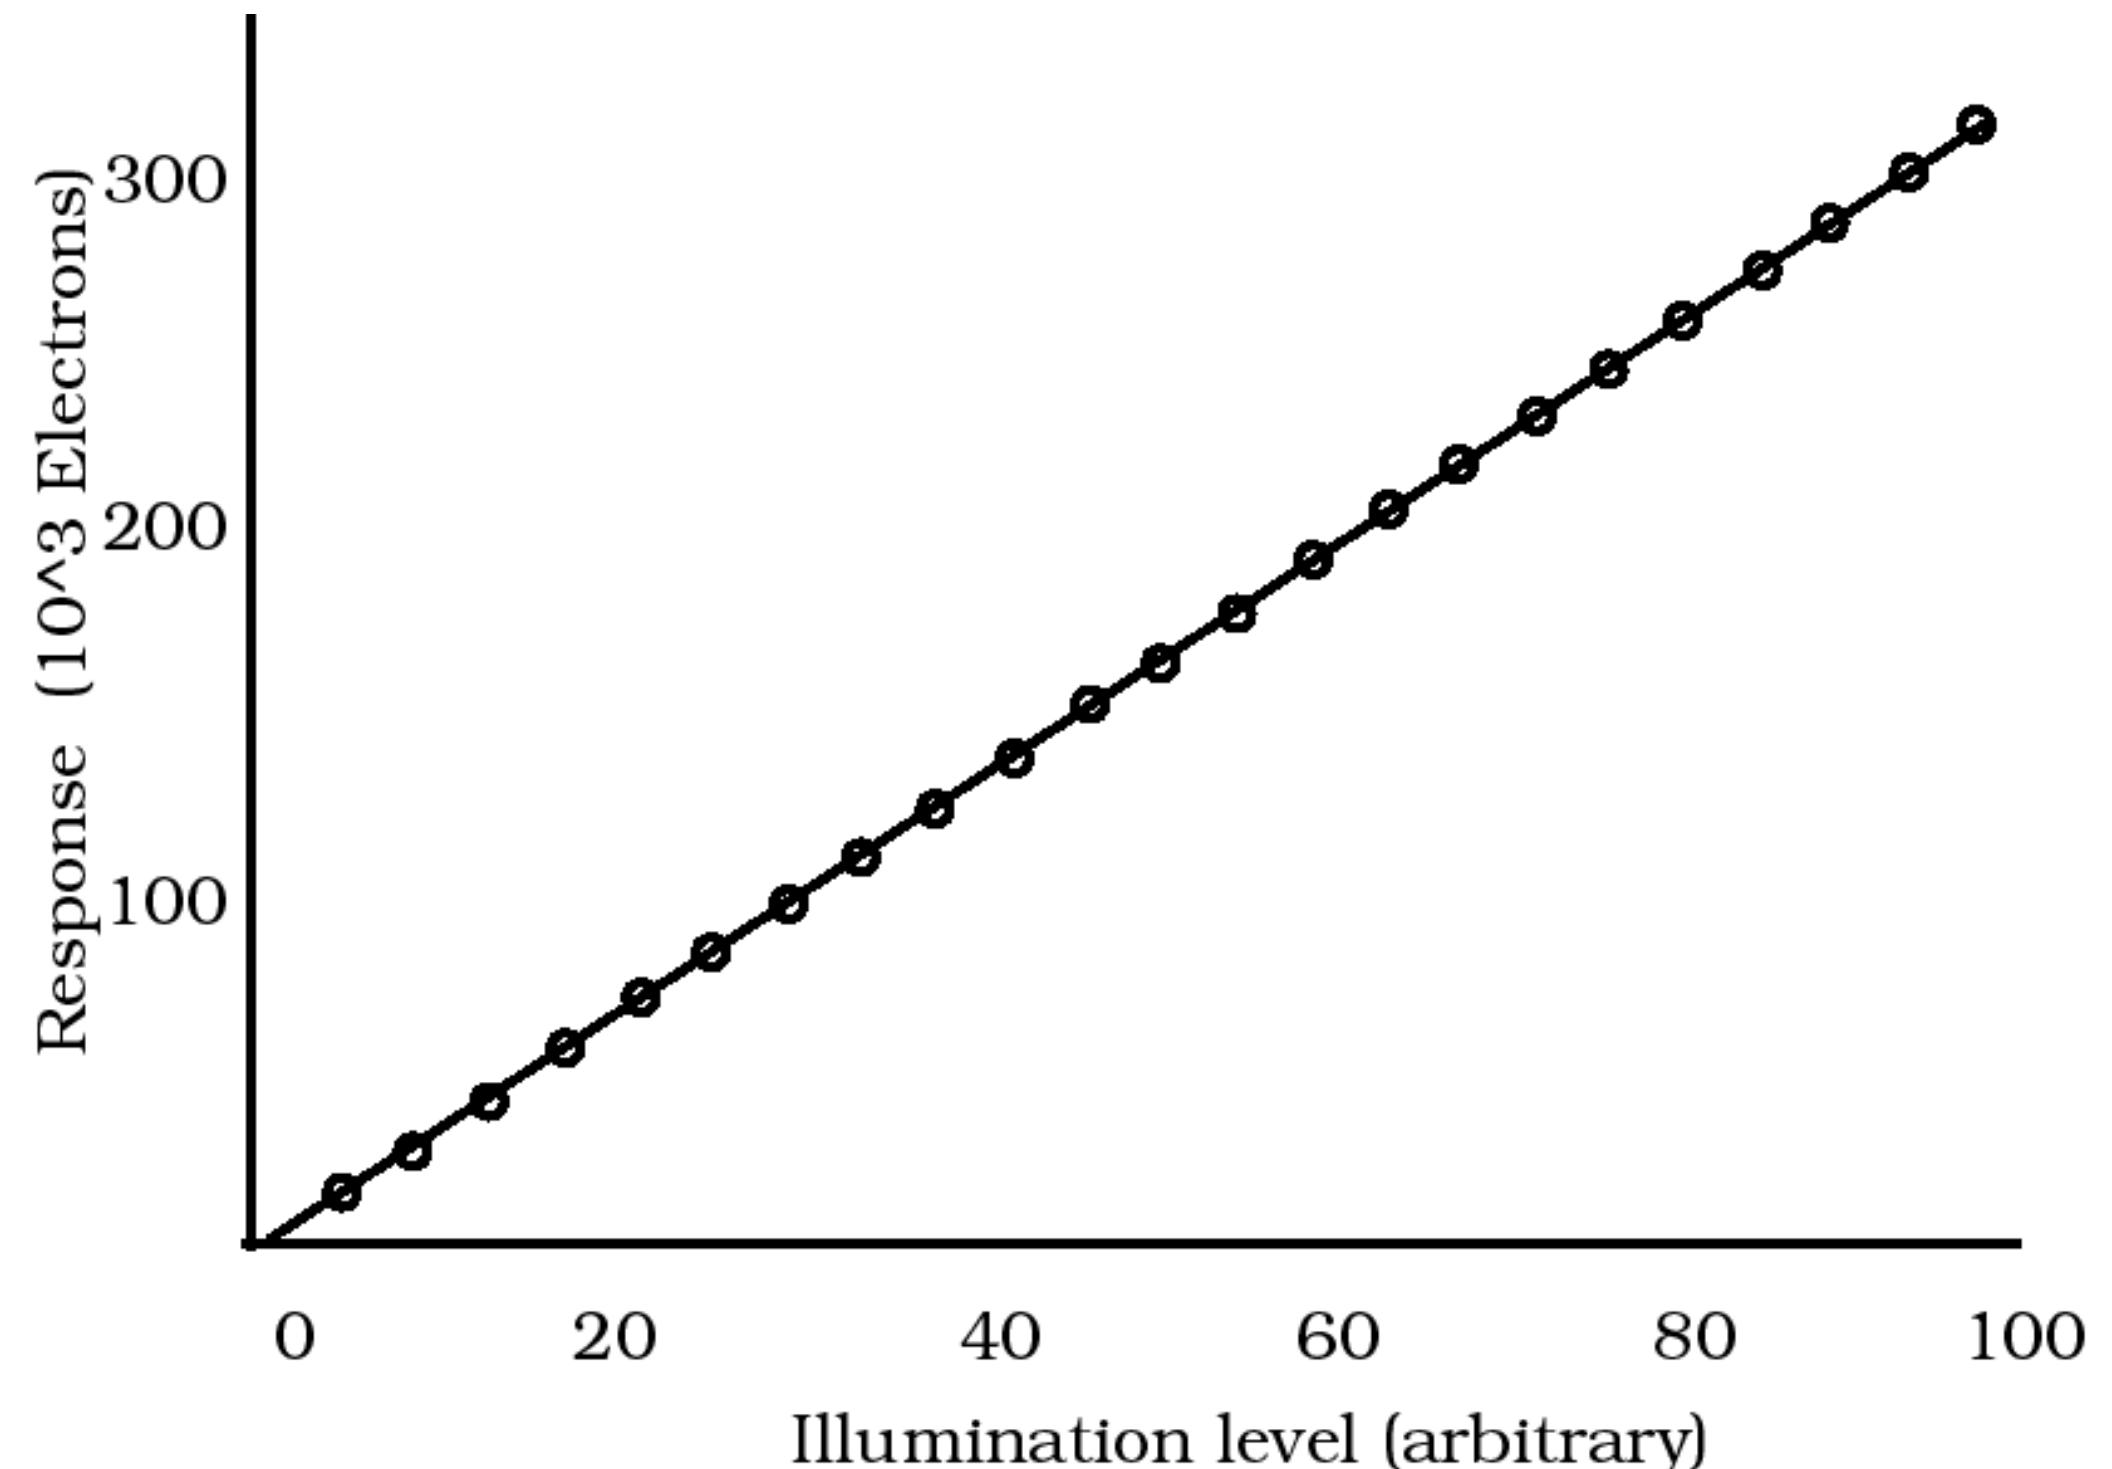
CMOS APS (active pixel sensor) pixel



CMOS response functions are linear

Photoelectric effect in silicon:

- **Response function from photons to electrons is linear**
- **May have some nonlinearity close to 0 due to noise and when close to pixel saturation**



(Epperson, P.M. et al. *Electro-optical characterization of the Tektronix TK5 ...*, *Opt Eng.*, 25, 1987)

Quantum efficiency

- Not all photons will produce an electron
 - Depends on quantum efficiency of the device

$$QE = \frac{\# \text{ electrons}}{\# \text{ photons}}$$

- Human vision: ~15%
- Typical digital camera: < 50%
- Best back-thinned CCD: > 90%
(e.g., telescope)

Sensing Color

Electromagnetic spectrum

Describes distribution of power (energy/time) by wavelength

Below: spectrum of various common light sources:

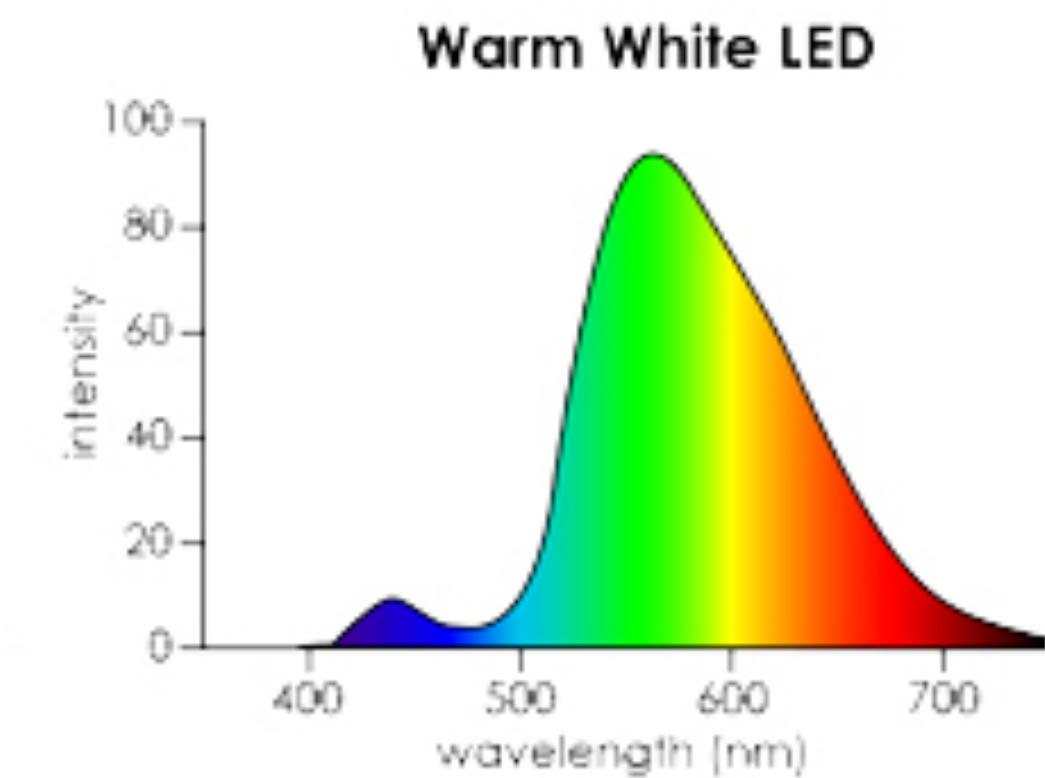
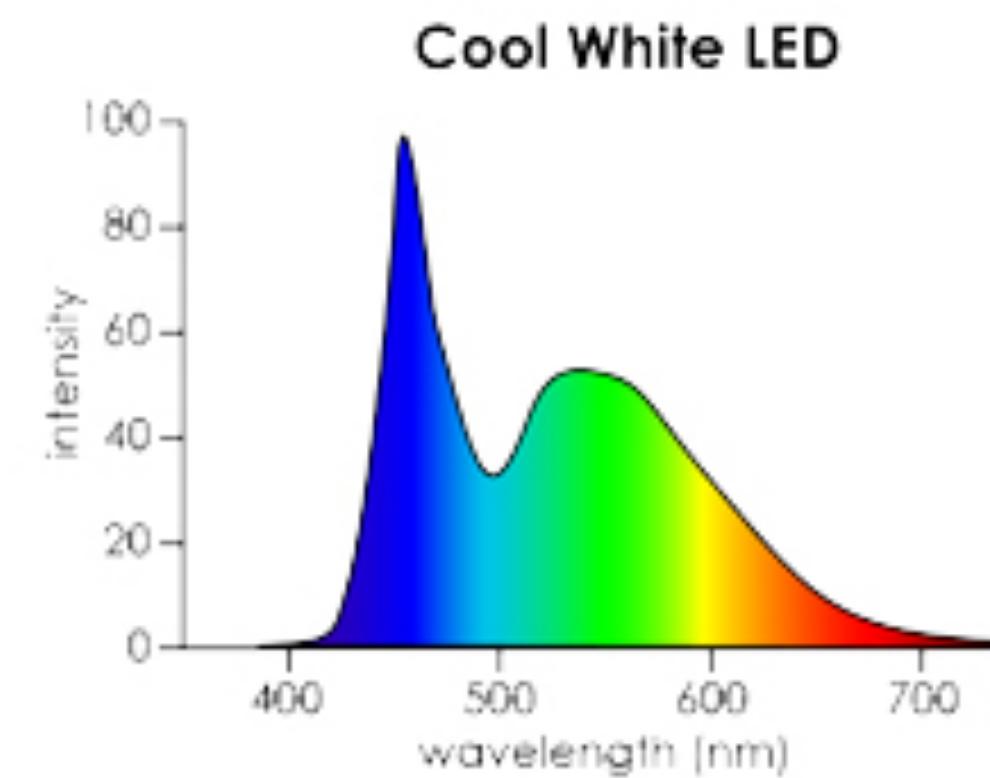
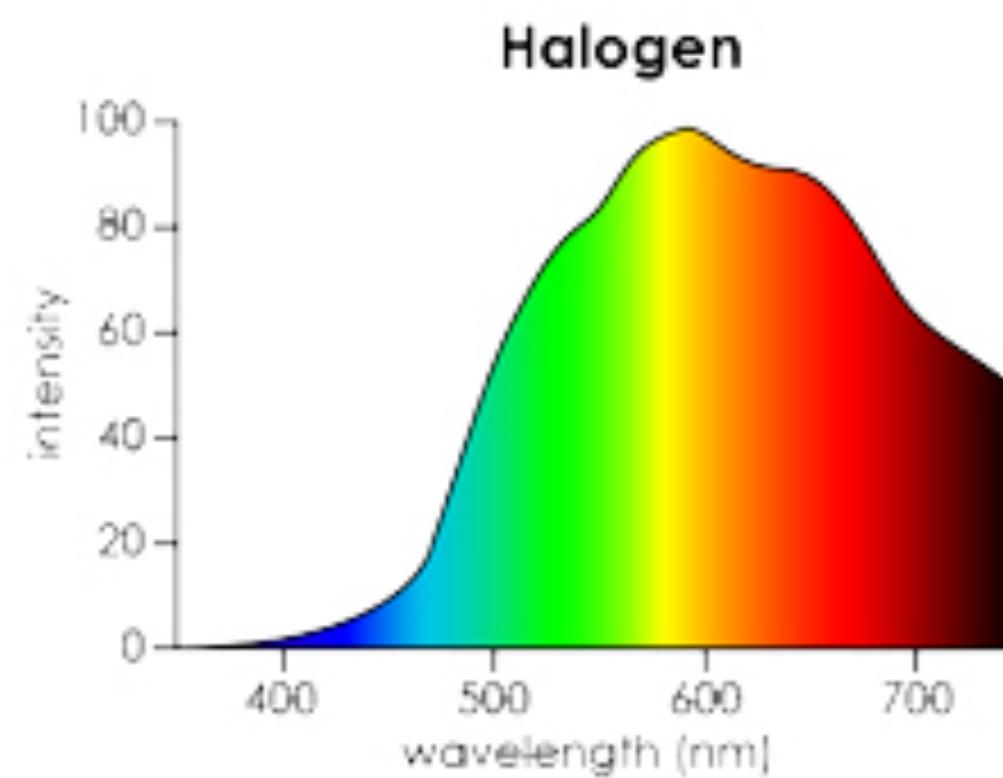
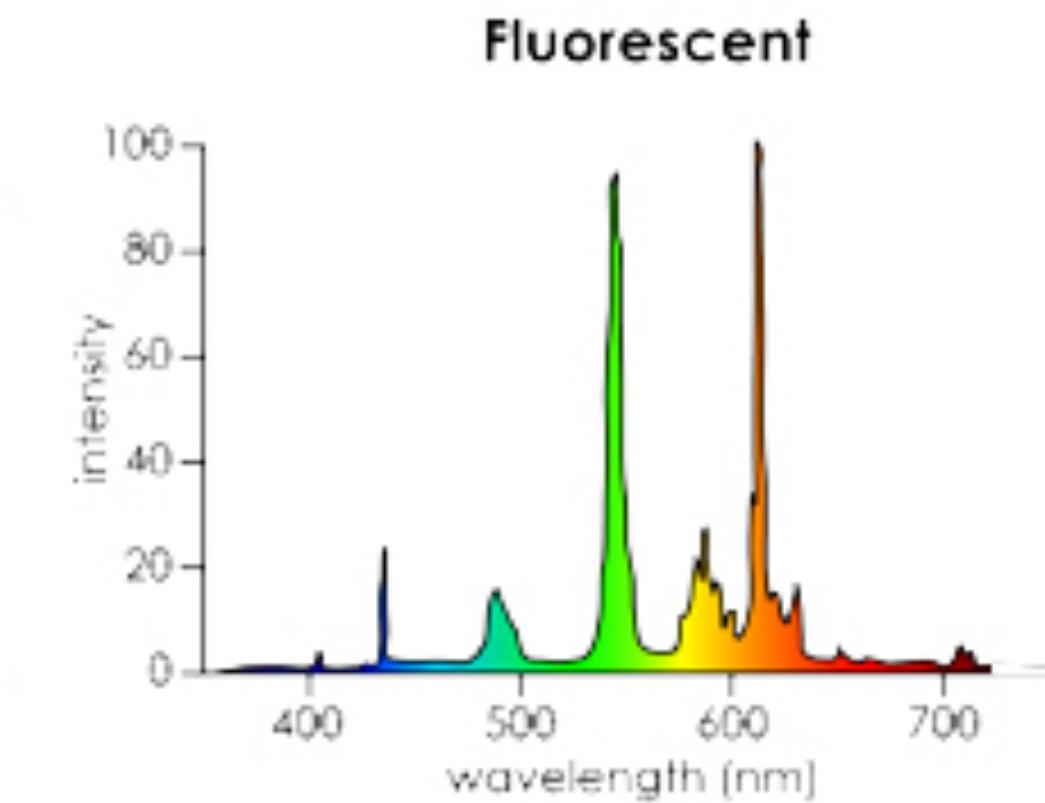
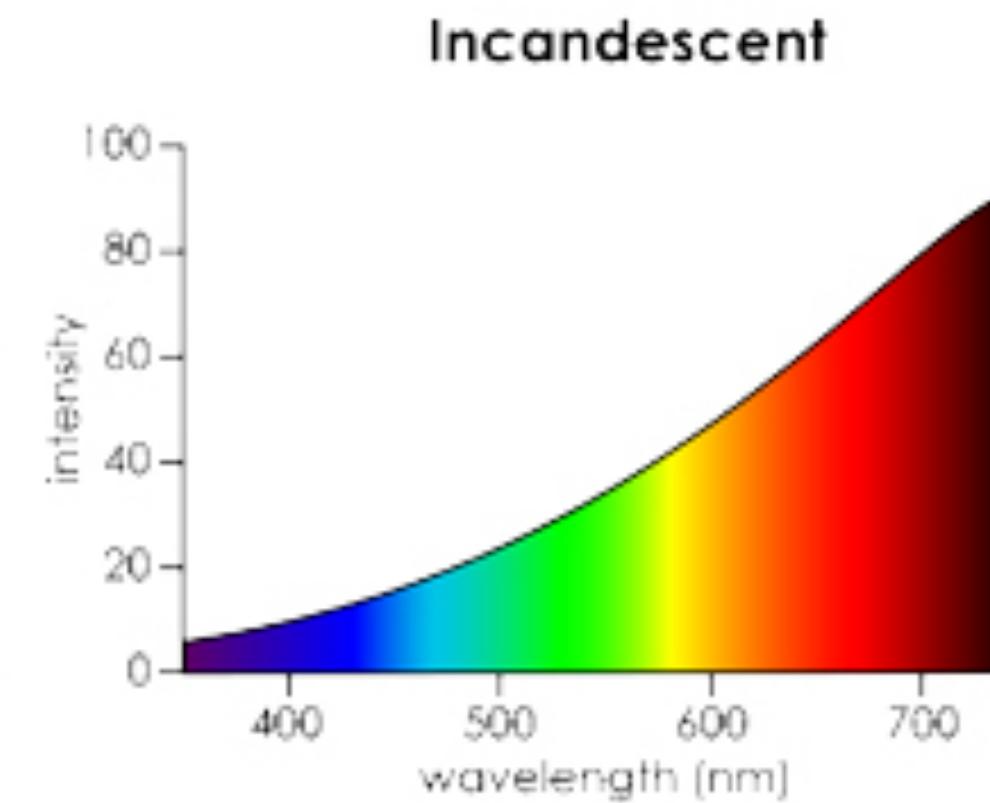
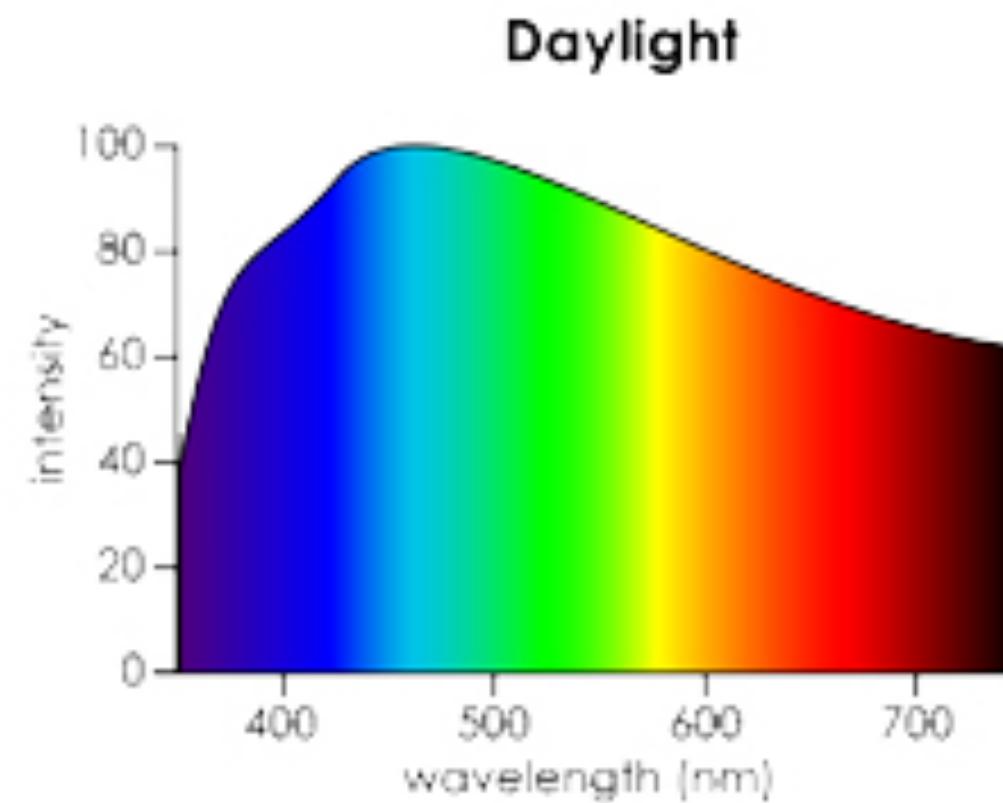


Figure credit:

Warm white vs. cool white



Image credit: (Oz Lighting: <https://www.ozlighting.com.au/blog/what-is-warm-white-versus-cool-white/>)

Photosensor response

- **Photosensor input: light**

- **Electromagnetic power distribution over wavelengths:** $\Phi(\lambda)$

- **Photosensor output: a “response” ... a number**

- **e.g., encoded in electrical signal**

- **Spectral response function:** $f(\lambda)$

- **Sensitivity of sensor to light of a given wavelength**
 - **Greater $f(\lambda)$ corresponds to more a efficient sensor (when $f(\lambda)$ is large, a small amount of light at wavelength λ will trigger a large sensor response)**

- **Total response of photosensor:**

$$R = \int_{\lambda} \Phi(\lambda) f(\lambda) d\lambda$$

Spectral response of cone cells in human eye

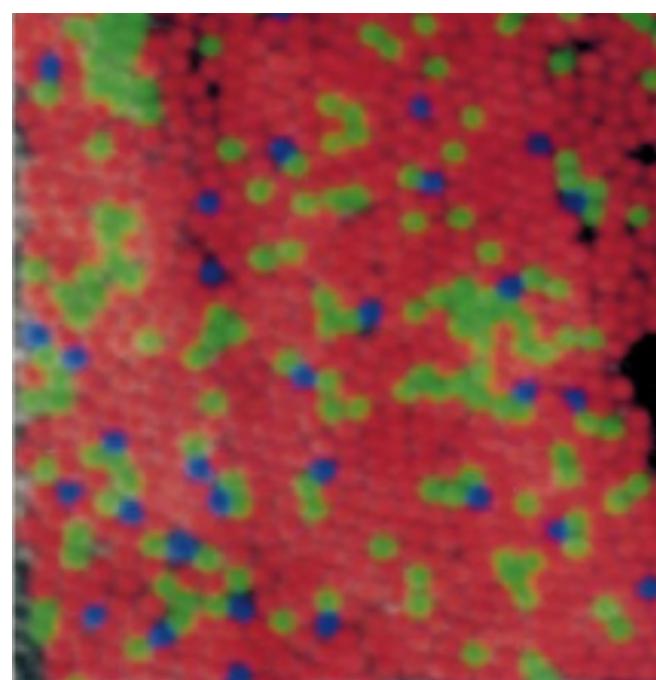
Three types of cells in eye responsible for color perception: S, M, and L cones (corresponding to peak response at short, medium, and long wavelengths)

Implication: the space of human-perceivable colors is three dimensional

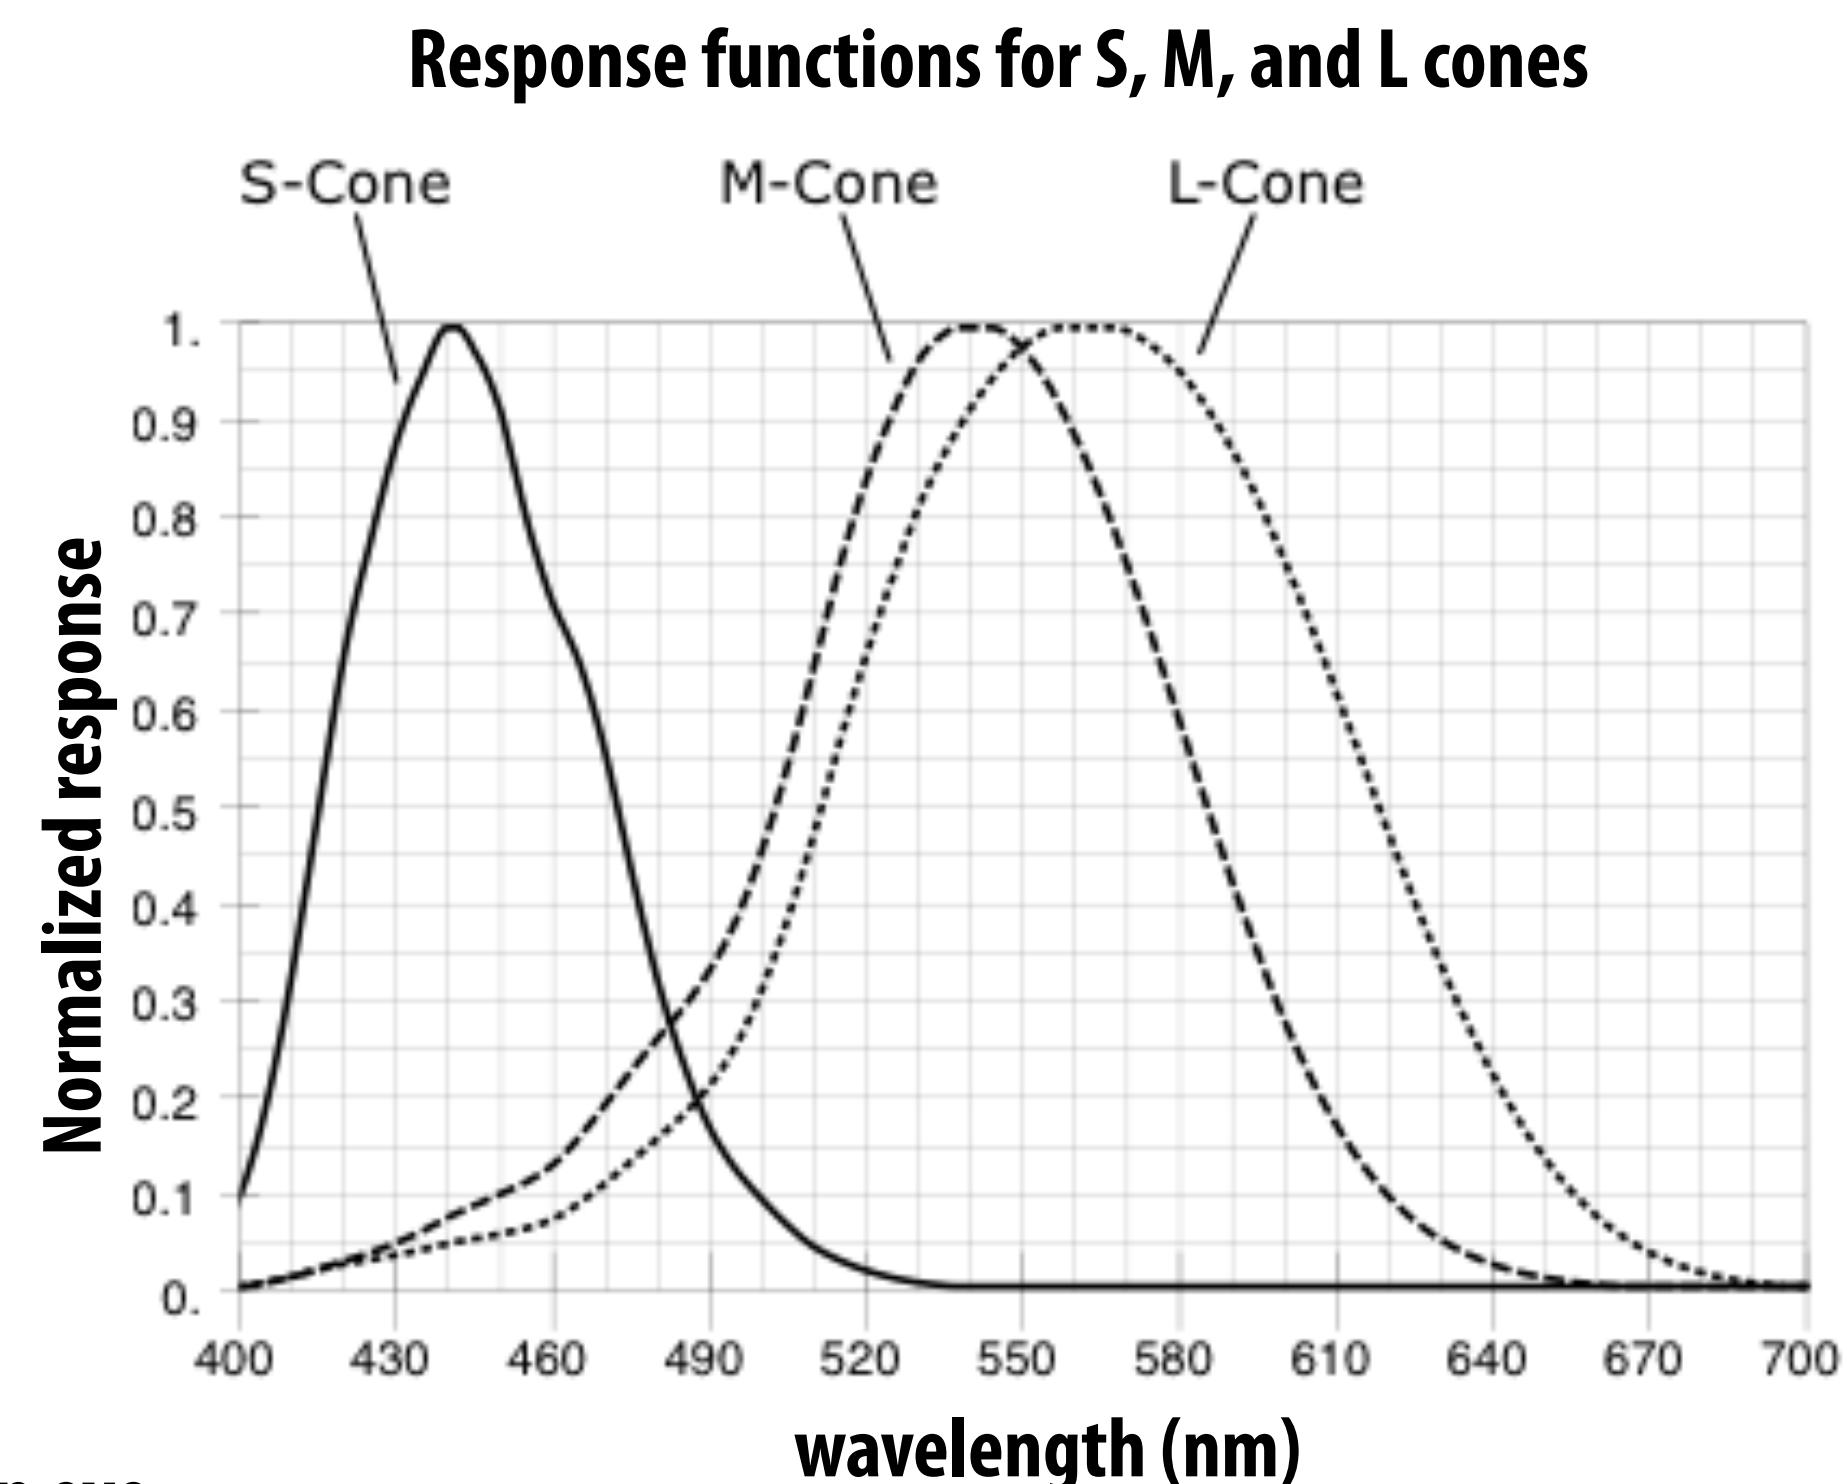
$$S = \int_{\lambda} \Phi(\lambda) S(\lambda) d\lambda$$

$$M = \int_{\lambda} \Phi(\lambda) M(\lambda) d\lambda$$

$$L = \int_{\lambda} \Phi(\lambda) L(\lambda) d\lambda$$

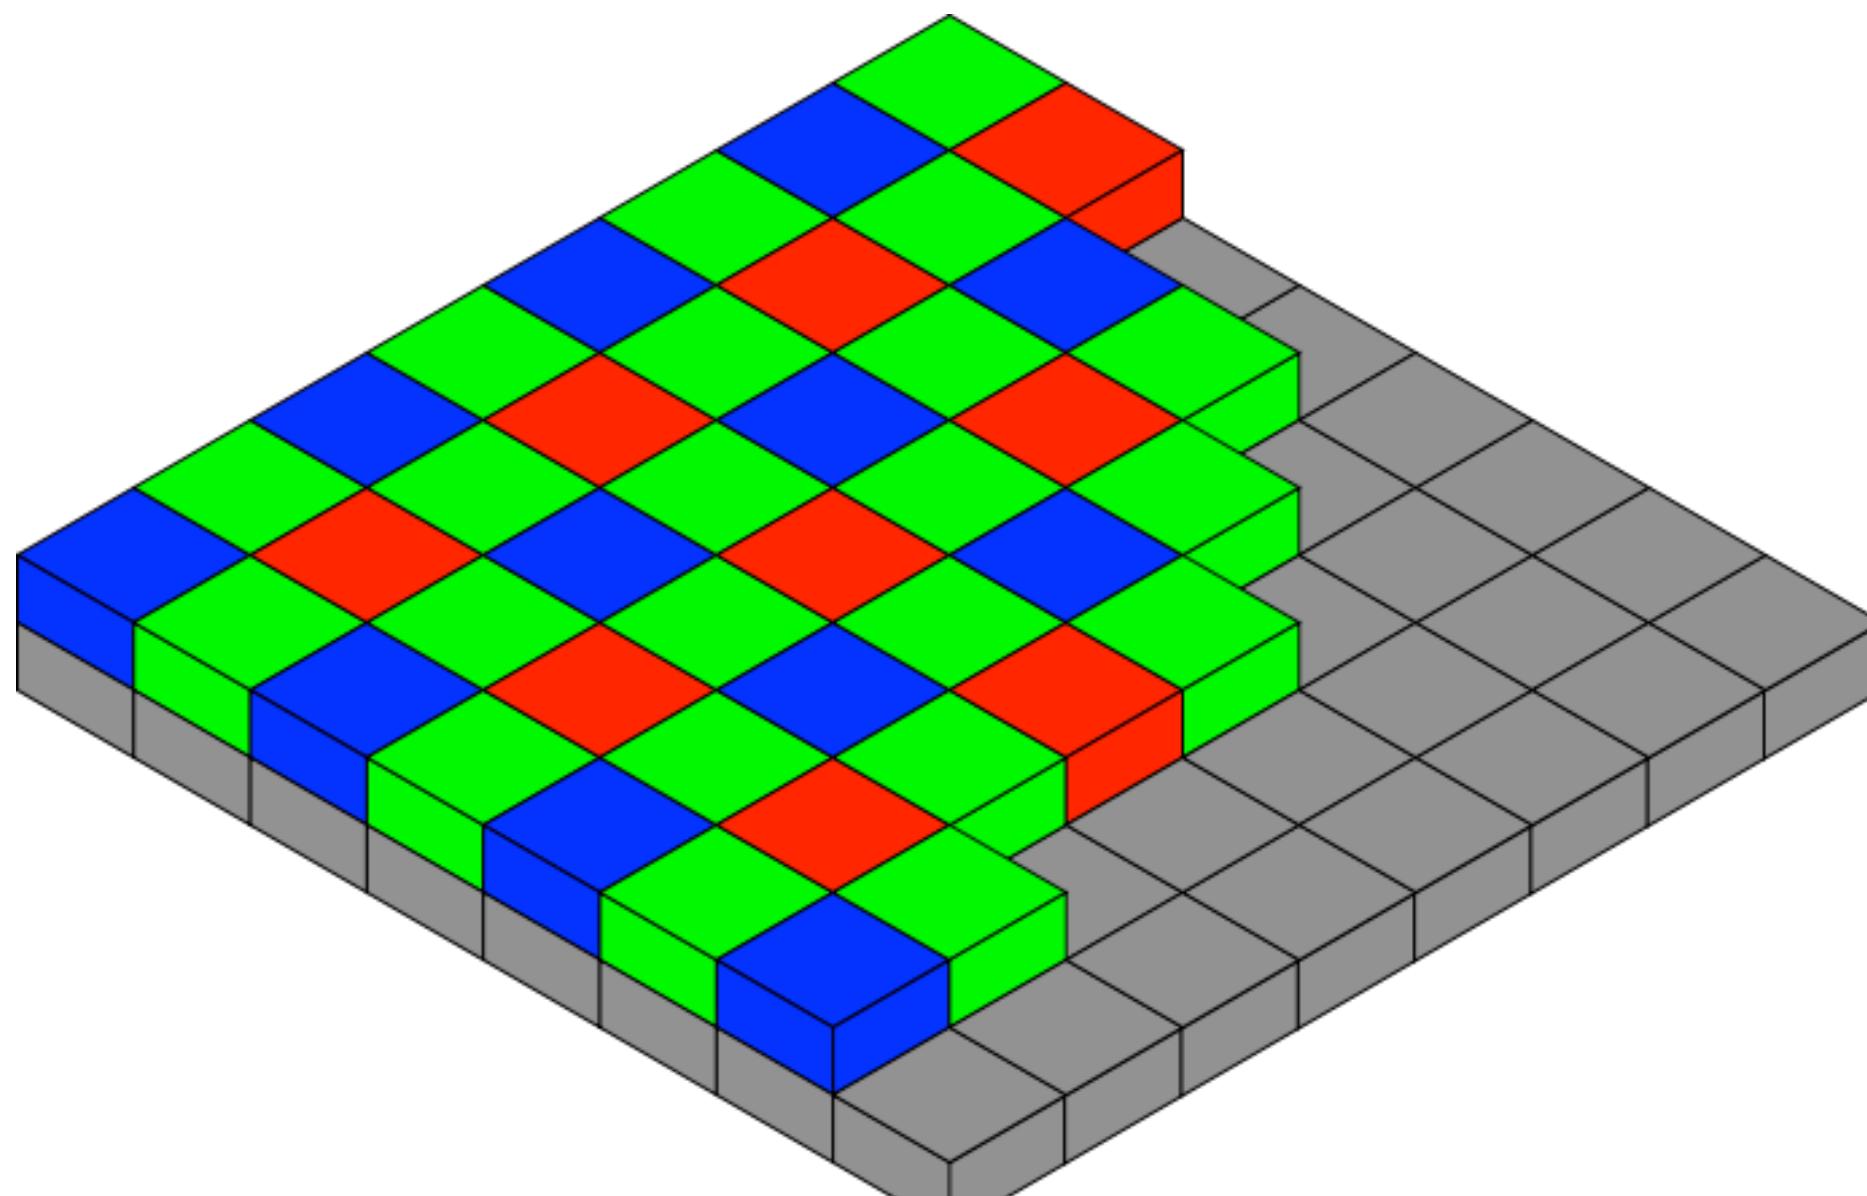


Uneven distribution of cone types in eye
~64% of cones are L cones, ~32% M cones



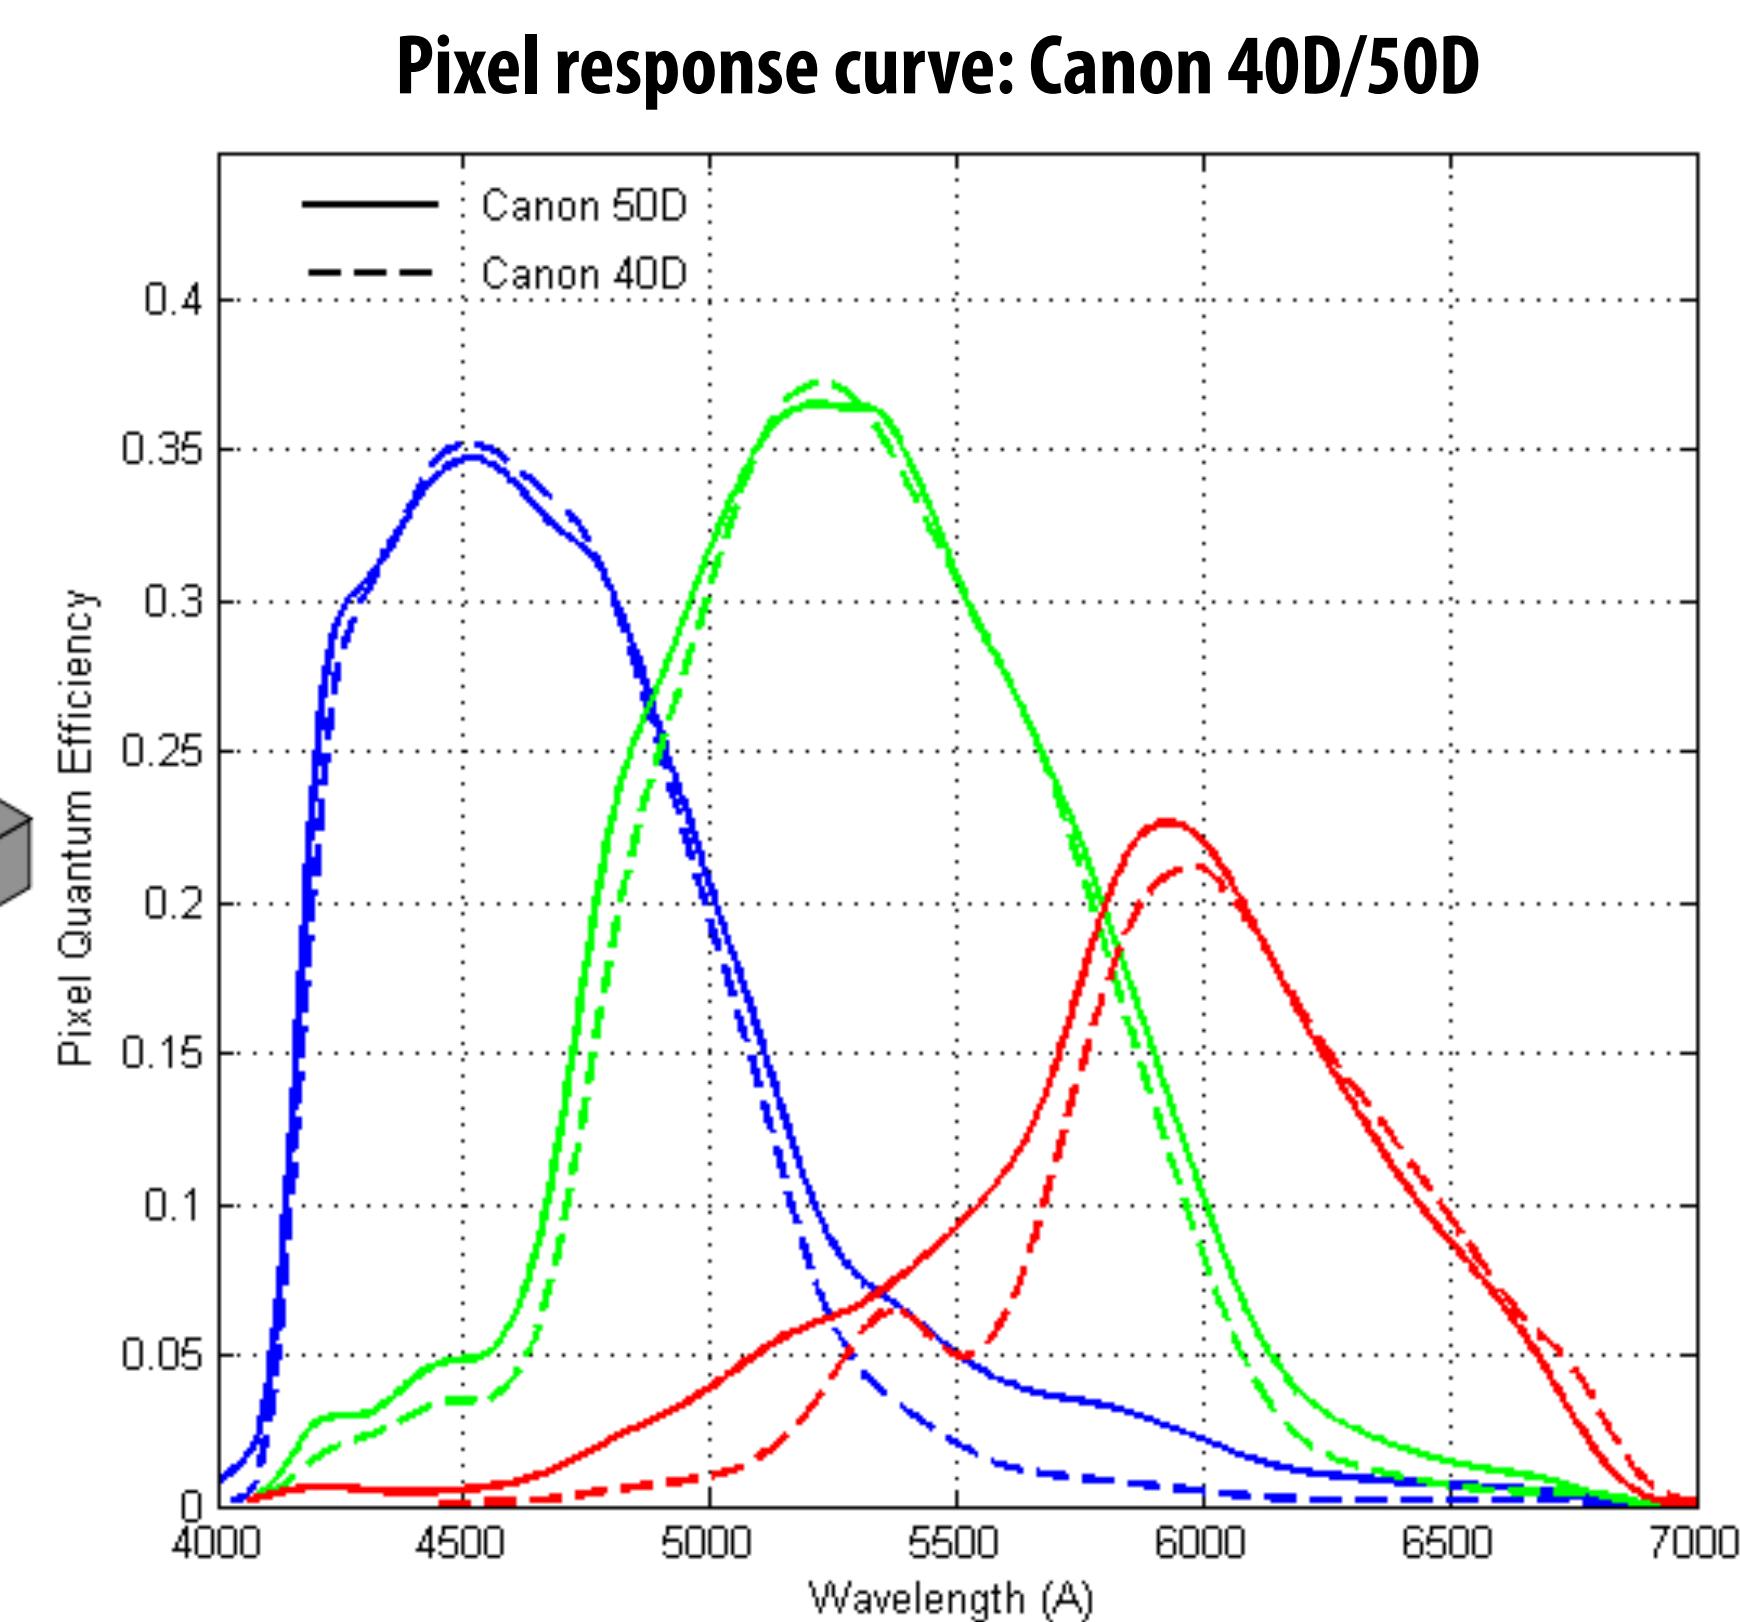
Color filter array (Bayer mosaic)

- Color filter array placed over sensor
- Result: different pixels have different spectral response (each pixel measures red, green, or blue light)
- 50% of pixels are green pixels



Traditional Bayer mosaic

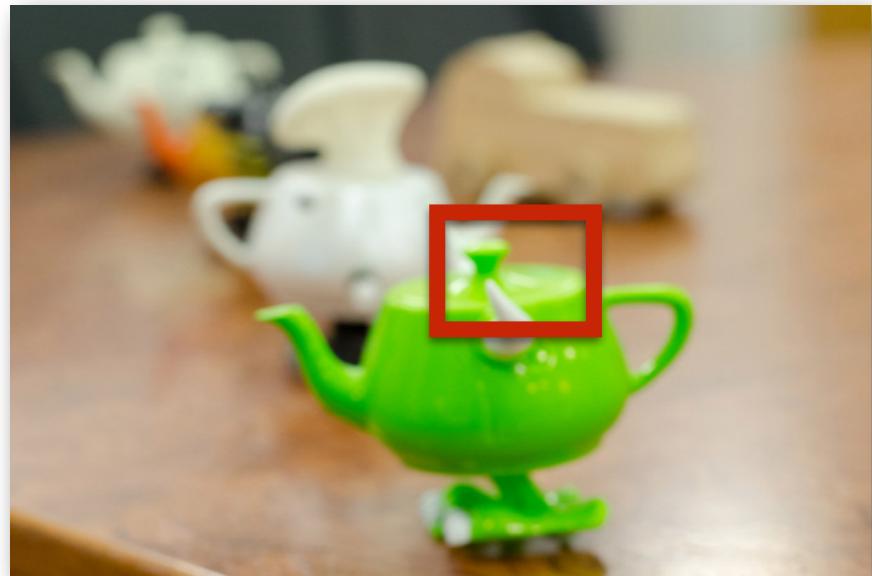
(other filter patterns exist: e.g., Sony's RGBE)



$$f(\lambda)$$

RAW sensor output (simulated data)

Light Hitting Sensor



RAW output of sensor





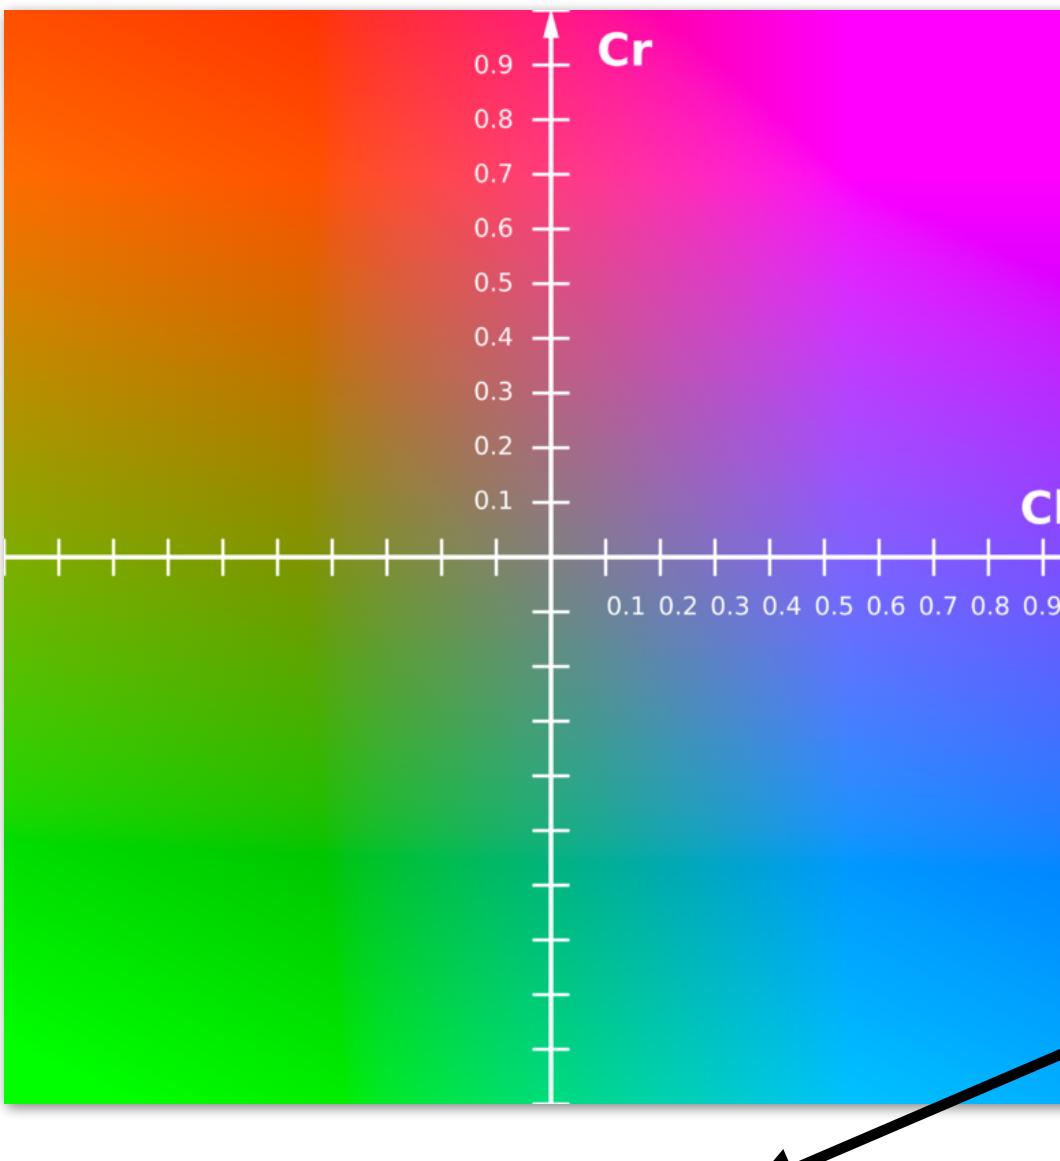
Another basis: Y'CbCr color space

Recall: color is represented as a 3D value

Y' = luma: perceived luminance

C_b = blue-yellow deviation from gray

C_r = red-cyan deviation from gray



“Gamma corrected” RGB
(primed notation indicates perceptual (non-linear) space)
We’ll describe what this means this later in the lecture.

Conversion from $R'G'B'$ to $Y'CbCr$:

$$Y' = 16 + \frac{65.738 \cdot R'_D}{256} + \frac{129.057 \cdot G'_D}{256} + \frac{25.064 \cdot B'_D}{256}$$

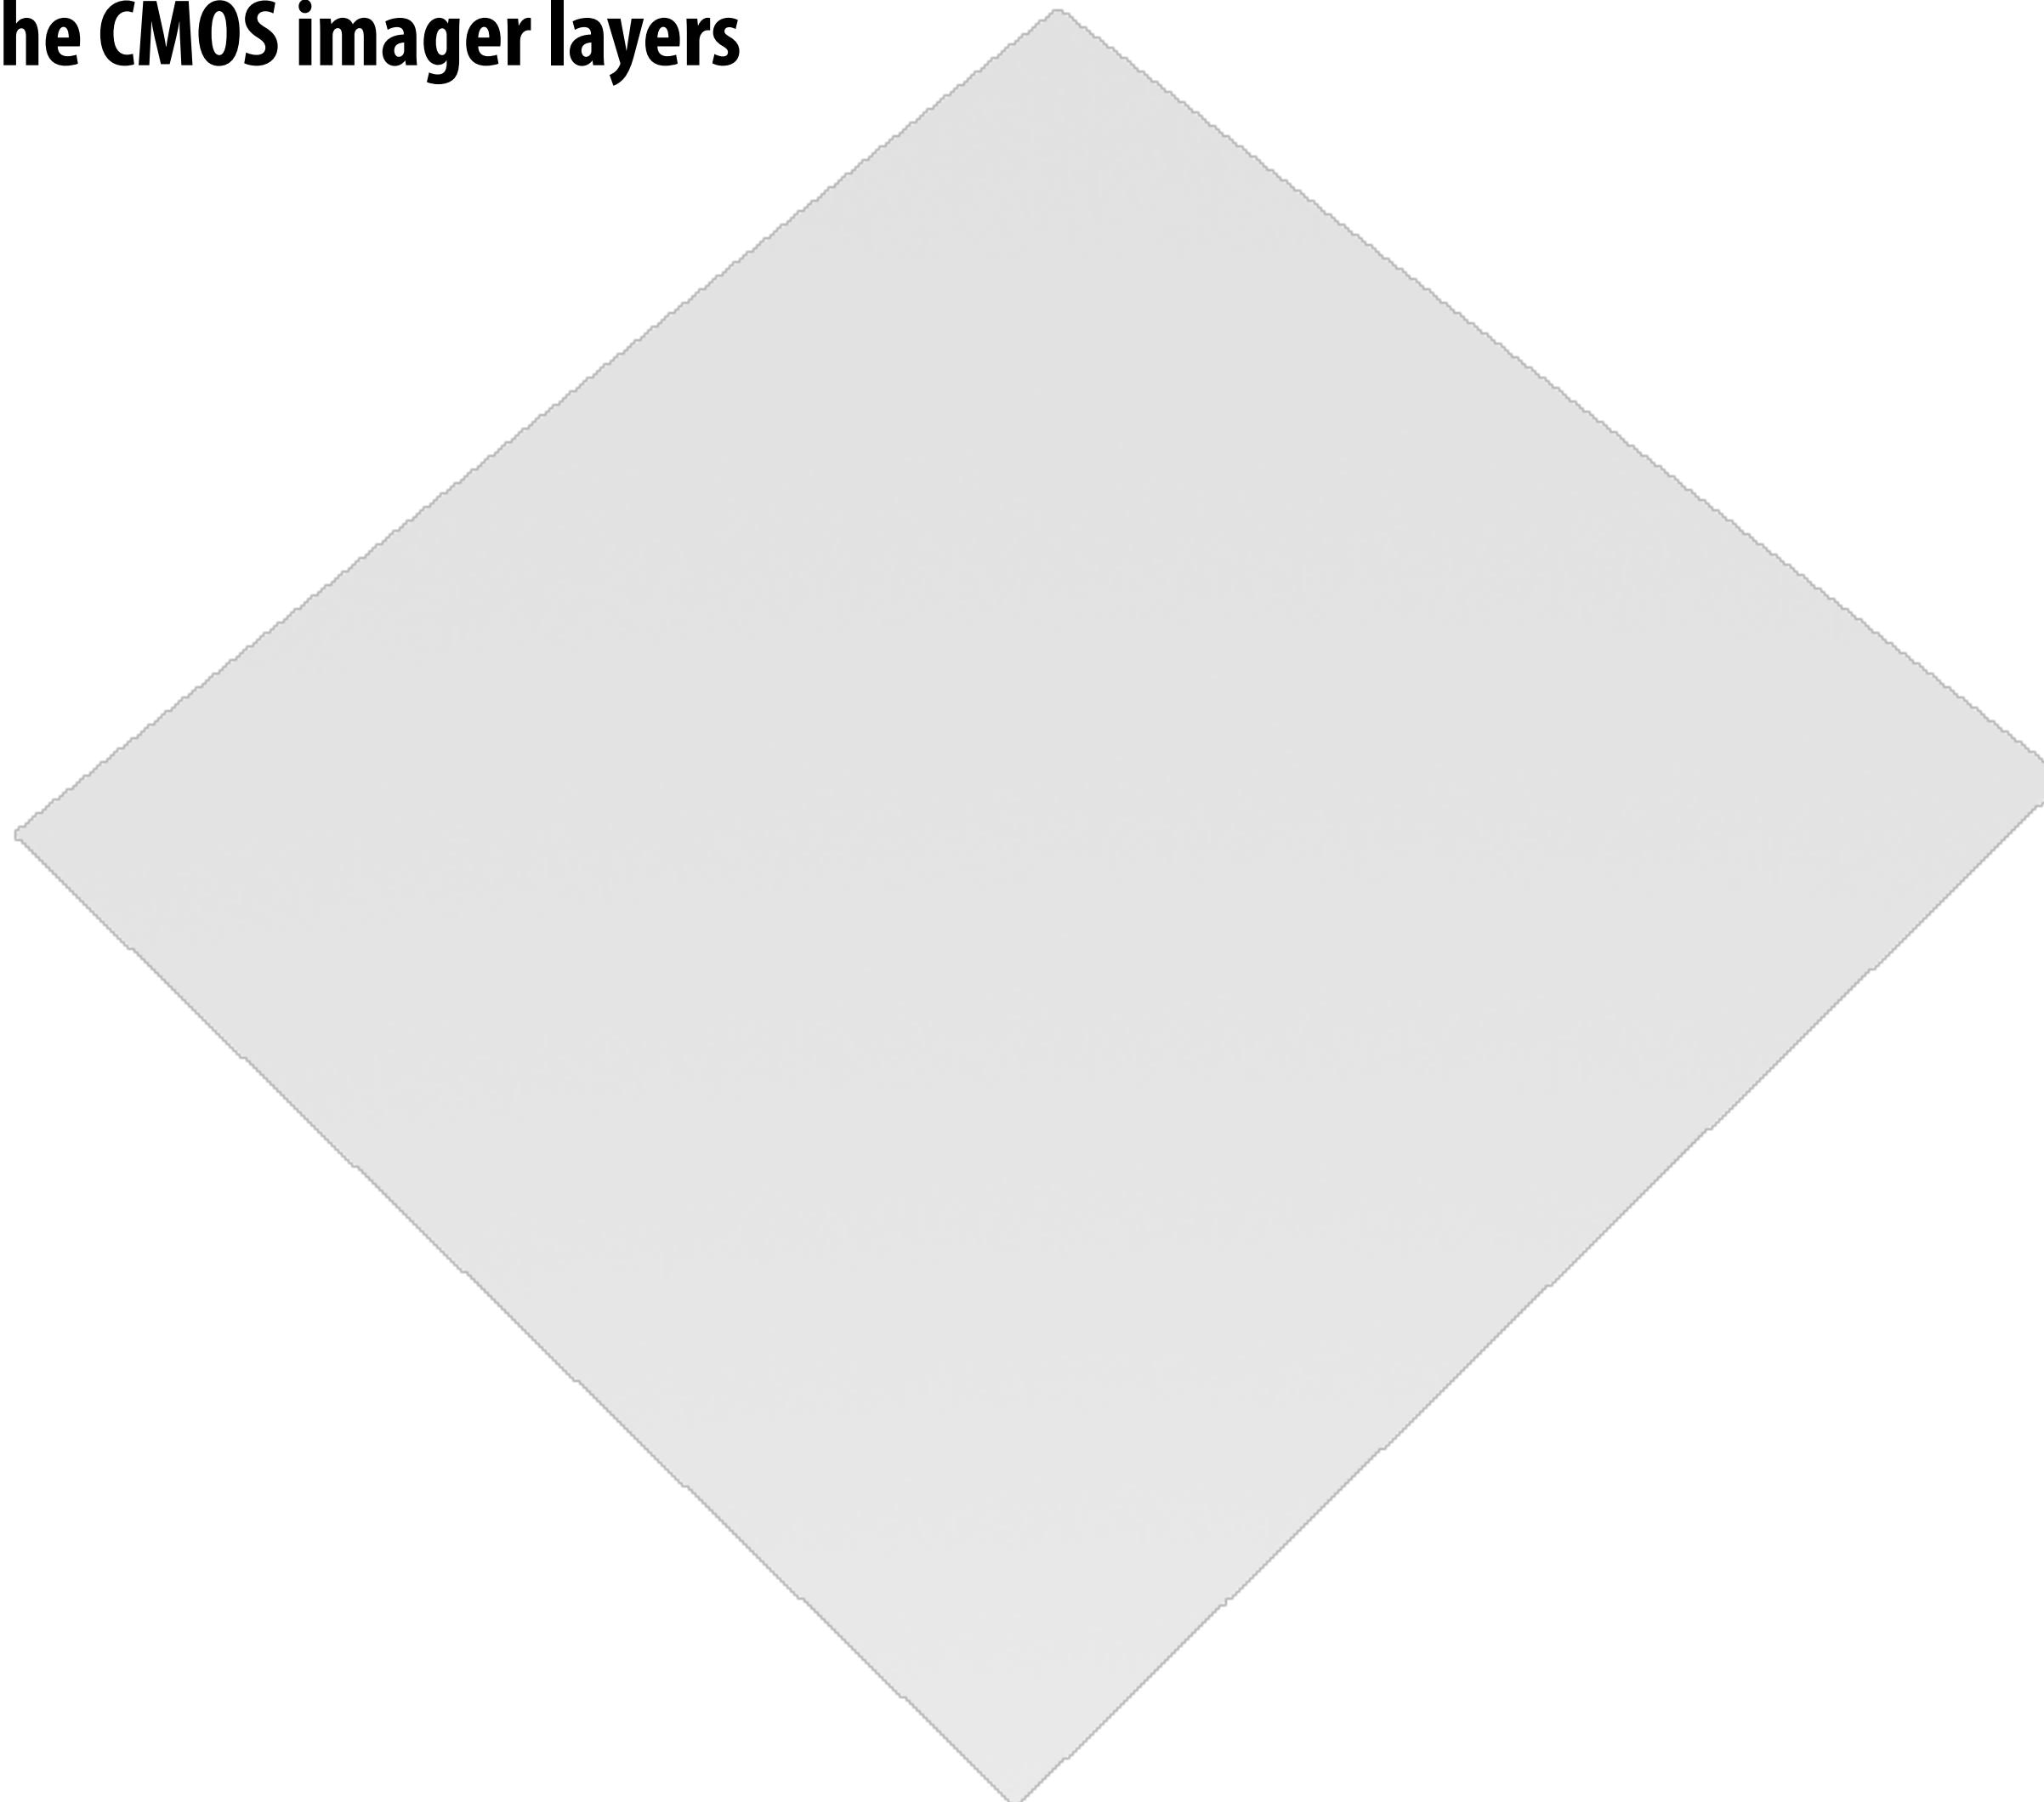
$$C_B = 128 + \frac{-37.945 \cdot R'_D}{256} - \frac{74.494 \cdot G'_D}{256} + \frac{112.439 \cdot B'_D}{256}$$

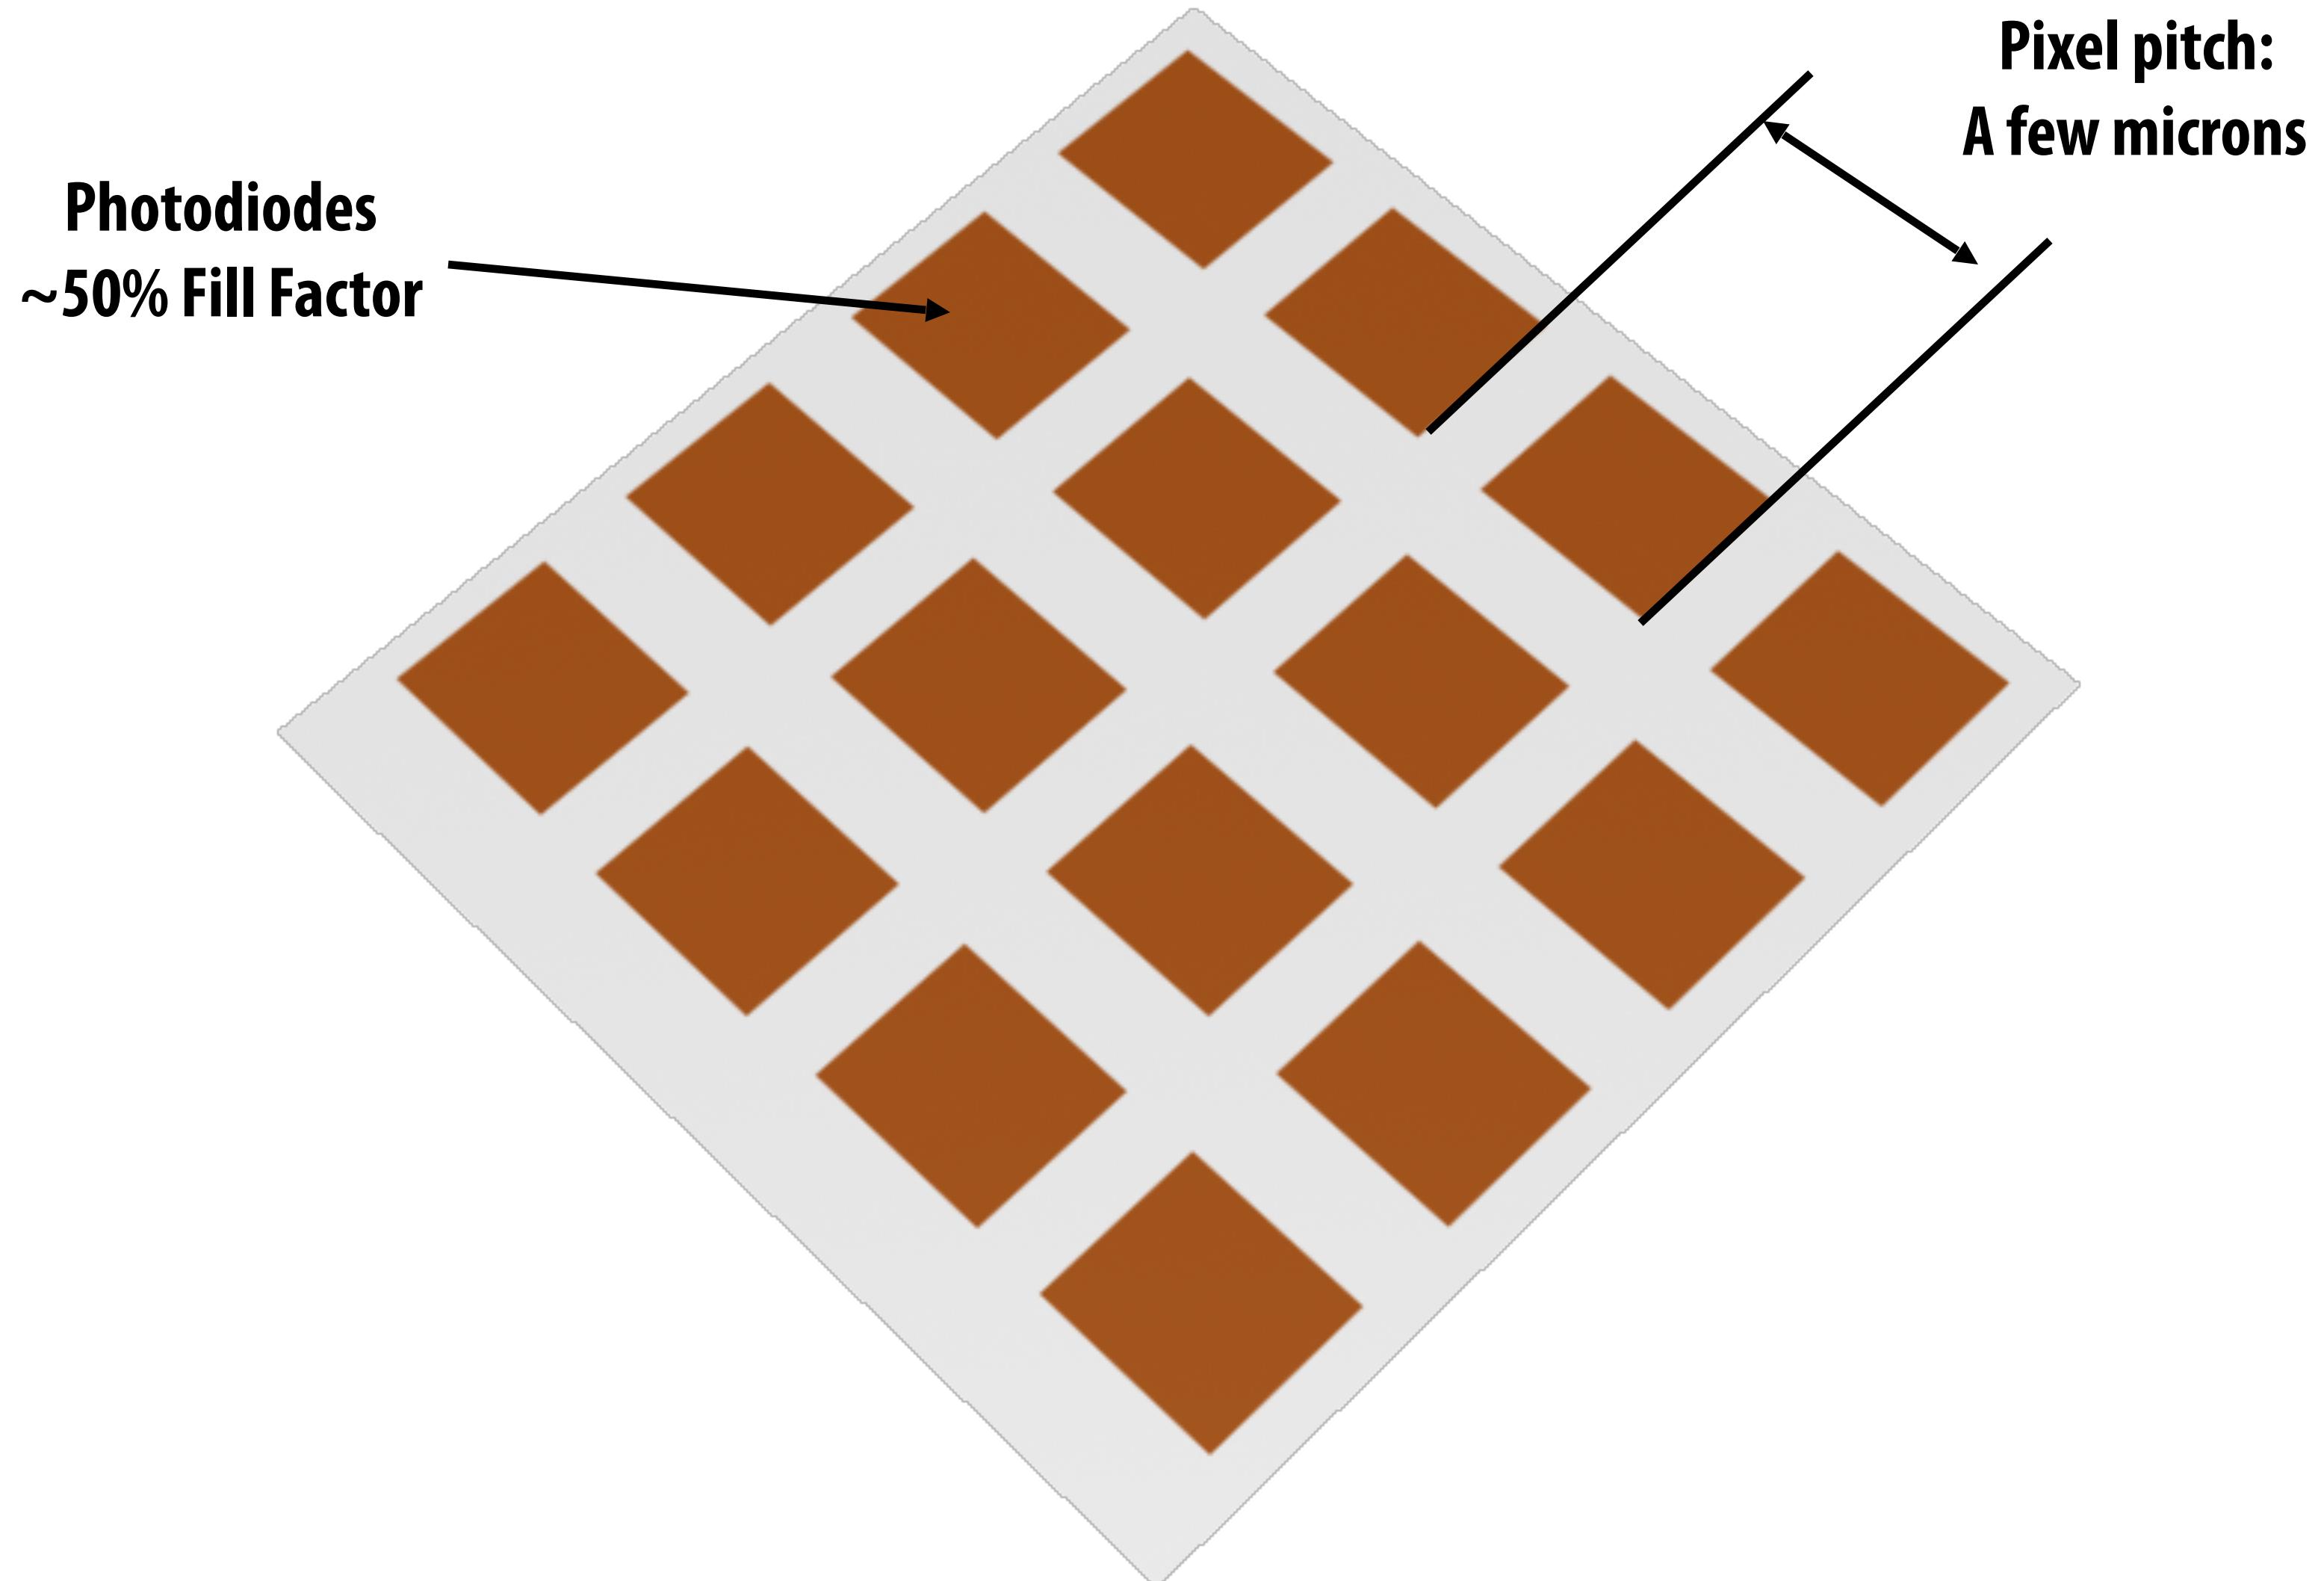
$$C_R = 128 + \frac{112.439 \cdot R'_D}{256} - \frac{94.154 \cdot G'_D}{256} - \frac{18.285 \cdot B'_D}{256}$$

CMOS Pixel Structure

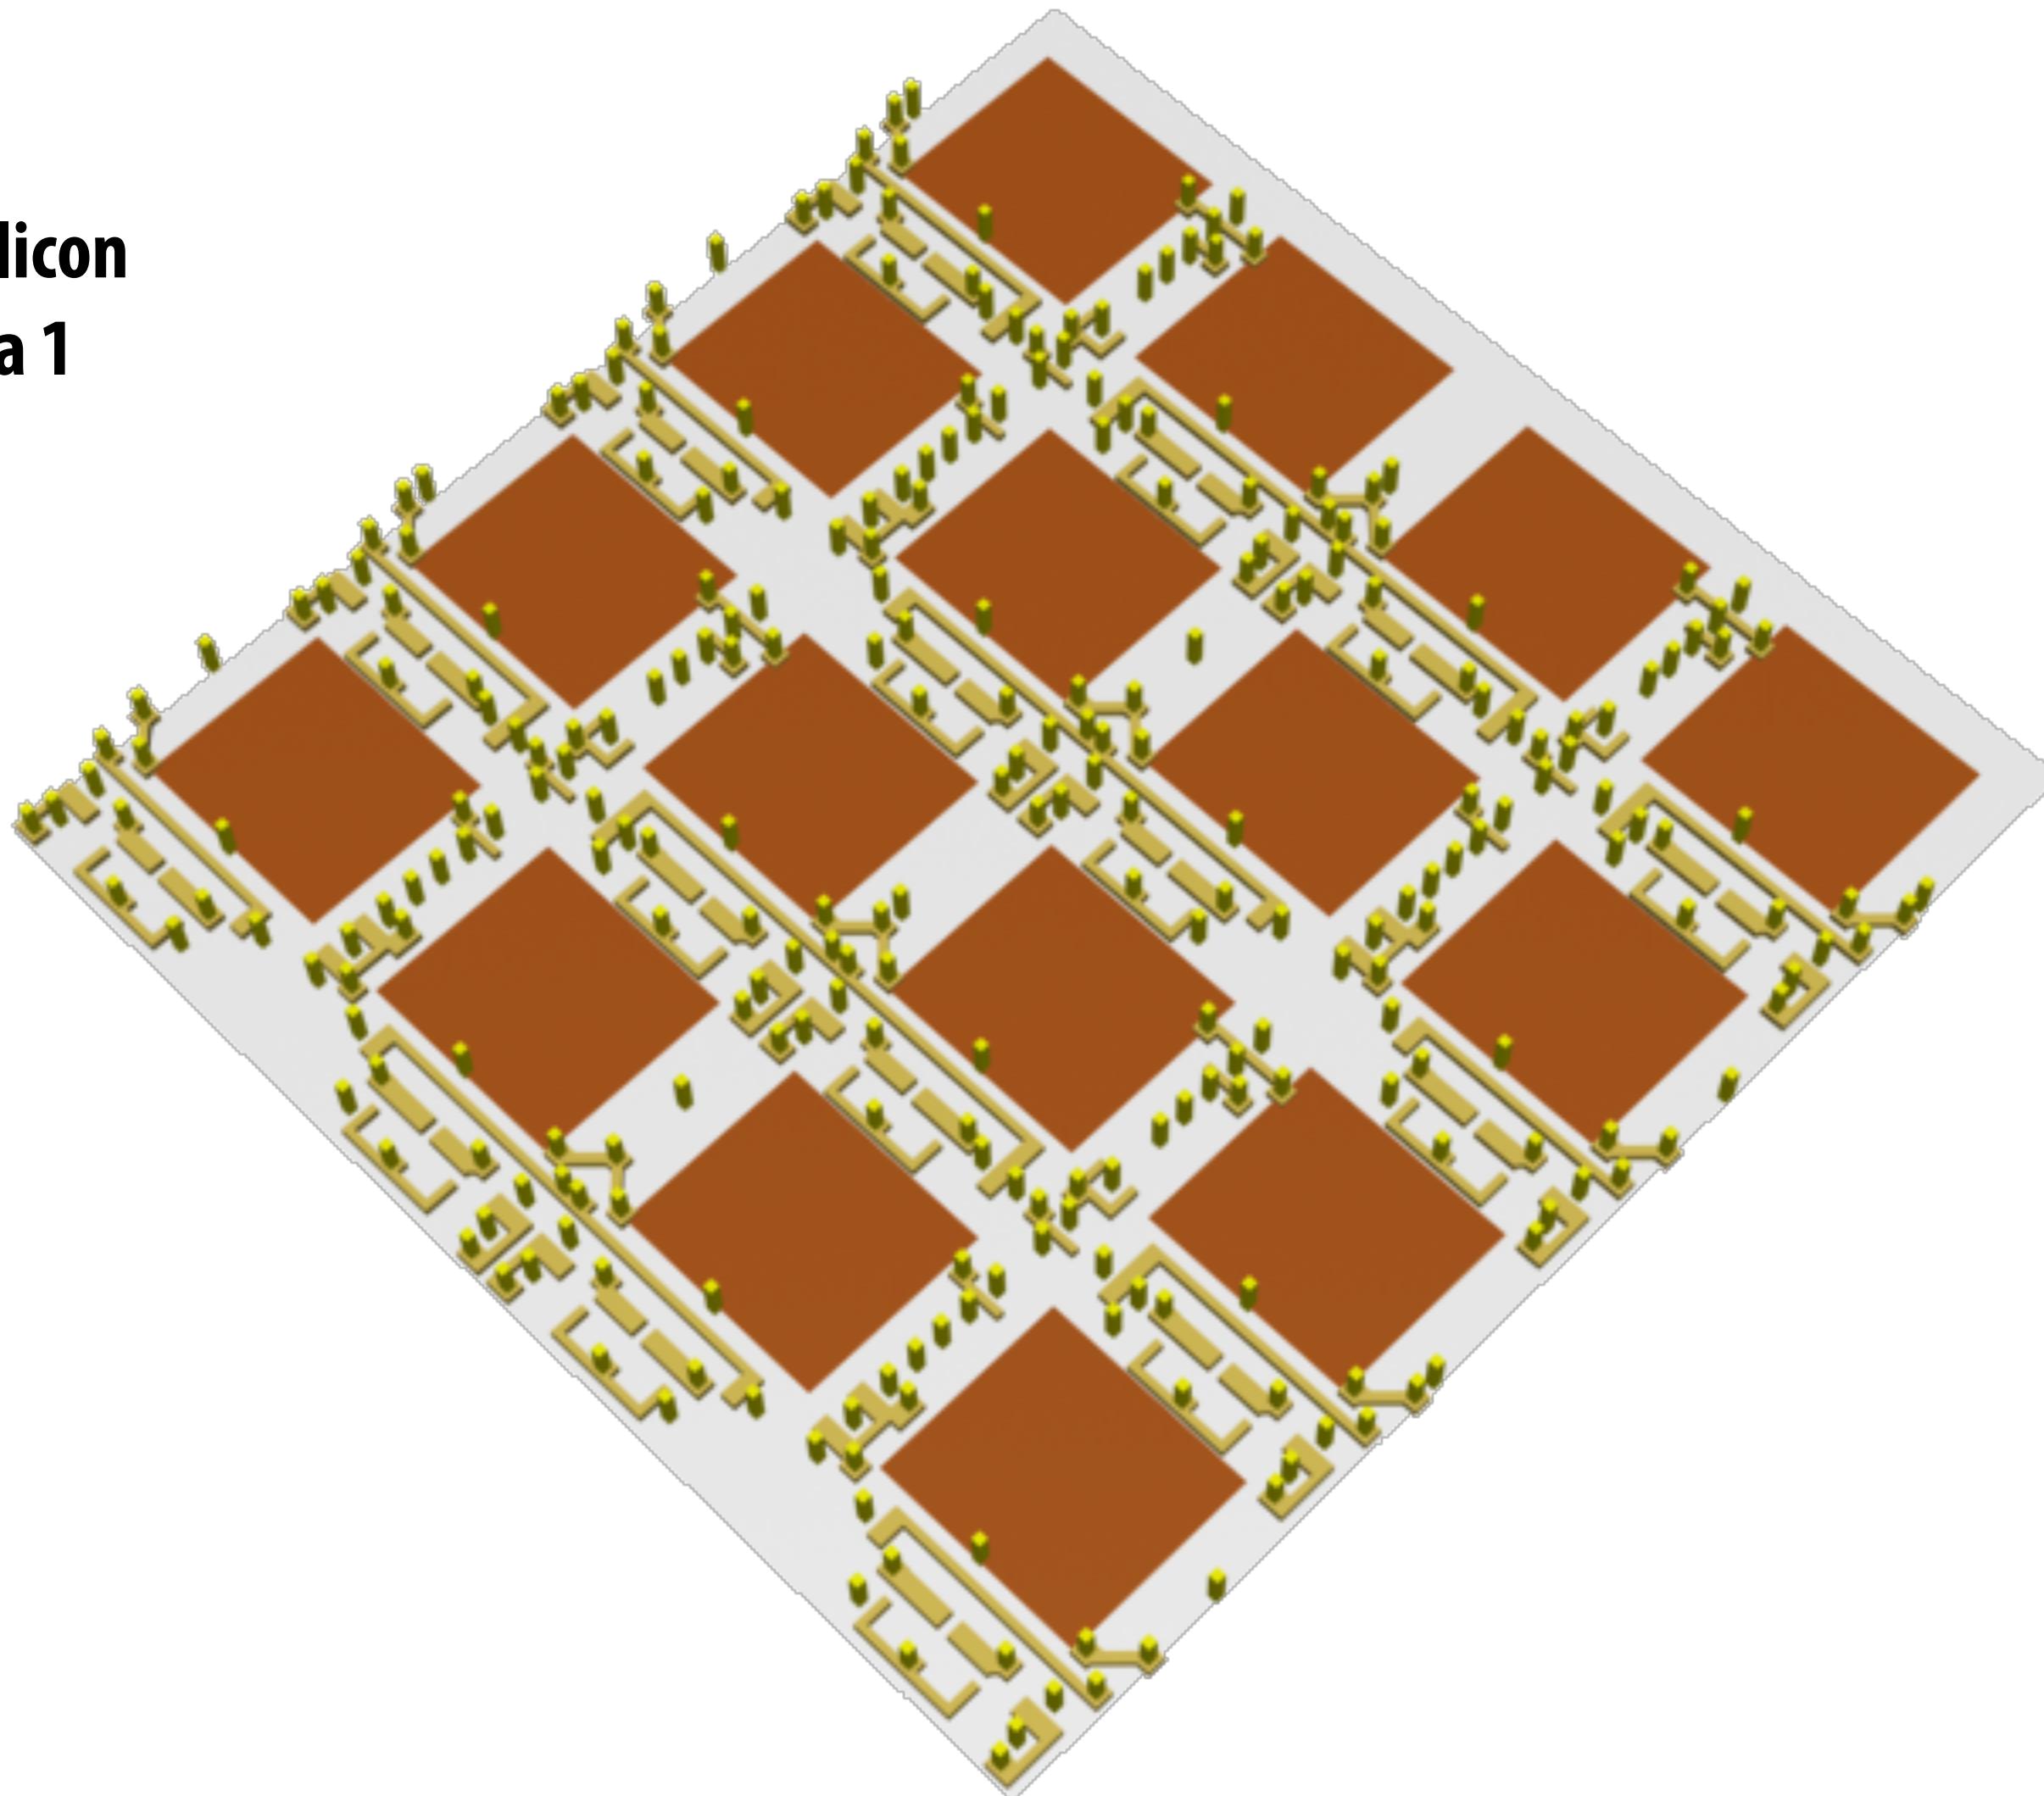
Front-side-illuminated (FSI) CMOS

Building up the CMOS imager layers

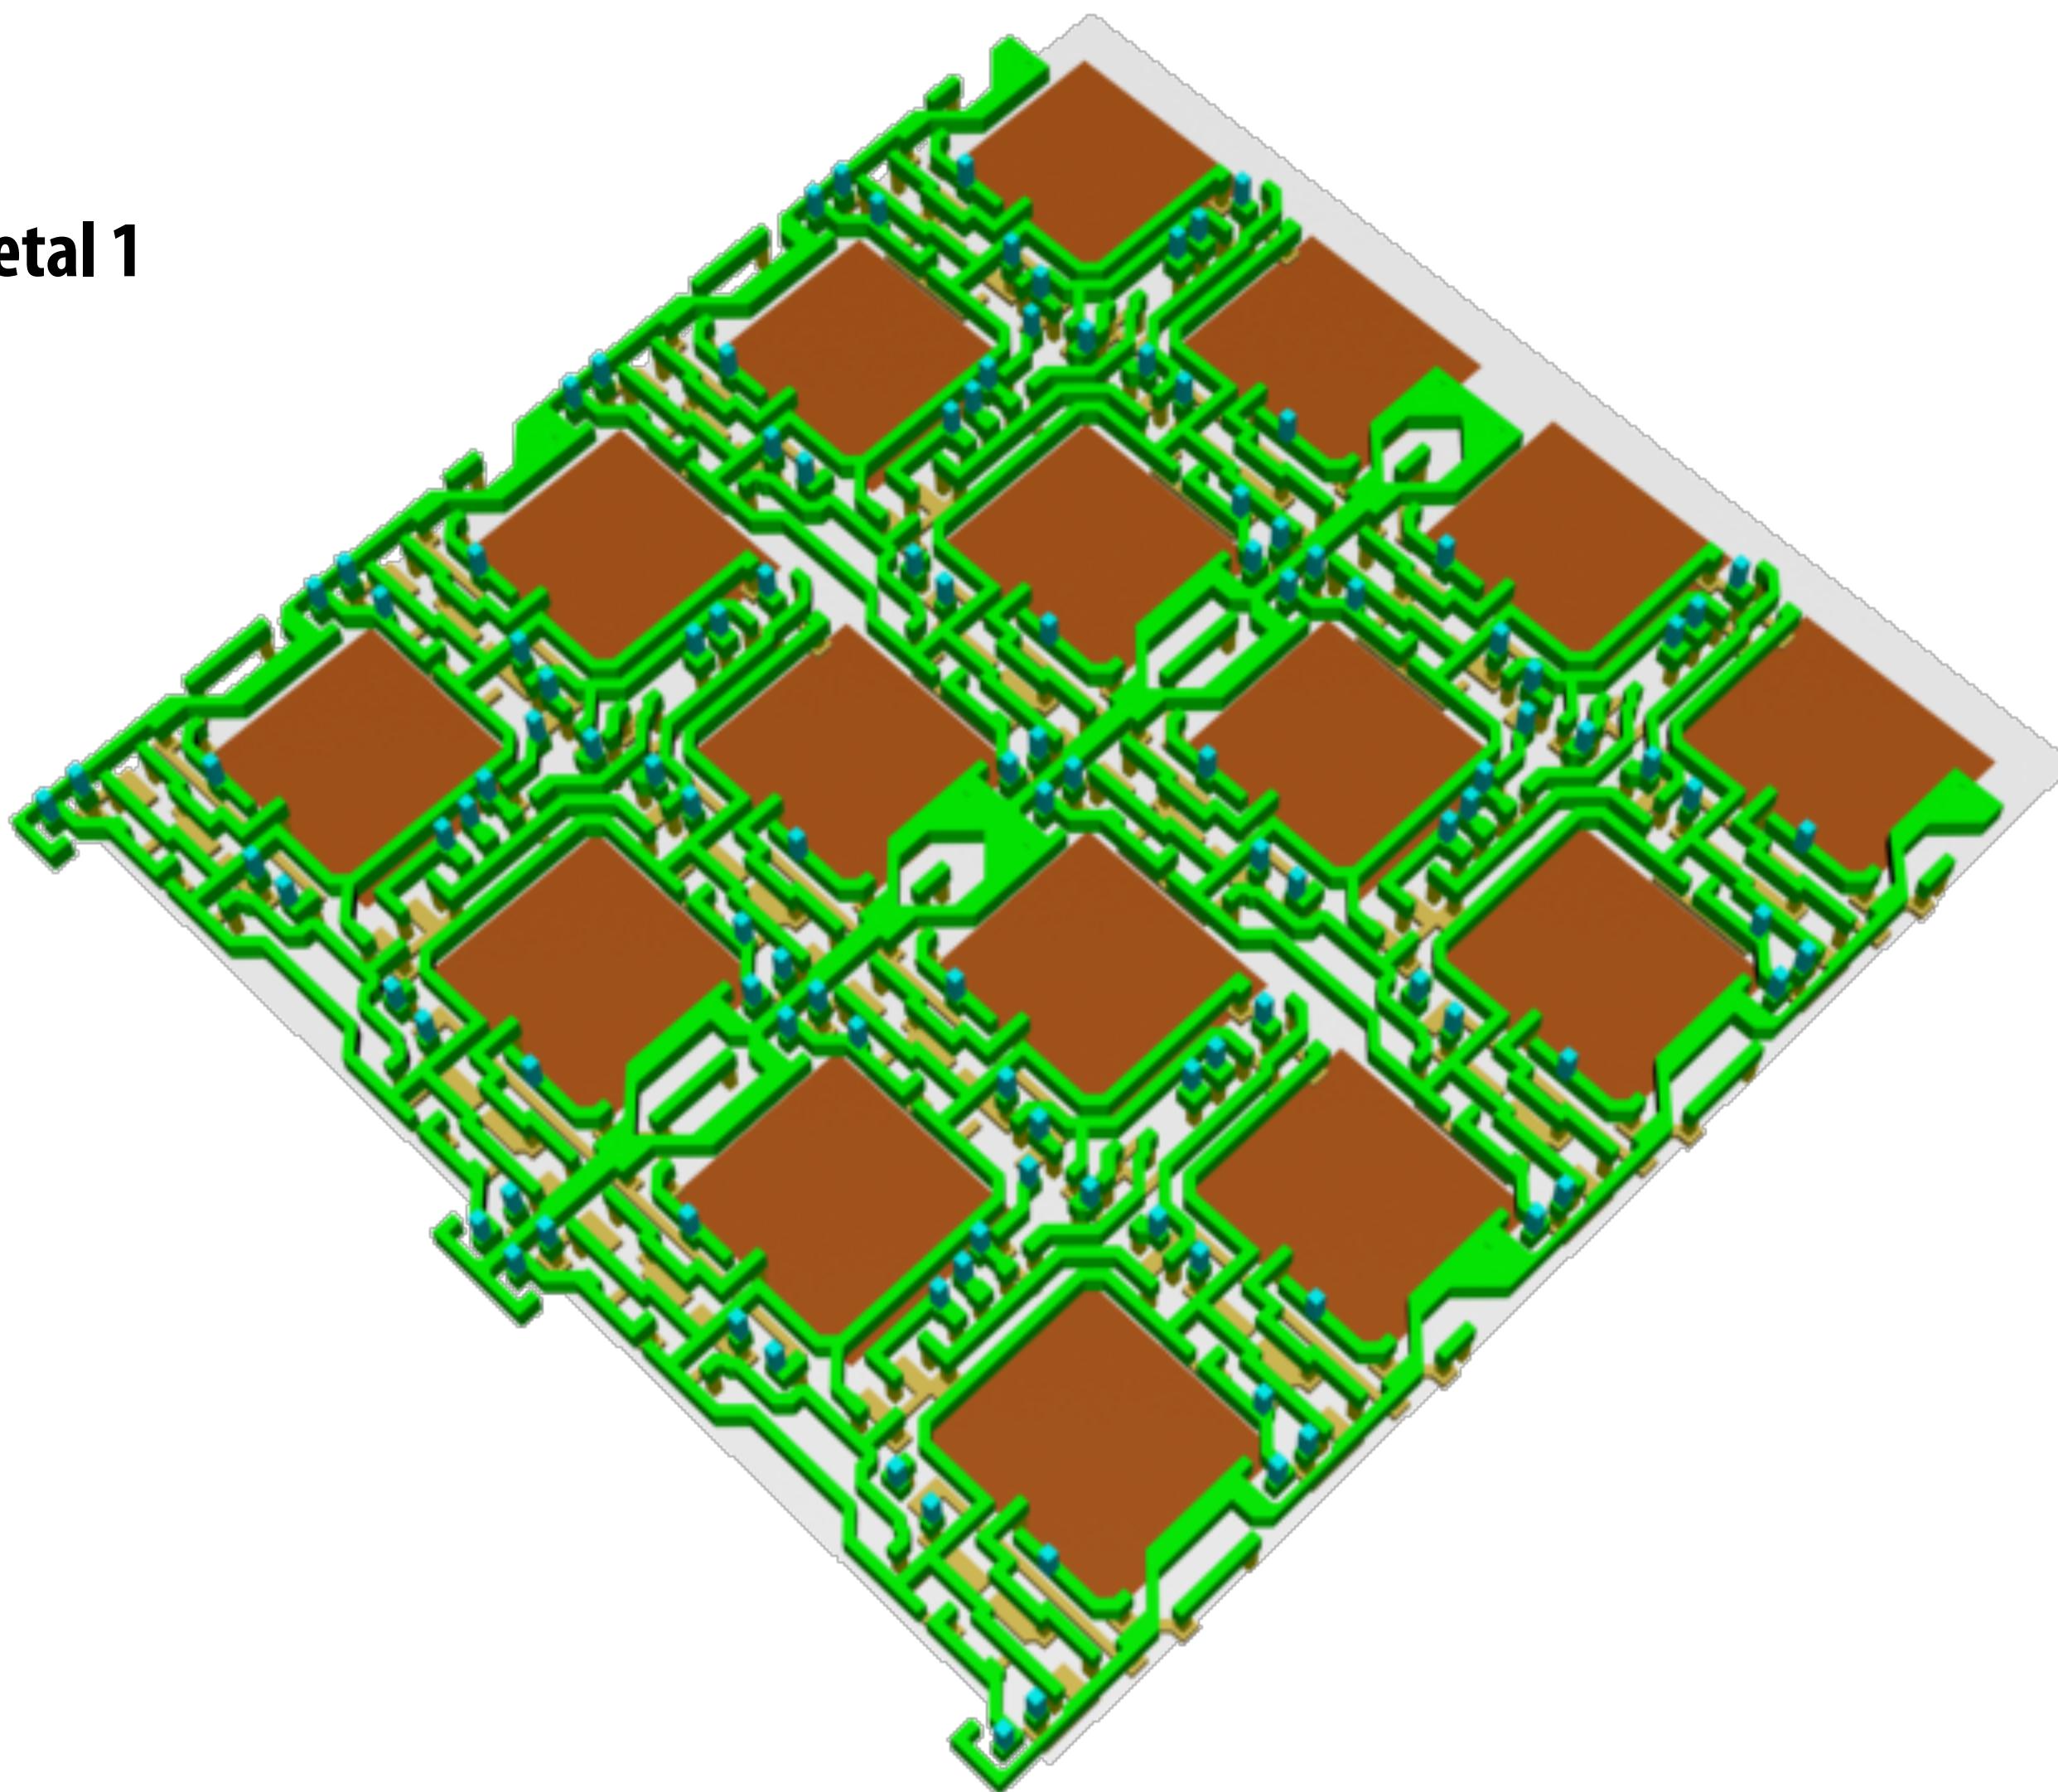




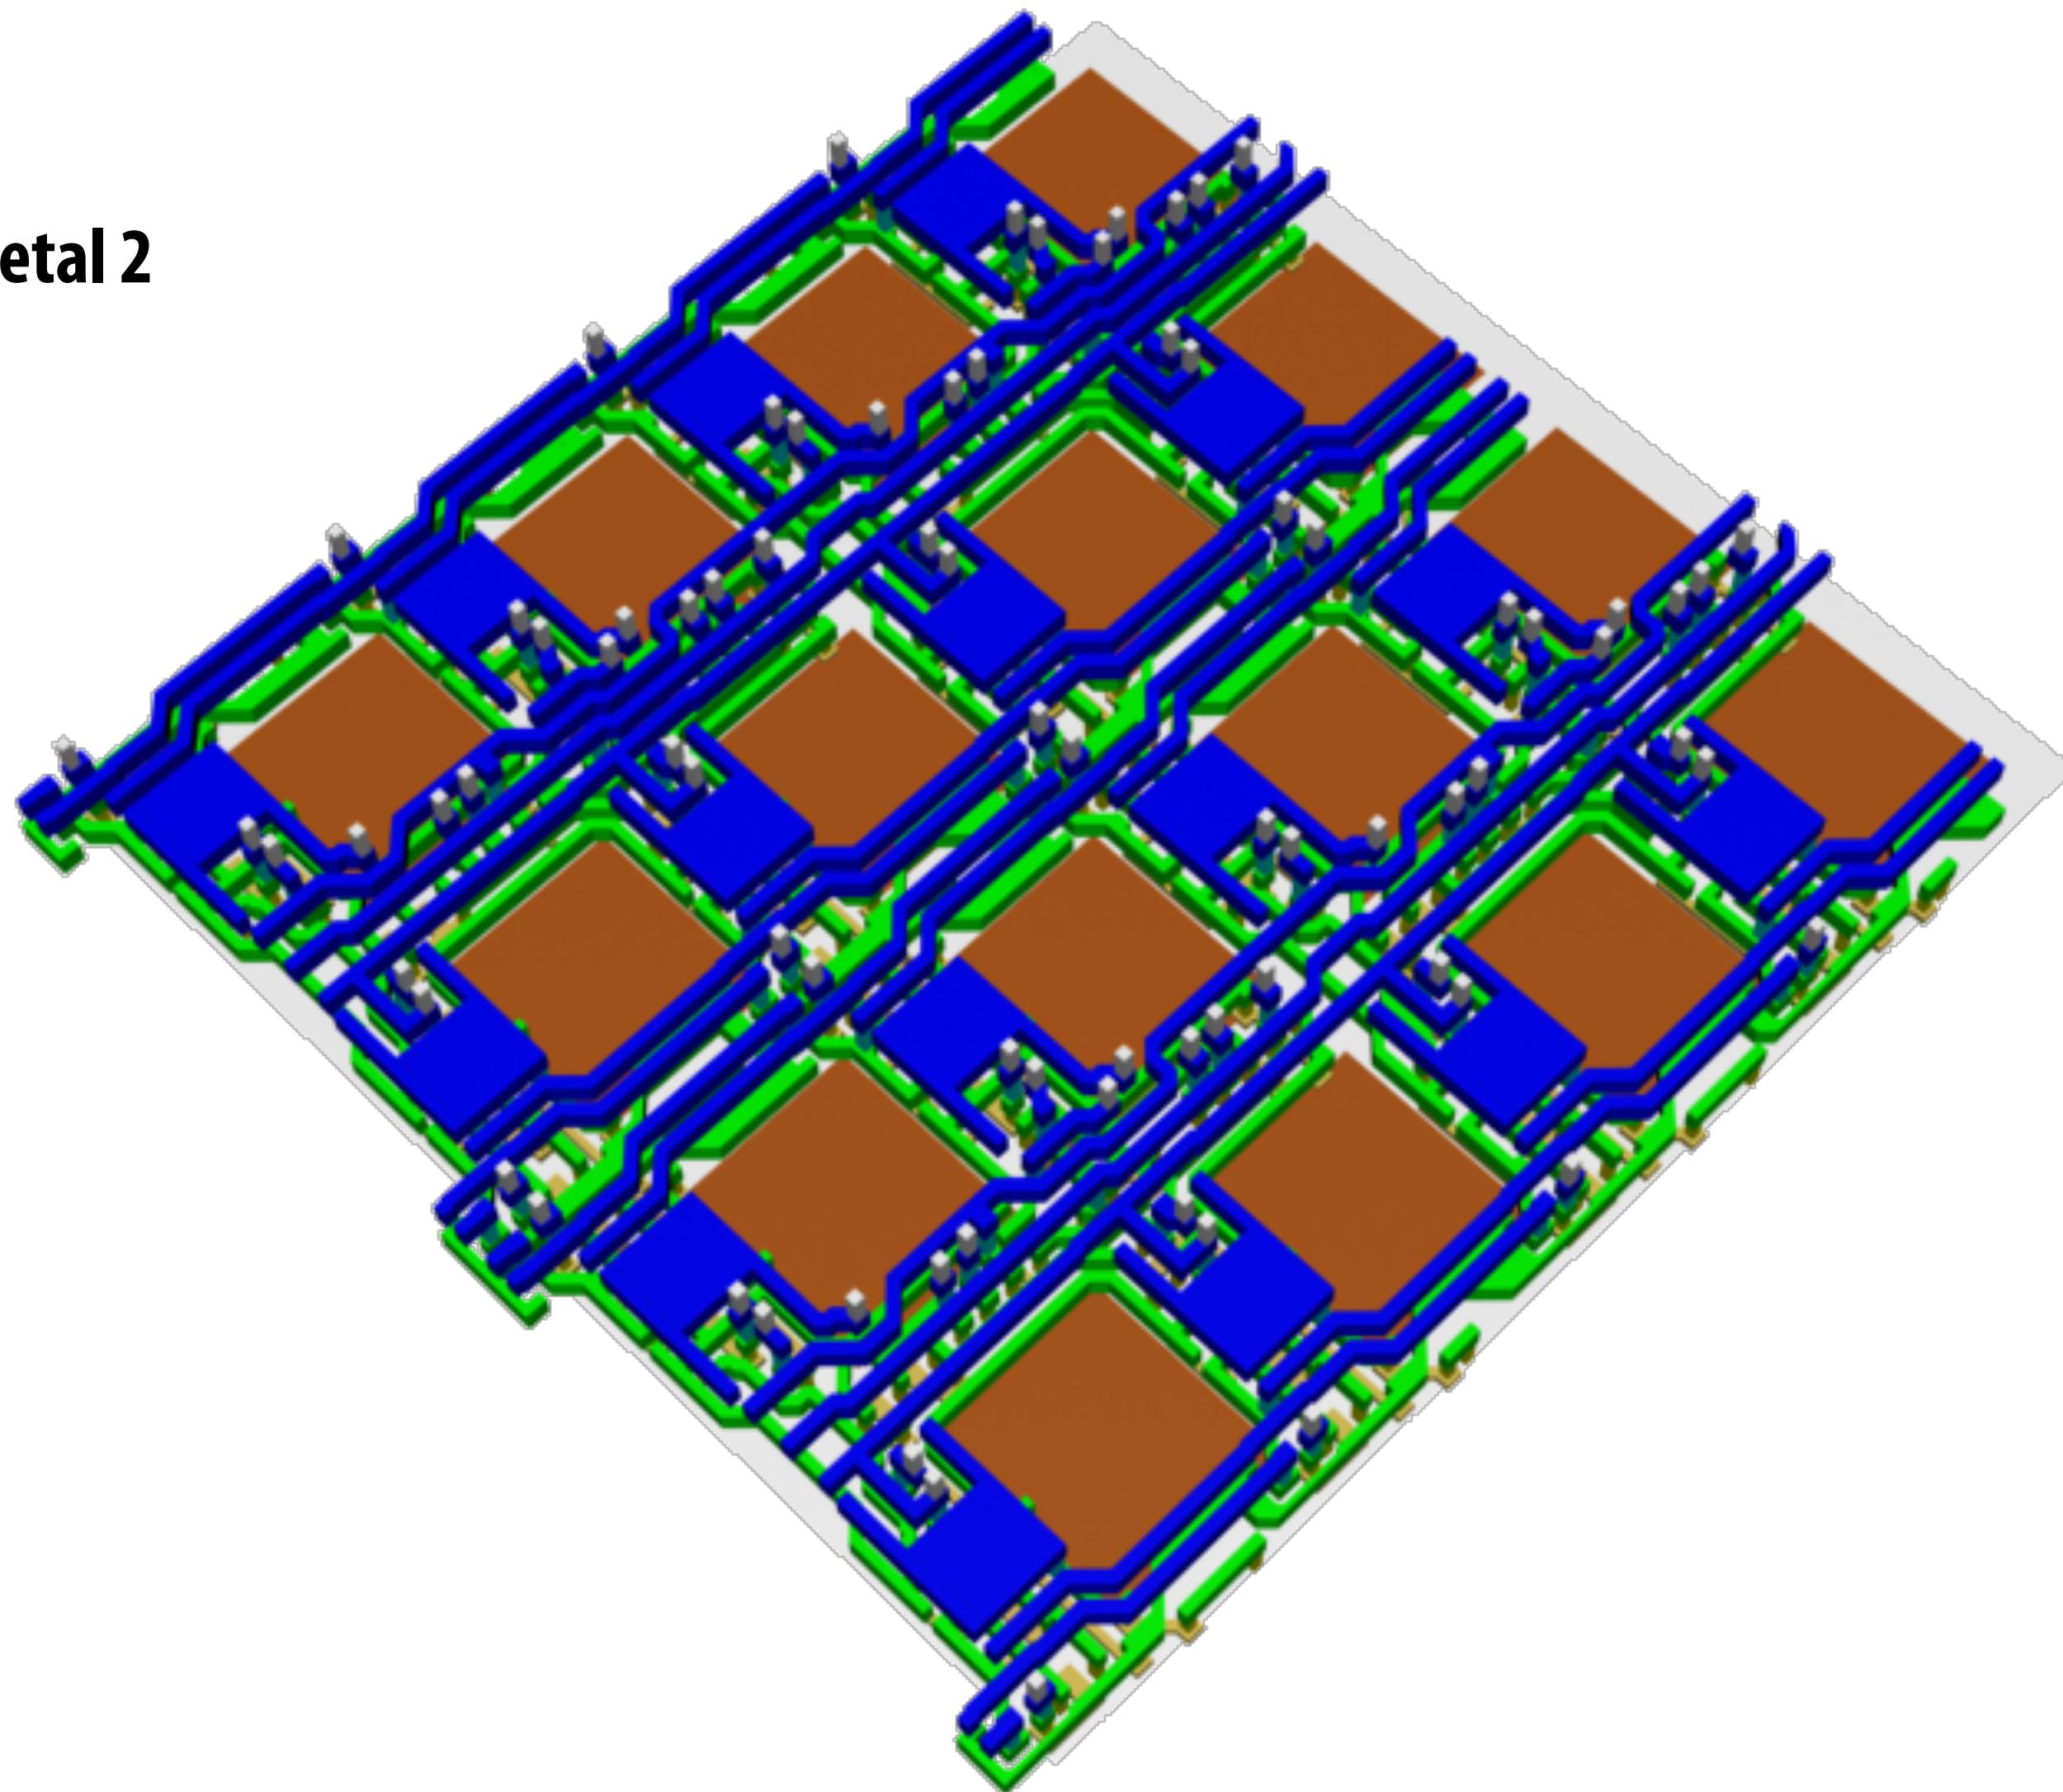
Polysilicon & Via 1



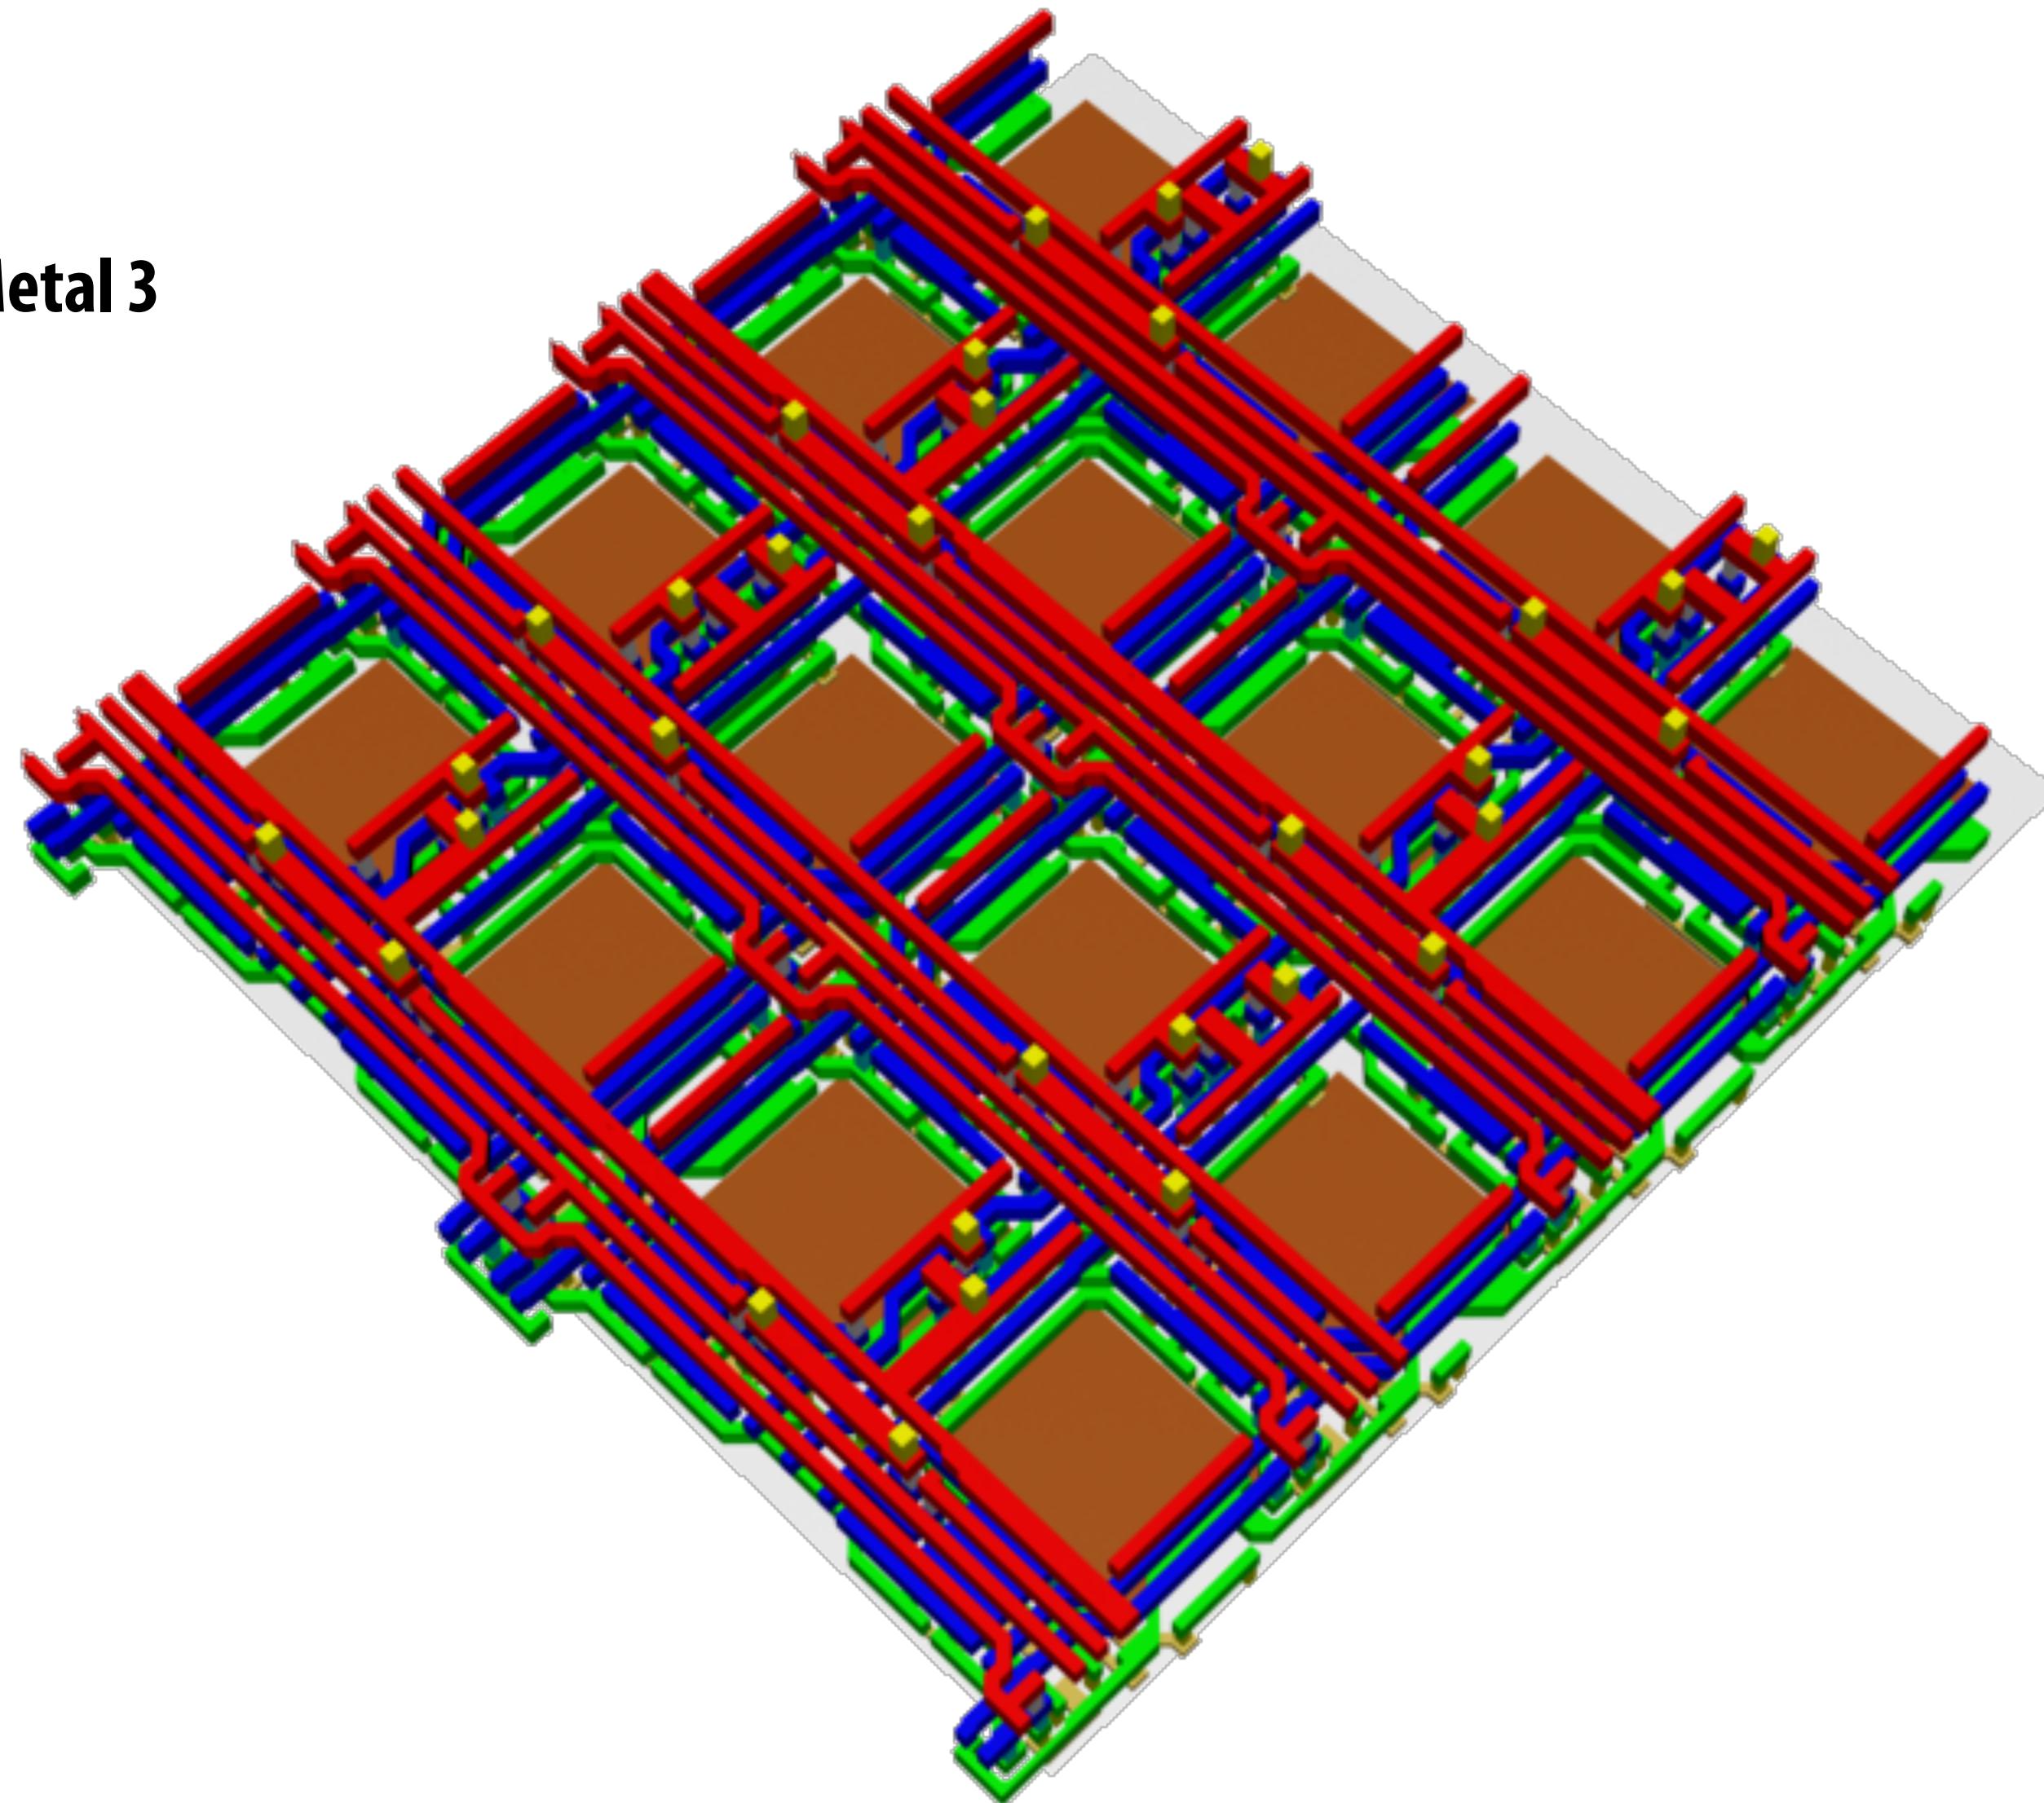
Metal 1



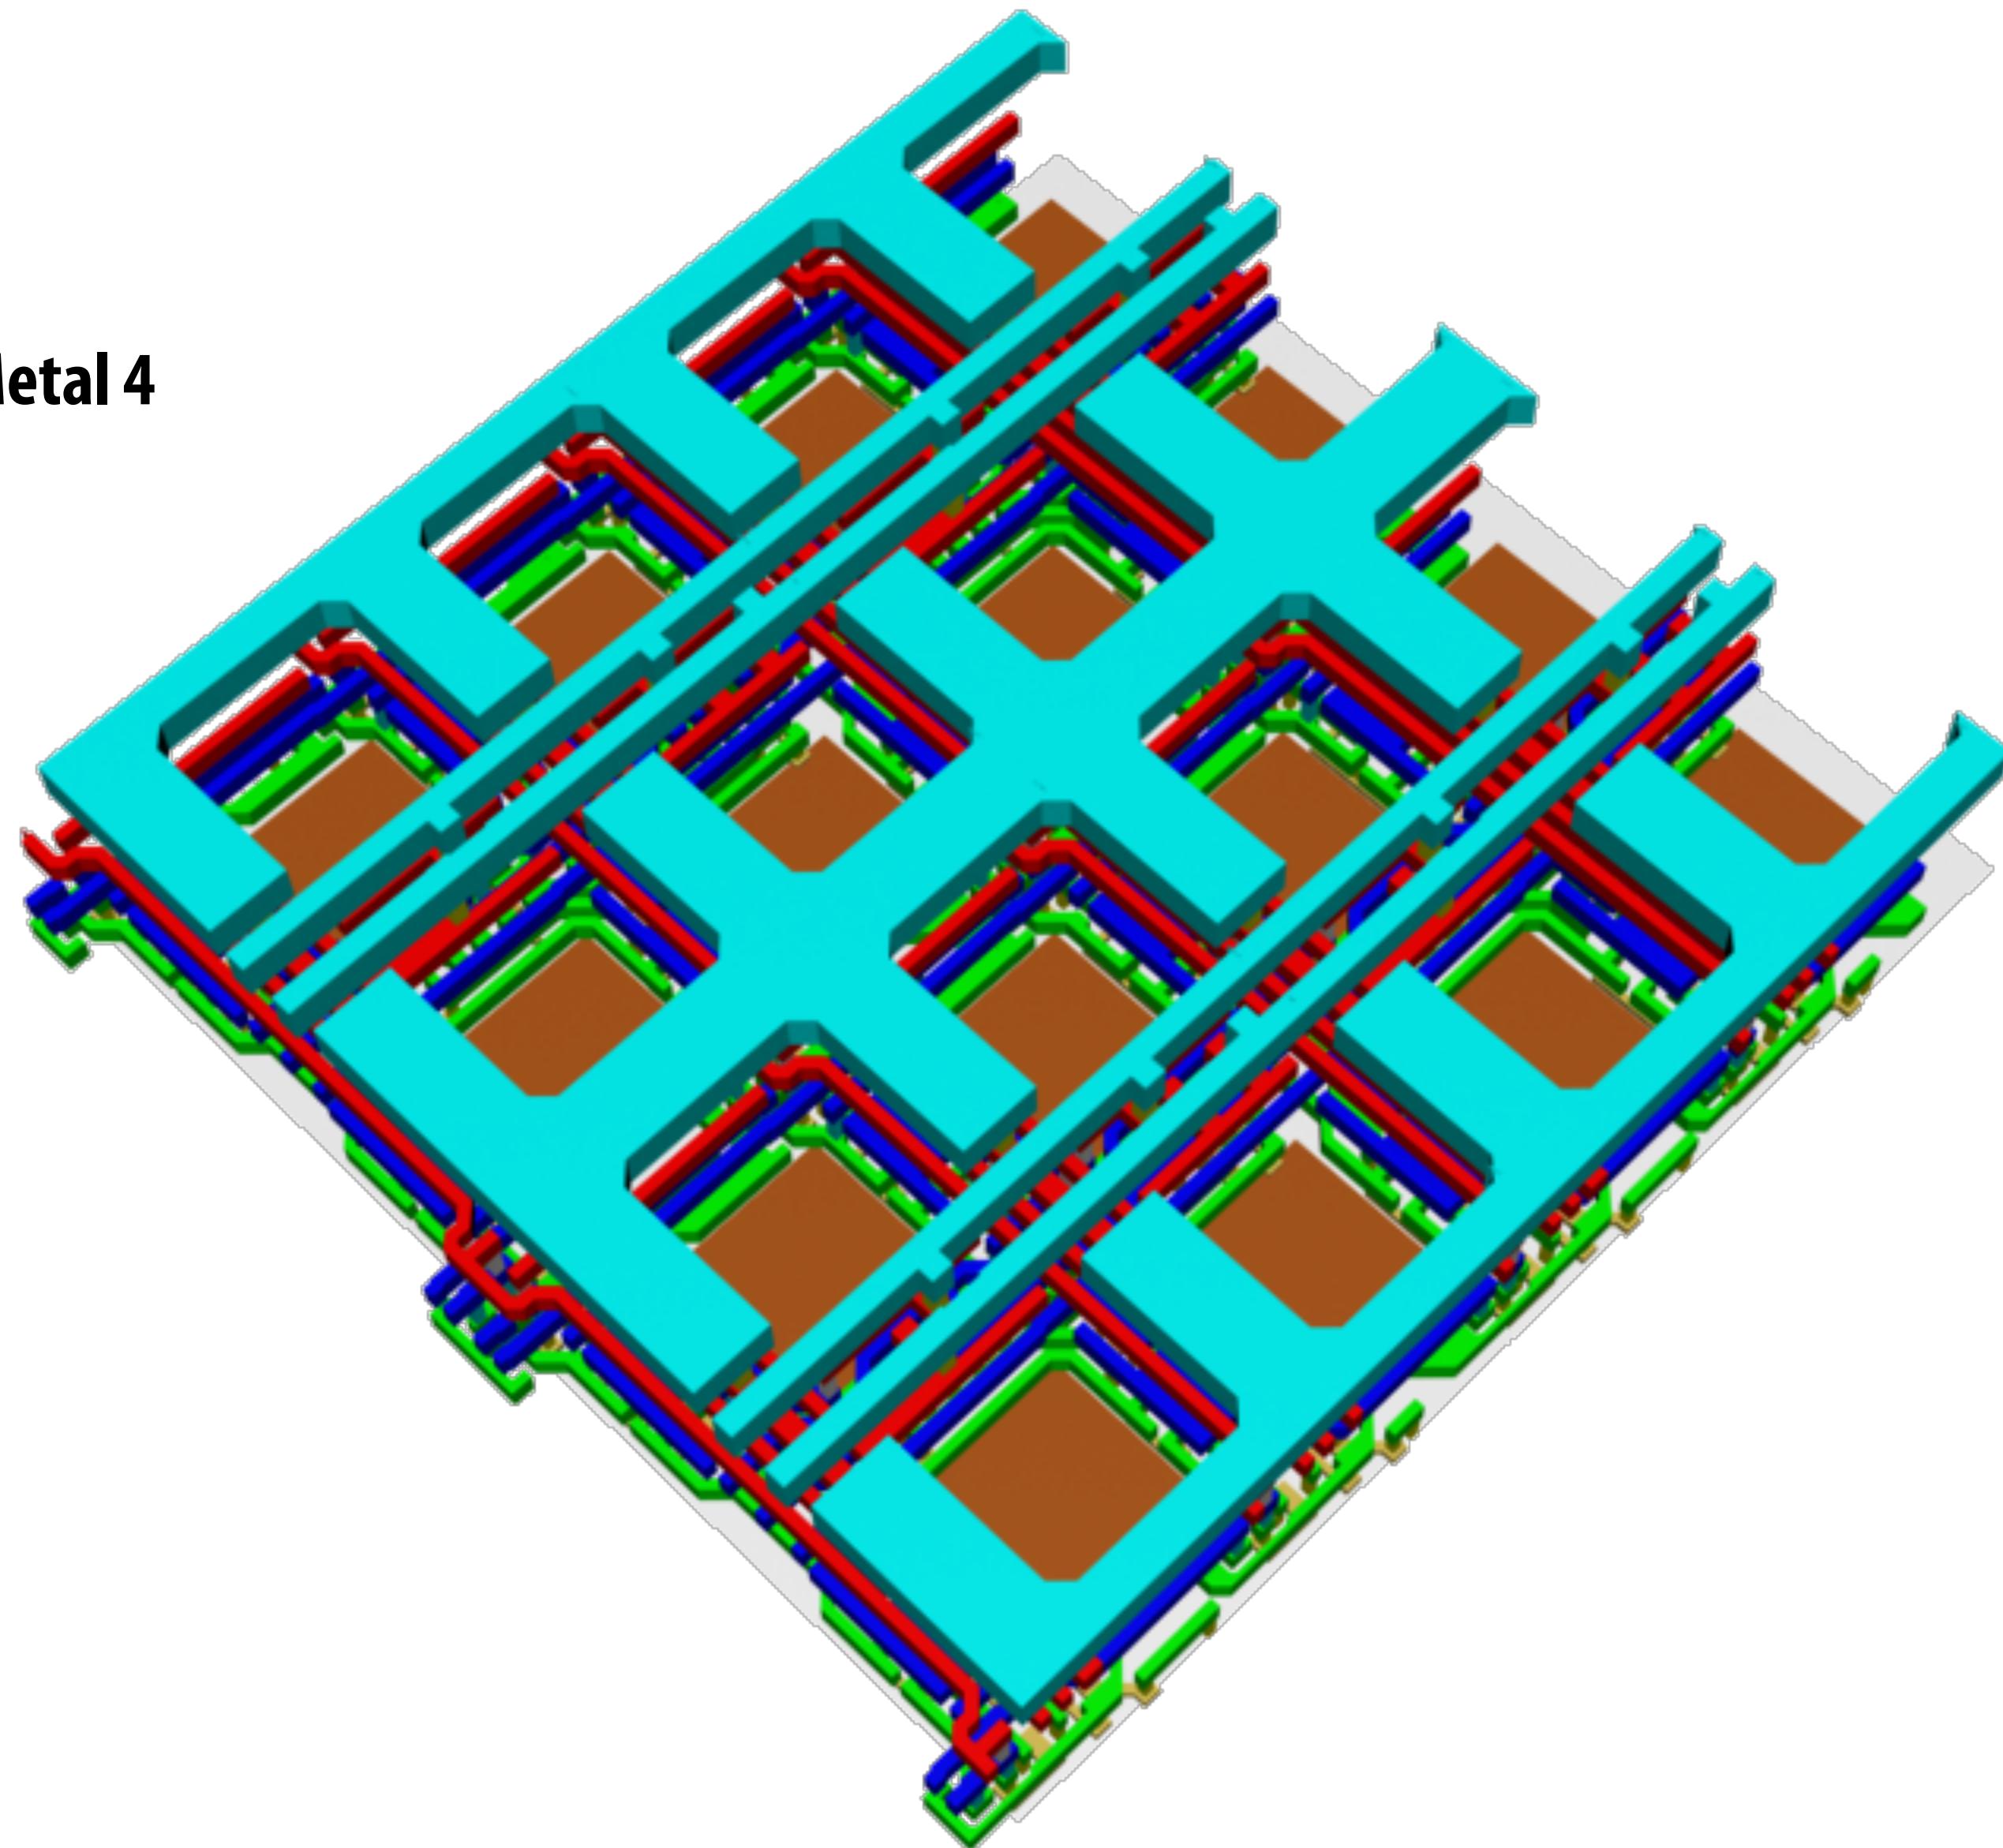
Metal 2



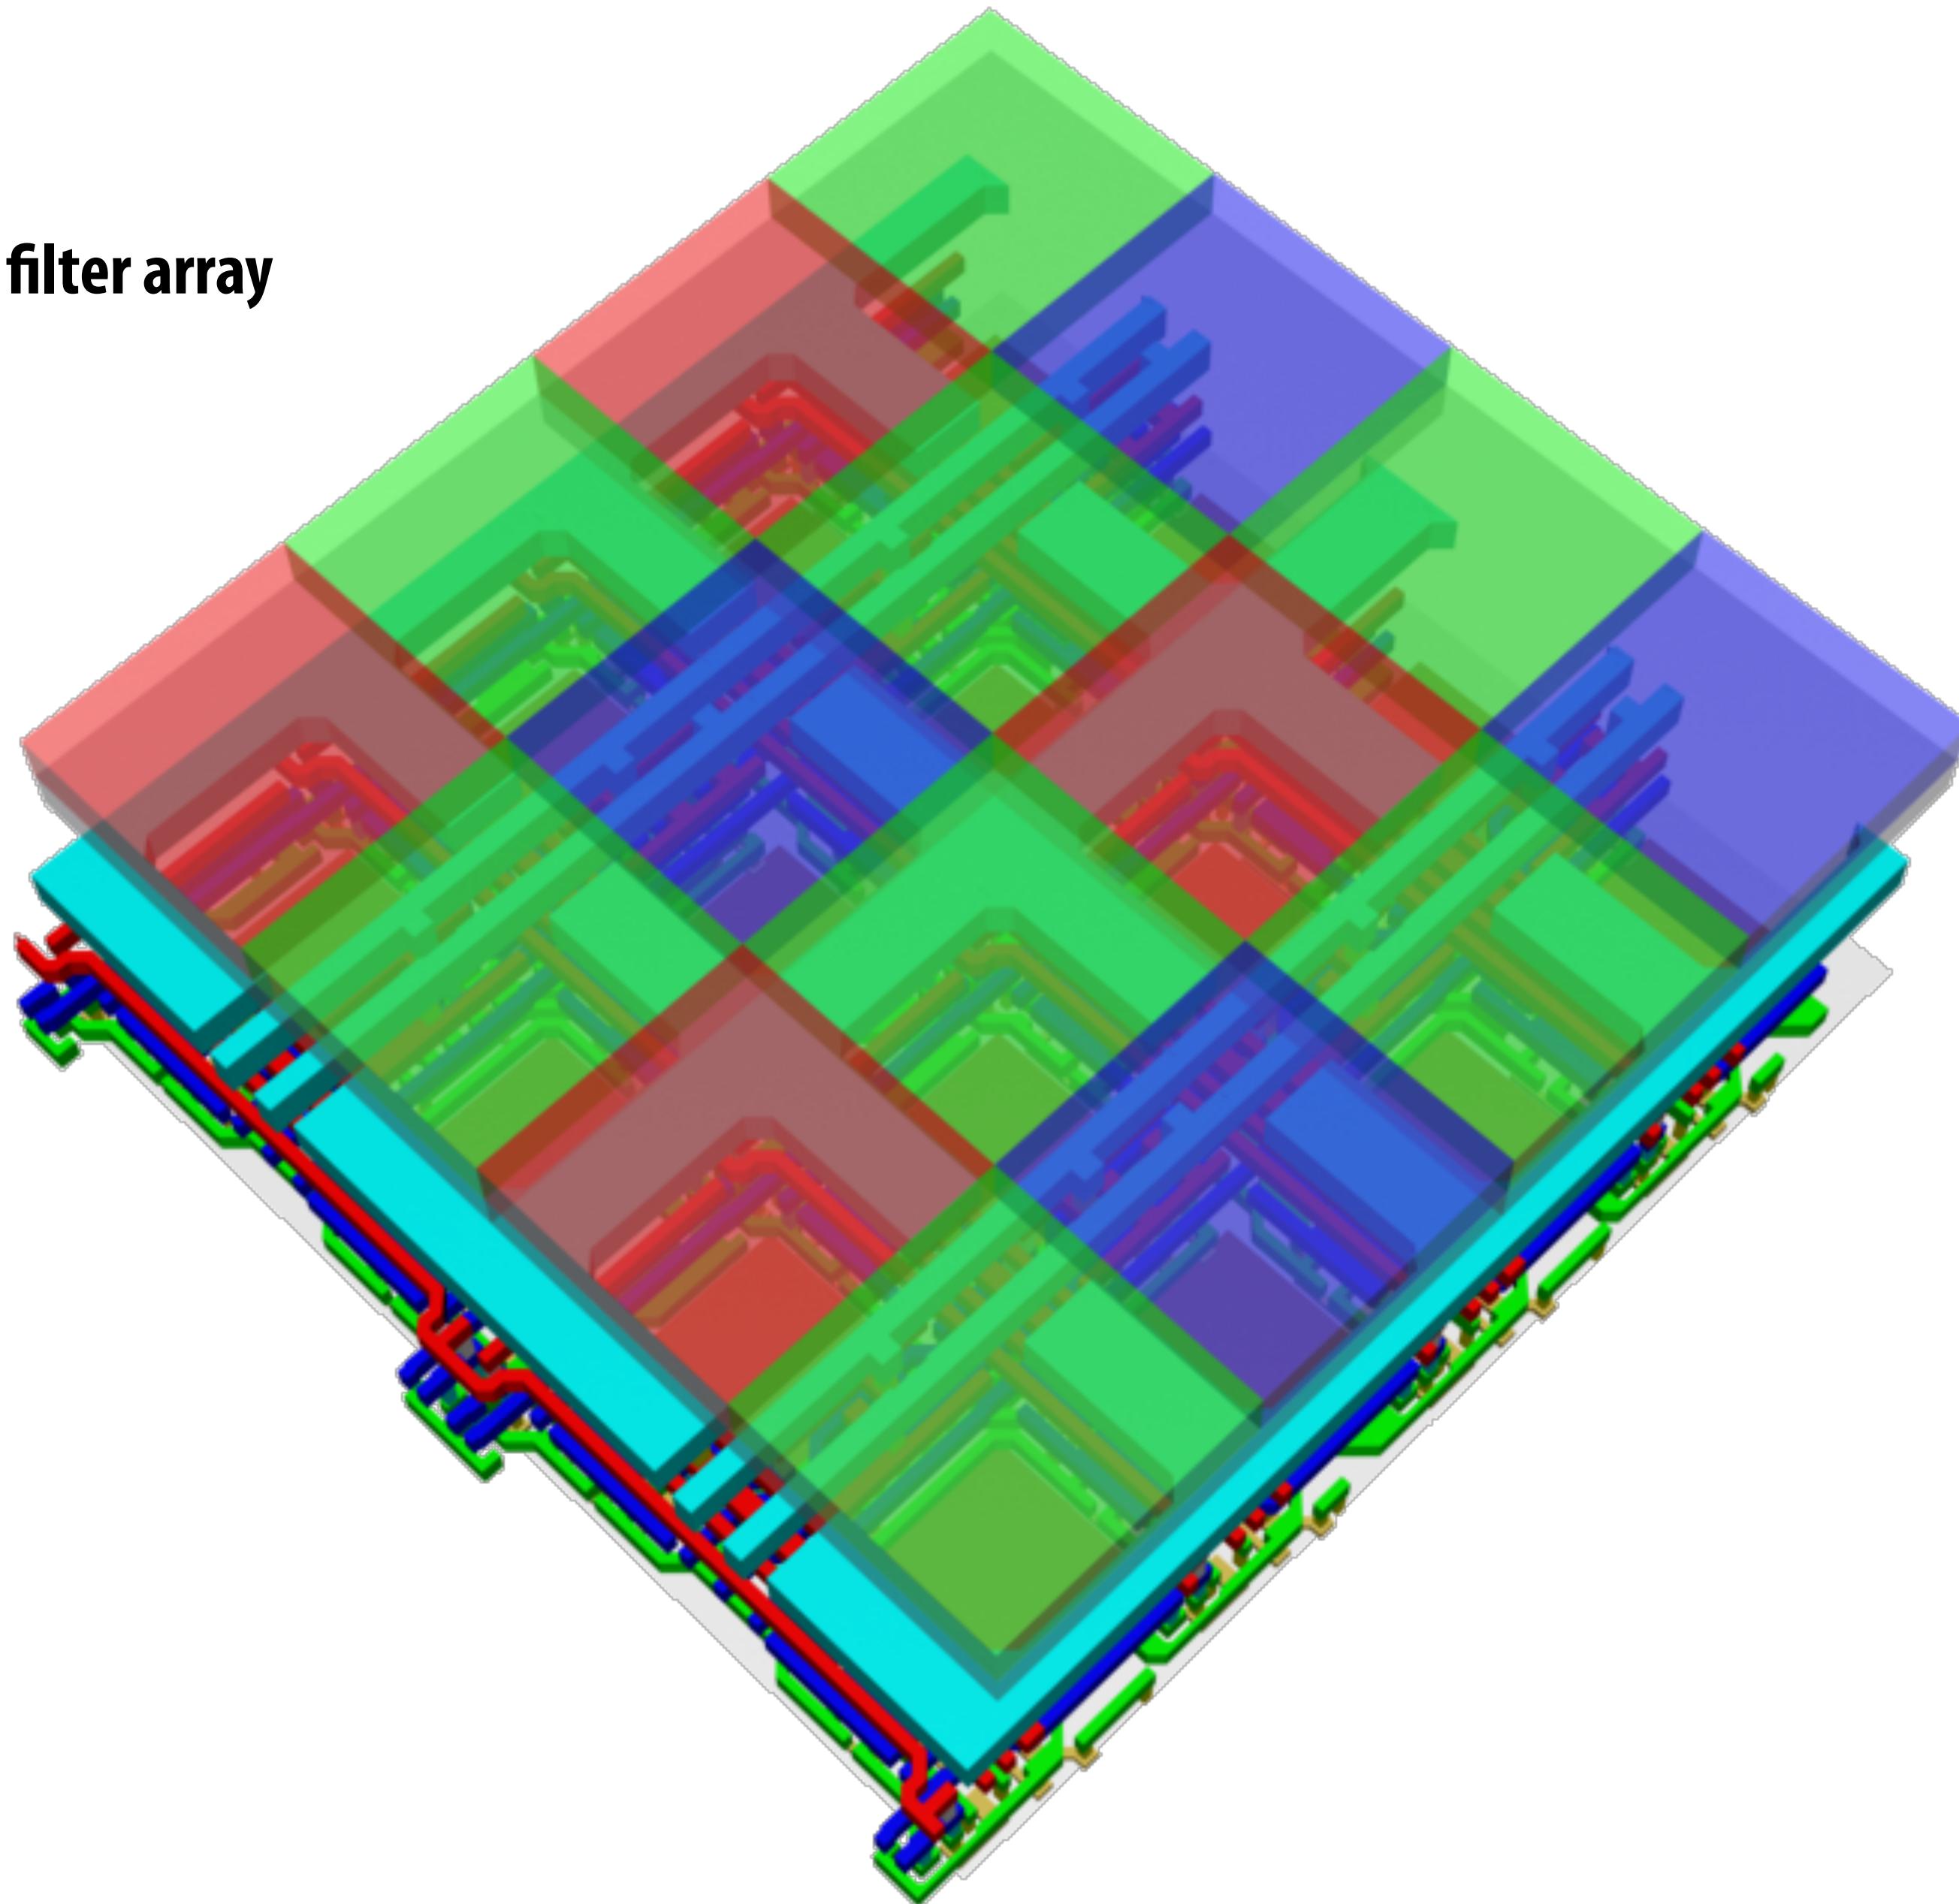
Metal 3



Metal 4

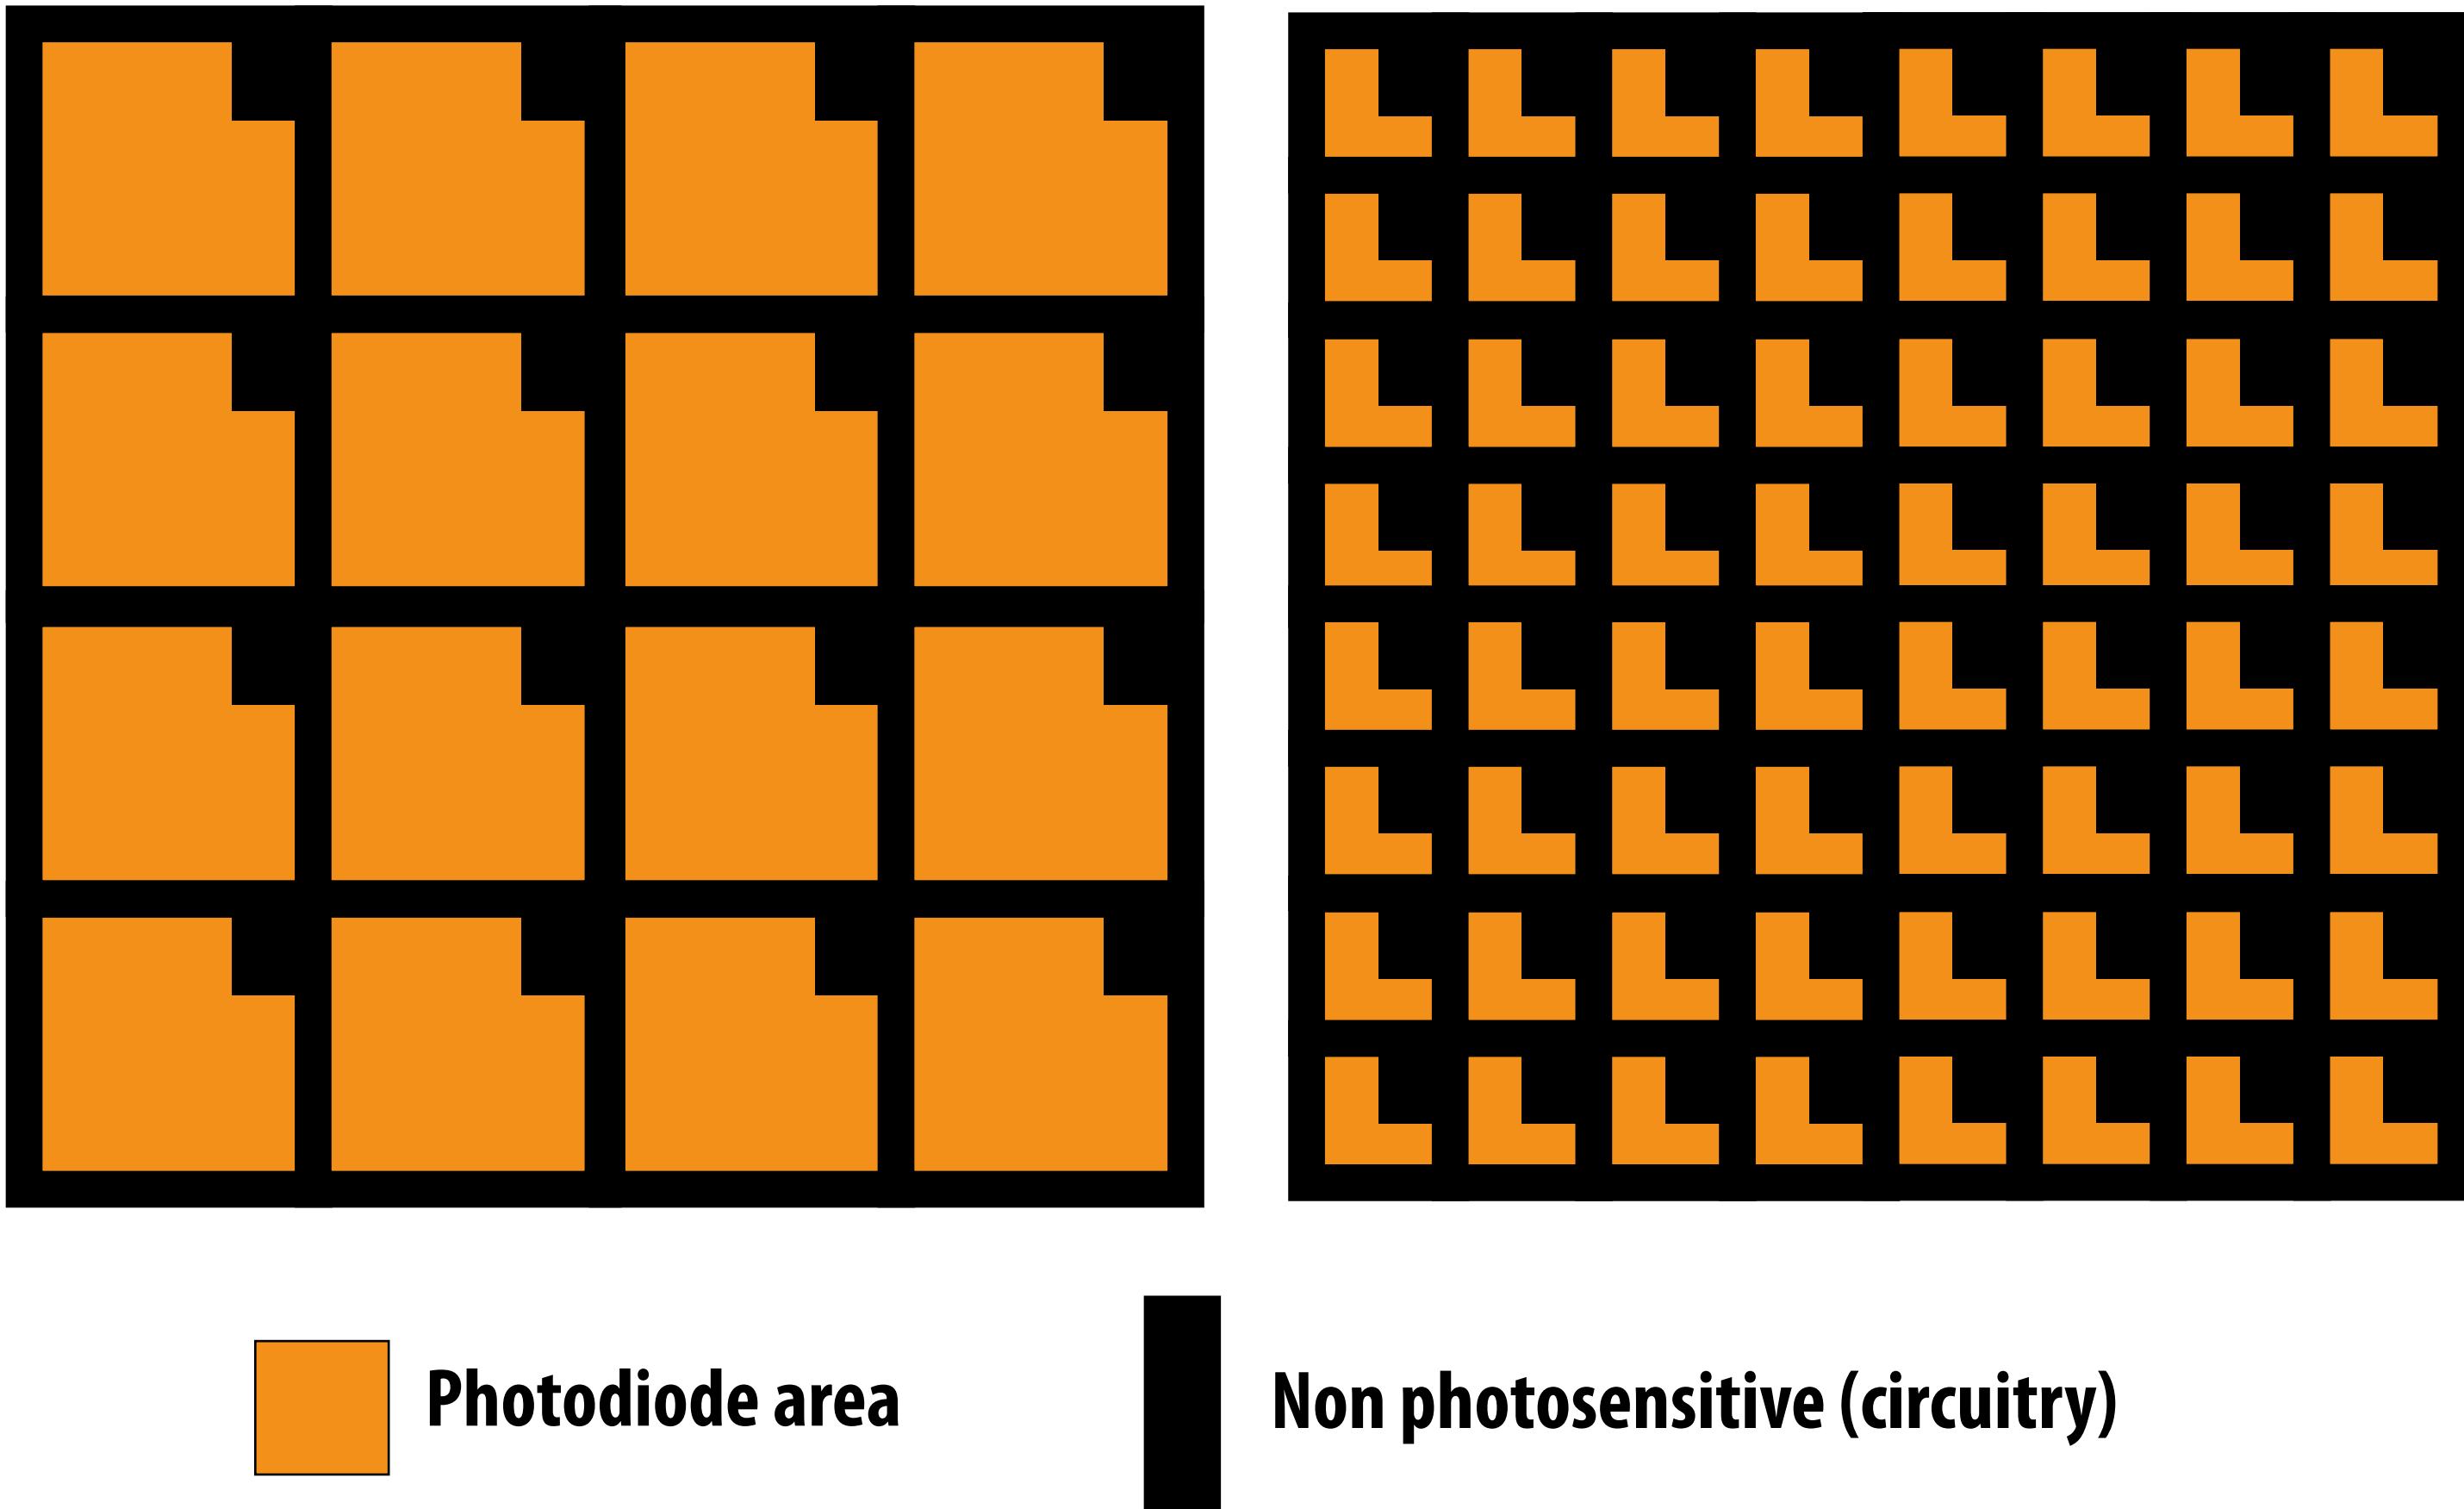


Color filter array

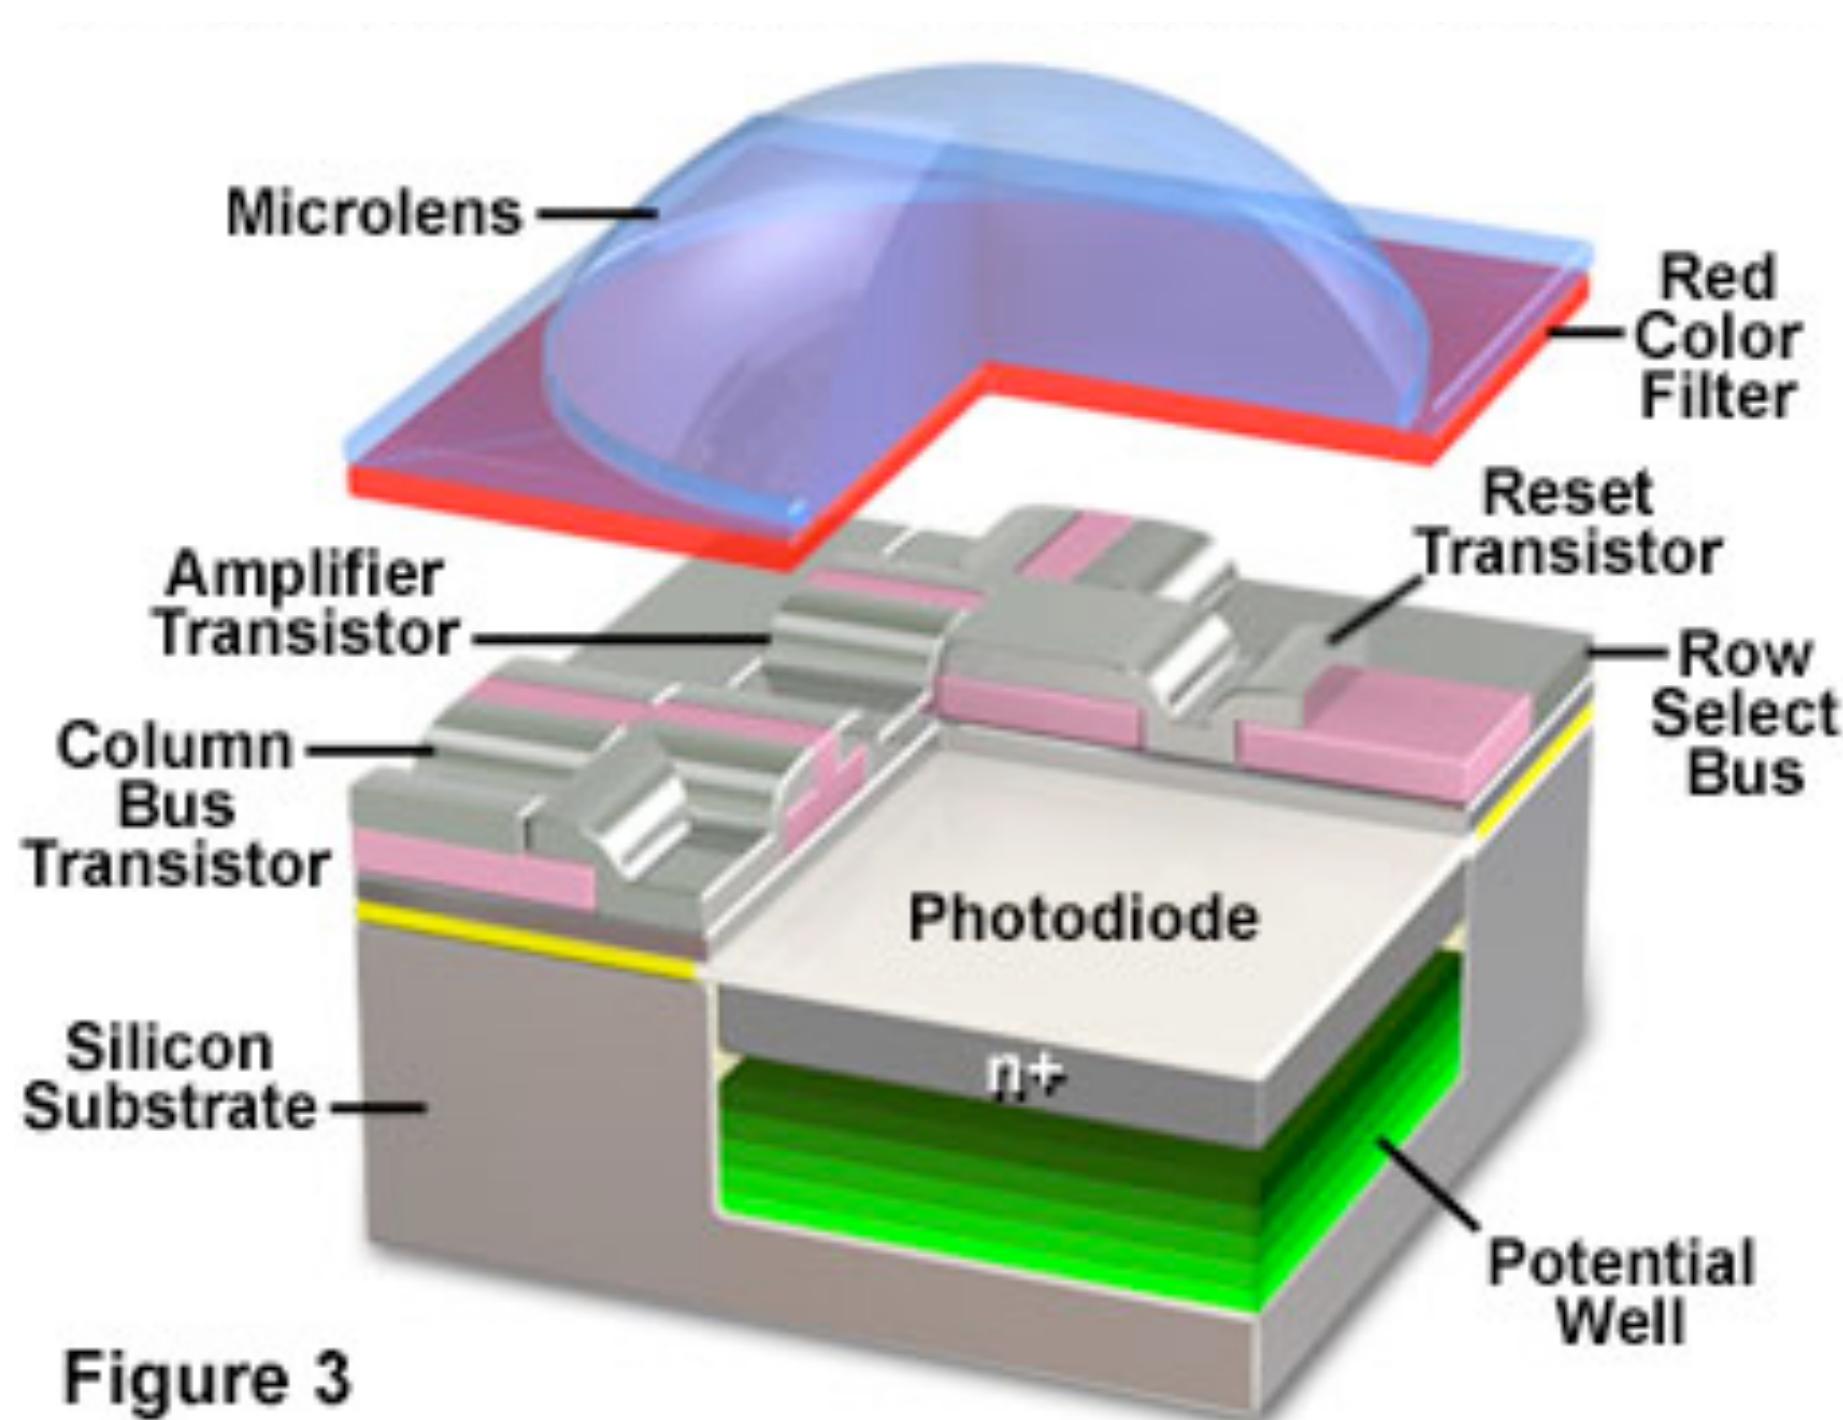


Pixel fill factor

Fraction of pixel area that integrates incoming light



CMOS sensor pixel

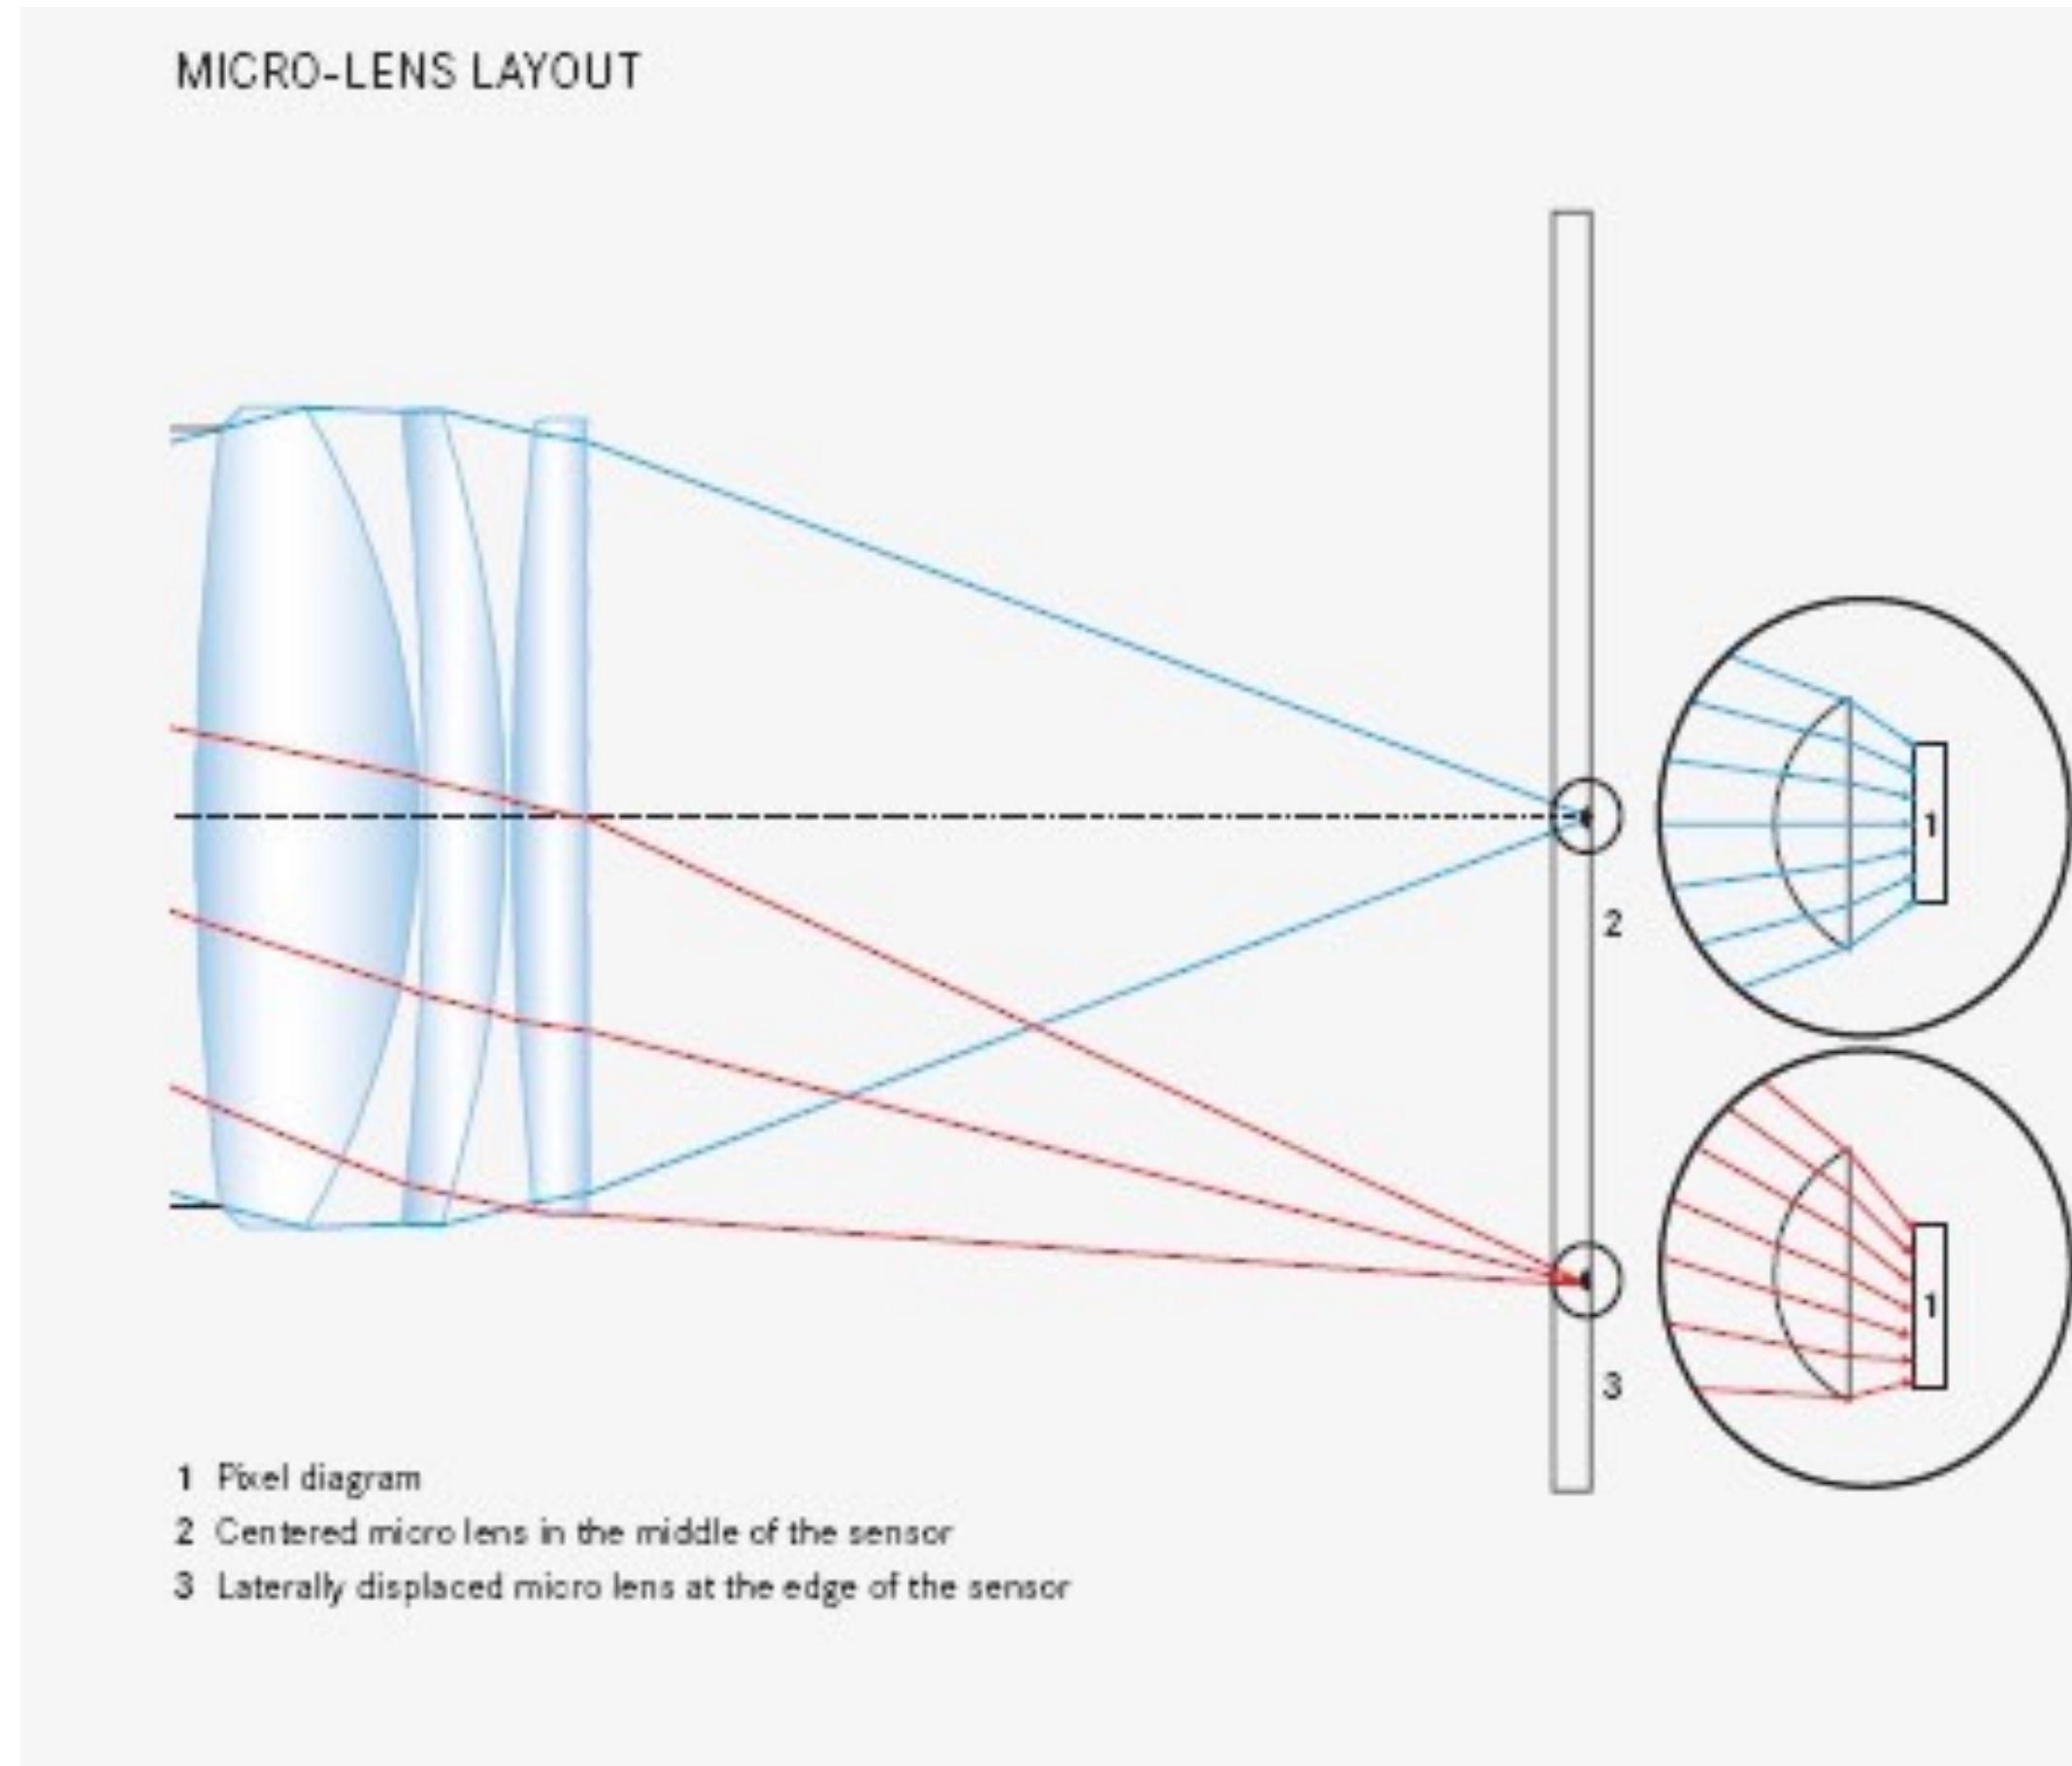


Color filter attenuates light

Microlens (a.k.a. lenslet) steers light toward photo-sensitive region (increases light-gathering capability)

Microlens also serves to prefilter signal. Why?

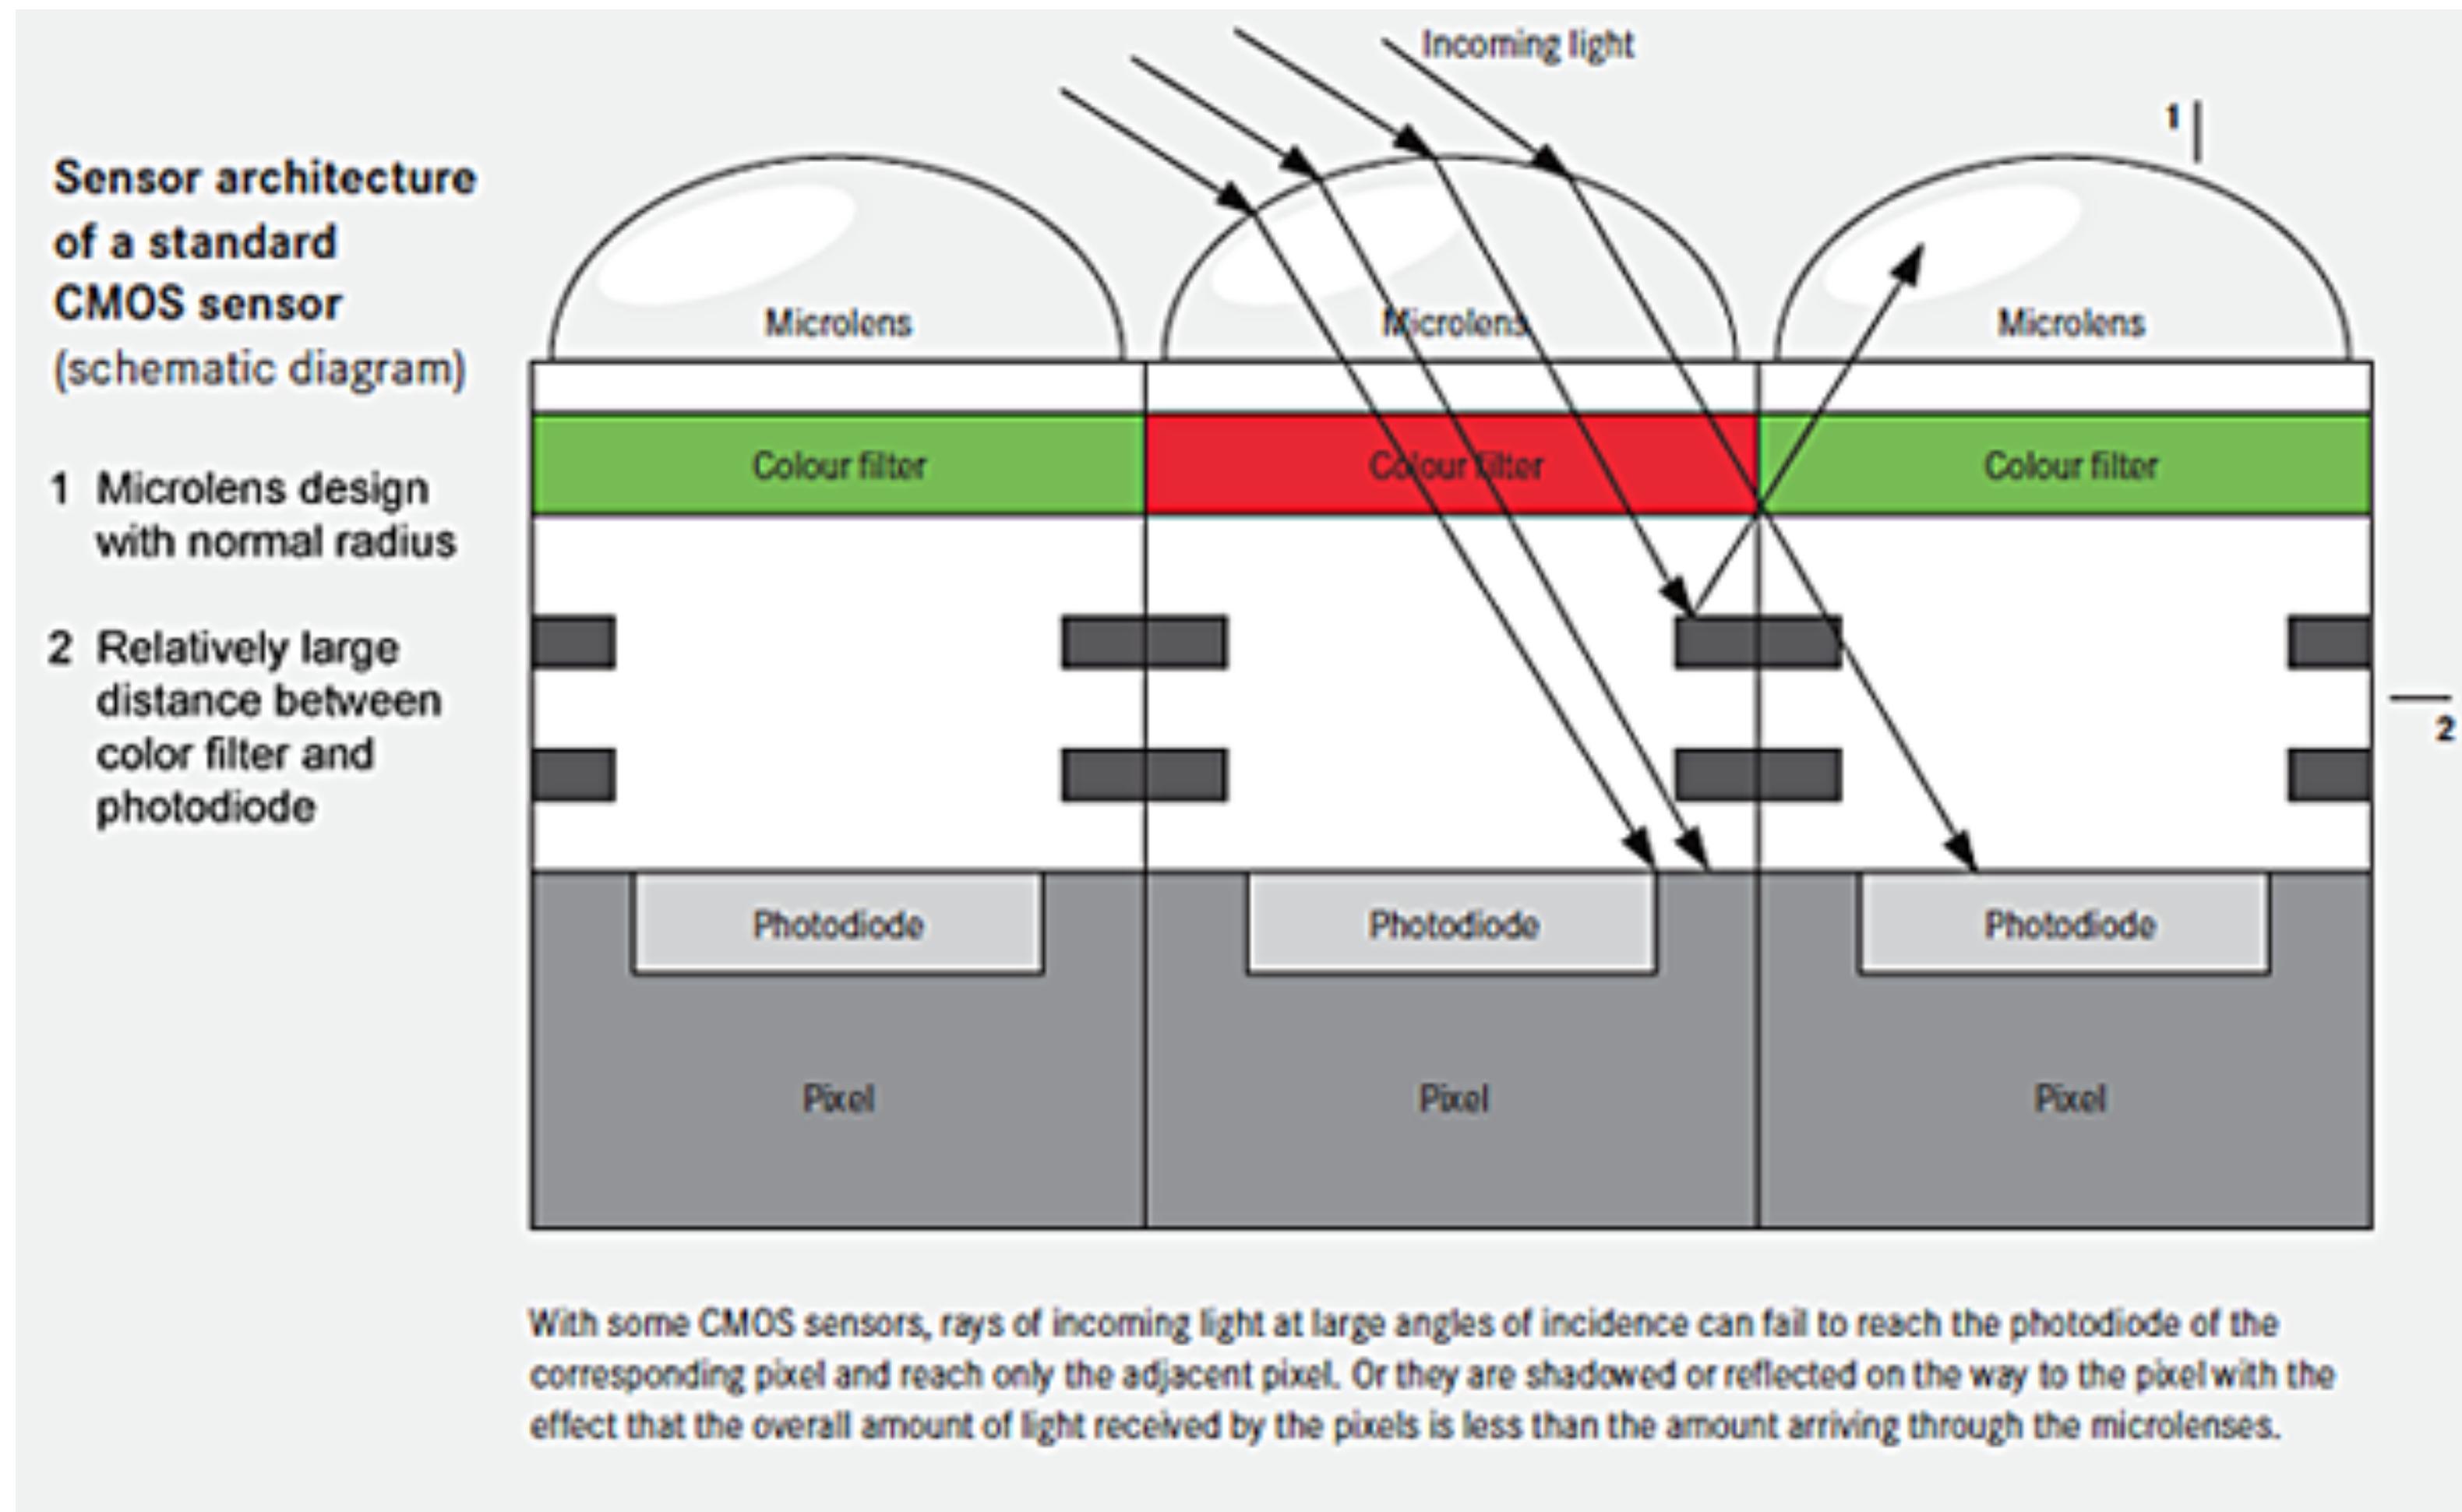
Using micro lenses to improve fill factor



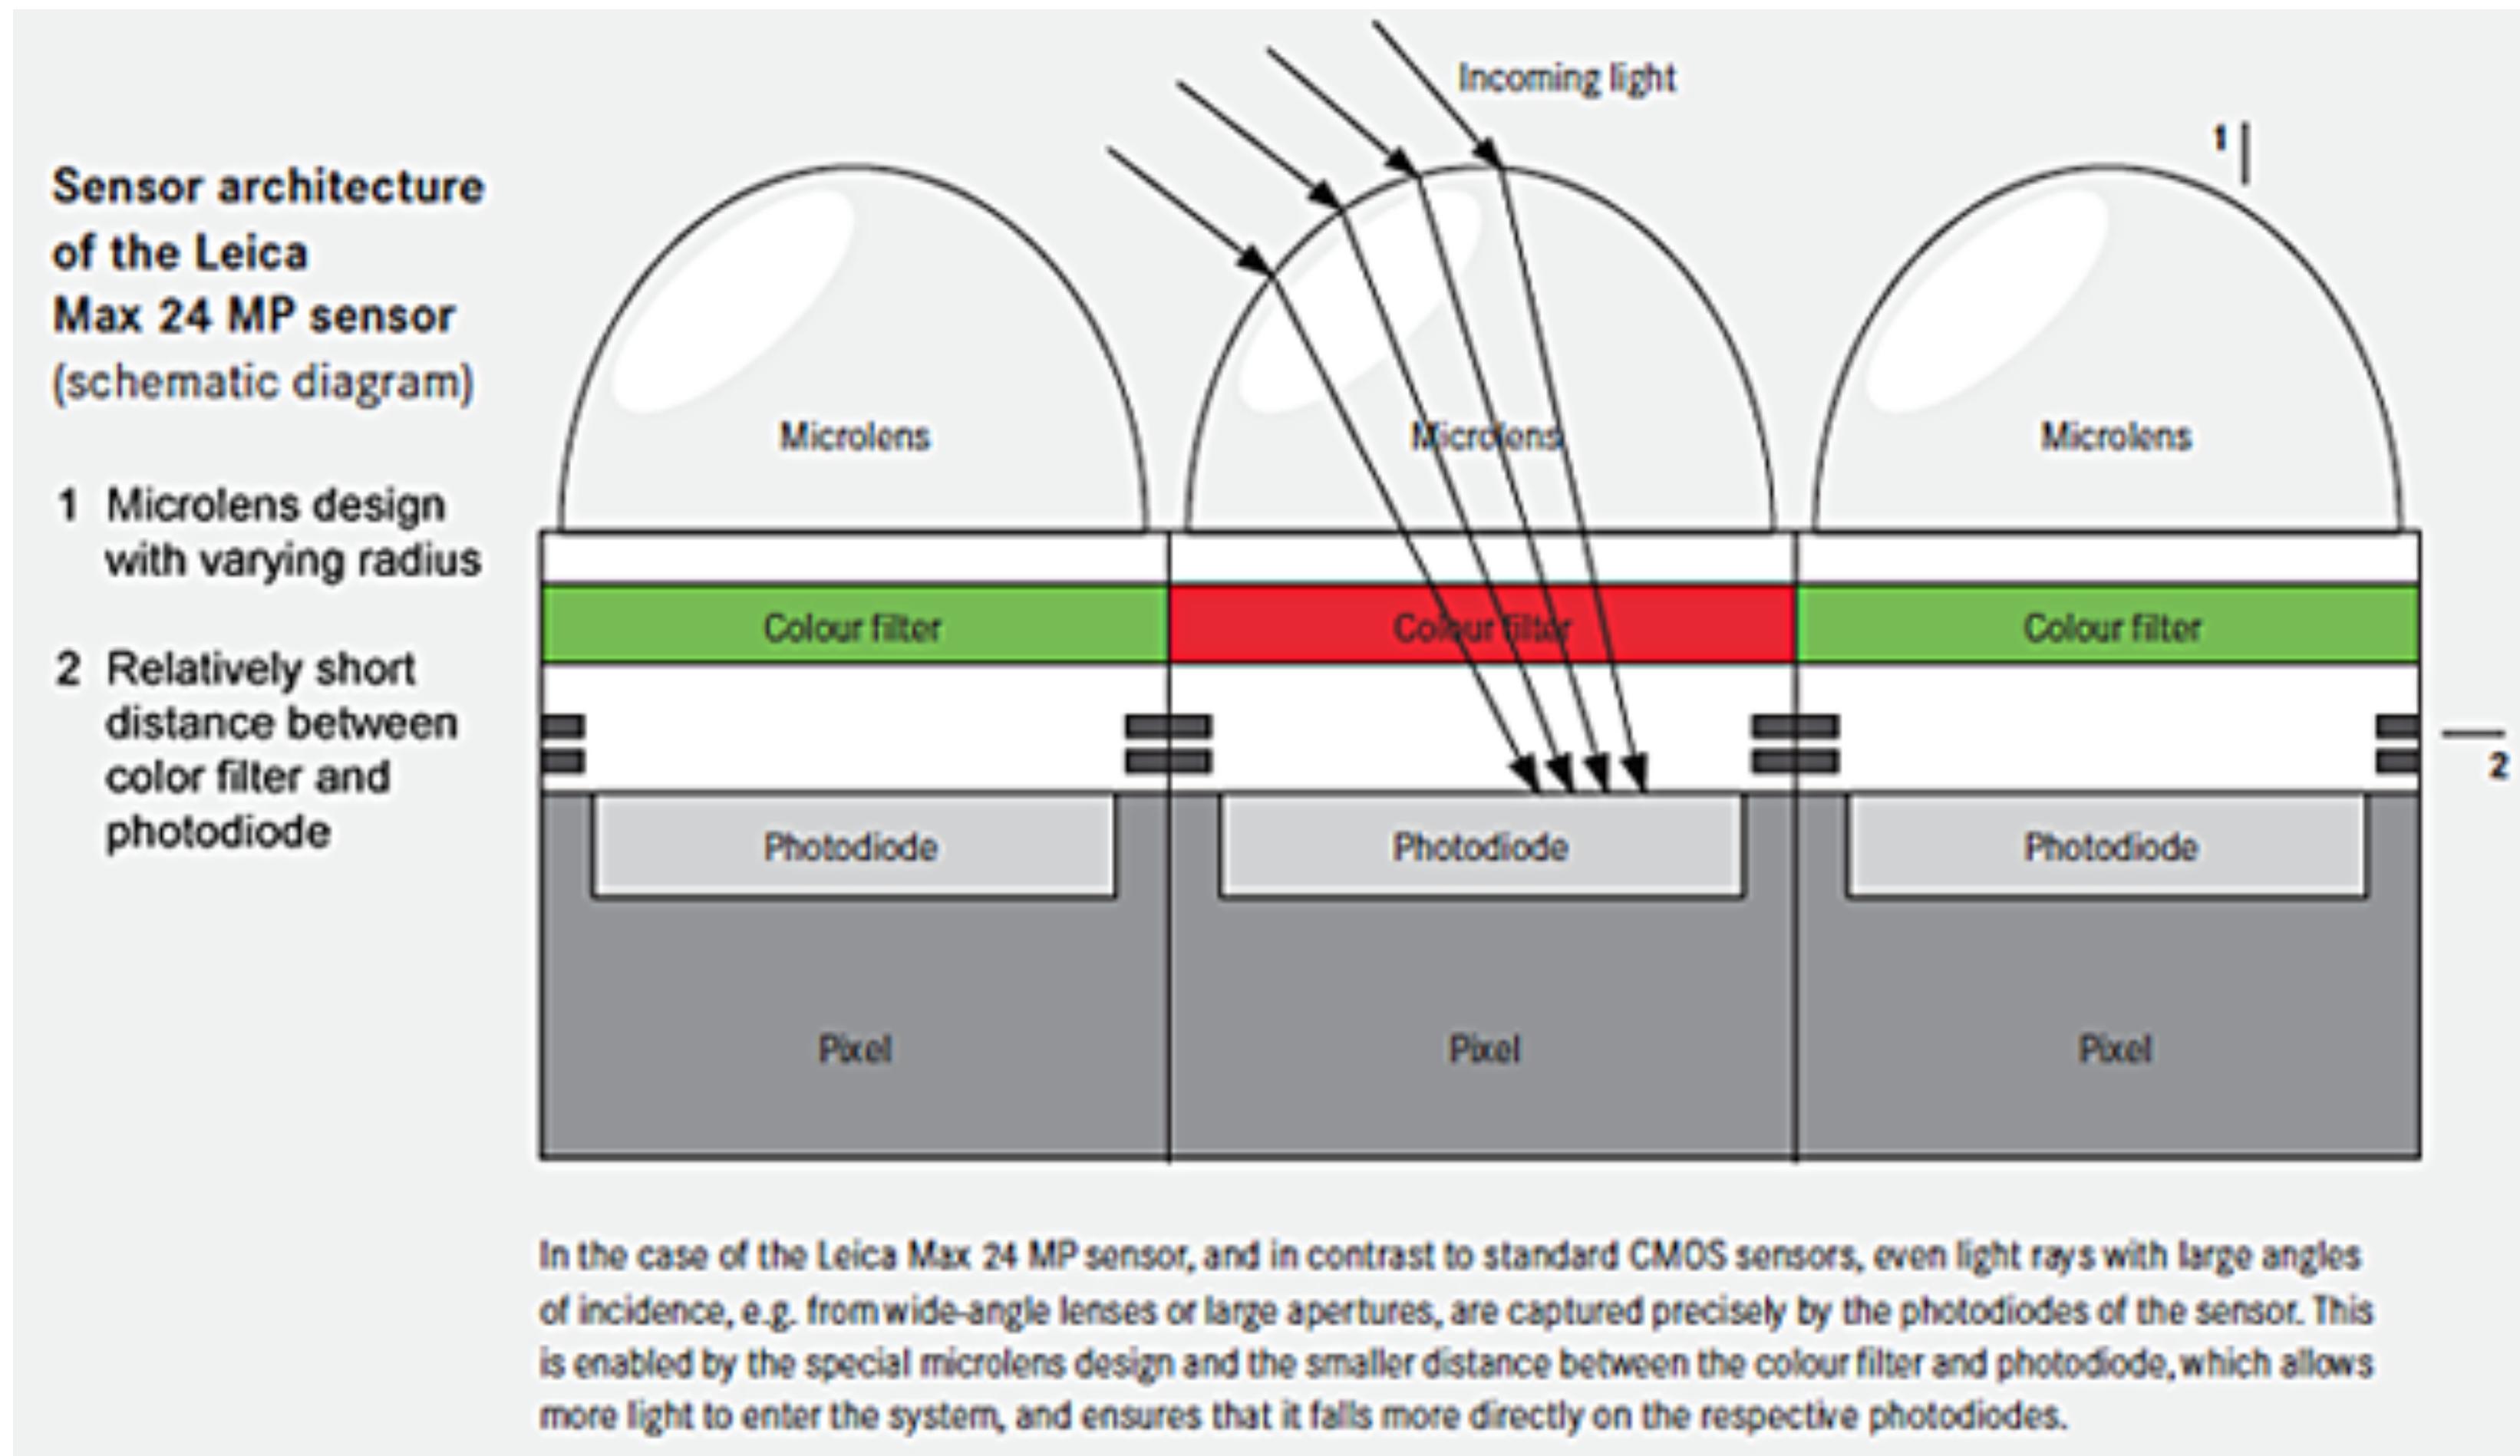
Leica M9

Shifted microlenses on M9 sensor.

Optical cross-talk



Pixel optics for minimizing cross-talk



In the case of the Leica Max 24 MP sensor, and in contrast to standard CMOS sensors, even light rays with large angles of incidence, e.g. from wide-angle lenses or large apertures, are captured precisely by the photodiodes of the sensor. This is enabled by the special microlens design and the smaller distance between the colour filter and photodiode, which allows more light to enter the system, and ensures that it falls more directly on the respective photodiodes.

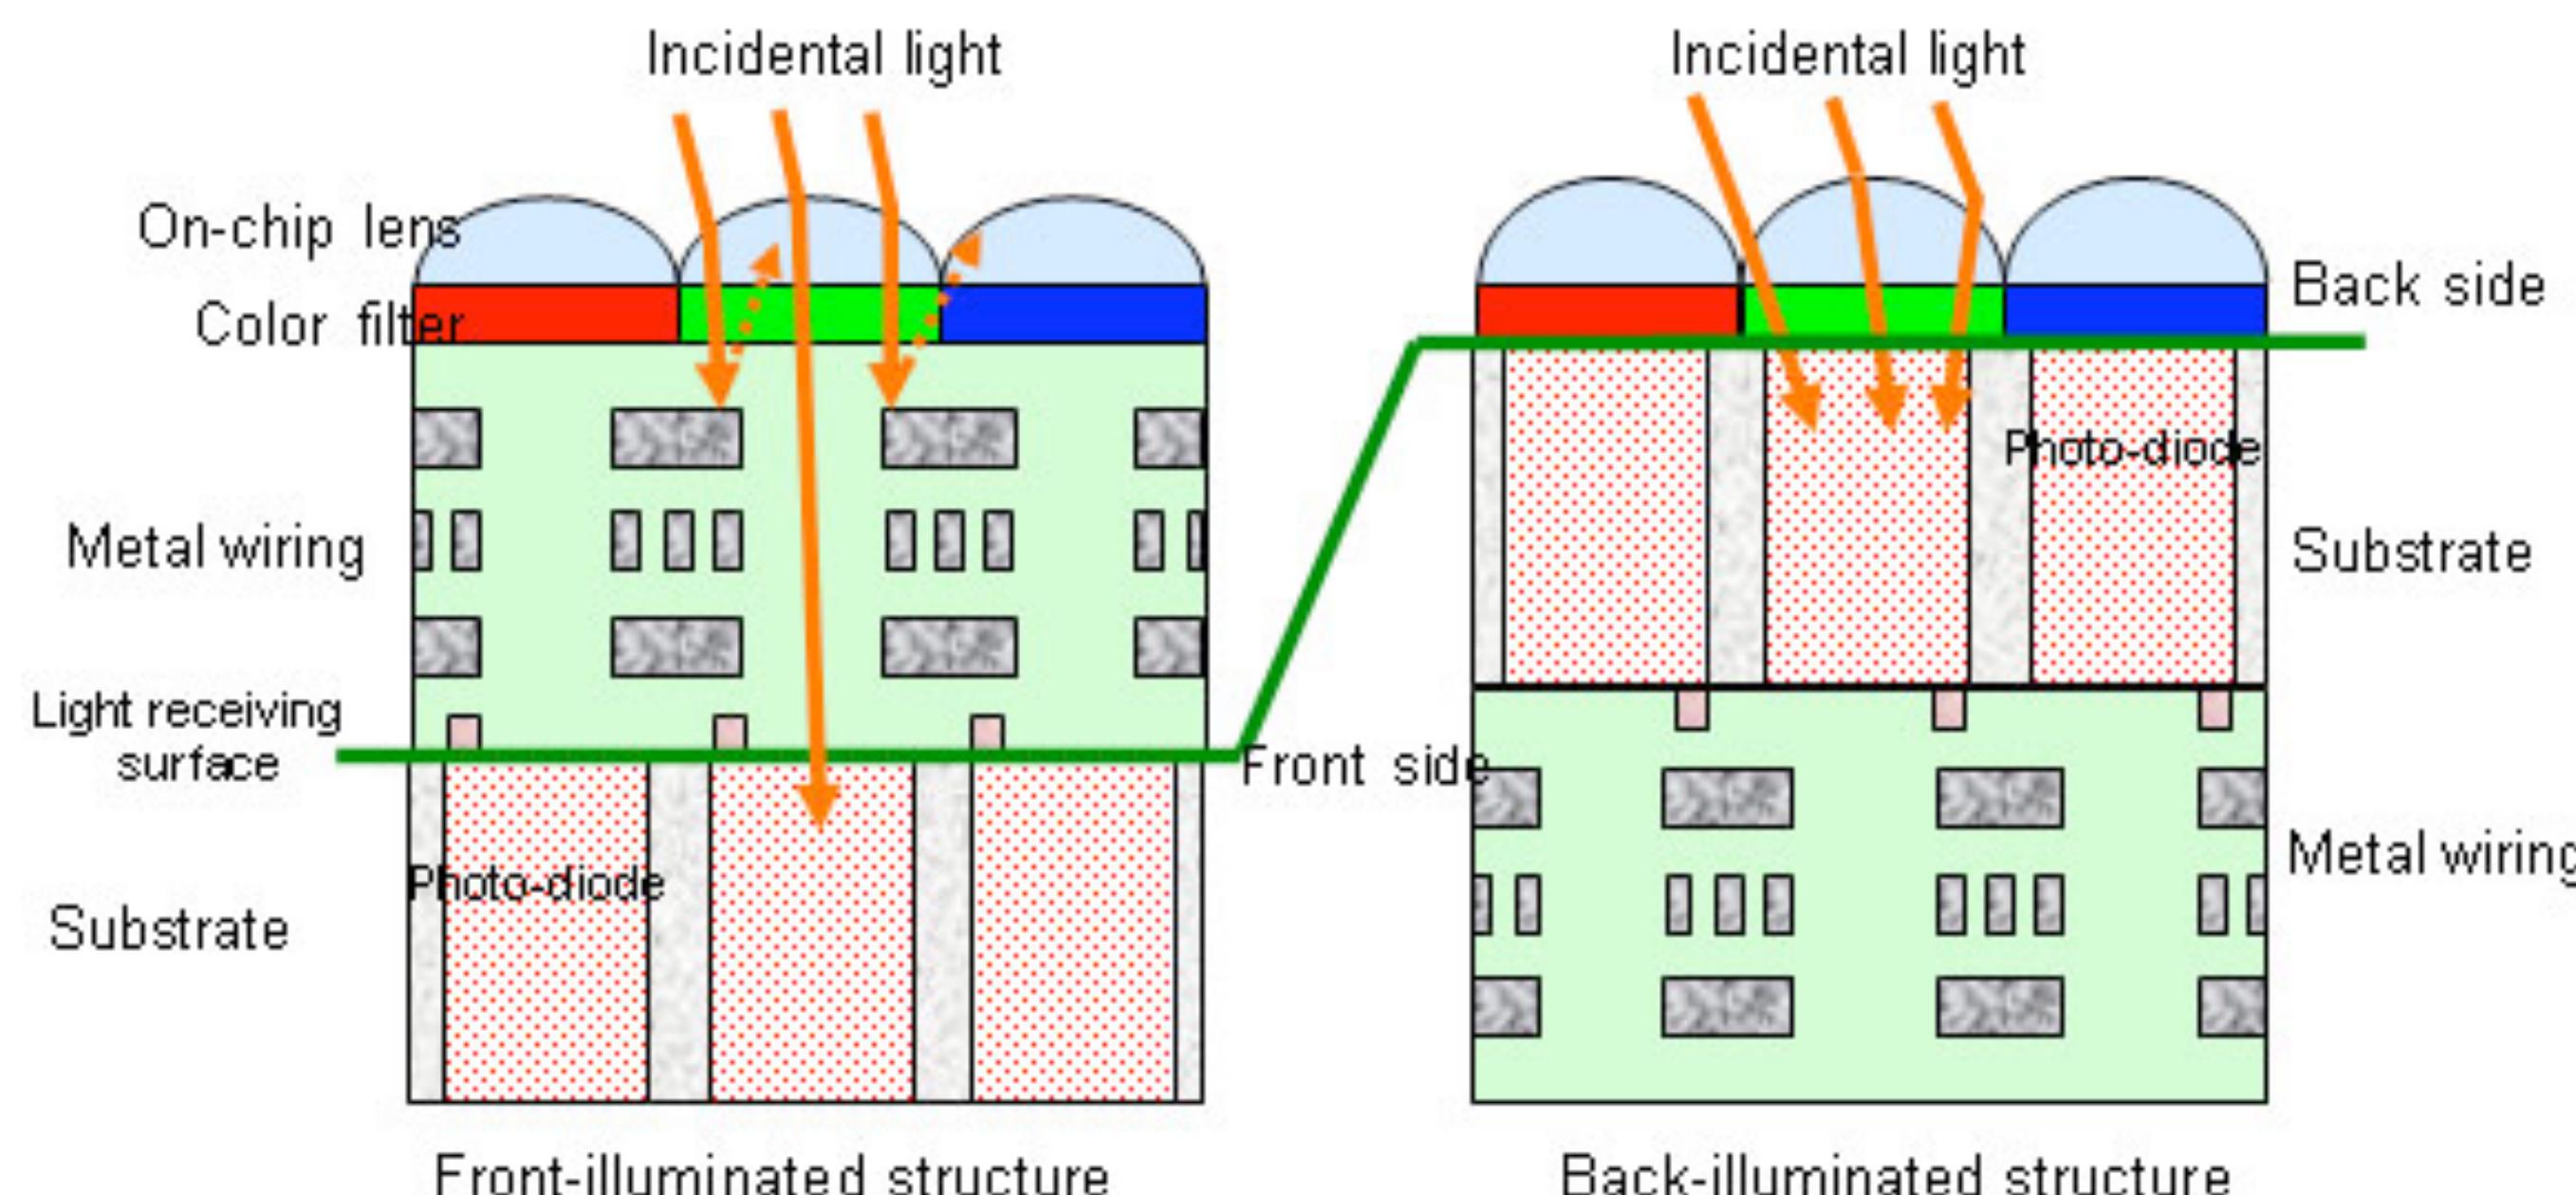
Slide credit: Ren Ng

<http://gmpphoto.blogspot.com/2012/09/the-new-leica-max-24mp-cmos-sensor.html>

Stanford CS348V, Winter 2018

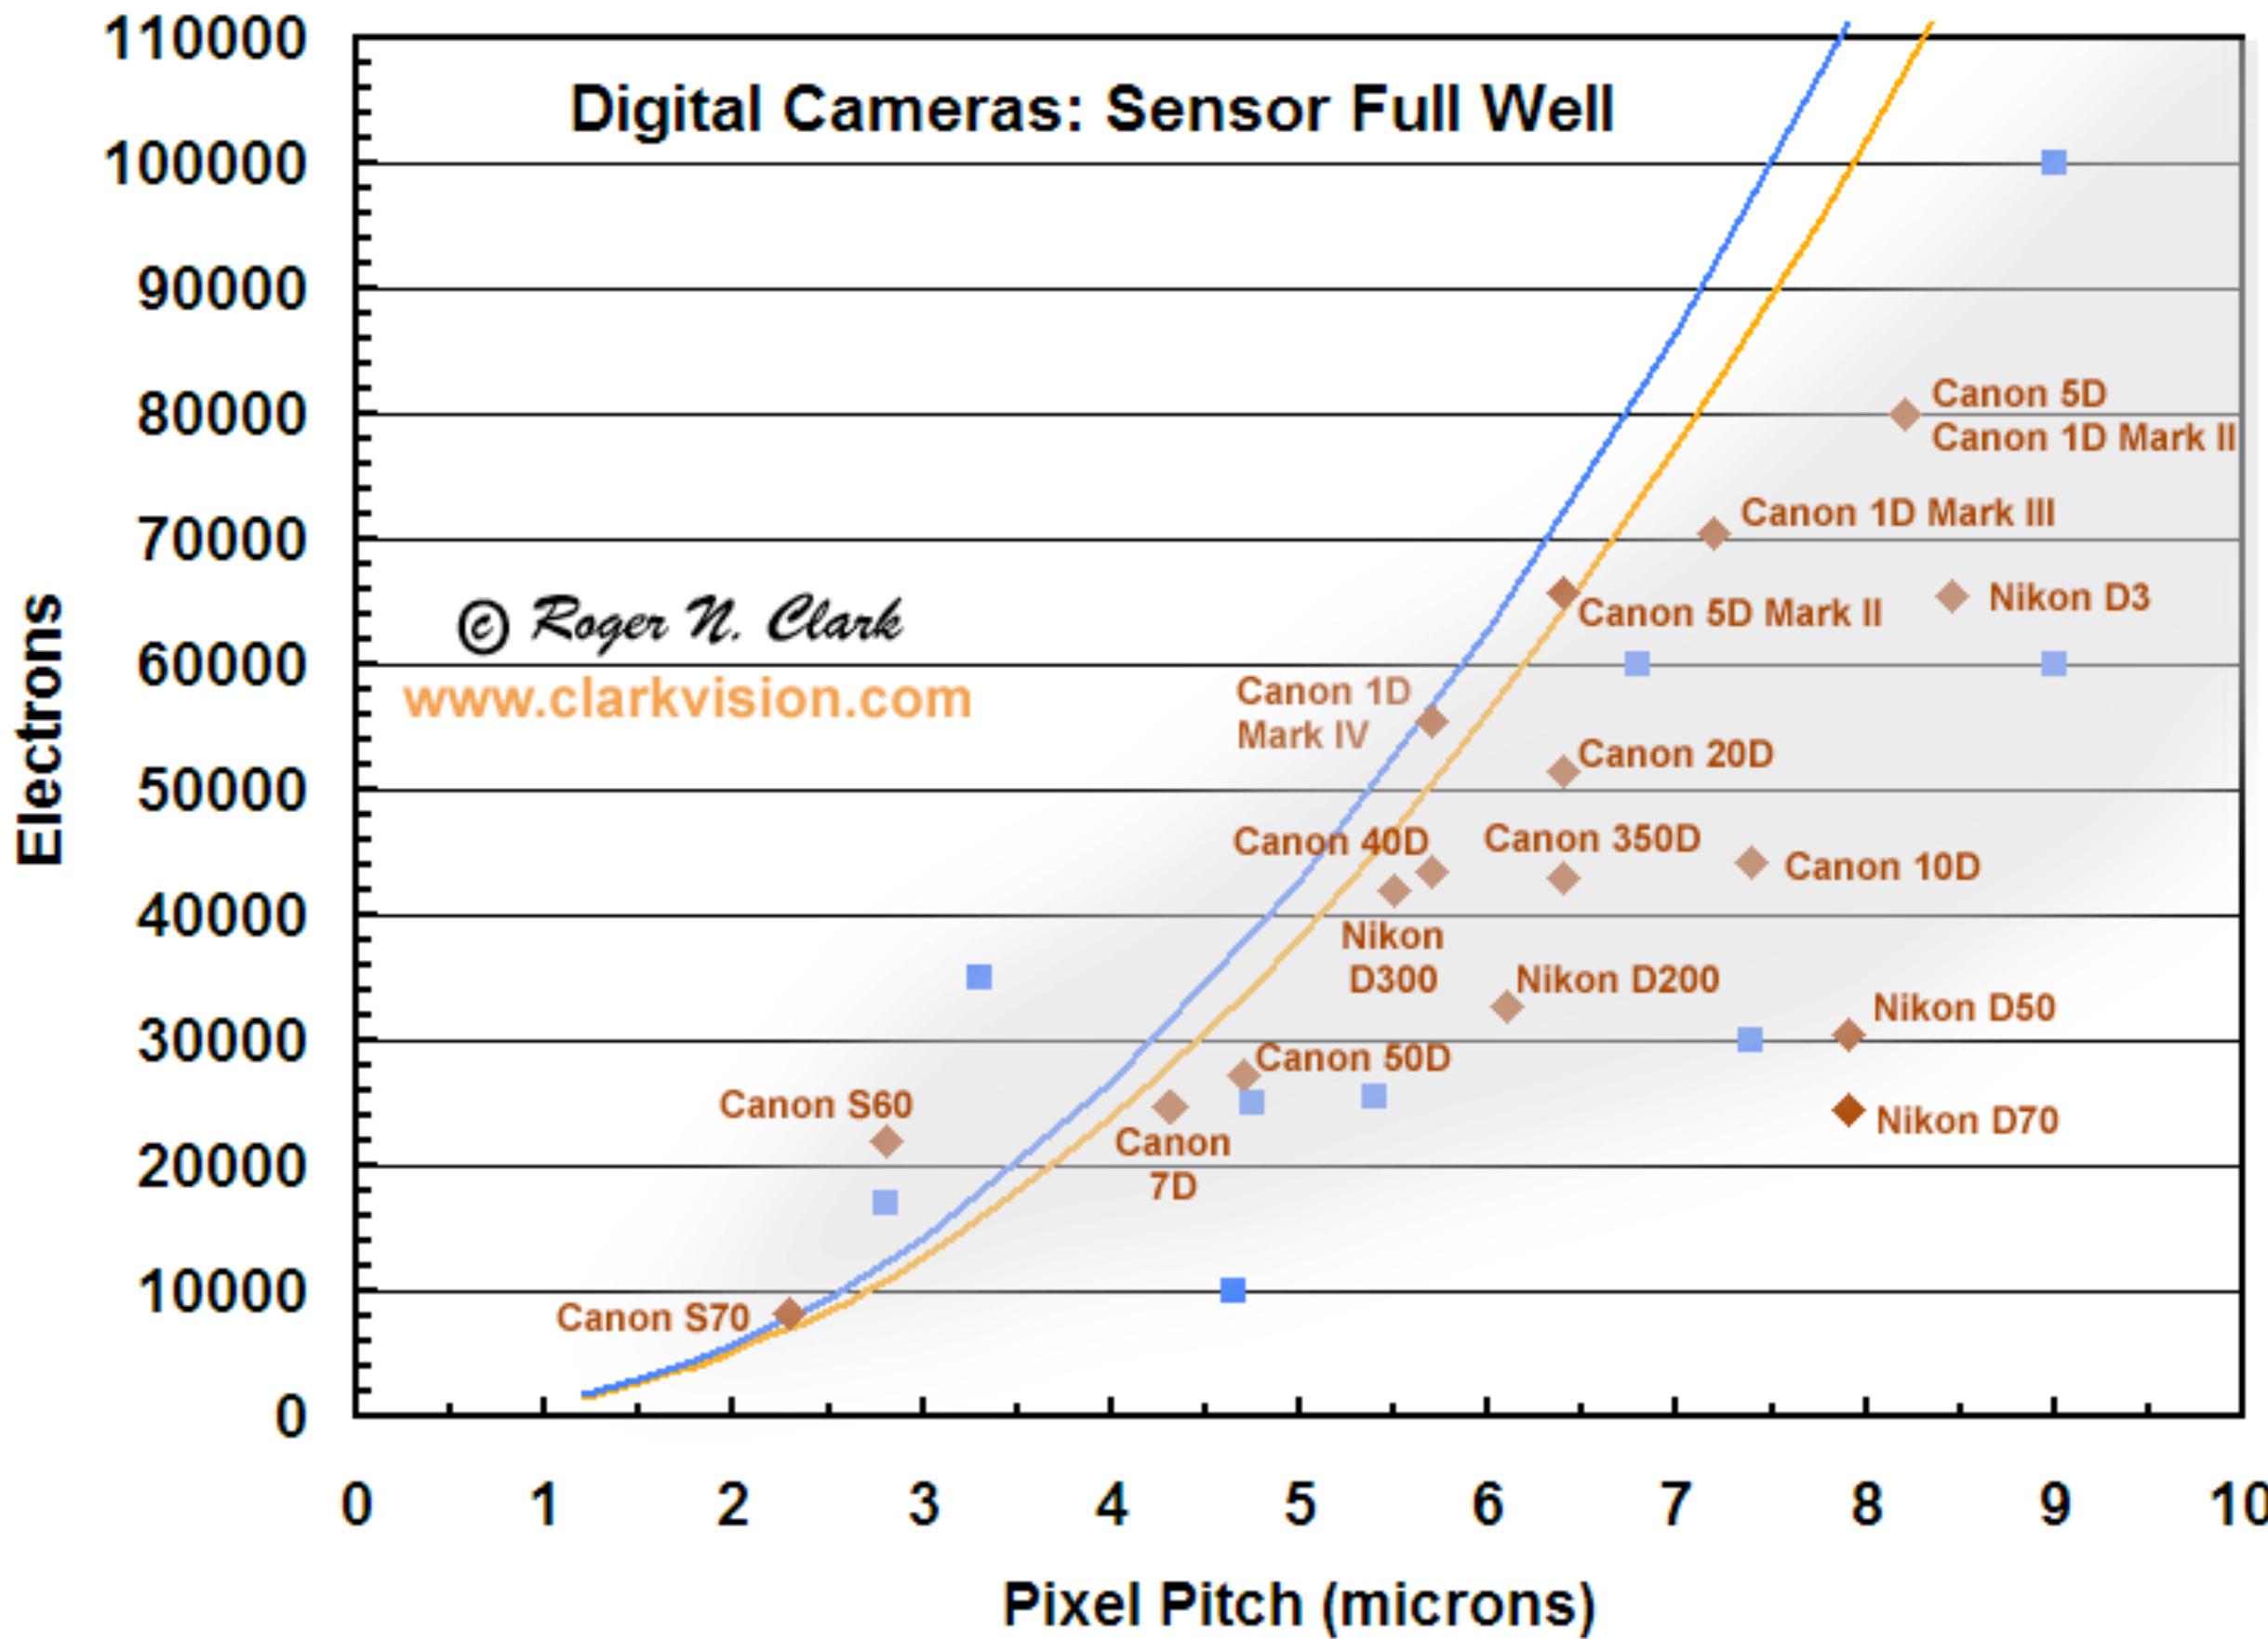
Backside illumination sensor

- Traditional CMOS: electronics block light
- Idea: move electronics underneath light gathering region
 - Increases fill factor
 - Reduces cross-talk due since photodiode closer to microns
 - Implication 1: better light sensitivity at fixed sensor size
 - Implication 2: equal light sensitivity at smaller sensor size (shrink sensor)



Full-well capacity

Pixel saturates when photon capacity is exceeded

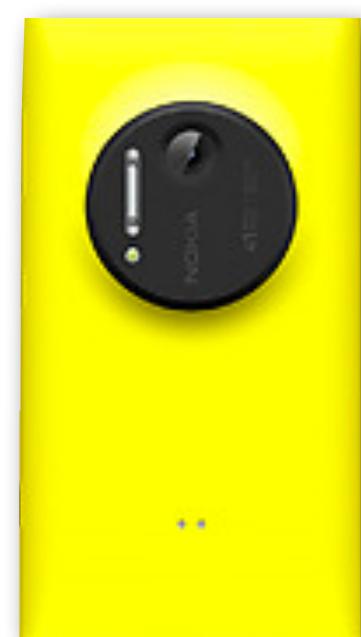
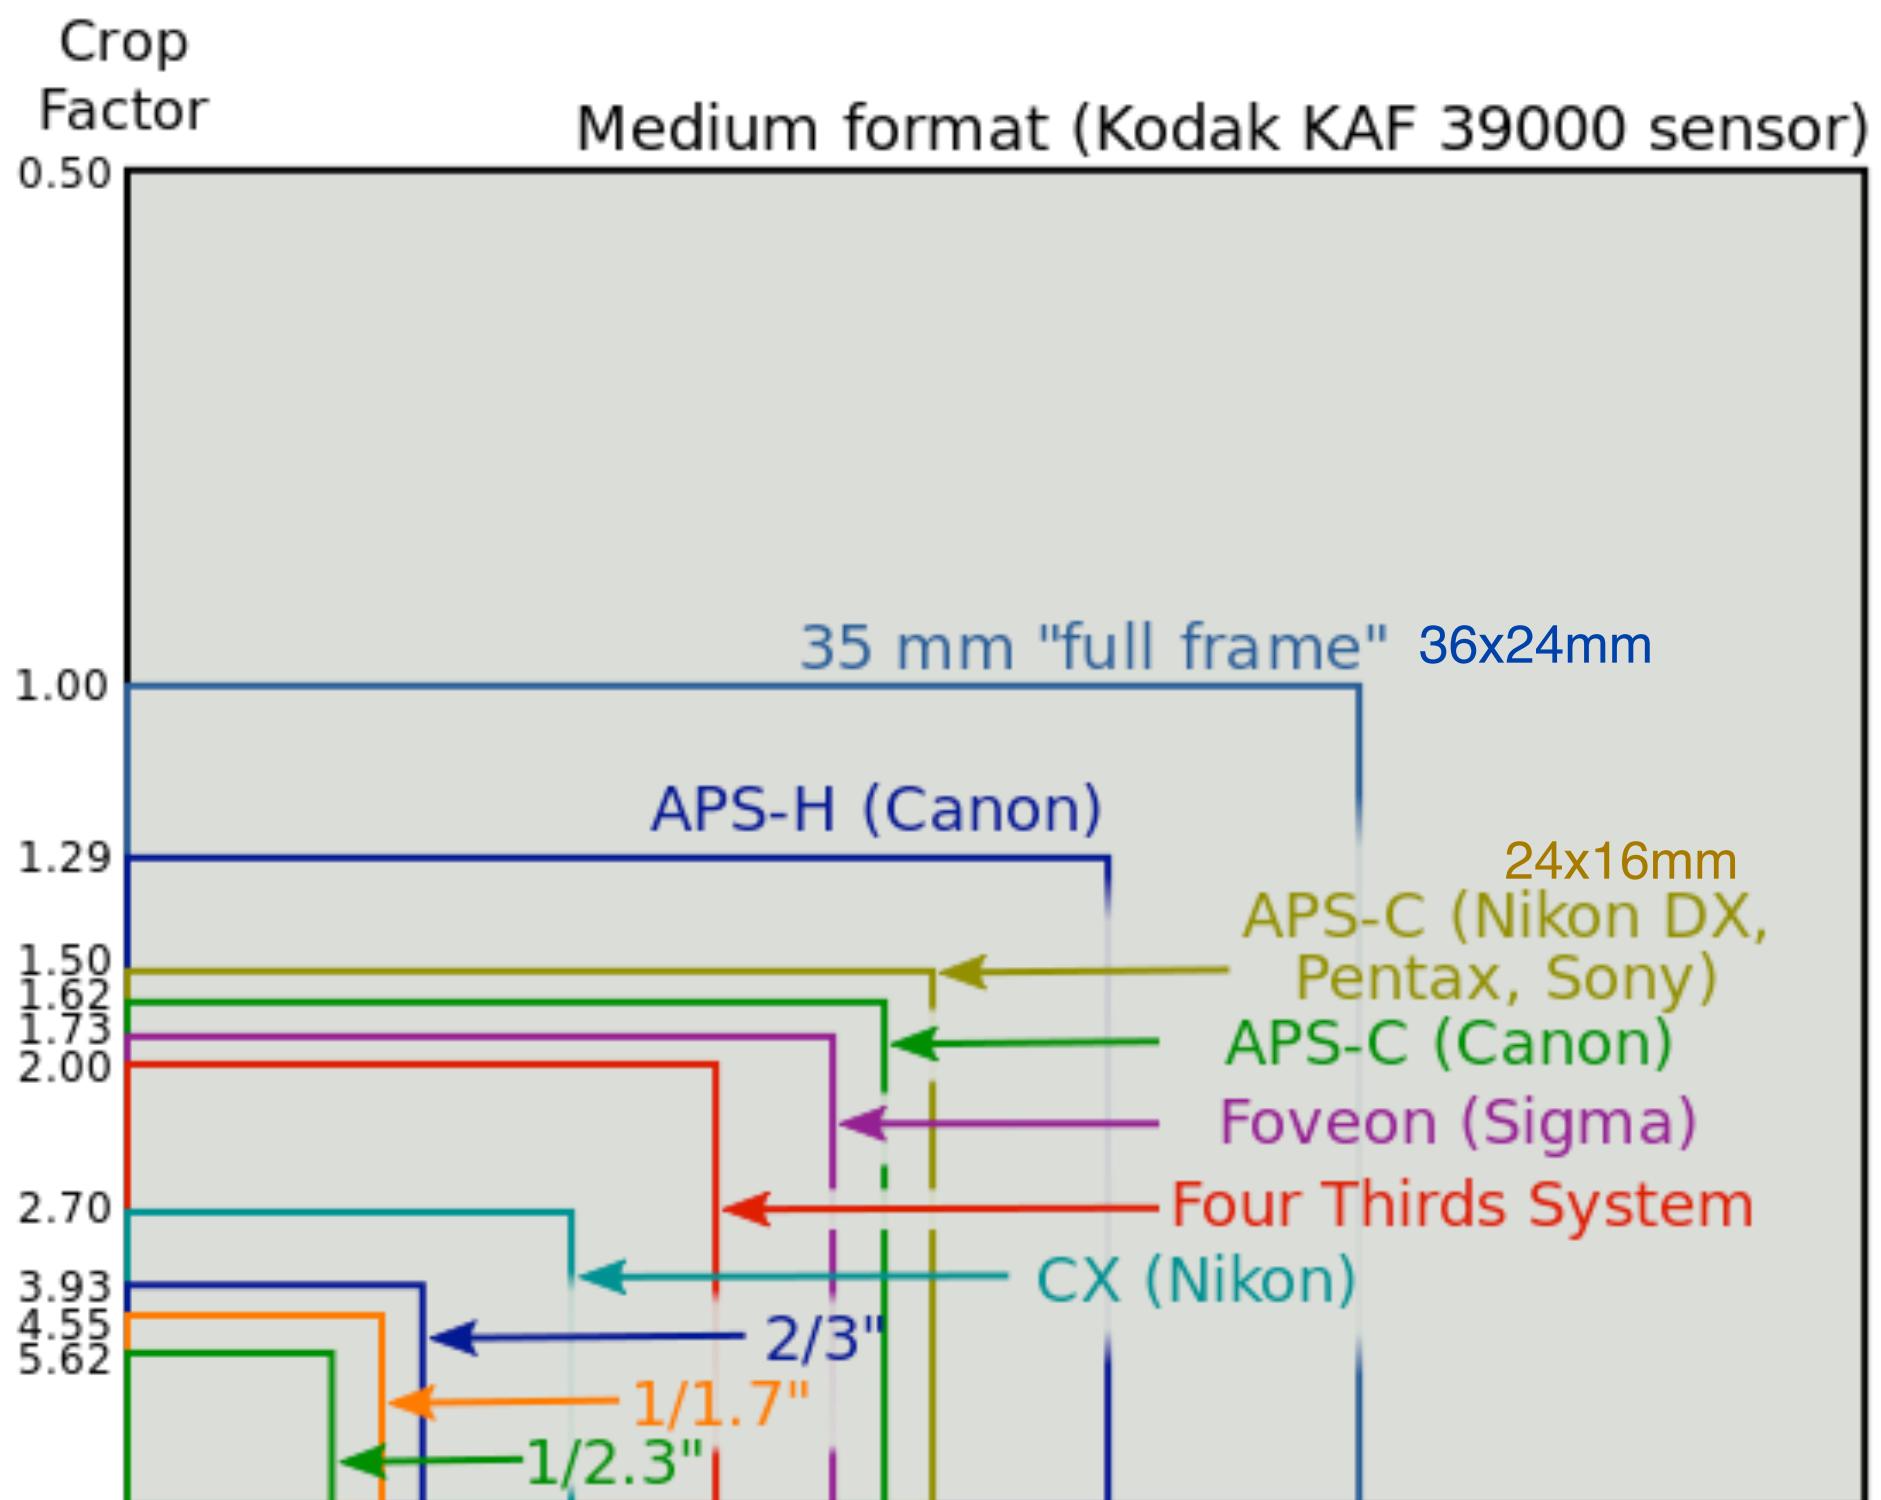


Oversaturated pixels



Bigger sensors = bigger pixels (or more pixels?)

- iPhone 6s (1.2 micron pixels, 12 MP)
- My Nikon D7000 (APS-C)
(4.8 micron pixels, 16 MP)
- Nikon D4 (full frame sensor)
(7.3 micron pixels, 16 MP)
- Implication: very high pixel count sensors
can be built with current CMOS technology
 - Full frame sensor with iPhone 6s pixel
size ~ 600 MP sensor



Nokia Lumia
(41 MP)

Steps in capturing an image

- 1. Clear sensor pixels**
- 2. Open camera's mechanical shutter (exposure begins)**
- 3. Optional: fire flash**
- 4. Close camera mechanical shutter (exposure ends)**
- 5. Read measurements off sensor**
 - For each row:**
 - Select row, read pixel for all columns in parallel**
 - Pass data stream through amplifier and ADC**

Electronic rolling shutter

Many cameras do not have a mechanical shutter
(e.g., smart-phone cameras)

- 1. Clear sensor pixels for row i (exposure begins)
- 2. Clear sensor pixels for row $i+1$ (exposure begins)
- ...
- 3. Read row i (exposure ends)
- 4. Read row $i+1$ (exposure ends)

Photo of red square, moving to right
over duration sensor is exposed



Each image row exposed for the same amount of time (same exposure)

Each image row exposed over different interval of time
(time offset determined by row read speed)

Rolling shutter effects

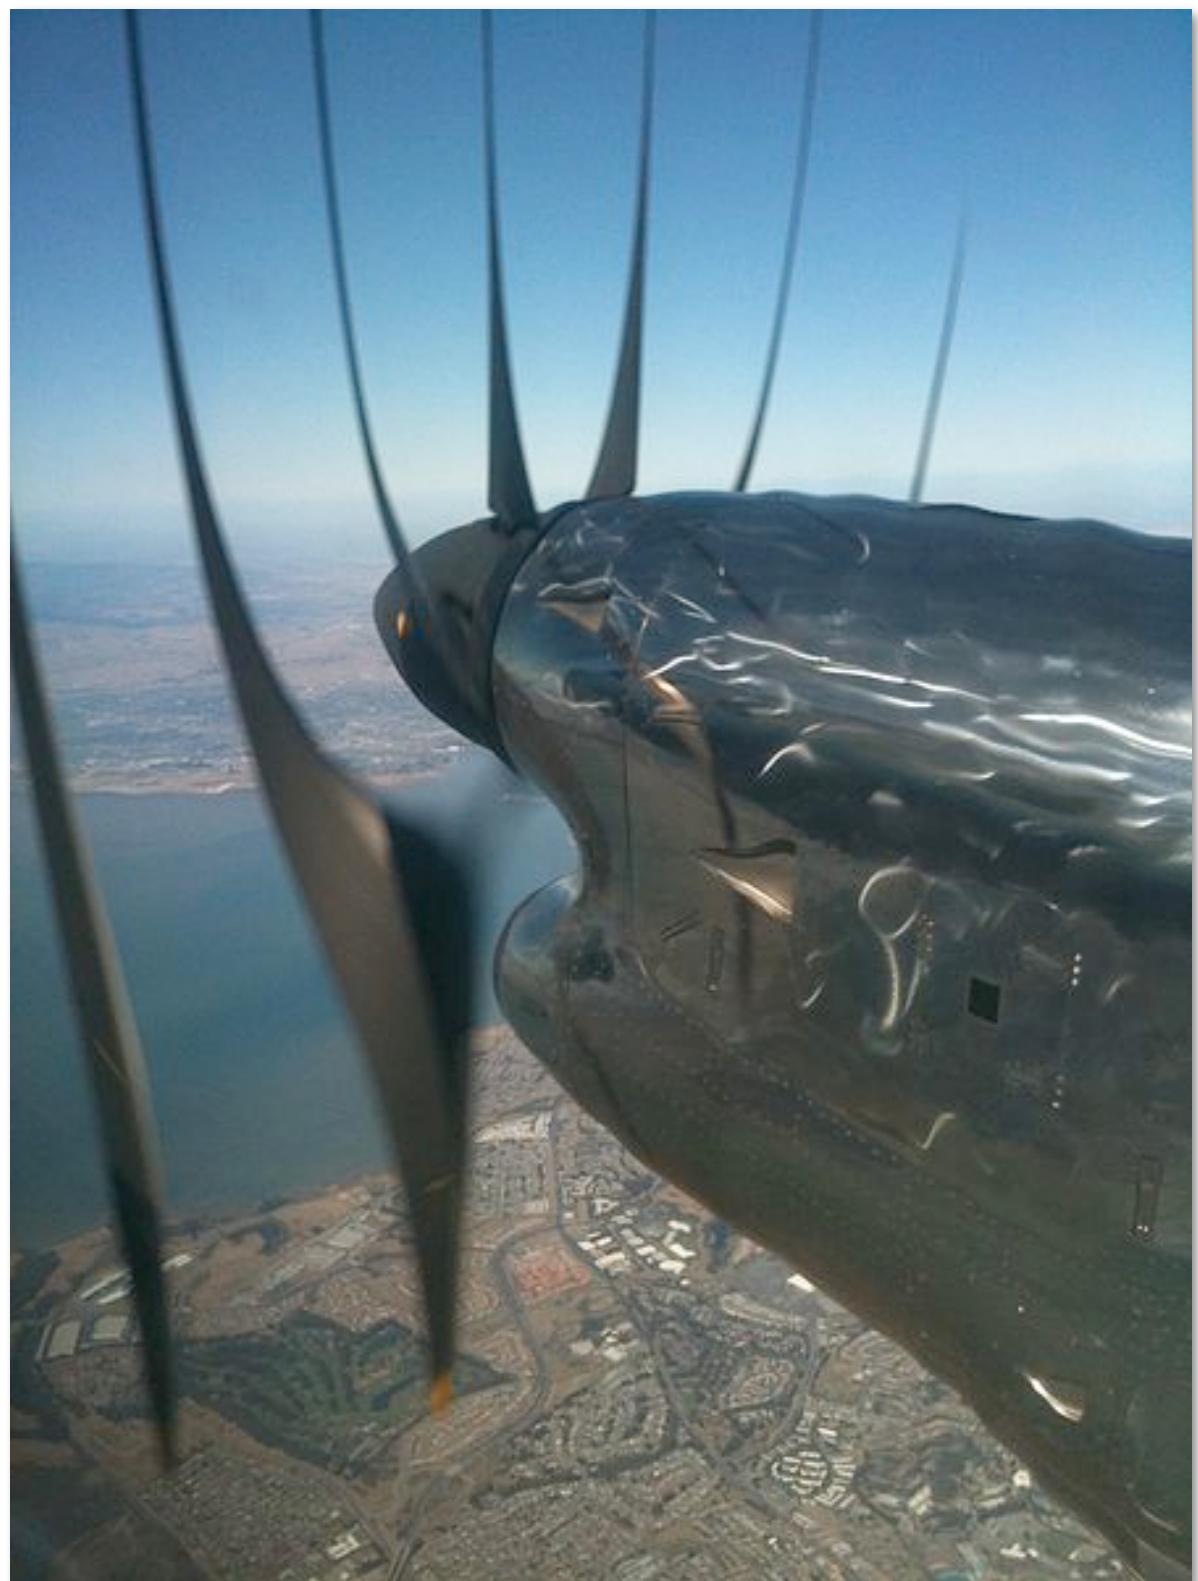


Image credit: Wikipedia

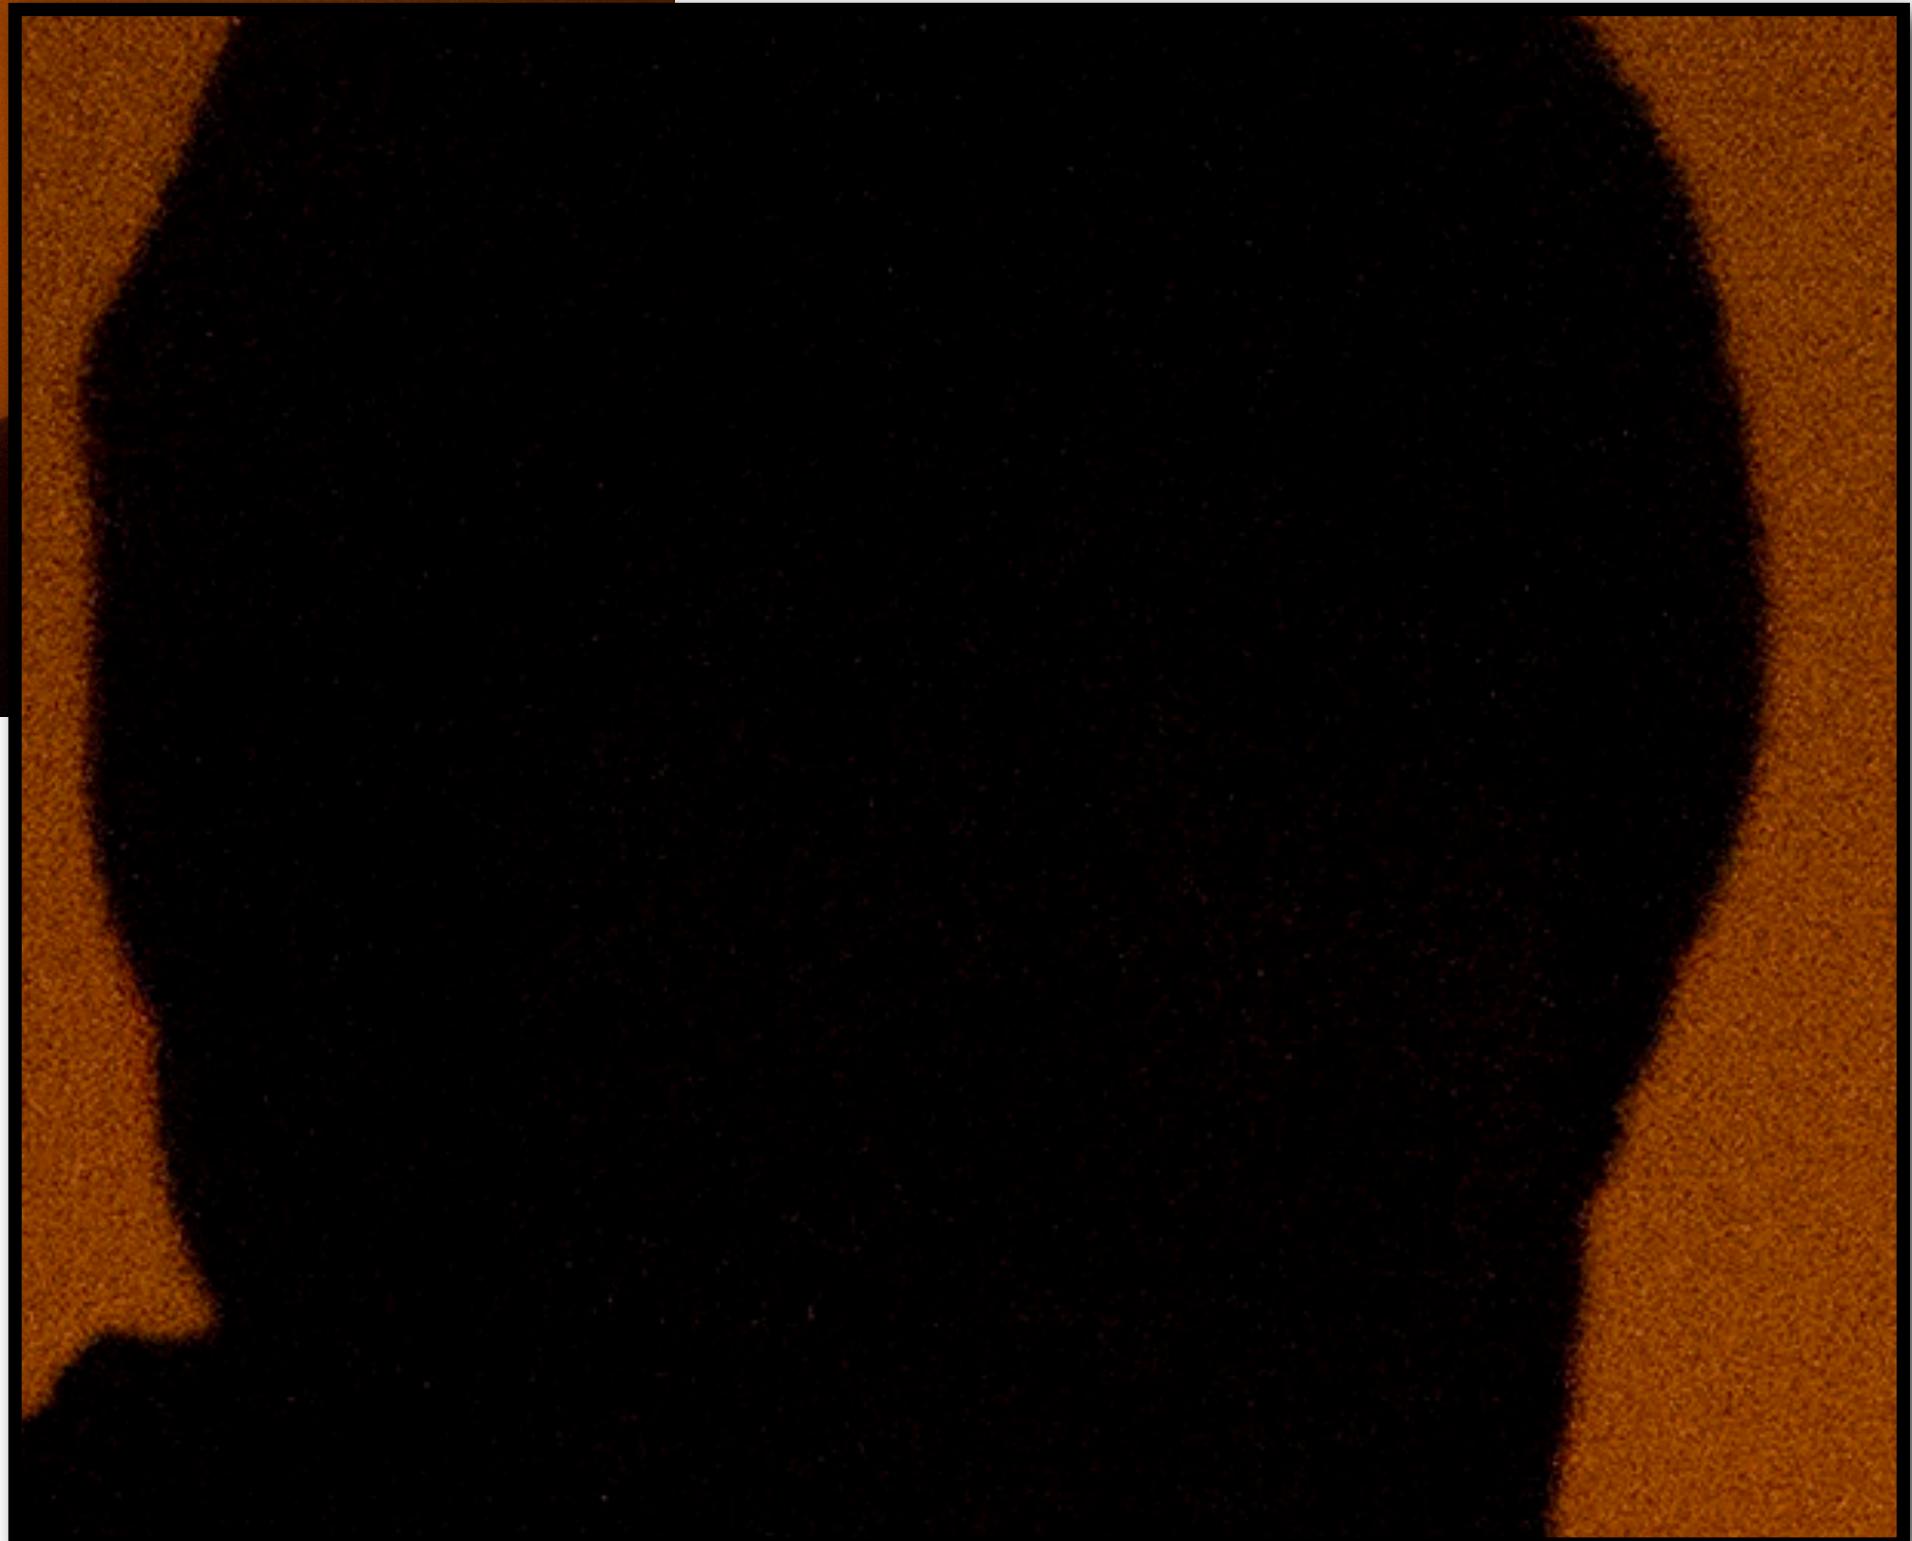


Image credit: Point Grey Research

Measurement noise



We've all been frustrated by noise in low-light photographs
(or in shadows in daytime images)



Measurement noise

■ Photon shot noise:

- Photon arrival rates feature poisson distribution
- Standard deviation = \sqrt{N}
- Signal-to-noise ratio (SNR): N/\sqrt{N}
- Brighter the signal the higher the SNR

■ Dark-shot noise



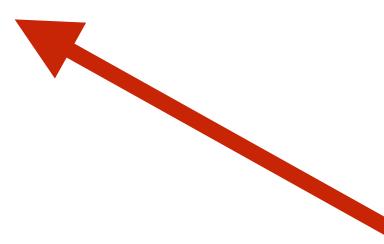
Addressed by: subtract dark image

- Due to leakage current in sensor
- Electrons dislodged due to thermal activity (increases exponentially with sensor temperature)

■ Non-uniformity of pixel sensitivity (due to manufacturing defects)

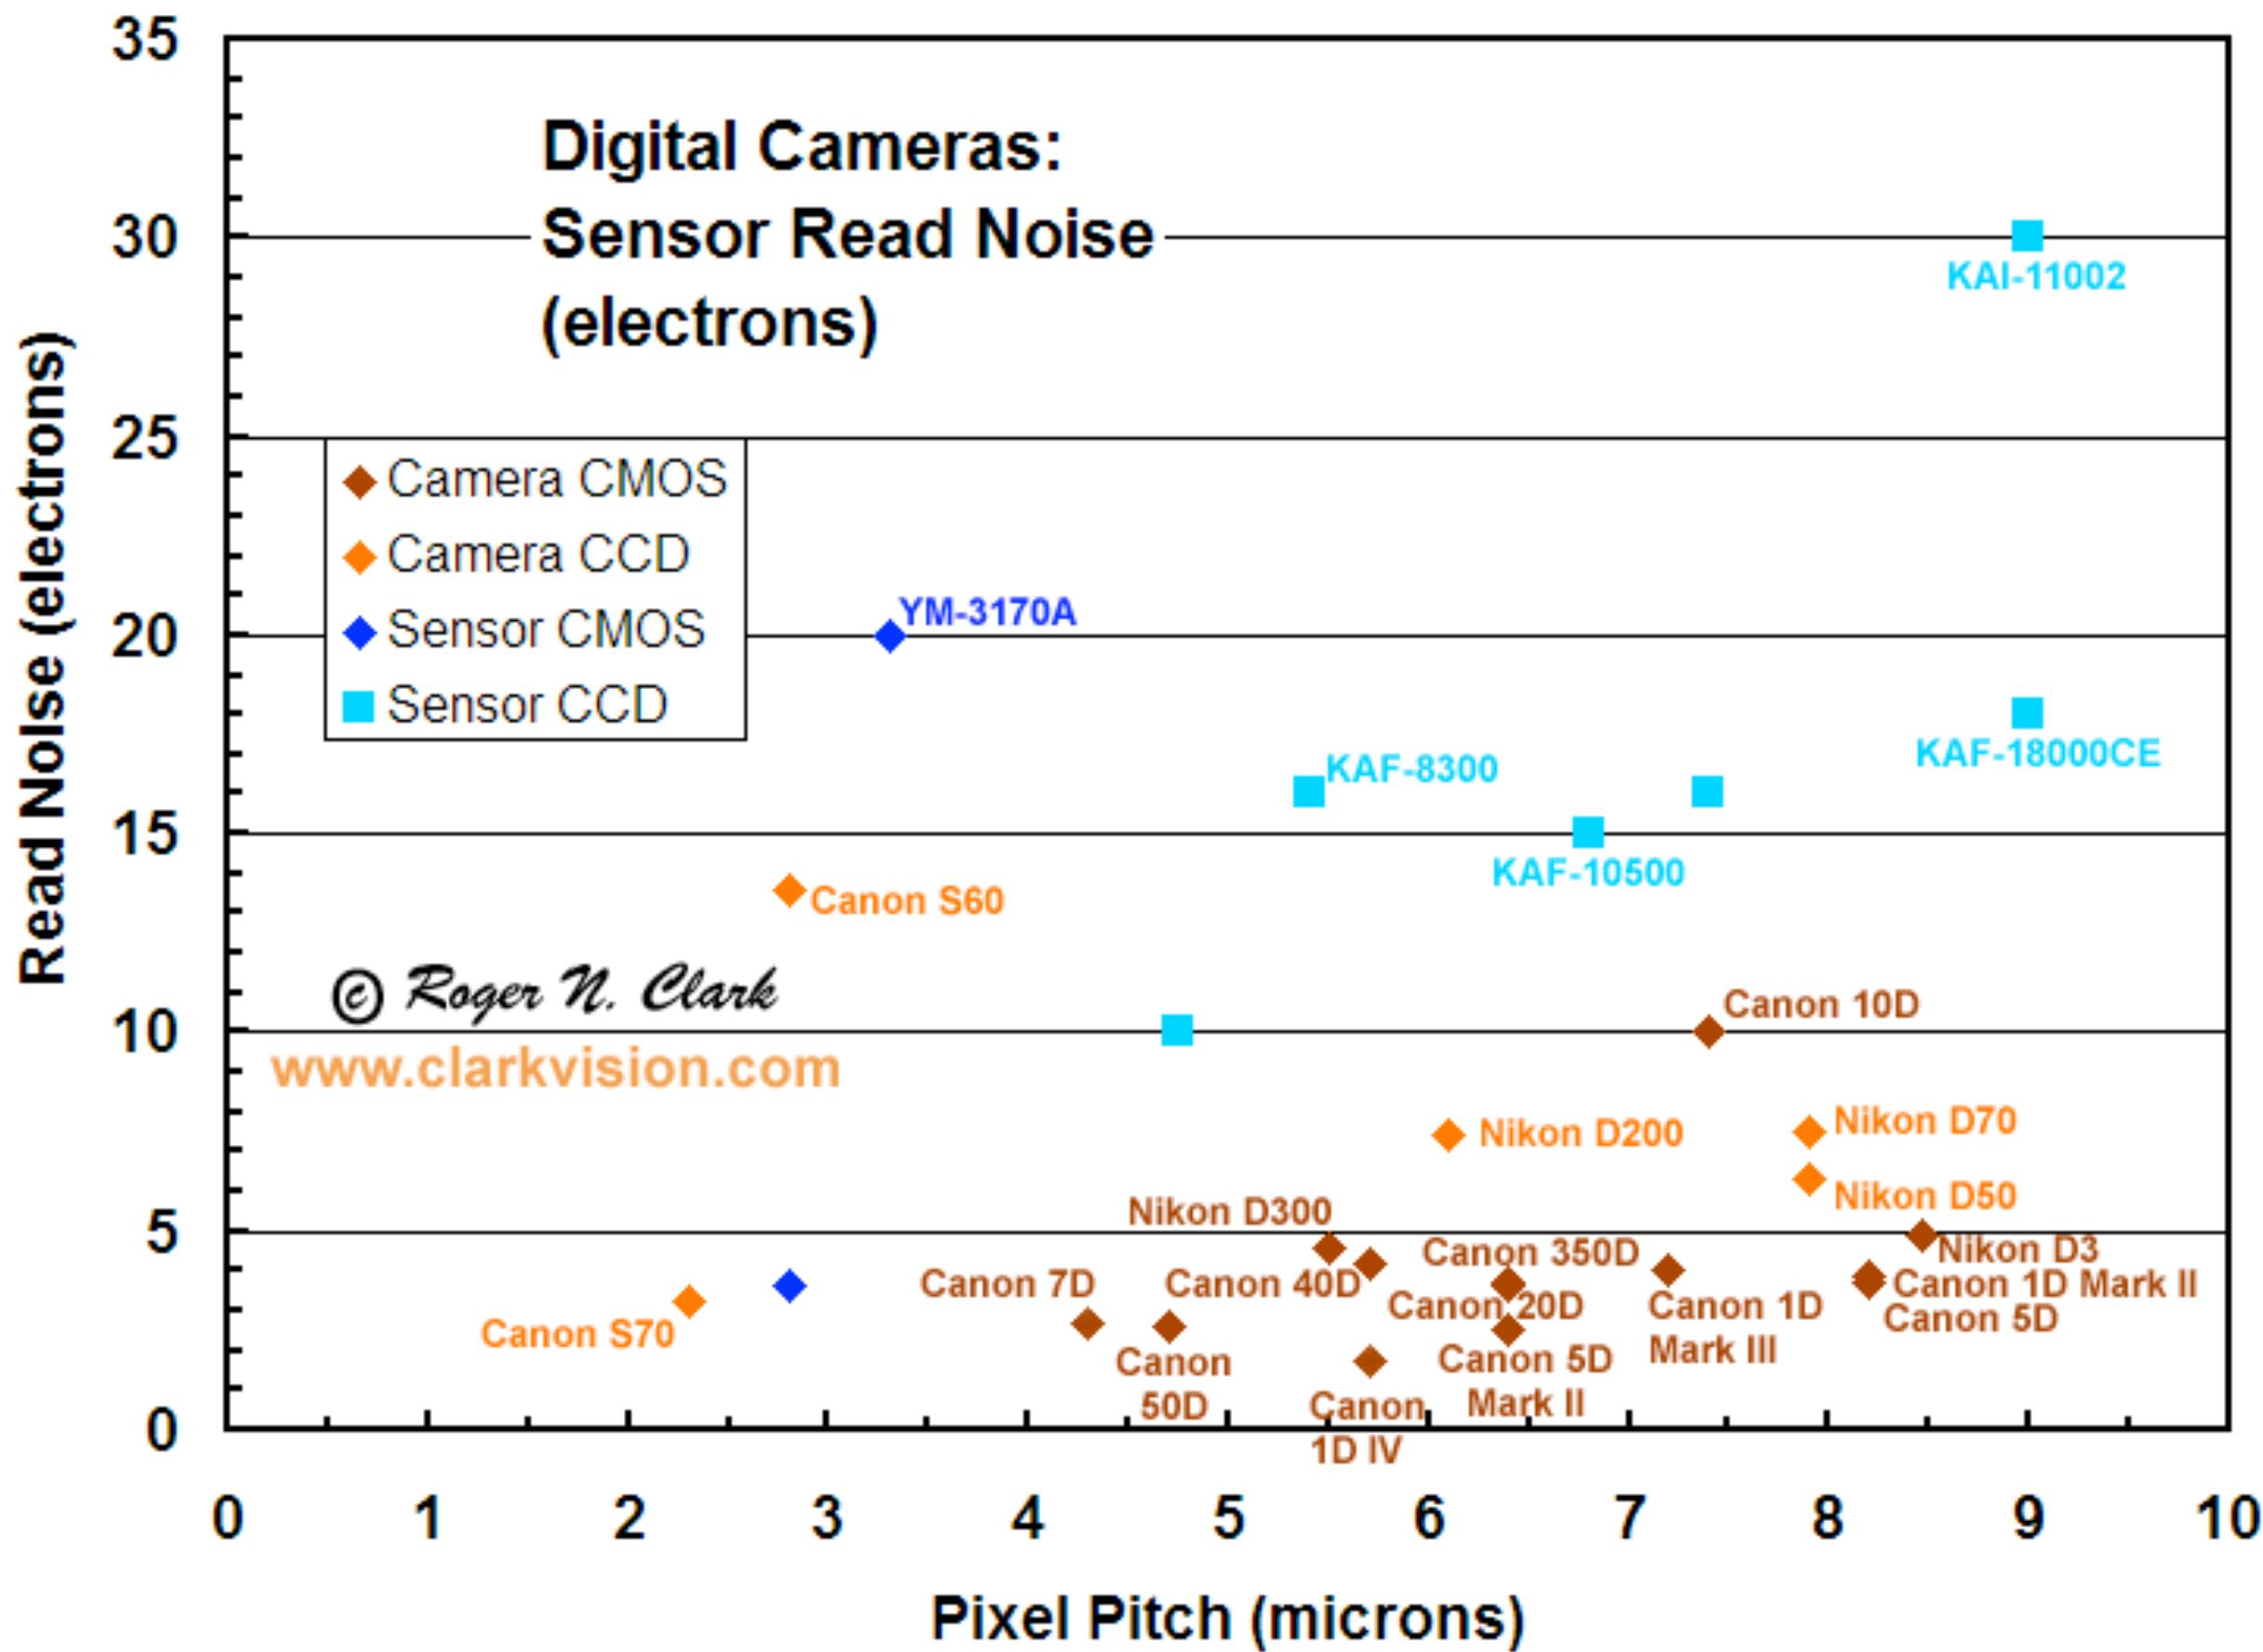
■ Read noise

- e.g., due to amplification / ADC



Addressed by: subtract flat field image
(e.g., image of gray wall)

Read noise

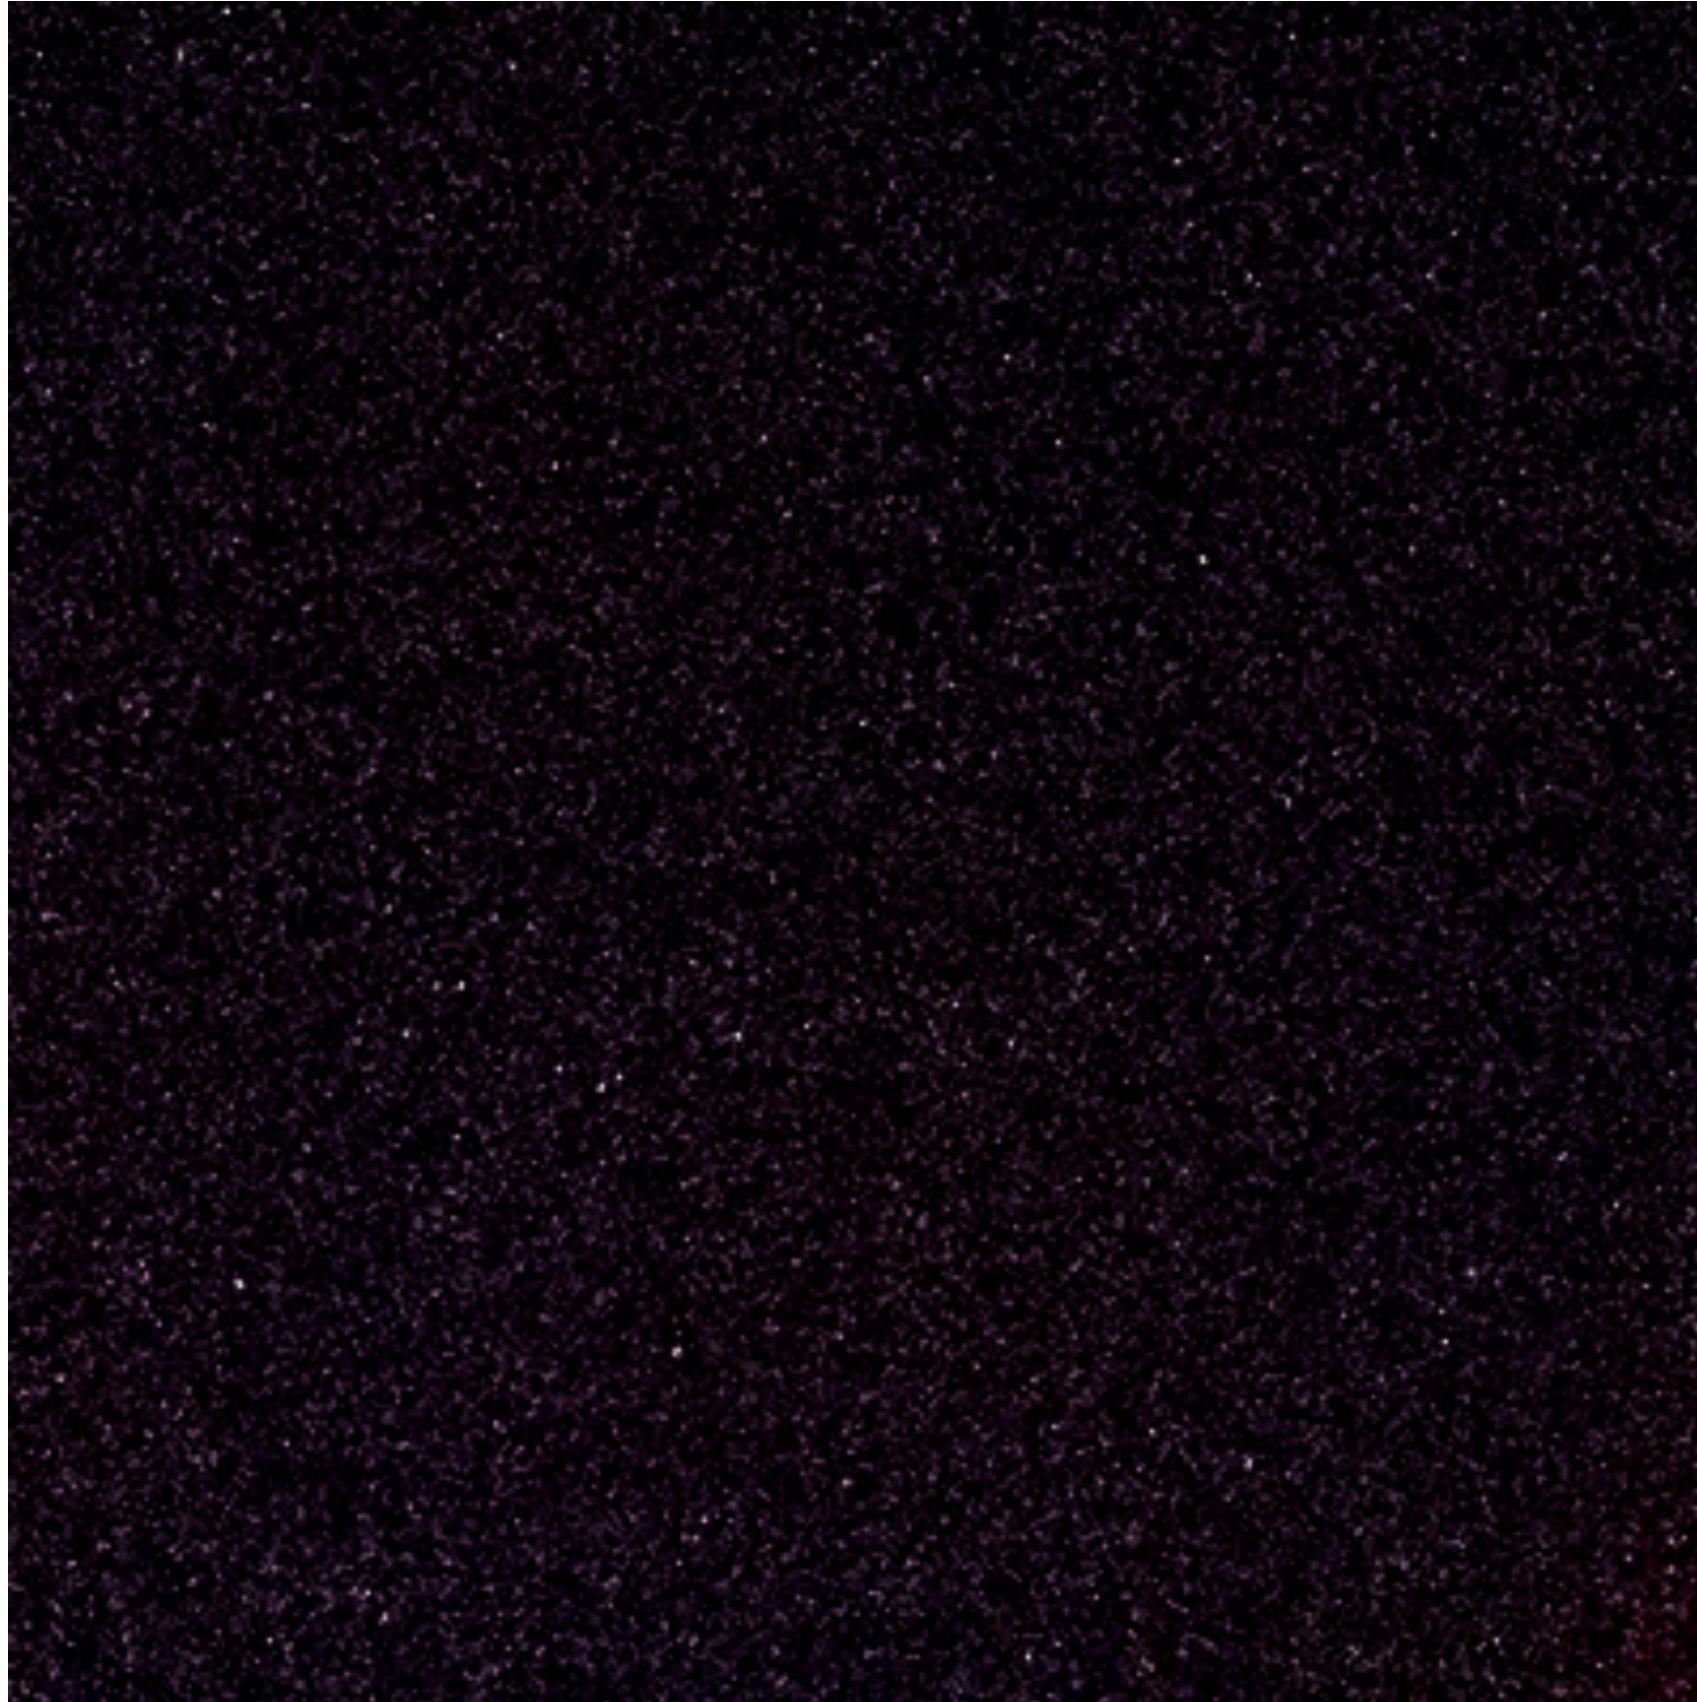


Read noise is largely independent of pixel size

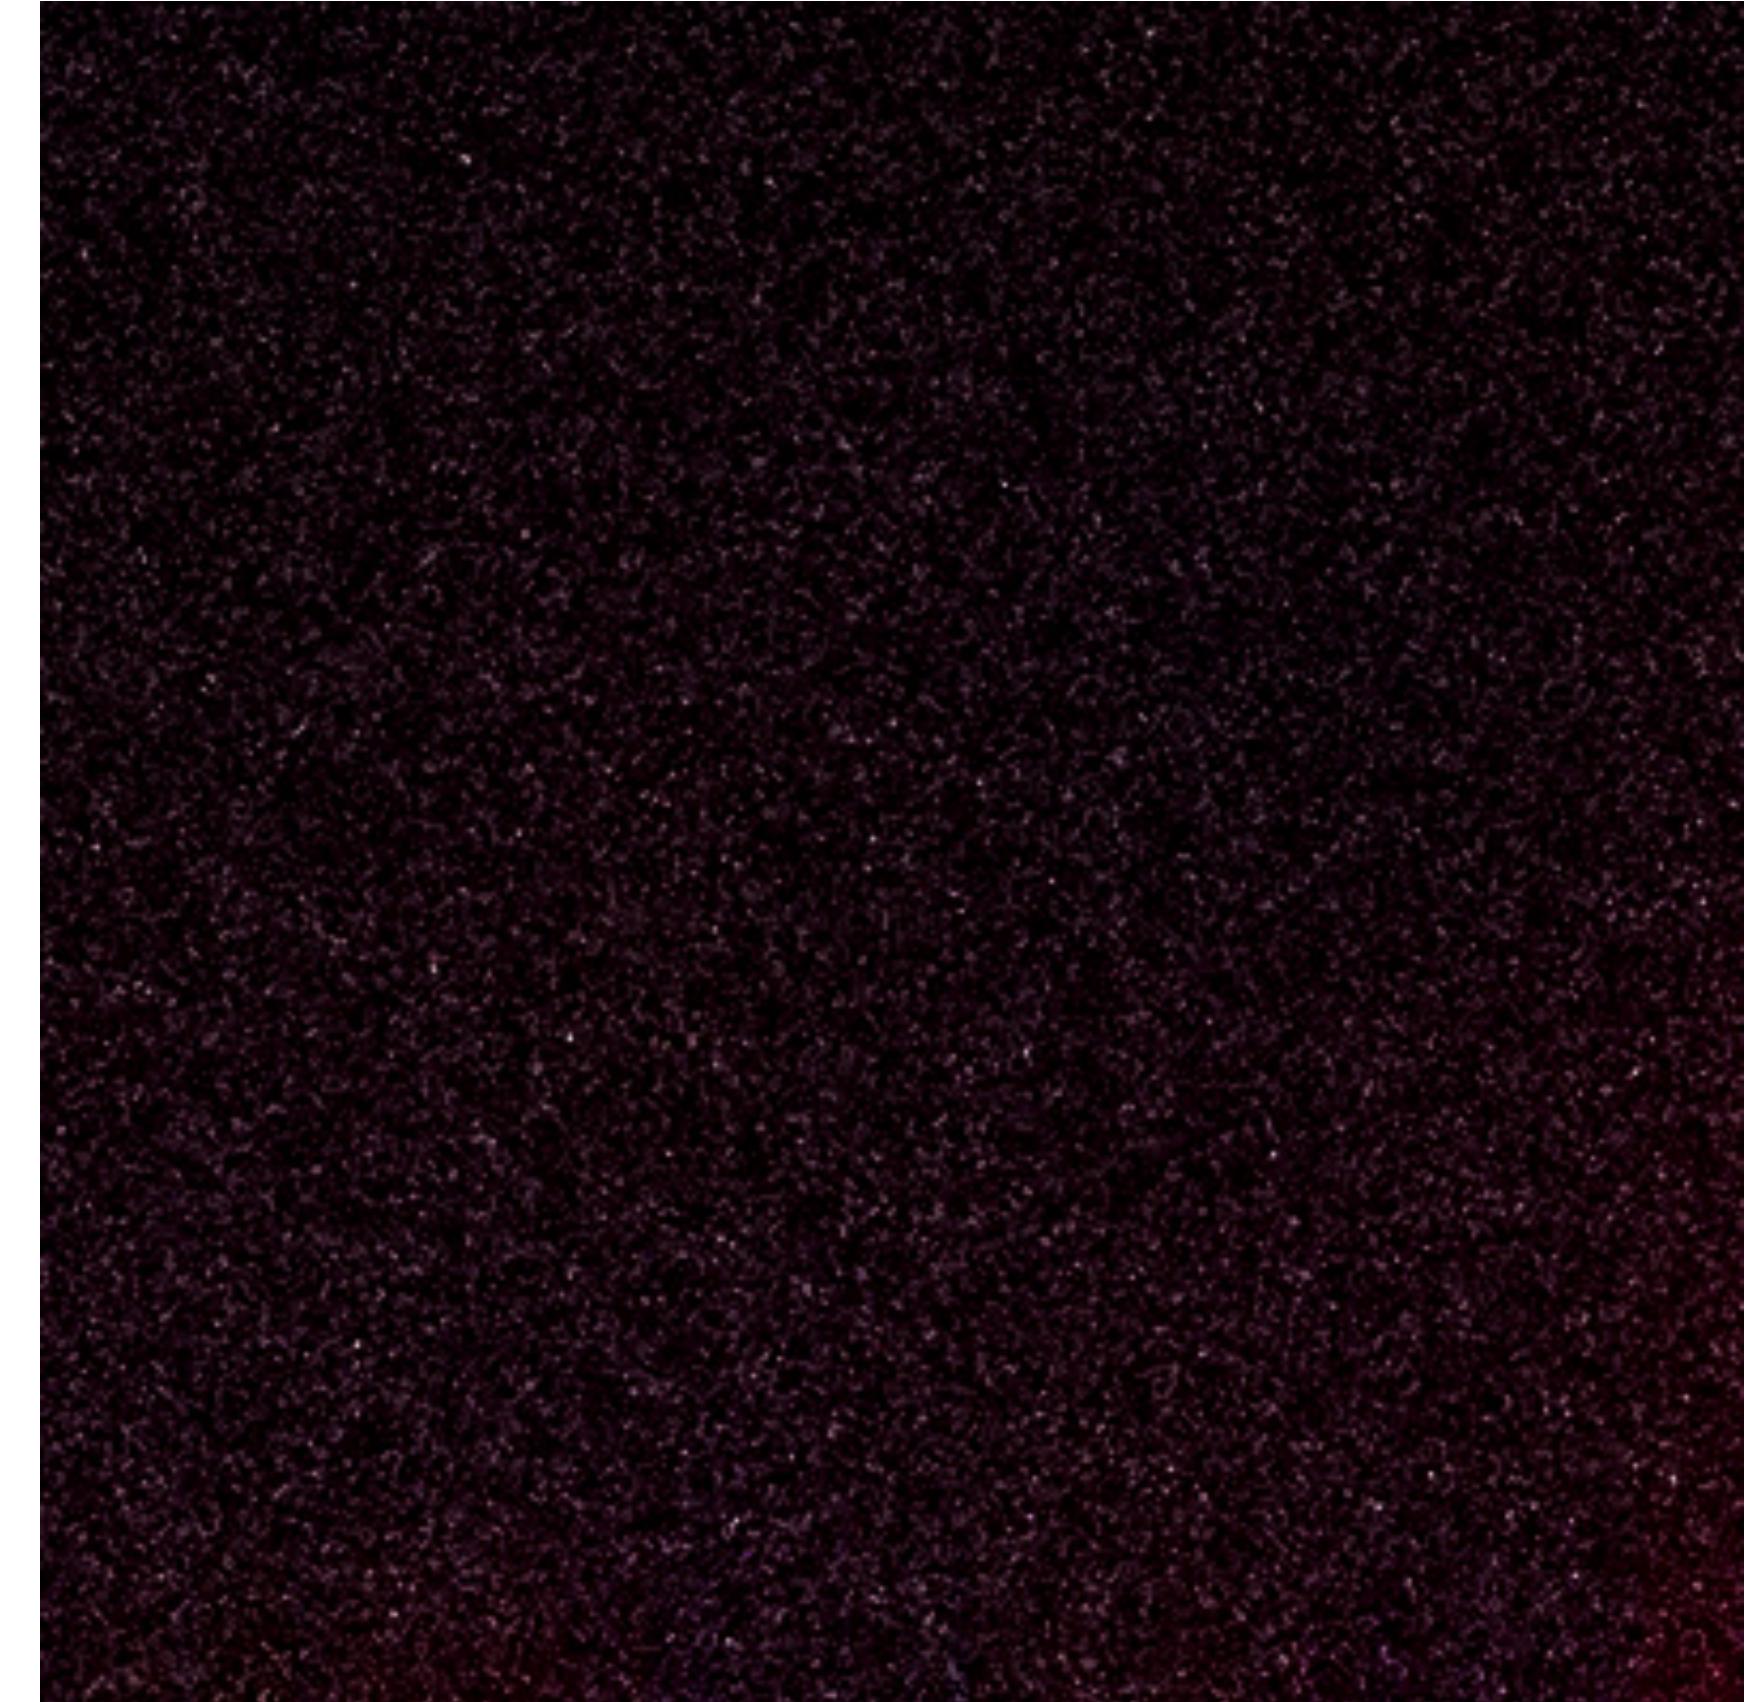
Large pixels + bright scene: noise determined largely by photon shot noise

Dark shot noise / read noise

Black image examples: Nikon D7000, High ISO



1/60 sec exposure



1 sec exposure

Maximize light gathering capability

■ Goal: increase signal-to-noise ratio

- Dynamic range of a pixel (ratio of brightest light measurable to dimmest light measurable) is determined by the noise floor (minimum signal) and the pixel's full-well capacity (maximum signal)

■ Big pixels

- Nikon D4: 7.3 μm
- iPhone 5s: 1.5 μm

■ Sensitive pixels

- Good materials
- High fill factor

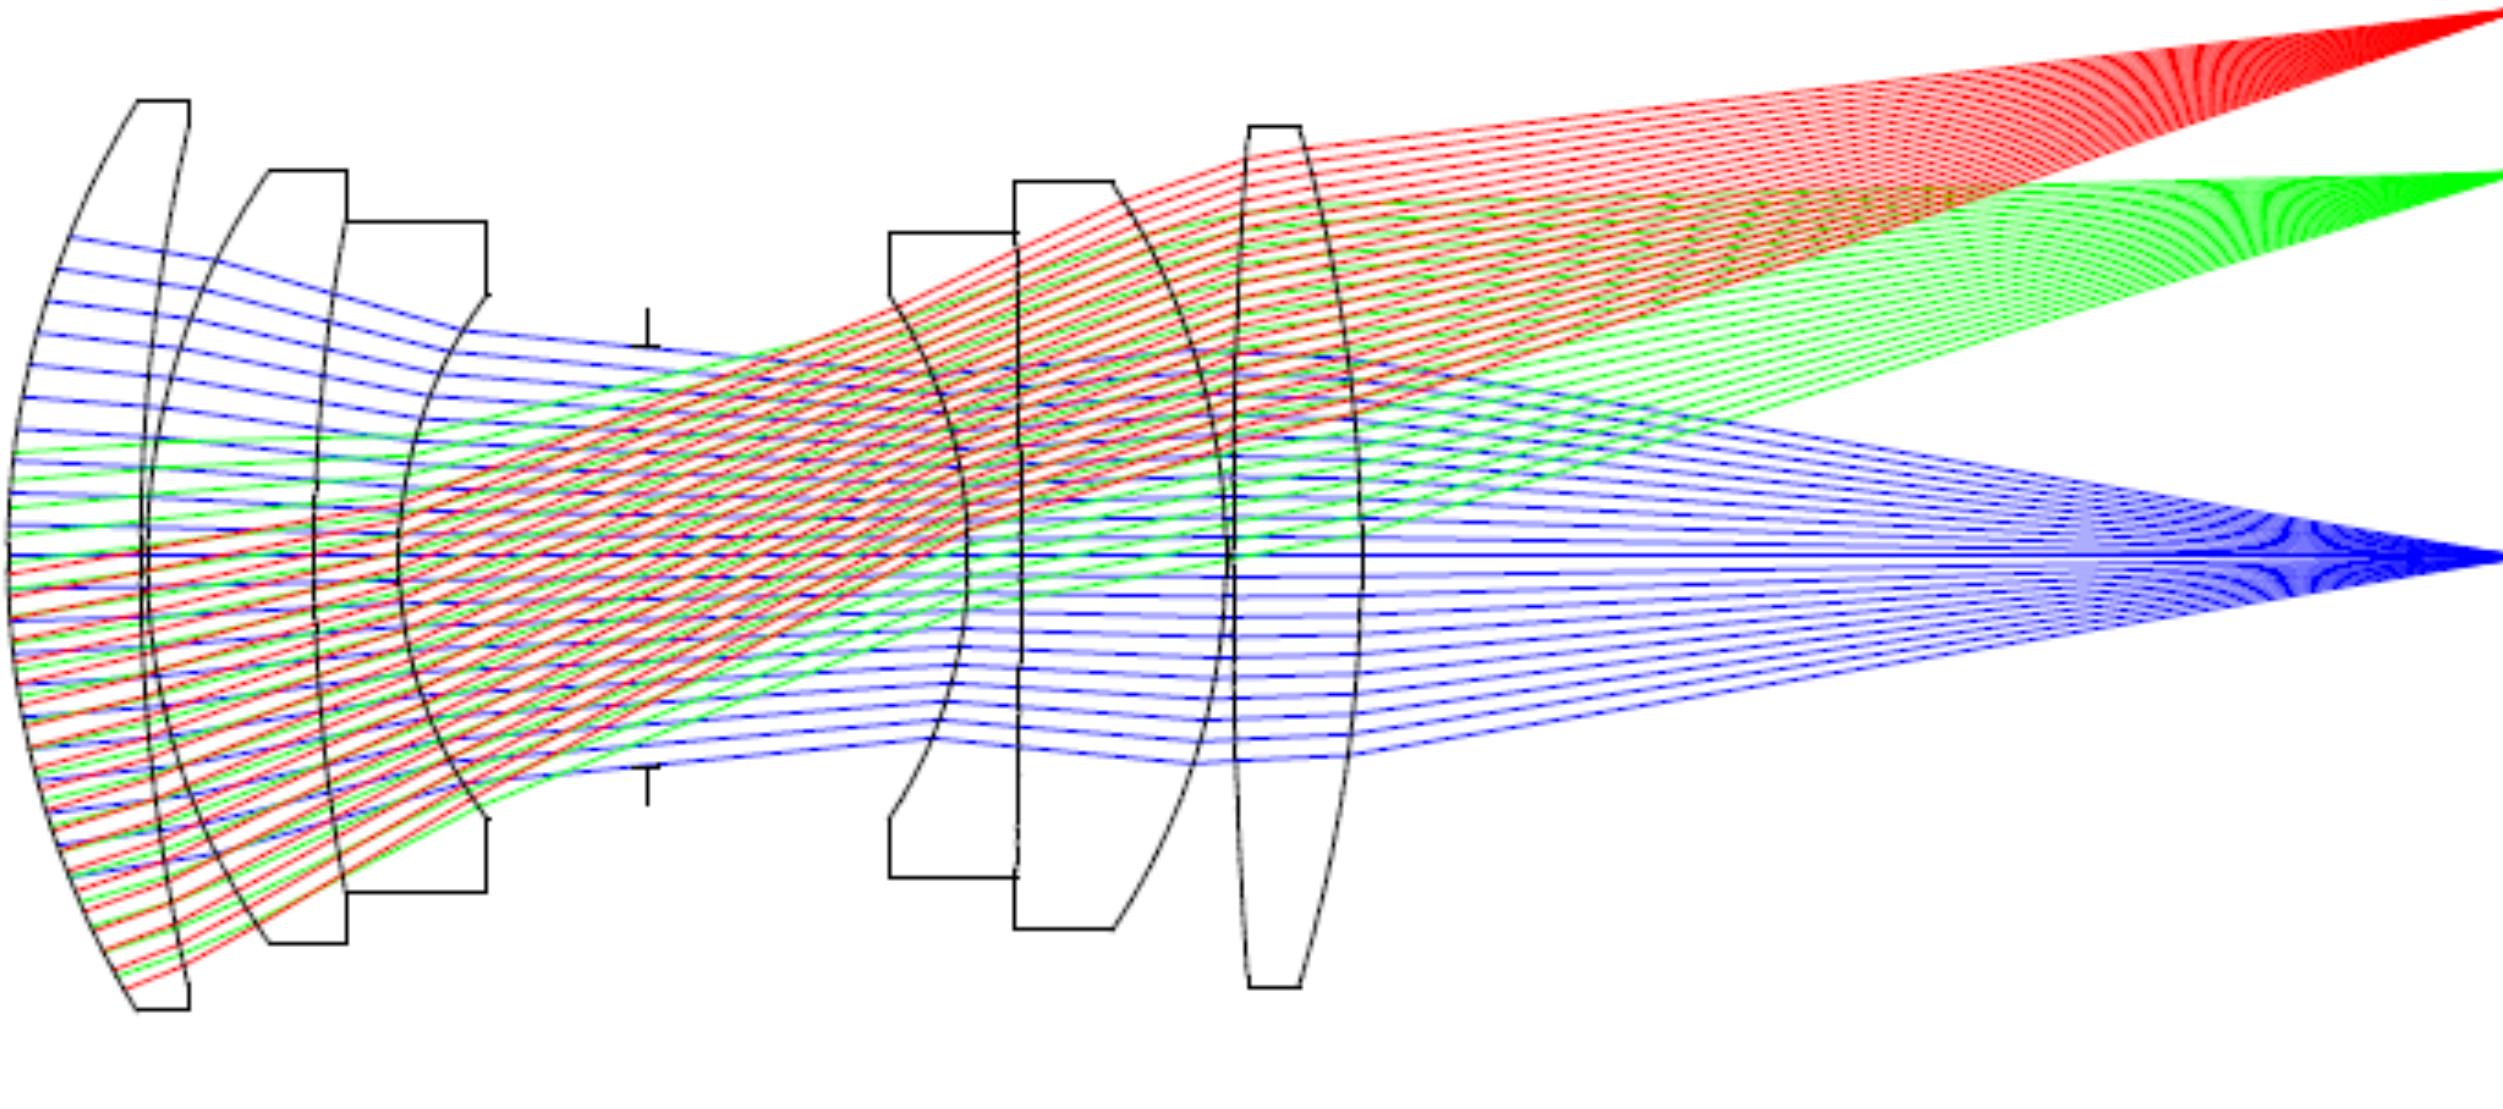
Vignetting

Image of white wall (Note: I contrast-enhanced the image to show effect)



Types of vignetting

Optical vignetting: less light reaches edges of sensor due to physical obstruction in lens



Pixel vignetting: light reaching pixel at an oblique angle is less likely to hit photosensitive region than light incident from straight above (e.g., obscured by electronics)

- **Microlens reduces pixel vignetting**

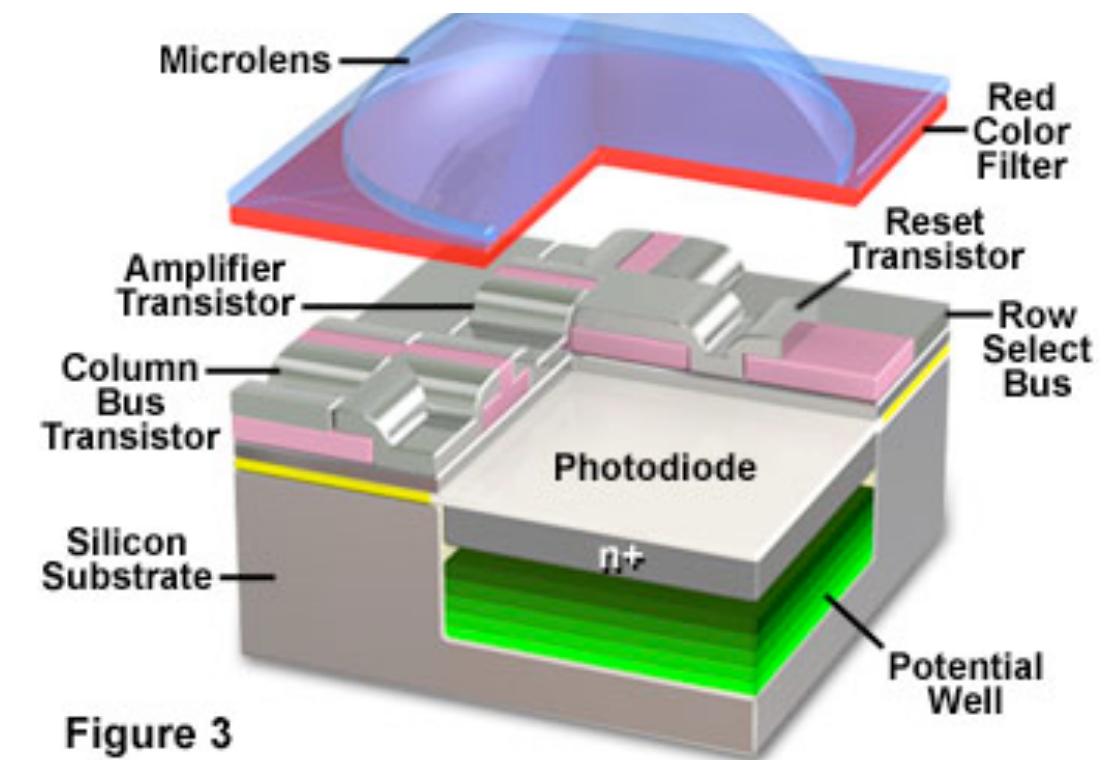


Figure 3

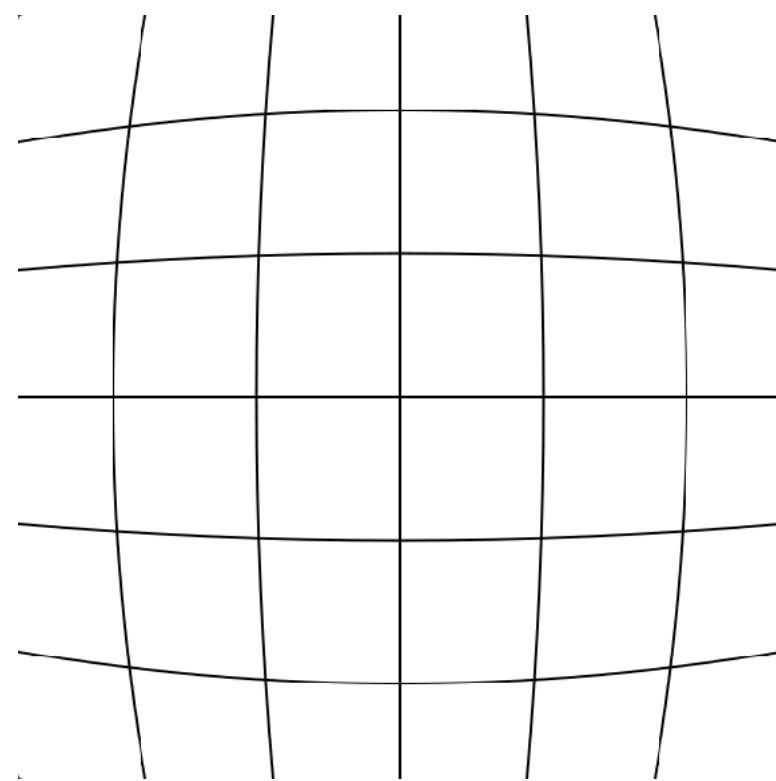
More challenges

■ Chromatic shifts over sensor

- Pixel light sensitivity changes over sensor due to interaction with microlens
(Index of refraction depends on wavelength, so some wavelengths are more likely to suffer from cross-talk or reflection. Ug!)

■ Dead pixels (stuck at white or black)

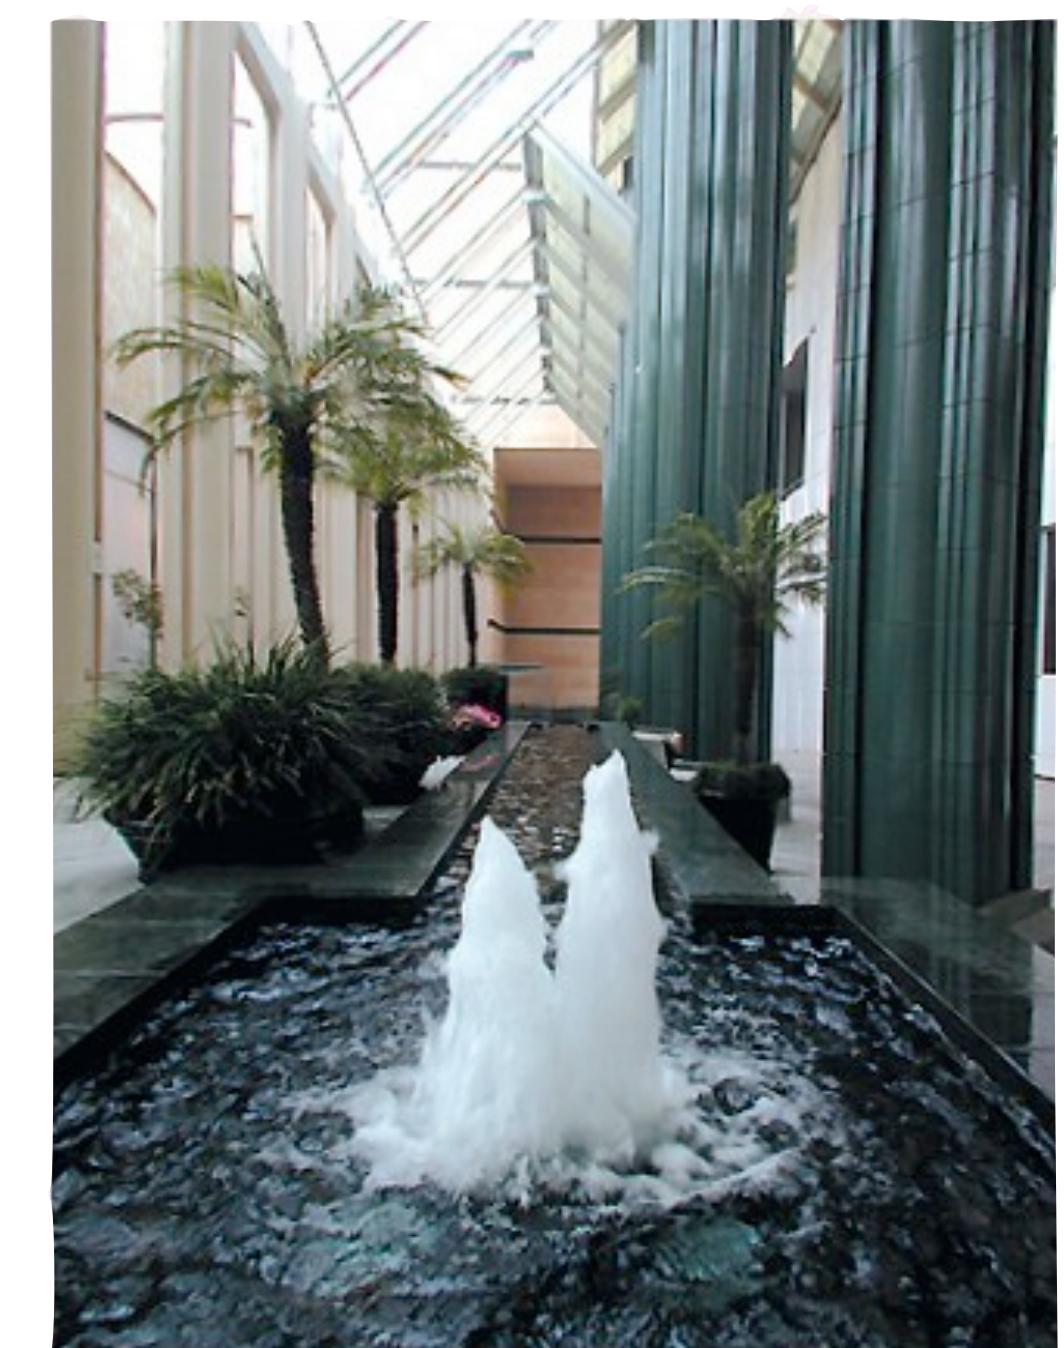
■ Lens distortion



Pincushion distortion



Captured Image



Corrected Image

Image credit: PCWorld

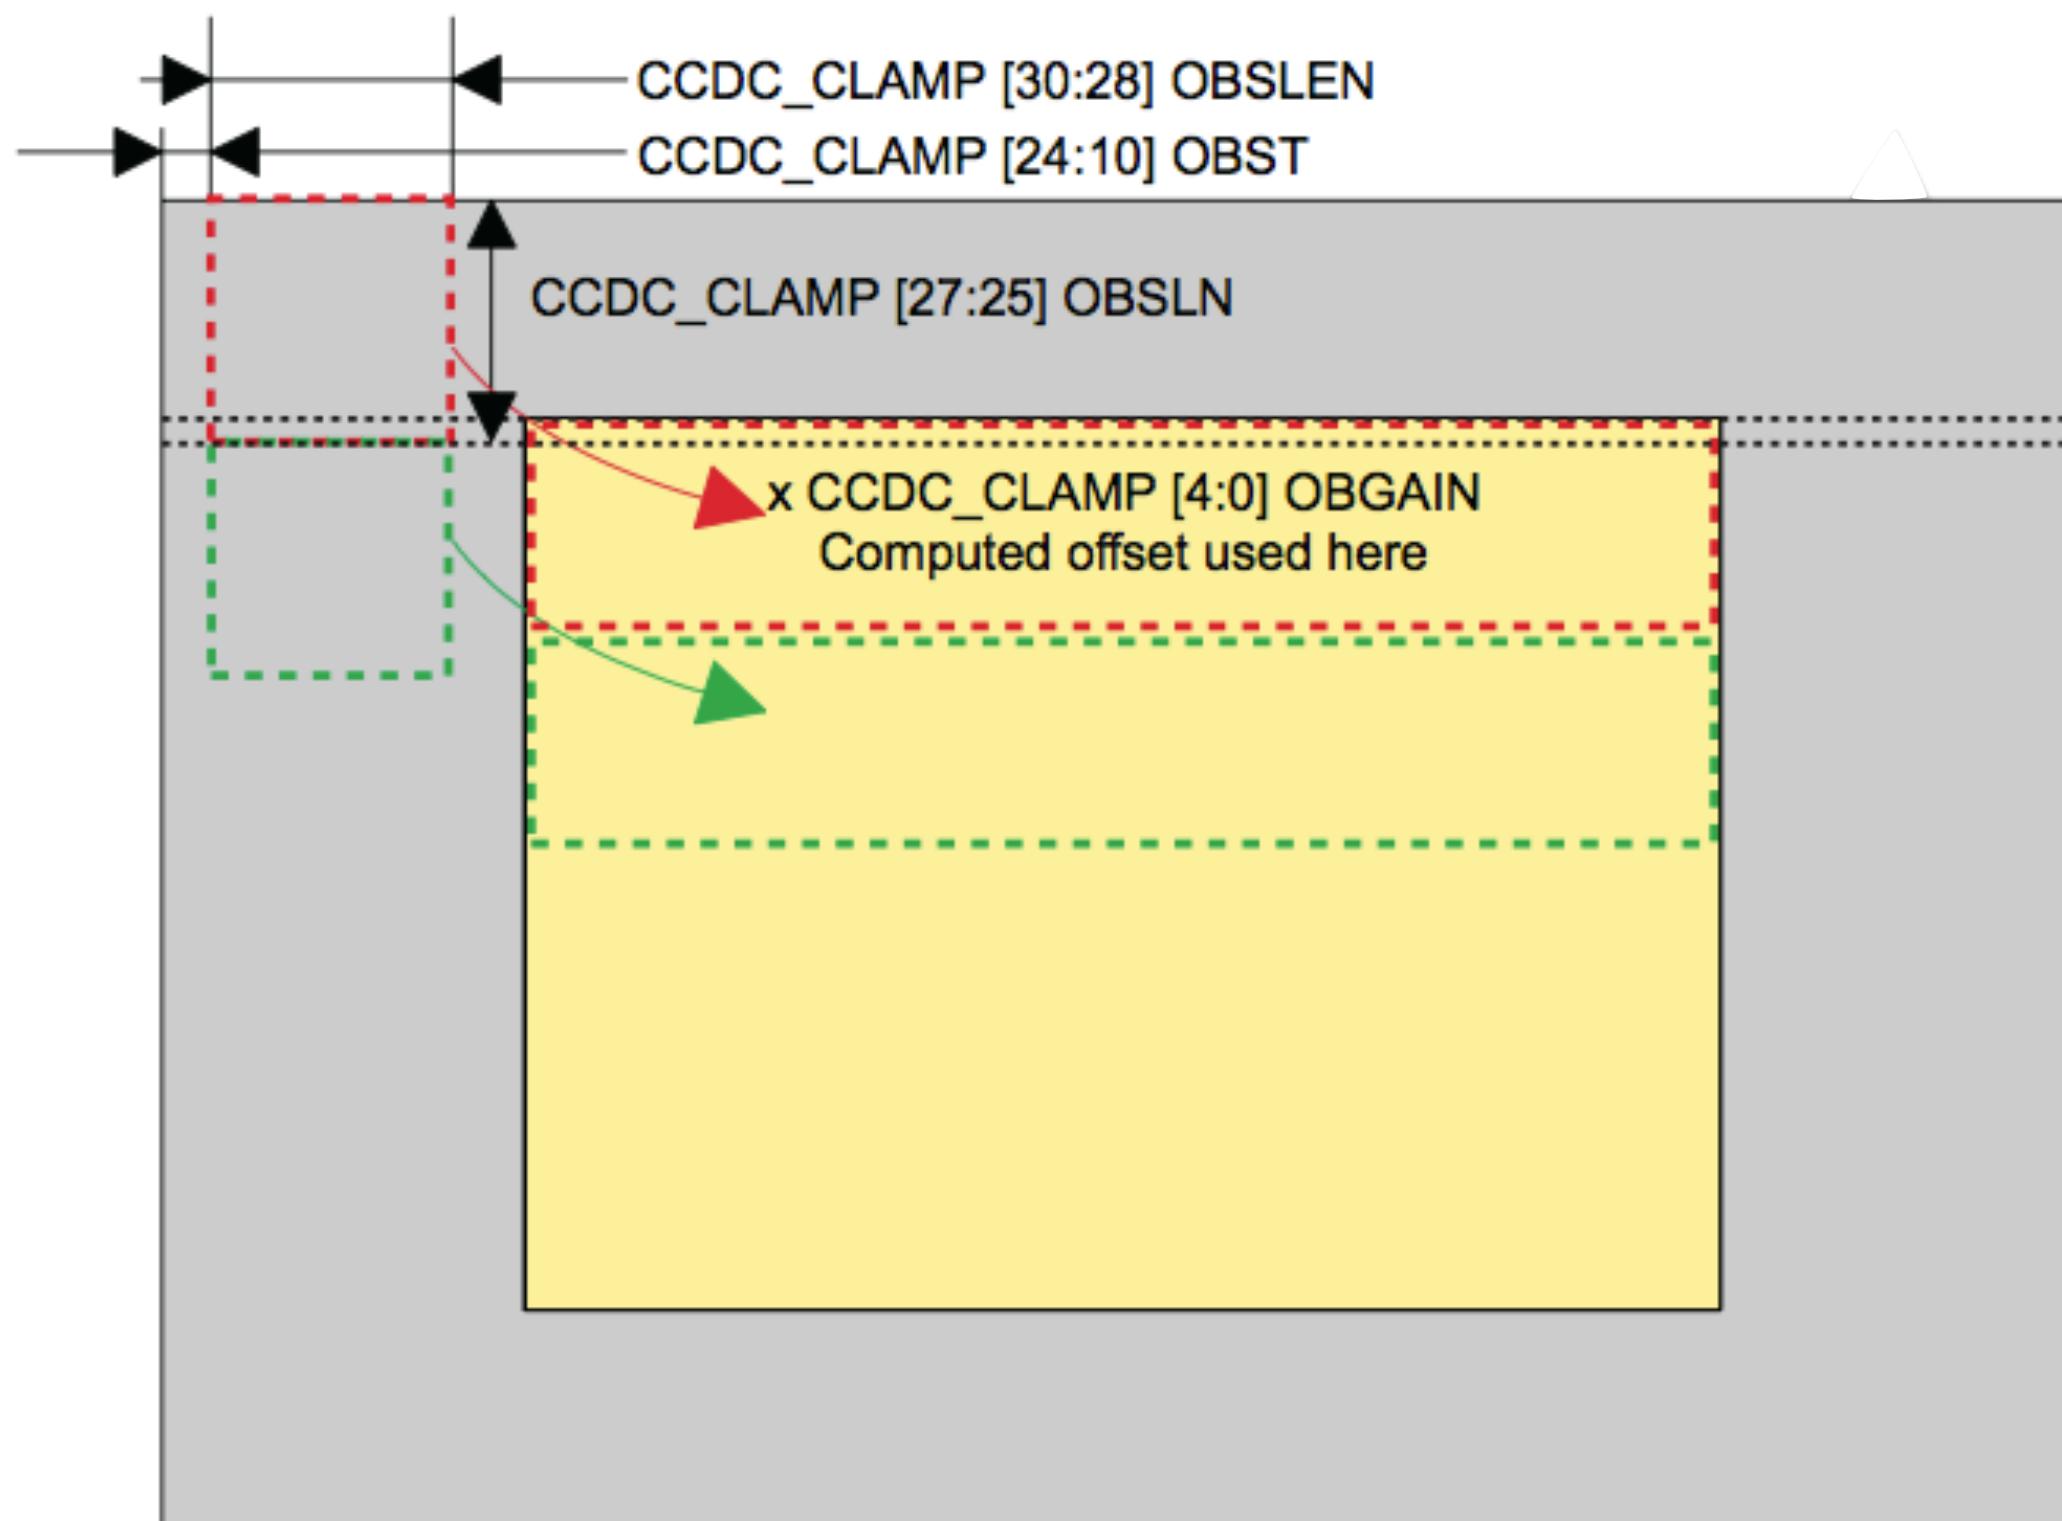
A simple RAW image processing pipeline

**Adopting terminology from Texas Instruments OMAP Image Signal Processor pipeline
(since public documentation exists)**

Assume: receiving 12 bits/pixel Bayer mosaiced data from sensor

Optical clamp: remove sensor offset bias

`output_pixel = input_pixel - [average of pixels from optically black region]`



**Remove bias due to sensor black level
(from nearby sensor pixels at time of shot)**

- Masked pixels
- Active pixels

Step 2: correct for defective pixels

- **Store LUT with known defective pixels**
 - e.g., determined on manufacturing line, during sensor calibration and test
- **Example correction methods**
 - Replace defective pixel with neighbor
 - Replace defective pixel with average of neighbors
 - Correct defect by subtracting known bias for the defect

```
output_pixel = (isdefectpixel(current_pixel_xy)) ?  
    average(previous_input_pixel, next_input_pixel) :  
    input_pixel;
```

Lens shading compensation

■ Correct for vignetting

- Recall good implementations will consider wavelength-dependent vignetting (that creates chromatic shift over the image)

■ Possible implementations:

- Use flat-field photo stored in memory
 - e.g., lower resolution buffer, upsampled on-the-fly
 - Use analytic function to model correction

```
offset = upsample_compensation_offset_buffer(current_pixel_xy);  
gain = upsample_compensation_gain_buffer(current_pixel_xy);  
  
output_pixel = offset + gain * input_pixel;
```

Optional dark-frame subtraction

- Similar computation to lens shading compensation

```
output_pixel = input_pixel - dark_frame[current_pixel_xy];
```

“Dark frame”
(result of exposure with shutter closed)

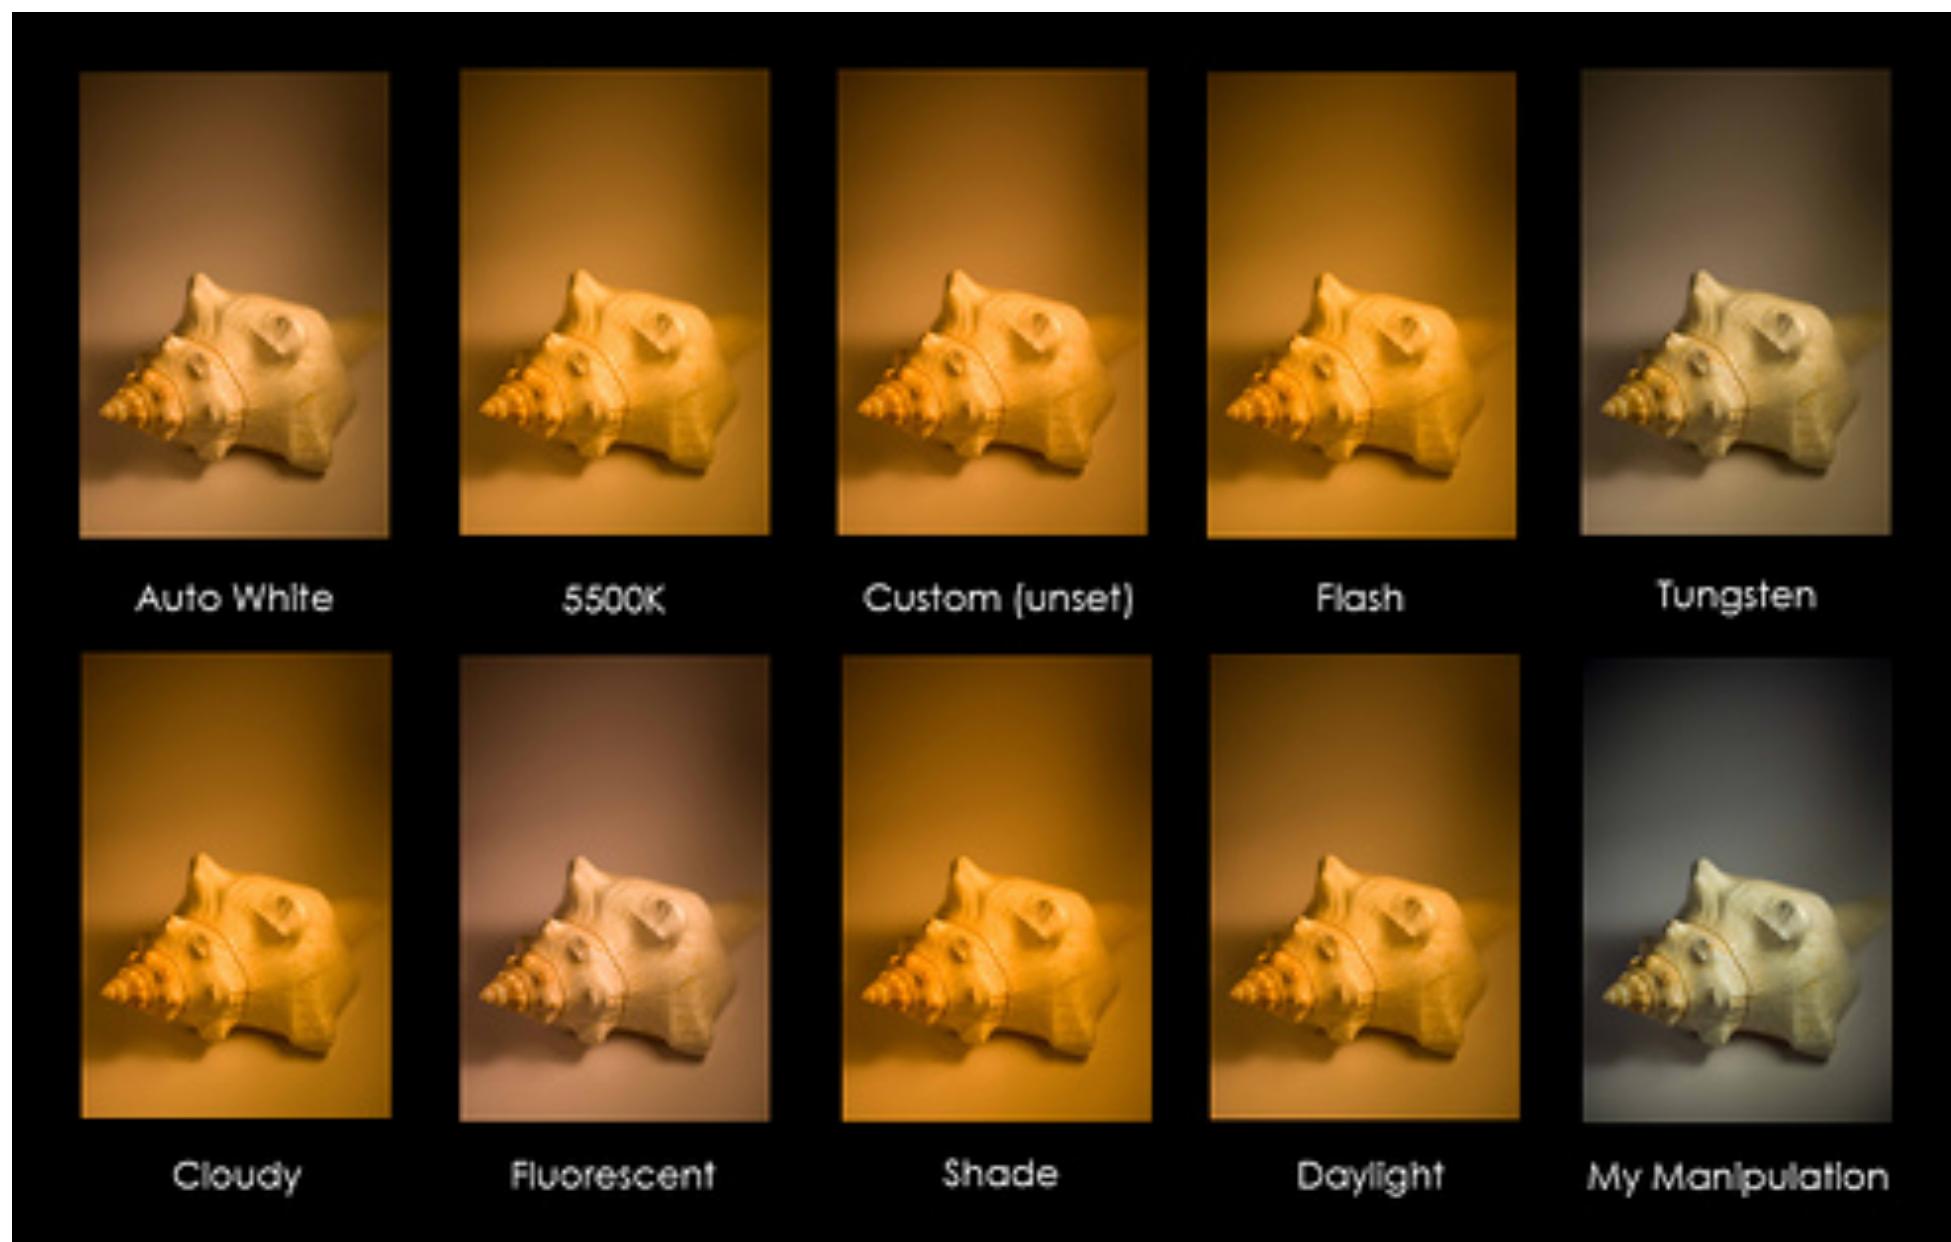


White balance

- Adjust relative intensity of rgb values (so neutral tones appear neutral)

```
output_pixel = white_balance_coeff * input_pixel
// note: in this example, white_balance_coeff is vec3
// (adjusts ratio of red-blue-green channels)
```

- White balance coefficients are determined based on analysis of image contents:
 - Simple auto-white balance algorithms
 - Gray world assumption: make average of all pixels in image gray
 - Find brightest region of image, make it white ([1,1,1])
- Modern cameras have sophisticated (heuristic-based) white-balance algorithms



Common data-driven solution:

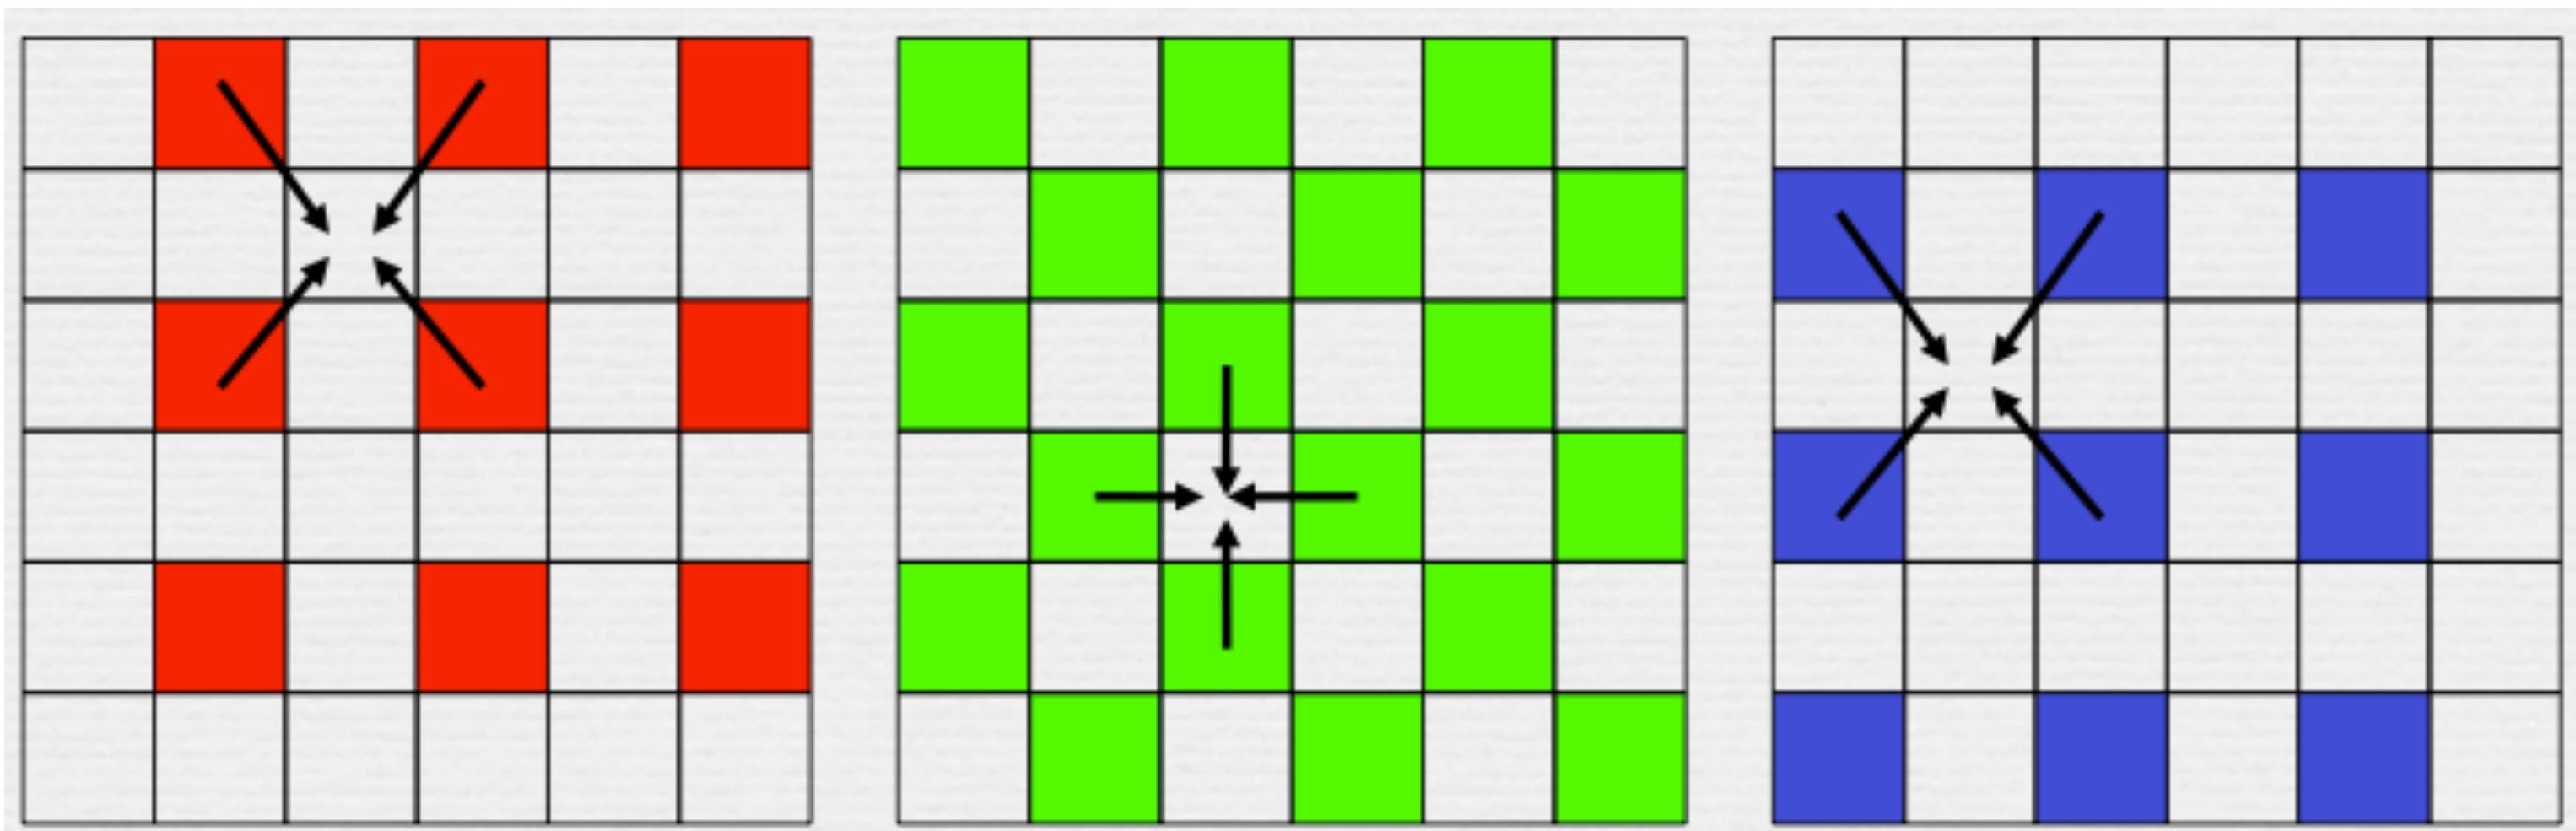
1. Compute features from input image (e.g., histogram)
2. Find similar images in database of images for which good white balance settings are known
3. Use white balance settings from database image

Image credit: basedigitalphotography.com

Stanford CS348V, Winter 2018

Demosiac

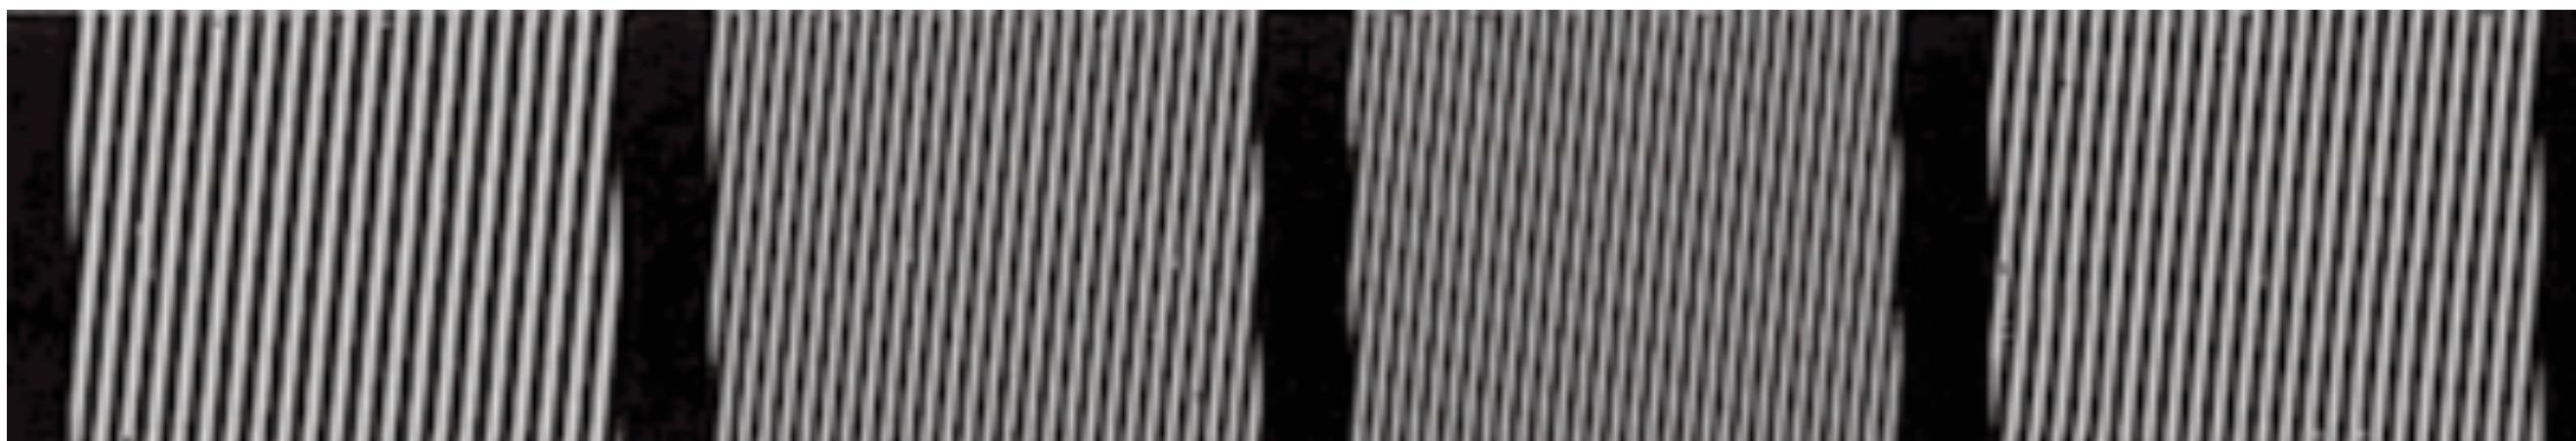
- Produce RGB image from mosaiced input image
- Basic algorithm: bilinear interpolation of mosaiced values (need 4 neighbors)
- More advanced algorithms:
 - Bicubic interpolation (wider filter support region... may overblur)
 - Good implementations attempt to find and preserve edges



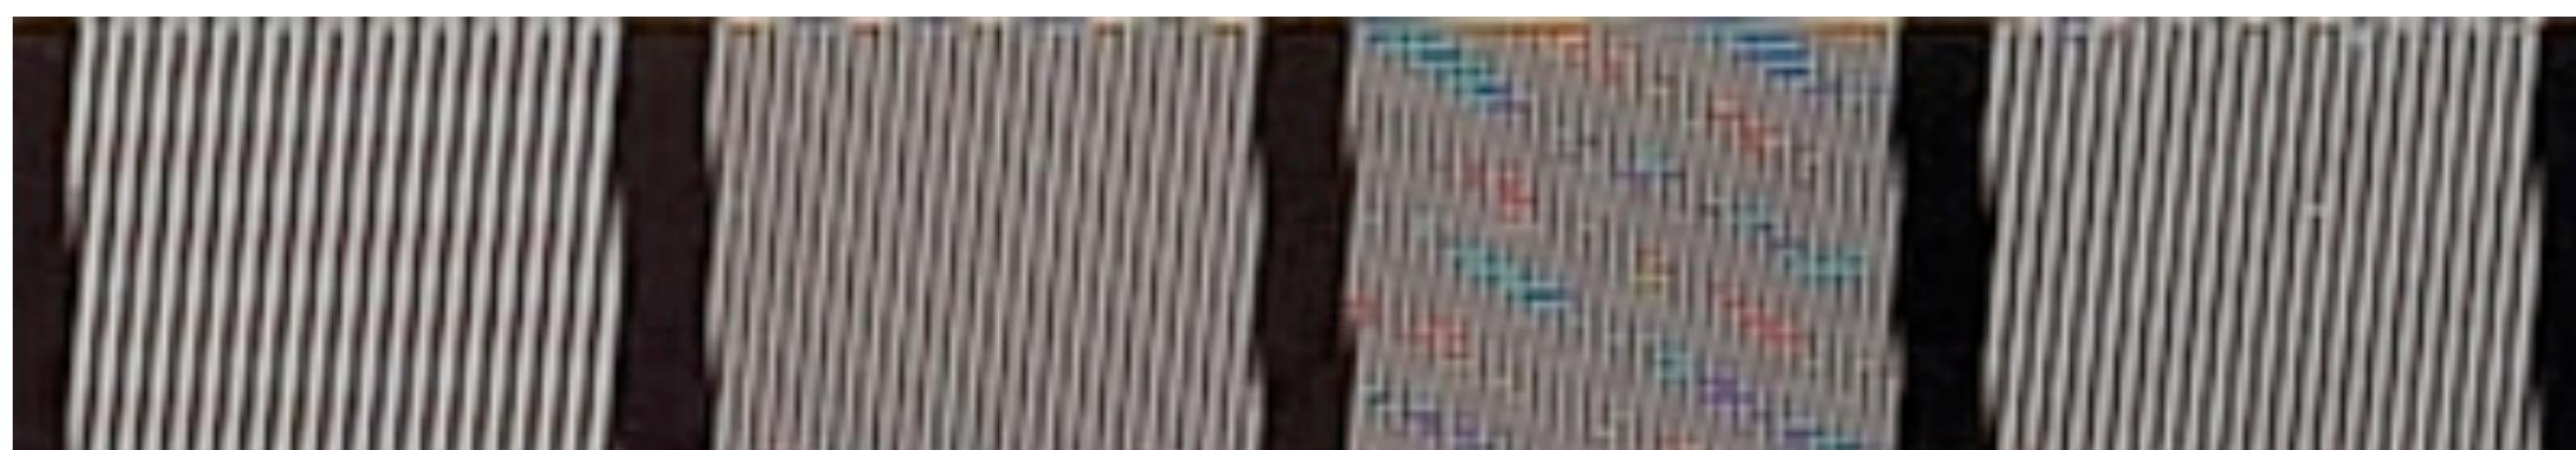
Demosaicing errors

■ Moire pattern color artifacts

- Common difficult case: fine diagonal black and white stripes
- Common solution:
 - Convert demosaiced value to YCbCr
 - Low-pass filter (blur) CbCr channels
 - Combine filtered CbCr with full resolution Y from sensor to get RGB

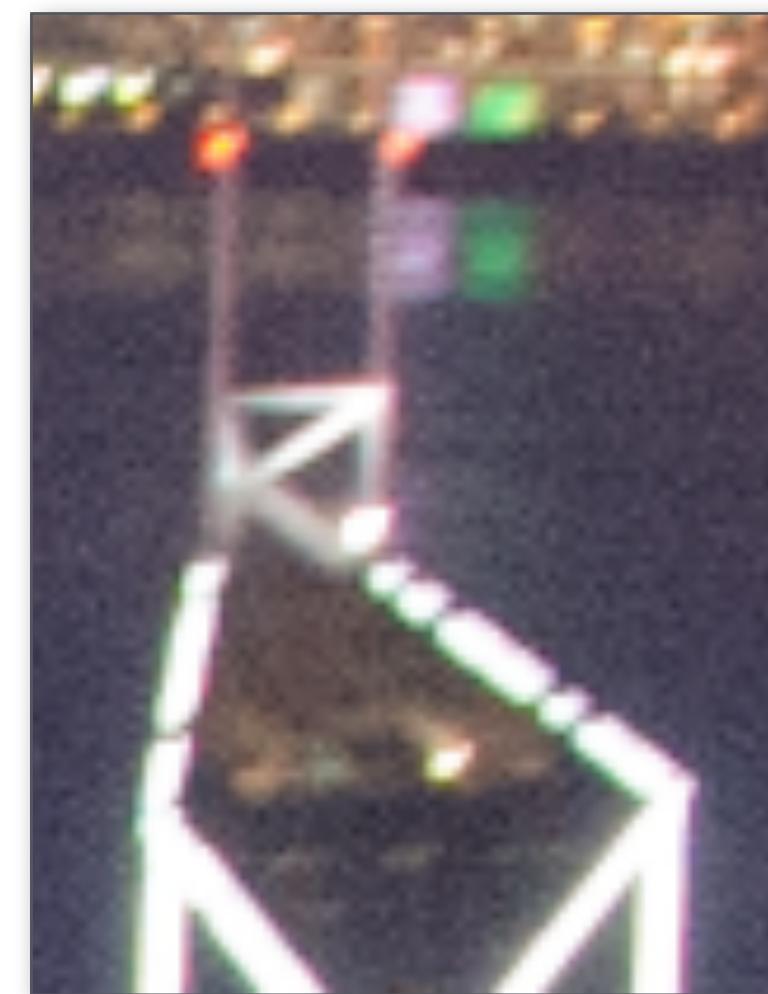
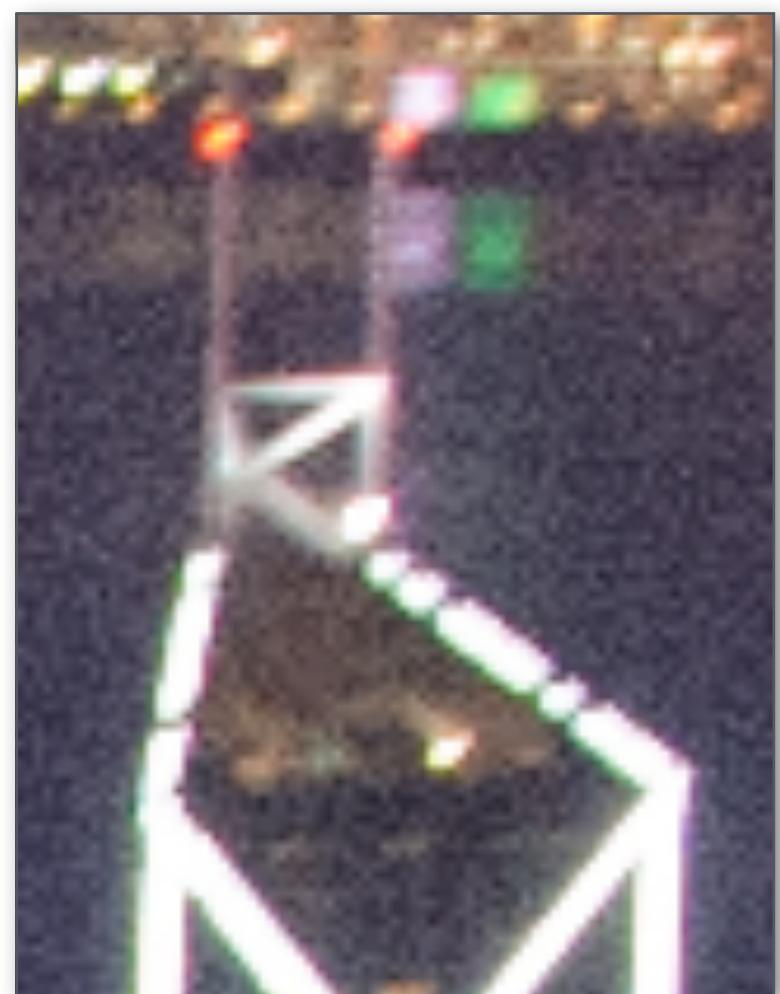
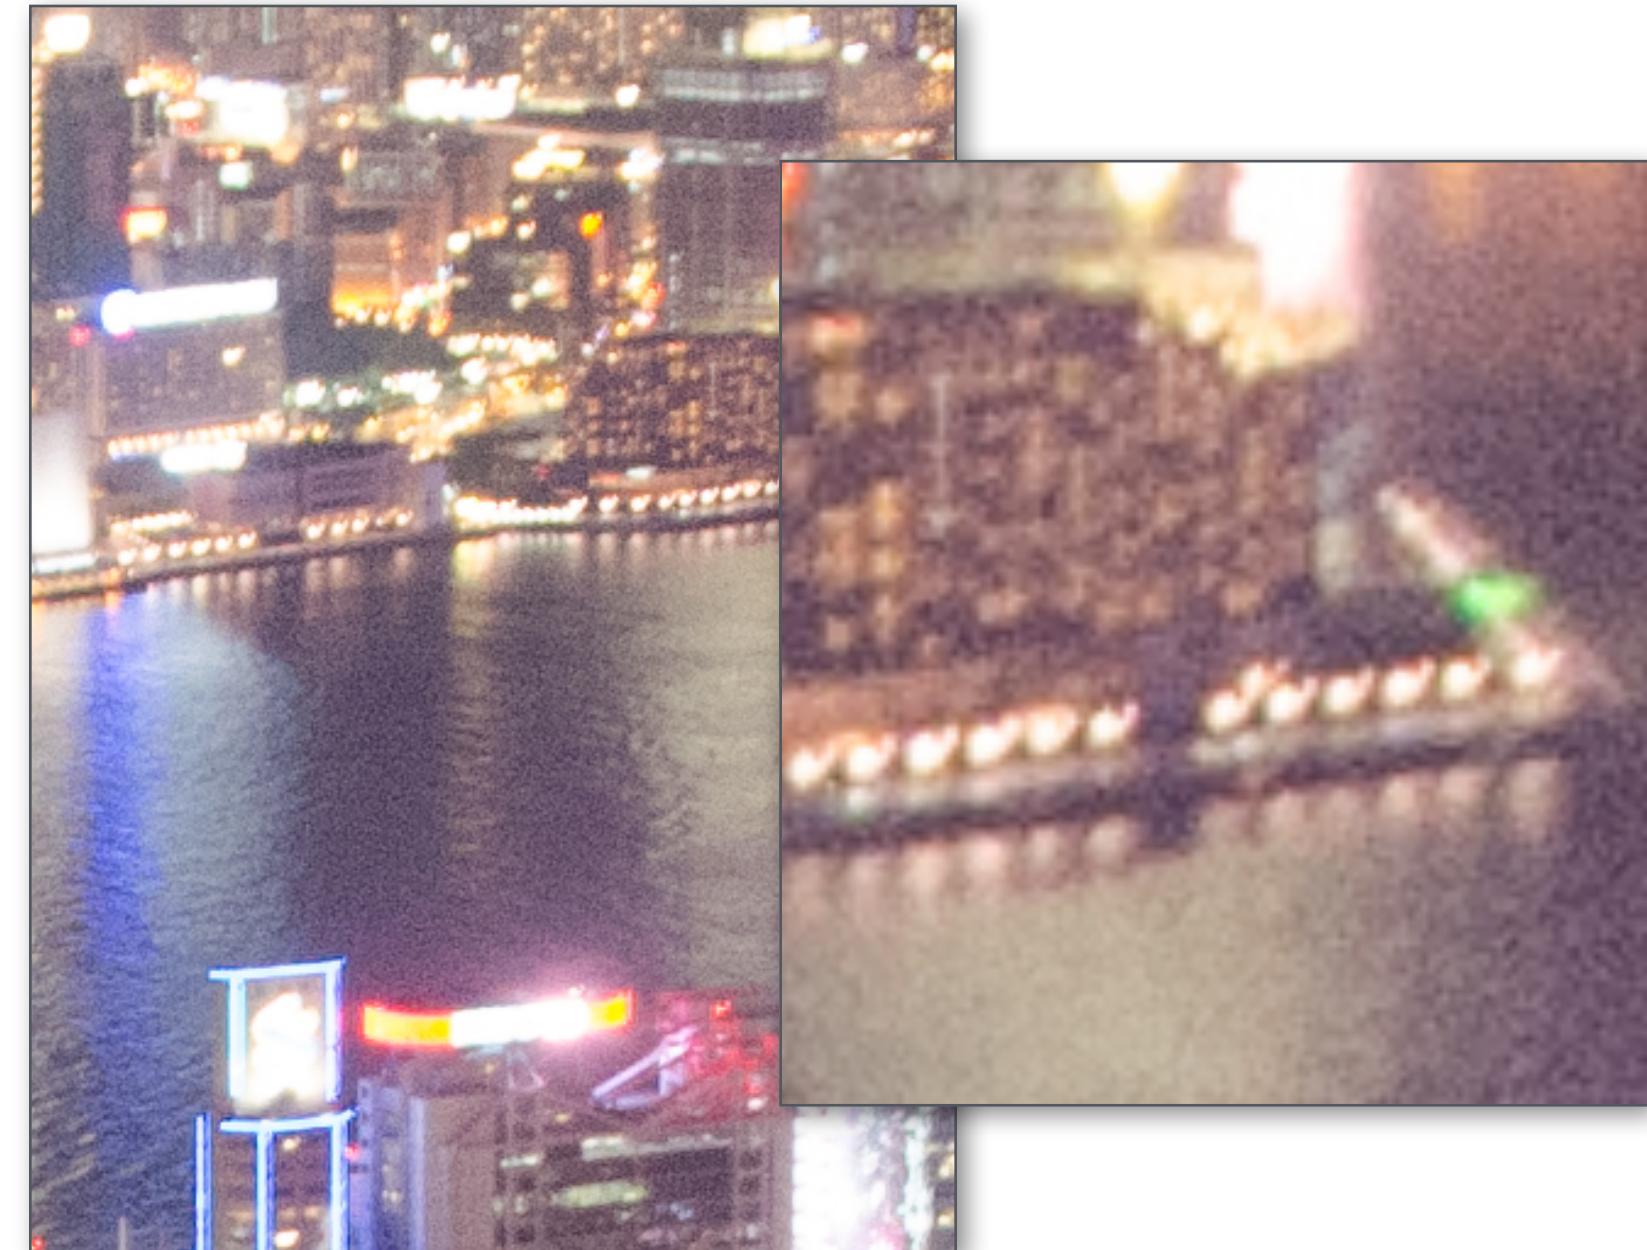


RAW data from sensor

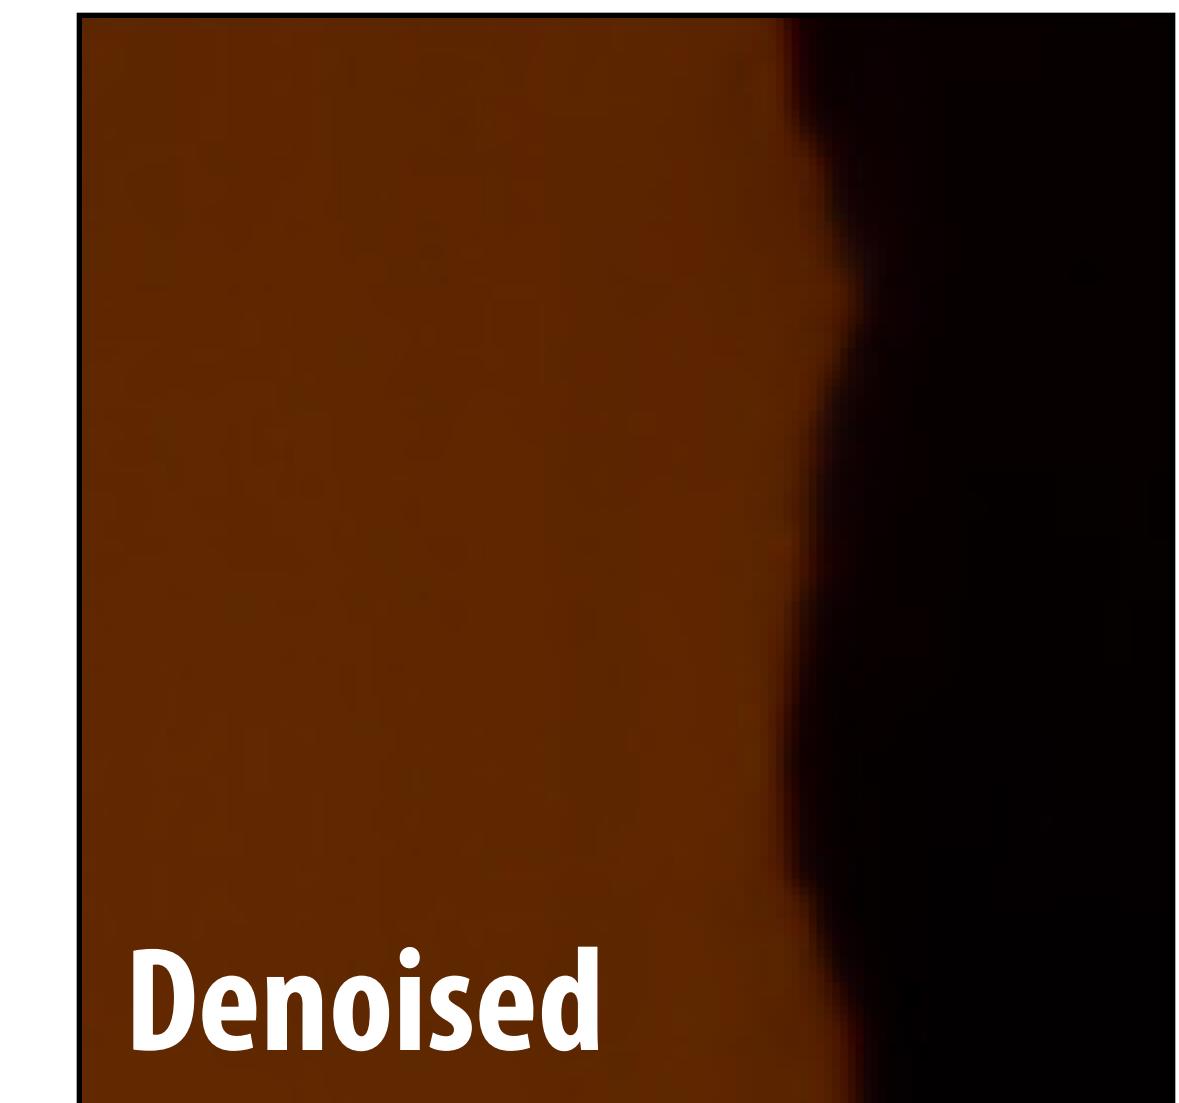
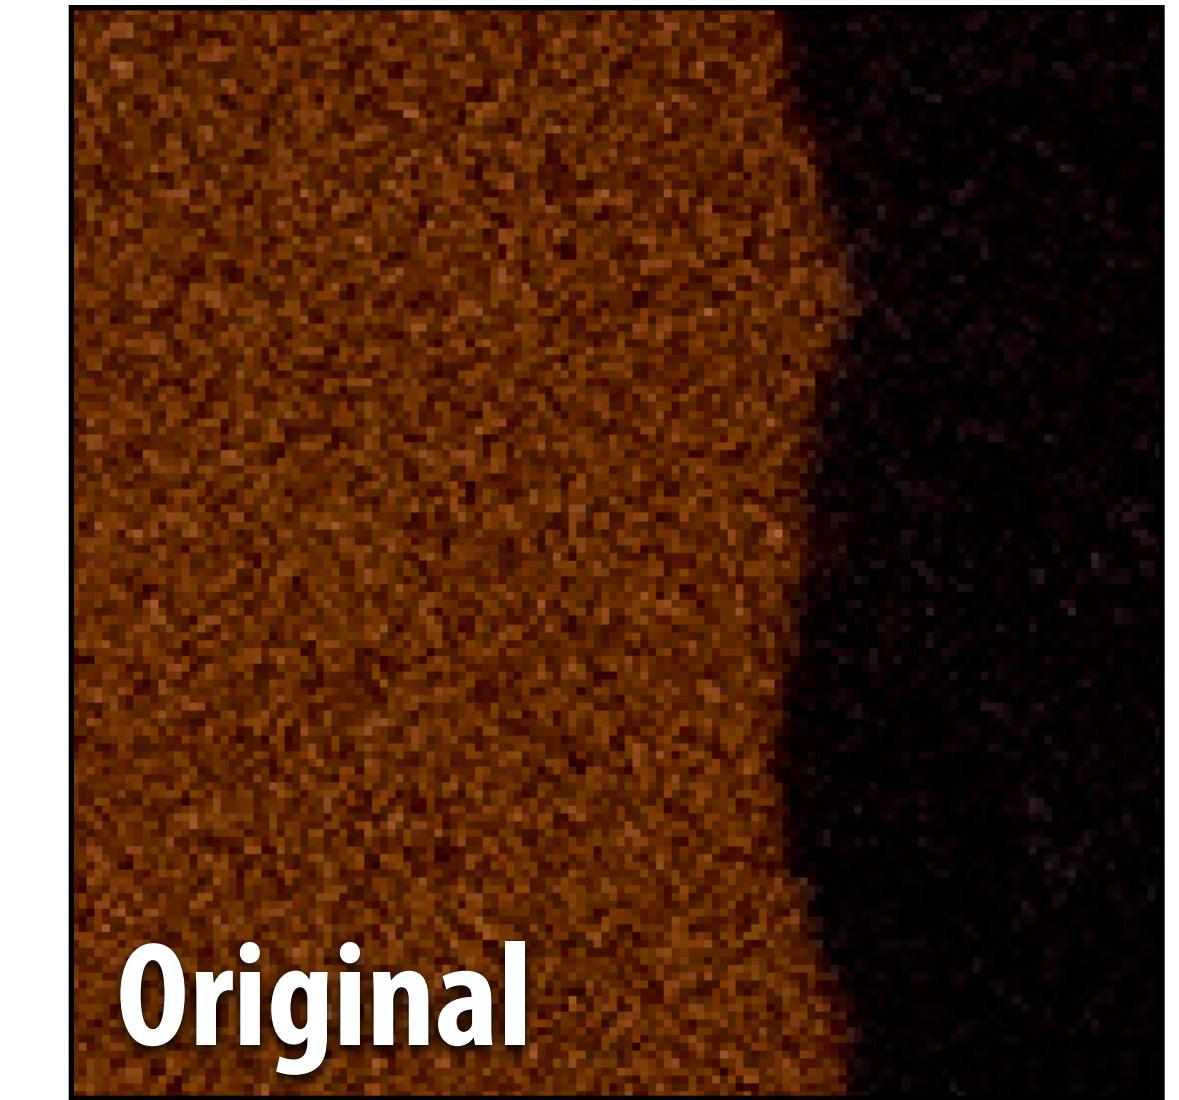
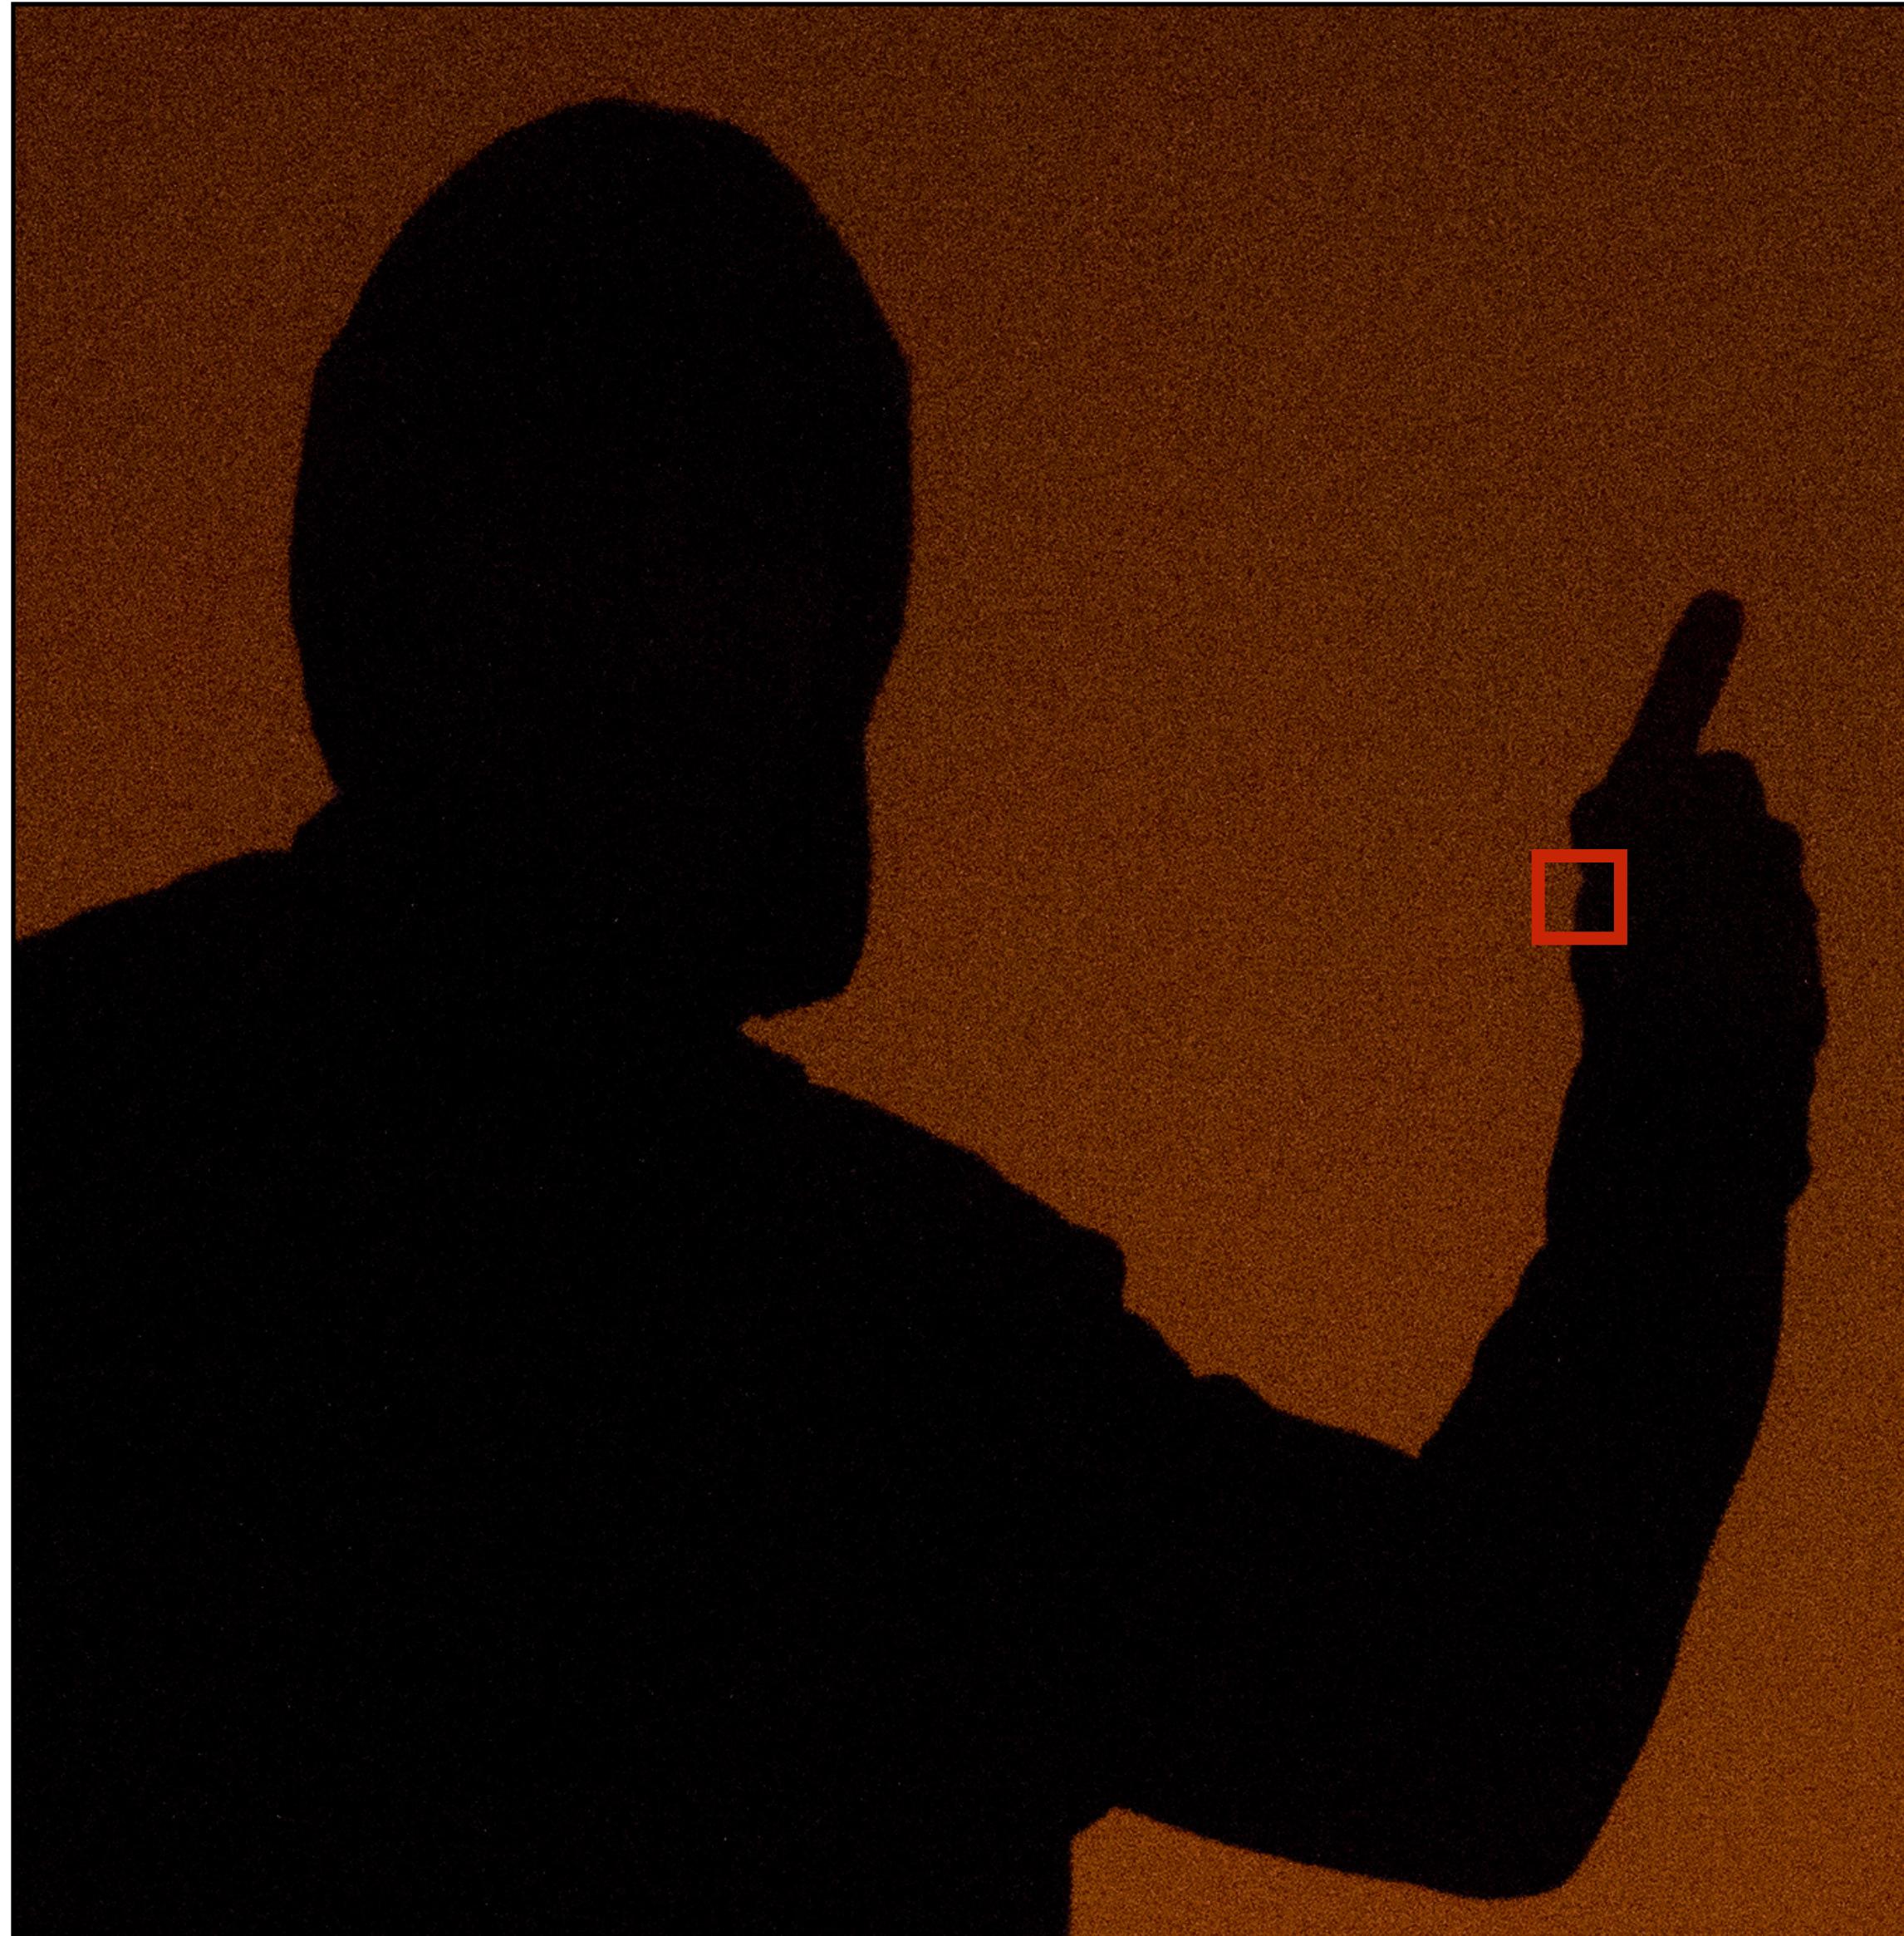


Demosaiced

Denoising: effect of downsizing on image noise

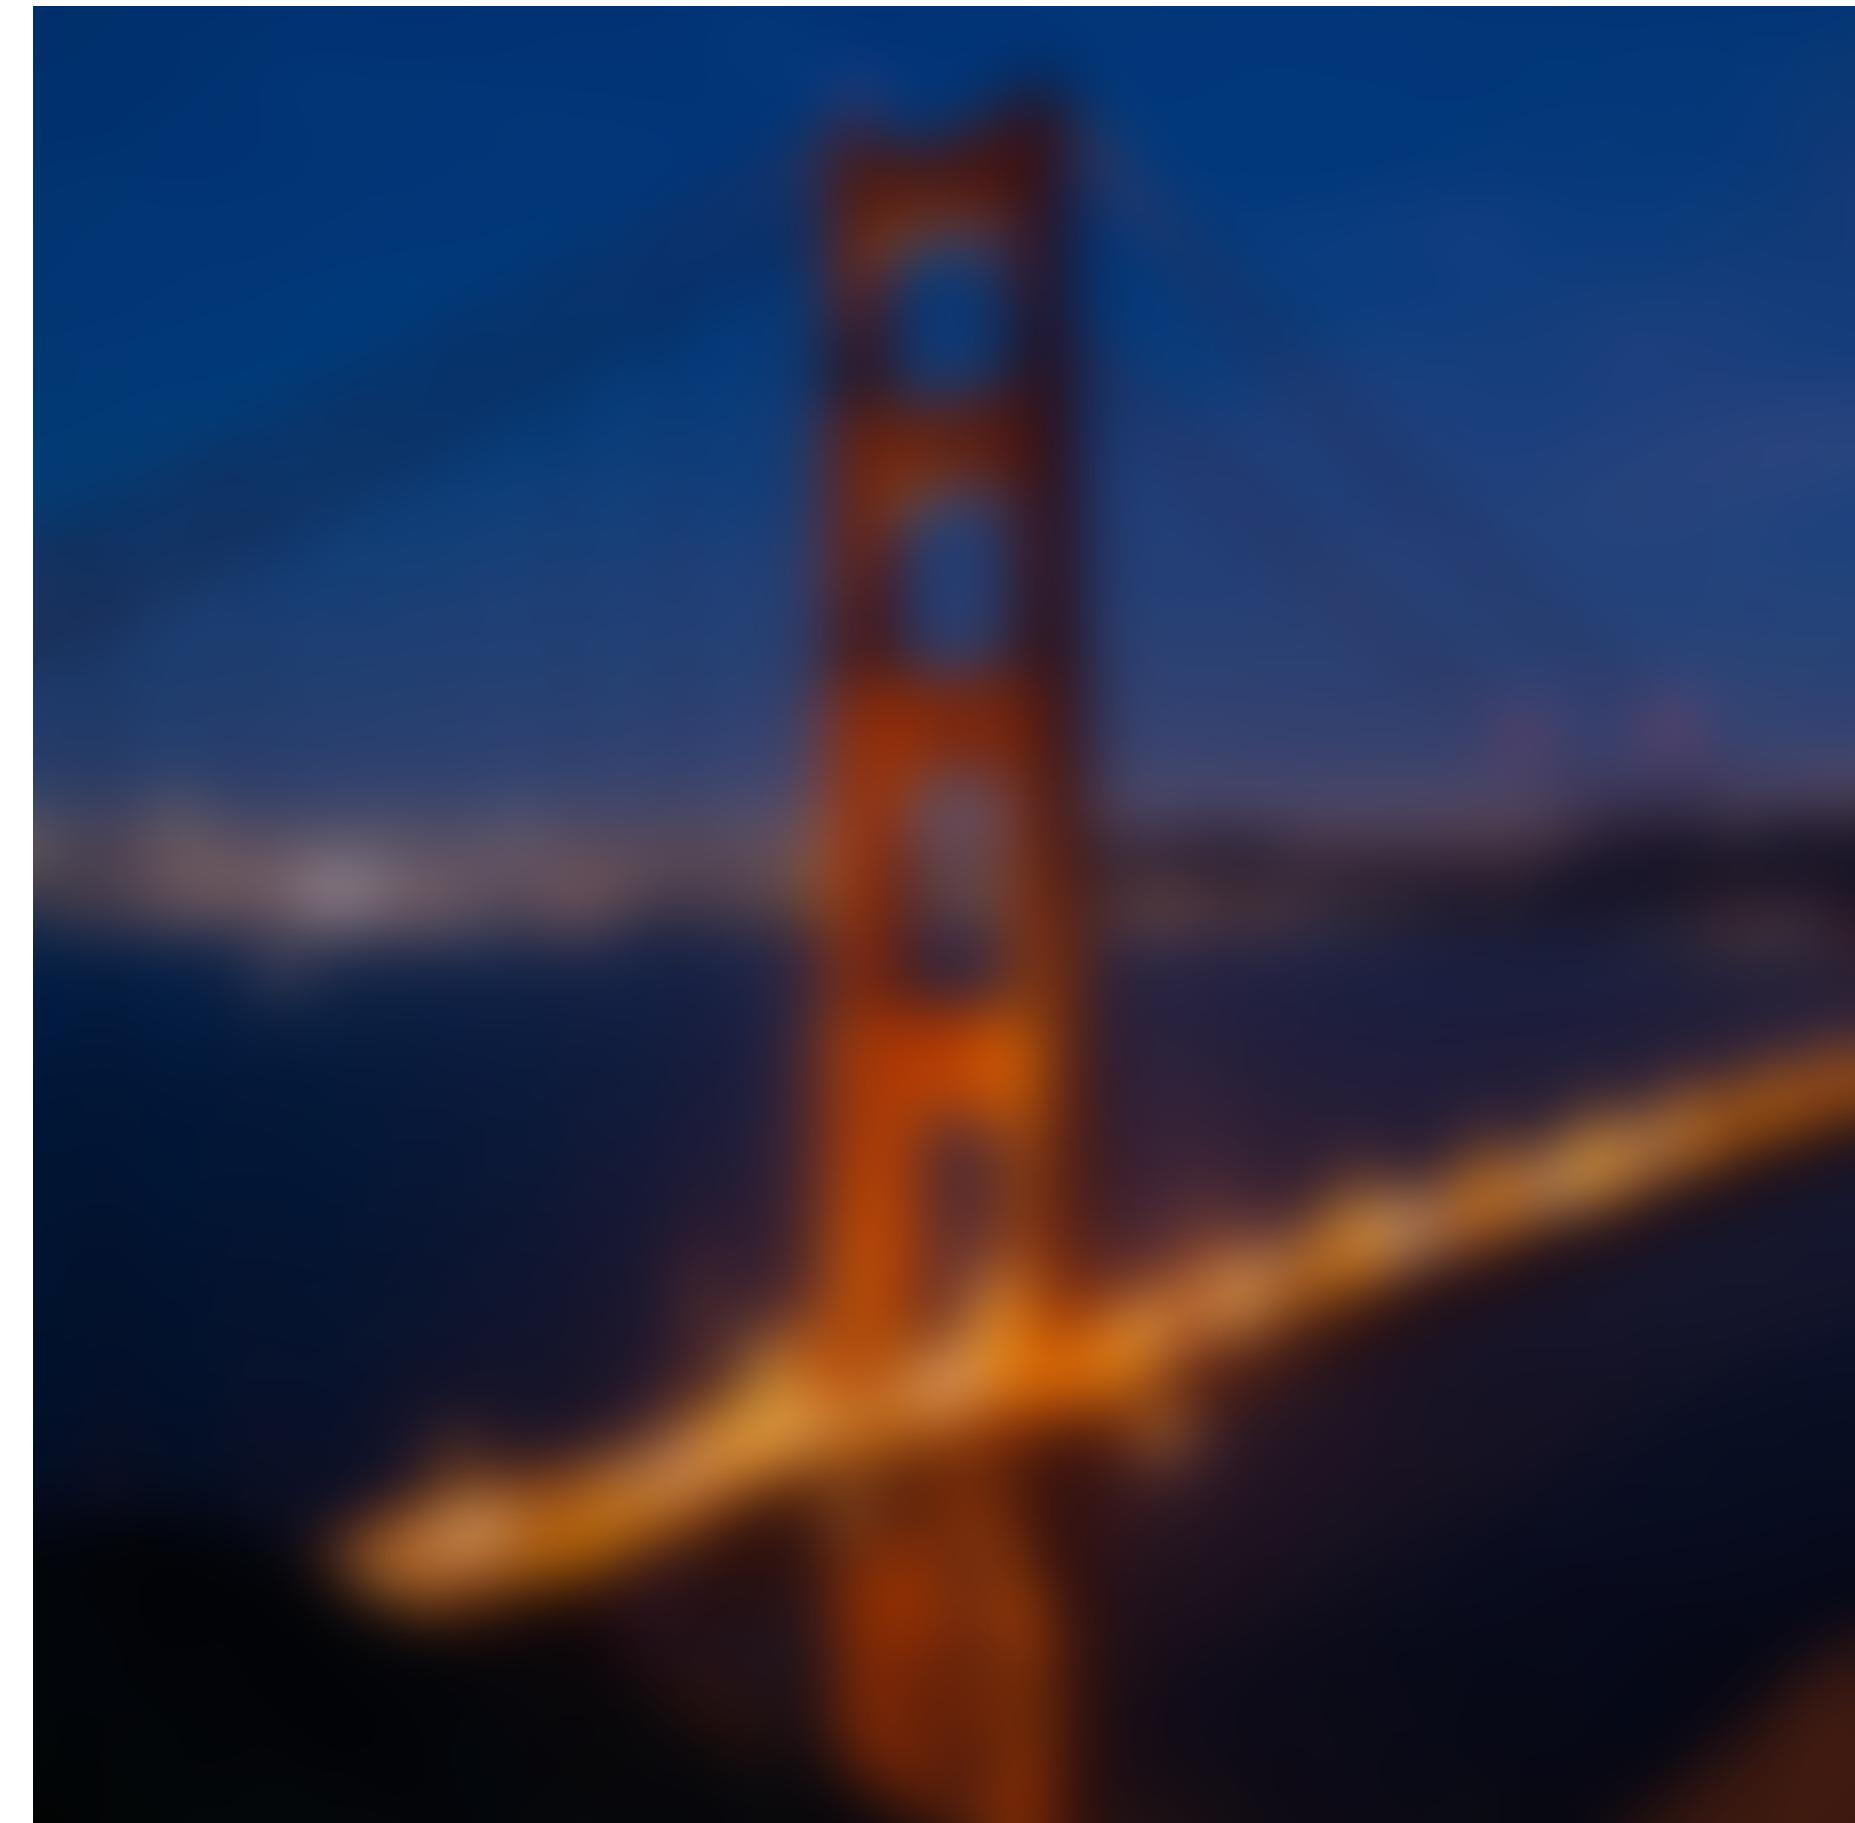


Denoising



Aside: image processing basics

Example image processing operations



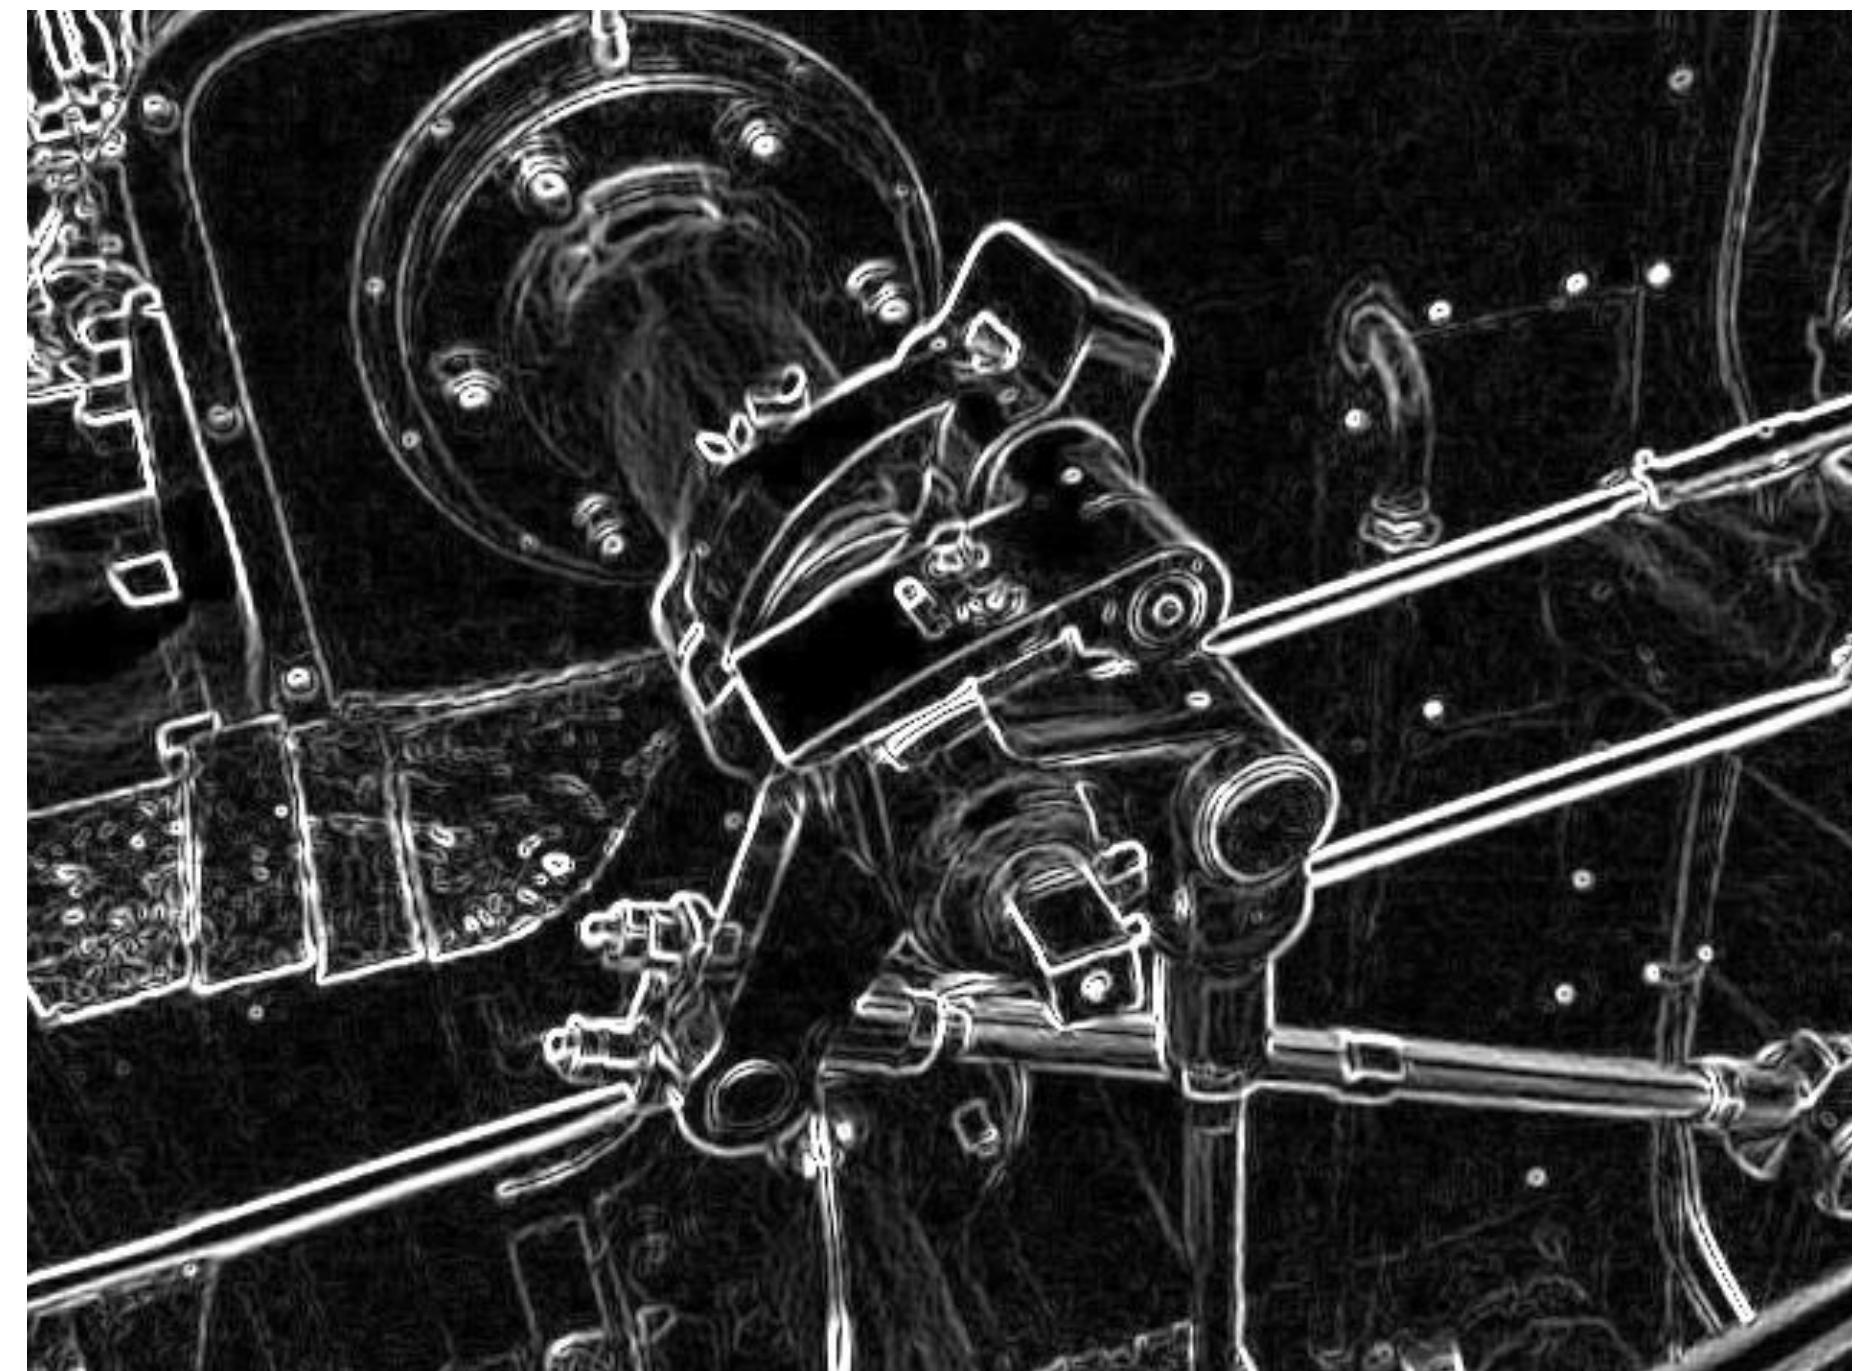
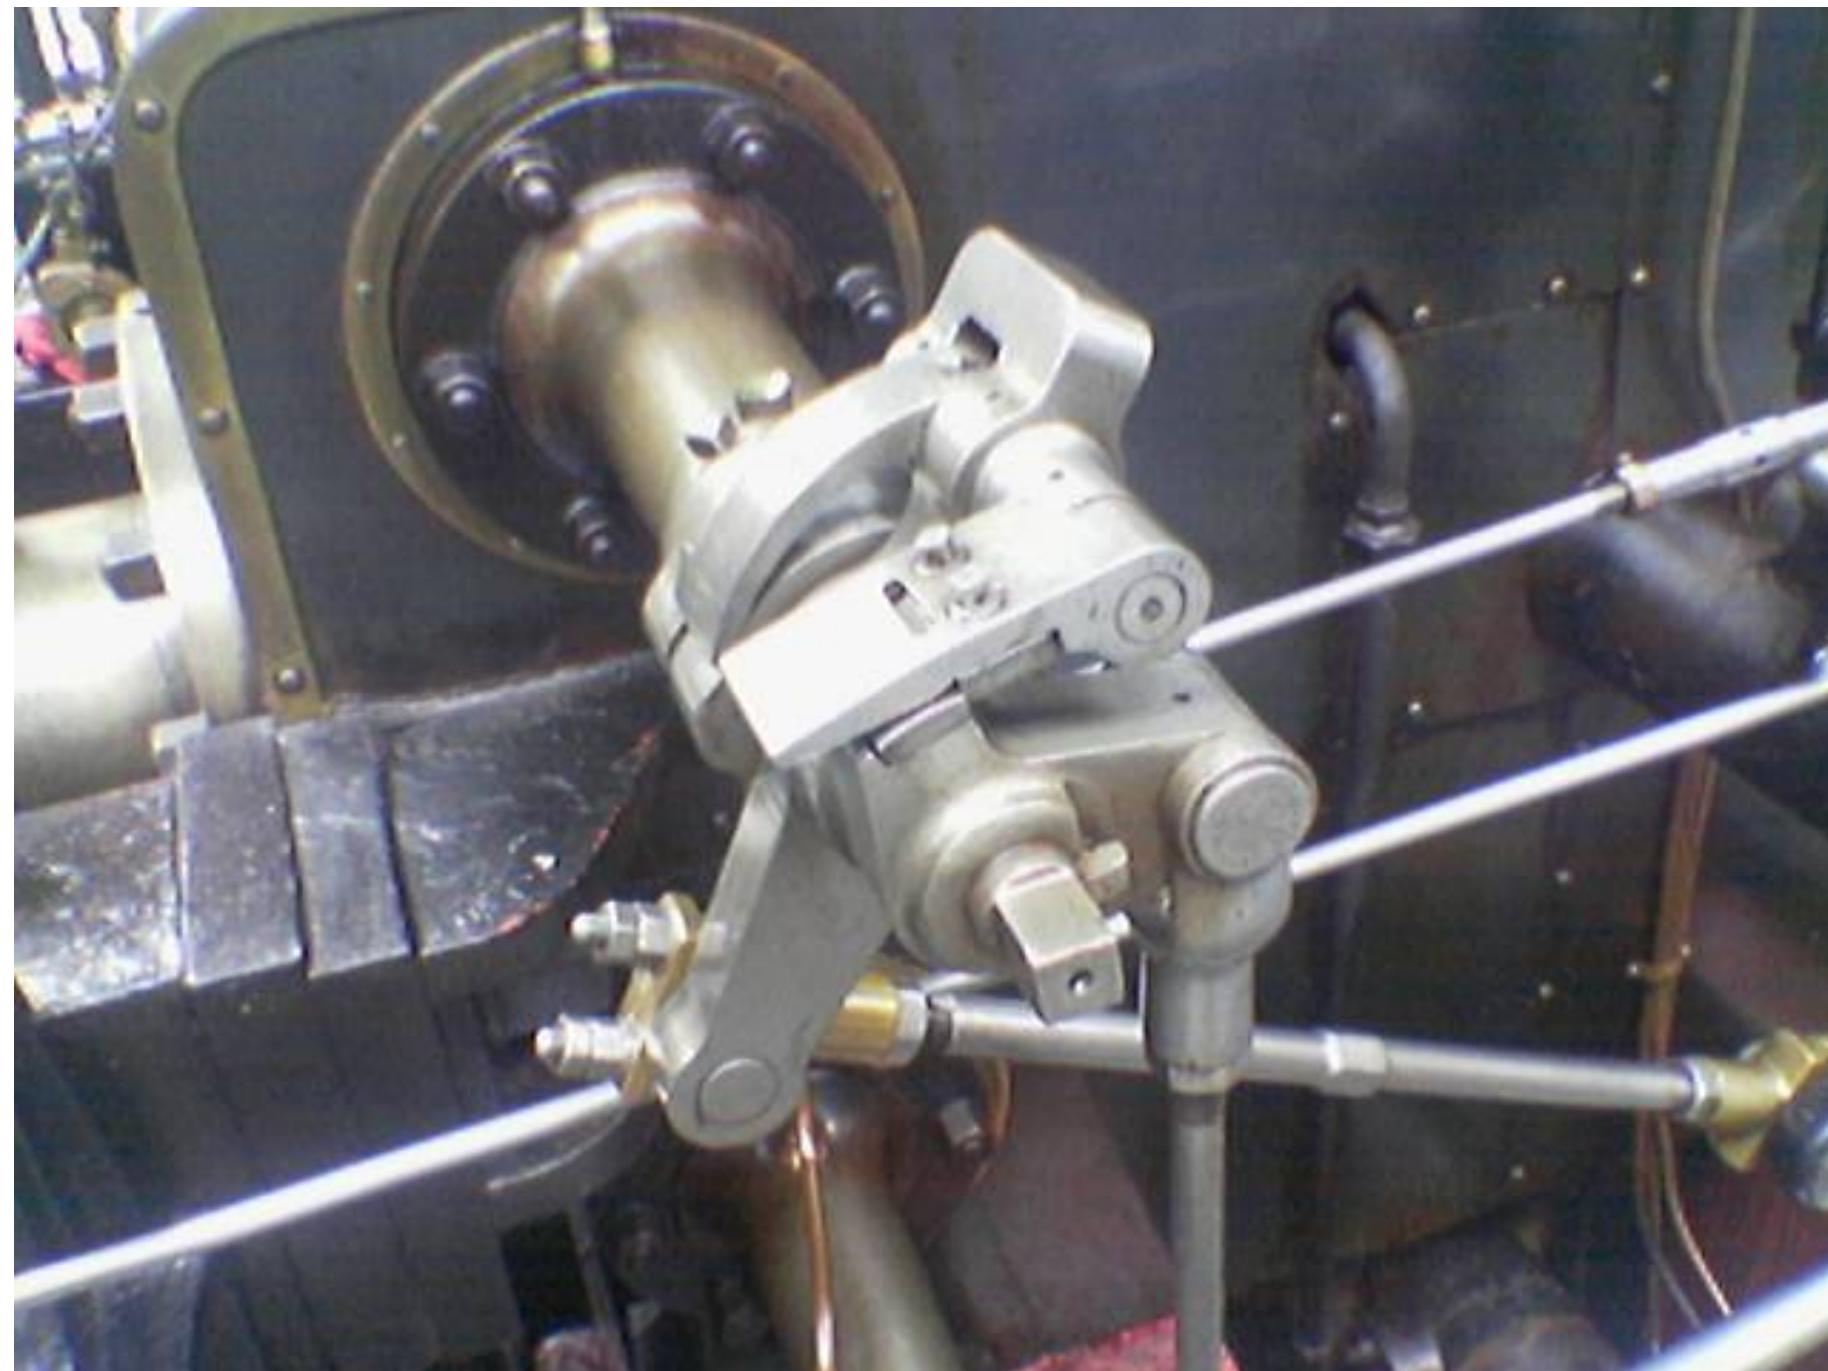
Blur

Example image processing operations



Sharpen

Edge detection



A “smarter” blur (doesn’t blur over edges)



Review: convolution

$$(f * g)(x) = \int_{-\infty}^{\infty} f(y)g(x - y)dy$$

output signal

filter

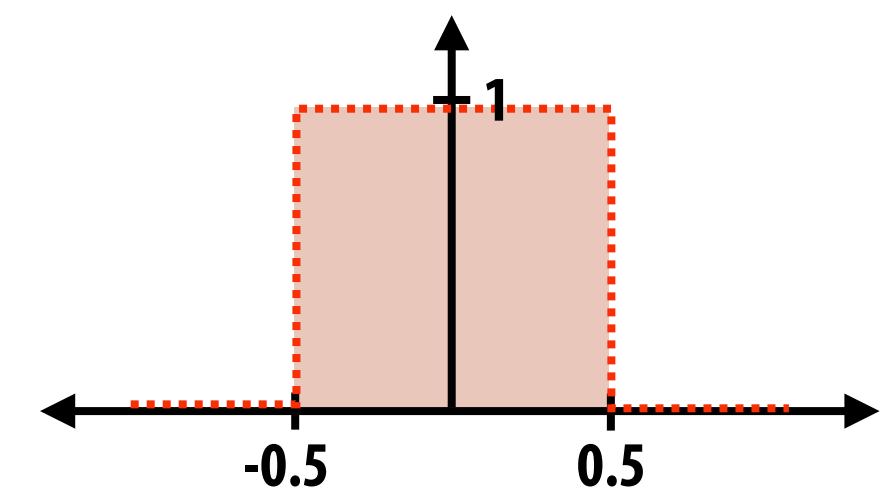
input signal

It may be helpful to consider the effect of convolution with the simple unit-area “box” function:

$$f(x) = \begin{cases} 1 & |x| \leq 0.5 \\ 0 & \text{otherwise} \end{cases}$$

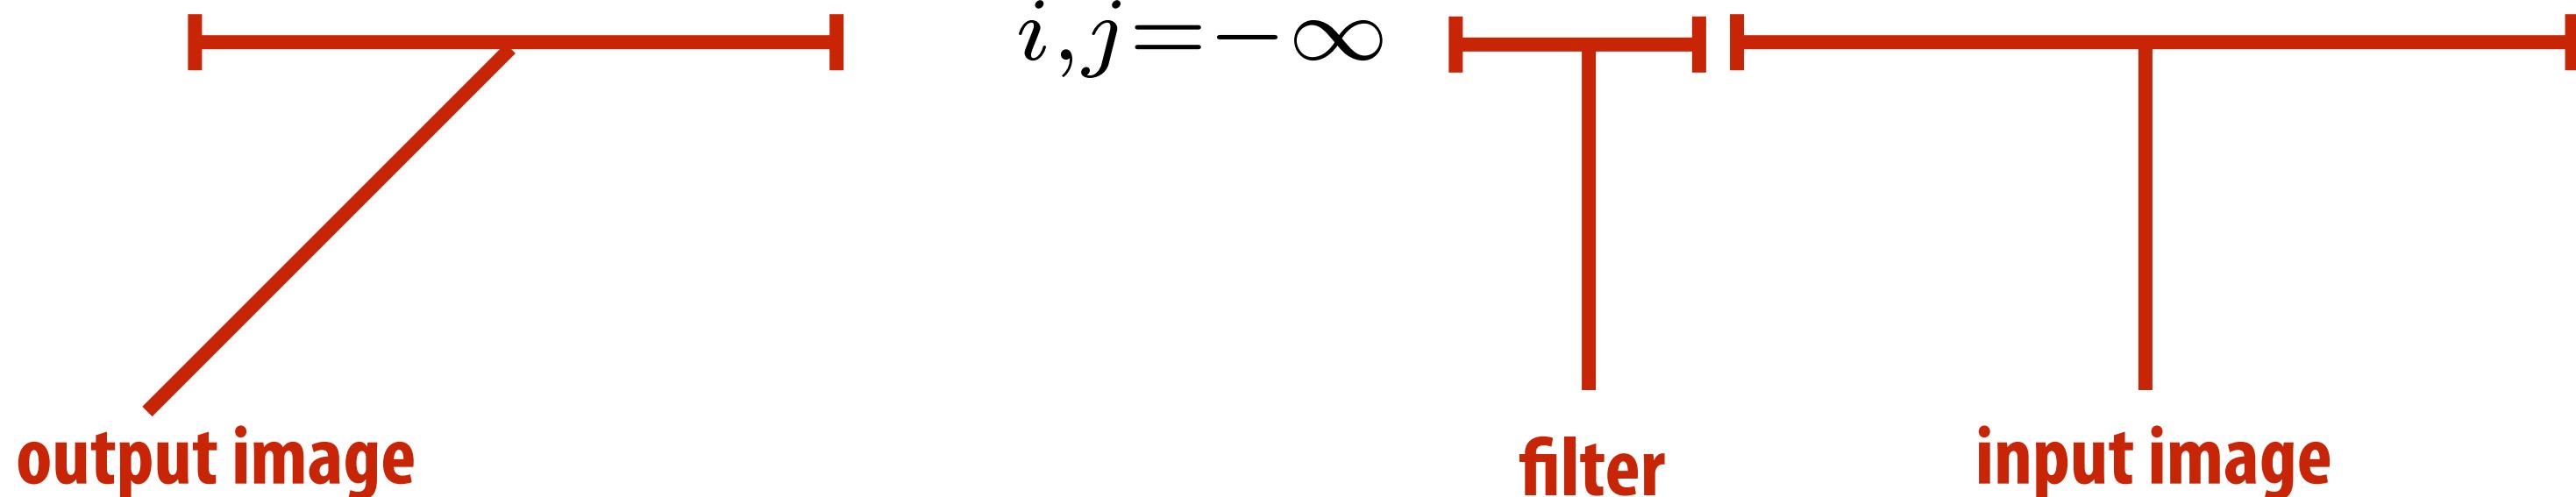
$$(f * g)(x) = \int_{-0.5}^{0.5} g(x - y)dy$$

$f * g$ is a “blurred” version of g



Discrete 2D convolution

$$(f * g)(x, y) = \sum_{i=0}^{\infty} f(i, j) I(x - i, y - j)$$



Consider $f(i, j)$ that is nonzero only when: $-1 \leq i, j \leq 1$

Then:

$$(f * g)(x, y) = \sum_{i,j=-1}^1 f(i, j) I(x - i, y - j)$$

And we can represent $f(i,j)$ as a 3×3 matrix of values where:

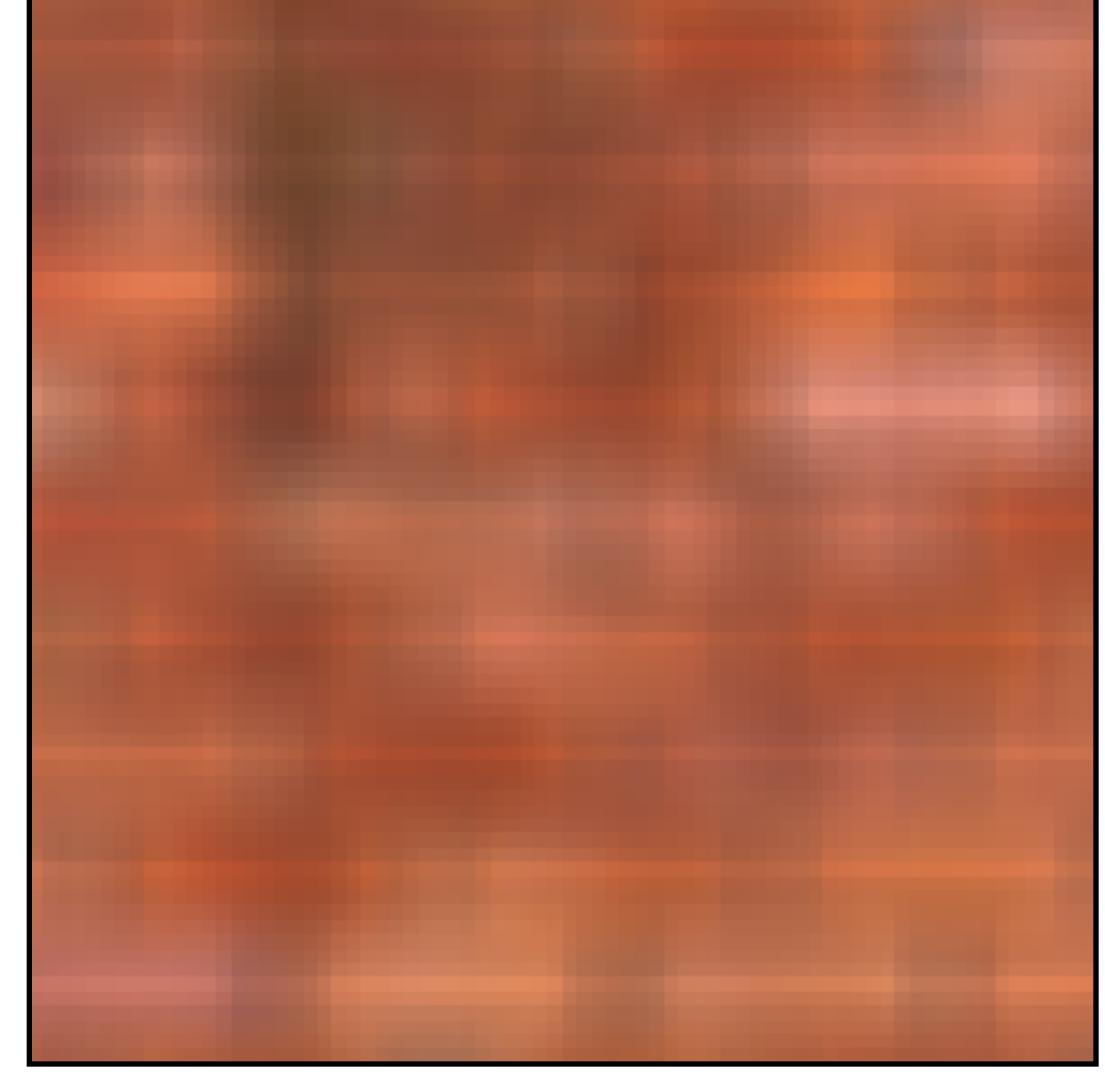
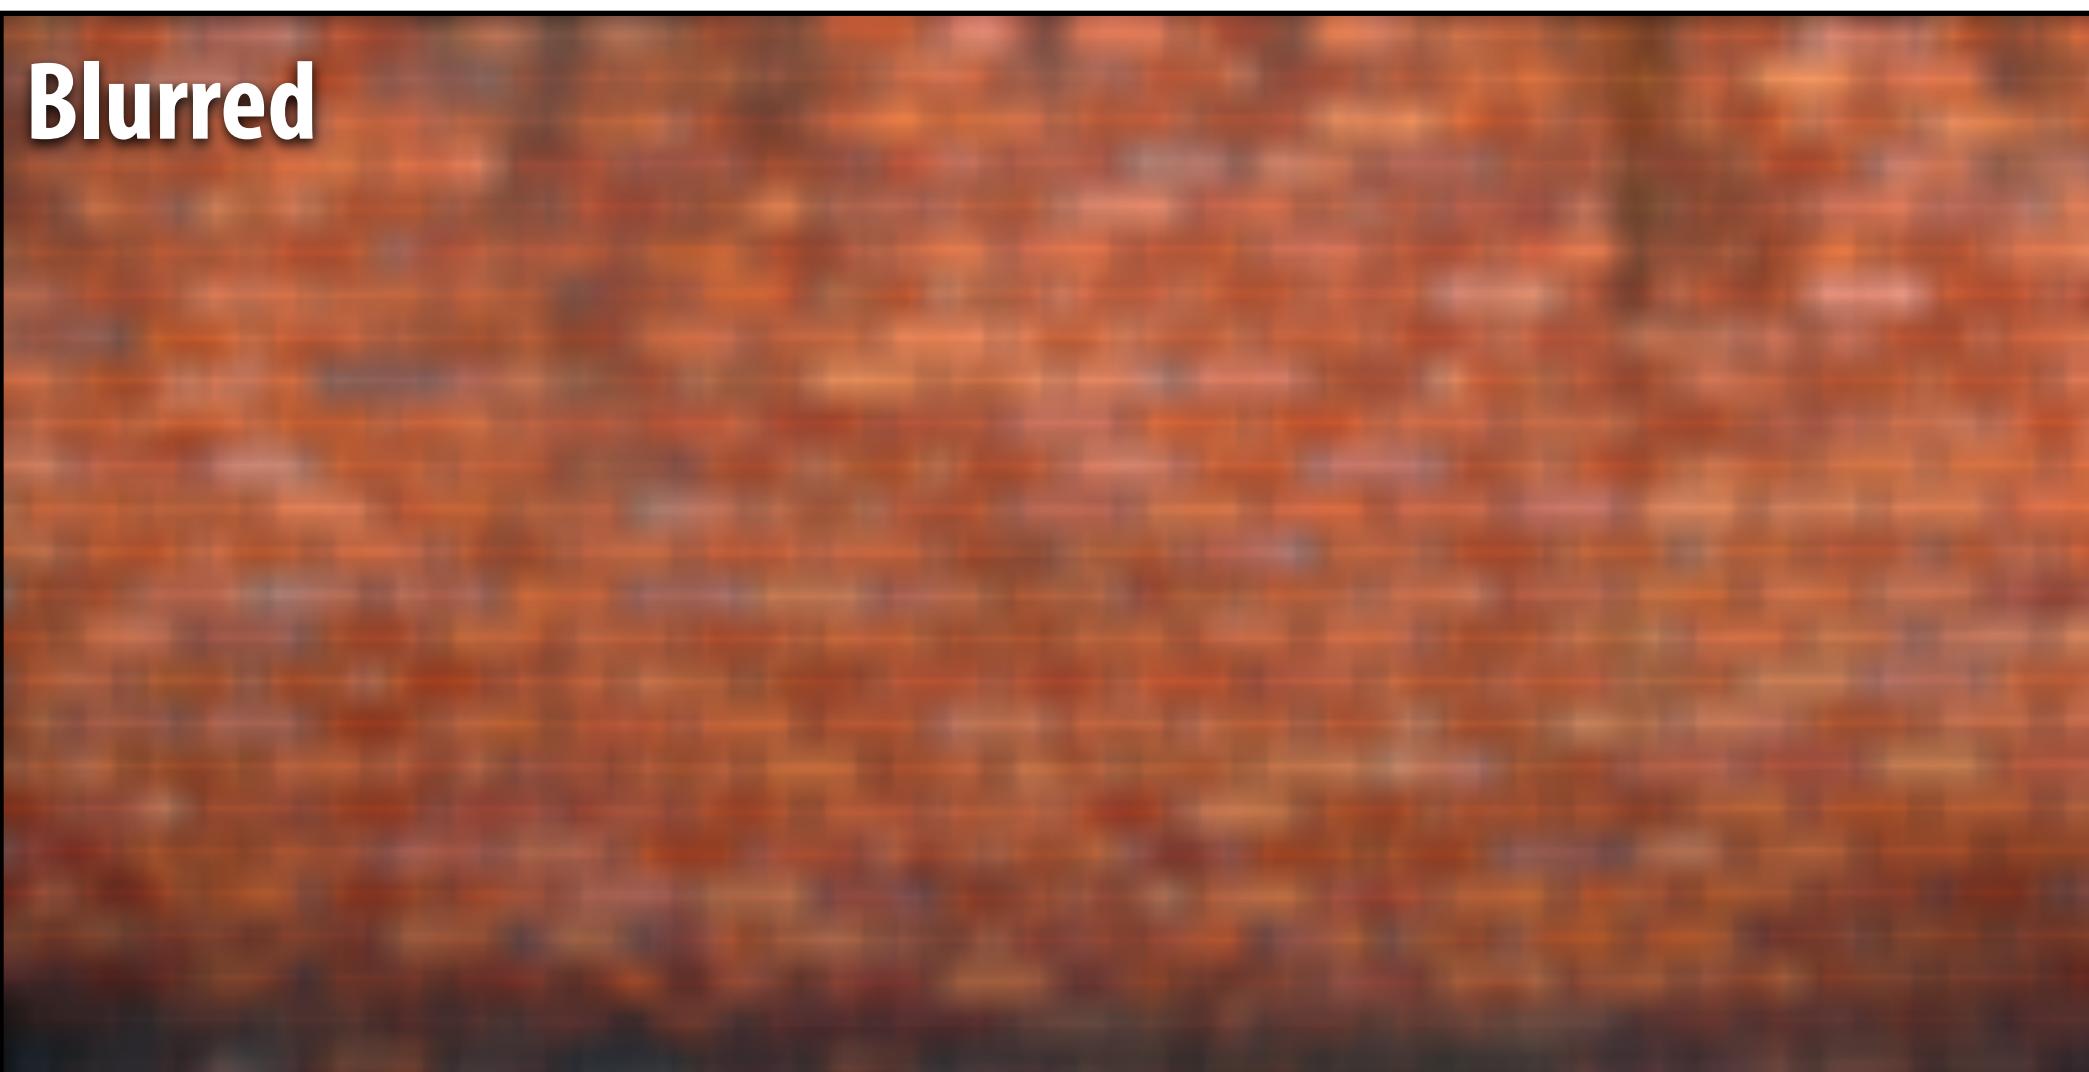
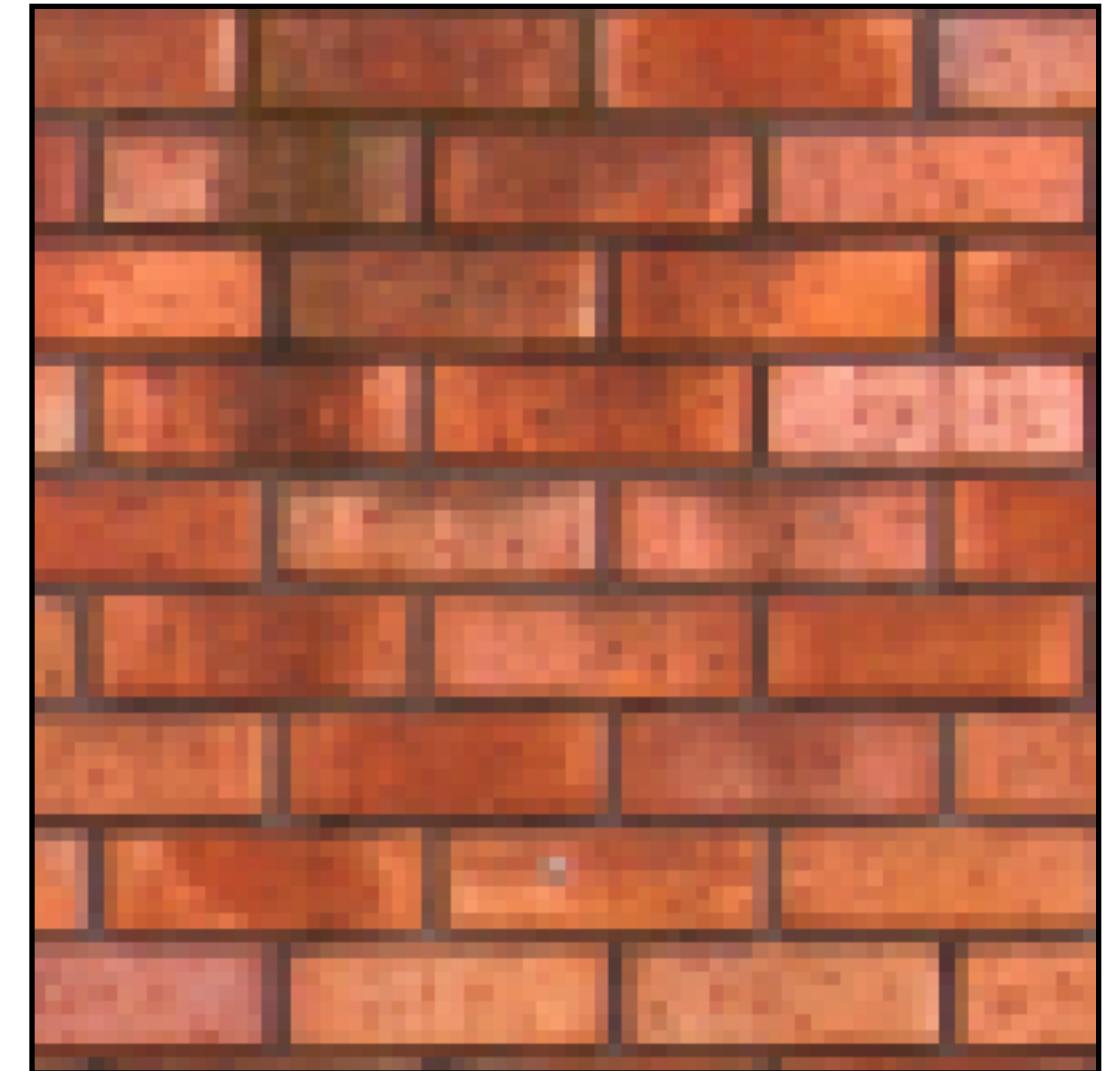
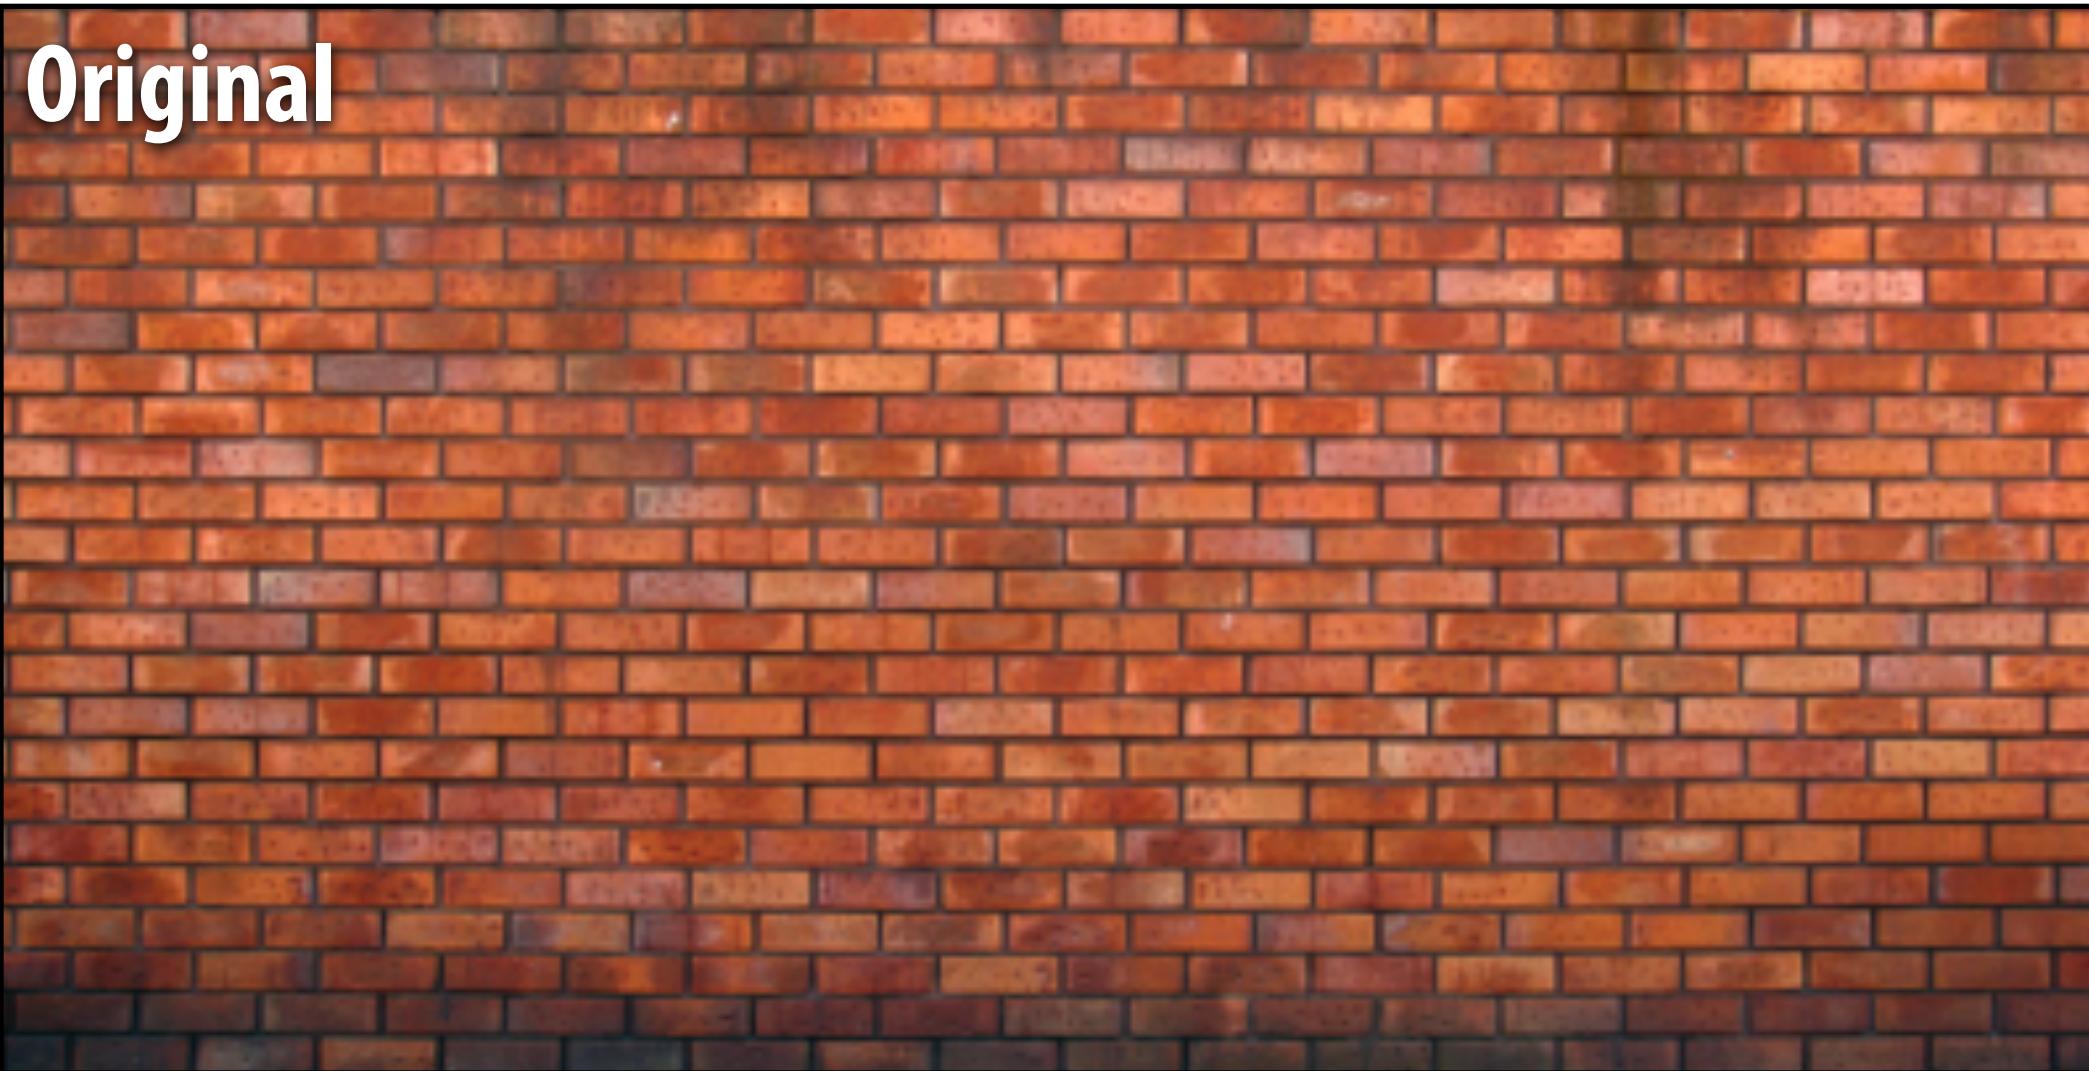
$$f(i, j) = \mathbf{F}_{i, j}$$

(often called: “filter weights”, “filter kernel”)

Simple 3x3 box blur

```
float input[(WIDTH+2) * (HEIGHT+2)];  
float output[WIDTH * HEIGHT]; ← For now: ignore boundary pixels and  
assume output image is smaller than  
input (makes convolution loop bounds  
much simpler to write)  
  
float weights[] = {1./9, 1./9, 1./9,  
                   1./9, 1./9, 1./9,  
                   1./9, 1./9, 1./9};  
  
for (int j=0; j<HEIGHT; j++) {  
    for (int i=0; i<WIDTH; i++) {  
        float tmp = 0.f; ← For now: ignore boundary pixels and  
assume output image is smaller than  
input (makes convolution loop bounds  
much simpler to write)  
        for (int jj=0; jj<3; jj++)  
            for (int ii=0; ii<3; ii++)  
                tmp += input[(j+jj)*(WIDTH+2) + (i+ii)] * weights[jj*3 + ii];  
        output[j*WIDTH + i] = tmp;  
    }  
}
```

7x7 box blur



Gaussian blur

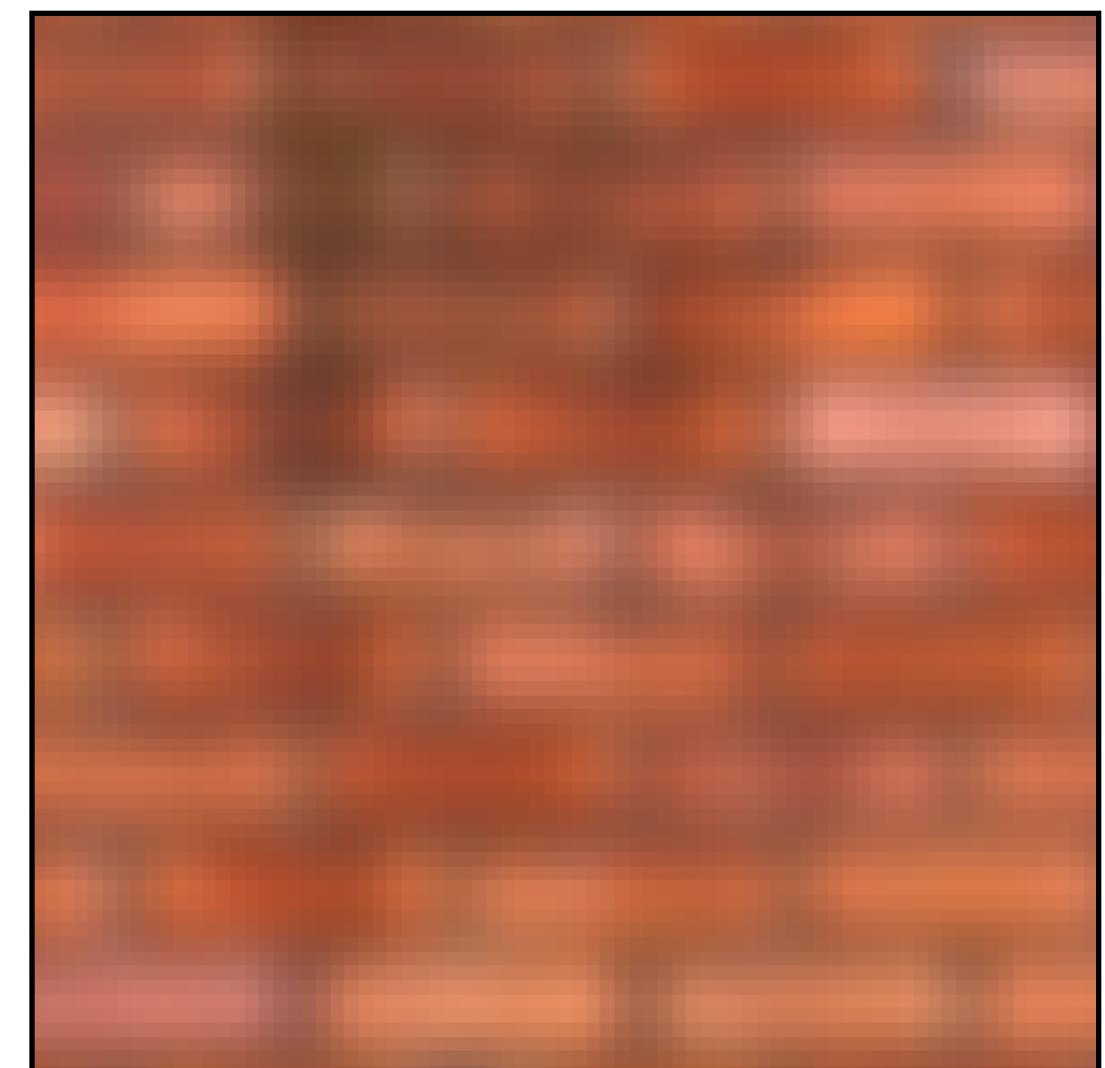
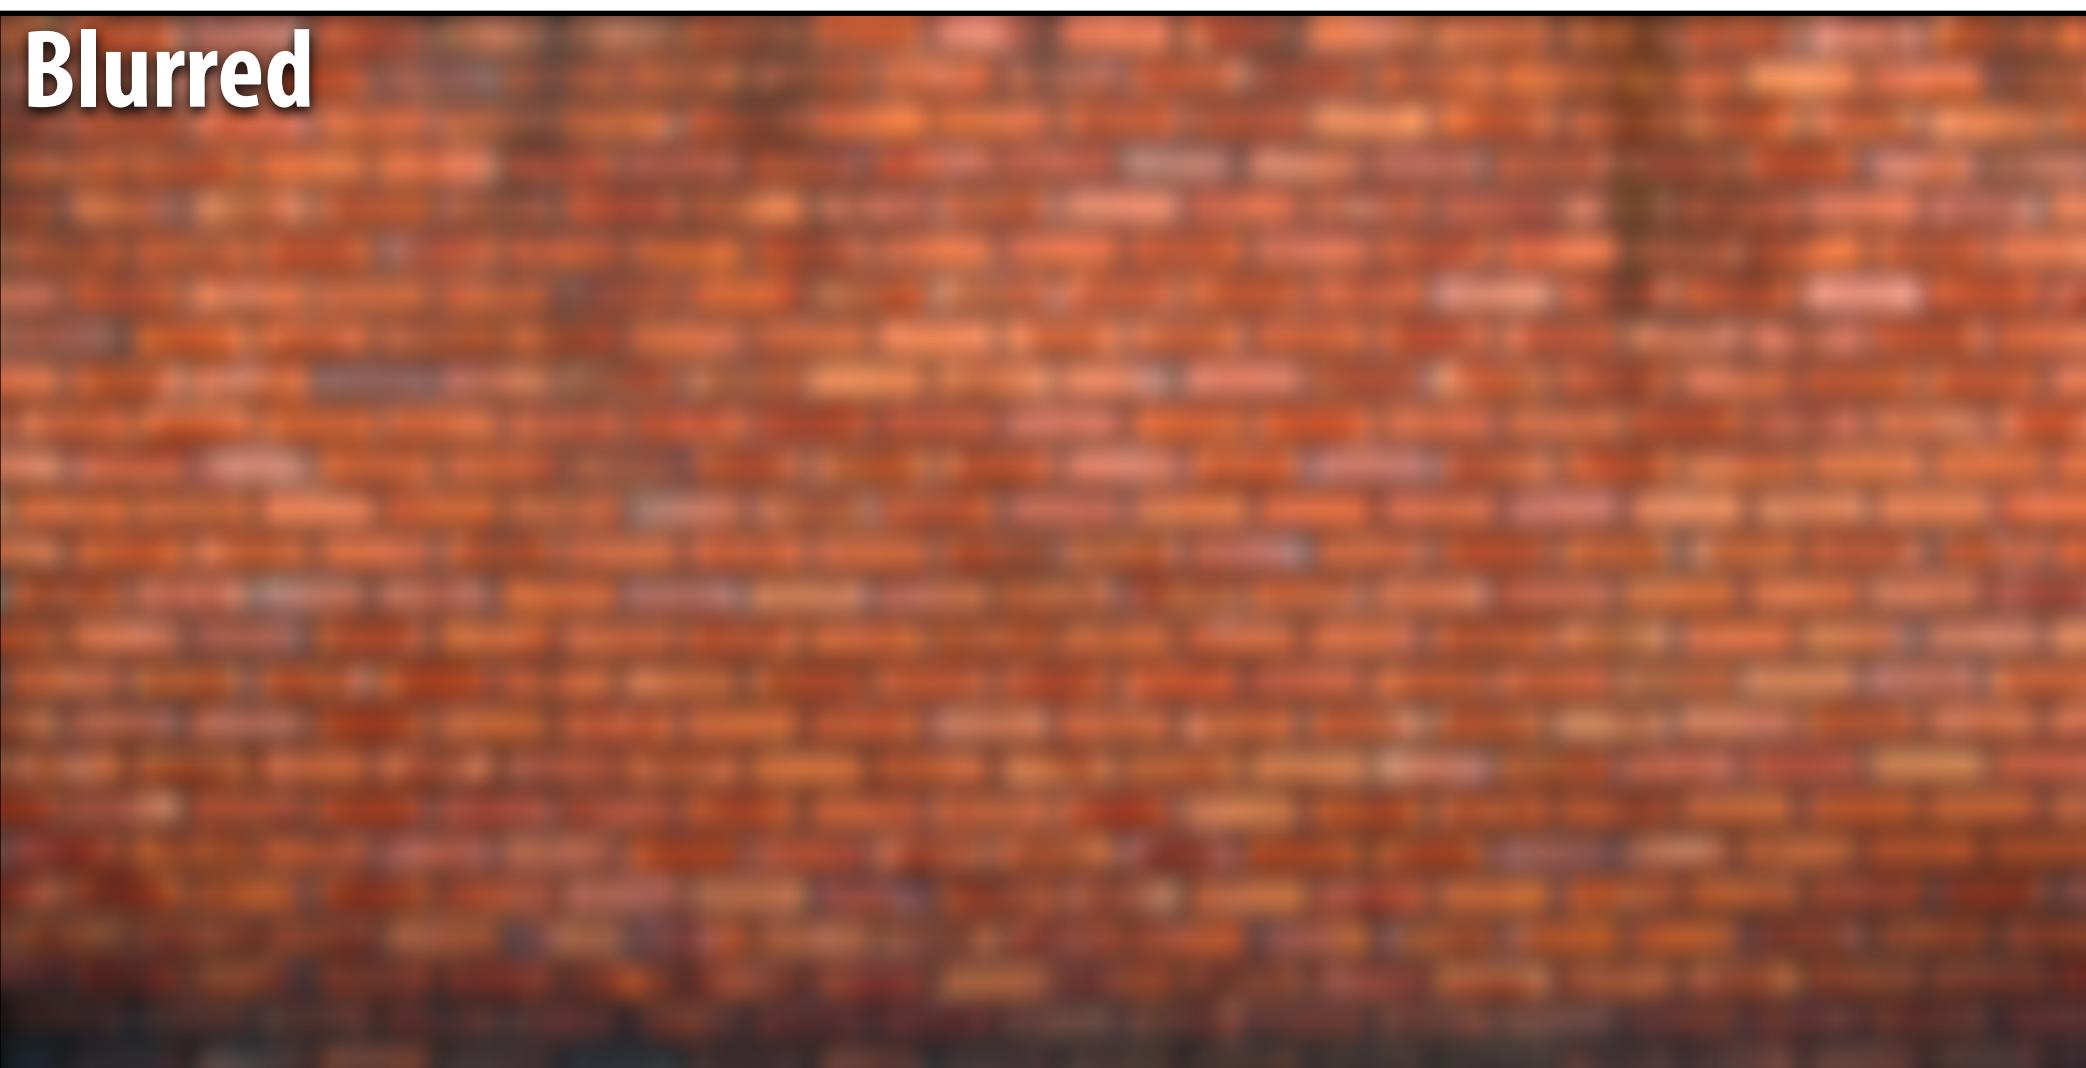
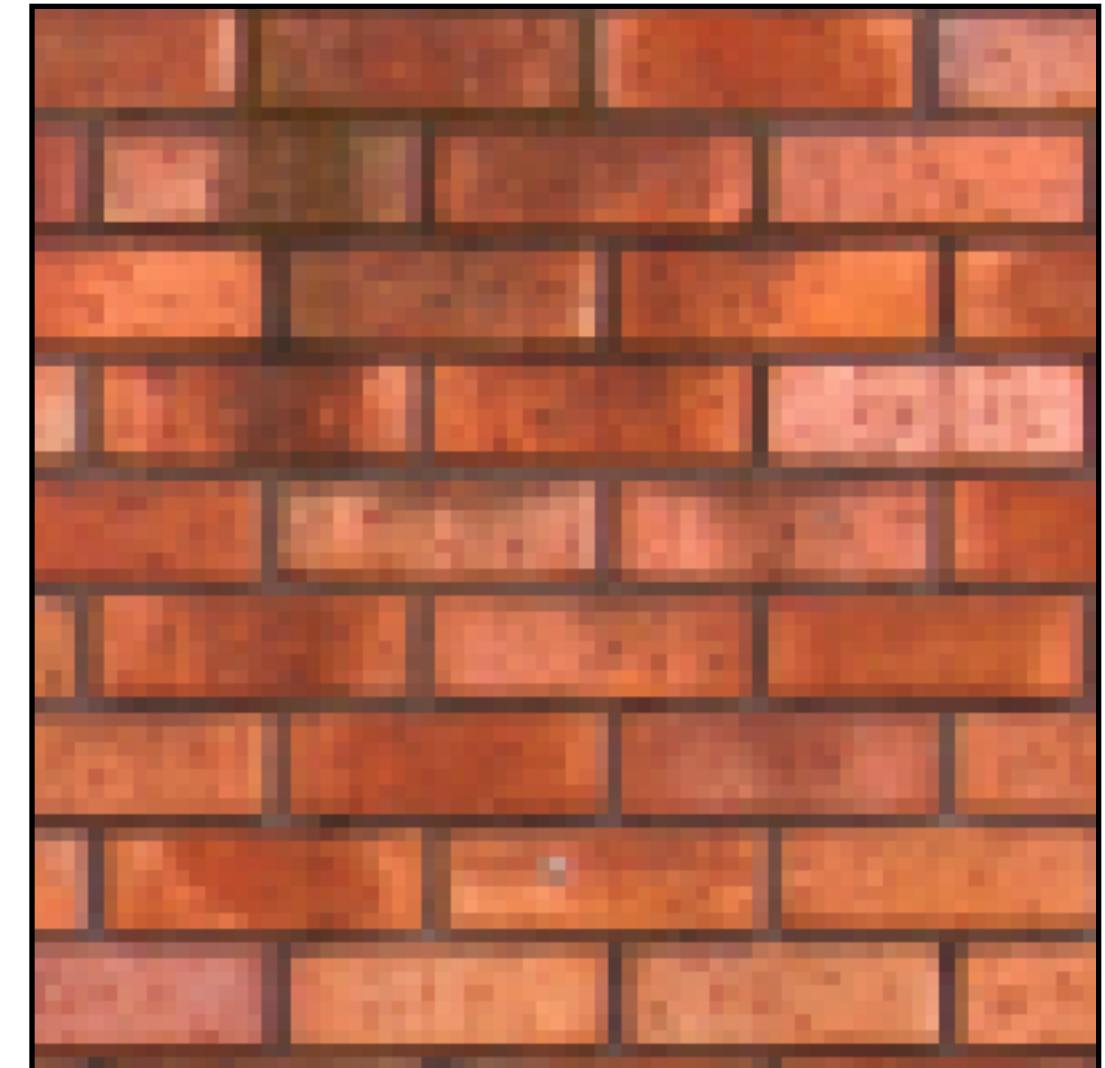
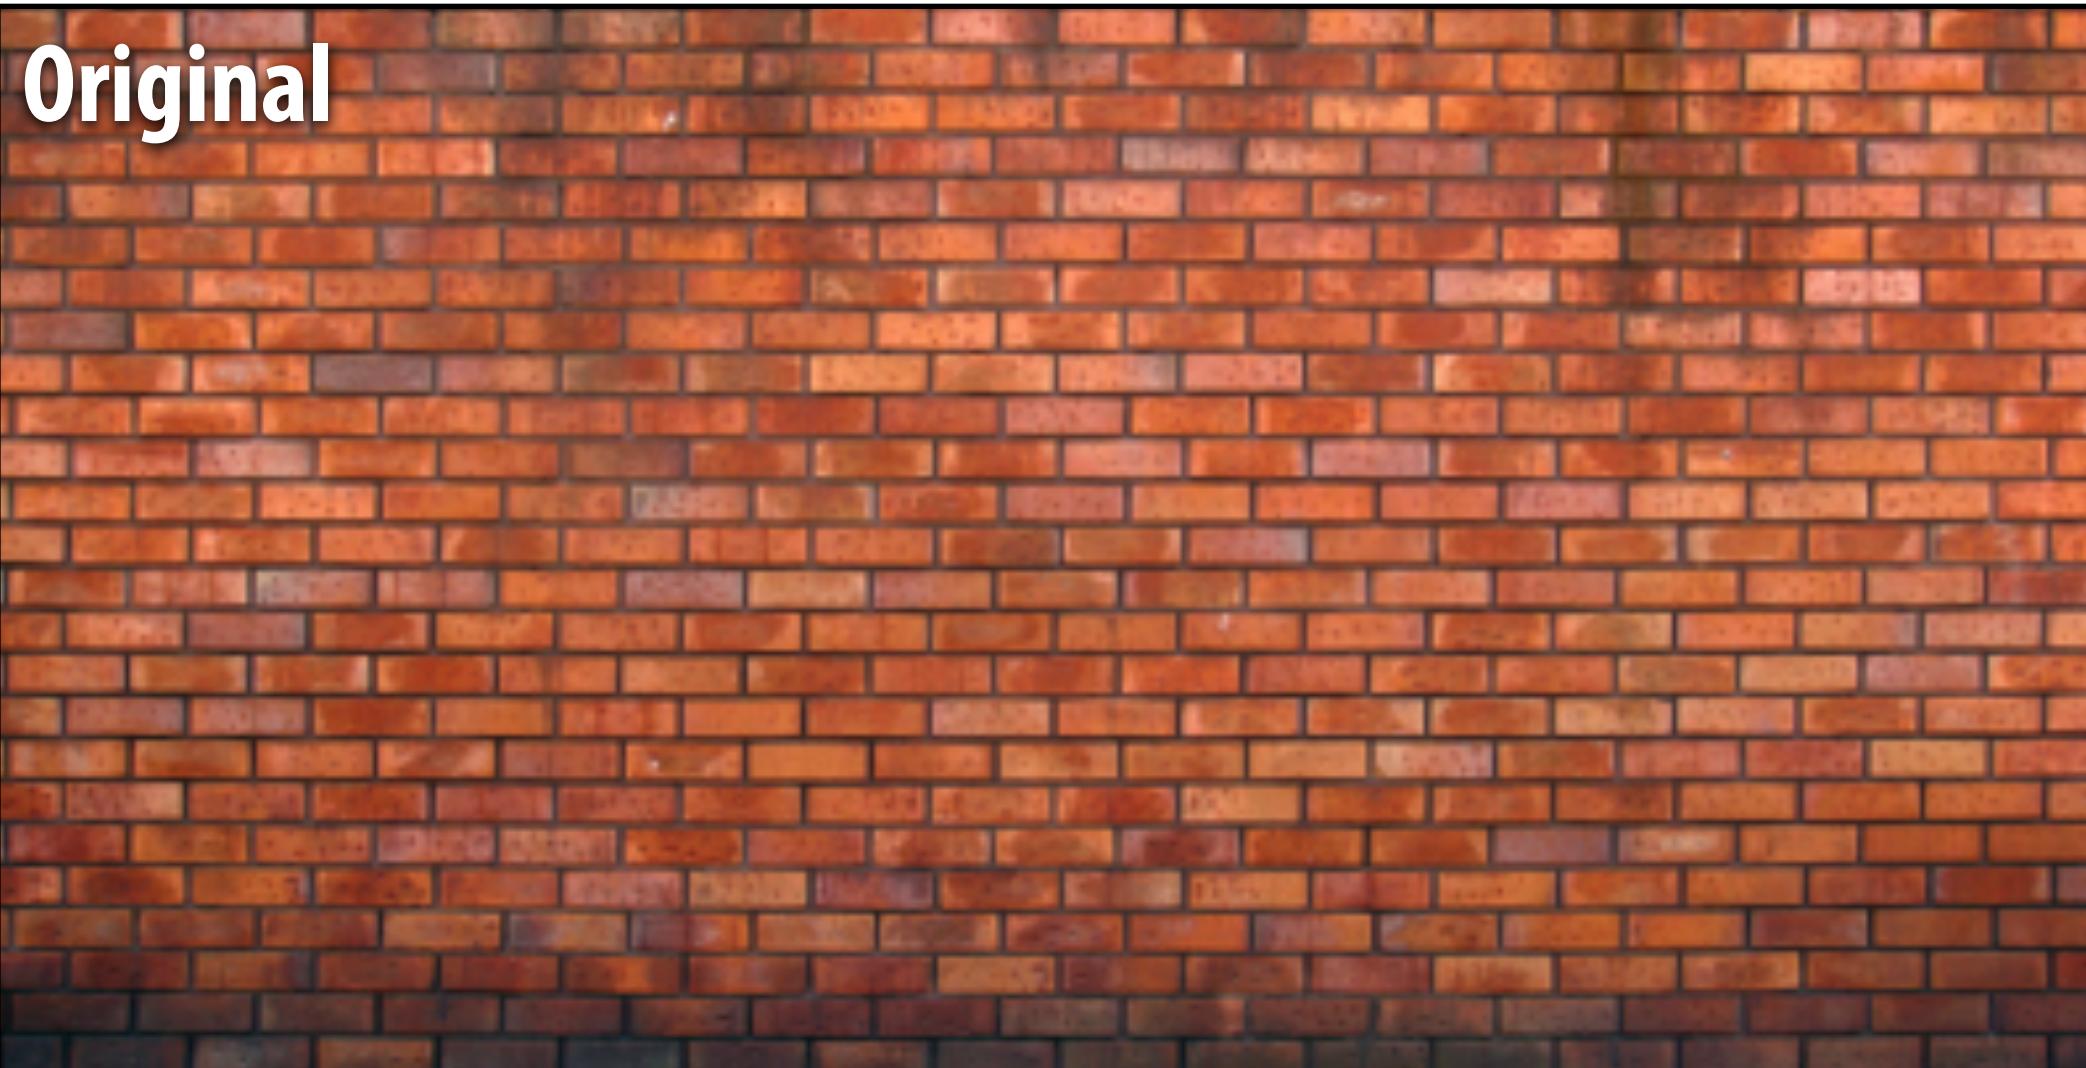
- Obtain filter coefficients from sampling 2D Gaussian

$$f(i, j) = \frac{1}{2\pi\sigma^2} e^{-\frac{i^2+j^2}{2\sigma^2}}$$

- Produces weighted sum of neighboring pixels (contribution falls off with distance)
 - In practice: truncate filter beyond certain distance for efficiency

$$\begin{bmatrix} .075 & .124 & .075 \\ .124 & .204 & .124 \\ .075 & .124 & .075 \end{bmatrix}$$

7x7 gaussian blur

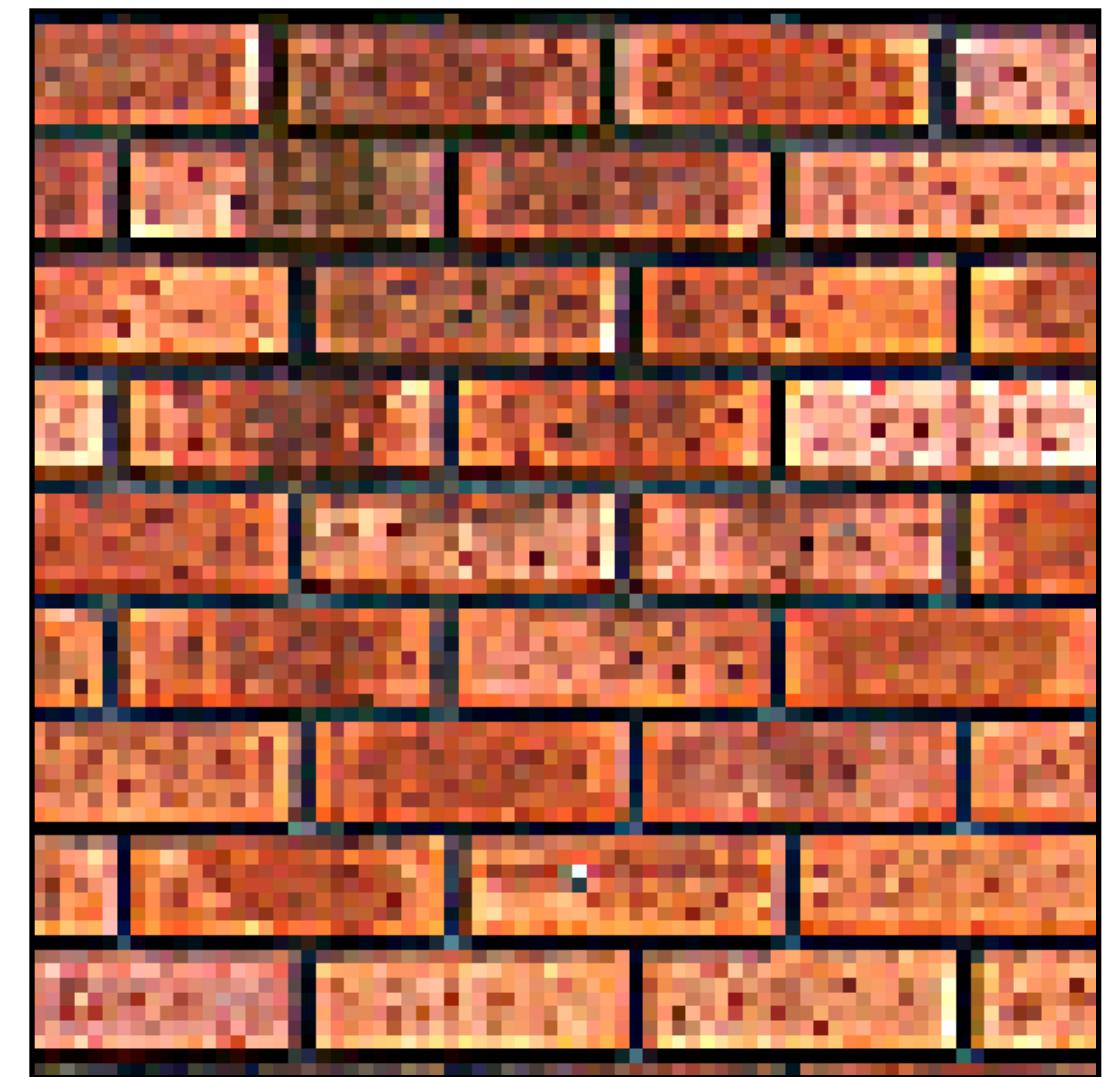
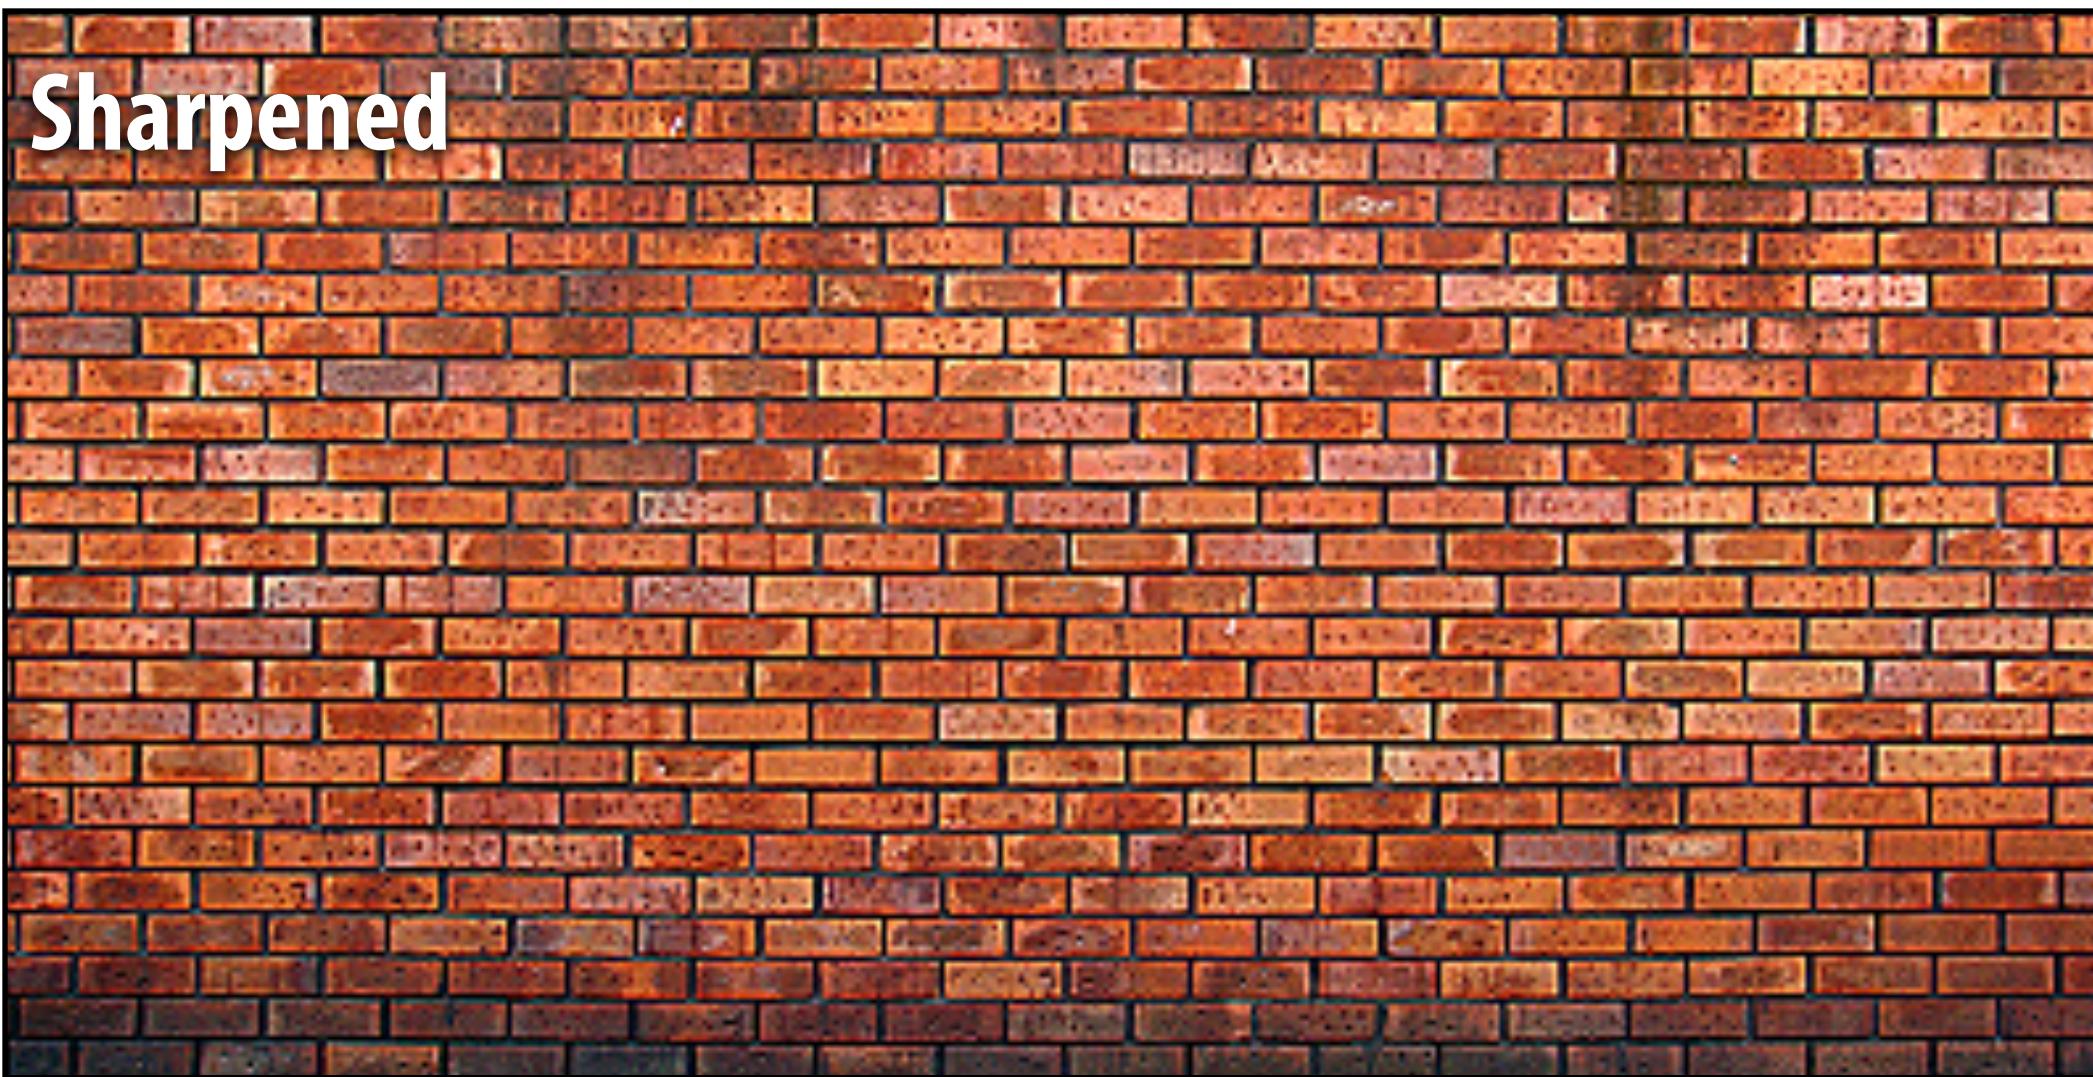
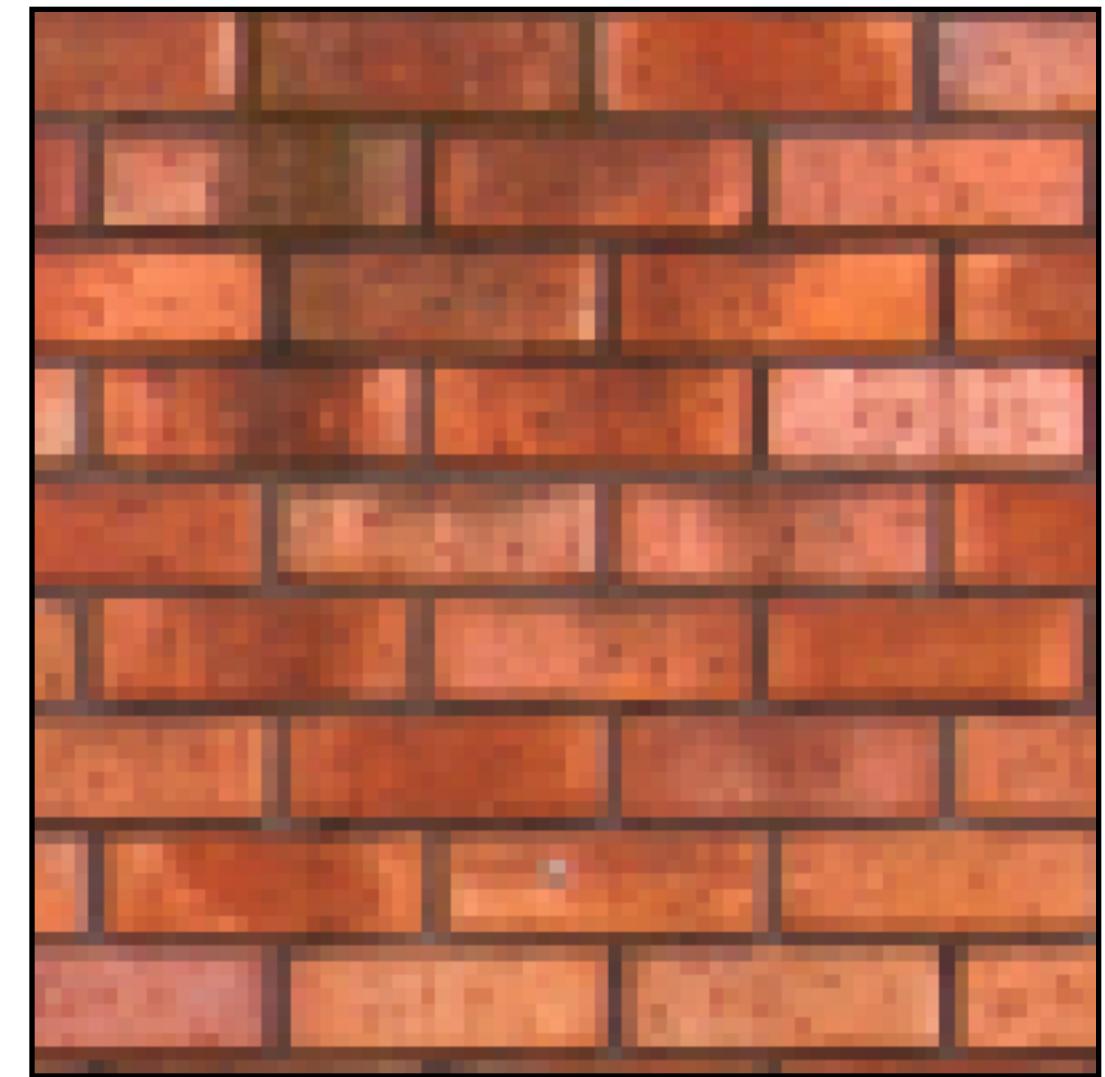
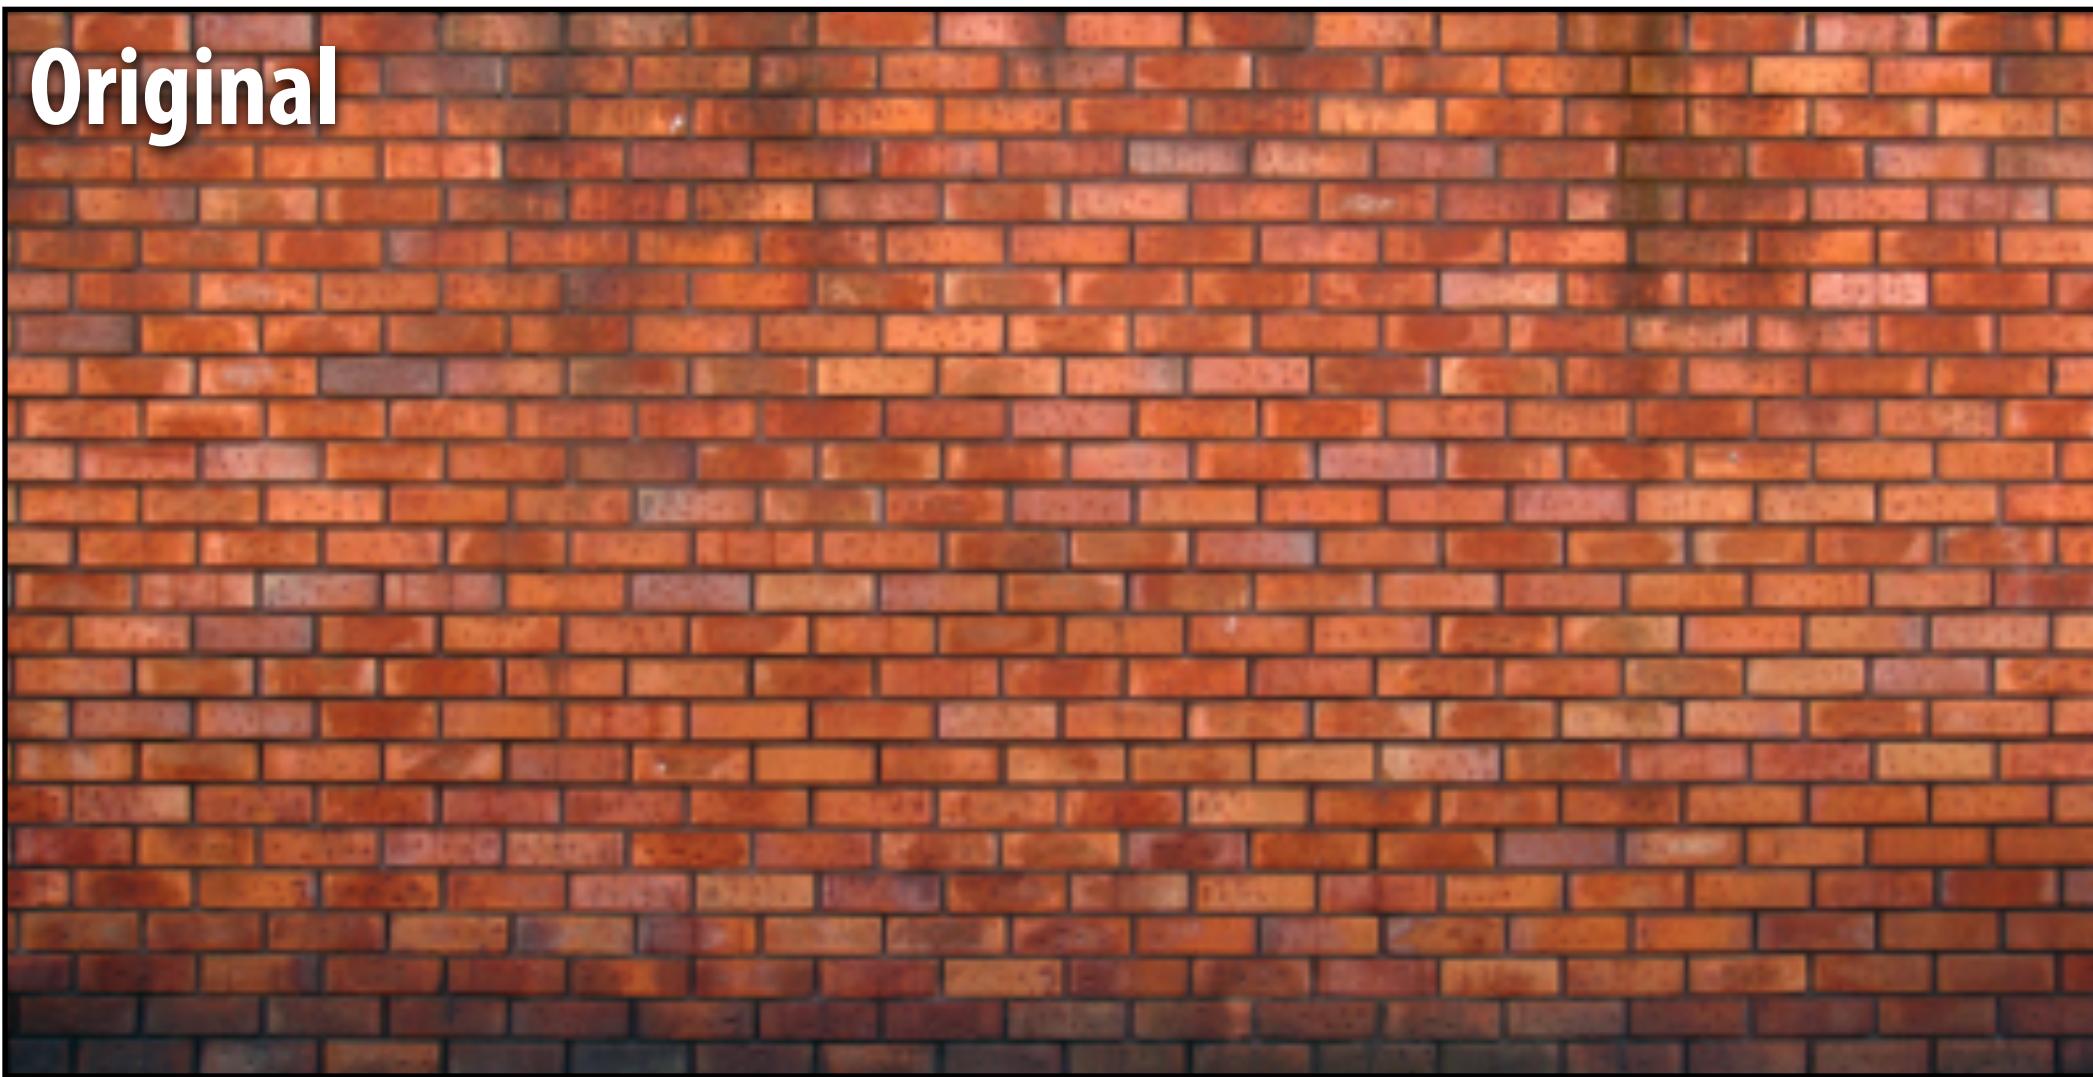


What does convolution with this filter do?

$$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$

Sharpens image!

3x3 sharpen filter



What does convolution with these filters do?

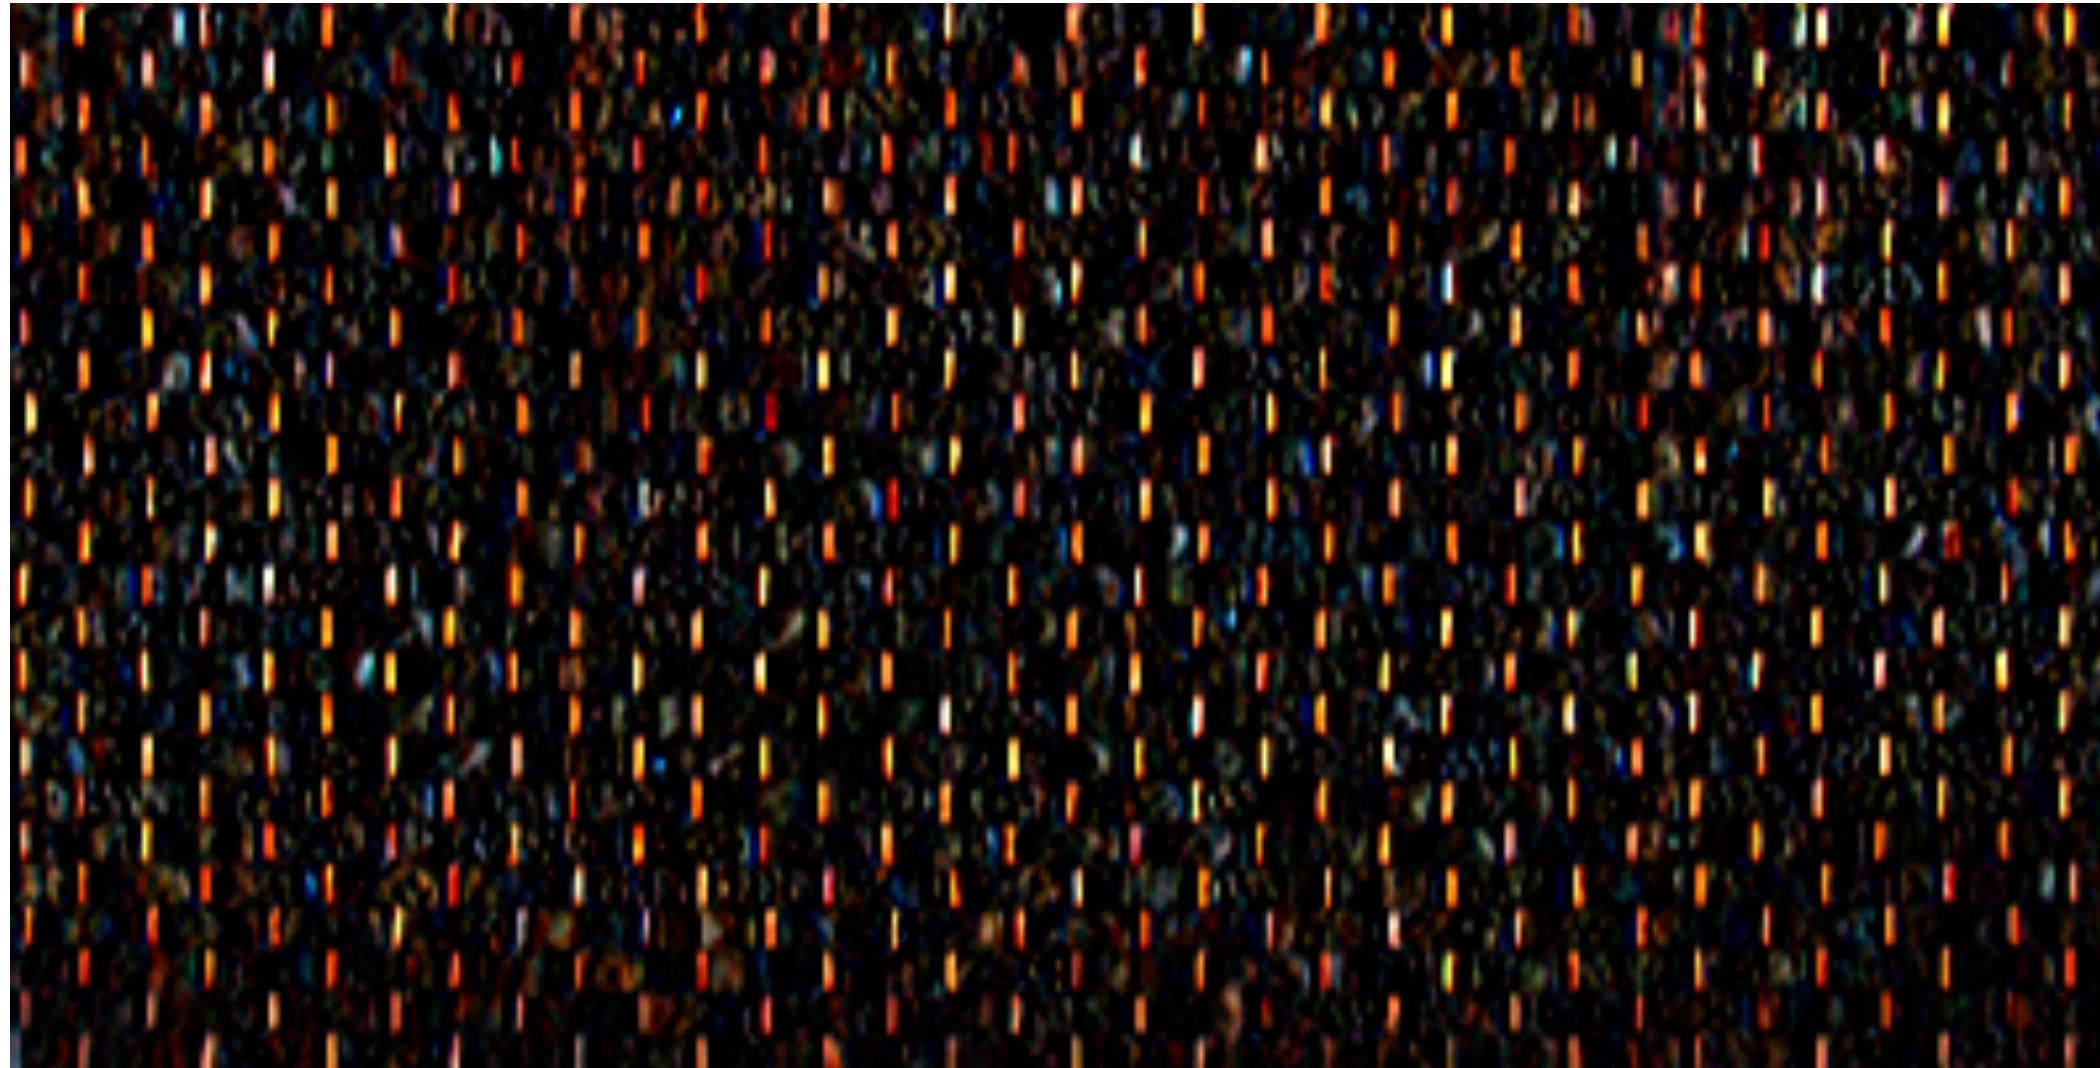
$$\begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

Extracts horizontal
gradients

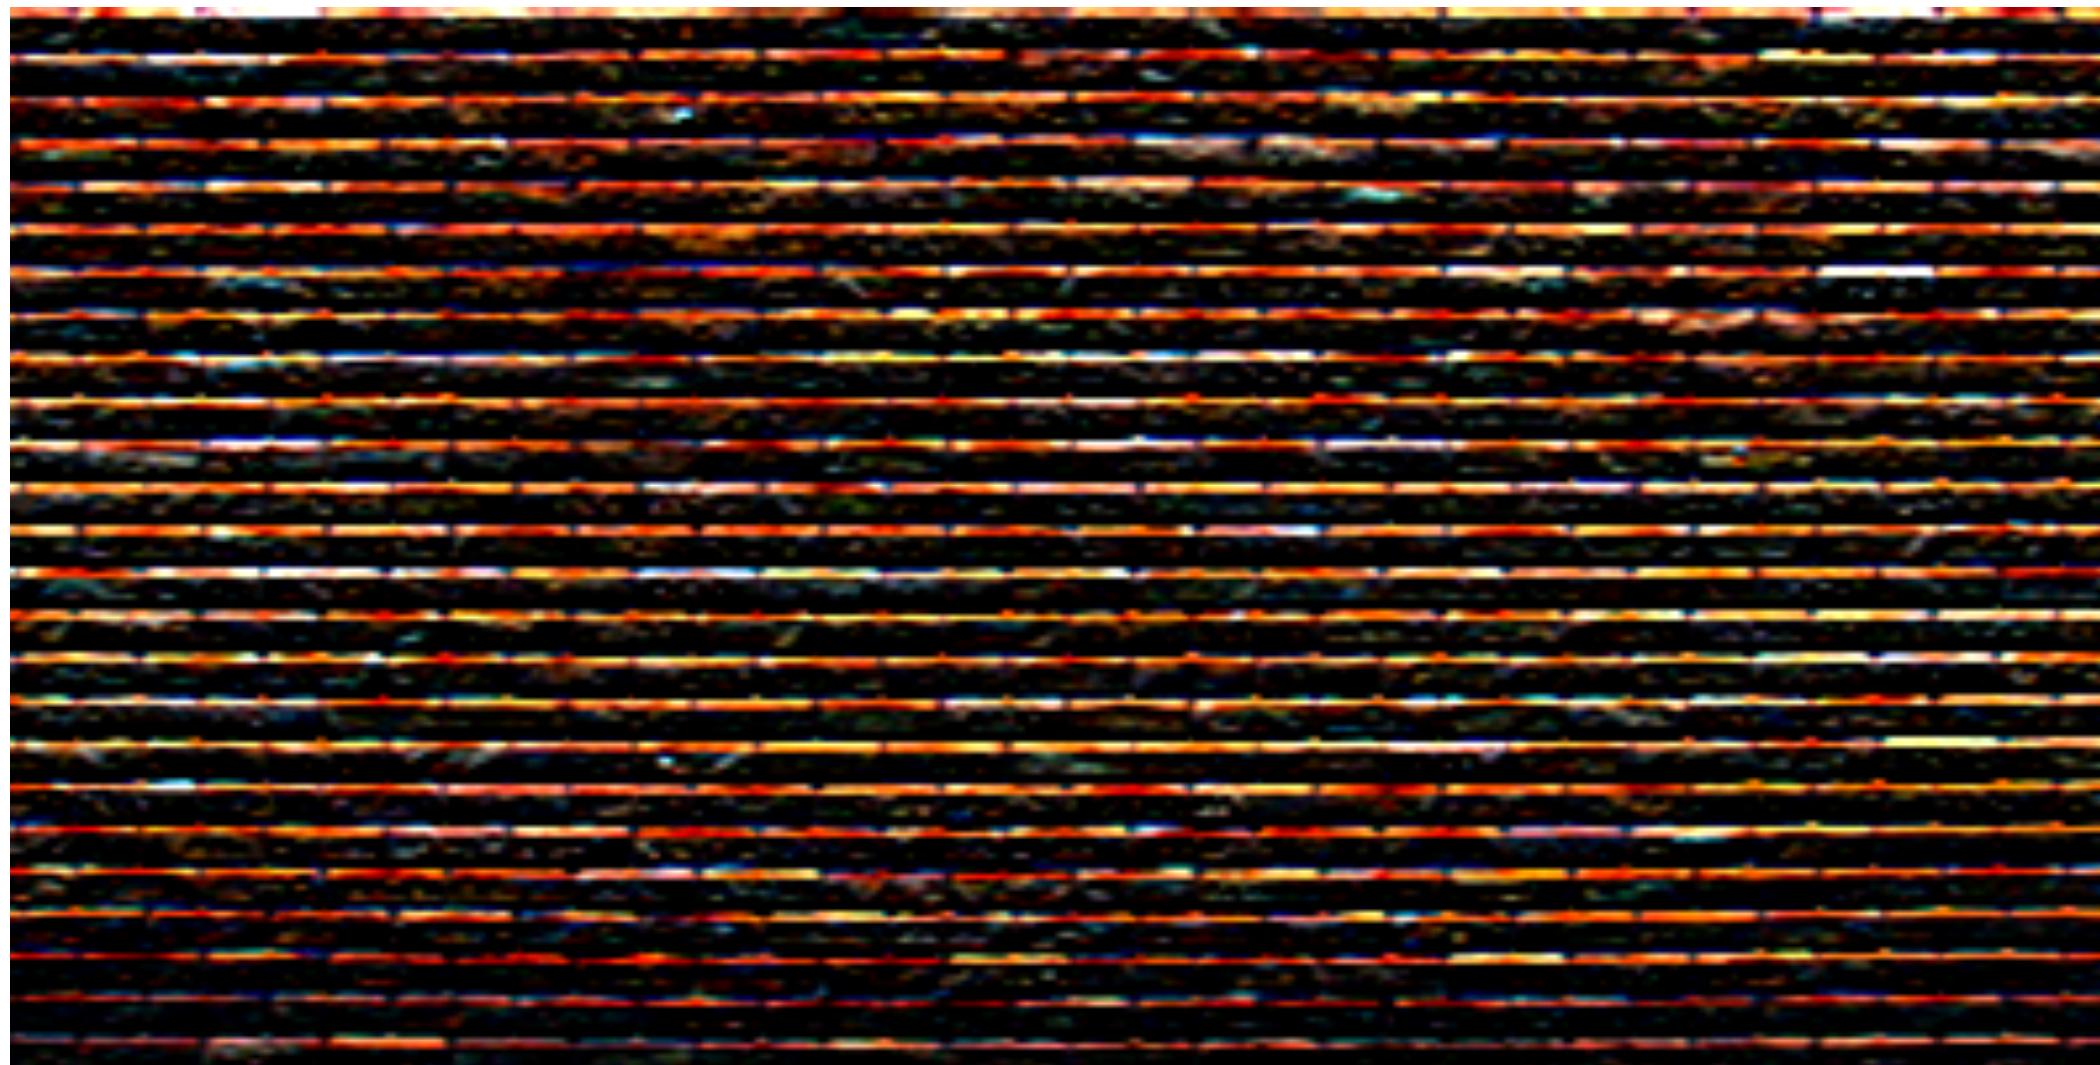
$$\begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$

Extracts vertical
gradients

Gradient detection filters



Horizontal gradients



Vertical gradients

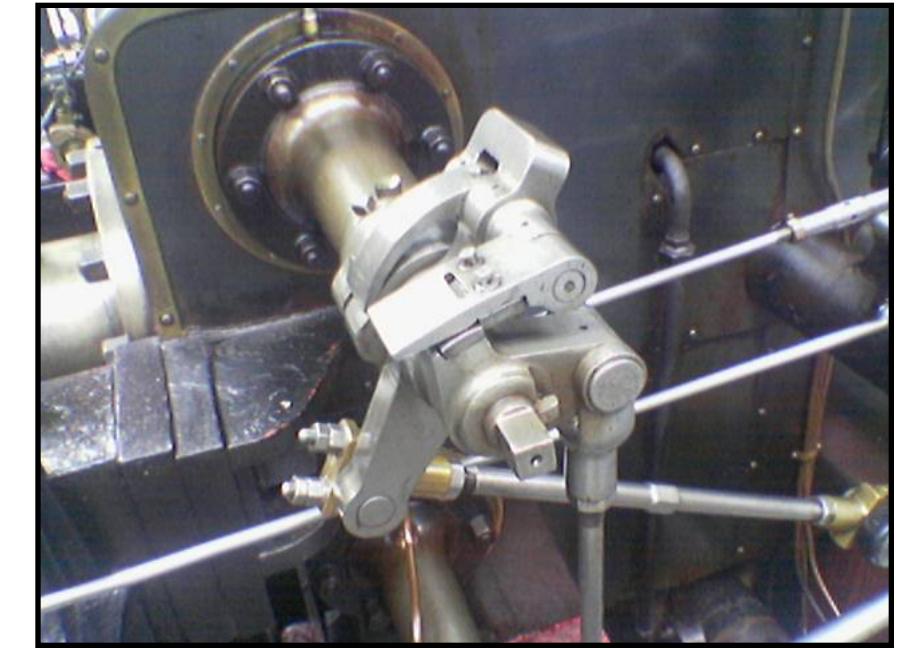
Note: you can think of a filter as a “detector” of a pattern, and the magnitude of a pixel in the output image as the “response” of the filter to the region surrounding each pixel in the input image (this is a common interpretation in computer vision)

Sobel edge detection

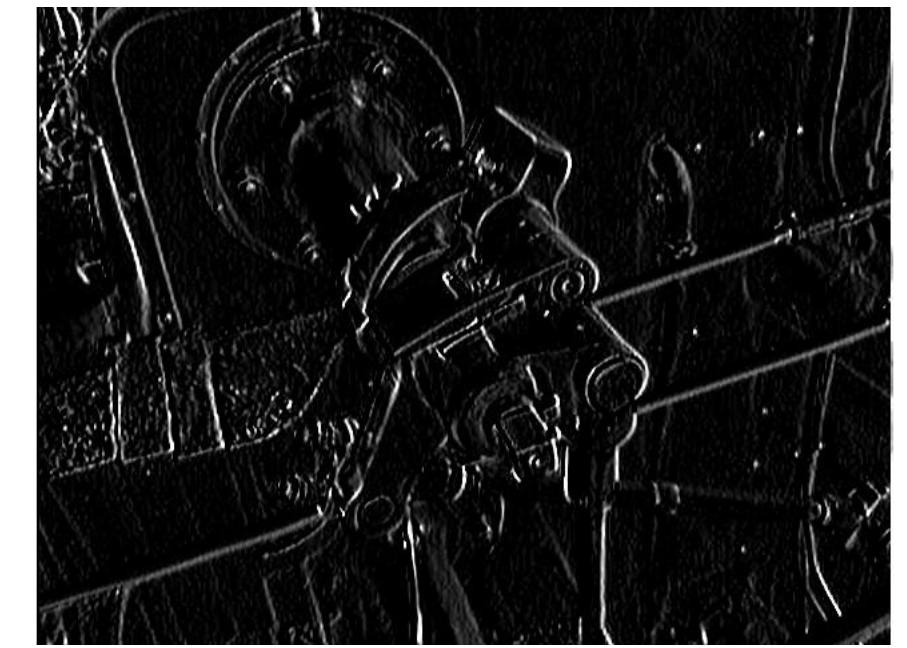
- Compute gradient response images

$$G_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix} * I$$

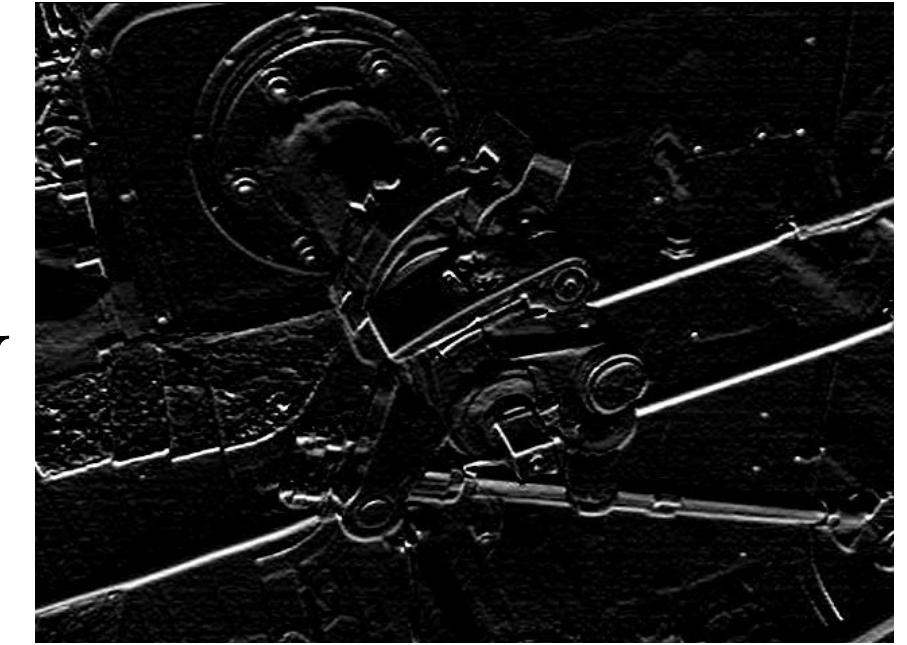
$$G_y = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix} * I$$



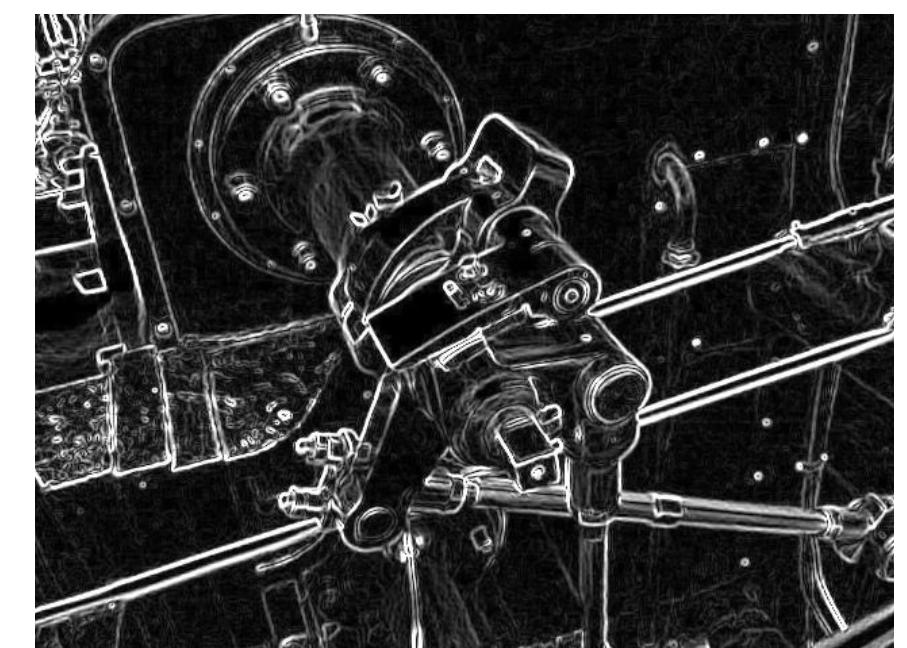
G_x



G_y



G



- Find pixels with large gradients

$$G = \sqrt{G_x^2 + G_y^2}$$

Pixel-wise operation on images

Cost of convolution with $N \times N$ filter?

```
float input[(WIDTH+2) * (HEIGHT+2)];  
float output[WIDTH * HEIGHT];
```

**In this 3x3 box blur example:
Total work per image = $9 \times \text{WIDTH} \times \text{HEIGHT}$**

```
float weights[] = {1./9, 1./9, 1./9,  
                   1./9, 1./9, 1./9,  
                   1./9, 1./9, 1./9};
```

```
for (int j=0; j<HEIGHT; j++) {  
    for (int i=0; i<WIDTH; i++) {  
        float tmp = 0.f;  
        for (int jj=0; jj<3; jj++)  
            for (int ii=0; ii<3; ii++)  
                tmp += input[(j+jj)*(WIDTH+2) + (i+ii)] * weights[jj*3 + ii];  
        output[j*WIDTH + i] = tmp;  
    }  
}
```

For $N \times N$ filter: $N^2 \times \text{WIDTH} \times \text{HEIGHT}$

Separable filter

- A filter is separable if can bee written as the outer product of two other filters. Example: a 2D box blur

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} = \frac{1}{3} \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} * \frac{1}{3} [1 \quad 1 \quad 1]$$

- Exercise: write 2D gaussian and vertical/horizontal gradient detection filters as product of 1D filters (they are separable!)
- Key property: 2D convolution with separable filter can be written as two 1D convolutions!

Implementation of 2D box blur via two 1D convolutions

```
int WIDTH = 1024
int HEIGHT = 1024;
float input[(WIDTH+2) * (HEIGHT+2)];
float tmp_buf[WIDTH * (HEIGHT+2)];
float output[WIDTH * HEIGHT];

float weights[] = {1./3, 1./3, 1./3};

for (int j=0; j<(HEIGHT+2); j++)
    for (int i=0; i<WIDTH; i++) {
        float tmp = 0.f;
        for (int ii=0; ii<3; ii++)
            tmp += input[j*(WIDTH+2) + i+ii] * weights[ii];
        tmp_buf[j*WIDTH + i] = tmp;
    }

for (int j=0; j<HEIGHT; j++) {
    for (int i=0; i<WIDTH; i++) {
        float tmp = 0.f;
        for (int jj=0; jj<3; jj++)
            tmp += tmp_buf[(j+jj)*WIDTH + i] * weights[jj];
        output[j*WIDTH + i] = tmp;
    }
}
```

Total work per image for NxN filter:
 $2N \times \text{WIDTH} \times \text{HEIGHT}$

Data-dependent filter (not a convolution)

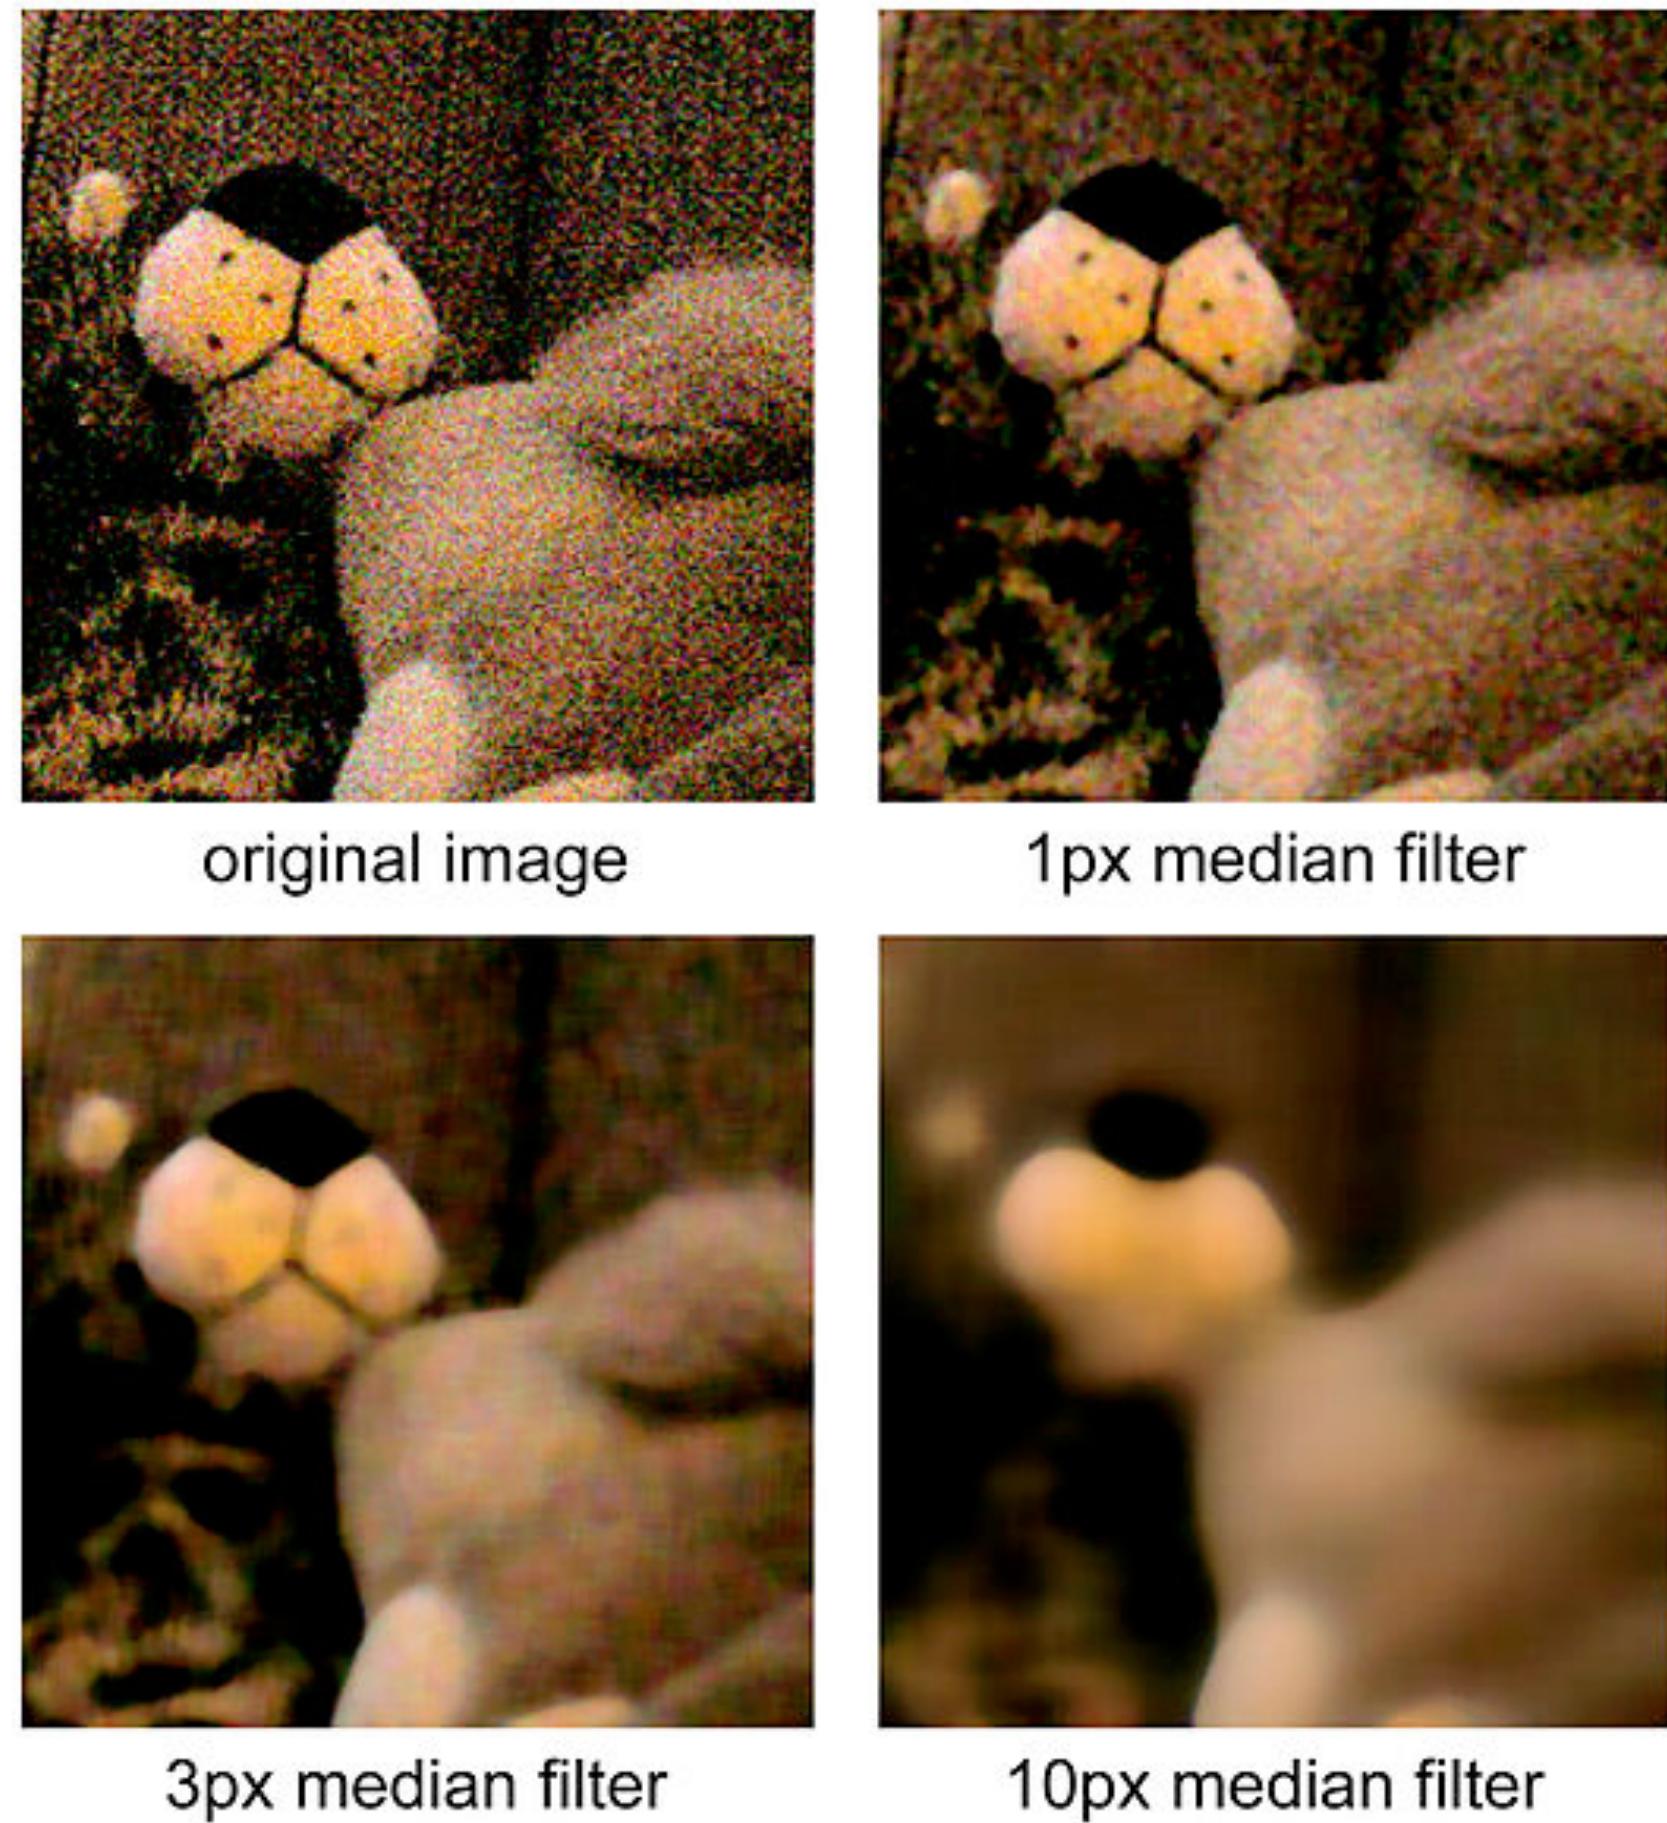
```
float input[(WIDTH+2) * (HEIGHT+2)];  
float output[WIDTH * HEIGHT];  
  
for (int j=0; j<HEIGHT; j++) {  
    for (int i=0; i<WIDTH; i++) {  
        float min_value = min( min(input[(j-1)*WIDTH + i], input[(j+1)*WIDTH + i]),  
                               min(input[j*WIDTH + i-1], input[j*WIDTH + i+1]) );  
        float max_value = max( max(input[(j-1)*WIDTH + i], input[(j+1)*WIDTH + i]),  
                               max(input[j*WIDTH + i-1], input[j*WIDTH + i+1]) );  
        output[j*WIDTH + i] = clamp(min_value, max_value, input[j*WIDTH + i]);  
    }  
}
```

**This filter clamps pixels to the min/max of its cardinal neighbors
(e.g., hot-pixel suppression)**

Median filter

- Replace pixel with median of its neighbors
 - Useful noise reduction filter: unlike gaussian blur, one bright pixel doesn't drag up the average for entire region
- Not linear, not separable
 - Filter weights are 1 or 0 (depending on image content)

```
uint8 input[(WIDTH+2) * (HEIGHT+2)];  
uint8 output[WIDTH * HEIGHT];  
for (int j=0; j<HEIGHT; j++) {  
    for (int i=0; i<WIDTH; i++) {  
        output[j*WIDTH + i] =  
            // compute median of pixels  
            // in surrounding 5x5 pixel window  
    }  
}  
}
```



- Basic algorithm for $N \times N$ support region:
 - Sort N^2 elements in support region, then pick median: $O(N^2 \log(N^2))$ work per pixel
 - Can you think of an $O(N^2)$ algorithm? What about $O(N)$?

5x5 median filter (N=5)

- $O(N^2)$ work-per-pixel solution for 8-bit pixel data (radix sort 8 bit-integer data)
 - Bin all pixels in support region, then scan histogram bins to find median

```
int WIDTH = 1024;
int HEIGHT = 1024;
uint8 input[(WIDTH+2) * (HEIGHT+2)];
uint8 output[WIDTH * HEIGHT];
int histogram[256];

for (int j=0; j<HEIGHT; j++) {
    for (int i=0; i<WIDTH; i++) {
        // construct histogram of support region
        for (int ii=0; ii<256; ii++)
            histogram[ii] = 0;
        for (int jj=0; jj<5; jj++)
            for (int ii=0; ii<5; ii++)
                histogram[input[(j+jj)*(WIDTH+2) + (i+ii)]]++;

        // scan the 256 bins to find median
        // median value of 5x5=25 elements is bin containing 13th value
        int count = 0;
        for (int ii=0; ii<256; ii++) {
            if (count + histogram[ii] >= 13)
                output[j*WIDTH + i] = uint8(ii);
            count += histogram[ii];
        }
    }
}
```

See Weiss [SIGGRAPH 2006] for
 $O(\lg N)$ work-per-pixel median filter
(incrementally updates histogram)

Bilateral filter



Example use of bilateral filter: removing noise while preserving image edges

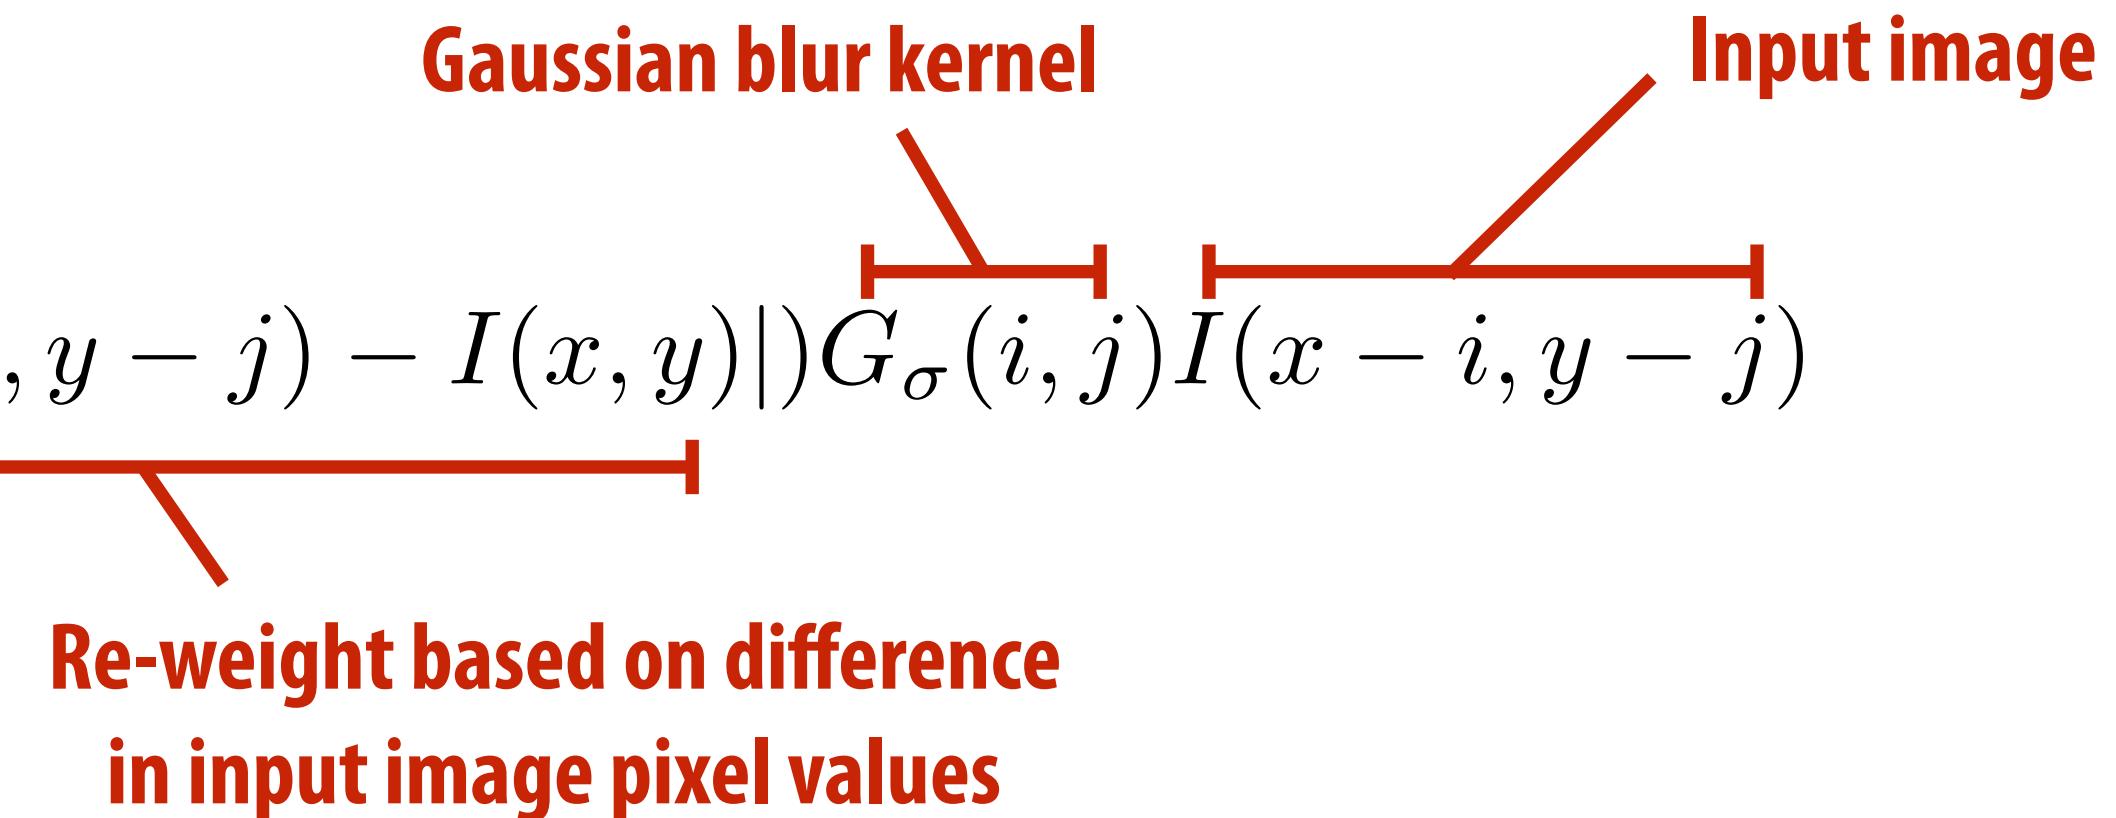
Bilateral filter

$$BF[I](p) = \frac{1}{W_p} \sum_{i,j} f(|I(x-i, y-j) - I(x, y)|) G_\sigma(i, j) I(x-i, y-j)$$

Normalization

For all pixels in support region of Gaussian kernel

$$W_p = \sum_{i,j} f(|I(x-i, y-j) - I(x, y)|) G_\sigma(i, j) I(x-i, y-j)$$

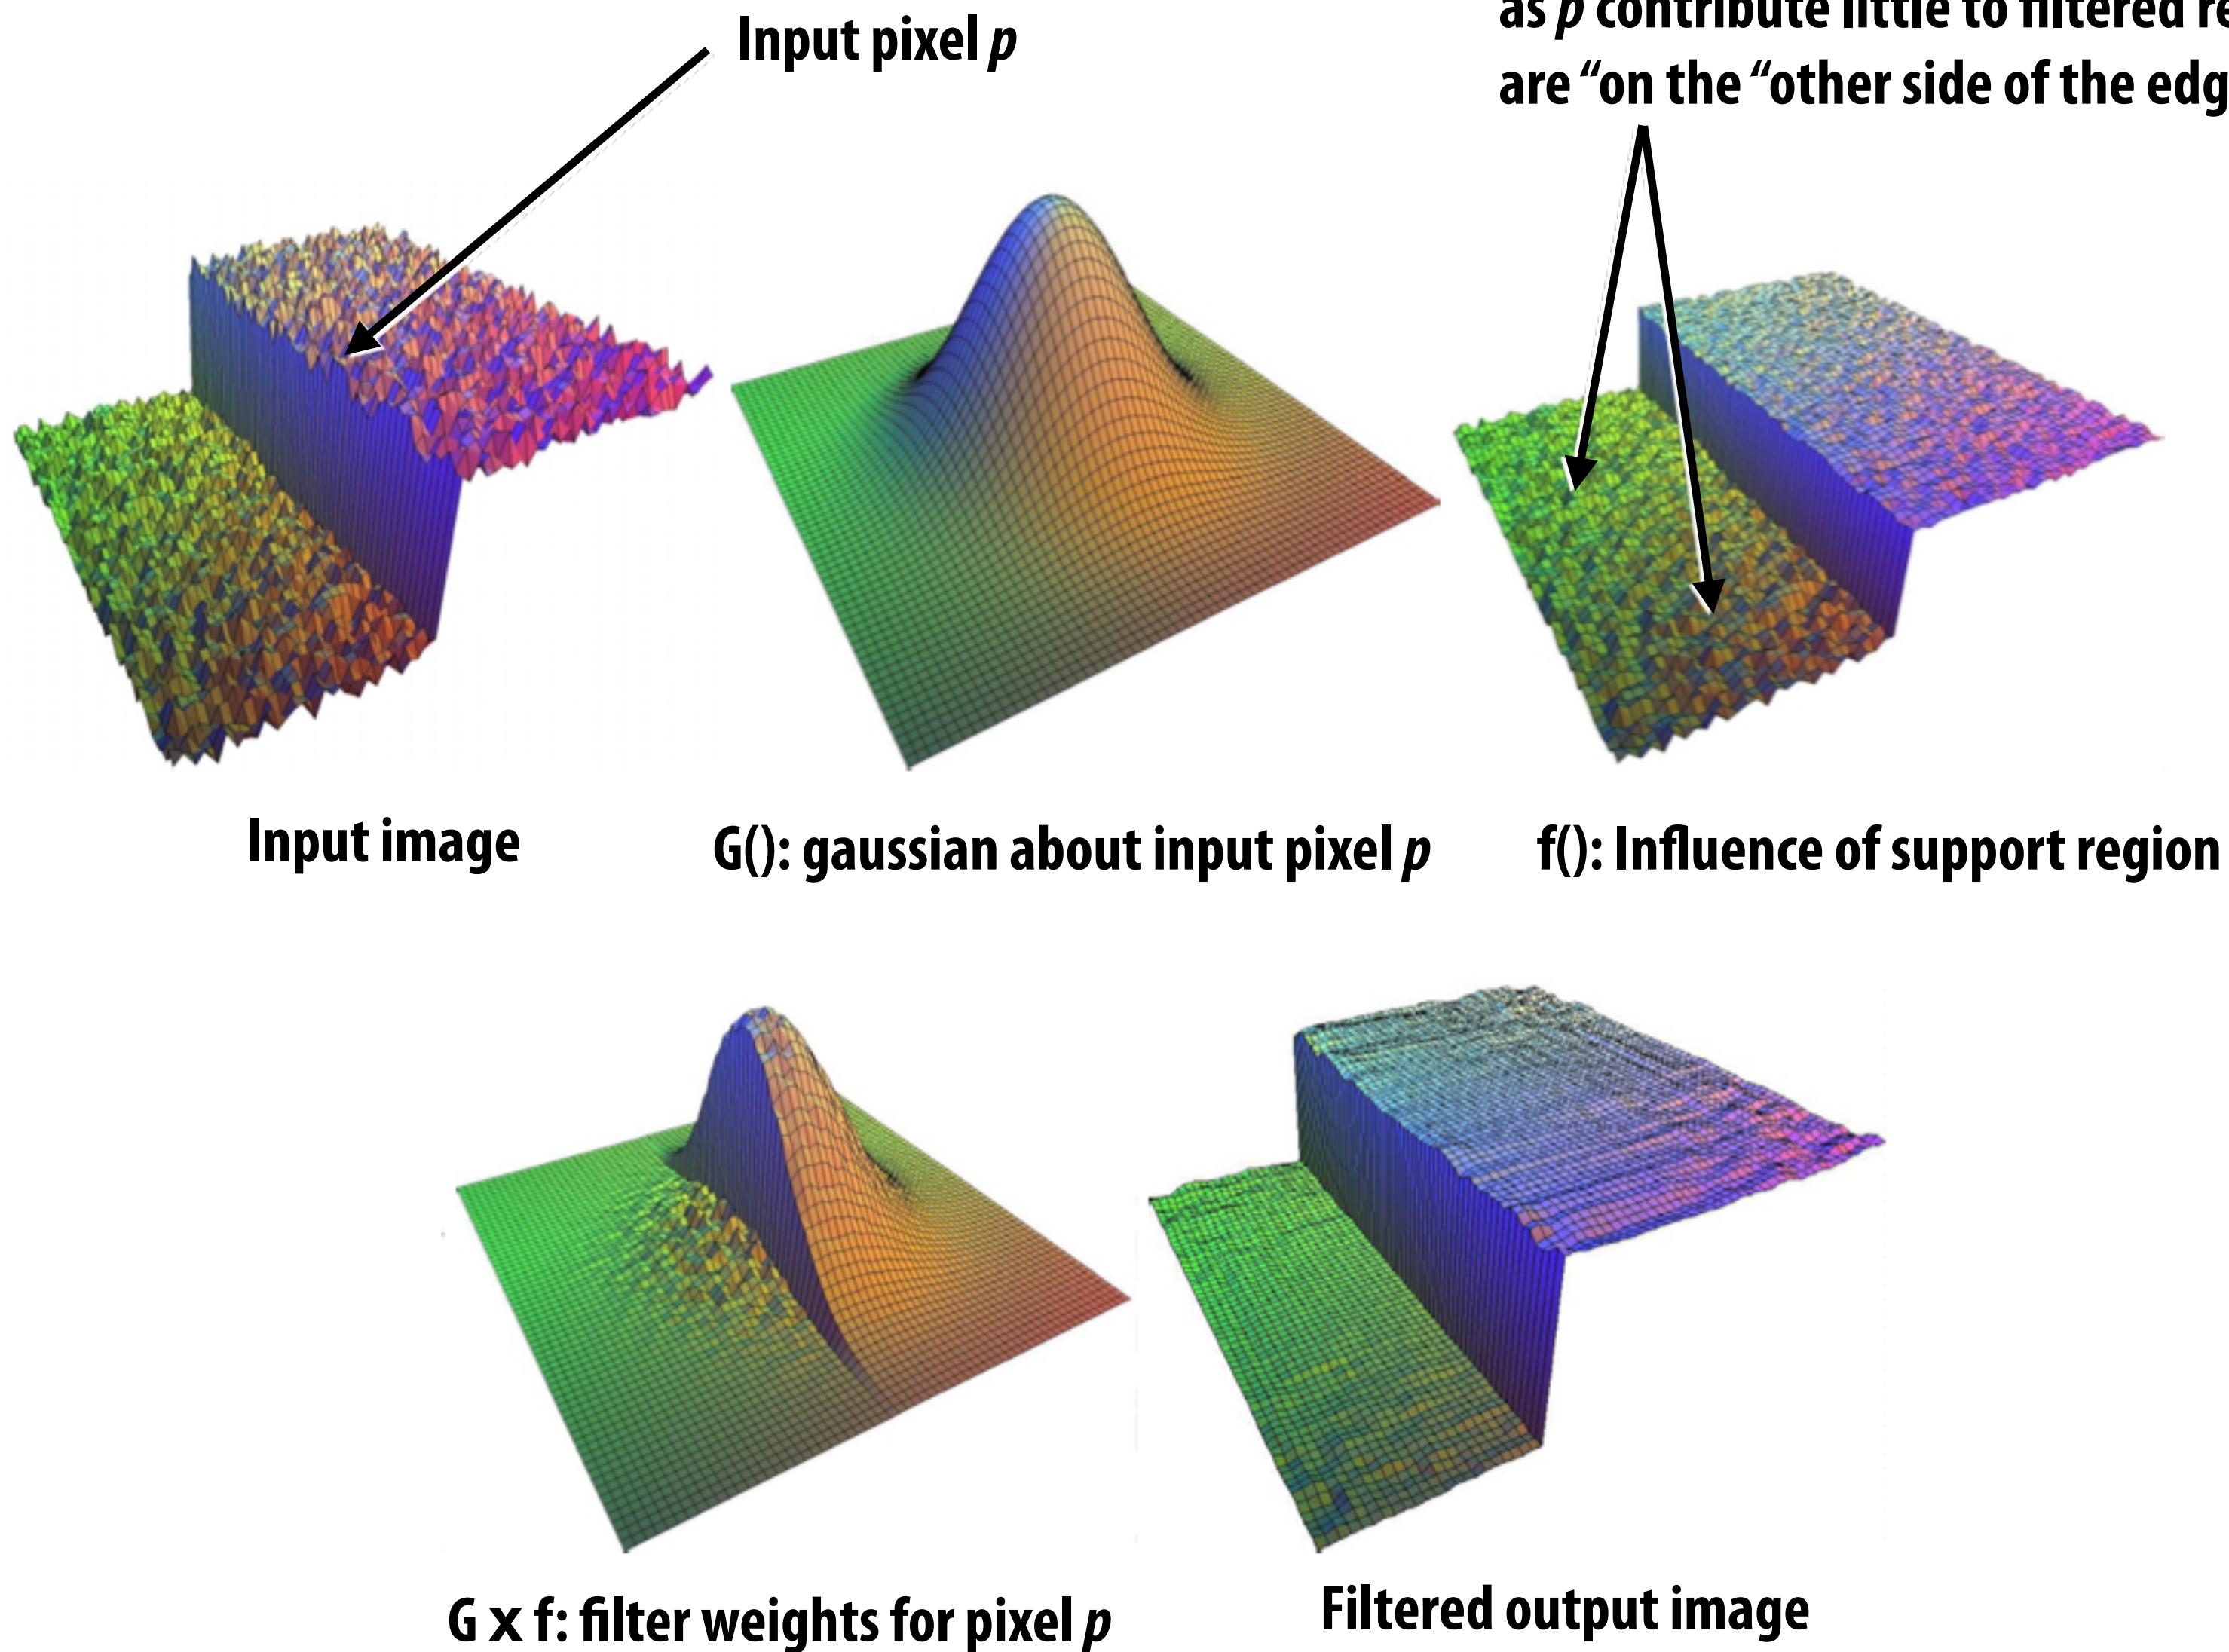


- The bilateral filter is an “edge preserving” filter: down-weight contribution of pixels on the “other side” of strong edges. $f(x)$ defines what “strong edge means”
- Spatial distance weight term $f(x)$ could itself be a gaussian
 - Or very simple: $f(x) = 0$ if $x > \text{threshold}$, 1 otherwise

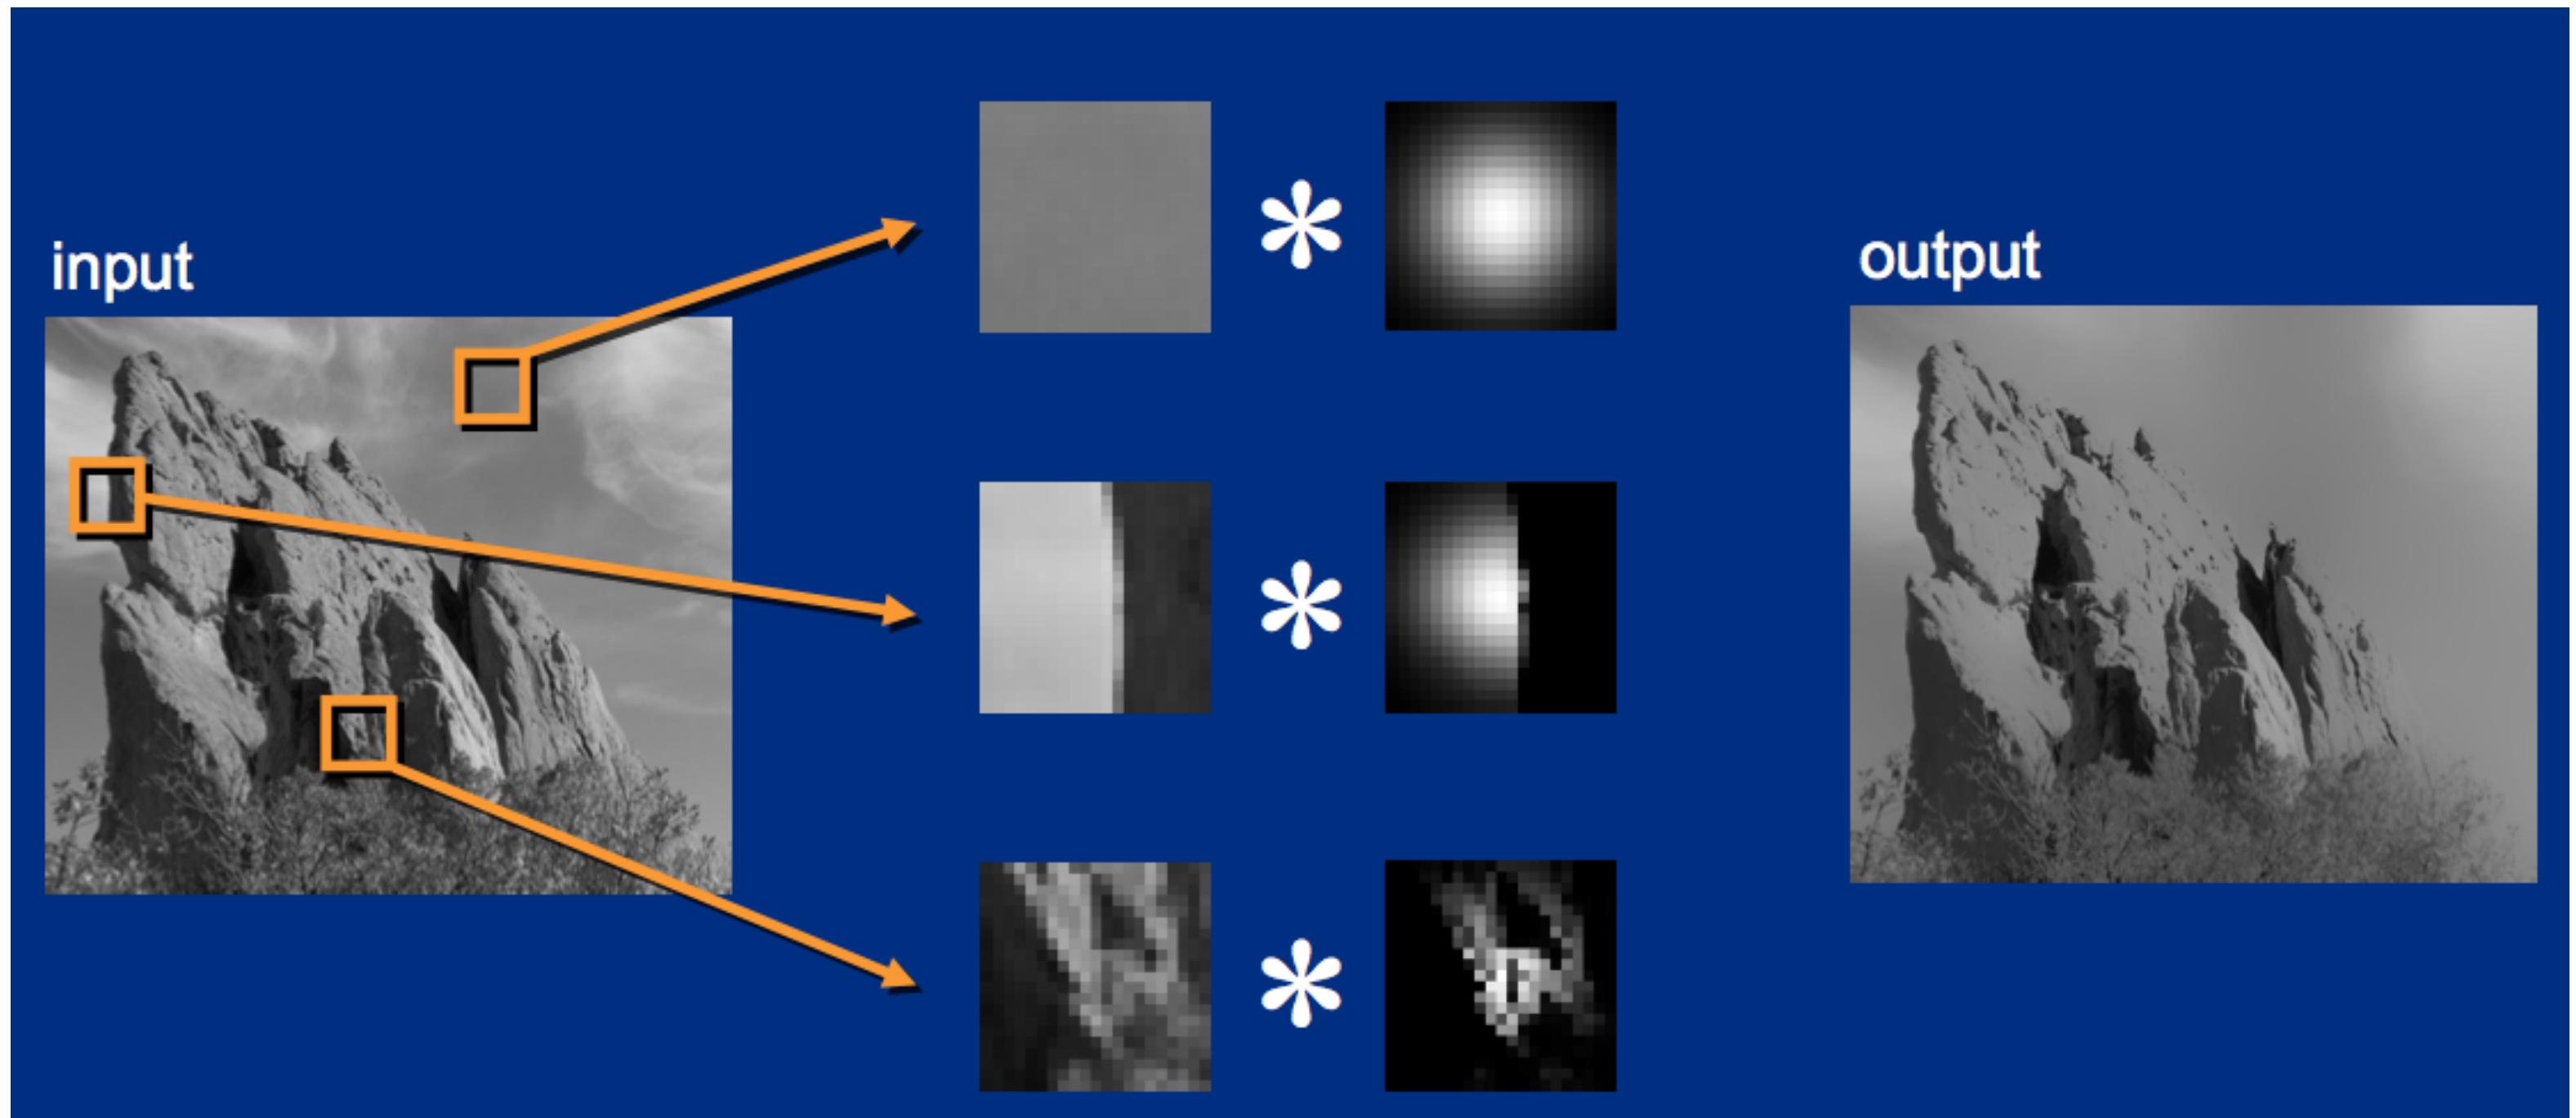
Value of output pixel (x,y) is the weighted sum of all pixels in the support region of a truncated gaussian kernel

But weight is combination of spatial distance and input image pixel intensity difference.
(non-linear filter: like the median filter, the filter’s weights depend on input image content)

Bilateral filter



Bilateral filter: kernel depends on image content



See Paris et al. [ECCV 2006] for a fast approximation to the bilateral filter

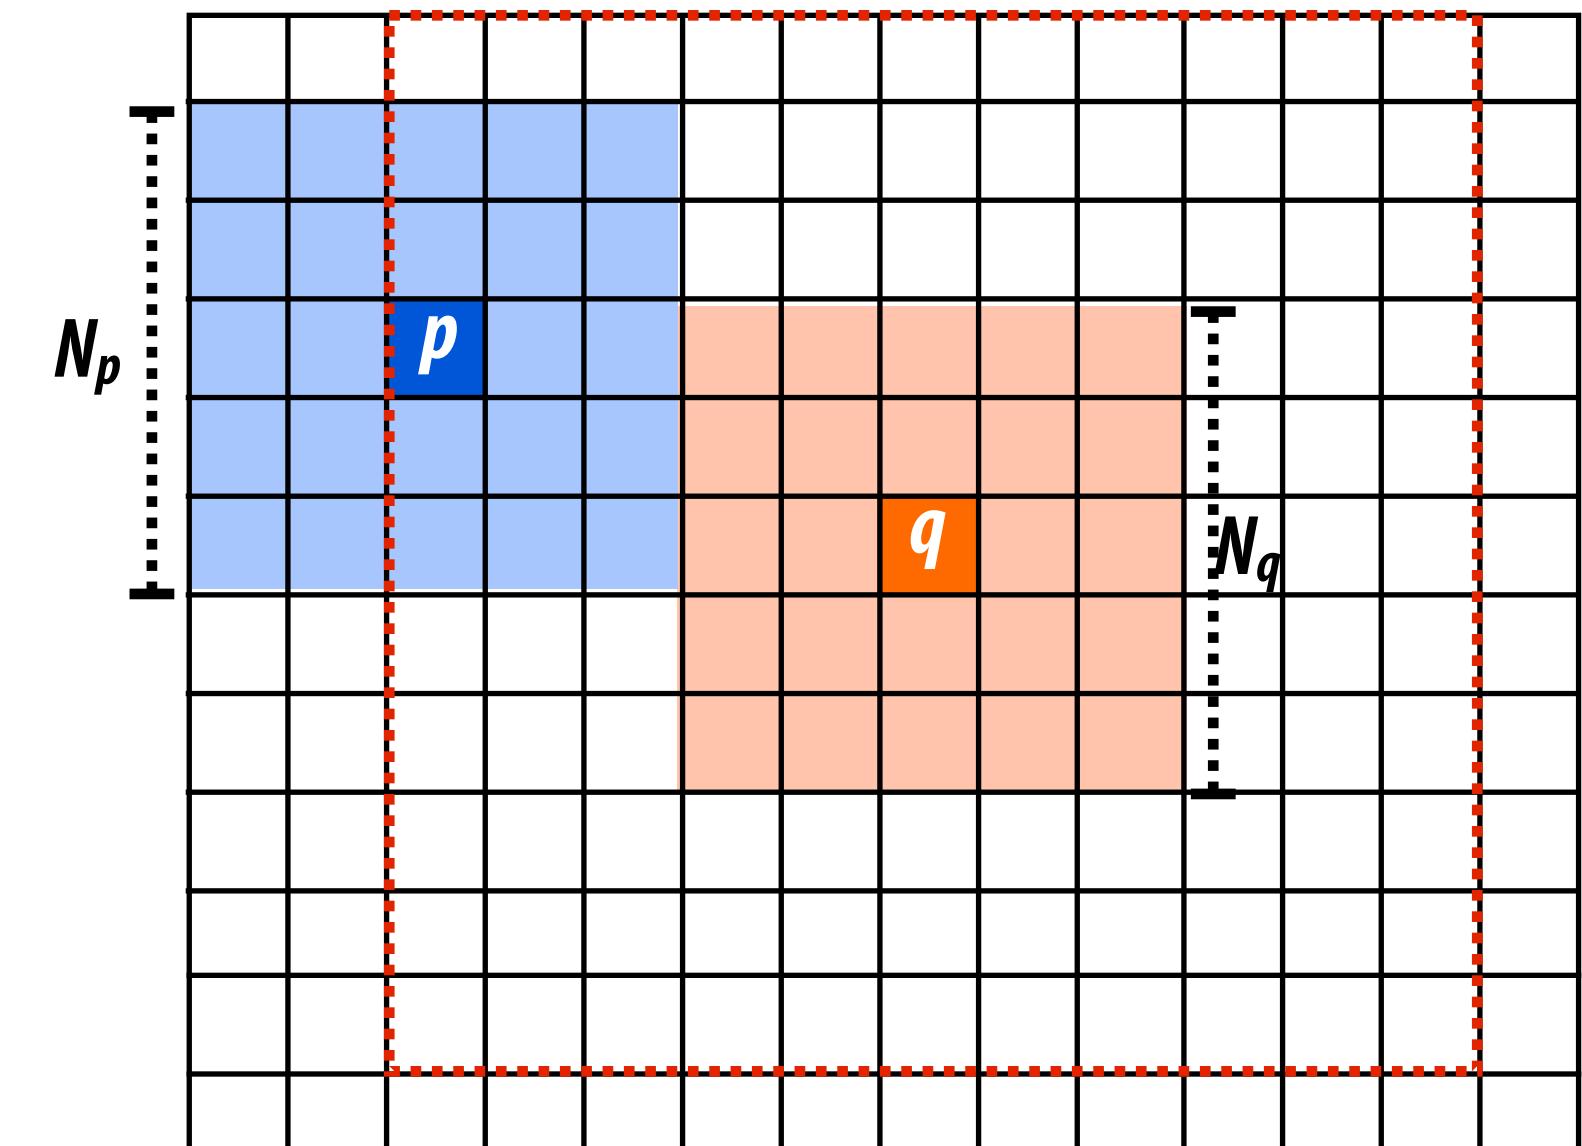
**Question: describe a type of edge the bilateral filter will not respect
(it will blur across these edges)**

Denoising using non-local means

- Main assumption: images have repeating texture
- Main idea: replace pixel with average value of nearby pixels that have a similar surrounding region

$$\text{NL}[I](p) = \sum_{q \in S} w(p, q) I(q)$$

$$w(p, q) = \frac{1}{C_p} e^{-\frac{\|N_p - N_q\|^2}{h^2}}$$



- N_p and N_q are vectors of pixel values in square window around pixels p and q (highlighted regions in figure)
- Difference between N_p and N_q = “similarity” of surrounding regions (here: L2 distance)
- C_p is a normalization constant to ensure weights sum to one for pixel p .
- S is the search region (given by dotted red line in figure)

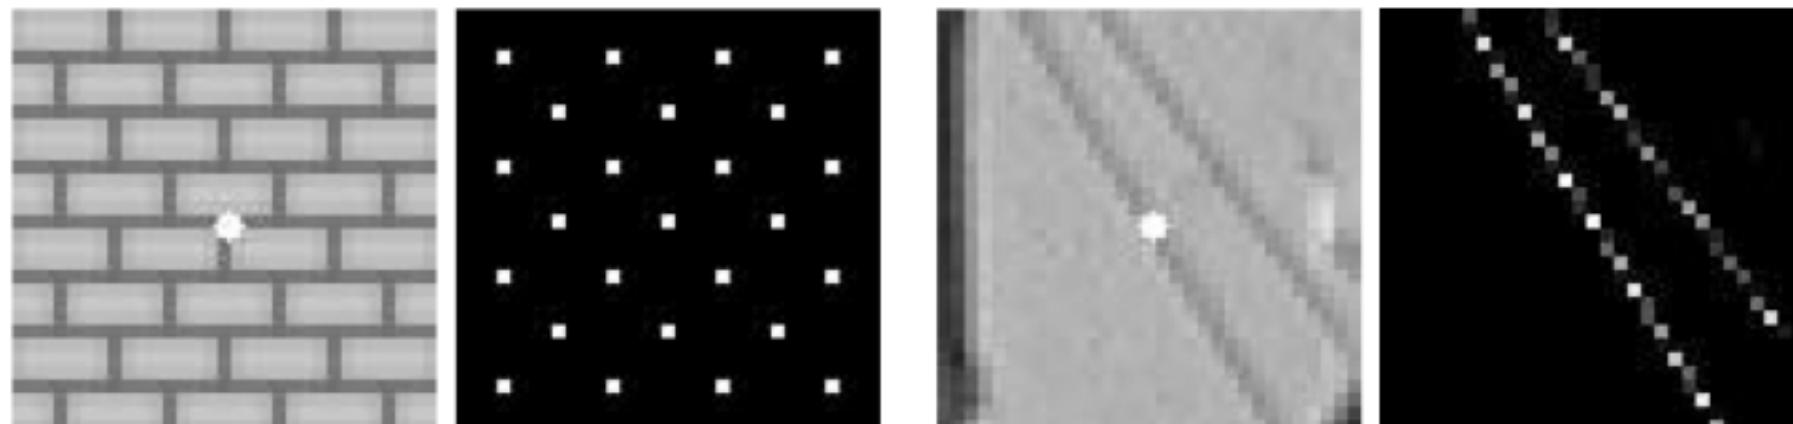
Denoising using non-local means

- Large weight for input pixels that have similar neighborhood as p
 - Intuition: “filtered result is the average of pixels like this one”
 - In example below-right: $q1$ and $q2$ have high weight, $q3$ has low weight



(A)

(B)

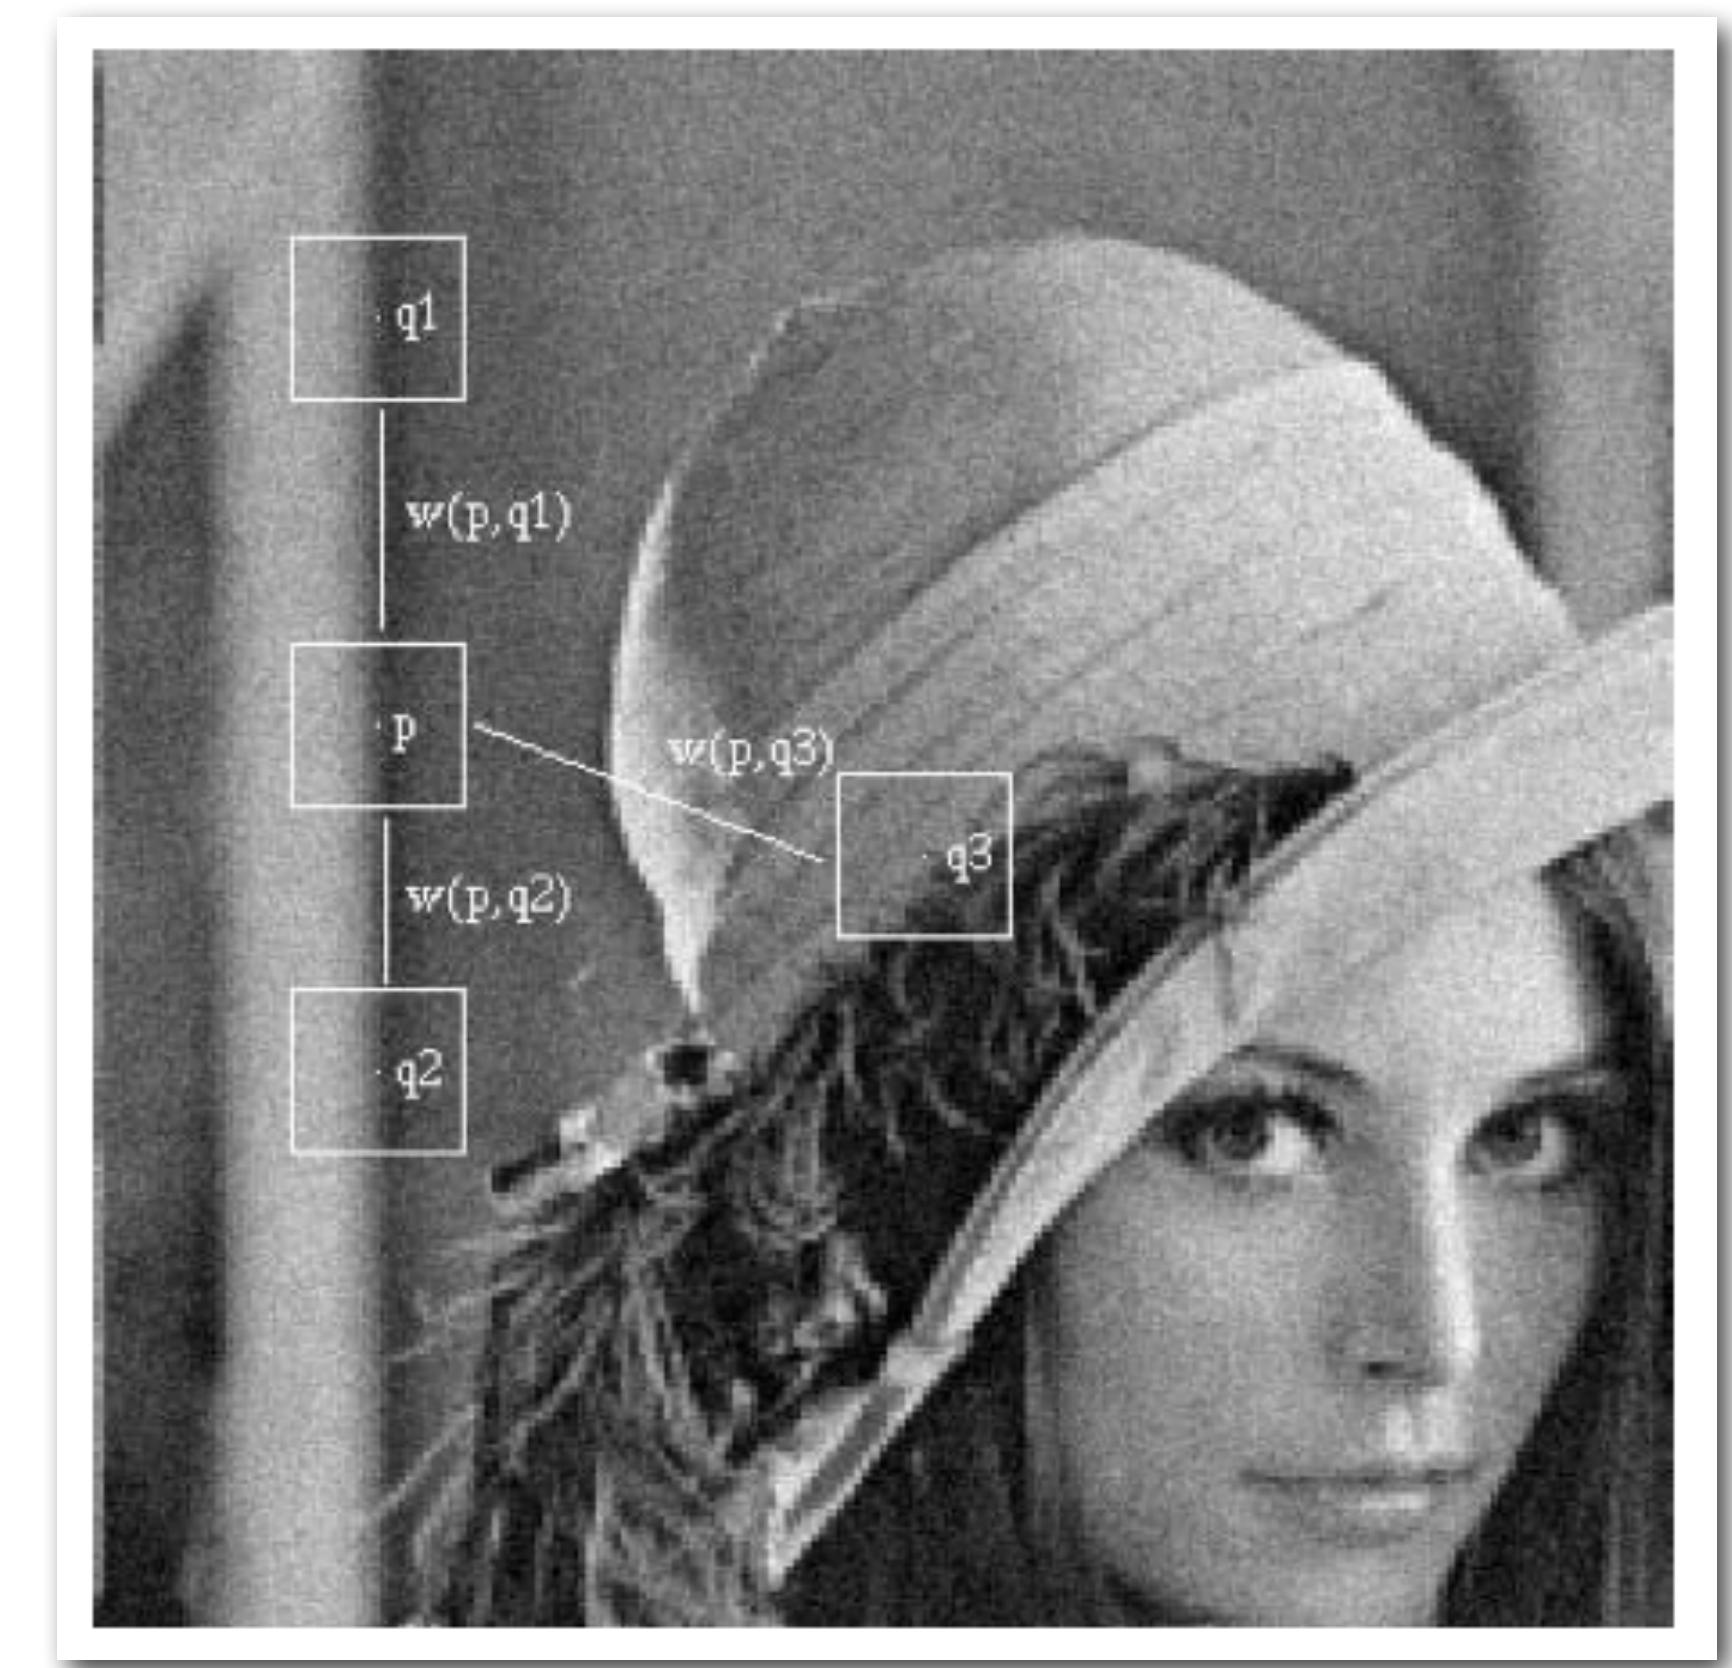


(C)

(D)

In each image pair above:

- Image at left shows the pixel to denoise.
- Image at right shows weights of pixels in 21x21-pixel kernel support window.



**End of aside on image processing basics
(back to our simple camera pipeline)**

Color-space conversion

- Measurements of sensor depend on sensor's spectral response
 - Response depends on bandwidths filtered by color filter array
- Convert representation to sensor-independent basis: e.g., sRGB
 - 3×3 matrix multiplication

```
output_rgb_pixel = COLOR_CONVERSION_MATRIX * input_rgb_pixel
```

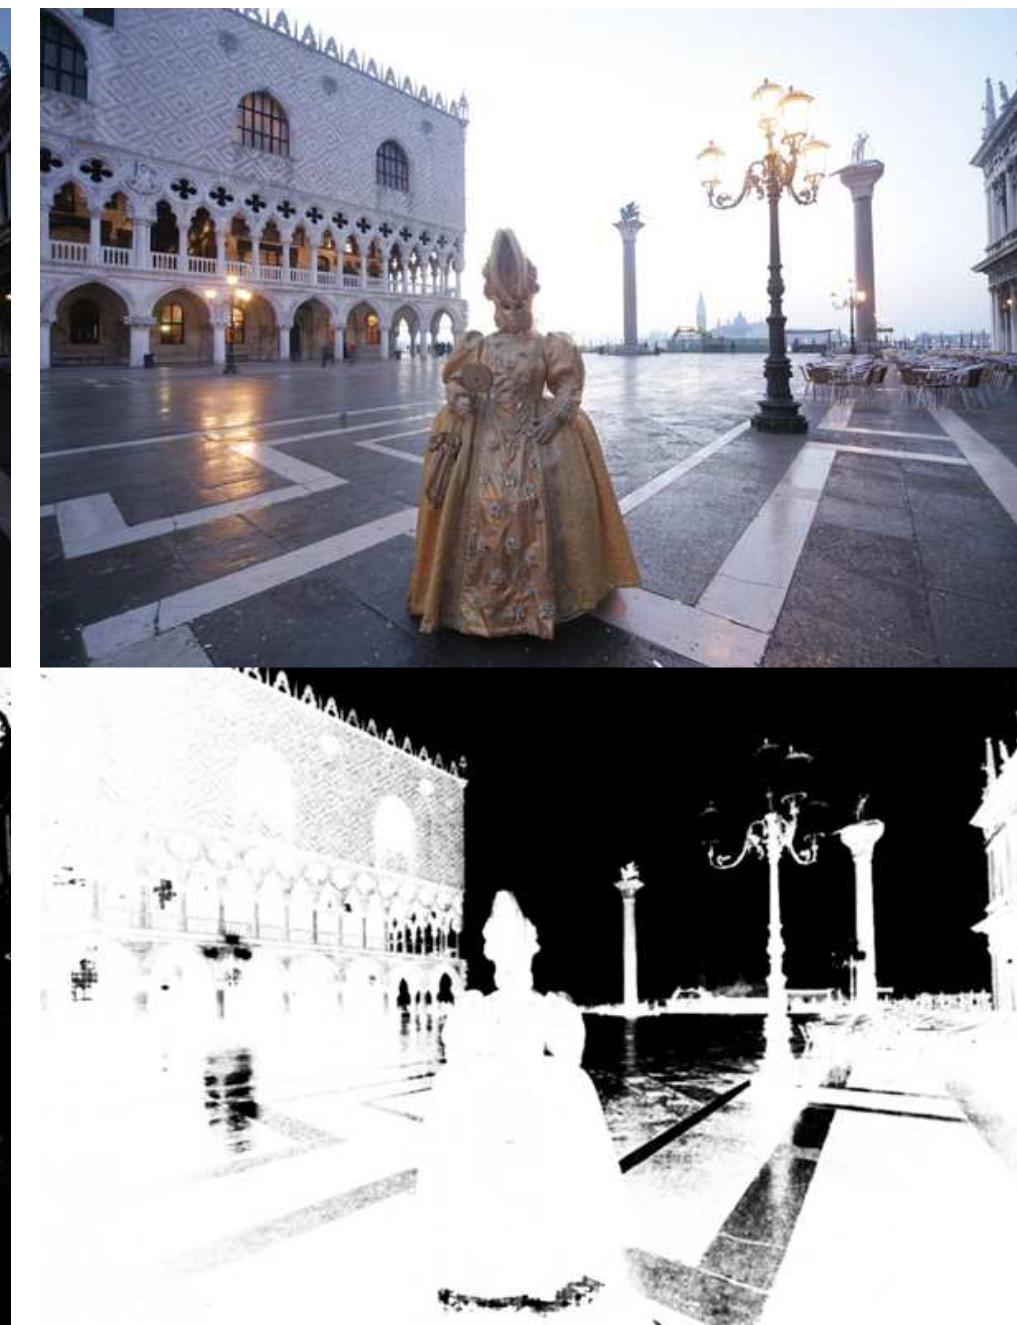
Represented in standard
color space (e.g., RGB)

Represented in
sensor-specific basis

Note: modern pipelines may perform more sophisticated (and non-linear) color transformations at this point. e.g., use a big lookup table to adjust certain tones (e.g., make sky-like tones bluer)

Local-tone adjustment

Weights



Improve picture's aesthetics by locally
adjusting contrast, boosting dark
regions, decreasing bright regions
(more details in later lecture)

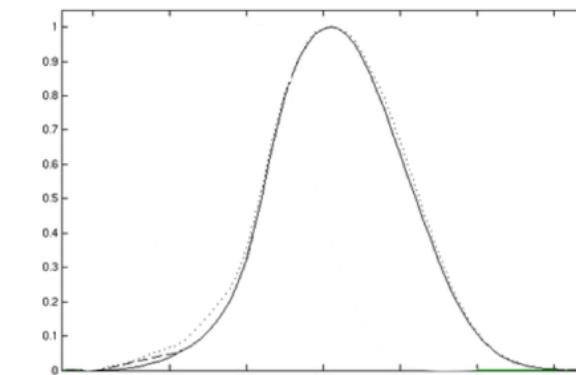
Combined image
(unique weights per pixel)



Lightness (perceived brightness) aka luma

$$\text{Lightness } (L^*) \xleftarrow{?} \text{Luminance } (Y) = \int_{\lambda} \text{Spectral sensitivity of eye (eye's response curve)} * \text{Radiance (energy spectrum from scene)}$$

(Perceived by brain) (Response of eye)



The graph shows a bell-shaped curve representing the spectral sensitivity of the human eye. The x-axis is labeled with values 400, 450, 500, 550, 600, 650, and 700, representing nanometers. The y-axis ranges from 0 to 1. The curve peaks at approximately 550 nm with a value of 1.0, and tapers off towards the ends of the visible spectrum.

Dark adapted eye: $L^* \propto Y^{0.4}$

Bright adapted eye: $L^* \propto Y^{0.5}$

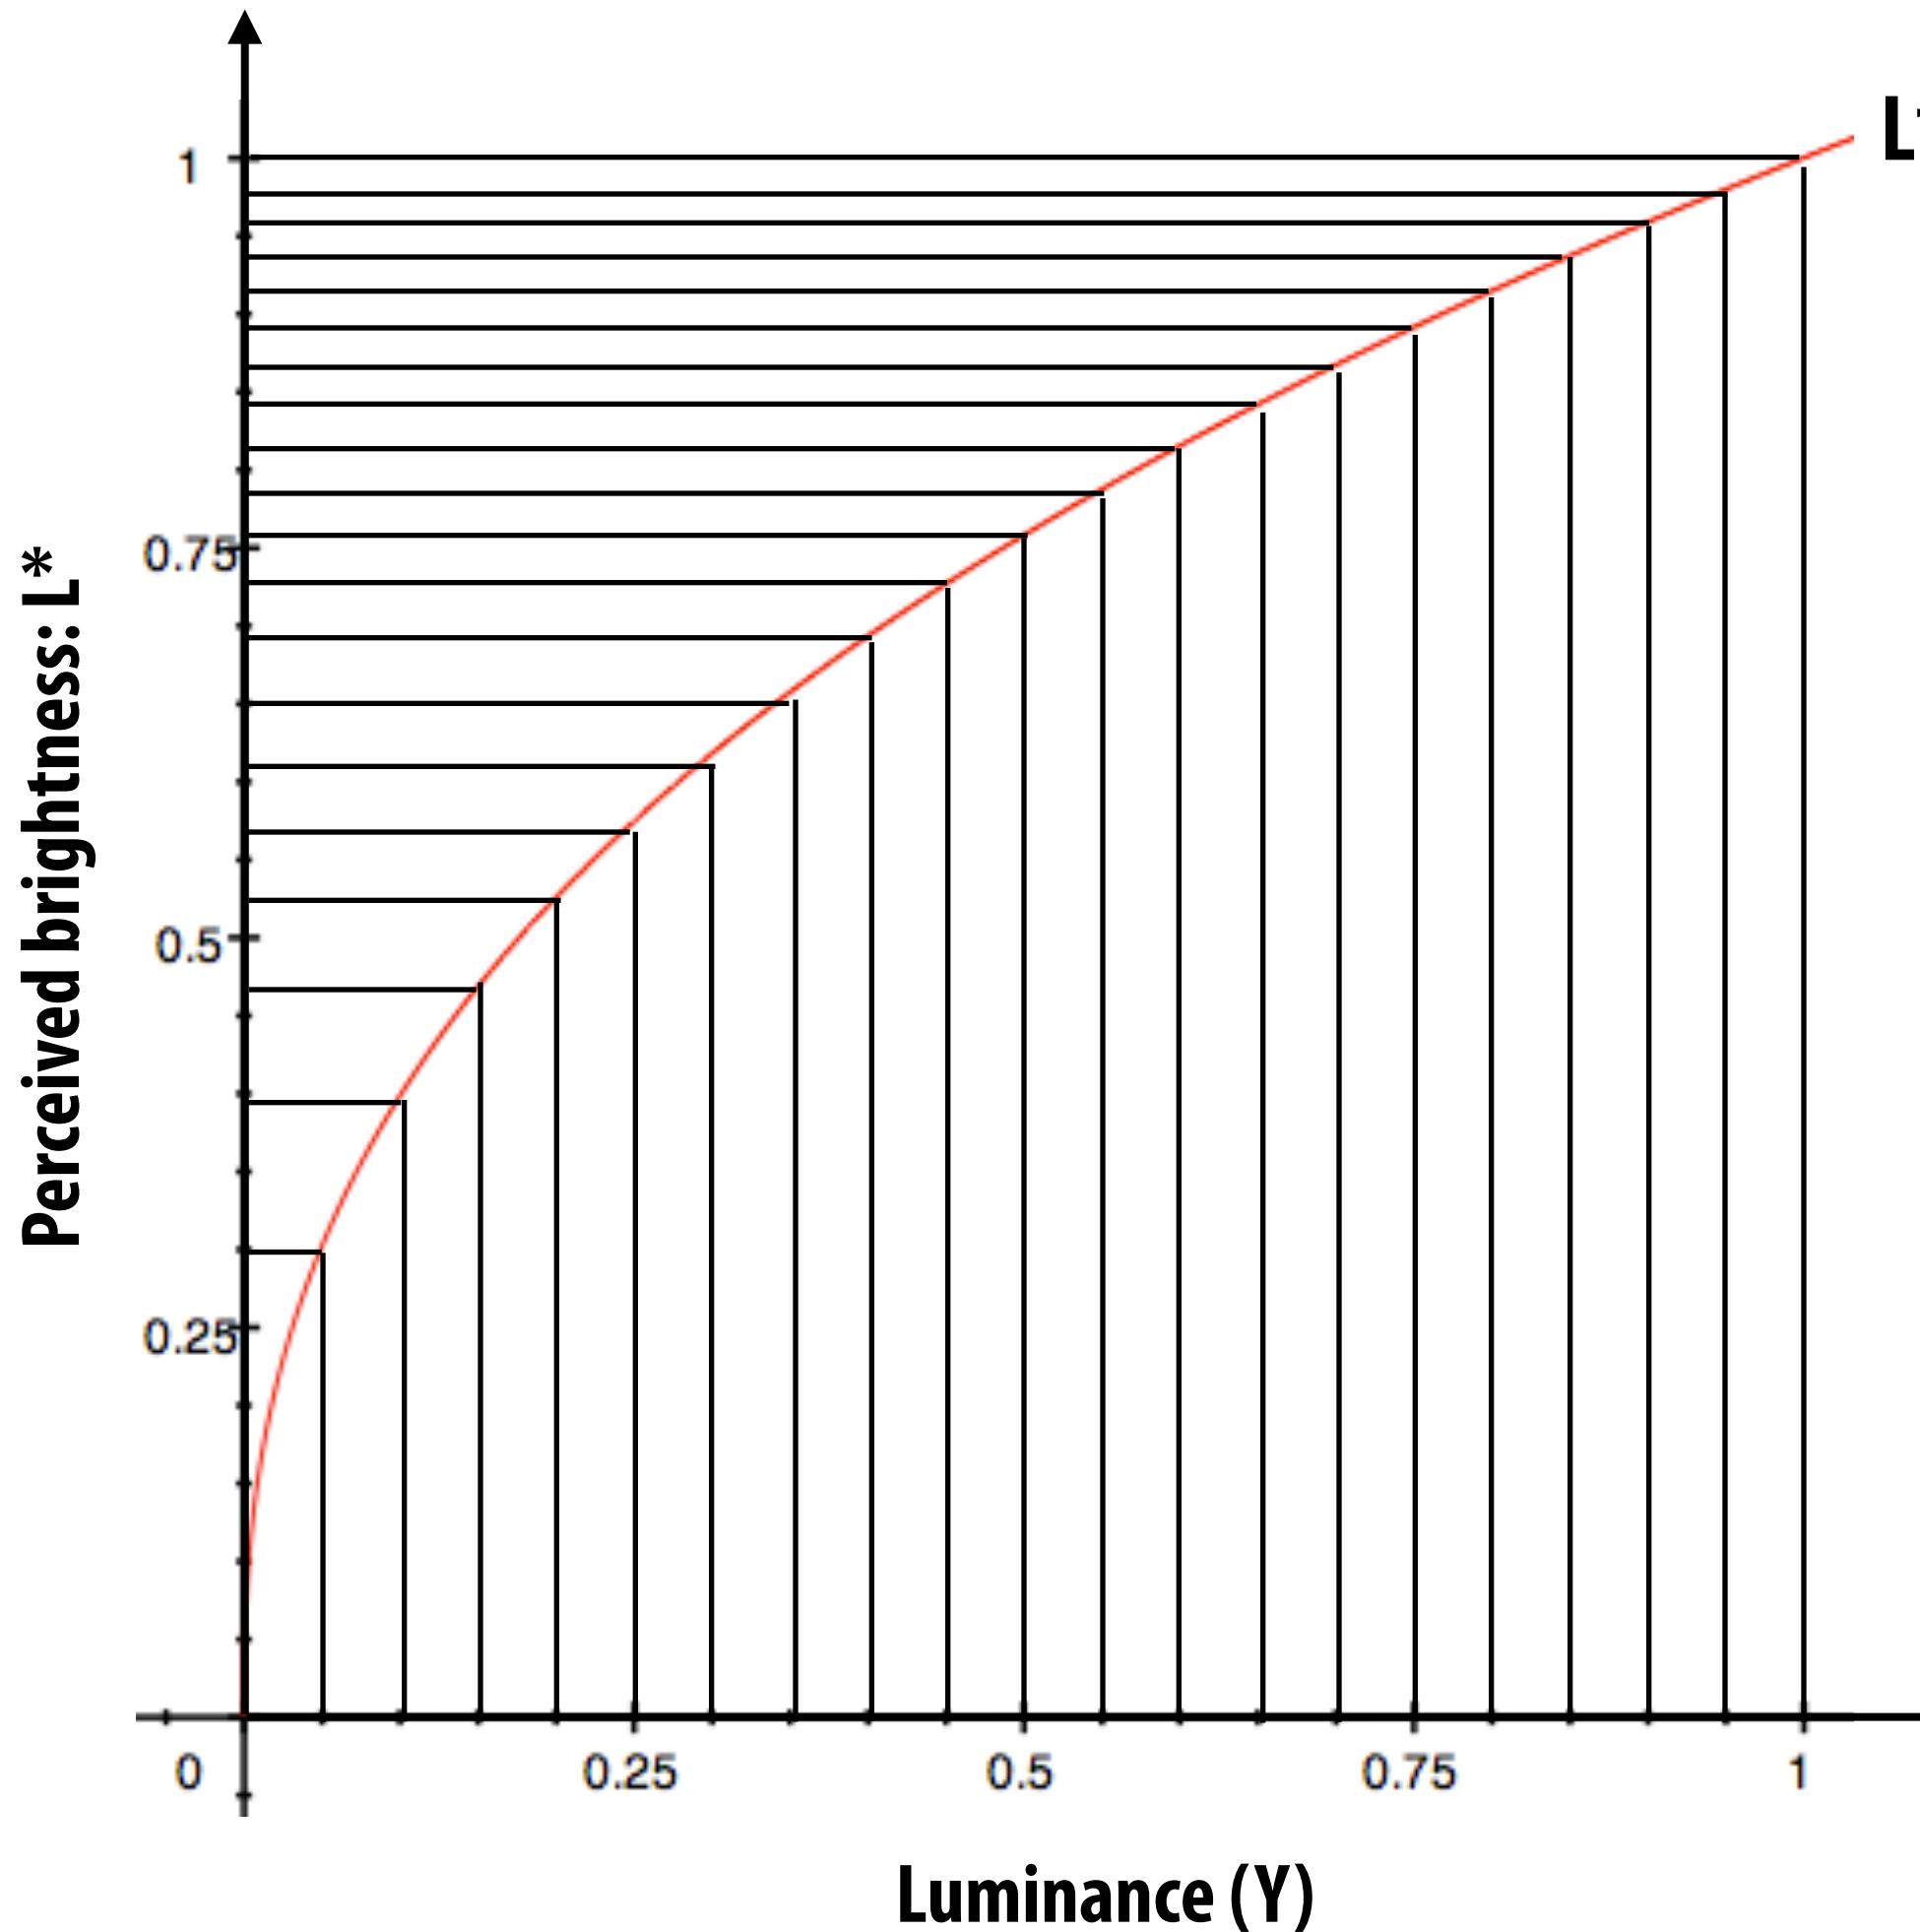
In a dark room, you turn on a light with luminance: Y_1

You turn on a second light that is identical to the first. Total output is now: $Y_2 = 2Y_1$

Total output appears $2^{0.4} = 1.319$ times brighter to dark-adapted human

Note: Lightness (L^*) is often referred to as luma (Y')

Consider an image with pixel values encoding luminance (linear in energy hitting sensor)



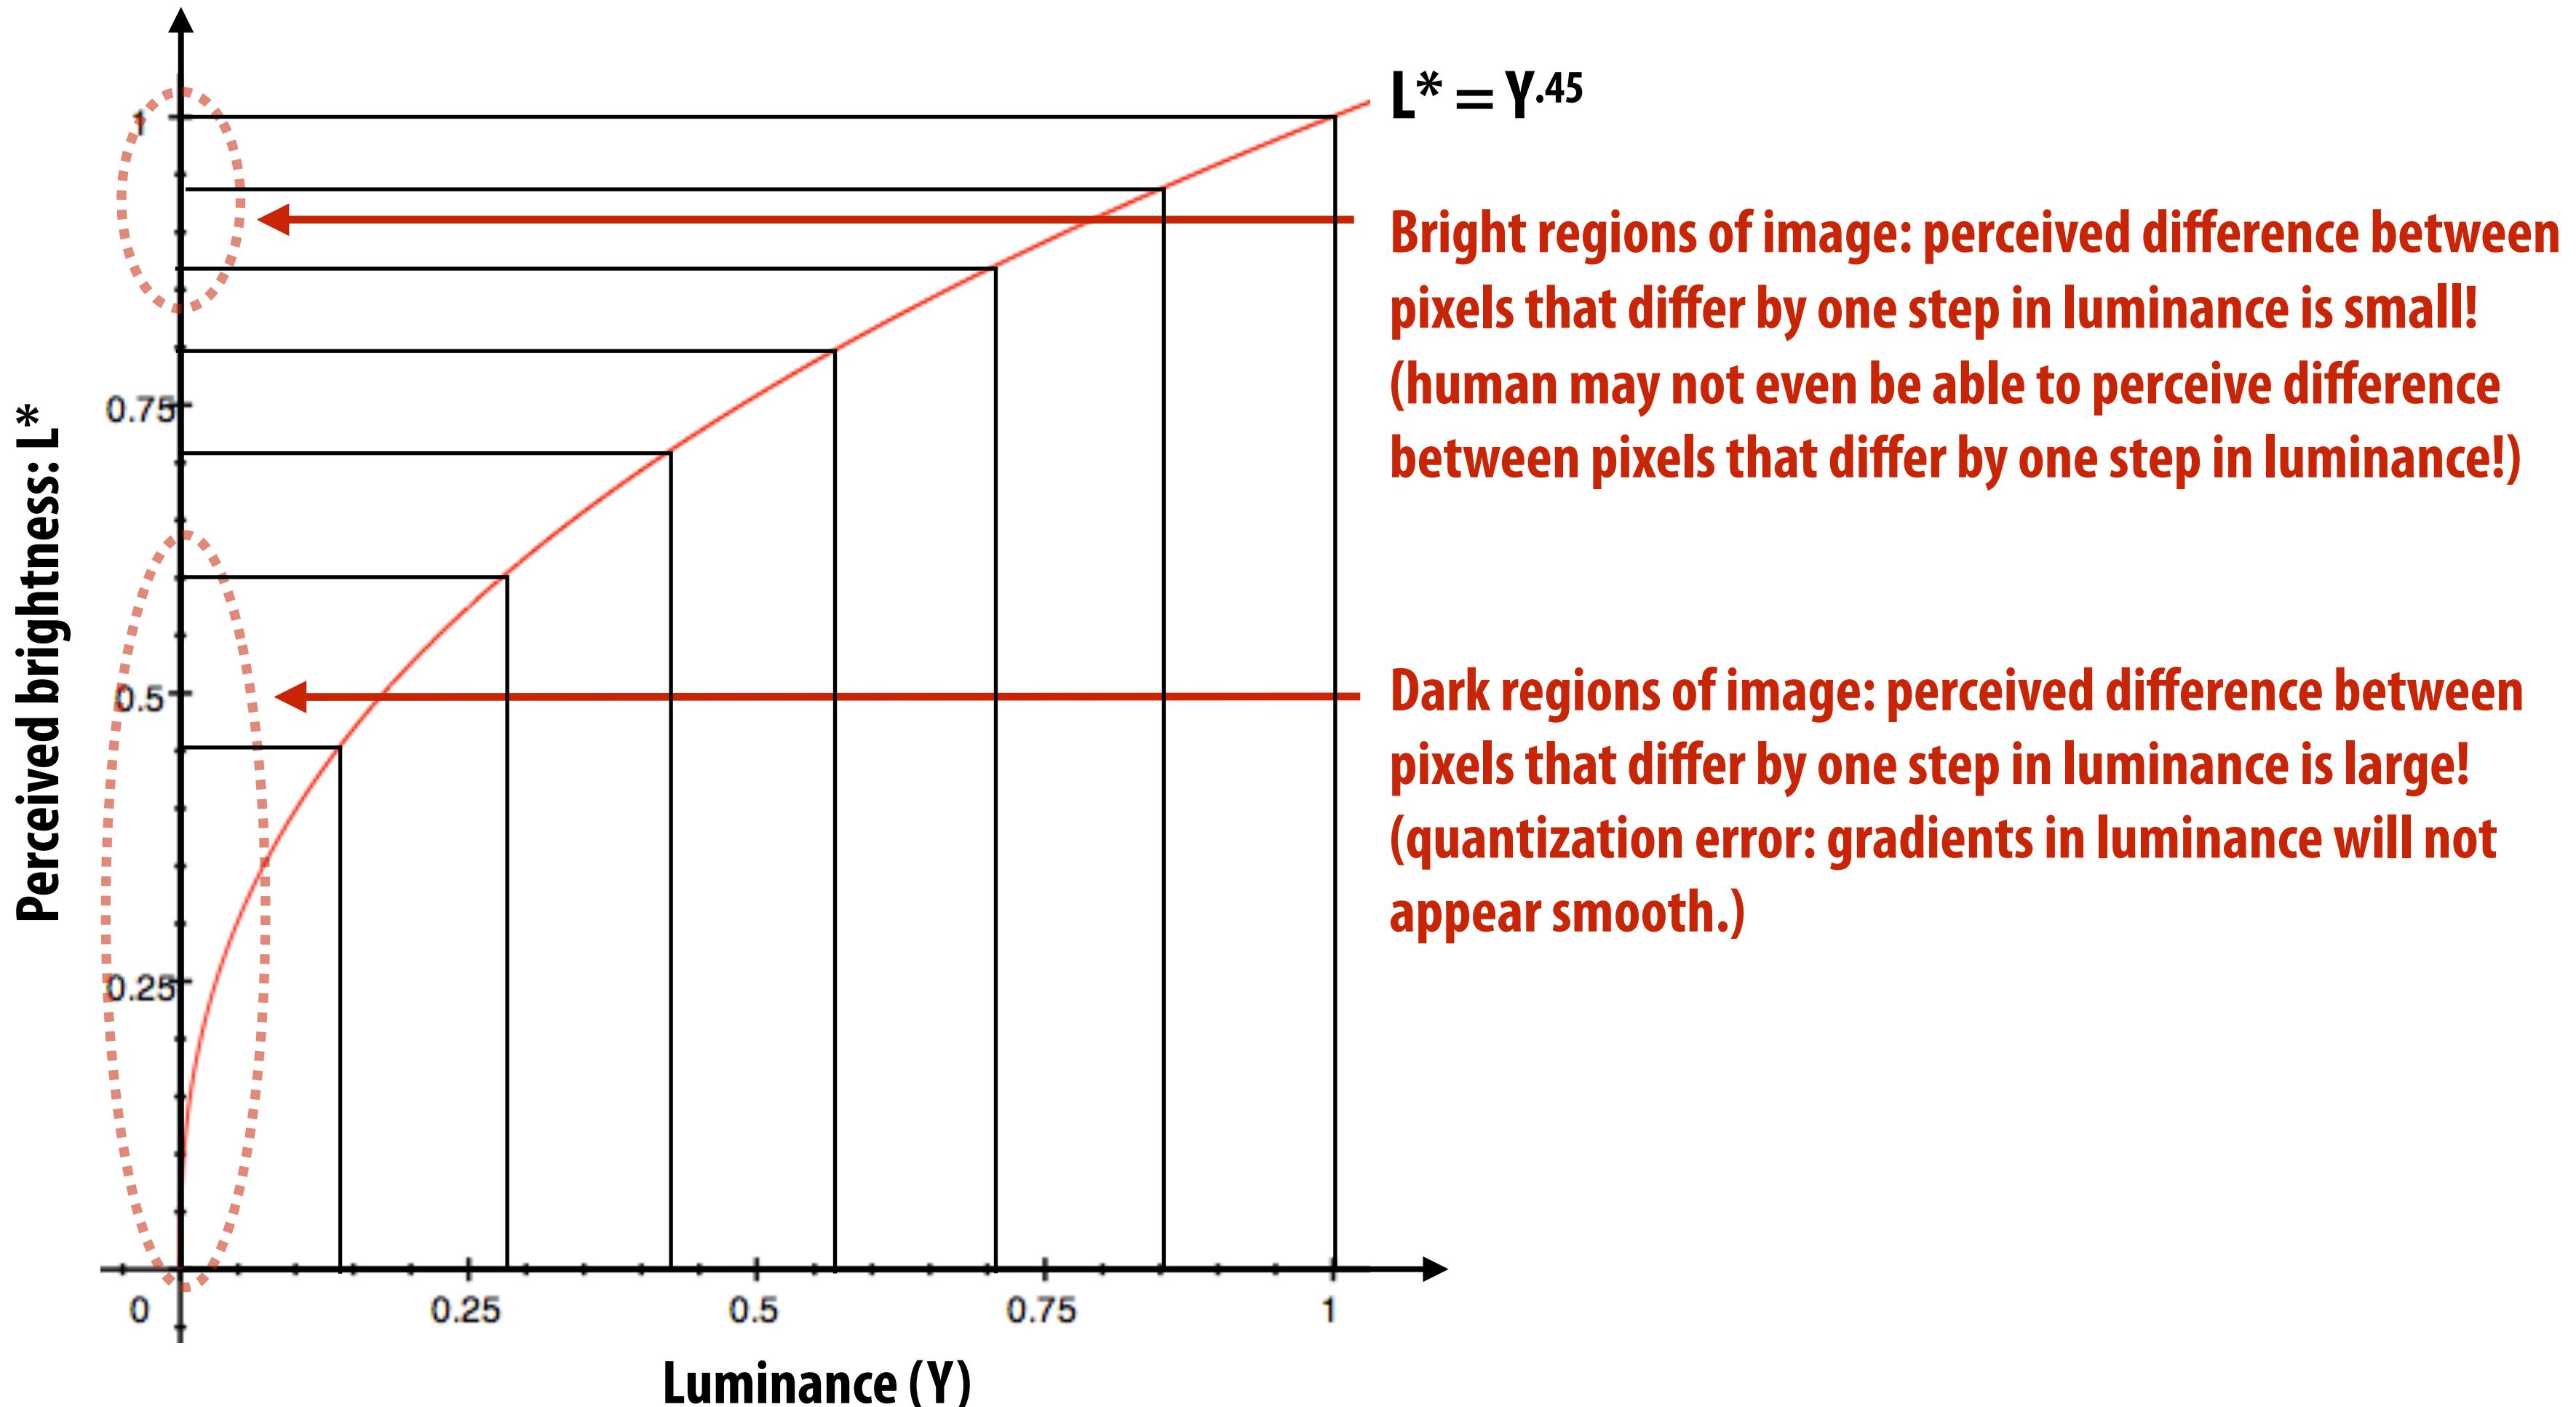
**Consider 12-bit sensor pixel:
Can represent 4096 unique luminance values
in output image**

**Values are ~ linear in luminance since they
represent the sensor's response**

Problem: quantization error

Many common image formats store 8 bits per channel (256 unique values)

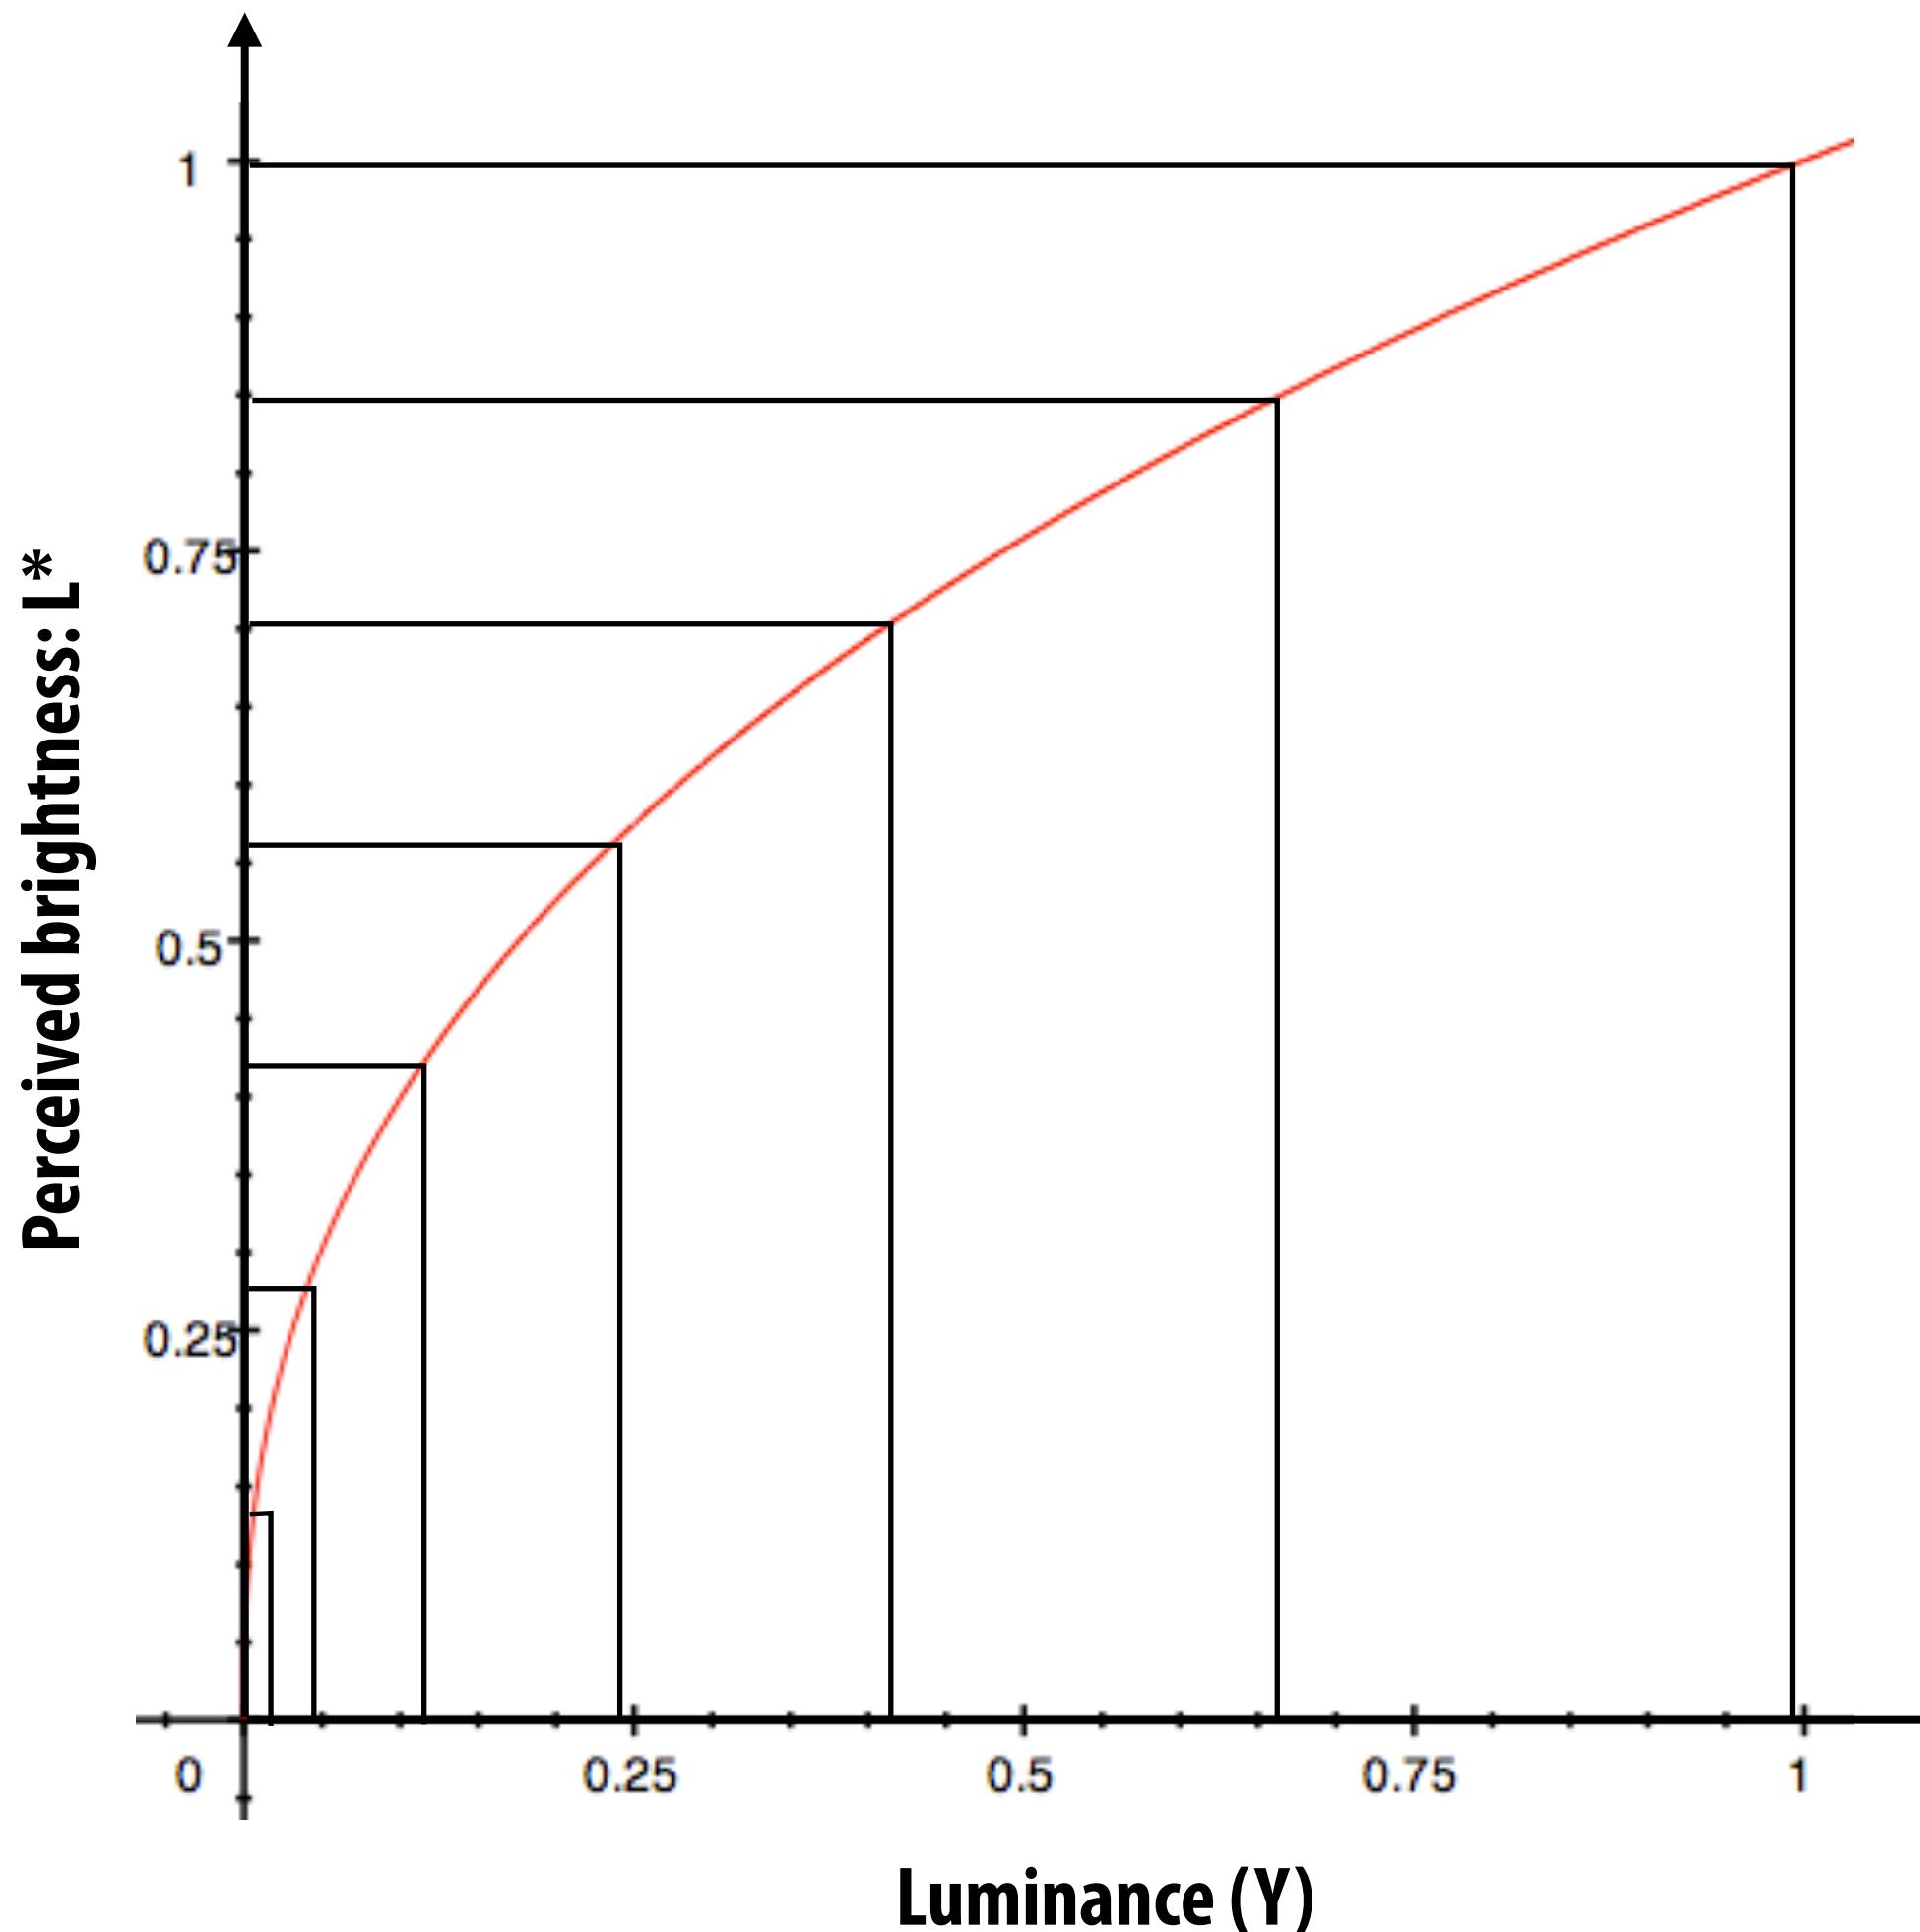
Insufficient precision to represent brightness in darker regions of image



Rule of thumb: human eye cannot differentiate <1% differences in luminance

Store lightness, not luminance

Idea: distribute representable pixel values evenly with respect to perceived brightness, not evenly in luminance (make more efficient use of available bits)



Solution: pixel stores $Y^{0.45}$

Must compute $(\text{pixel_value})^{2.2}$ prior to display on LCD

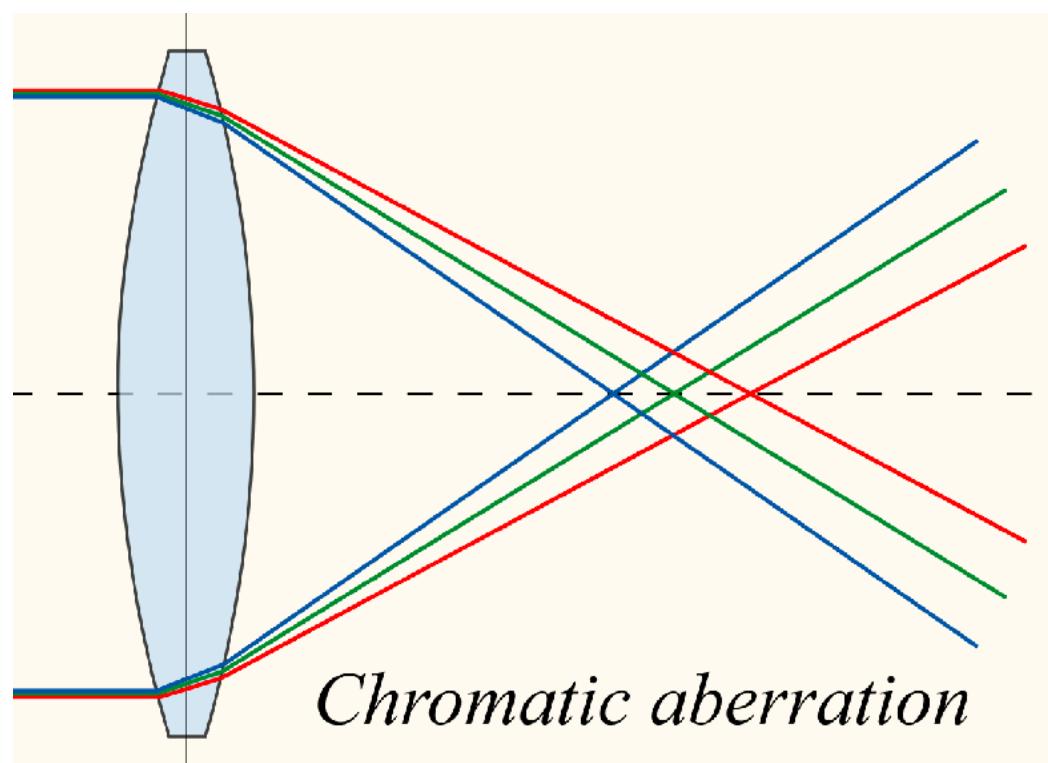
Warning: must take caution with subsequent pixel processing operations once pixels are encoded in a space that is not linear in luminance.

e.g., When adding images should you add pixel values that are encoded as lightness or as luminance?

Summary: simplified image processing pipeline

- **Correct for sensor bias (using measurements of optically black pixels)**
- **Correct pixel defects**
- **Vignetting compensation**
- **Dark-frame subtract (optional)**
- **White balance**
 - 12-bits per pixel
 - 1 intensity per pixel
 - Pixel values linear in energy
- **Demosaic**
- **Denoise / sharpen, etc.**
- **Color Space Conversion**
 - 3x12-bits per pixel
 - RGB intensity per pixel
 - Pixel values linear in energy
- **Local tone mapping**
- **Gamma Correction (non-linear mapping)**
- **Color Space Conversion (Y'CbCr)**
 - 3x8-bits per pixel
 - Pixel values **perceptually** linear

Chromatic aberration correction



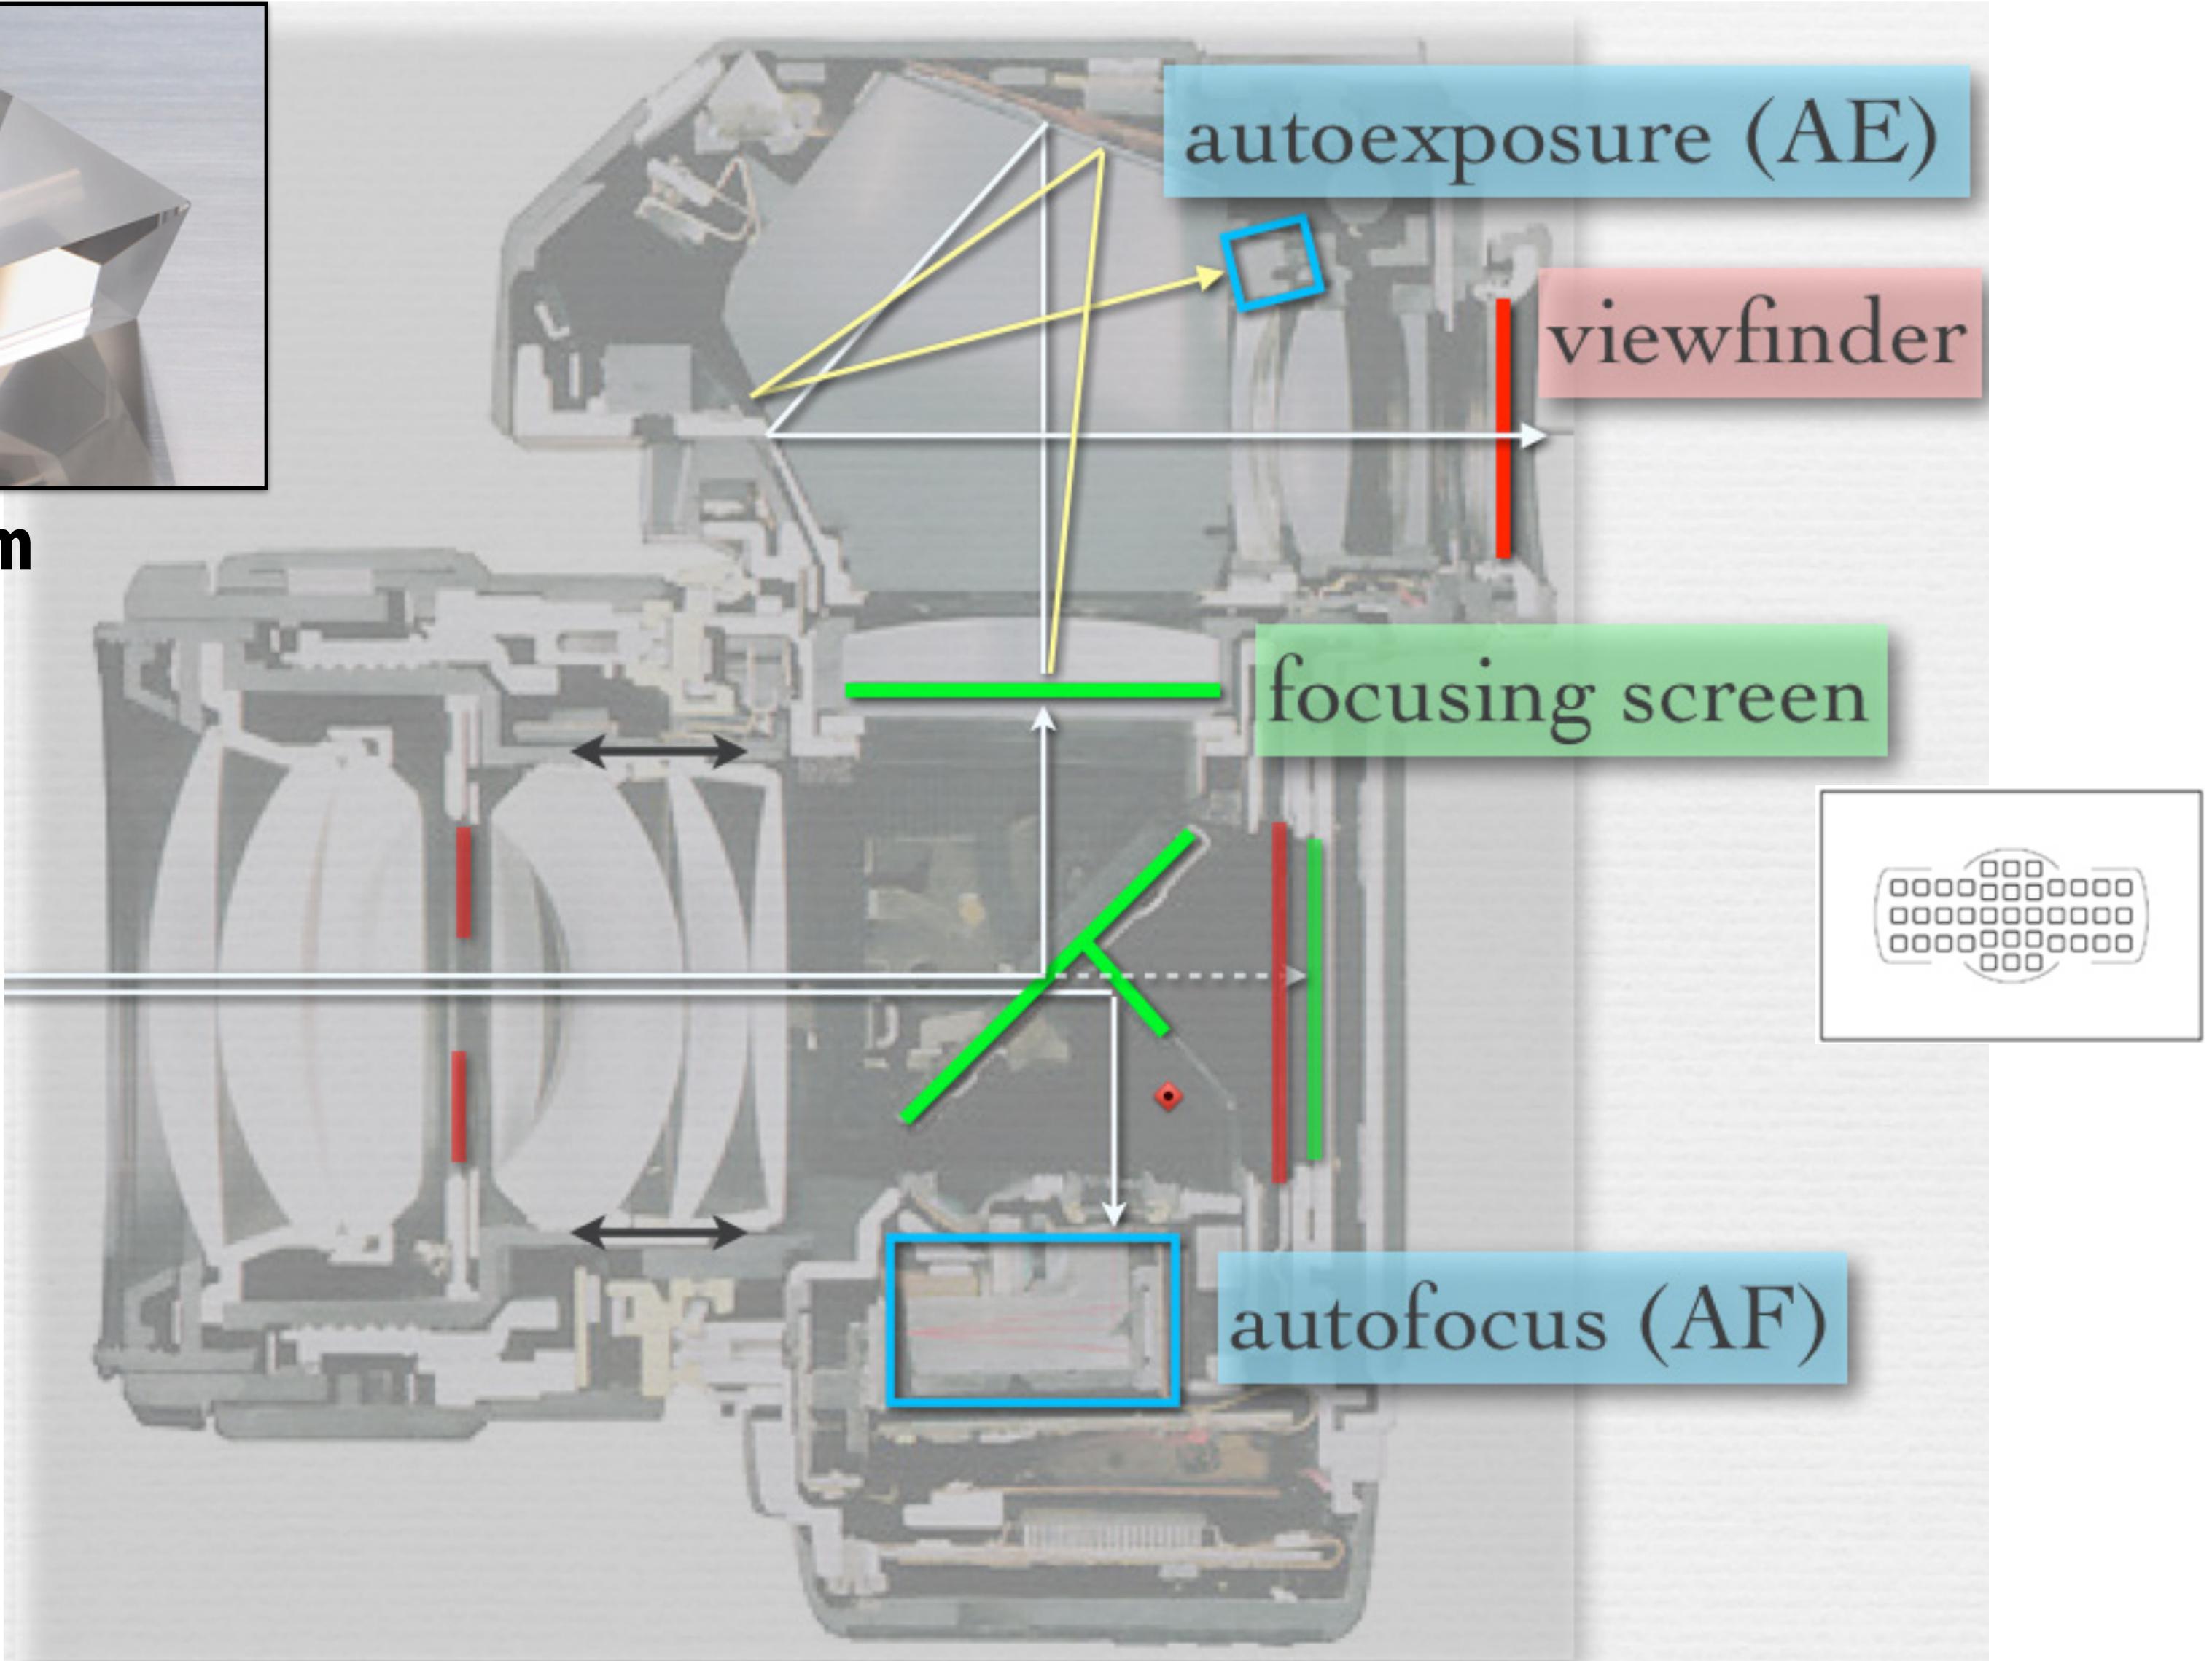
Autofocus demos

- **Phase-detection auto focus**
 - Common in SLRs
- **Contrast-detection auto focus**
 - Point-and-shoots, smart-phone cameras

SLR Camera

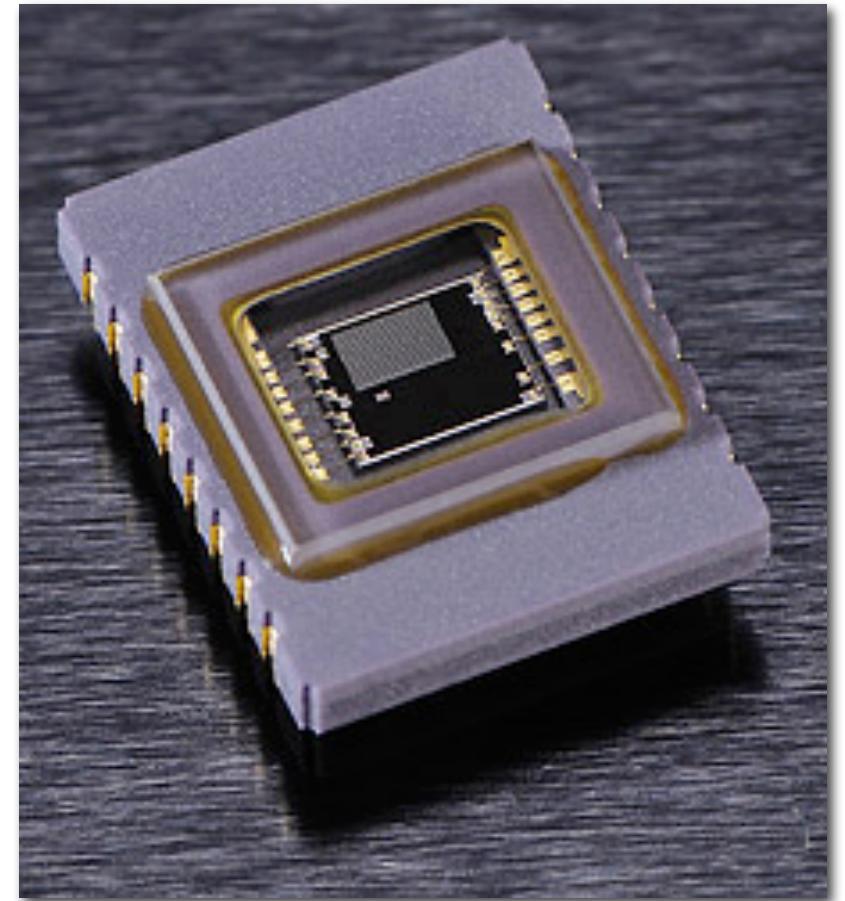
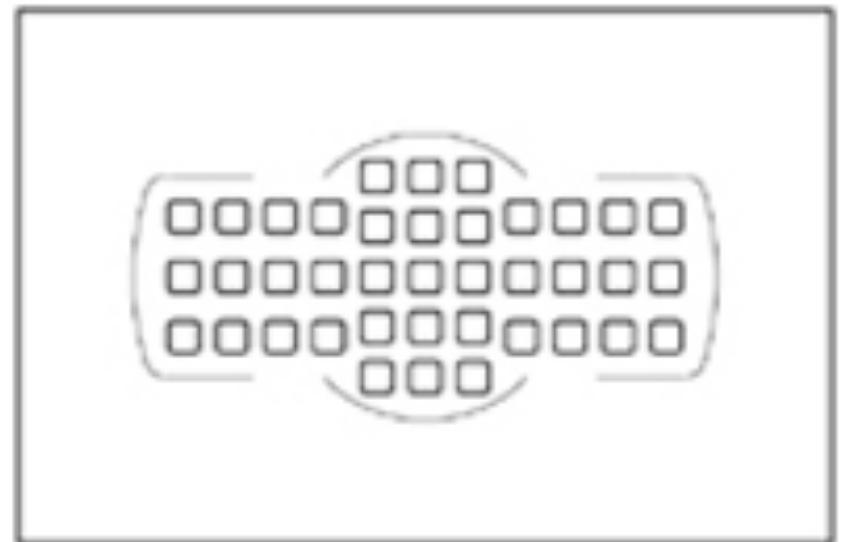


Pentaprism

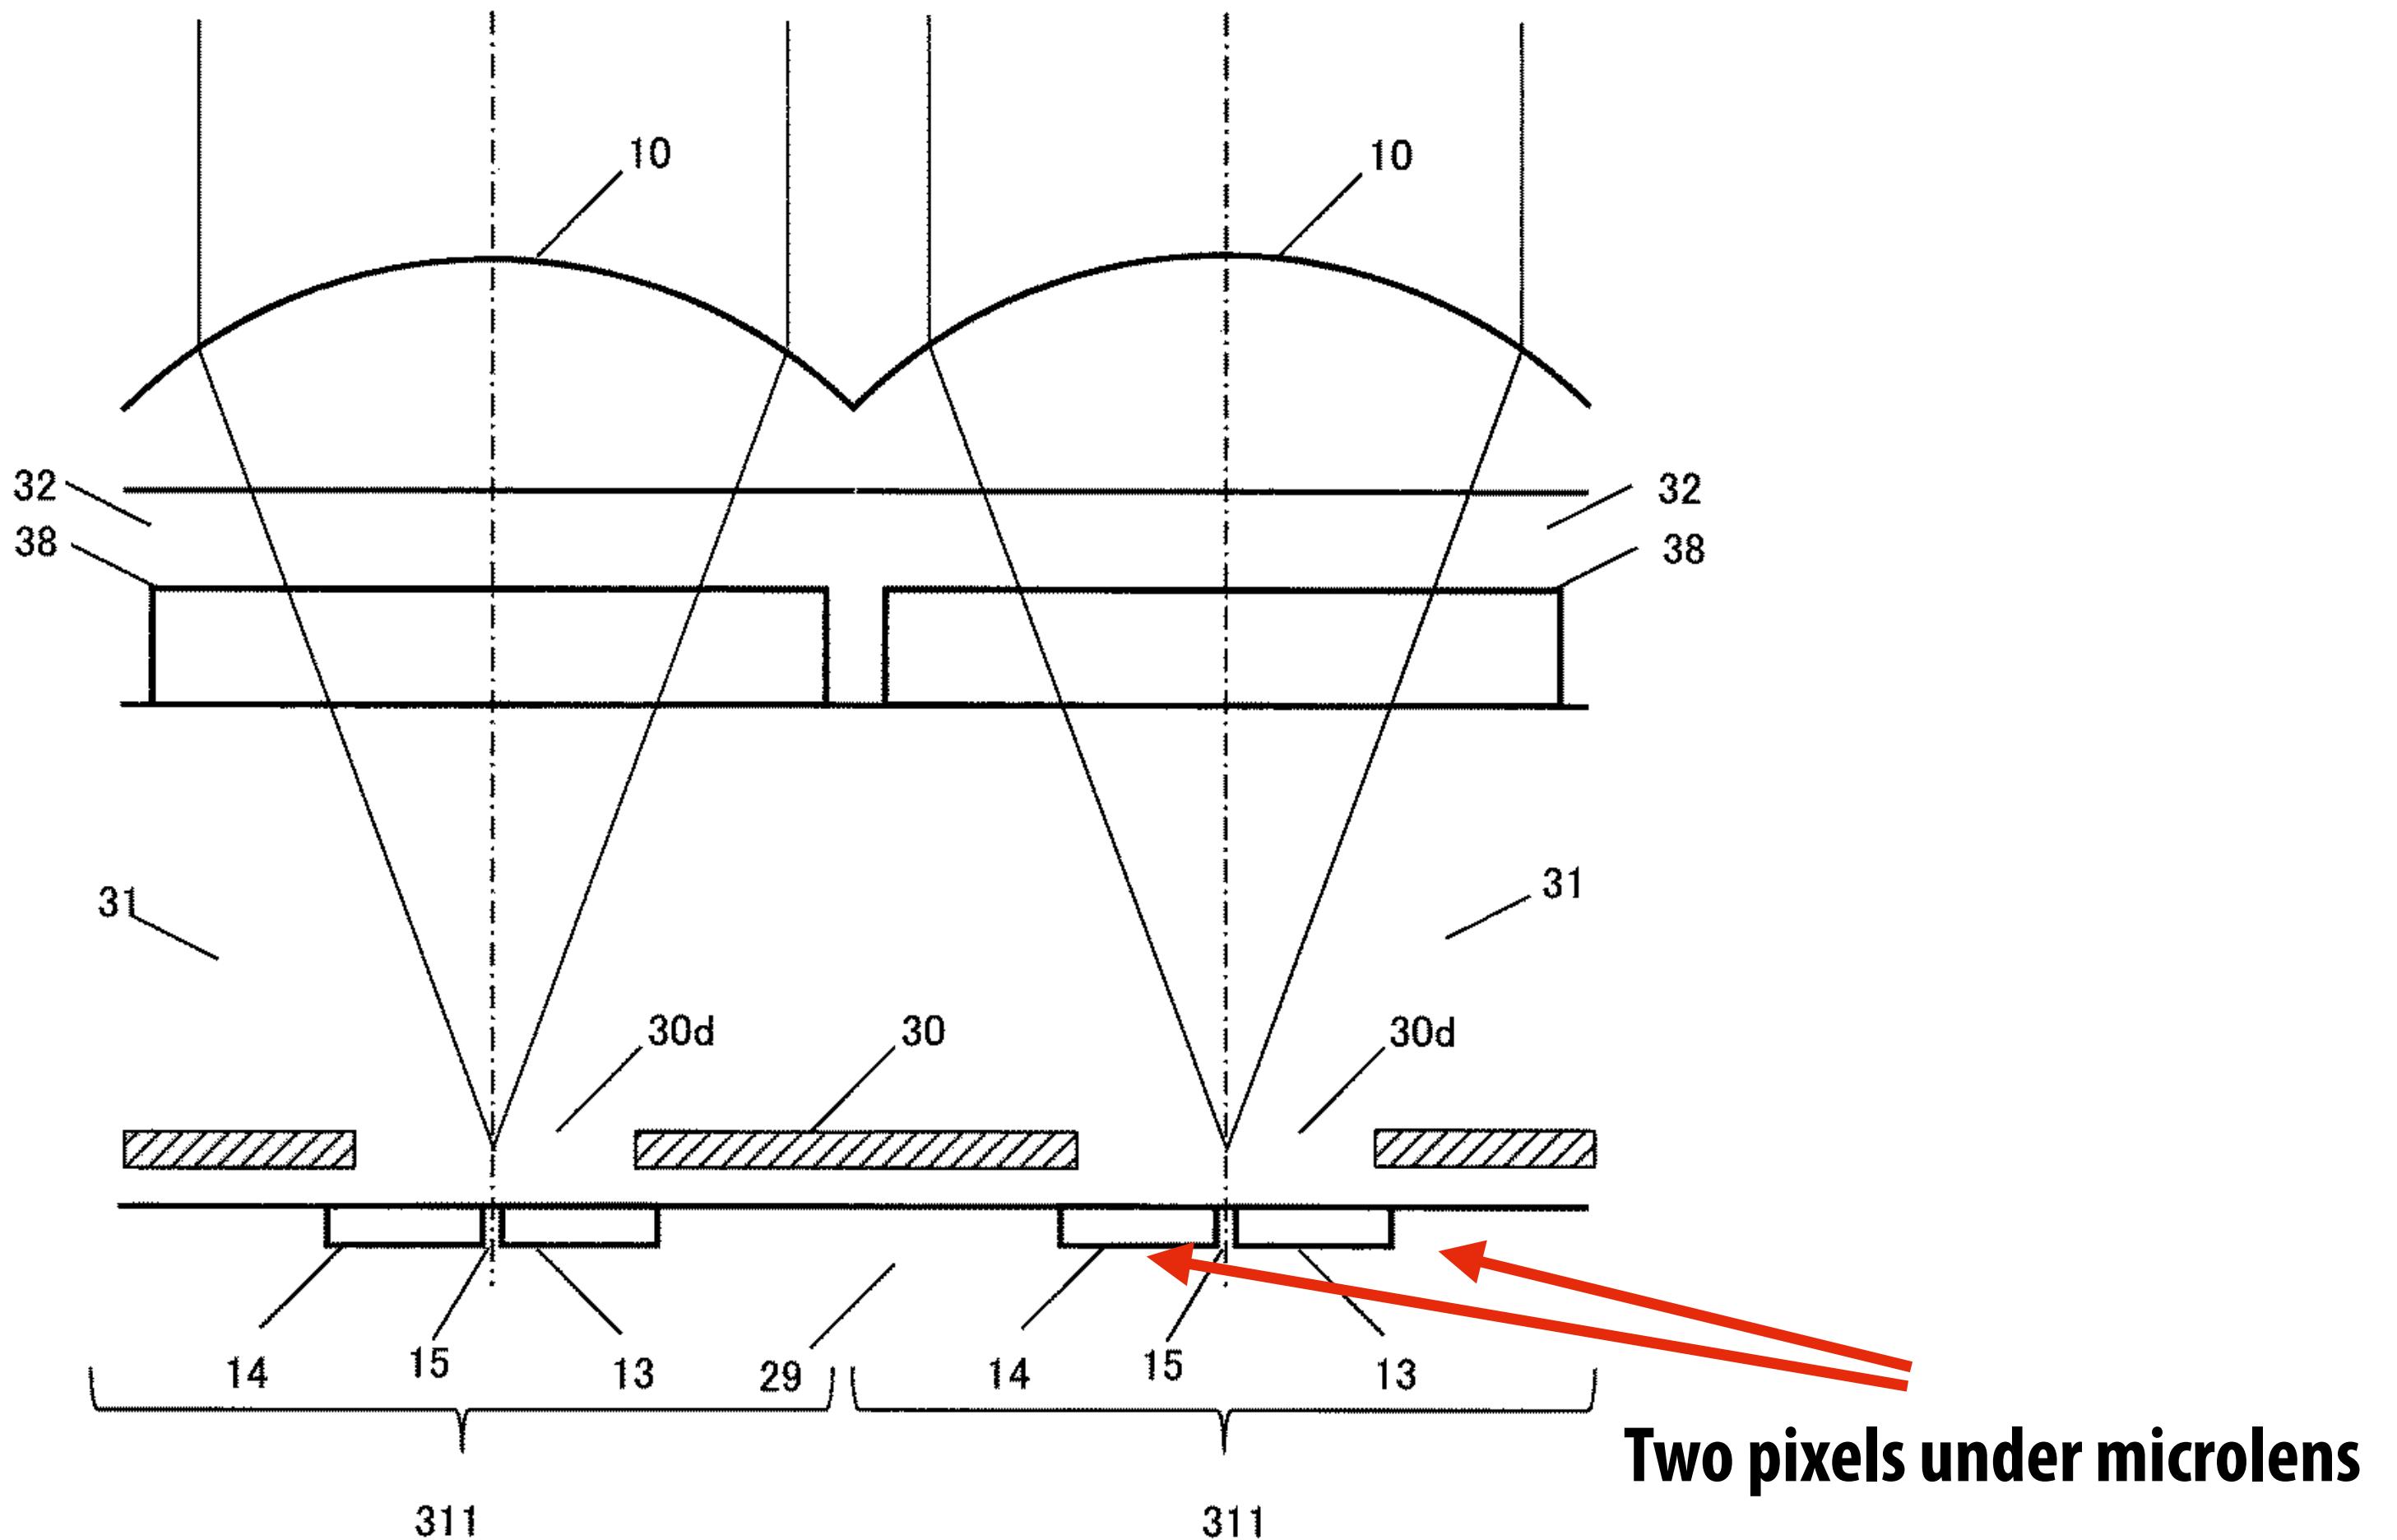


Nikon D7000

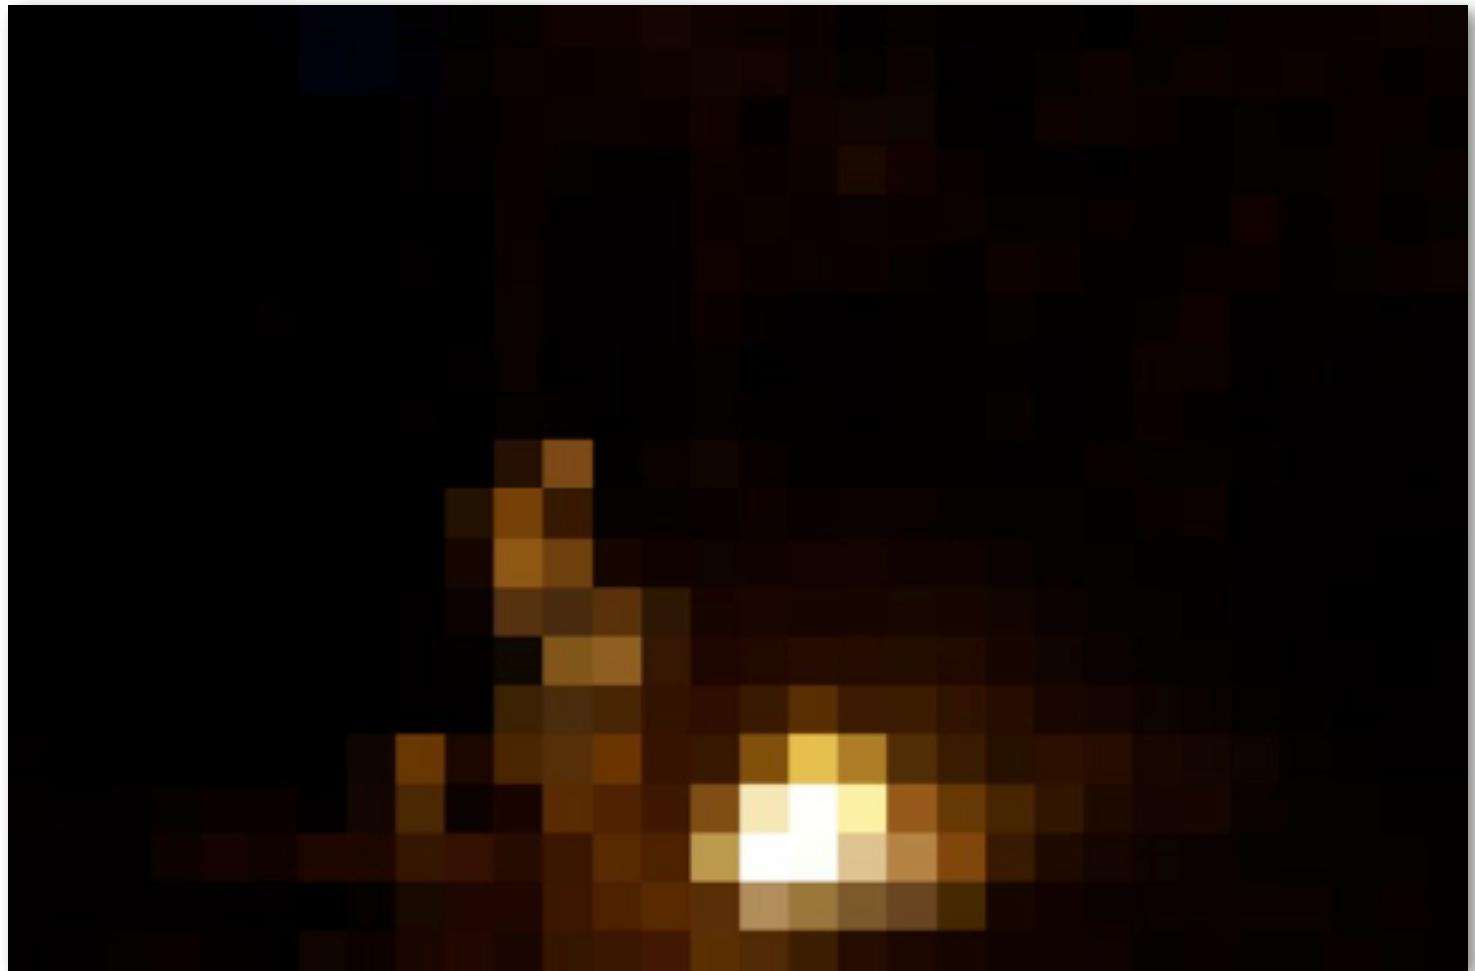
- **Auto-focus sensor: 39 regions**
- **Metering sensor: 2K pixels**
 - Auto-exposure
 - Auto-white-balance
 - Subject tracking to aid focus (predicts movement)
- **Shutter lag ~ 50ms**



Split pixel sensor



Auto exposure



Low resolution metering sensor capture

**Metering sensor pixels are large
(higher dynamic range than main sensor)**

How do we set exposure?

What if a camera doesn't have a separate metering sensor?

AF/AE summary

- **DSLRs have additional sensing/processing hardware to assist with the “3A’s” (auto-focus, auto-exposure, auto-white-balance)**
 - Phase-detection AF: optical system directs light to AF sensor
 - Example: Nikon metering sensor: large pixels to avoid over-saturation
- **Point-and-shoots/smartphone cameras make these measurements by performing image processing operations on data from the main sensor**
 - Contrast-detection AF: search for lens position that produces large image gradients
 - Exposure metering: if pixels are saturating, meter again with lower exposure
- **In general, implementing AF/AE/AWB is an image understanding problem (“computer vision”)**
 - Understand the scene well enough to set the camera’s image capture and image processing parameters to best approximate the image a human would perceive
 - As processing/sensing capability increases, algorithms are becoming more sophisticated

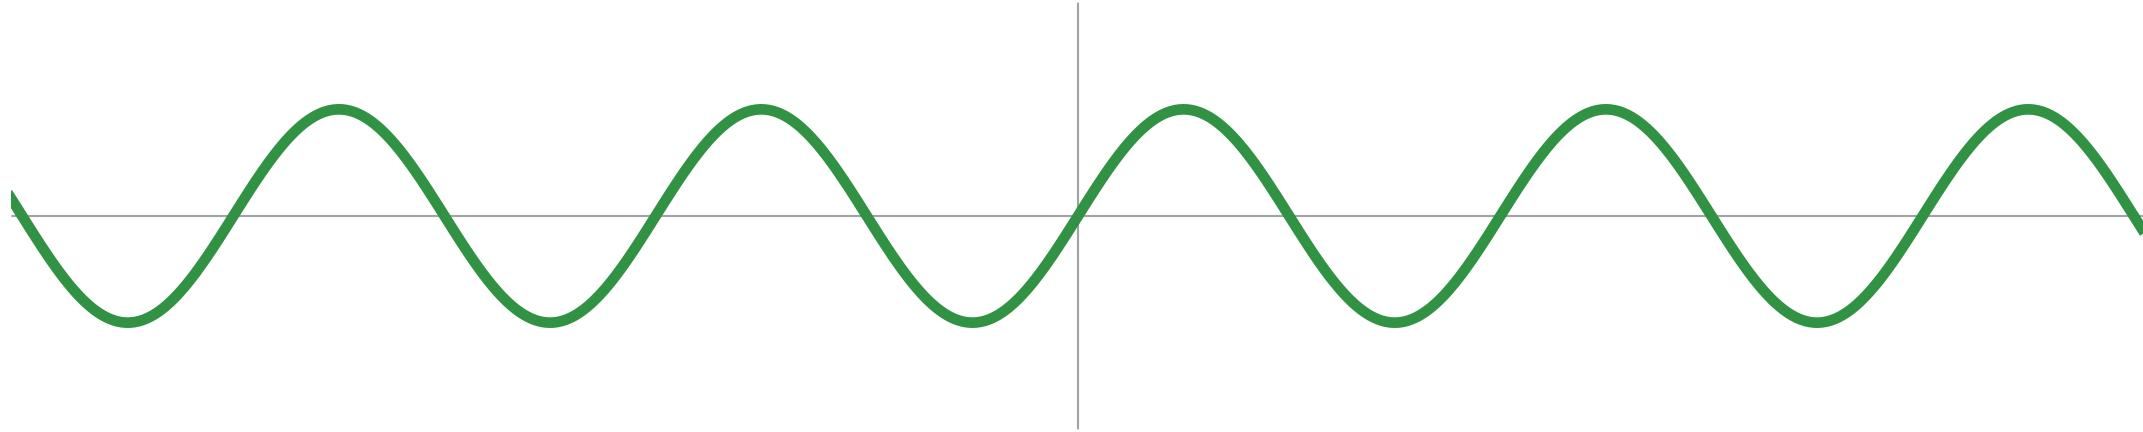
Takeways

- The values of pixels in picture you see on screen are quite different than the values output by the photosensor in a modern digital camera.
- The sequence of operations we discussed today is carried out at high frame rates by special purpose hardware in most cameras today
- In the coming lectures we'll discuss more advanced image processing operations in use in modern camera pipelines
 - Local contrast enhancement, advanced denoising, high-dynamic range imaging, etc.
 - Growing sophistication and diversity of techniques suggests that image-signal processors will likely become increasingly programmable (e.g., Qualcomm's Hexagon DSP)

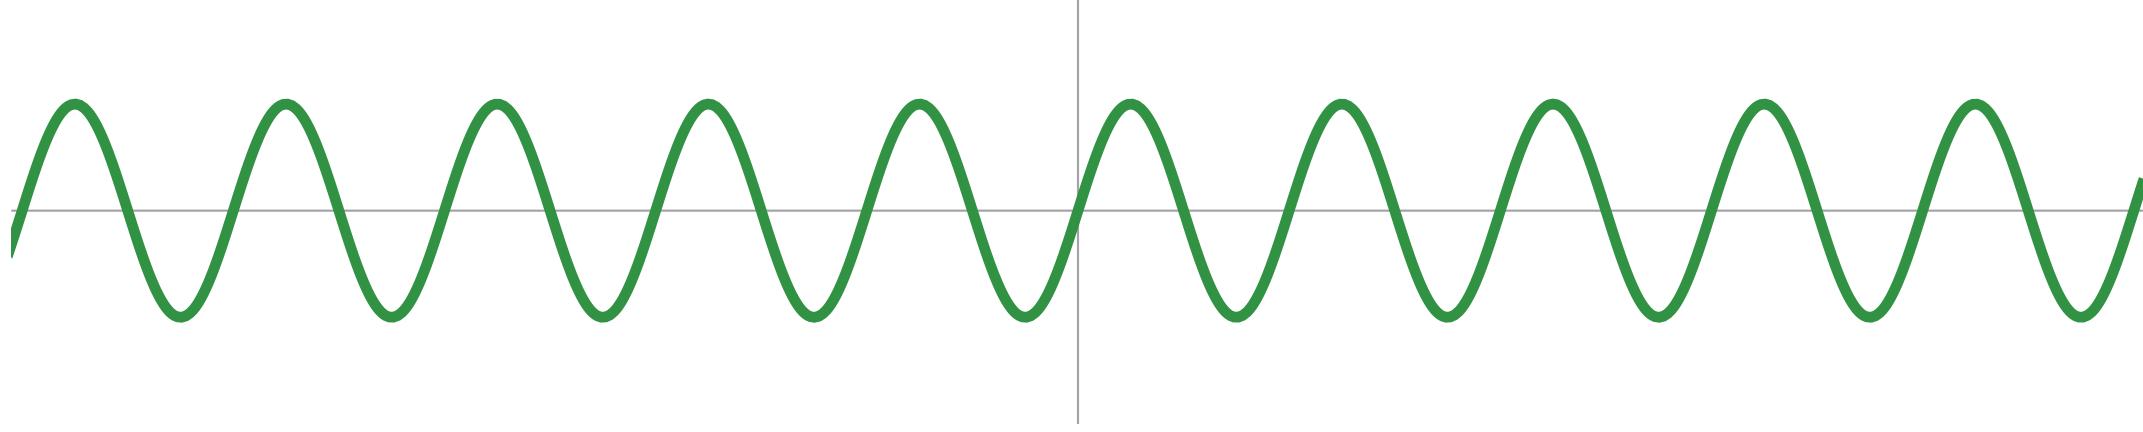
Frequency interpretation of images

Representing sound as a superposition of frequencies

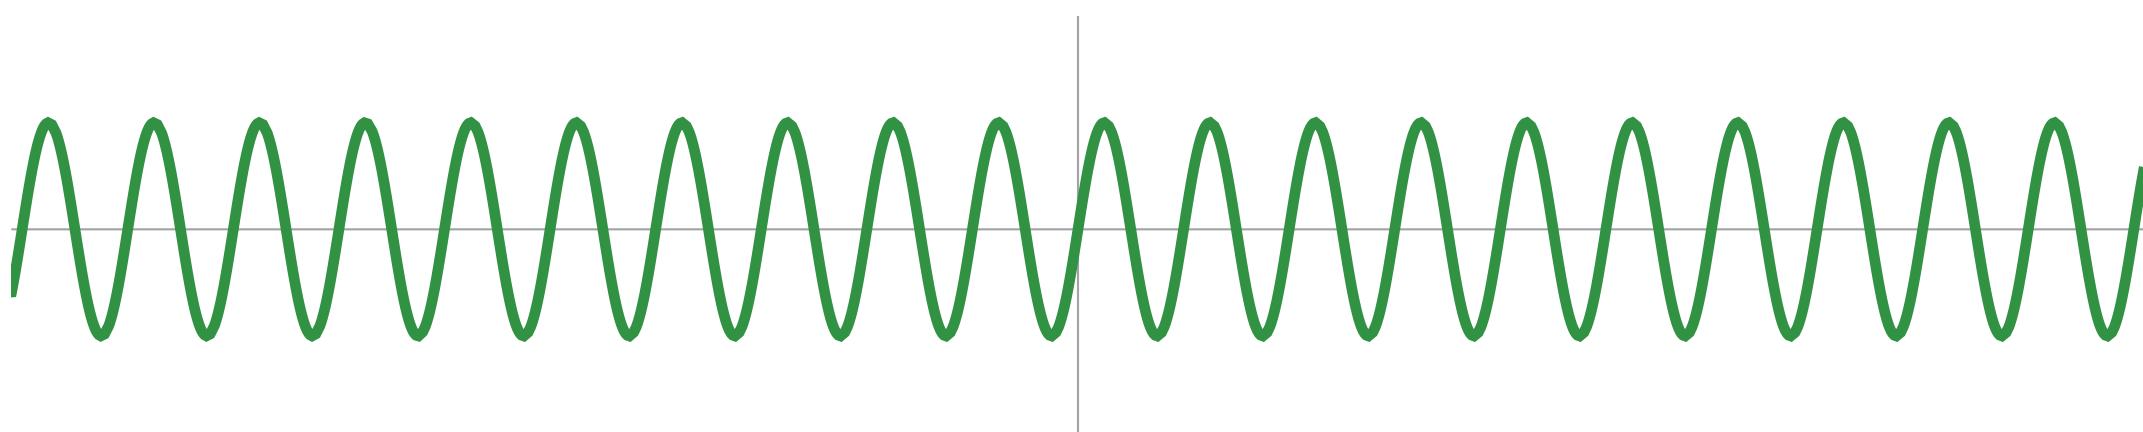
$$f_1(x) = \sin(\pi x)$$



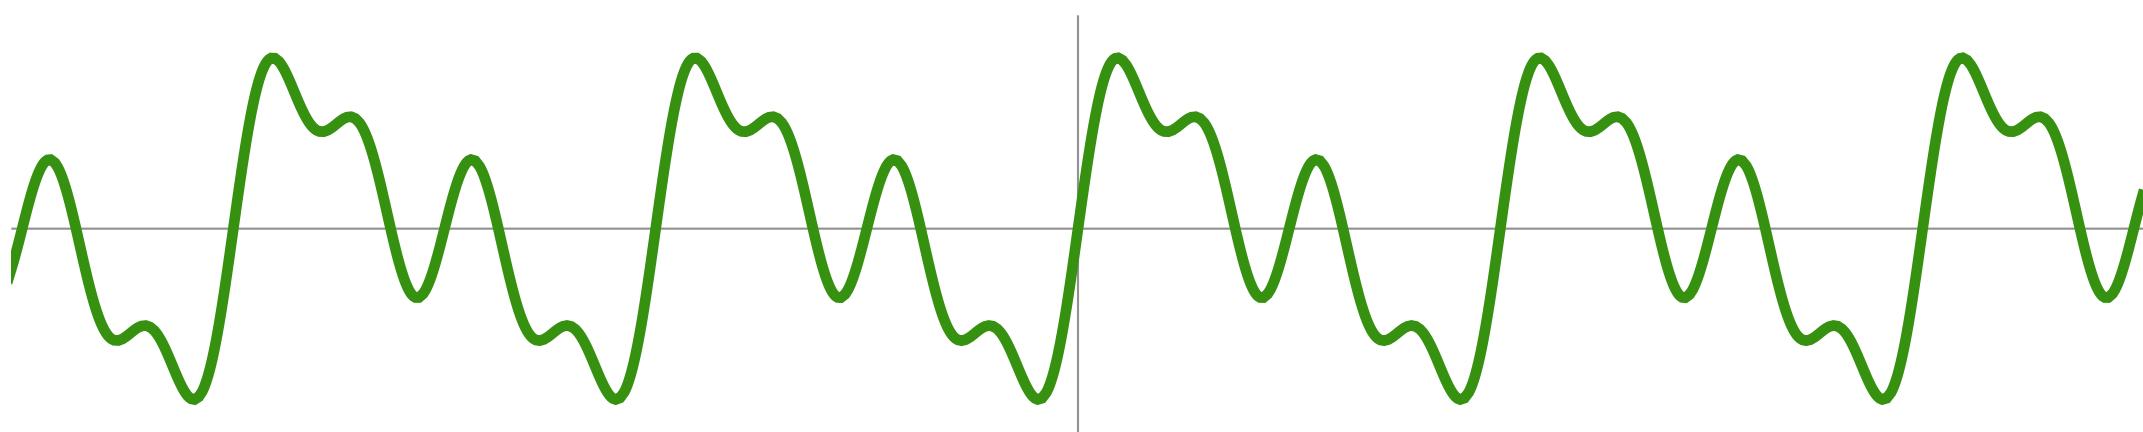
$$f_2(x) = \sin(2\pi x)$$



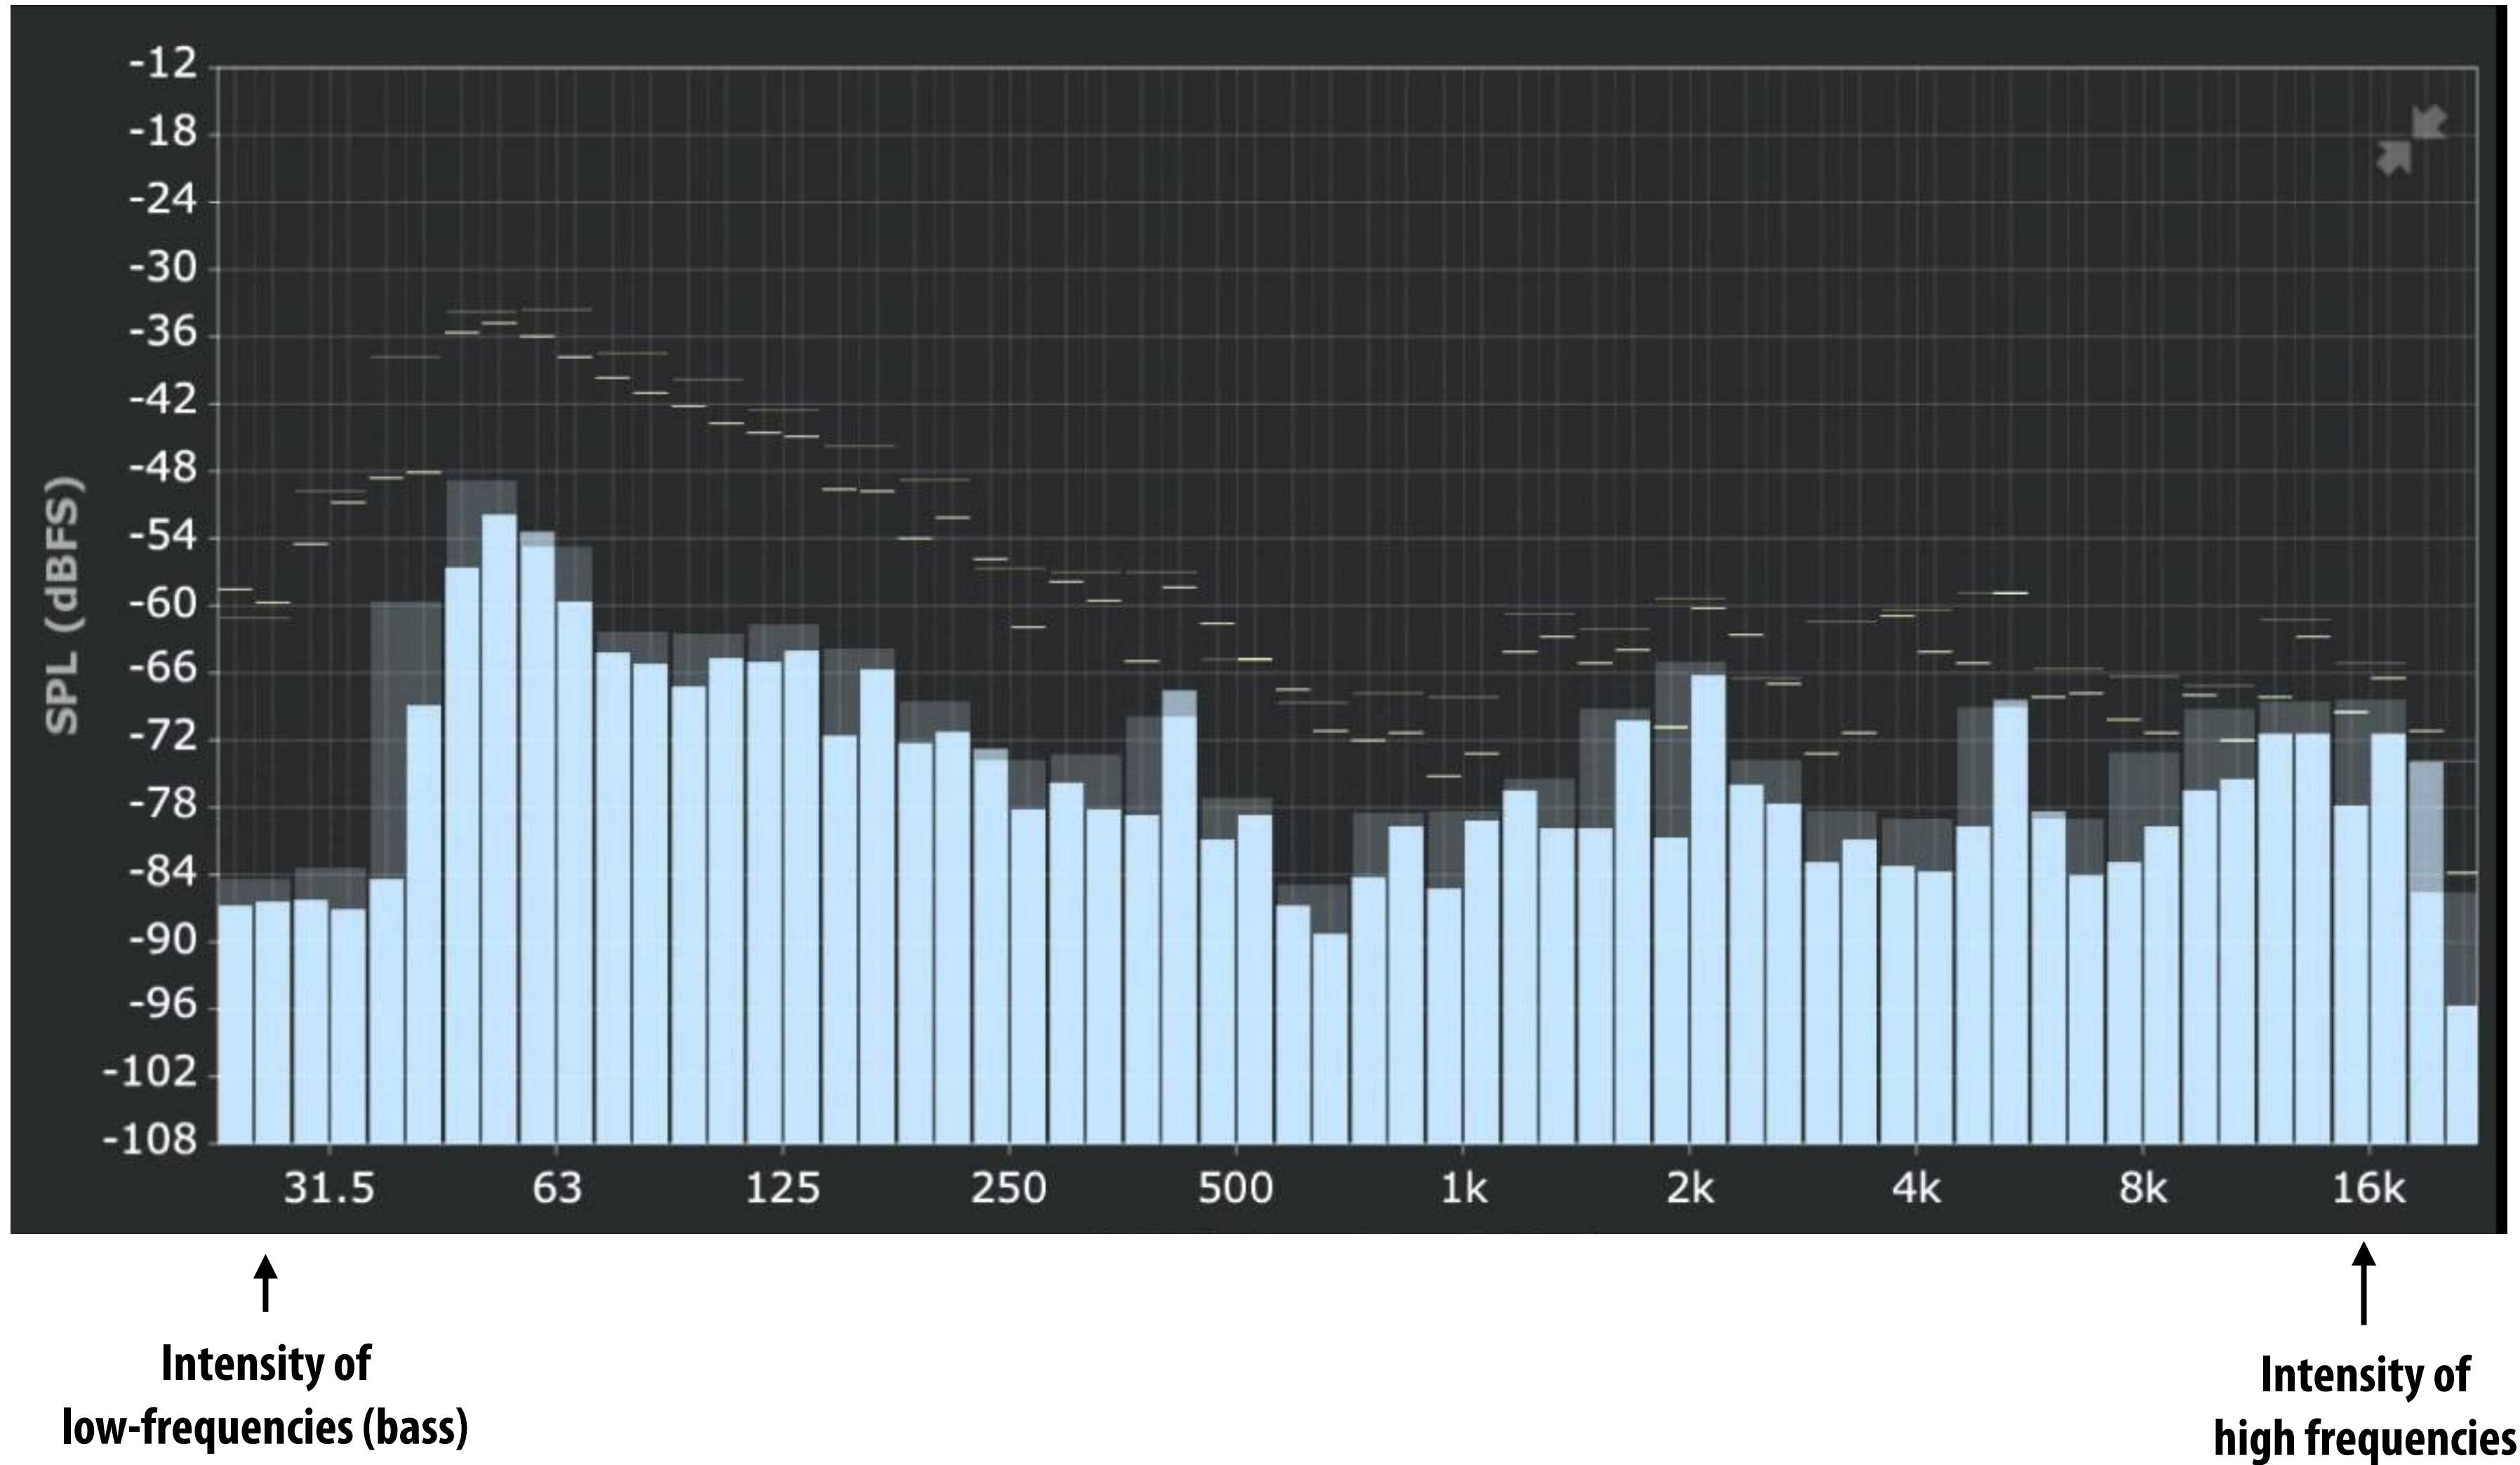
$$f_4(x) = \sin(4\pi x)$$



$$f(x) = f_1(x) + 0.75 f_2(x) + 0.5 f_4(x)$$



Audio spectrum analyzer: representing sound as a sum of its constituent frequencies



↑
Intensity of
low-frequencies (bass)

↑
Intensity of
high frequencies

Fourier transform

- Convert representation of signal from spatial/temporal domain to frequency domain by projecting signal into its component frequencies

$$\begin{aligned} f(\xi) &= \int_{-\infty}^{\infty} f(x) e^{-2\pi i x \xi} dx \\ &= \int_{-\infty}^{\infty} f(x) (\cos(2\pi \xi x) - i \sin(2\pi \xi x)) dx \end{aligned}$$

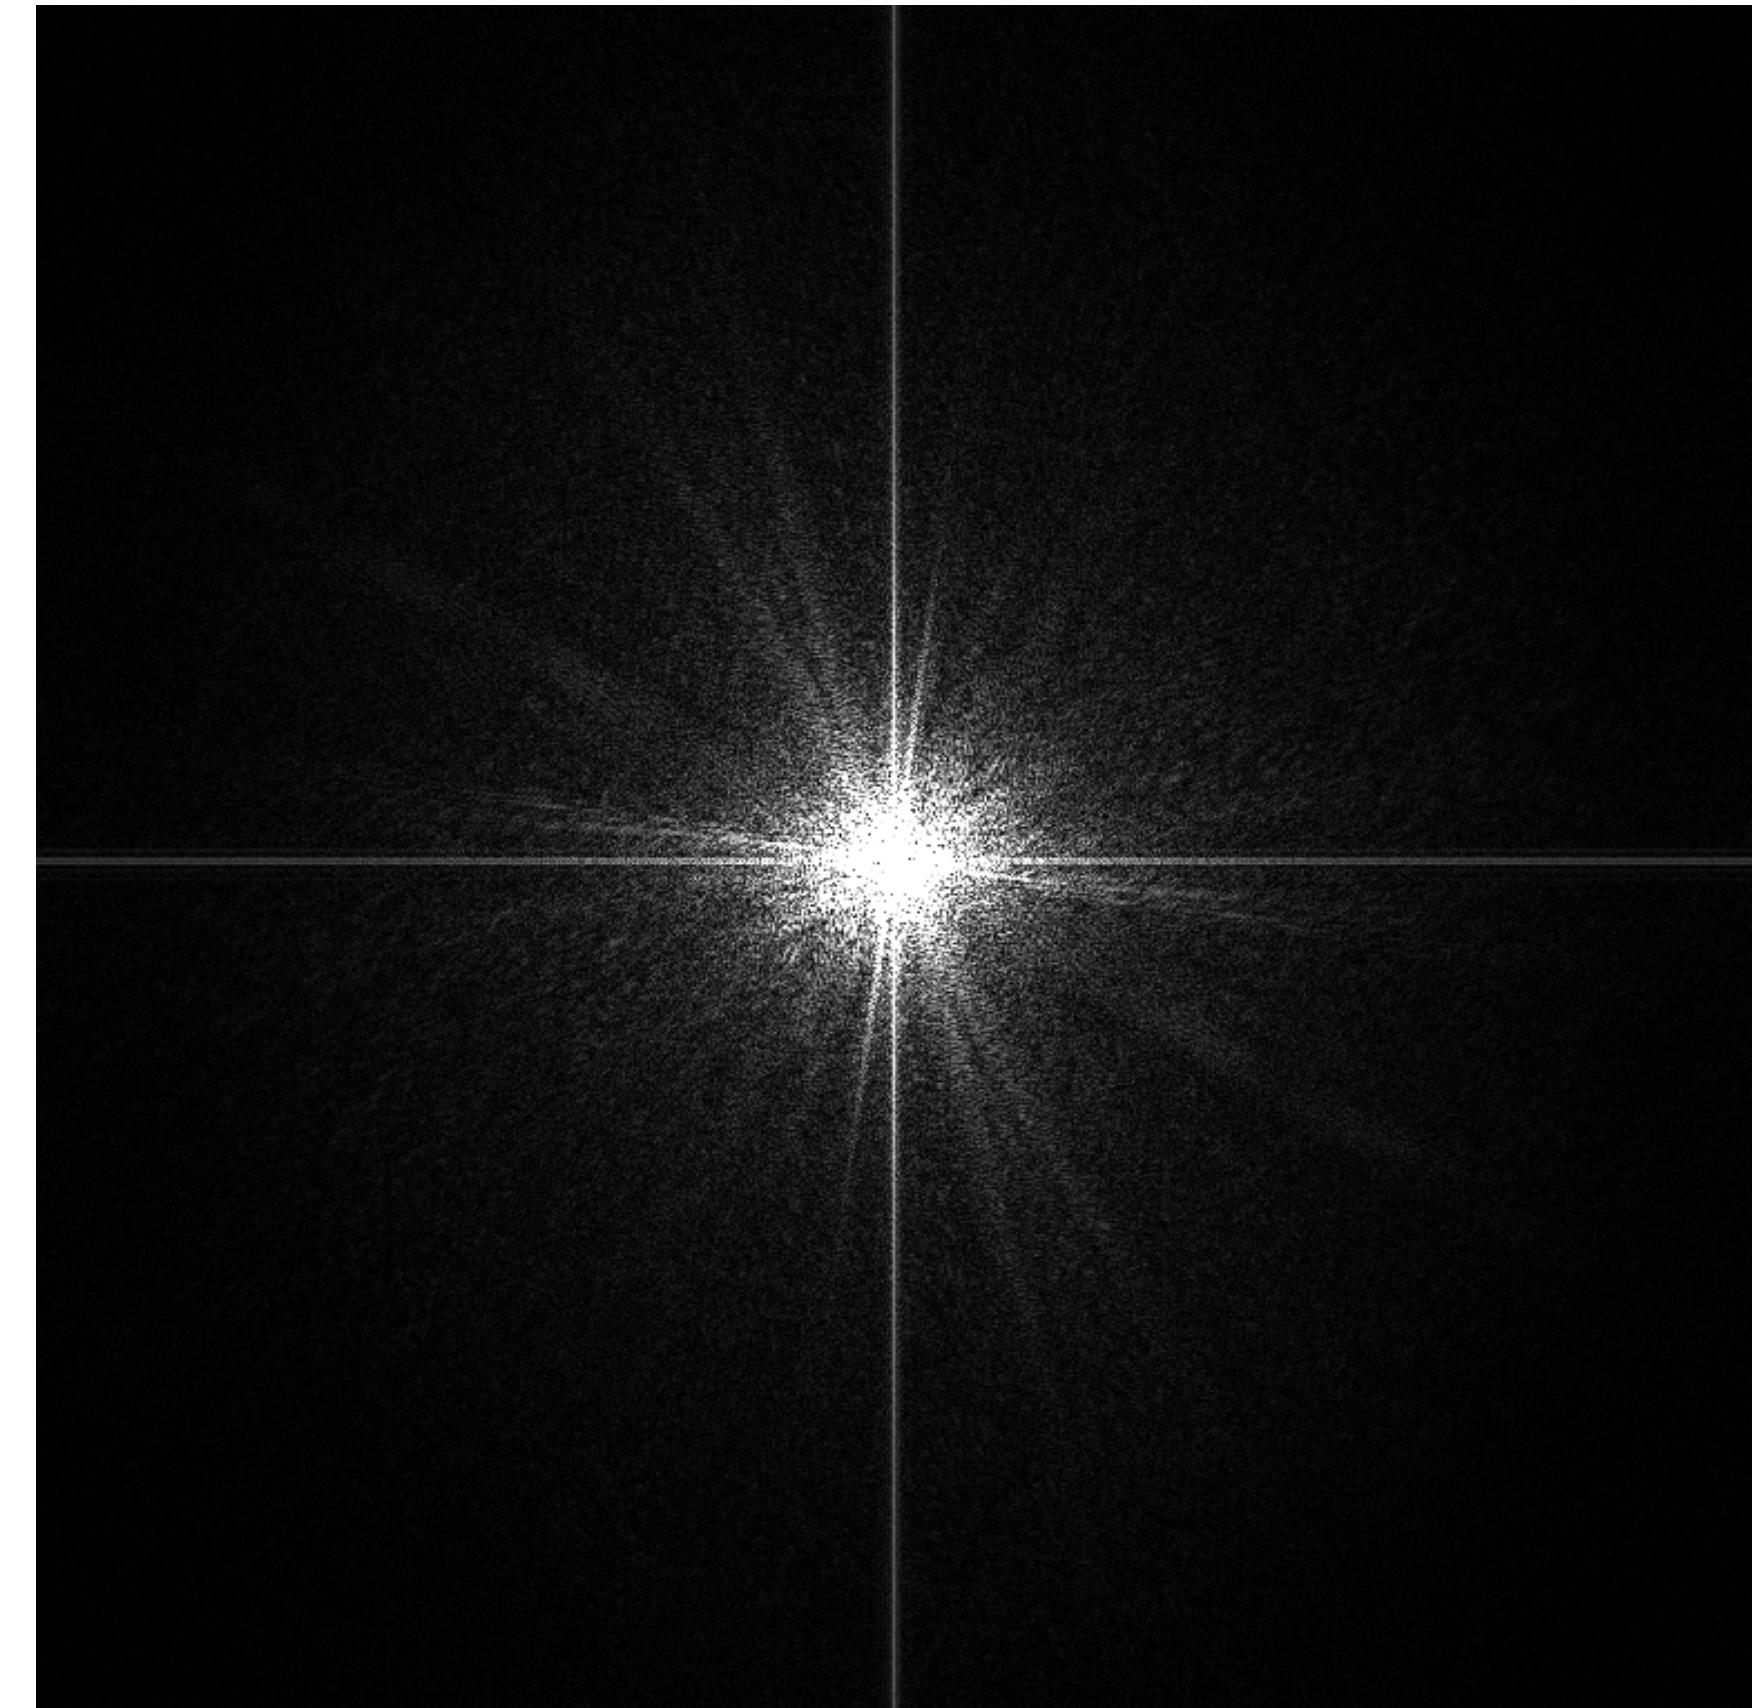
- 2D form:

$$f(u, v) = \iint f(x, y) e^{-2\pi i (ux + vy)} dx dy$$

Visualizing the frequency content of images



Spatial domain result

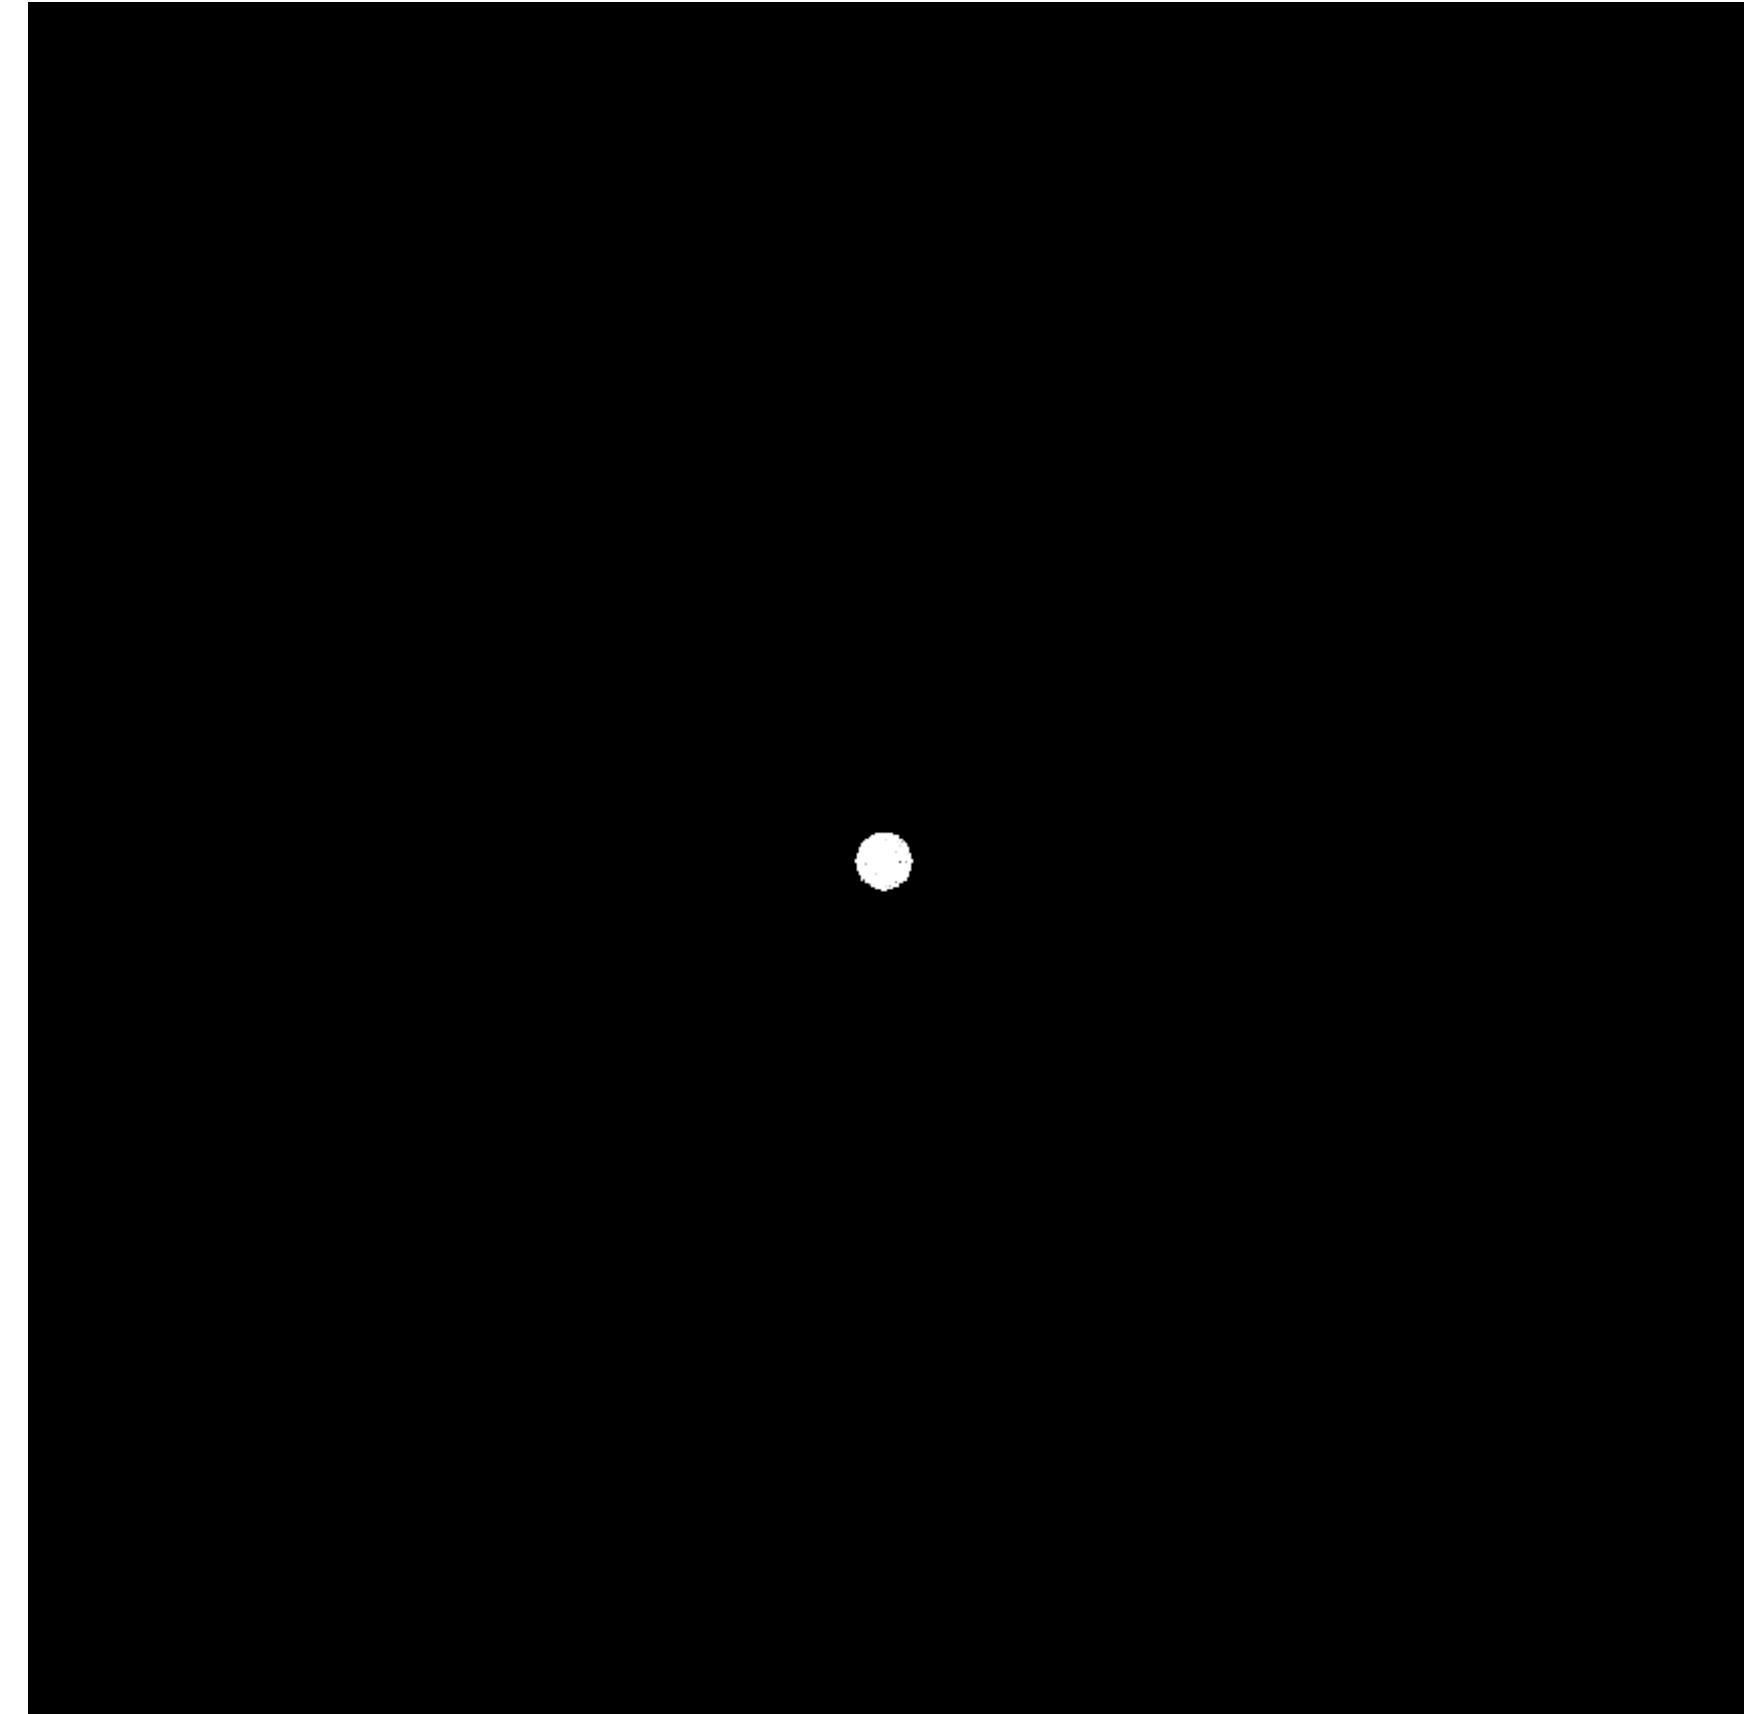


Spectrum

Low frequencies only (smooth gradients)



Spatial domain result

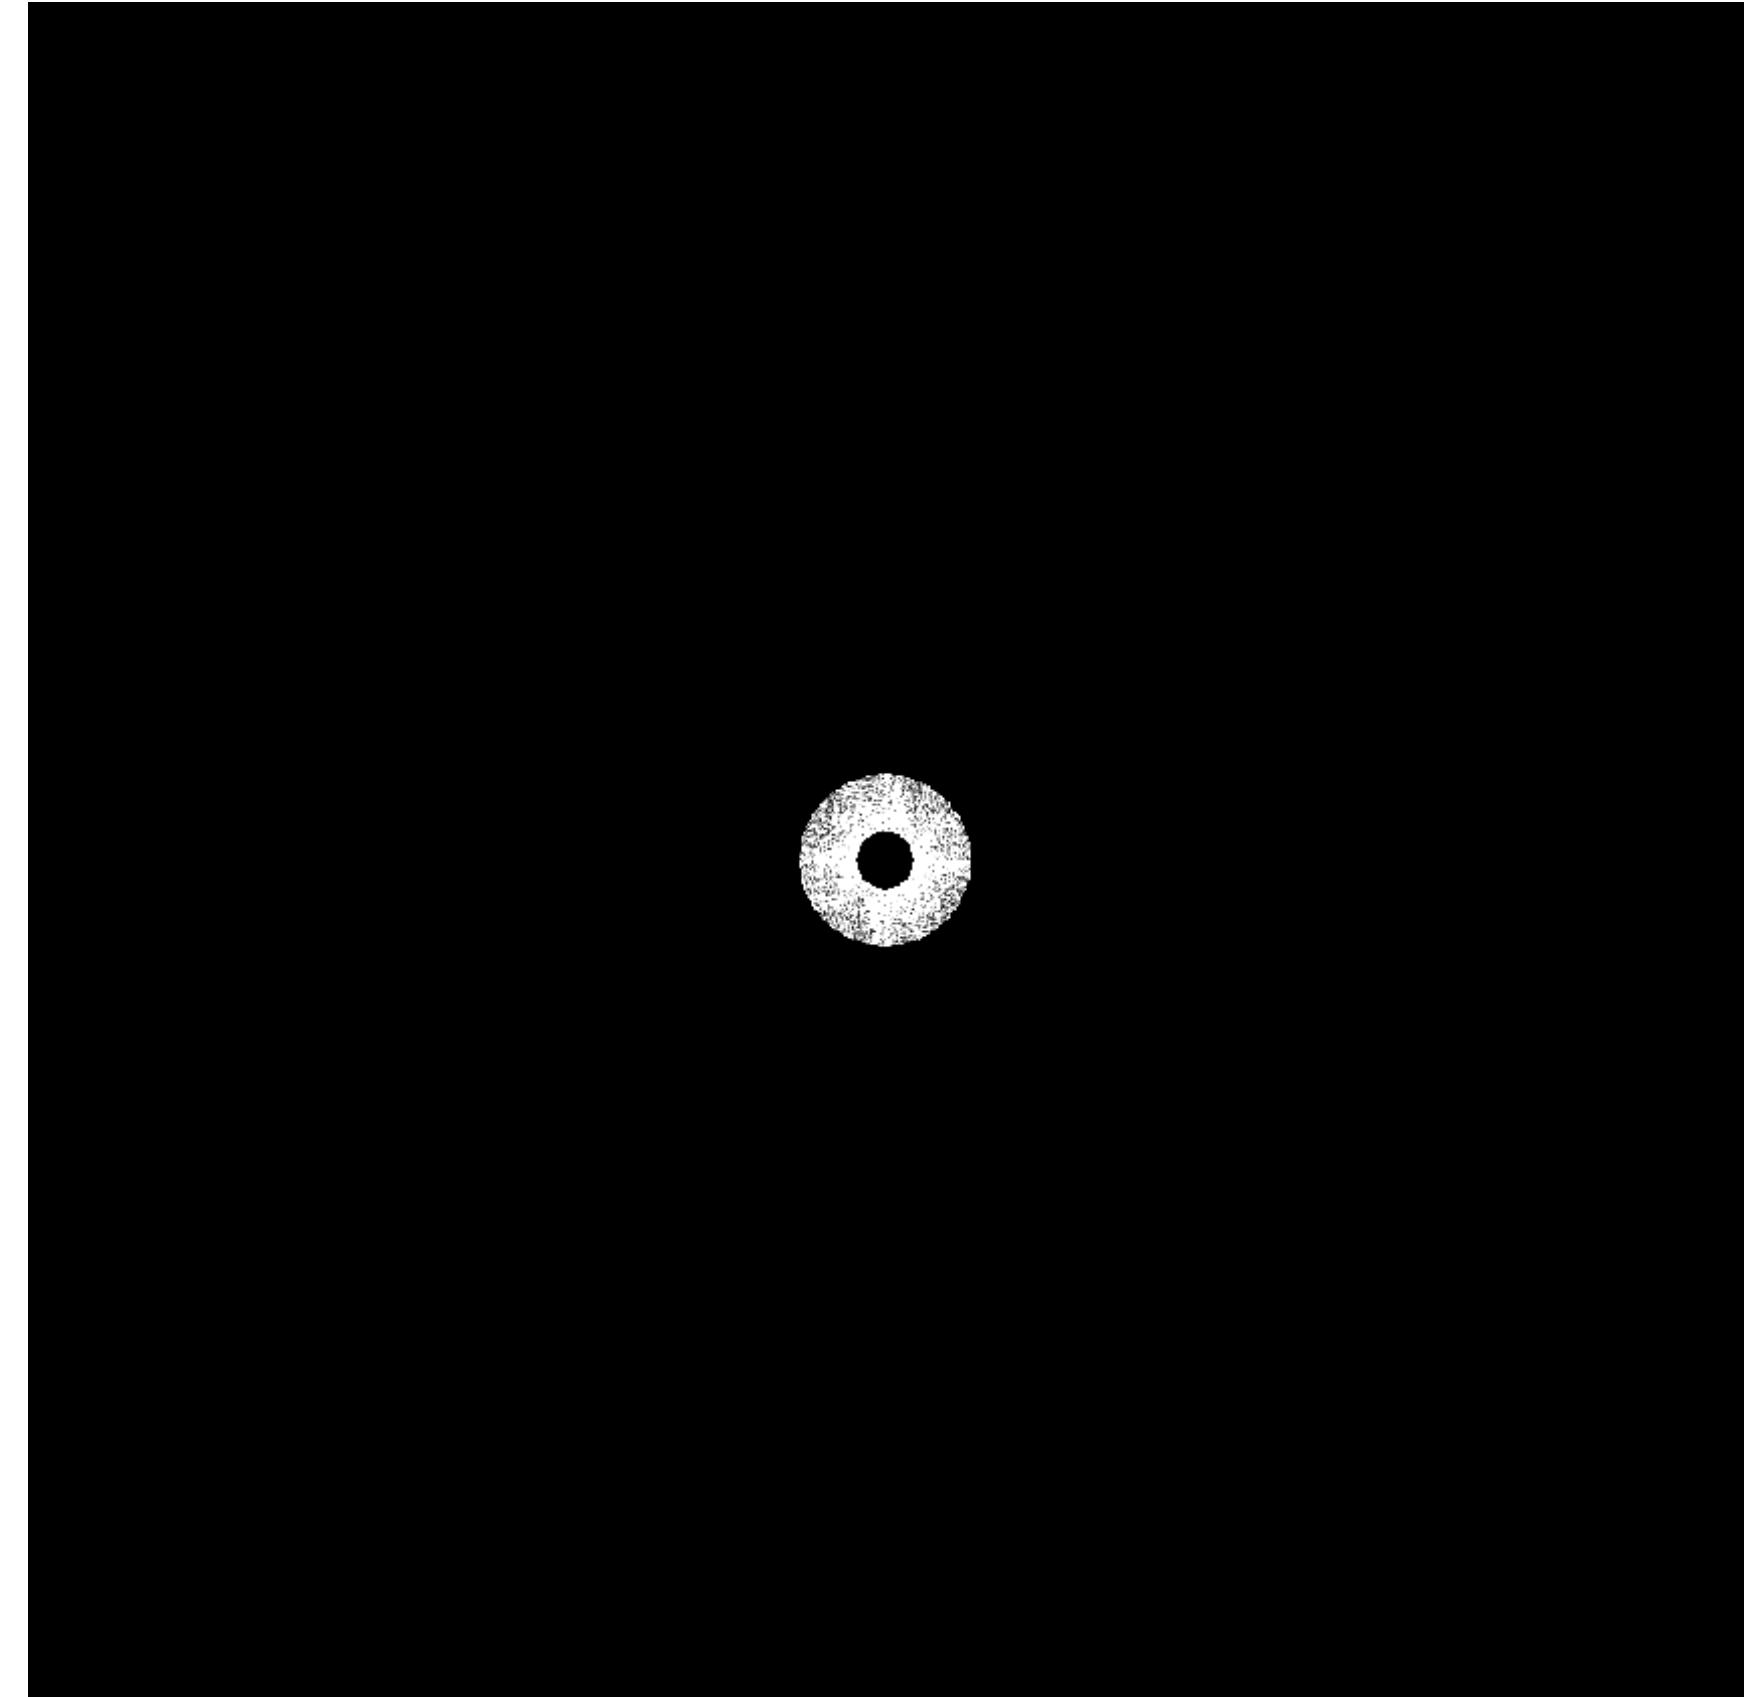


Spectrum (after low-pass filter)
All frequencies above cutoff have 0 magnitude

Mid-range frequencies



Spatial domain result

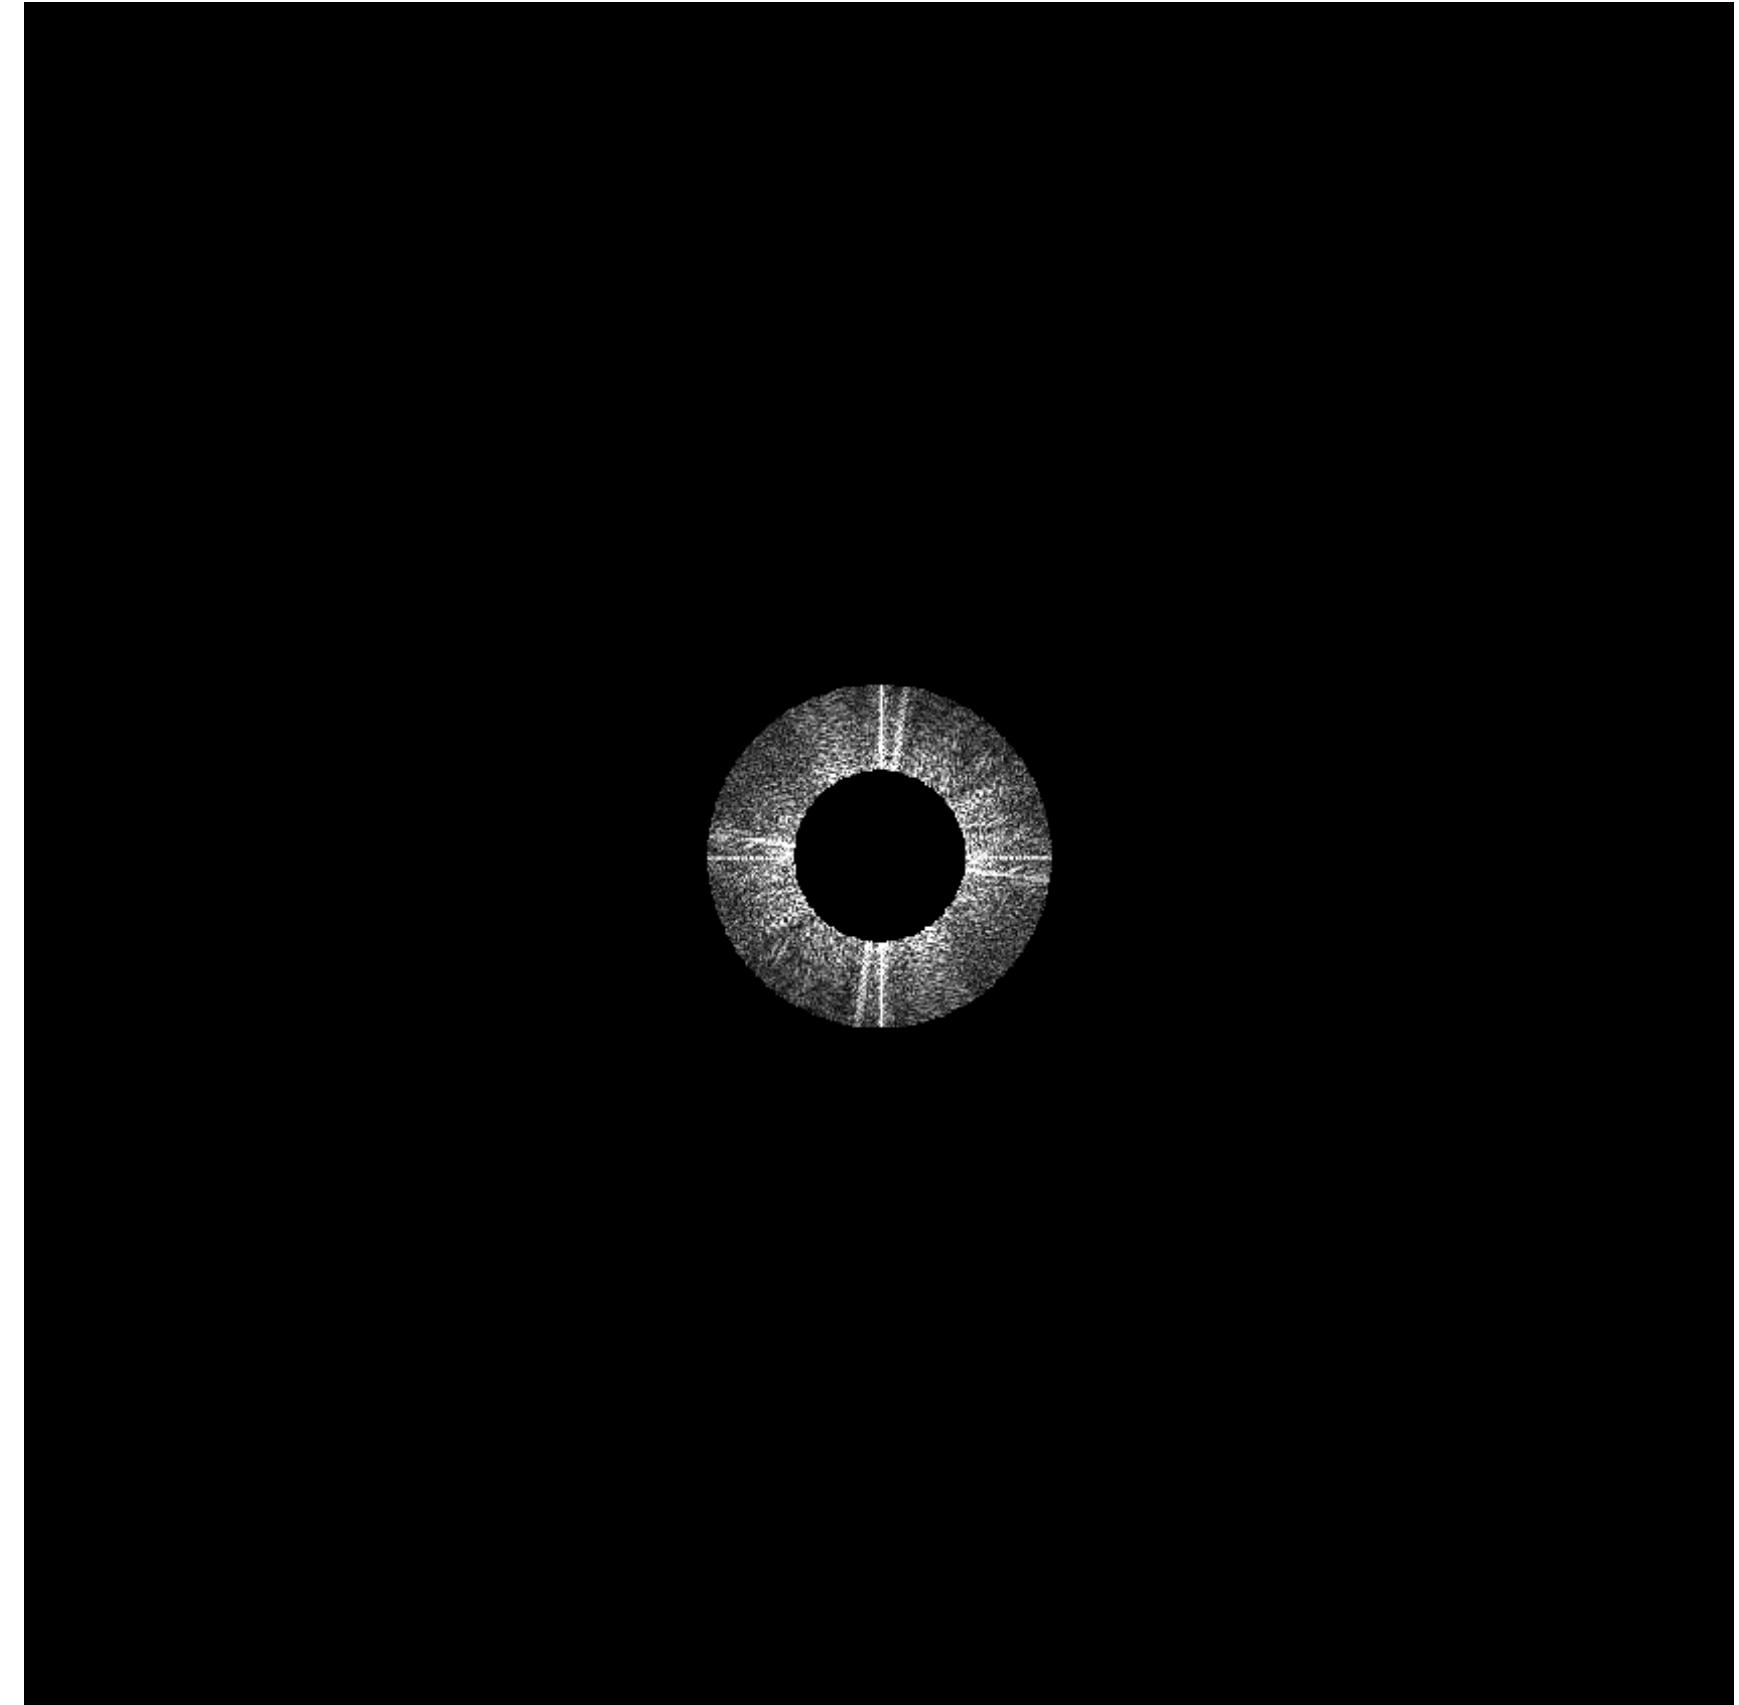


Spectrum (after band-pass filter)

Mid-range frequencies



Spatial domain result

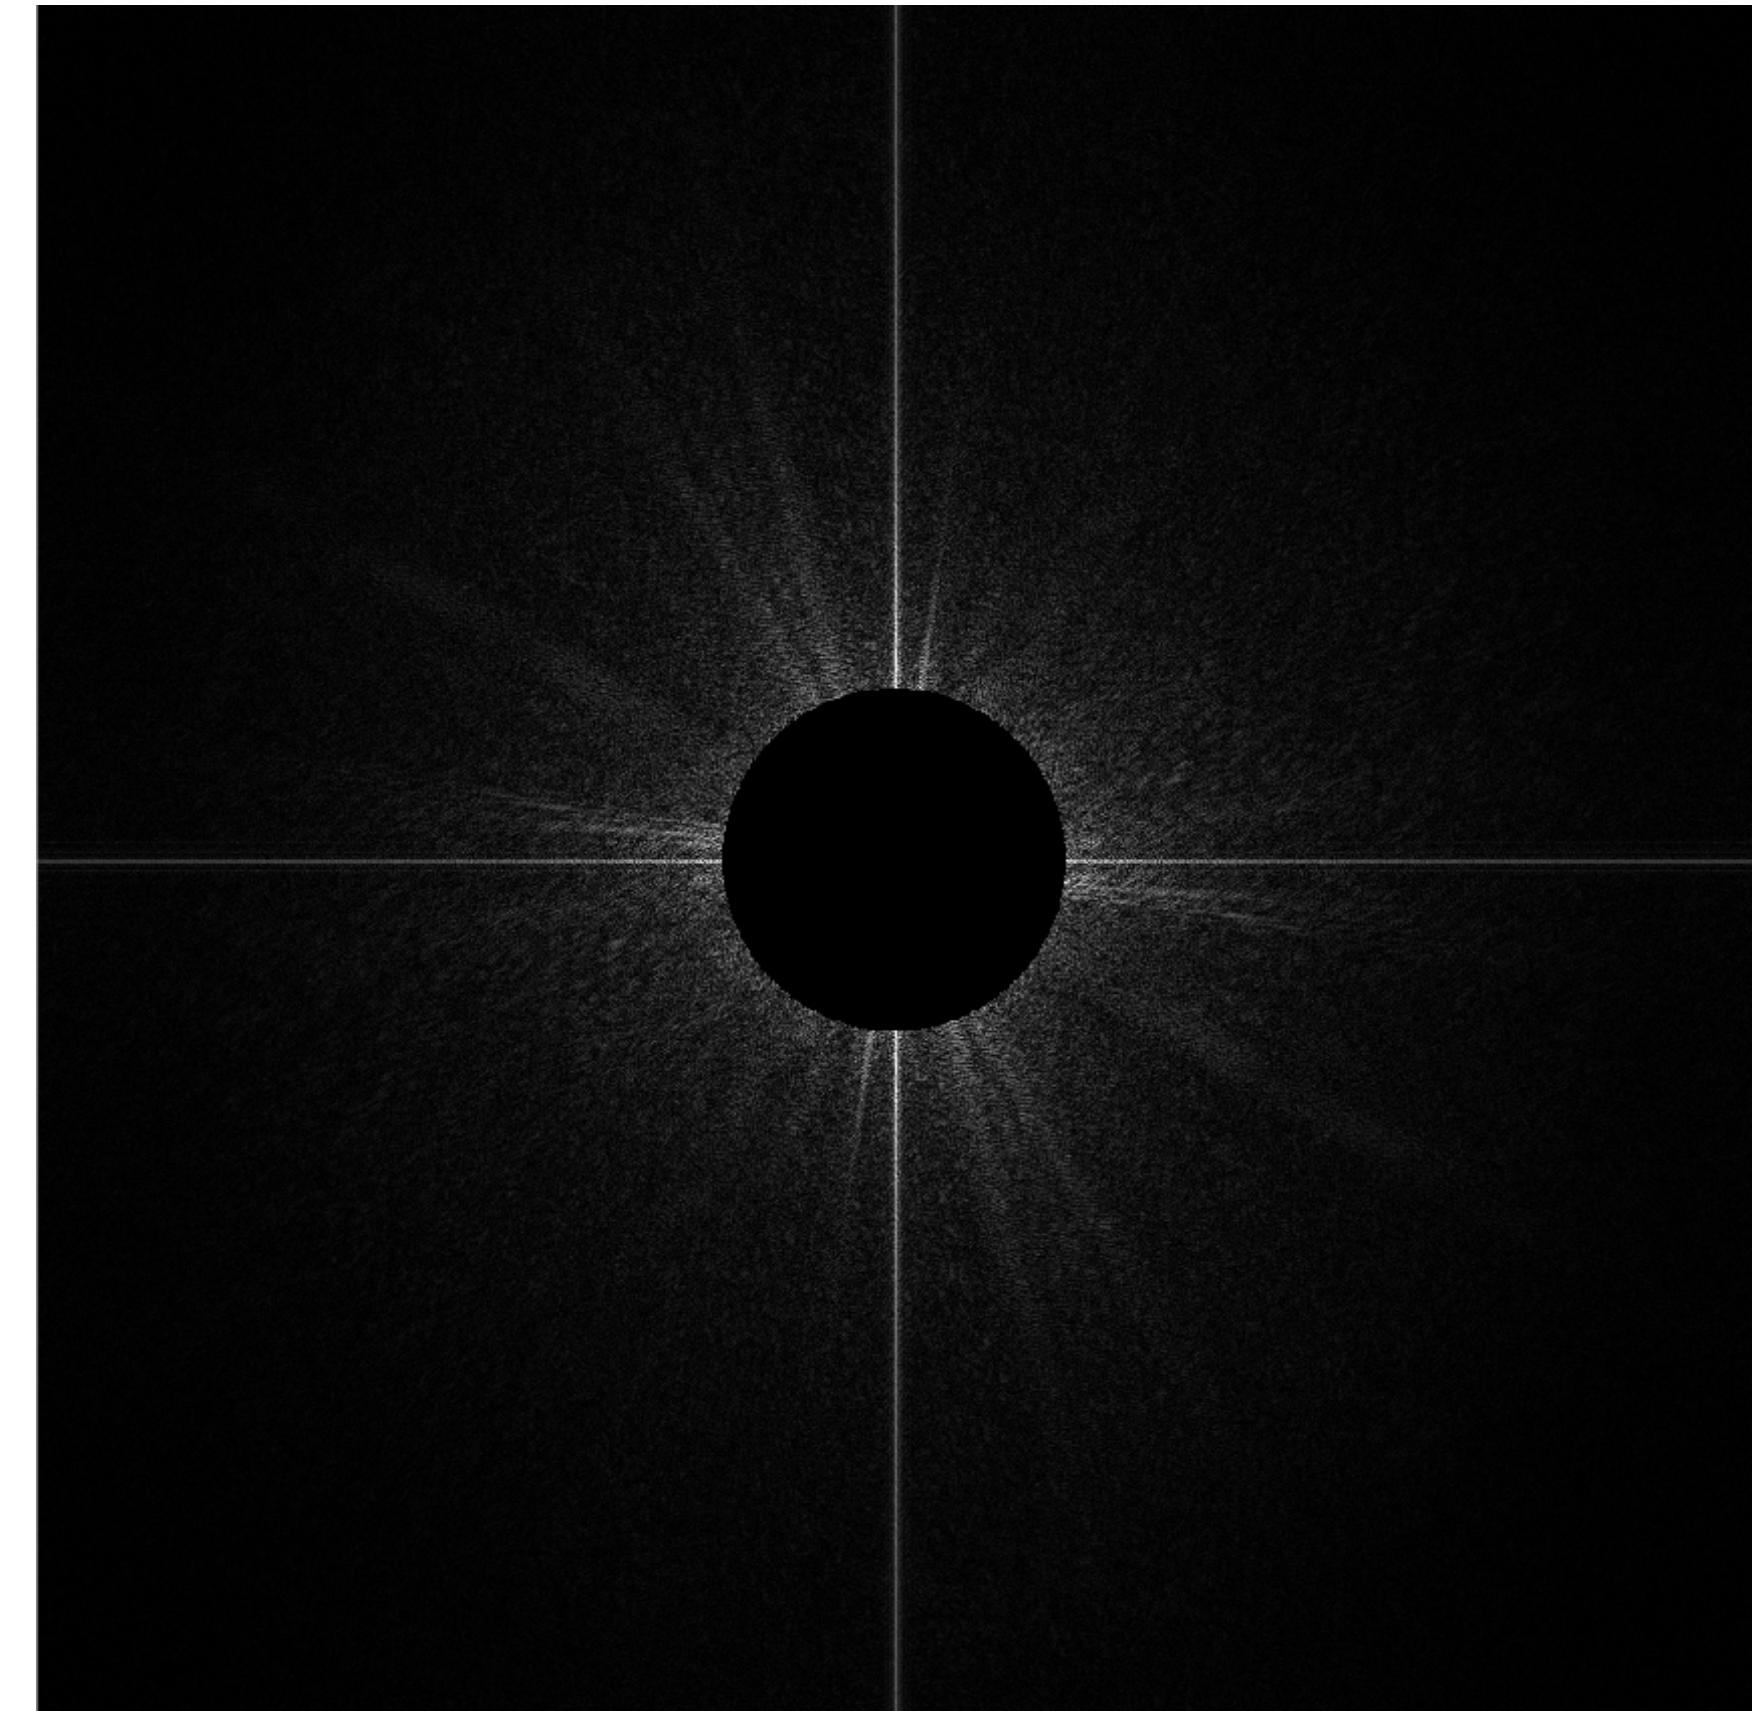


Spectrum (after band-pass filter)

High frequencies (edges)



**Spatial domain result
(strongest edges)**



**Spectrum (after high-pass filter)
All frequencies below threshold
have 0 magnitude**

An image as a sum of its frequency components



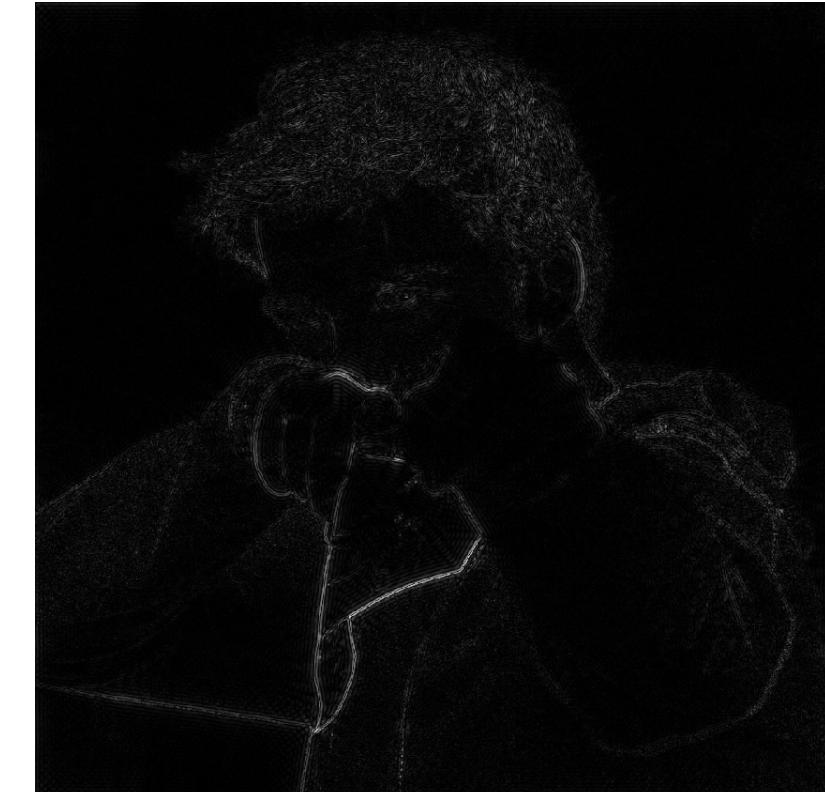
+



+



+



=



Another (linear) sharpening filter

$$\text{blurred} = g * I$$

$$\text{fine} = I - \text{blurred} \xleftarrow{\hspace{10em}} \text{Extract high frequencies}$$

$$\text{sharpened} = I + 0.5 \times \text{fine} \xleftarrow{\hspace{10em}} \text{Boost high frequencies}$$

**But what if we'd like to localize image edits
both in space and in frequency?**

Downsample

- **Step 1: Remove high frequencies**
- **Step 2: Sparsely sample pixels (below: every other pixel)**

```
float input[(2*WIDTH+2) * (2*HEIGHT+2)];
float output[WIDTH * HEIGHT];

float weights[] = {1/64, 3/64, 3/64, 1/64,      // 4x4 blur (approx Gaussian)
                  3/64, 9/64, 9/64, 3/64,
                  3/64, 9/64, 9/64, 3/64,
                  1/64, 3/64, 3/64, 1/64};

for (int j=0; j<HEIGHT; j++) {
    for (int i=0; i<WIDTH; i++) {
        float tmp = 0.f;
        for (int jj=0; jj<3; jj++)
            for (int ii=0; ii<3; ii++)
                tmp += input[(2*j+jj)*(WIDTH+2) + (2*i+ii)] * weights[jj*3 + ii];
        output[j*WIDTH + i] = tmp;
    }
}
```

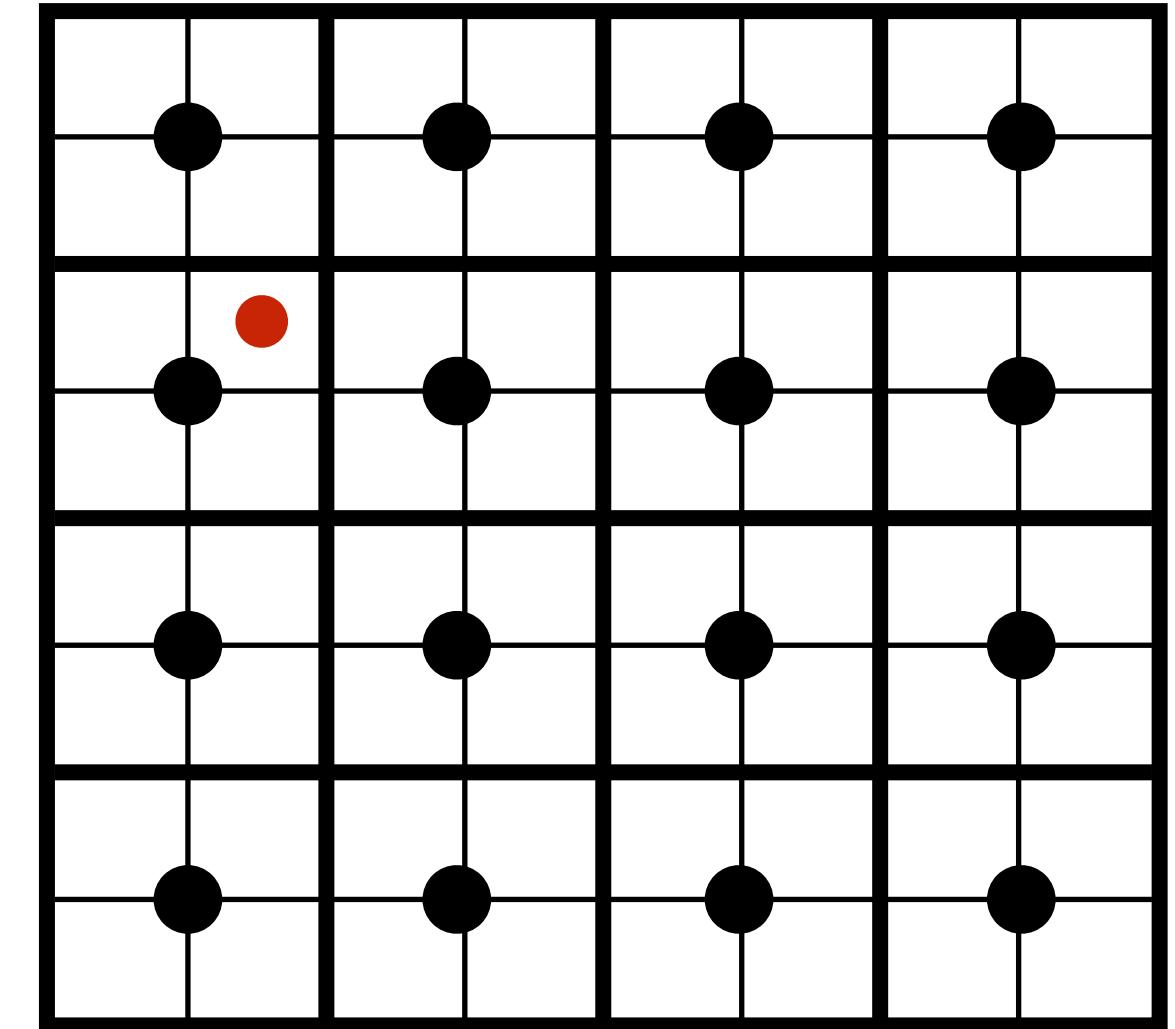
Upsample

Via bilinear interpolation of samples from low resolution image

```
float input[WIDTH * HEIGHT];
float output[2*WIDTH * 2*HEIGHT];

for (int j=0; j<2*HEIGHT; j++) {
    for (int i=0; i<2*WIDTH; i++) {
        int row = j/2;
        int col = i/2;
        float w1 = (i%2) ? .75 : .25;
        float w2 = (j%2) ? .75 : .25;

        output[j*2*WIDTH + i] = w1 * w2 * input[row*WIDTH + col] +
            (1.0-w1) * w2 * input[row*WIDTH + col+1] +
            w1 * (1-w2) * input[(row+1)*WIDTH + col] +
            (1.0-w1)*(1.0-w2) * input[(row+1)*WIDTH + col+1];
    }
}
```



Gaussian pyramid



$G_0 = \text{image}$

Each image in pyramid contains increasingly low-pass filtered signal

down() = downsample operation



$G_1 = \text{down}(G_0)$



$G_2 = \text{down}(G_1)$



Gaussian pyramid



G₀

Gaussian pyramid



G₁

Gaussian pyramid



G₂

Gaussian pyramid



G₃

Gaussian pyramid



G₄

Gaussian pyramid



G₅

Laplacian pyramid



$$L_0 = G_0 - \text{up}(G_1)$$

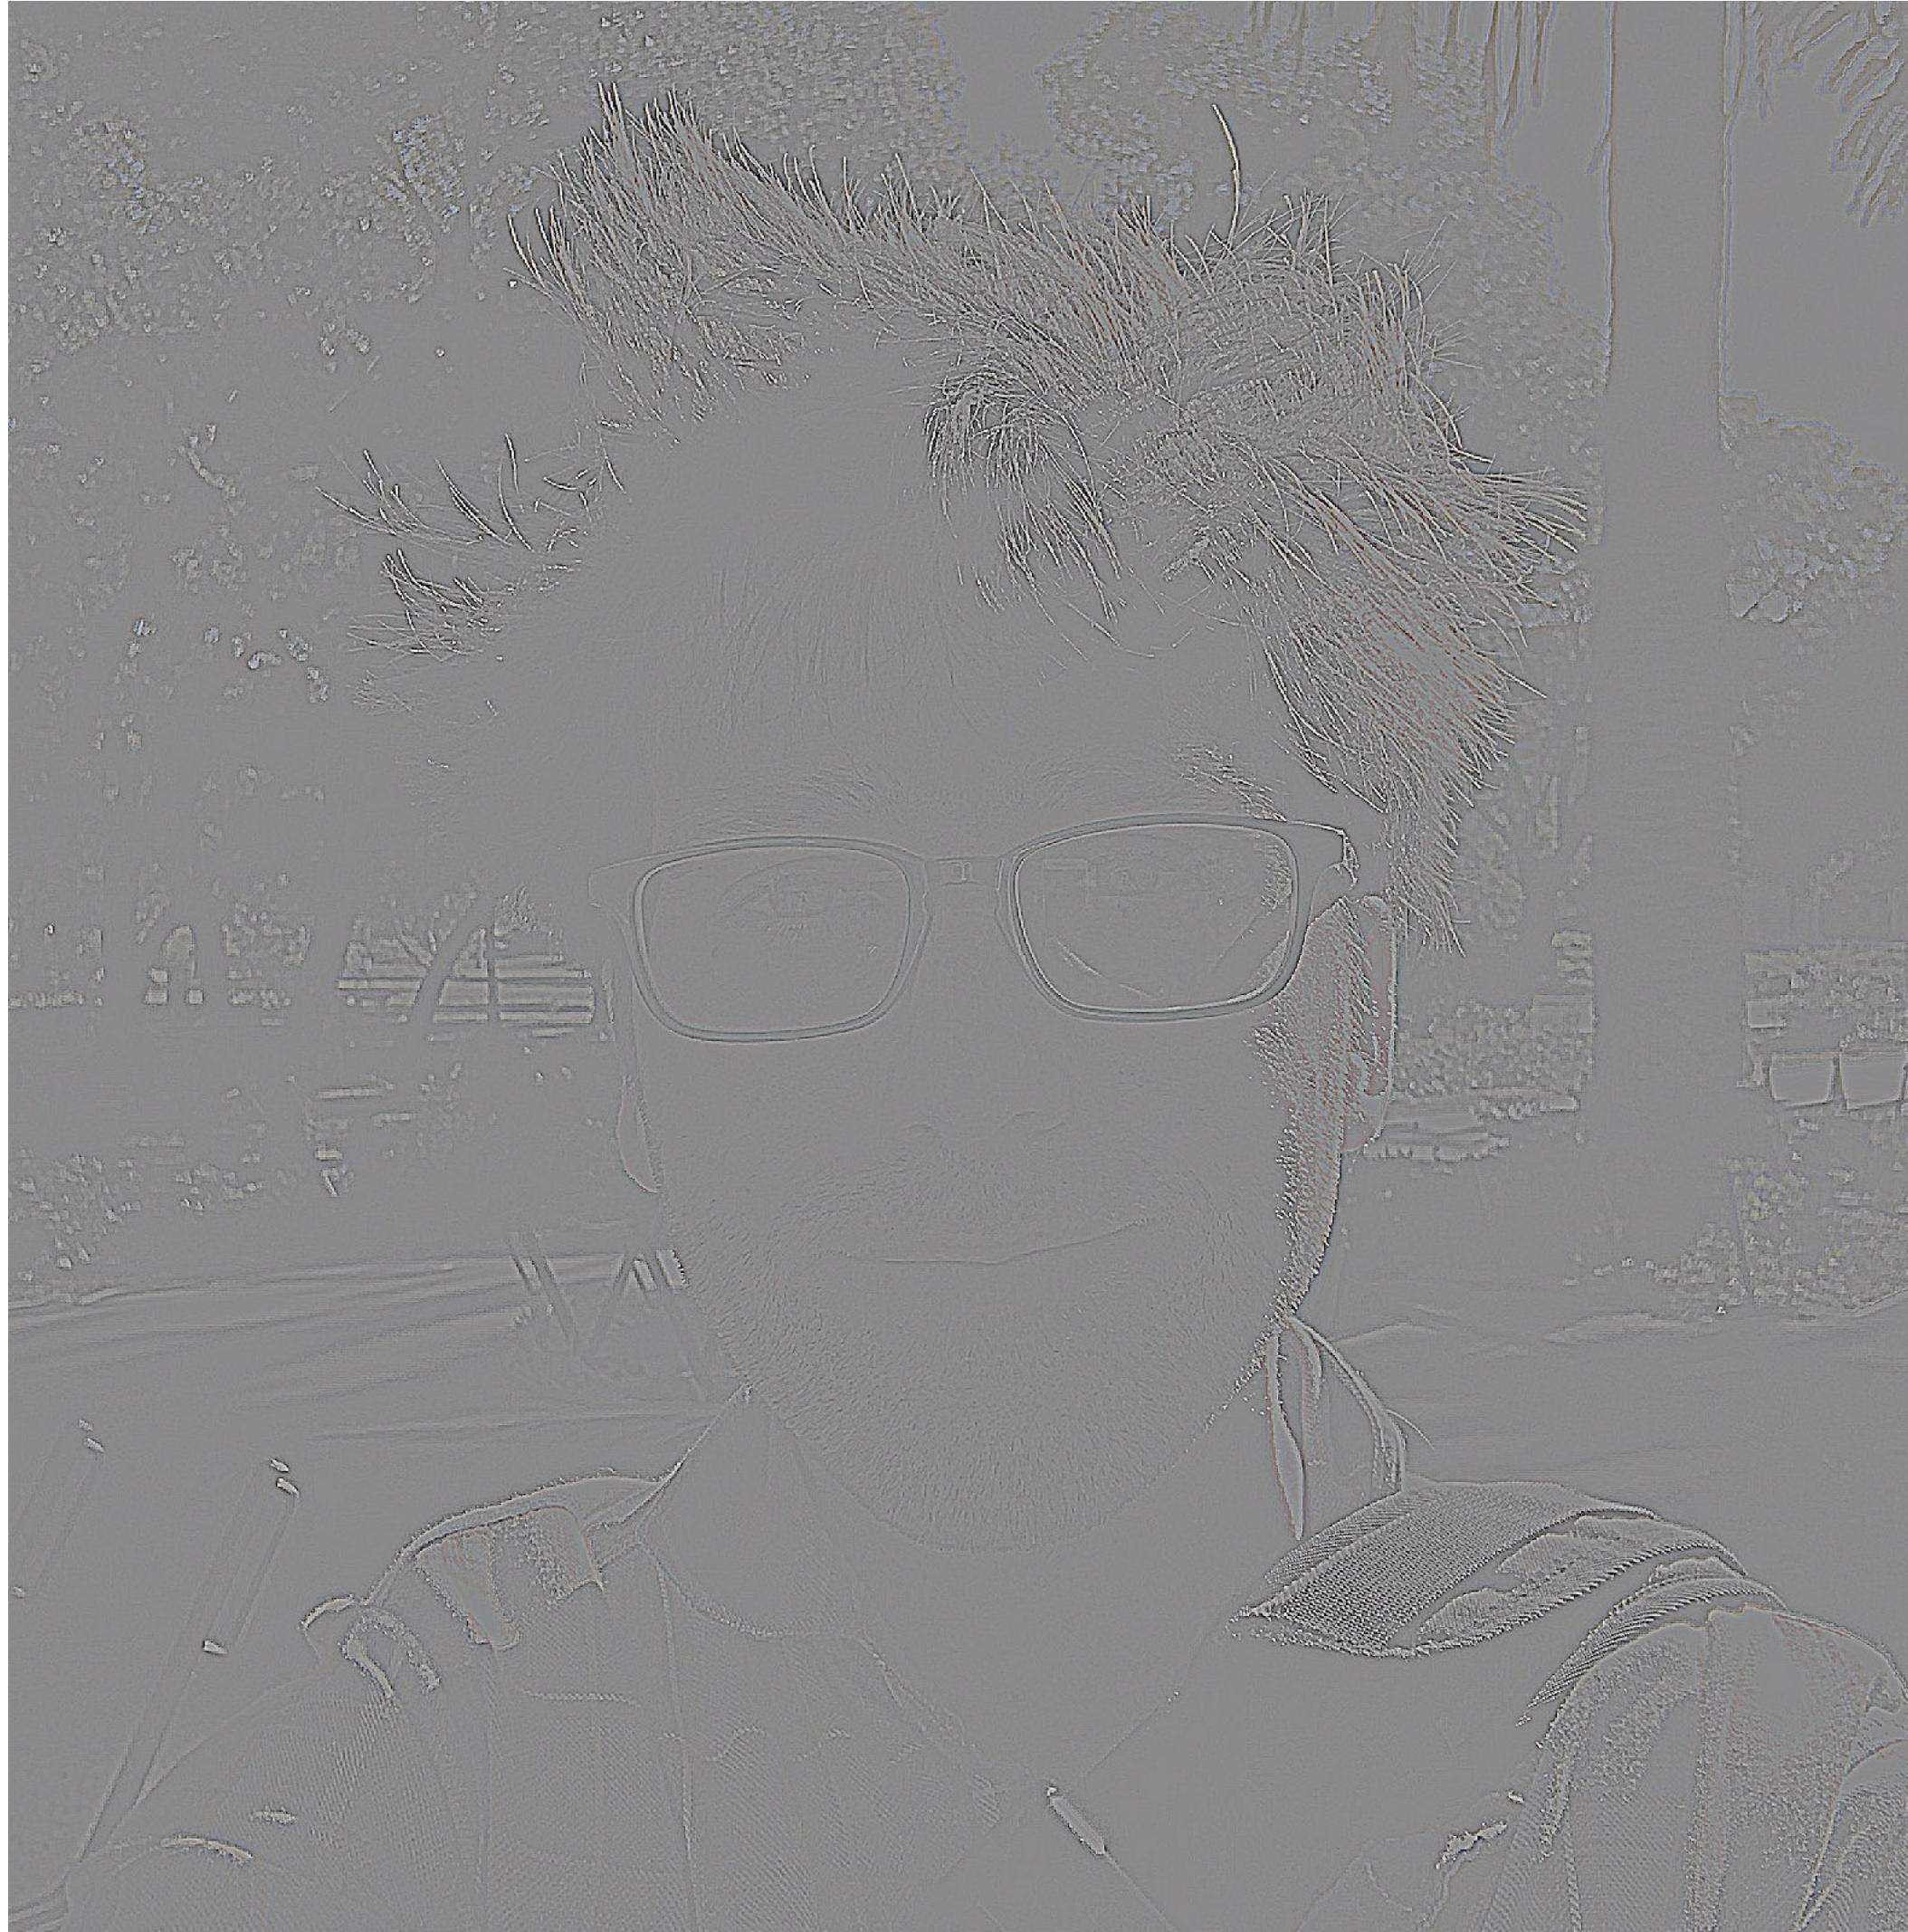


$$G_1 = \text{down}(G_0)$$

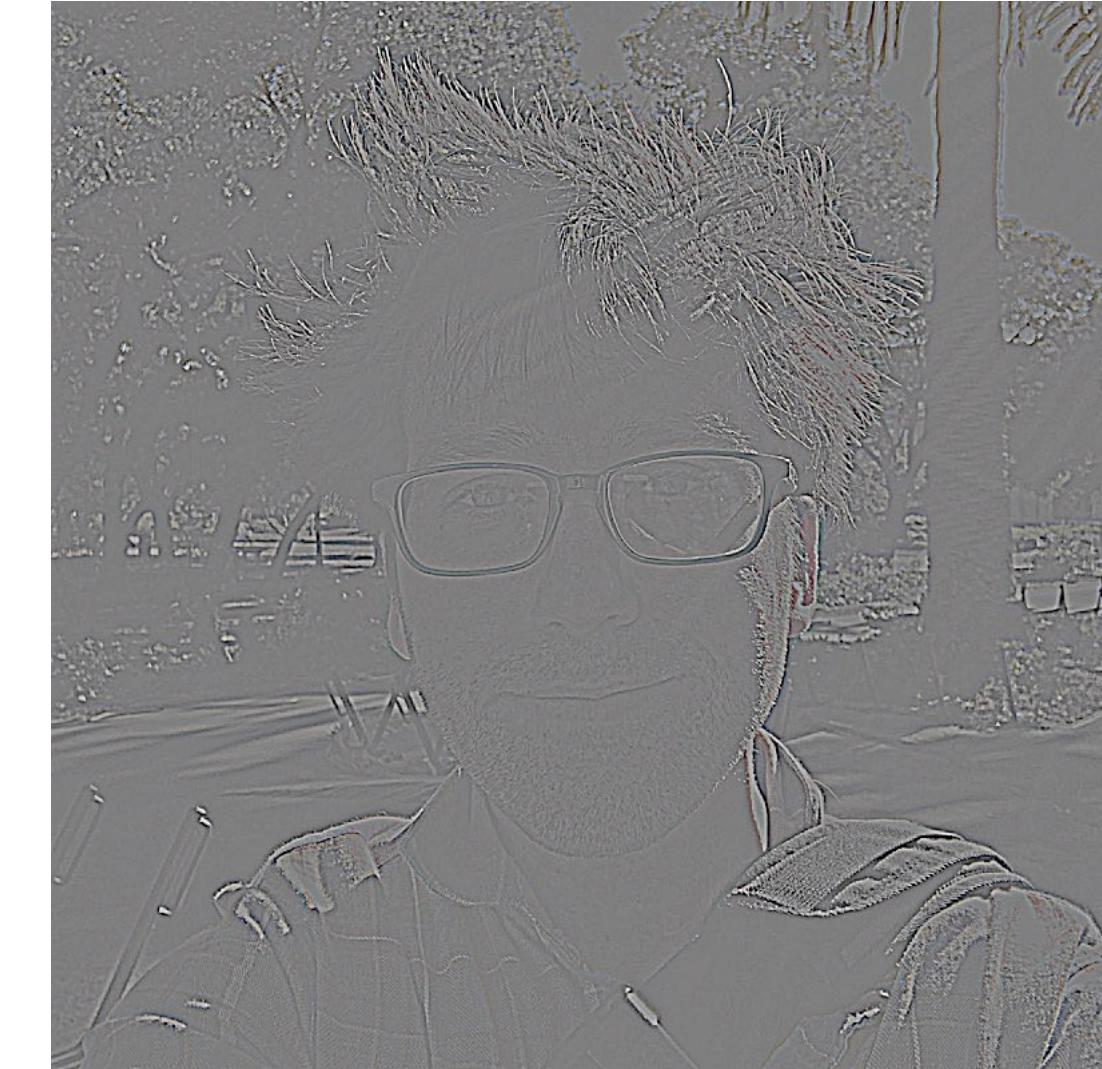
$$G_0$$

Each image in Laplacian pyramid represents increasingly high frequency content from image

Laplacian pyramid



$$L_0 = G_0 - up(G_1)$$

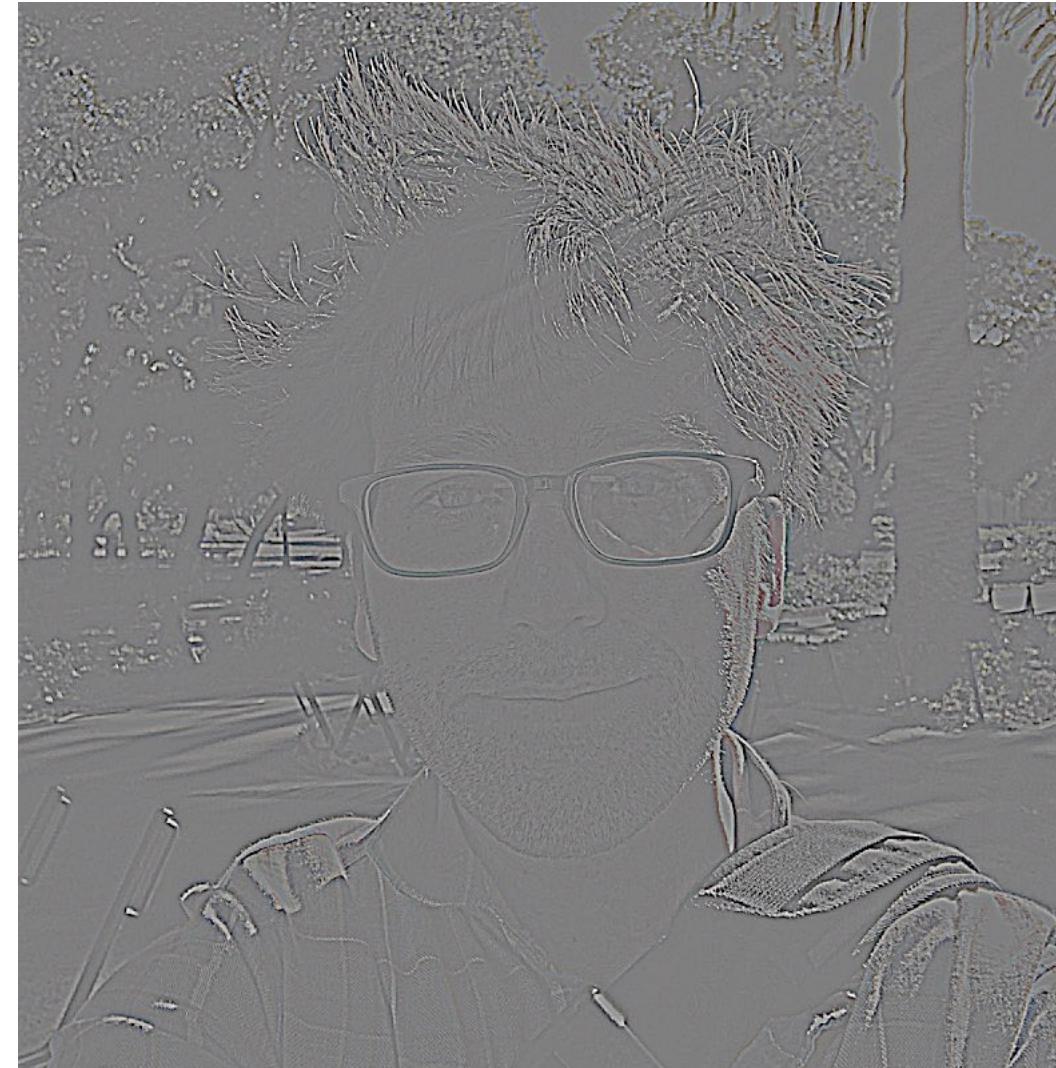


$$L_1 = G_1 - up(G_2)$$

Laplacian pyramid



$$L_0 = G_0 - up(G_1)$$



$$L_1 = G_1 - up(G_2)$$



$$L_2 = G_2 - up(G_3)$$

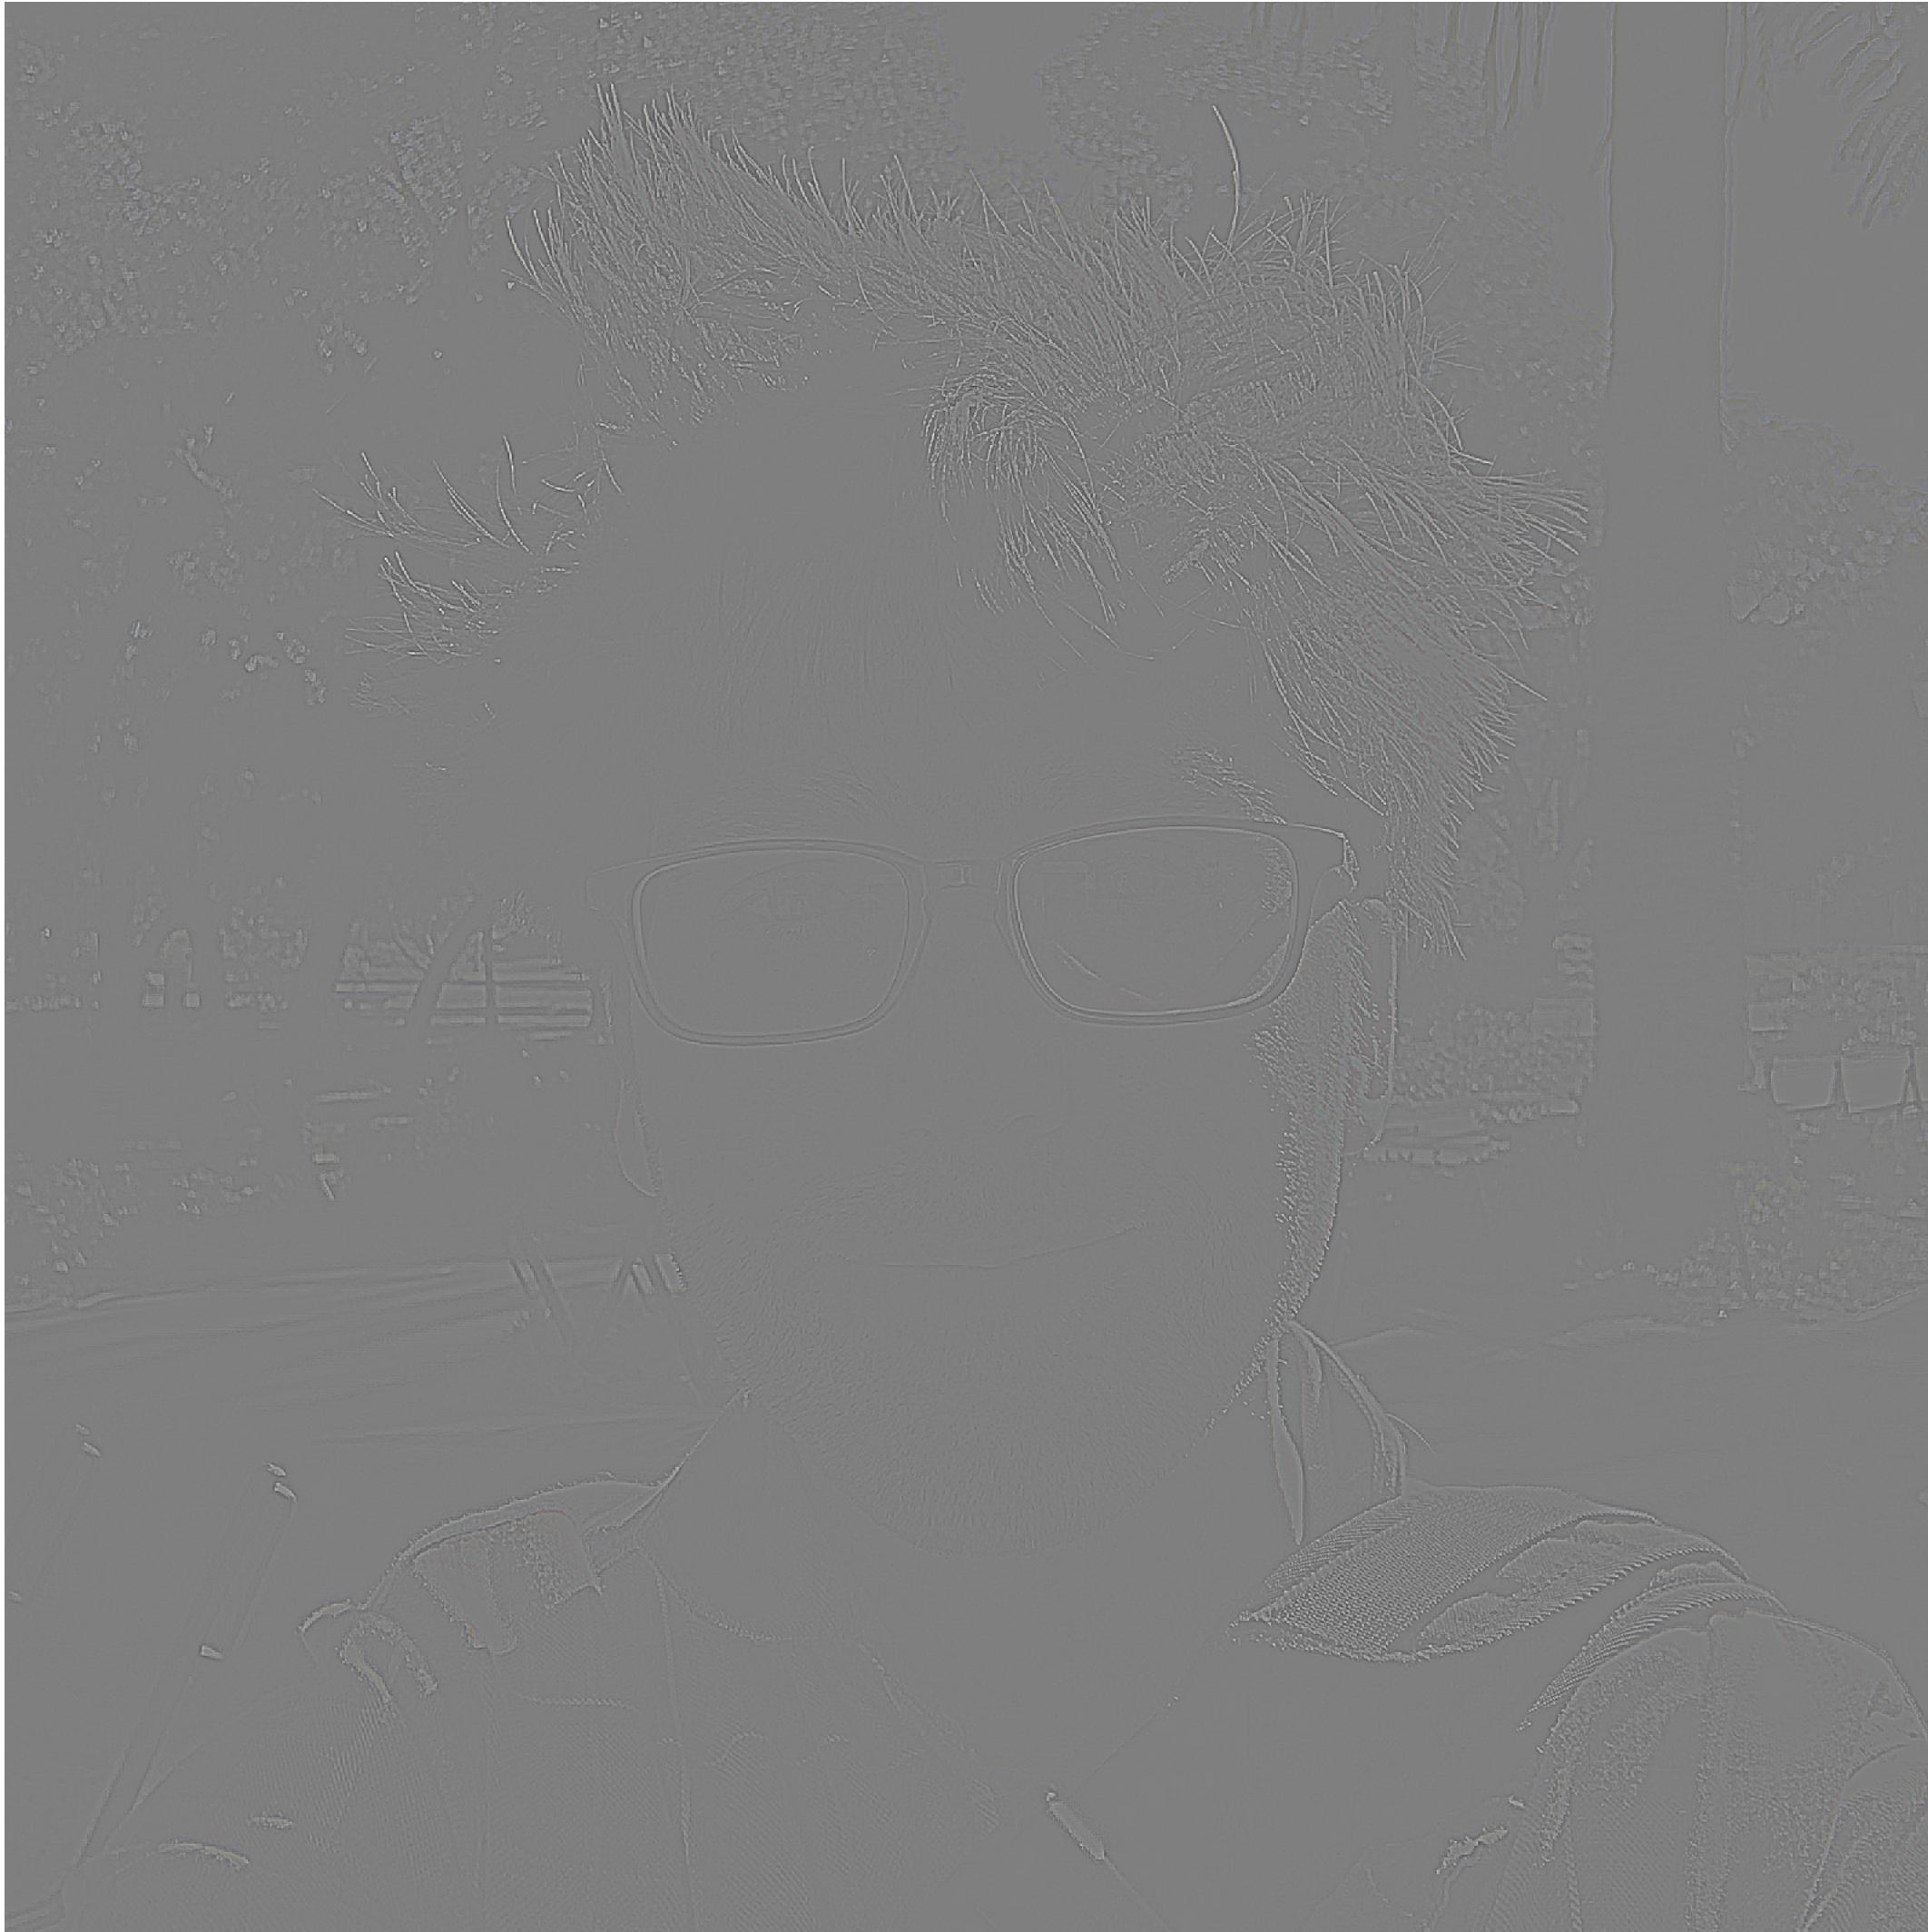


$$L_4 = G_4$$

$$L_3 = G_3 - up(G_4)$$

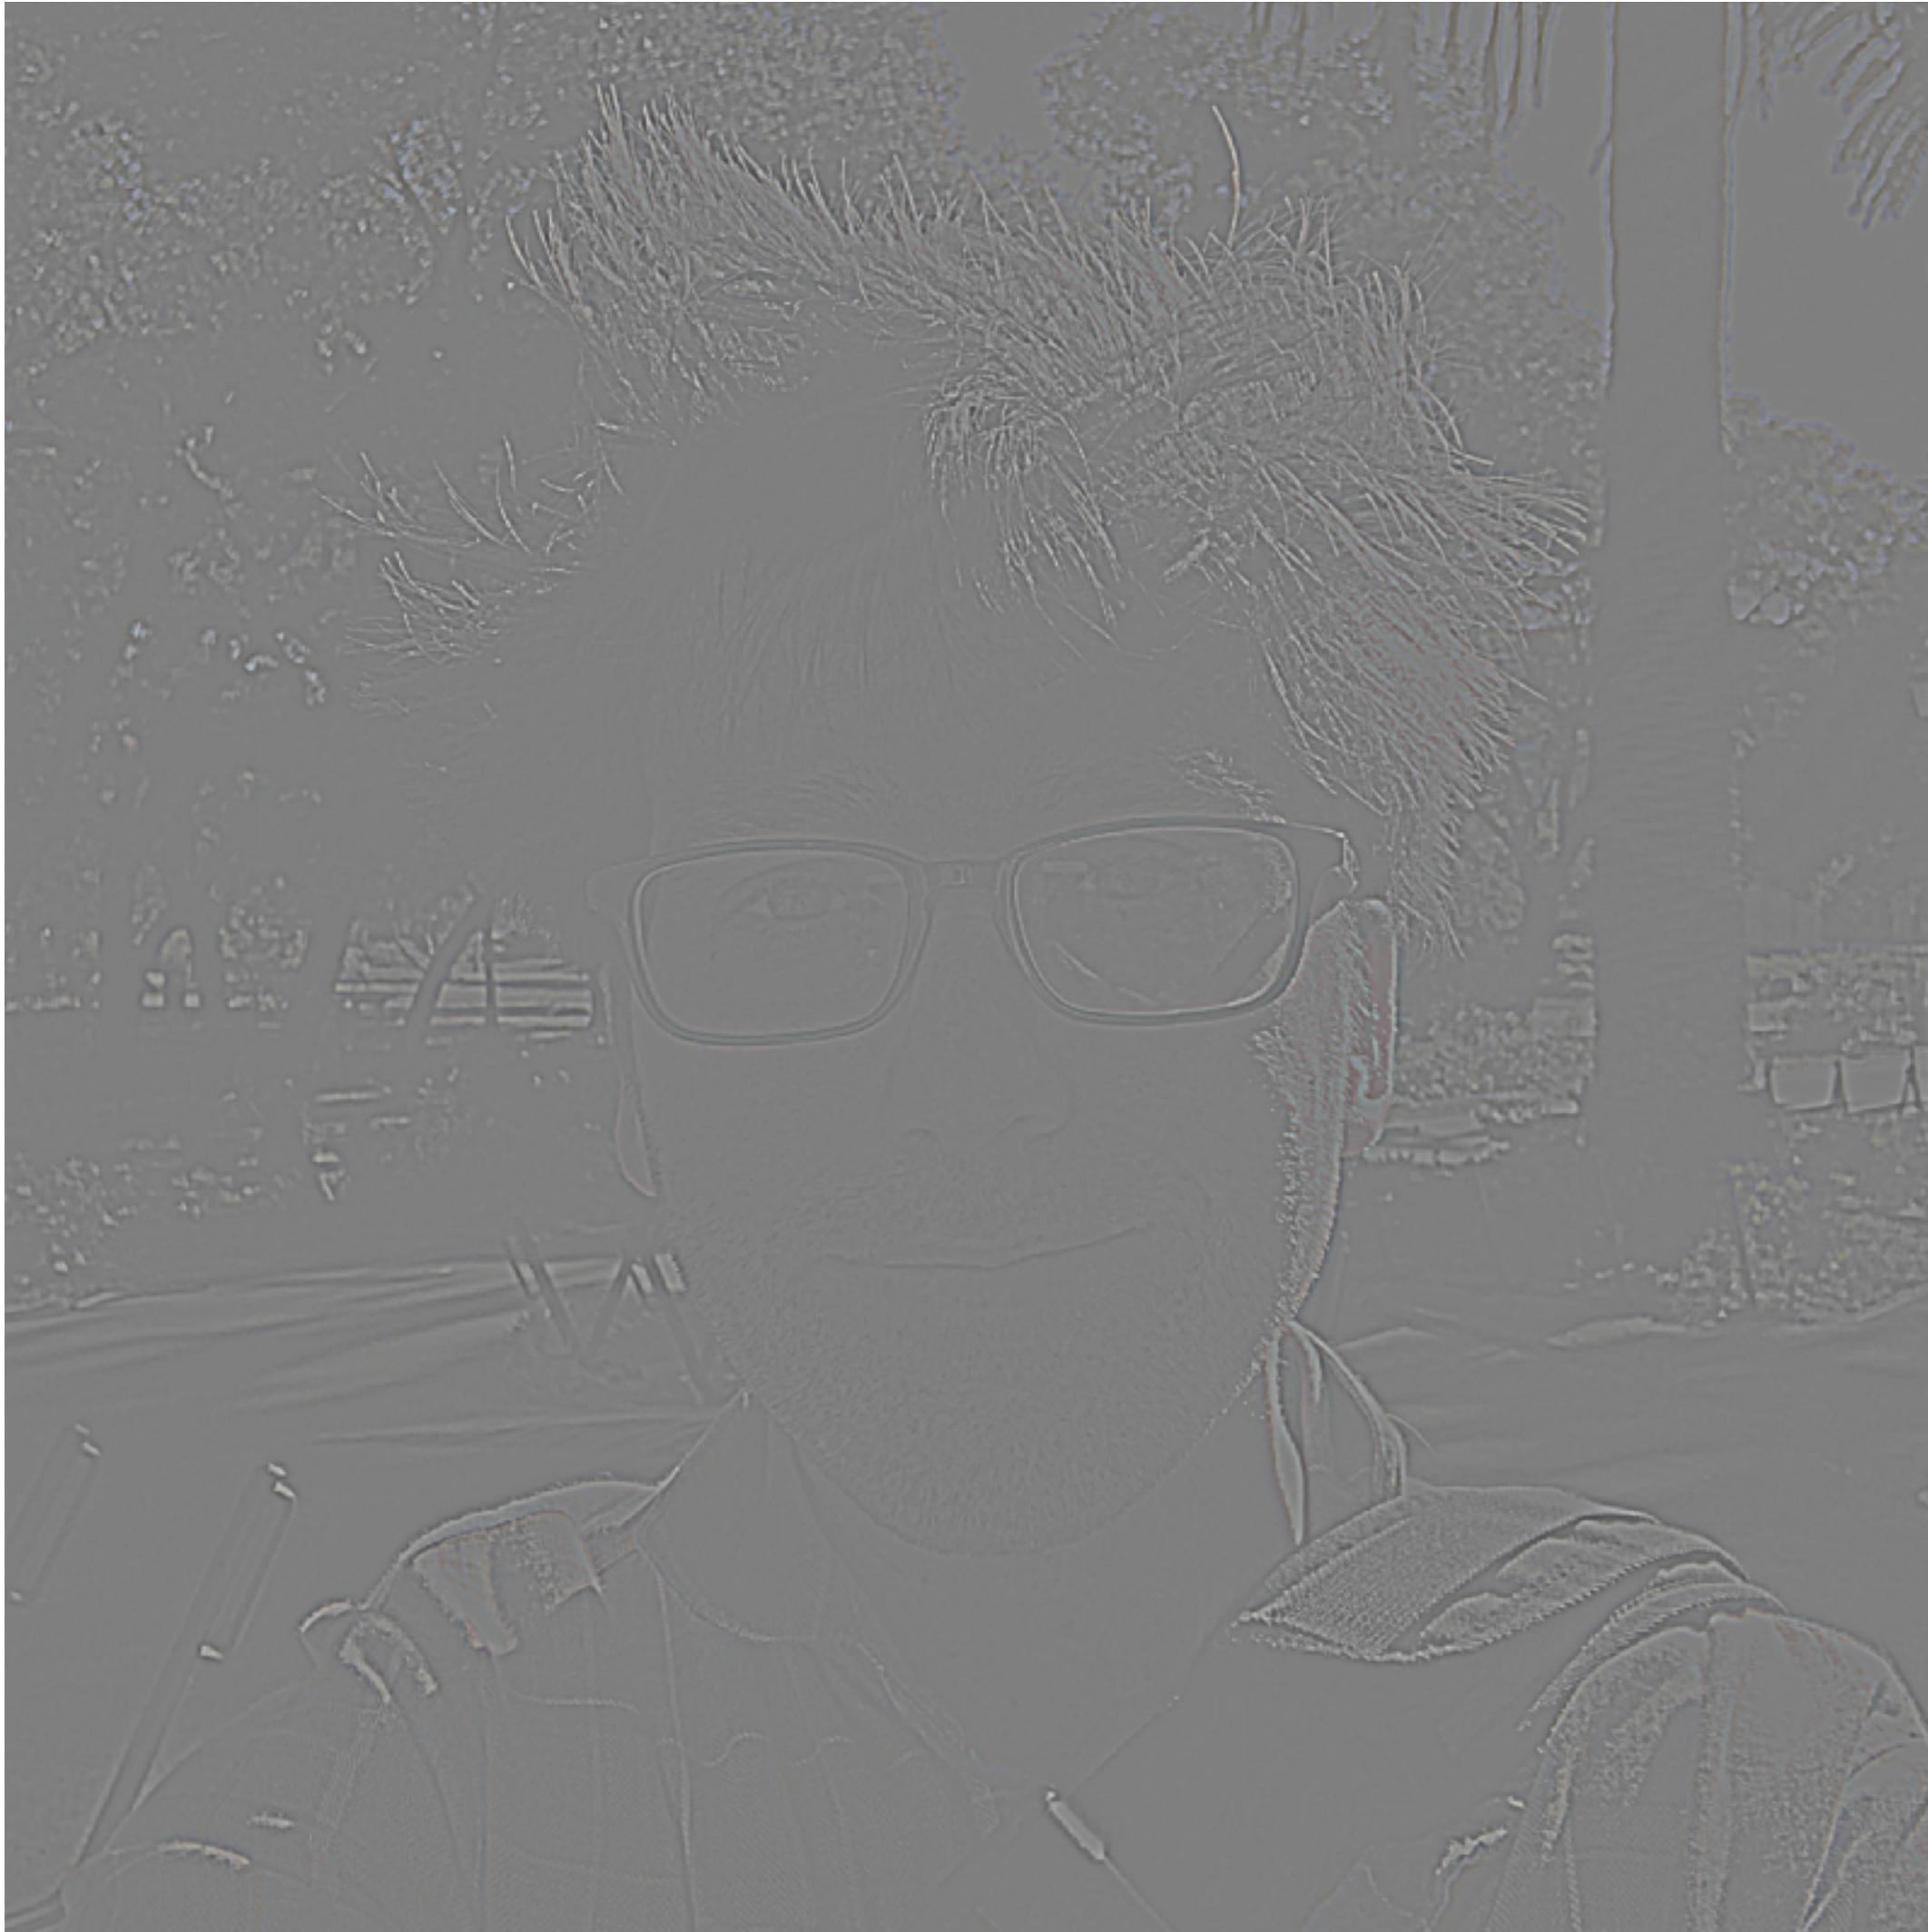
**Question: how do you
reconstruct original image
from its Laplacian pyramid?**

Laplacian pyramid



$$\mathbf{L}_0 = \mathbf{G}_0 - \mathbf{up}(\mathbf{G}_1)$$

Laplacian pyramid



$$L_1 = G_1 - up(G_2)$$

Laplacian pyramid



$$L_2 = G_2 - up(G_3)$$

Laplacian pyramid



$$L_3 = G_3 - up(G_4)$$

Laplacian pyramid



$$L_4 = G_4 - up(G_5)$$

Laplacian pyramid

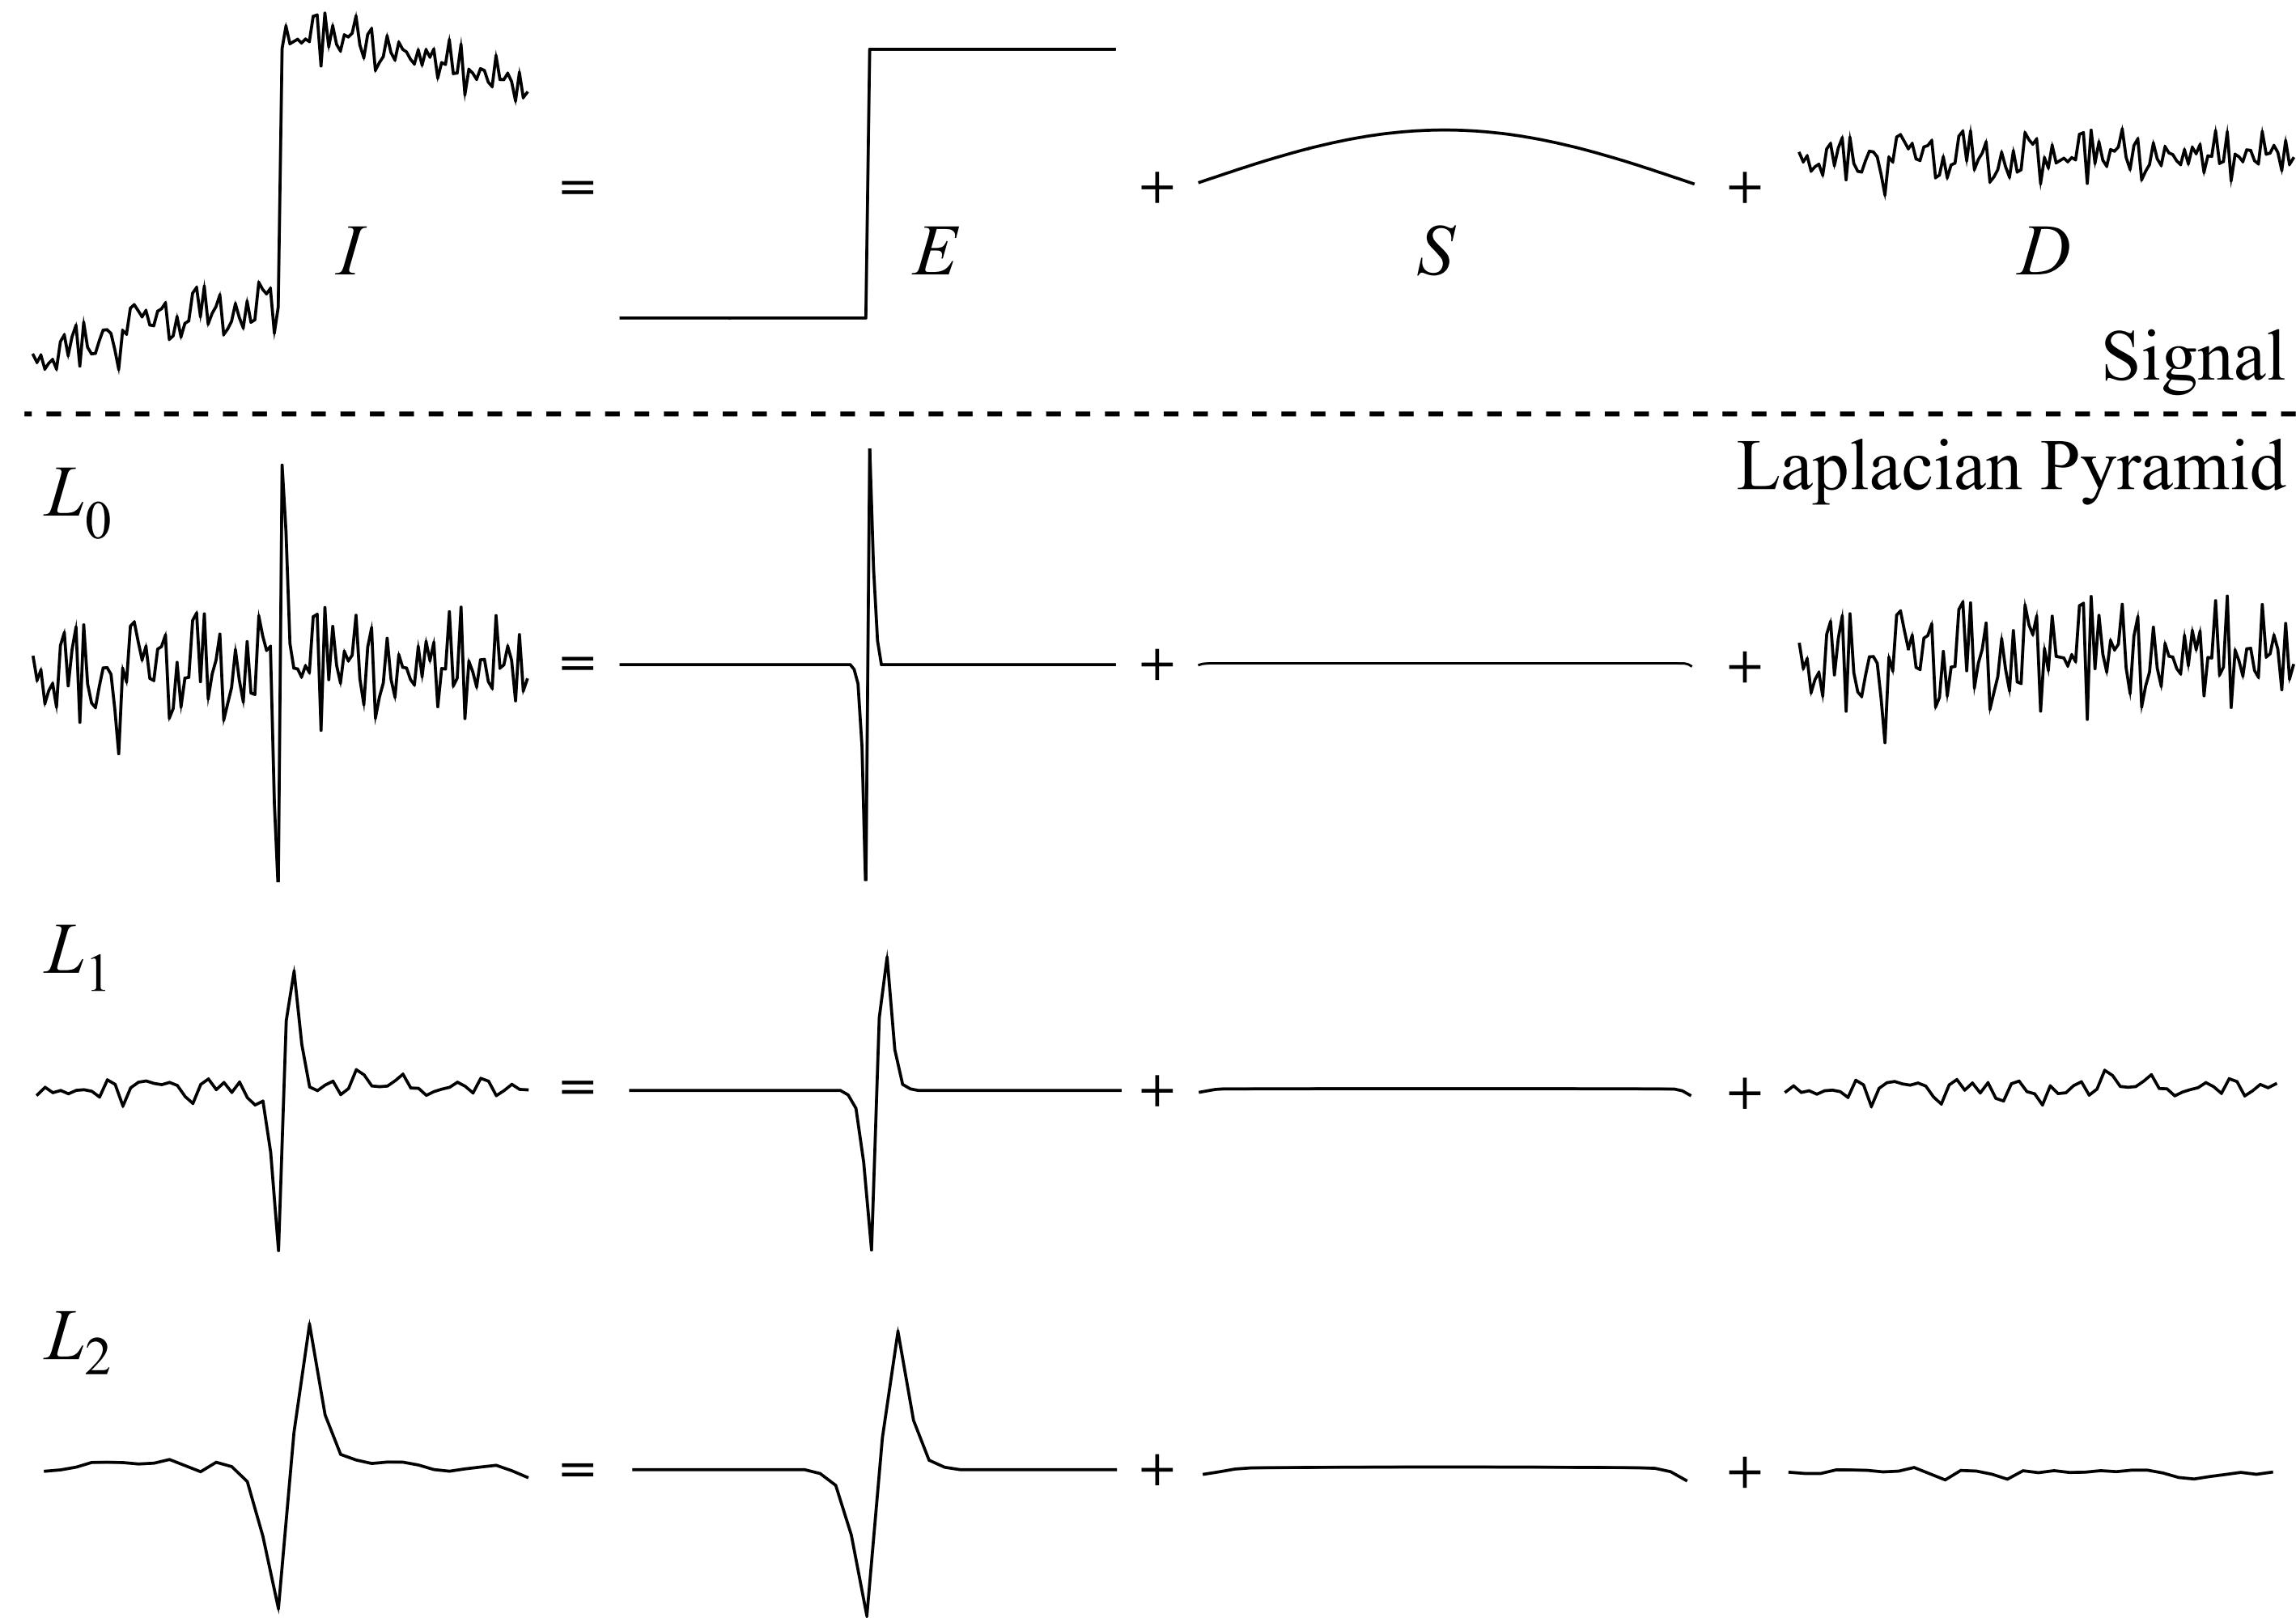


$$\mathbf{L}_5 = \mathbf{G}_5$$

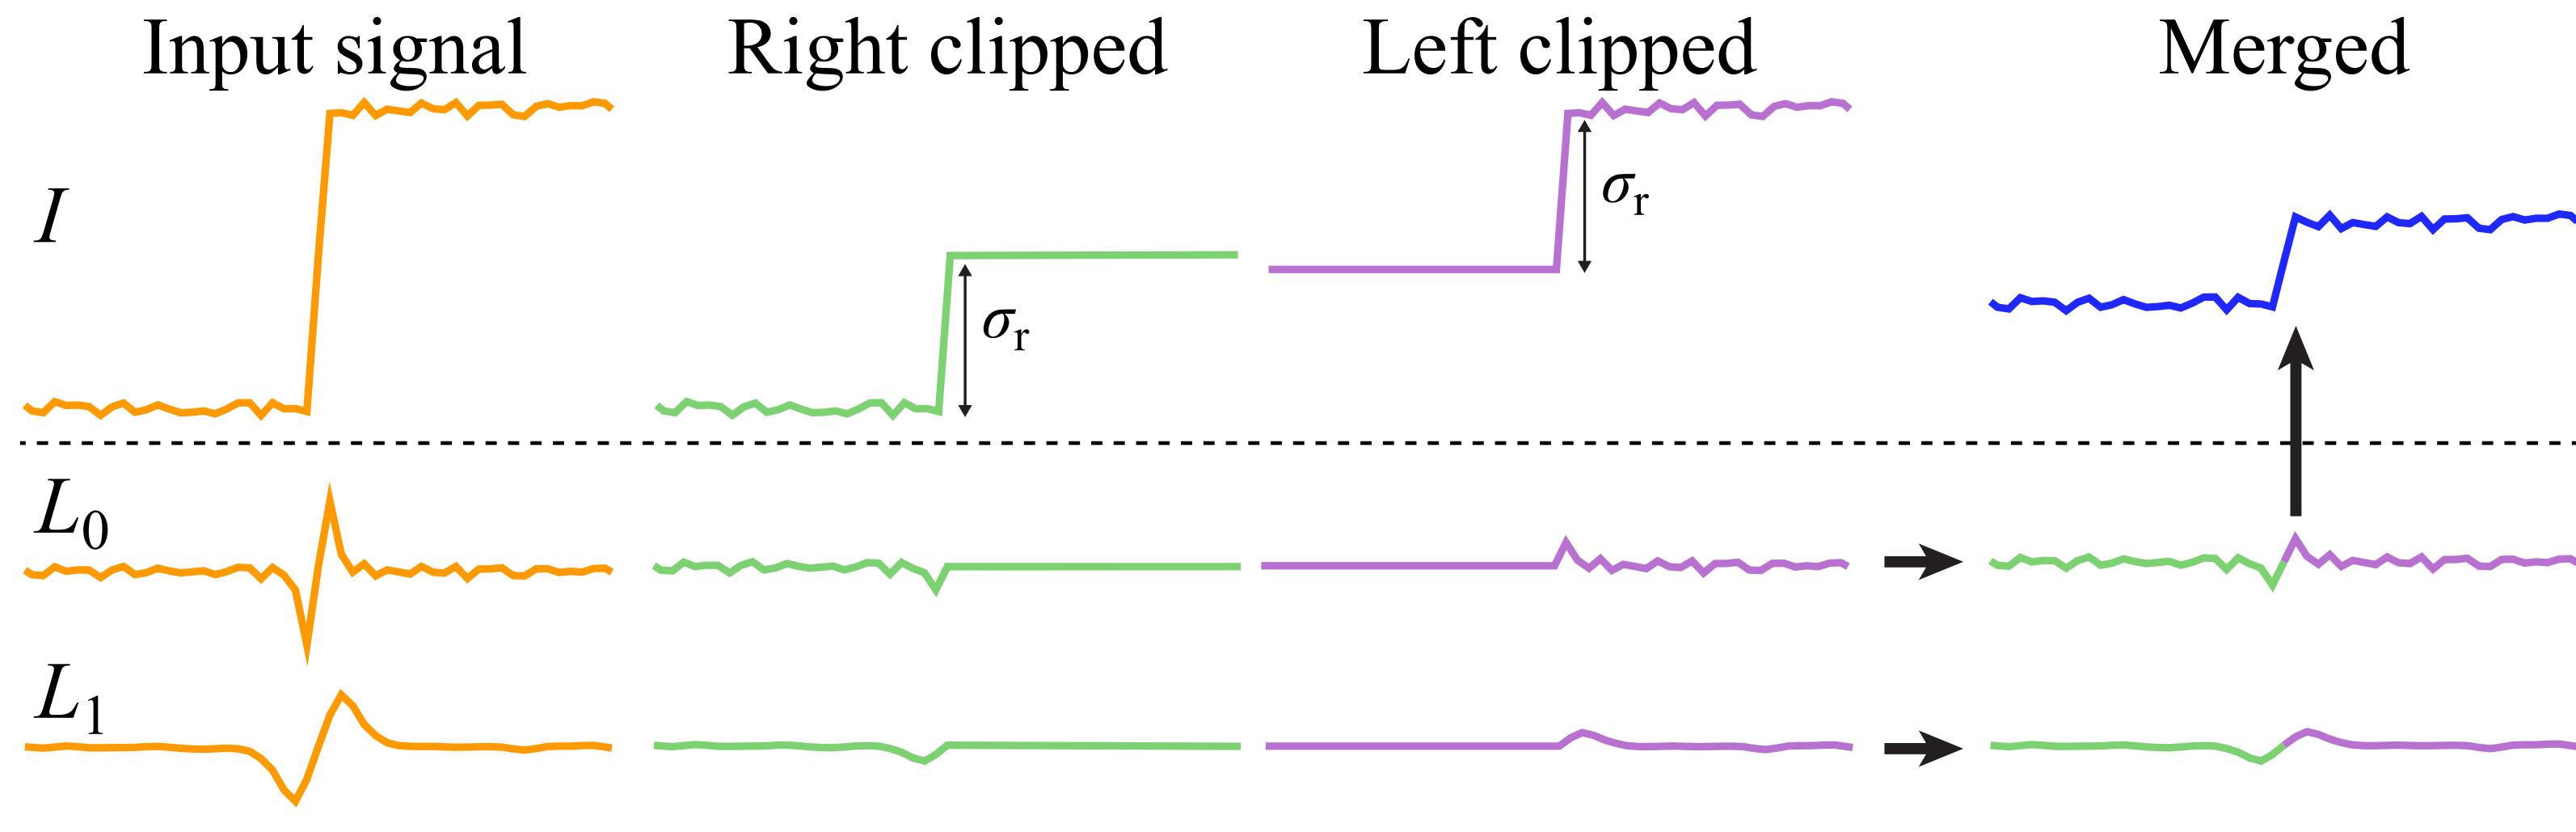
Summary

- **Gaussian and Laplacian pyramids are image representations where each pixel maintains information about frequency content in a region of the image**
- **$G_i(x,y)$ — frequencies up to limit given by i**
- **$L_i(x,y)$ — frequencies added to G_{i+1} to get G_i**
- **Notice: to boost band of frequencies in image around pixel (x,y) , increase coefficient $L_i(x,y)$ in Laplacian pyramid**

Reducing range of an edge (without reducing detail)



Reducing range of an edge (without reducing detail)



Local Laplacian filtering:

Compute $G_i(x, y)$ — semi-local statistic around (x, y) (less global for larger i)

Modify original image I based on $G_i(x, y)$ to get $I' = f(I)$

Here: $f()$ clamps values that are far from $G_i(x, y)$

Compute Laplacian pyramid from I' : $L(I')$

Set coefficient $L_i(x, y)$ of output Laplacian pyramid to $L(I')_i(x, y)$

Collapse output Laplacian pyramid to obtain final image.

Local laplacian filtering example



Original Image



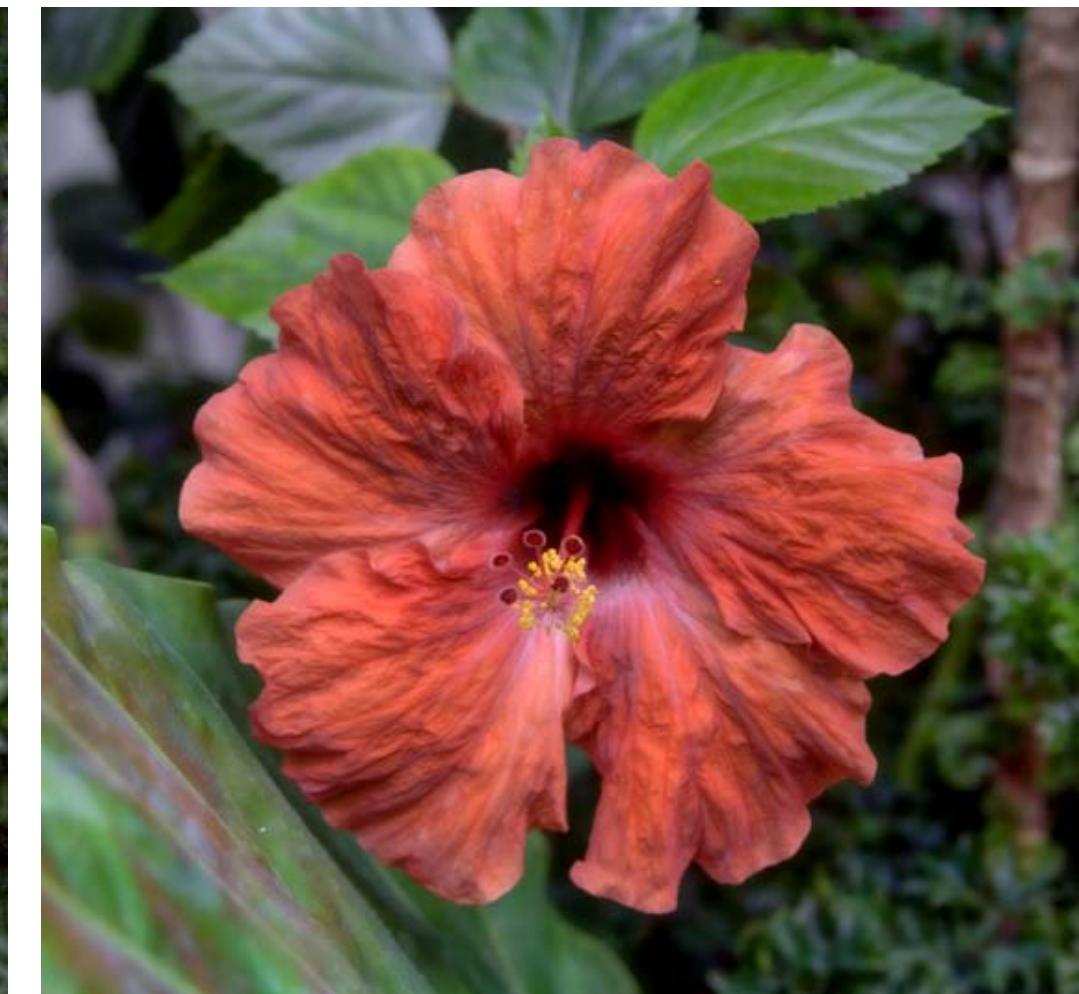
Detail removed



Detail increased
(all levels)



Detail increased
(lowest two levels)



Detail increased
(level 3 and higher)

Local laplacian filtering example



(a) input HDR image tone-mapped with a simple gamma curve (details are compressed)



(b) our pyramid-based tone mapping, set to preserve details without increasing them



(c) our pyramid-based tone mapping, set to strongly enhance the contrast of details

Image processing workload characteristics

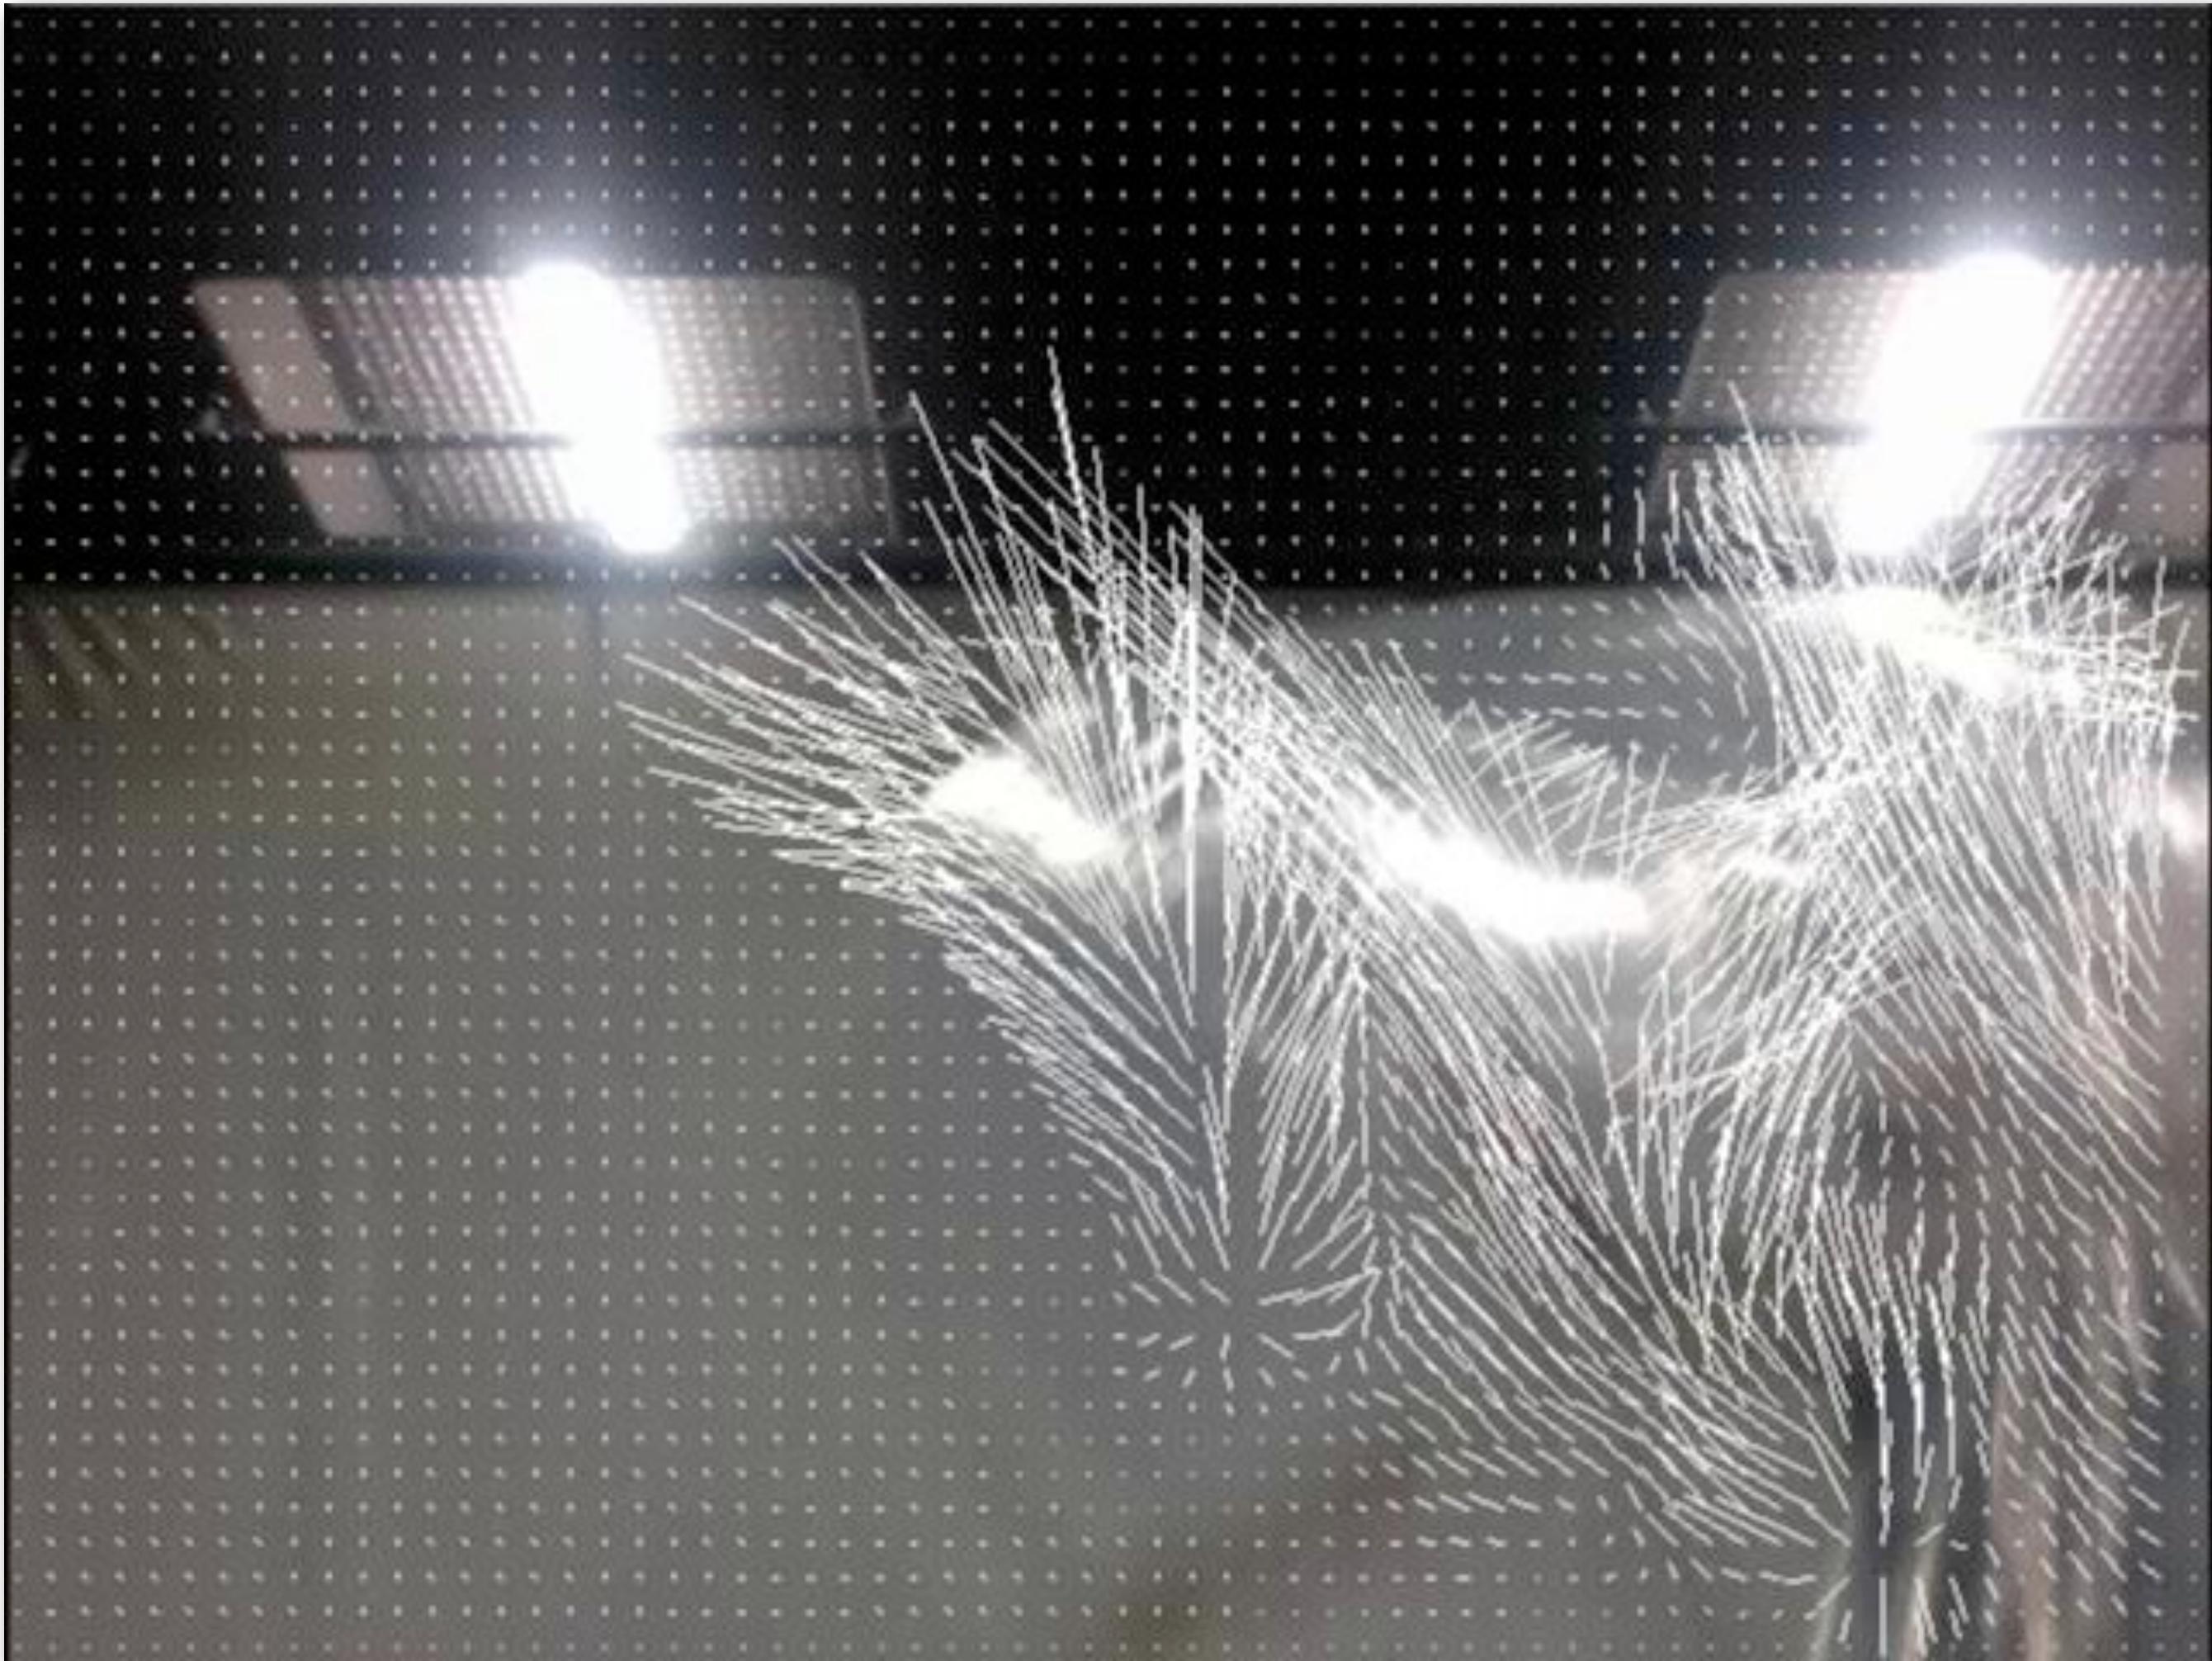
- “Pointwise operations”:
 - $\text{output_pixel} = f(\text{input_pixel})$
- “Stencil” computations (e.g., convolution, demosaic, etc.)
 - output pixel (x,y) depends on fixed-size local region of input around (x,y)
- Lookup tables
- Multi-resolution operations (upsampling/downsampling)
- Long sequences of operations

Bonus slides:

Estimating motion using optical flow

Optical flow

Goal: determine 2D screen-space velocity of visible objects in image



Optical flow

- Given image A (at time t) and image B (at time $t + \Delta t$) compute the per-pixel motion needed to correspond the two images
- Major assumption 1: “brightness constancy”
 - The appearance of a scene surface point that is visible in both images A and B is the same in both images

$$I(x, y, t) = I(x + \Delta x, y + \Delta y, t + \Delta t)$$

← The point observed at (x, y) at time t
moves to $(x + \Delta x, y + \Delta y)$ at $t + \Delta t$
(and has the same appearance in both situations)

Taylor expansion

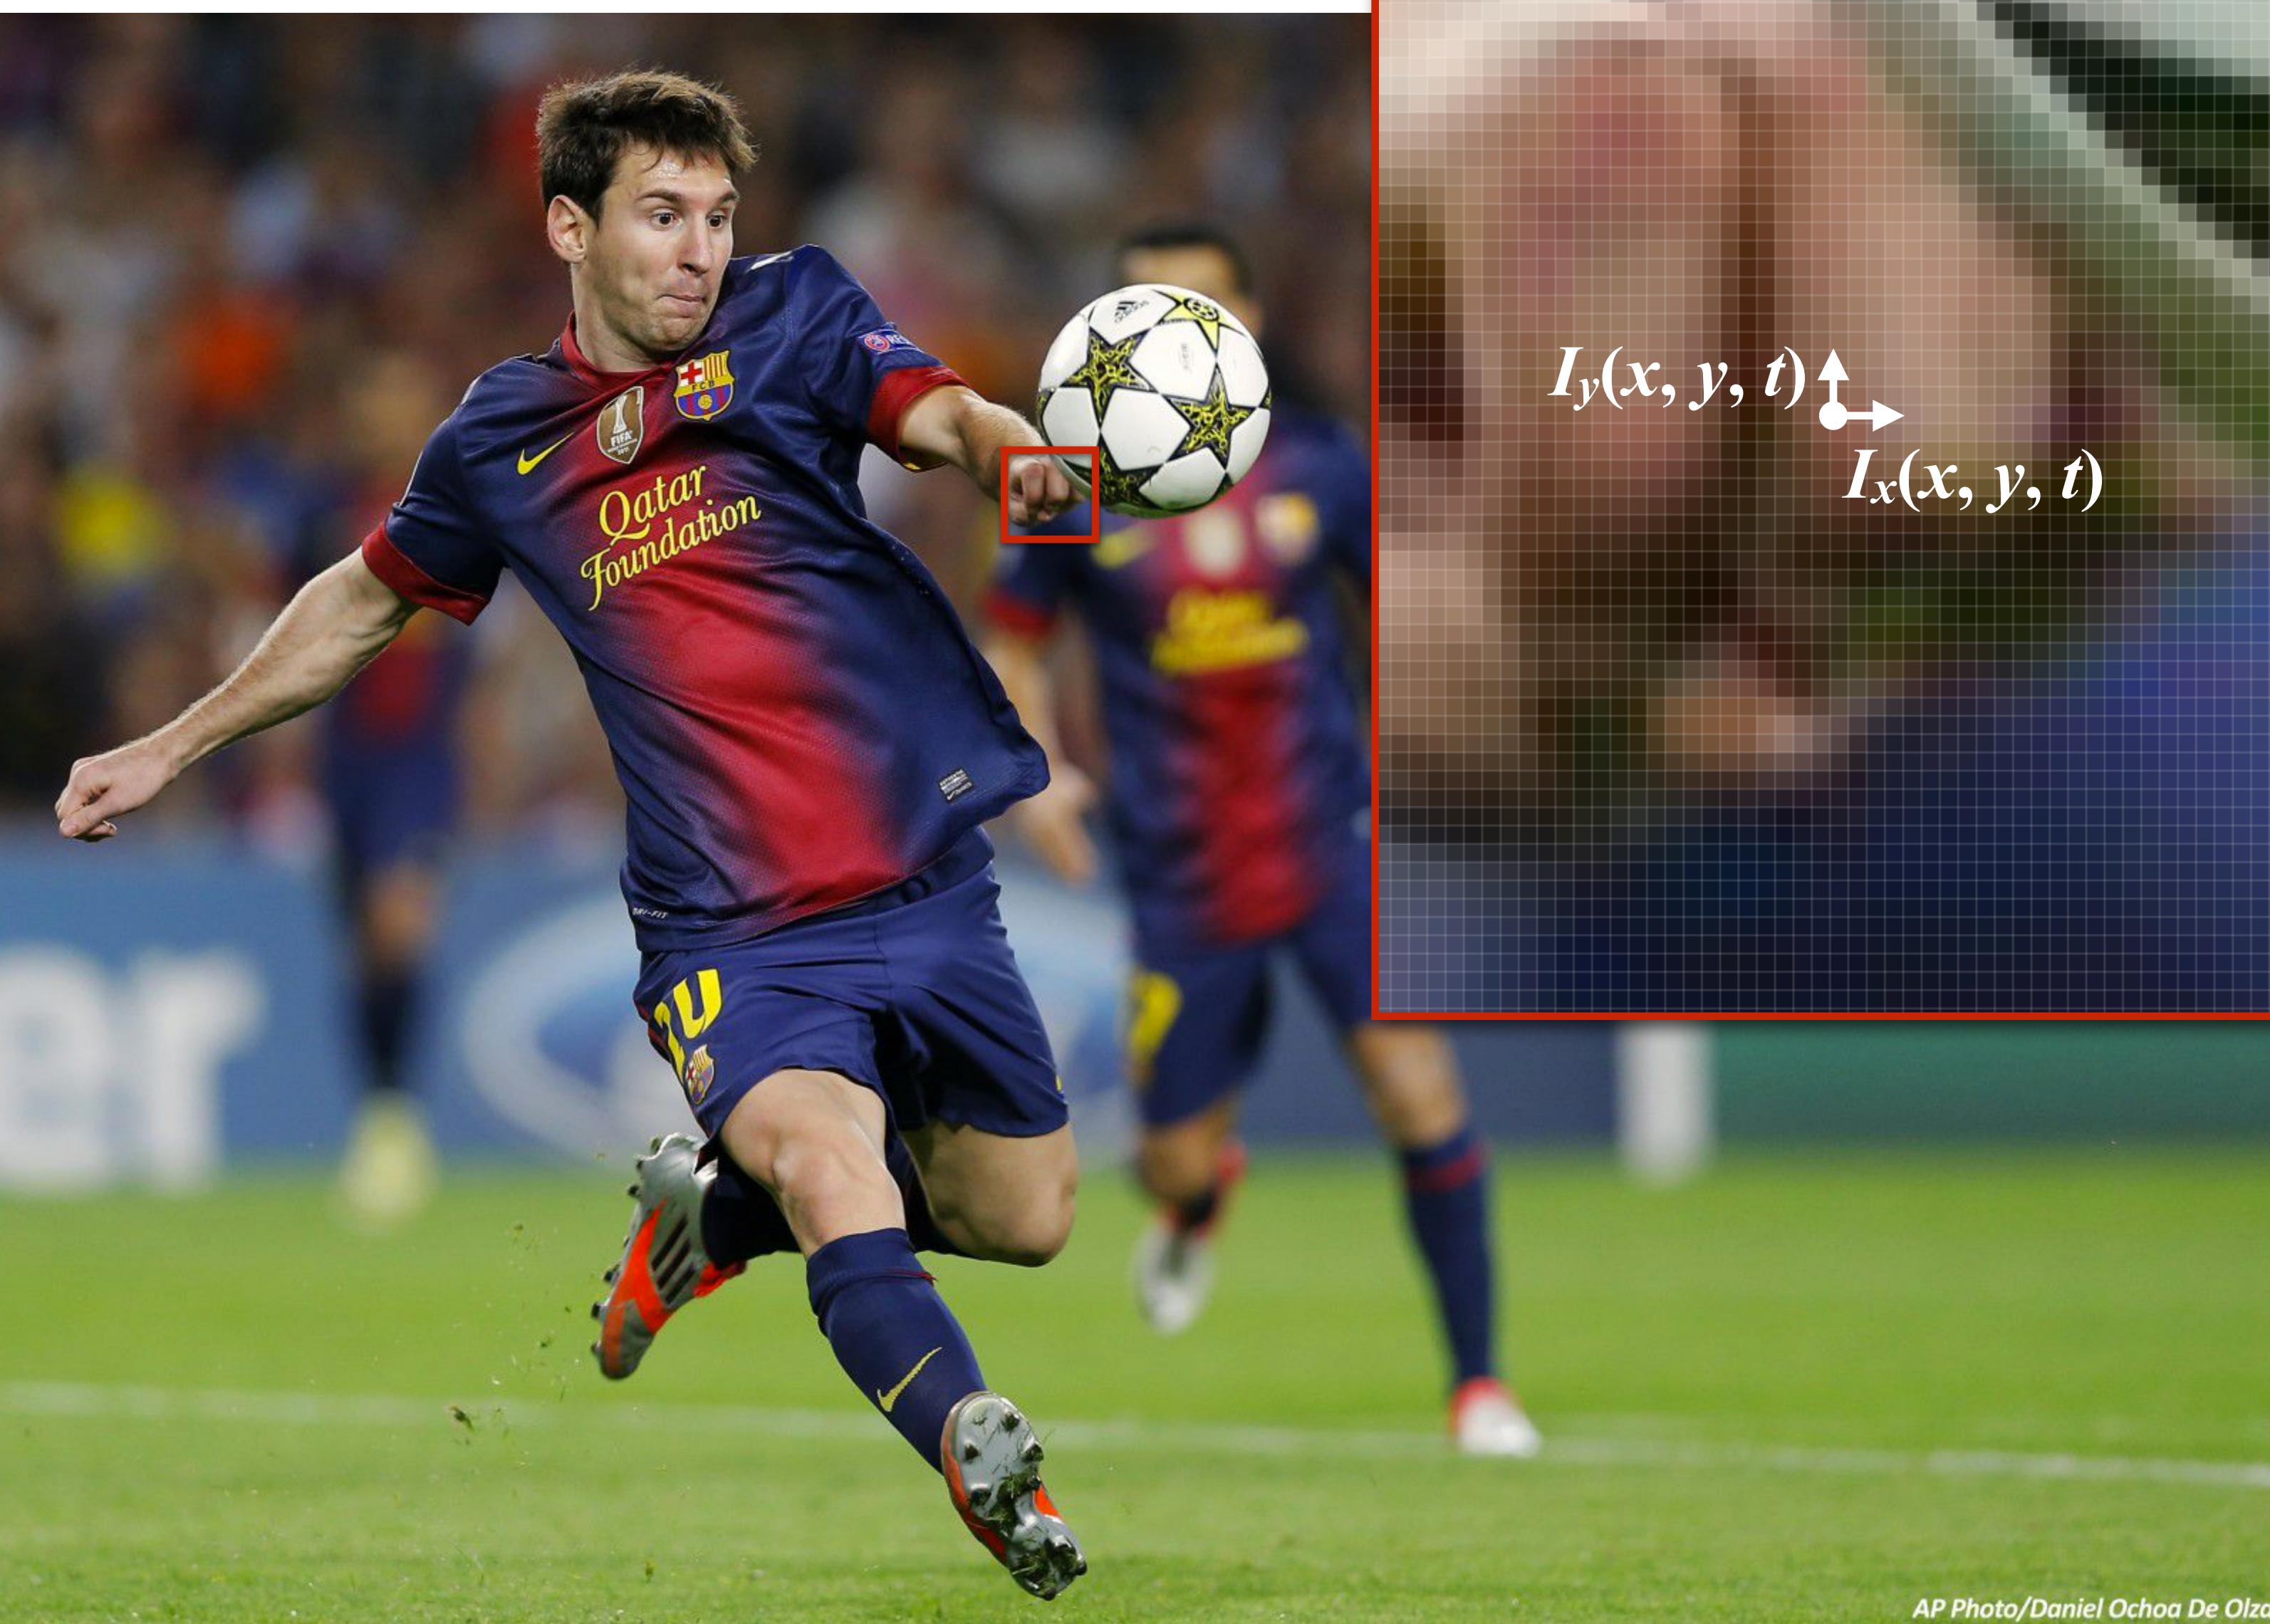
$$I(x + \Delta x, y + \Delta y, t + \Delta t) = I(x, y, t) + I_x(x, y, t)\Delta x + I_y(x, y, t)\Delta y + I_t(x, y, t)\Delta t + \text{higher order terms}$$

So...

$$I(x, y, t) \approx I(x, y, t) + I_x(x, y, t)\Delta x + I_y(x, y, t)\Delta y + I_t(x, y, t)\Delta t$$

$$I_x(x, y, t)\Delta x + I_y(x, y, t)\Delta y + I_t(x, y, t)\Delta t = 0$$

↑
The observed change in pixel (x, y)
Is due to object motion at point by $(\Delta x, \Delta y)$



AP Photo/Daniel Ochoa De Olza



Gradient-constraint equation for a pixel is underconstrained

Gradient-constraint equation is insufficient to solve for motion

One equation, two unknowns: $(\Delta x, \Delta y)$

$$I_x(x, y, t)\Delta x + I_y(x, y, t)\Delta y + I_t(x, y, t)\Delta t = 0$$

Known: spatial image gradients in image A

Known: observed change in pixel (x, y) over consecutive frames

Major assumption 2: nearby pixels have similar motion (Lucas-Kanade)

$$I_x(x_0, y_0, t)\Delta x + I_y(x_0, y_0, t)\Delta y + I_t(x_0, y_0, t)\Delta t = 0$$

$$I_x(x_1, y_1, t)\Delta x + I_y(x_1, y_1, t)\Delta y + I_t(x_1, y_1, t)\Delta t = 0$$

$$I_x(x_2, y_2, t)\Delta x + I_y(x_2, y_2, t)\Delta y + I_t(x_2, y_2, t)\Delta t = 0$$

⋮

Now we have an overconstrained system, compute least squares solution

Weighted least-squares solution

Gradient-constraint equation is insufficient to solve for motion

One equation, two unknowns: $(\Delta x, \Delta y)$

$$I_x(x_0, y_0, t)\Delta x + I_y(x_0, y_0, t)\Delta y + I_t(x_0, y_0, t)\Delta t = 0$$

$$I_x(x_1, y_1, t)\Delta x + I_y(x_1, y_1, t)\Delta y + I_t(x_1, y_1, t)\Delta t = 0$$

$$I_x(x_2, y_2, t)\Delta x + I_y(x_2, y_2, t)\Delta y + I_t(x_2, y_2, t)\Delta t = 0$$

•
•
•

Compute weighted least squares solution by minimizing:

$$E(\Delta x, \Delta y) = \sum_{x_i, y_i} w(x_i, y_i, x, y) [I_x(x_i, y_i, t)\Delta x + I_y(x_i, y_i, t)\Delta y + I_t(x_i, y_i, t)\Delta t]^2$$

(x_i, y_i) are pixels in region around (x, y) .

**Weighting function $w()$ weights error contribution based on distance between (x_i, y_i) and (x, y)
e.g., Gaussian fall-off.**

Solving for motion

$E(\Delta x, \Delta y)$ minimized when derivatives are zero:

$$\frac{dE(\Delta x, \Delta y)}{d(\Delta x)} = \sum_{x_i, y_i} w(x_i, y_i, x, y) \left[I_x^2 \Delta x + I_x I_y \Delta y + I_x I_t \right] = 0$$

$$\frac{dE(\Delta x, \Delta y)}{d(\Delta y)} = \sum_{x_i, y_i} w(x_i, y_i, x, y) \left[I_y^2 \Delta y + I_x I_y \Delta x + I_y I_t \right] = 0$$

Rewrite, now solve the following linear system for $\Delta x, \Delta y$:

$$\begin{array}{l} \mathbf{A0} \\ \Delta x \sum_{x_i, y_i} w(x_i, y_i, x, y) I_x^2 + \Delta y \sum_{x_i, y_i} w(x_i, y_i, x, y) I_x I_y + \sum_{x_i, y_i} w(x_i, y_i, x, y) I_x I_t = 0 \\ \mathbf{A1} \\ \Delta x \sum_{x_i, y_i} w(x_i, y_i, x, y) I_x I_y + \Delta y \sum_{x_i, y_i} w(x_i, y_i, x, y) I_y^2 + \sum_{x_i, y_i} w(x_i, y_i, x, y) I_y I_t = 0 \end{array}$$

Precompute partial derivatives I_x, I_y, I_t from original images A and B

For each pixel (x, y) : evaluate **A0**, **B0**, **C0**, **A1**, **B1**, **C1**, then solve for $(\Delta x, \Delta y)$ at (x, y)

Optical flow, implemented in practice

Gradient-constraint equation makes a linear motion assumption

$$I(x, y, t) \approx I(x, y, t) + I_x(x, y, t)\Delta x + I_y(x, y, t)\Delta y + I_t(x, y, t)\Delta t$$

$$I_x(x, y, t)\Delta x + I_y(x, y, t)\Delta y + I_t(x, y, t)\Delta t = 0$$

The observed change in pixel (x,y)

Is due to object motion at point by ($\Delta x, \Delta y$)

- **Improvement: iterative techniques use this original flow field to compute higher order residuals (to estimate non-linear motion)**
- **Question: why is it important for optical flow implementation to be very efficient?**
 - **Hint: consider linear-motion assumption**