

Lecture 6:

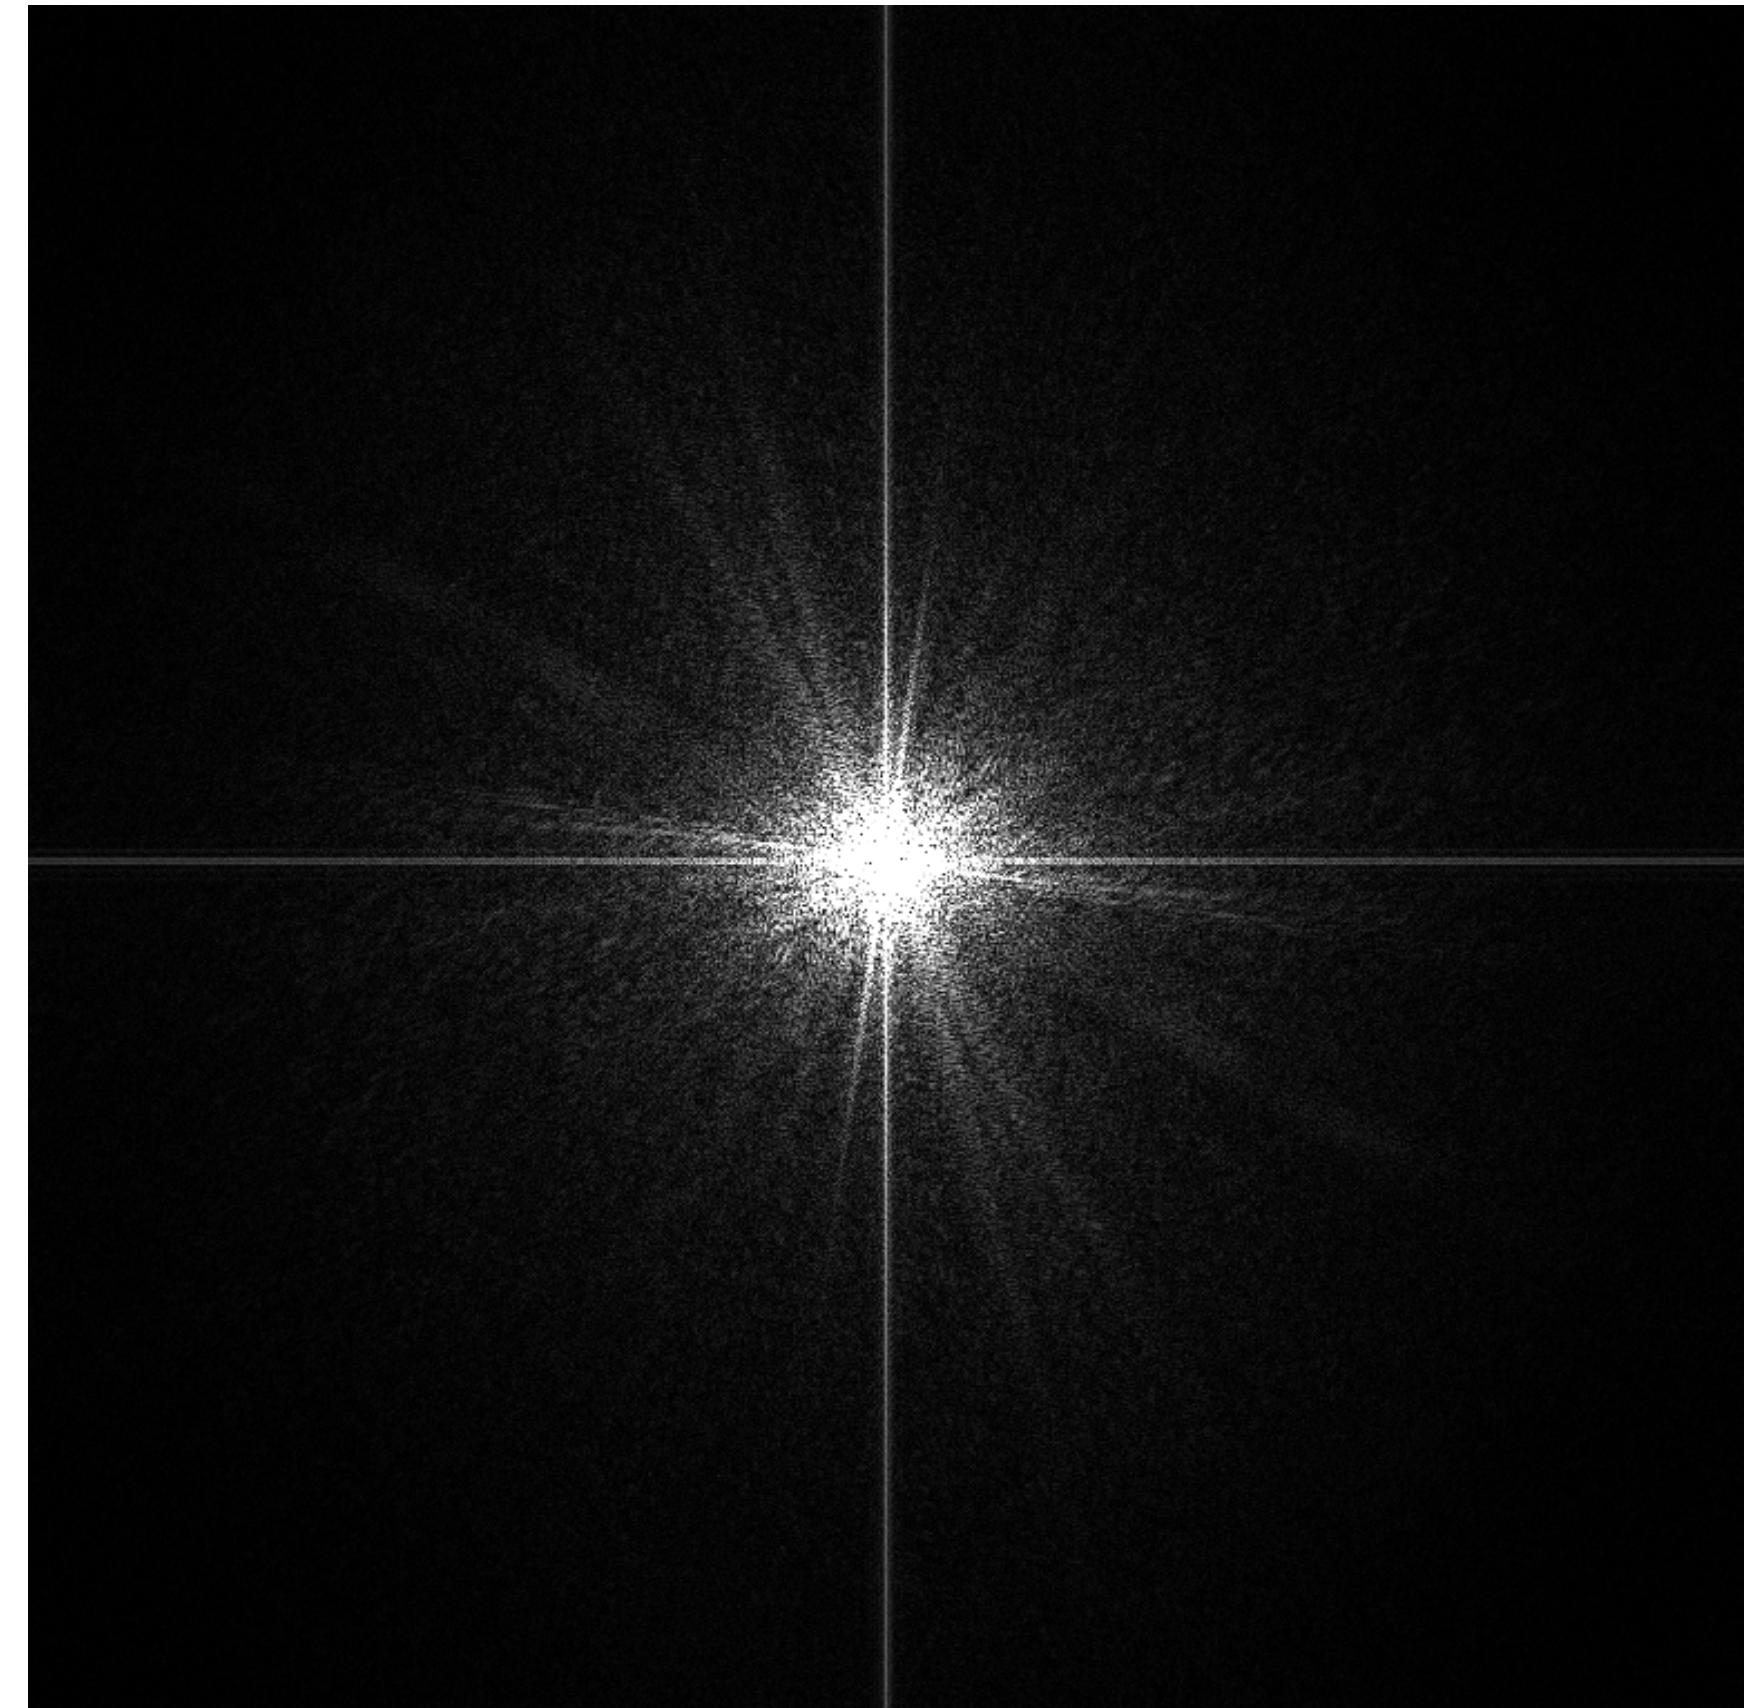
Lossy Image and Video Compression

**Visual Computing Systems
Stanford CS348V, Winter 2018**

Recall: frequency content of images



Spatial domain result

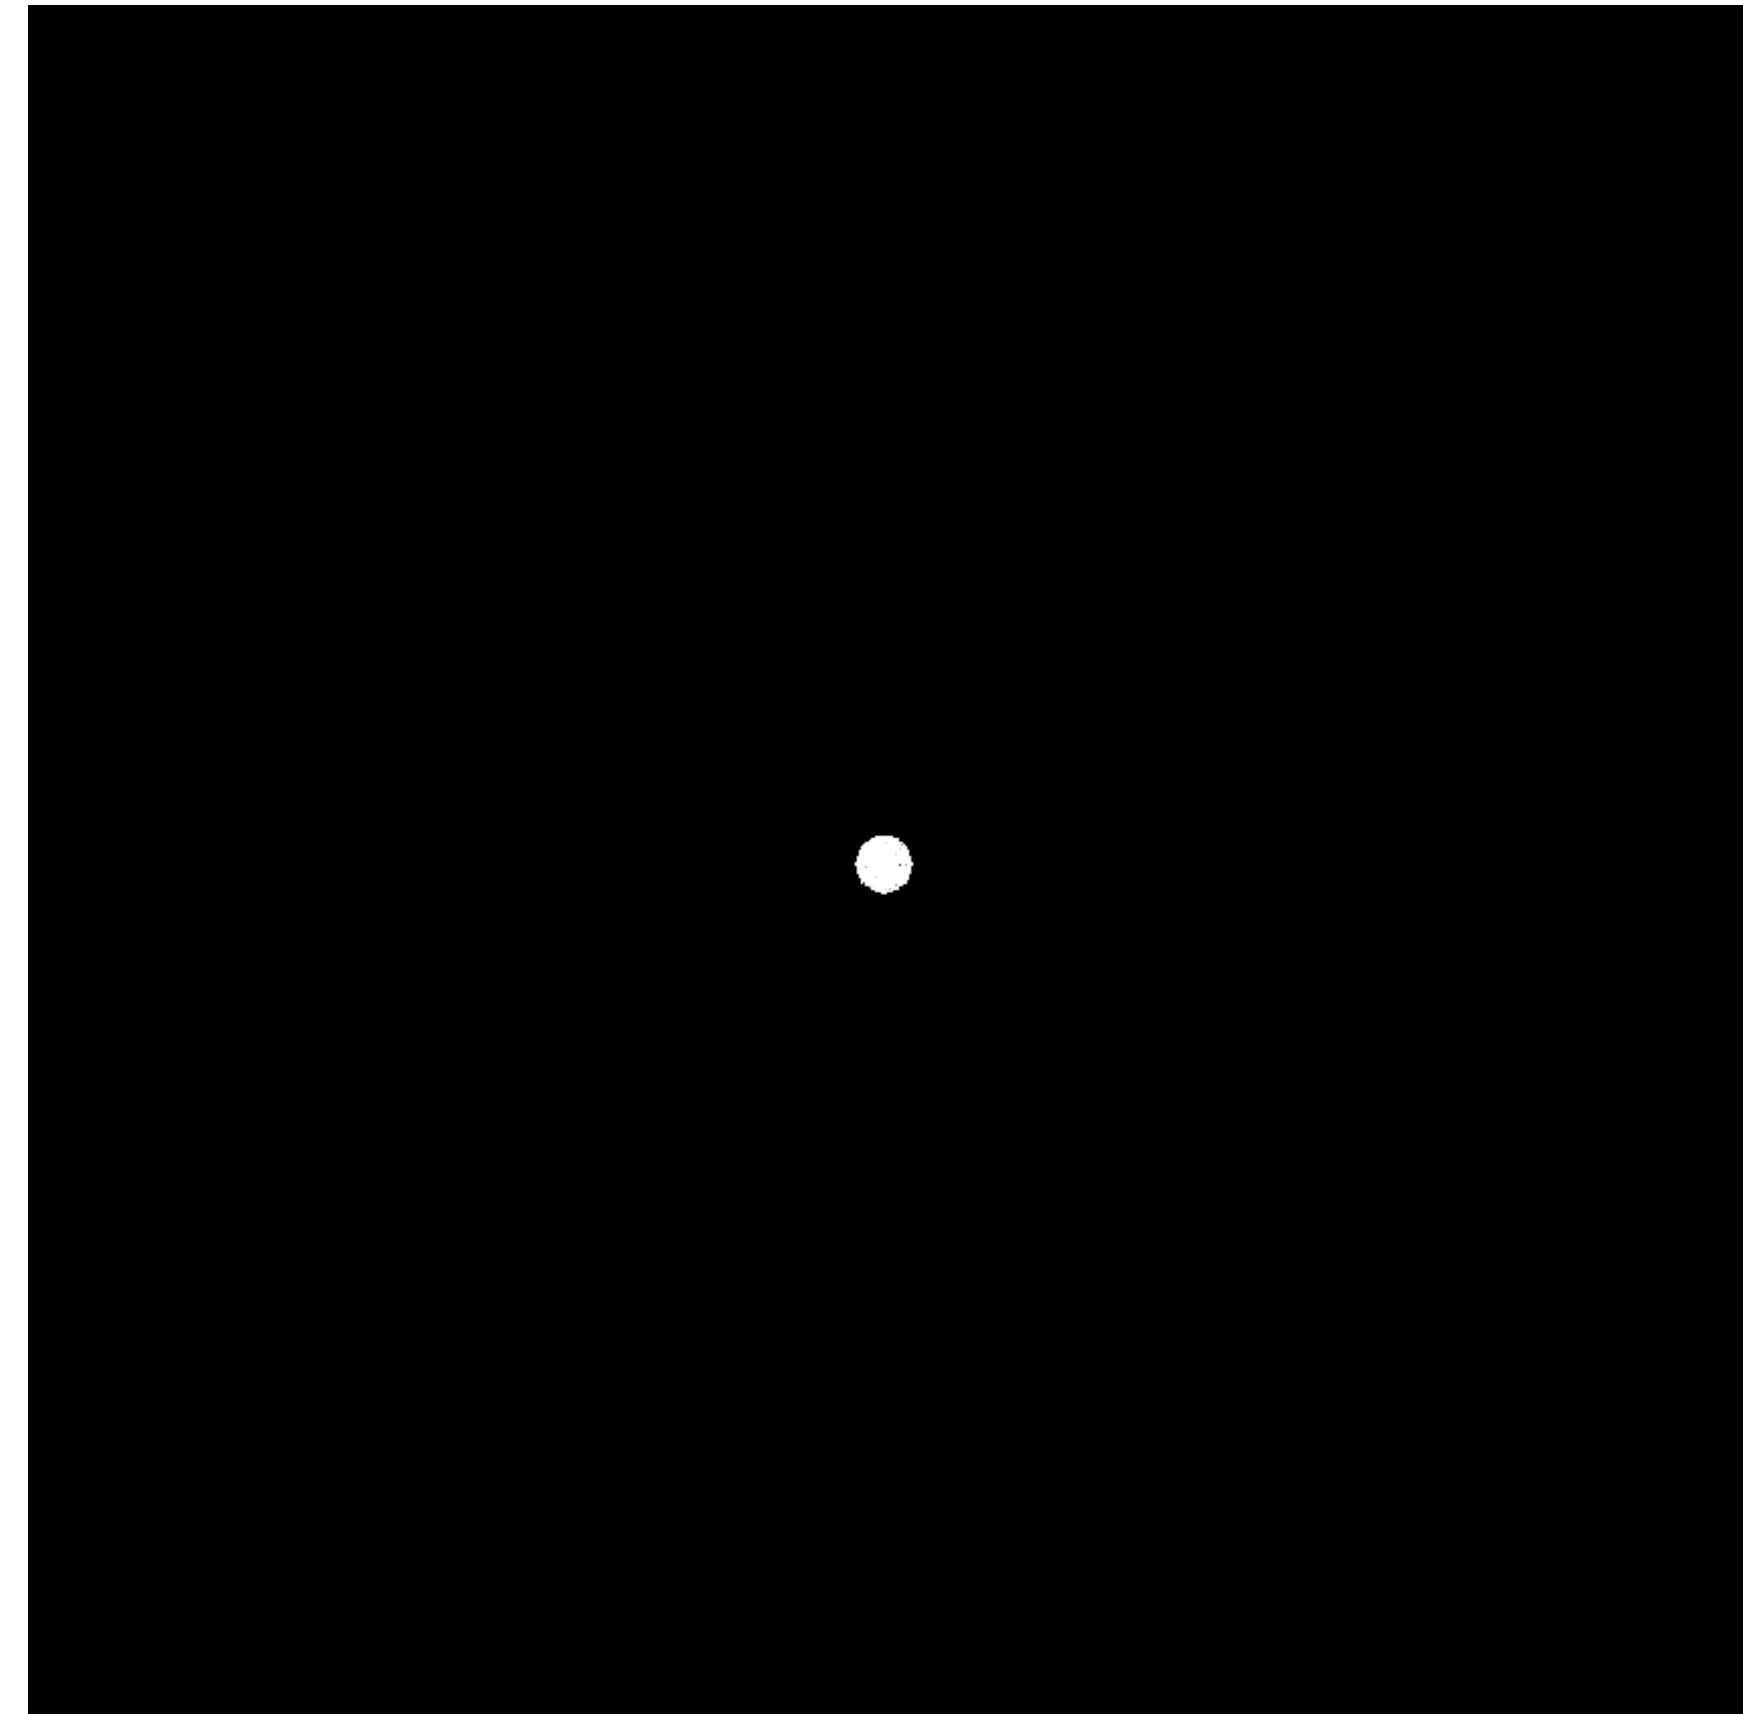


Spectrum

Recall: frequency content of images

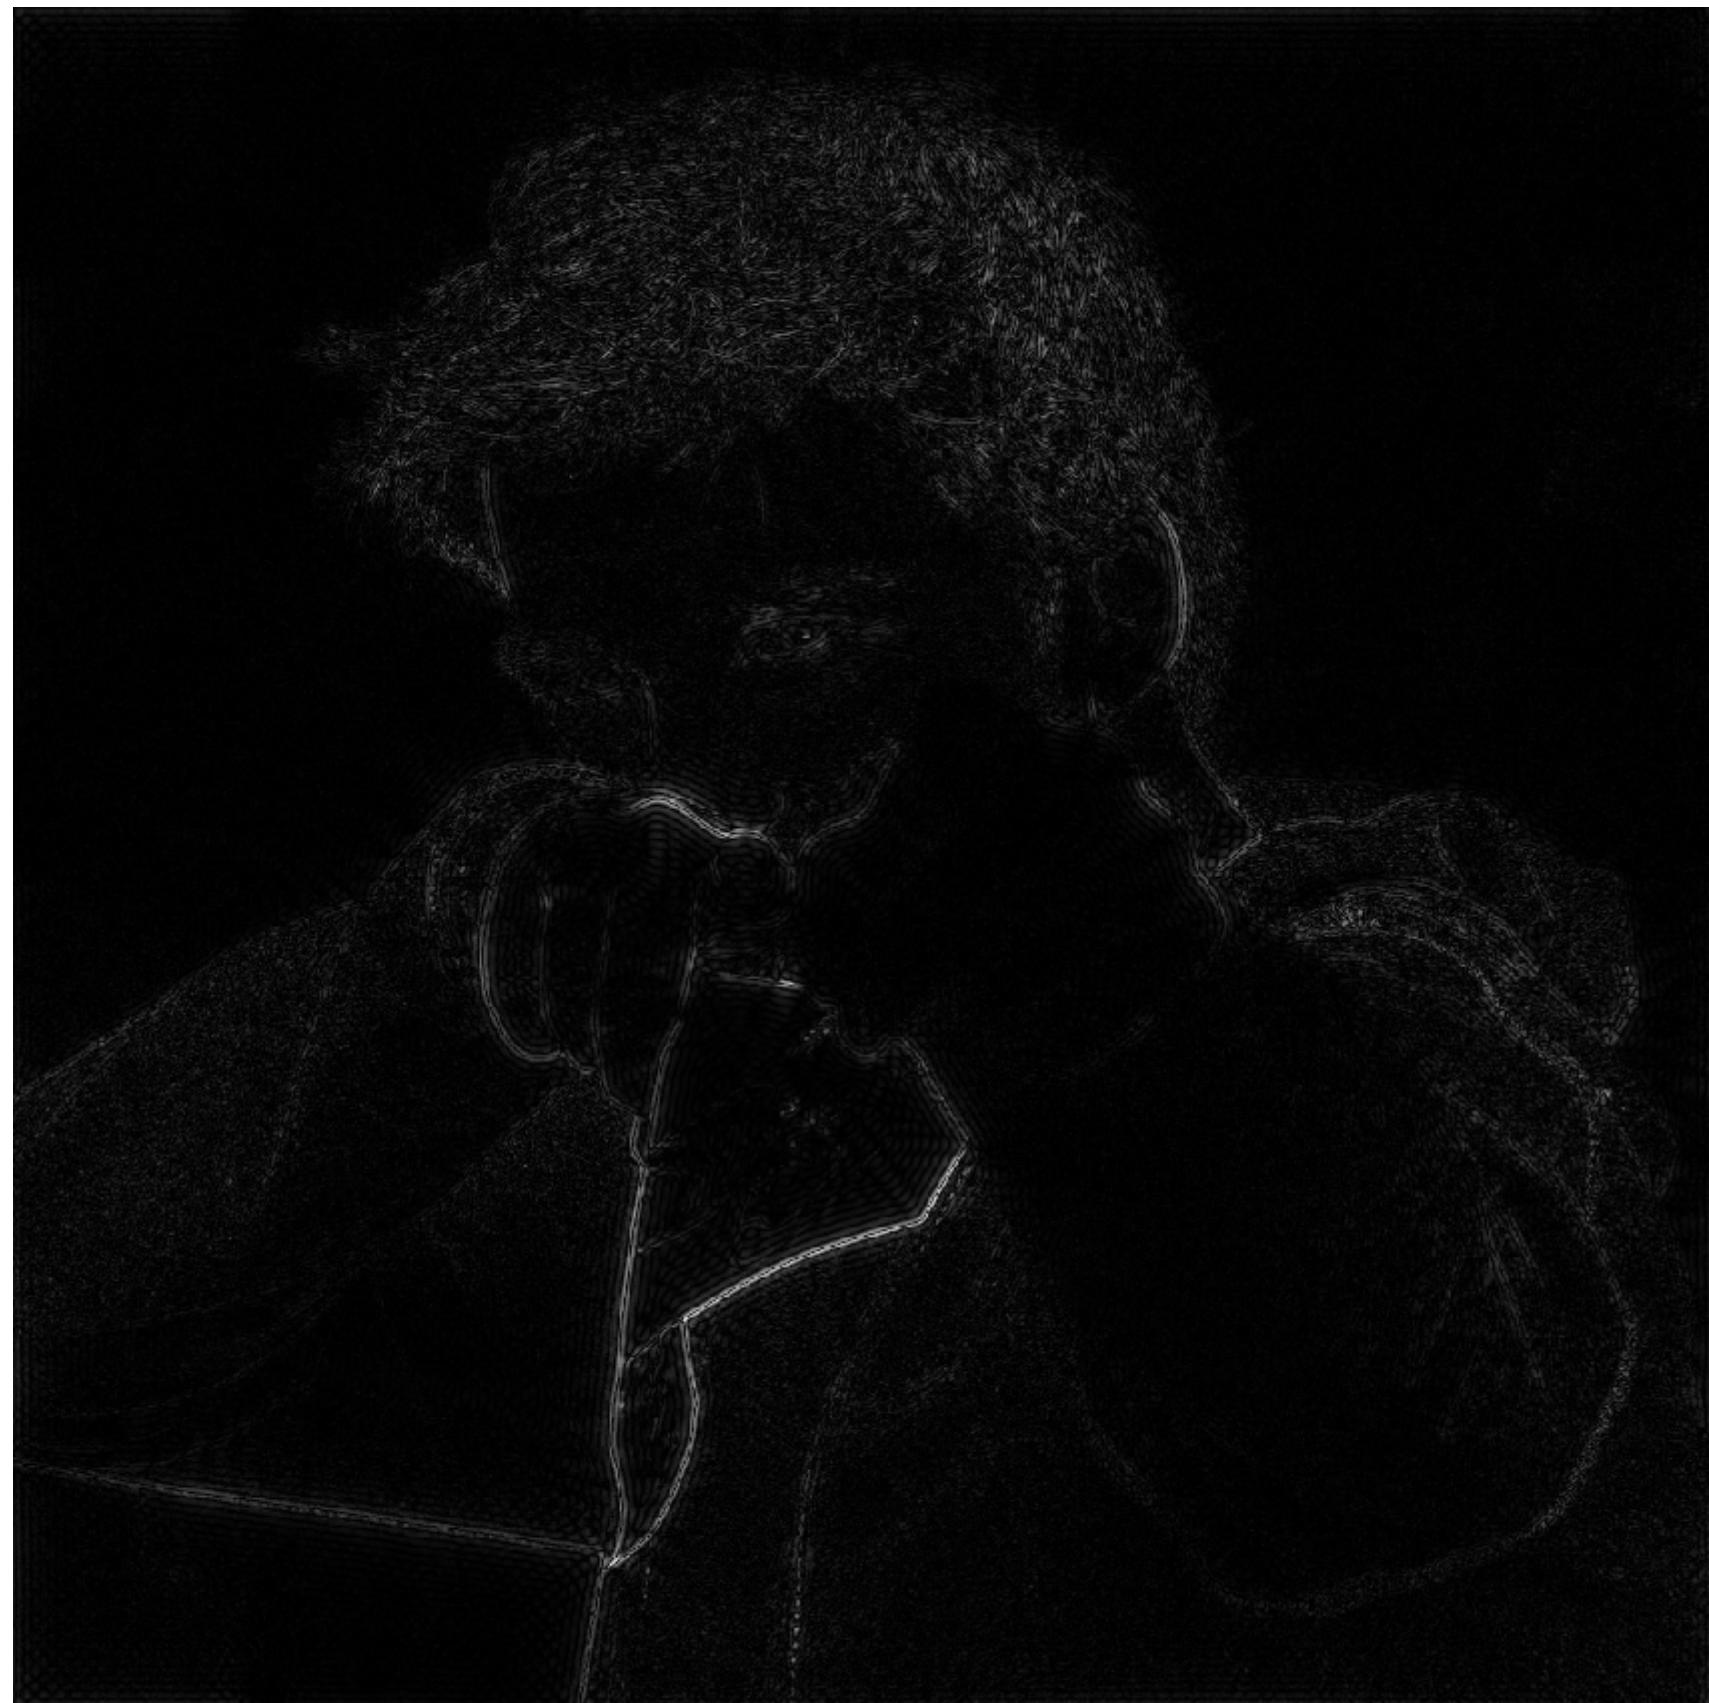


Spatial domain result

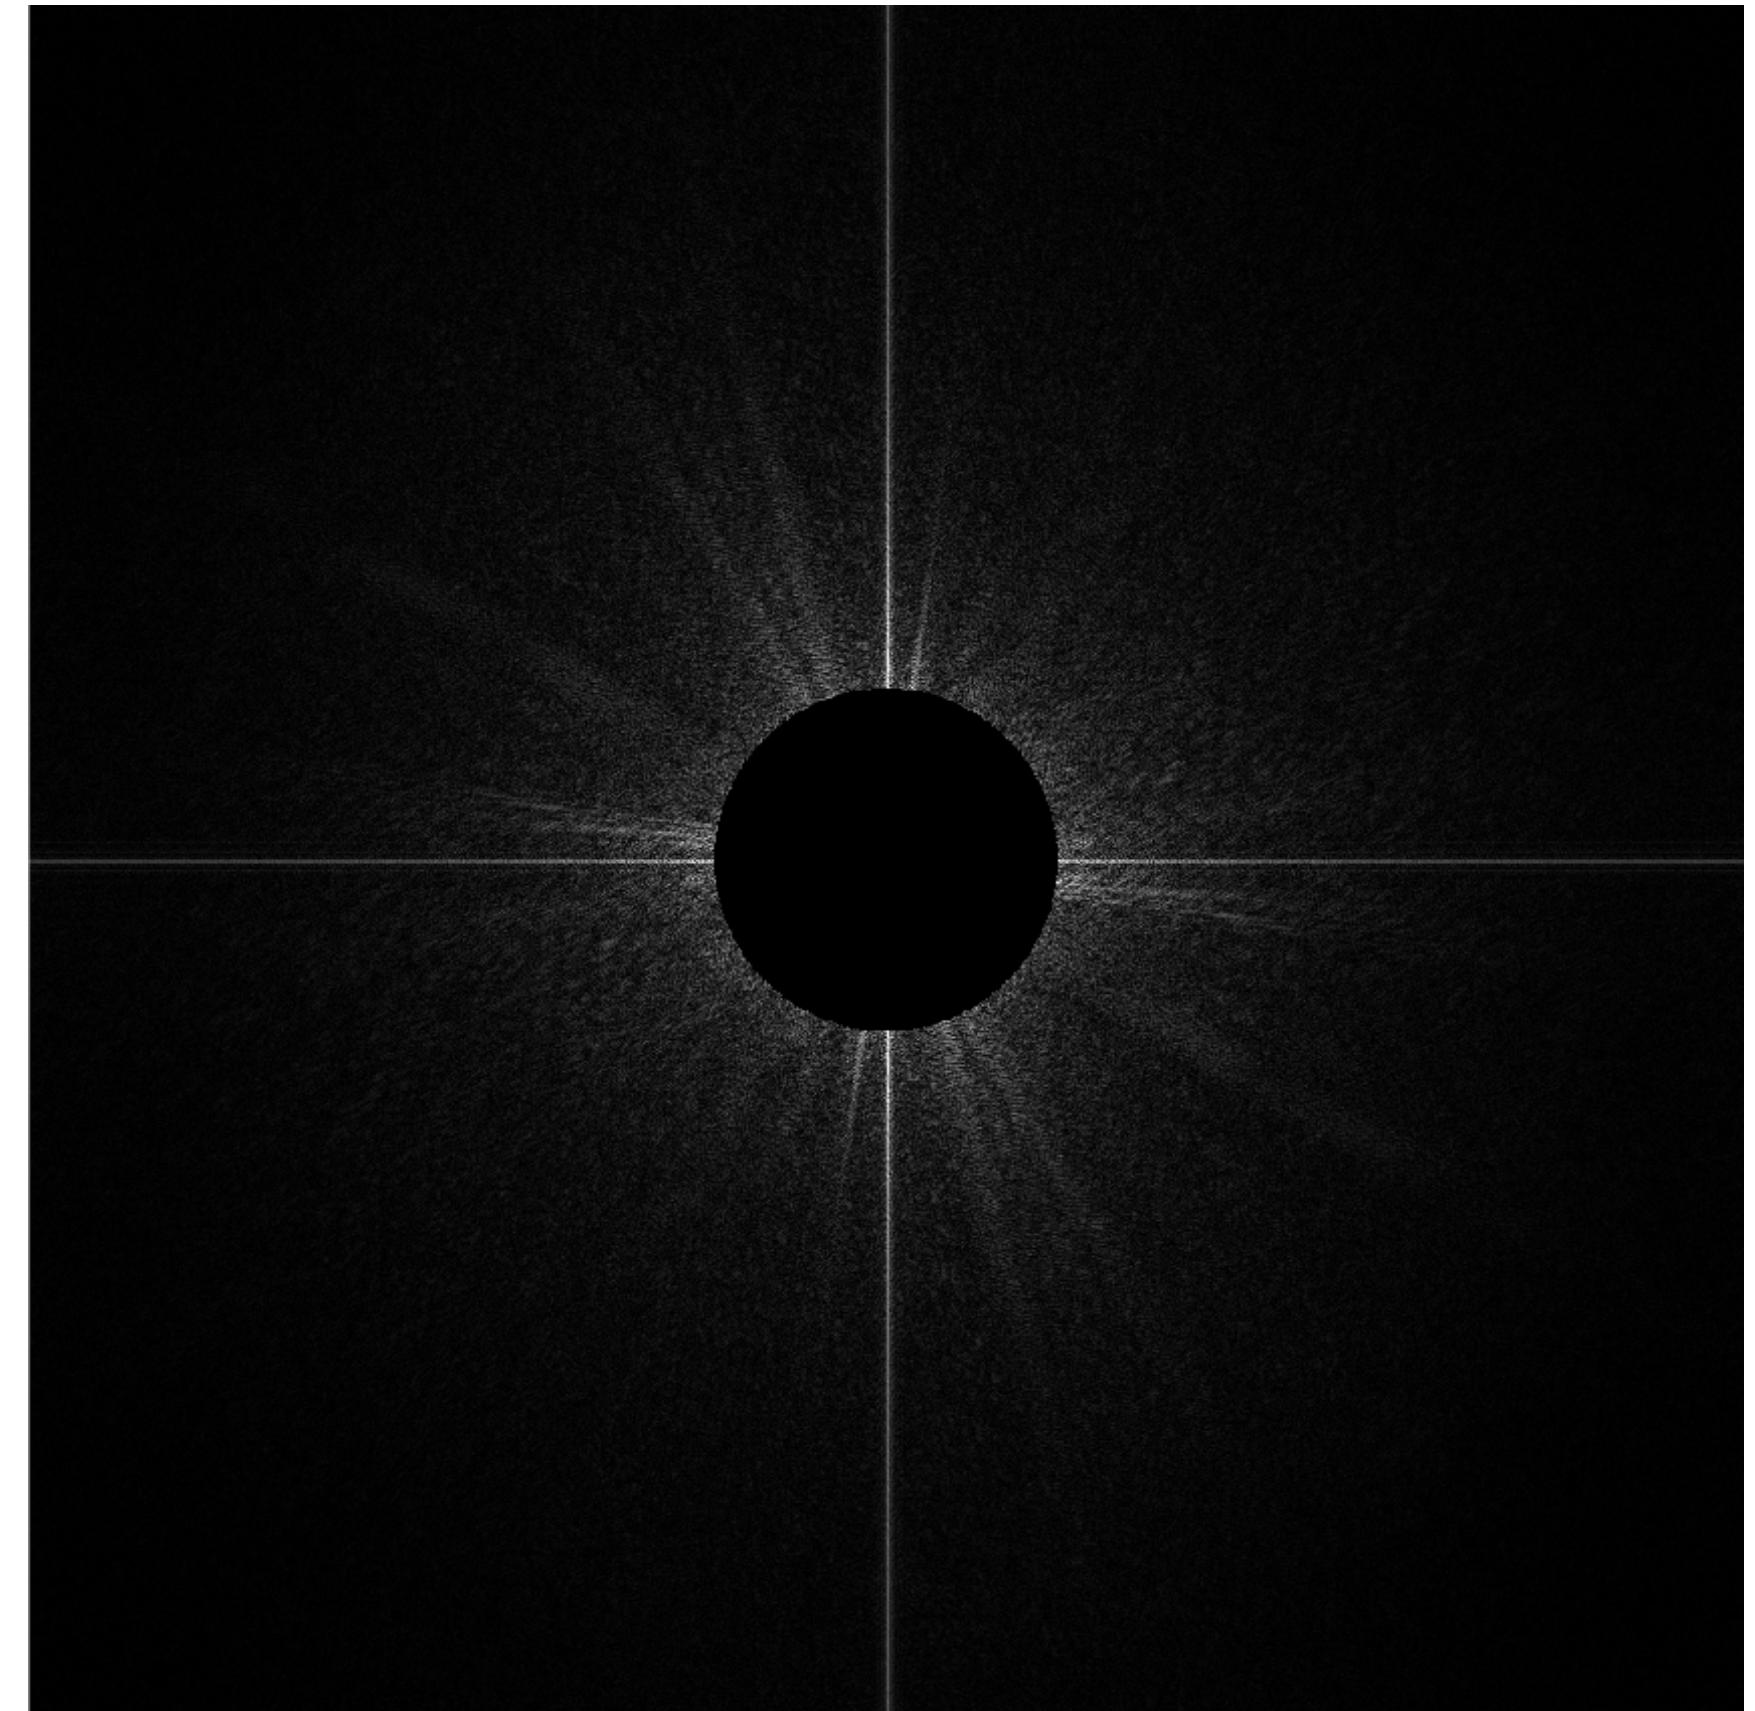


Spectrum (after low-pass filter)
All frequencies above cutoff have 0 magnitude

Recall: frequency content of images



**Spatial domain result
(strongest edges)**



**Spectrum (after high-pass filter)
All frequencies below threshold
have 0 magnitude**

A photograph of Santa Cruz



A photograph of Santa Cruz (with noise added)



A photograph of Santa Cruz (with more noise added)



JPEG compression: the big ideas

- **Low-frequency content is predominant in images of the real world**

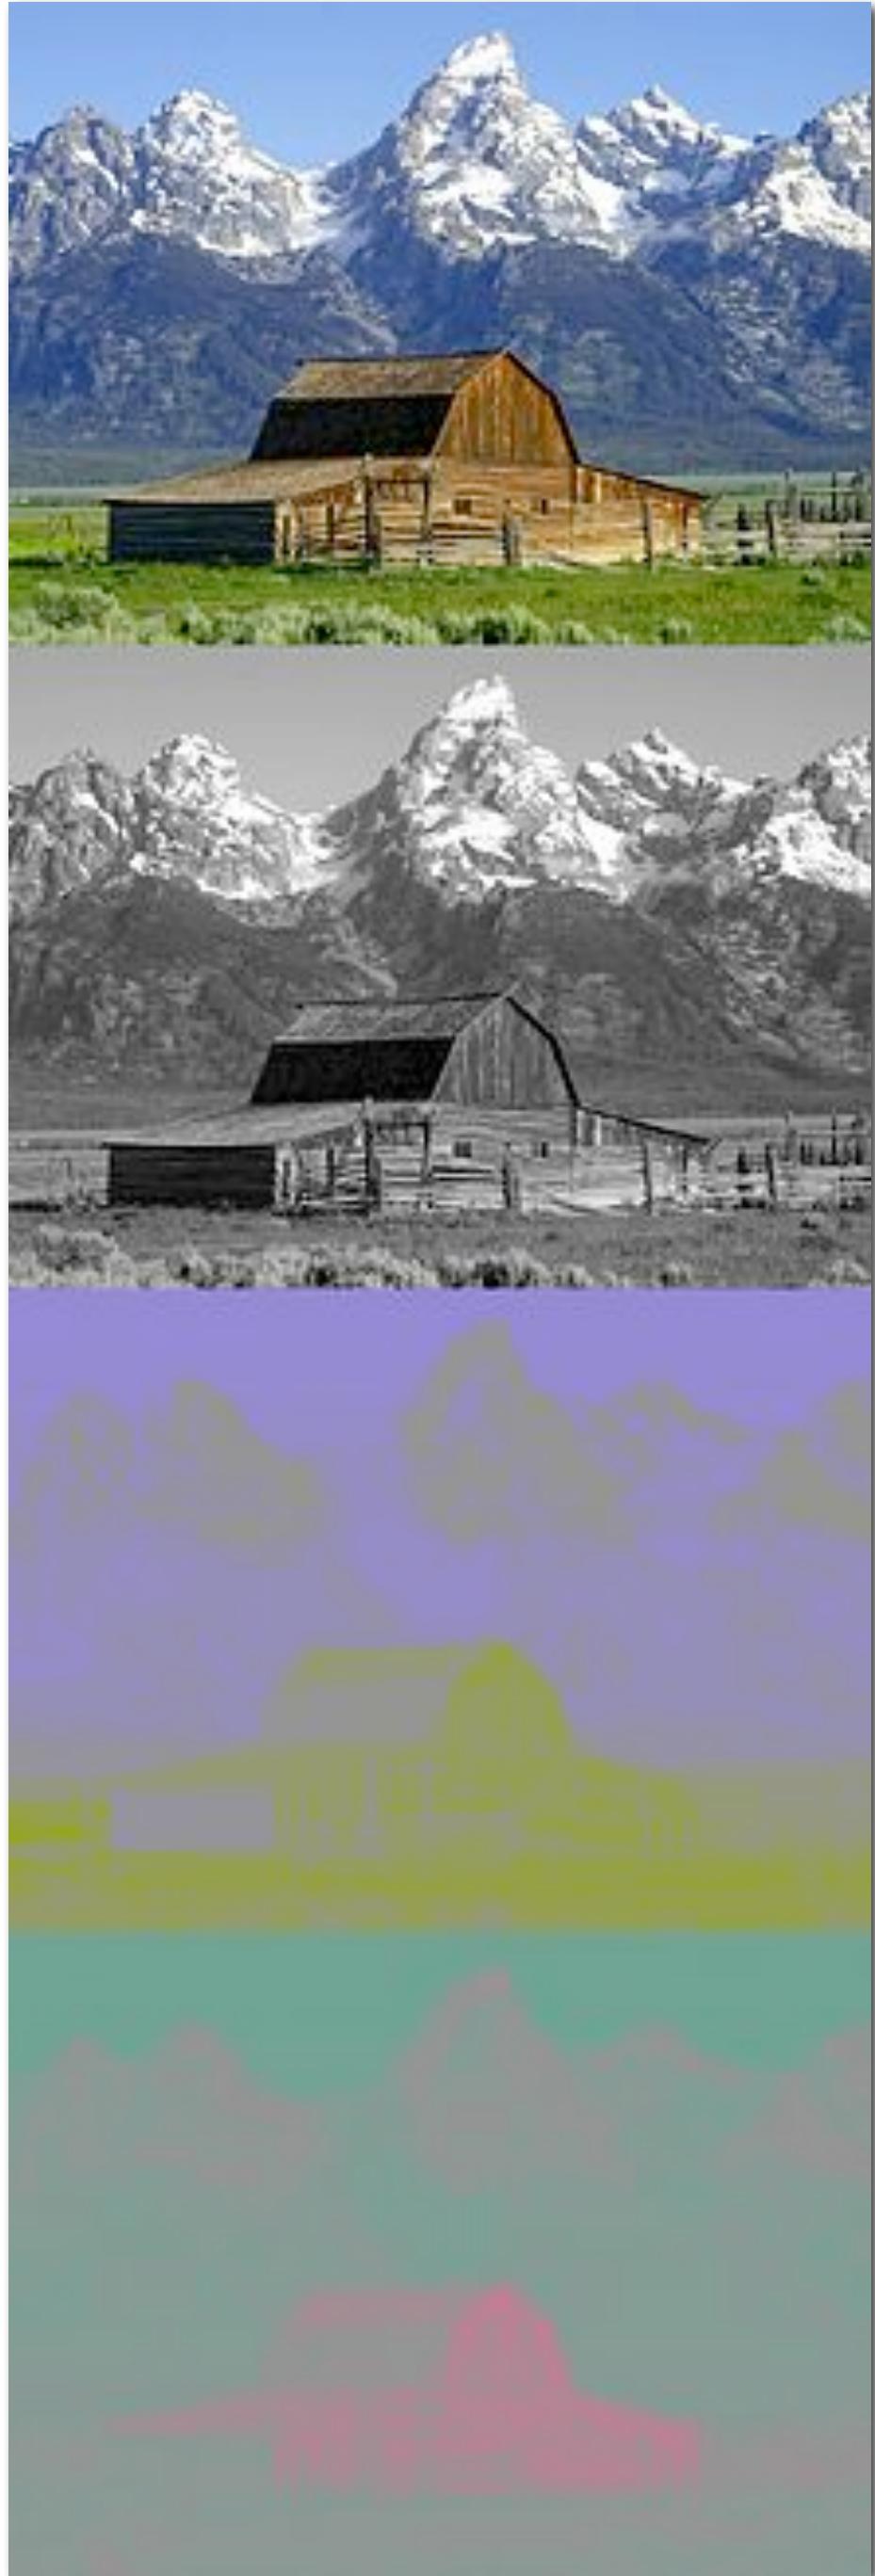


Therefore, it is often acceptable for a compression scheme to introduce errors in high-frequency components of the image.

- **The human visual system is:**

- **less sensitive to high frequency sources of error**
- **less sensitive to detail in chromaticity than in luminance**



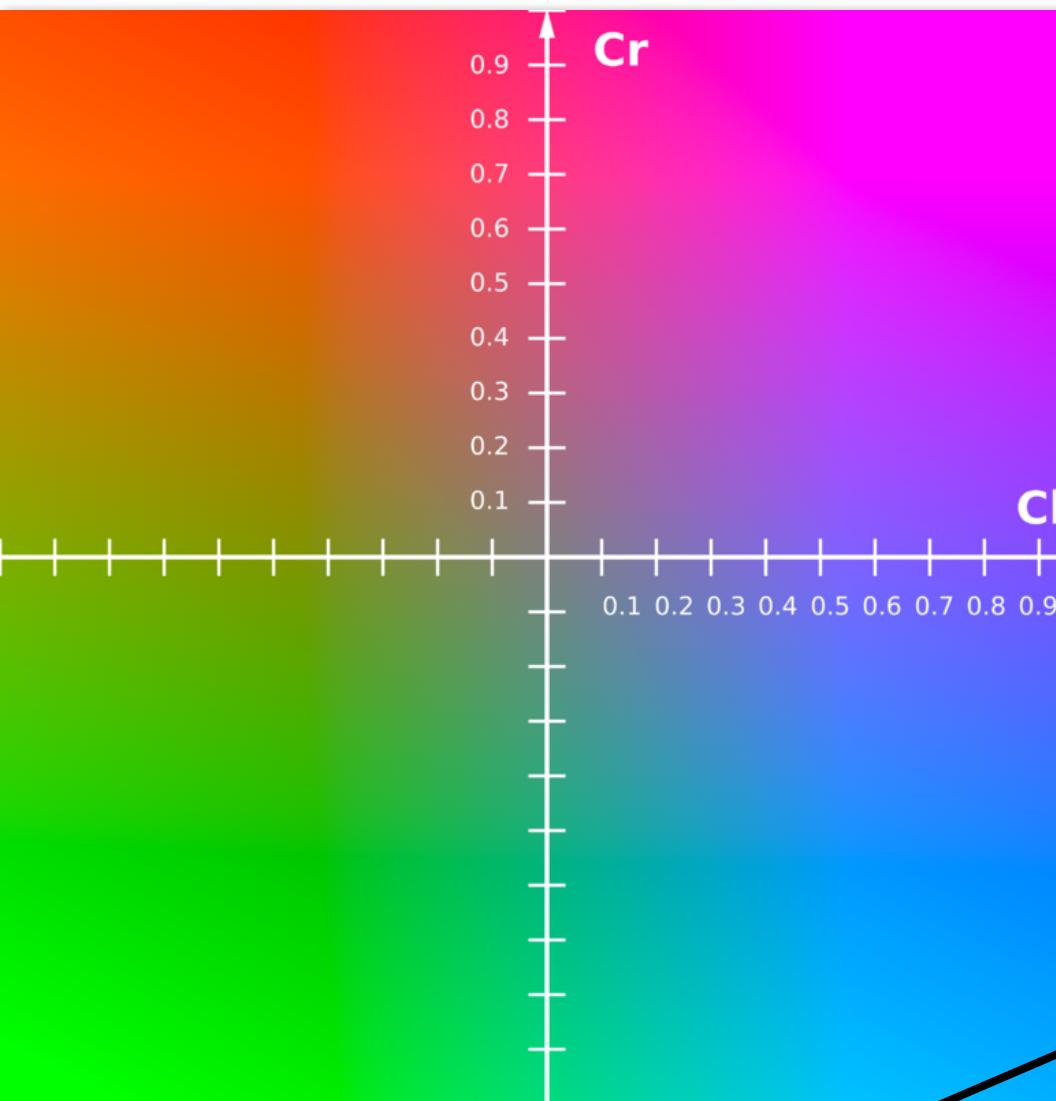


Y'CbCr color space

Y' = luma: perceived luminance (non-linear)

Cb = blue-yellow deviation from gray

Cr = red-cyan deviation from gray



Non-linear RGB
(primed notation indicates
perceptual (non-linear) space)

Conversion from R'G'B' to Y'CbCr:

$$Y' = 16 + \frac{65.738 \cdot R'_D}{256} + \frac{129.057 \cdot G'_D}{256} + \frac{25.064 \cdot B'_D}{256}$$

$$C_B = 128 + \frac{-37.945 \cdot R'_D}{256} - \frac{74.494 \cdot G'_D}{256} + \frac{112.439 \cdot B'_D}{256}$$

$$C_R = 128 + \frac{112.439 \cdot R'_D}{256} - \frac{94.154 \cdot G'_D}{256} - \frac{18.285 \cdot B'_D}{256}$$

Example: compression in Y'CbCr



Original picture of Kayvon

Example: compression in Y'CbCr



**Contents of CbCr color channels downsampled by a factor of 20 in each dimension
(400x reduction in number of samples)**

Example: compression in Y'CbCr



Full resolution sampling of luma (Y')

Example: compression in Y'CbCr



**Reconstructed result
(looks pretty good)**

Chroma subsampling

Y'CbCr is an efficient representation for storage (and transmission) because Y' can be stored at higher resolution than CbCr without significant loss in perceived visual quality

Y'_{00} Cb_{00} Cr_{00}	Y'_{10}	Y'_{20} Cb_{20} Cr_{20}	Y'_{30}
Y'_{01} Cb_{01} Cr_{01}	Y'_{11}	Y'_{21} Cb_{21} Cr_{21}	Y'_{31}

Y'_{00} Cb_{00} Cr_{00}	Y'_{10}	Y'_{20} Cb_{20} Cr_{20}	Y'_{30}
Y'_{01}	Y'_{11}	Y'_{21}	Y'_{31}

4:2:2 representation:

Store Y' at full resolution

**Store Cb, Cr at full vertical resolution,
but only half horizontal resolution**

4:2:0 representation:

Store Y' at full resolution

**Store Cb, Cr at half resolution in both
dimensions**

X:Y:Z notation:

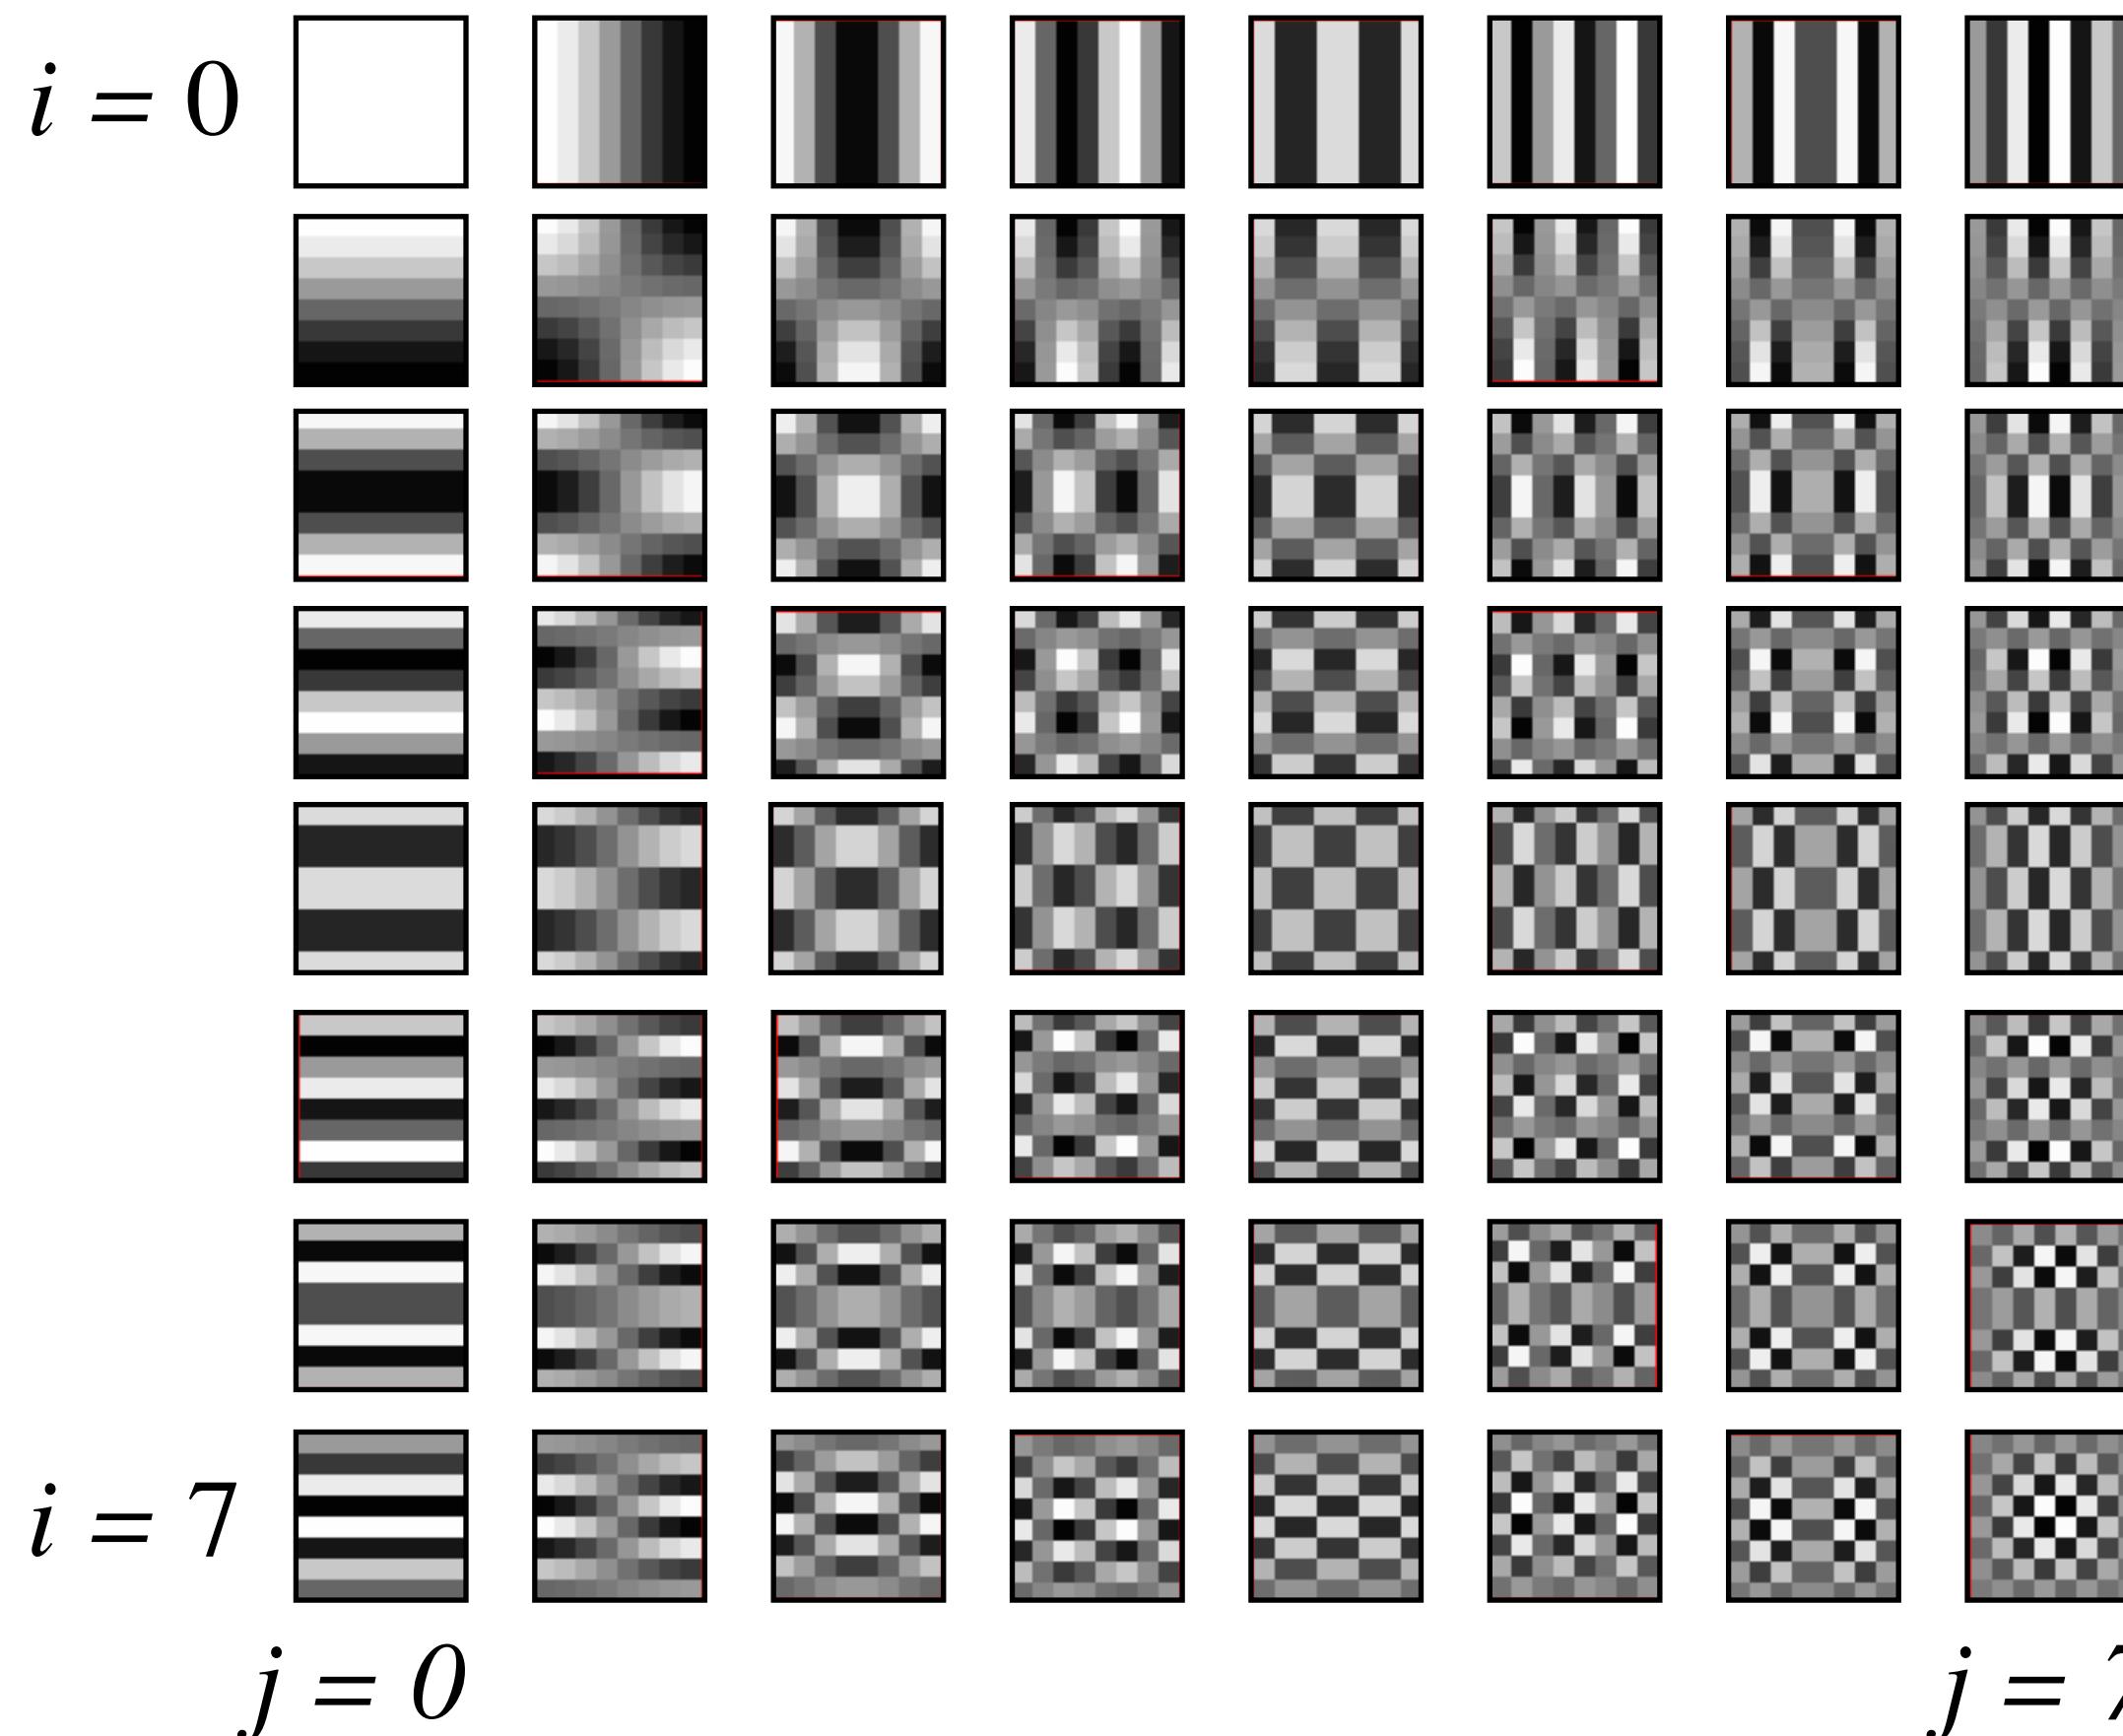
X = width of block

Y = number of chroma samples in first row

Z = number of chroma samples in second row

Apply discrete cosine transform (DCT) to each 8x8 block of image values

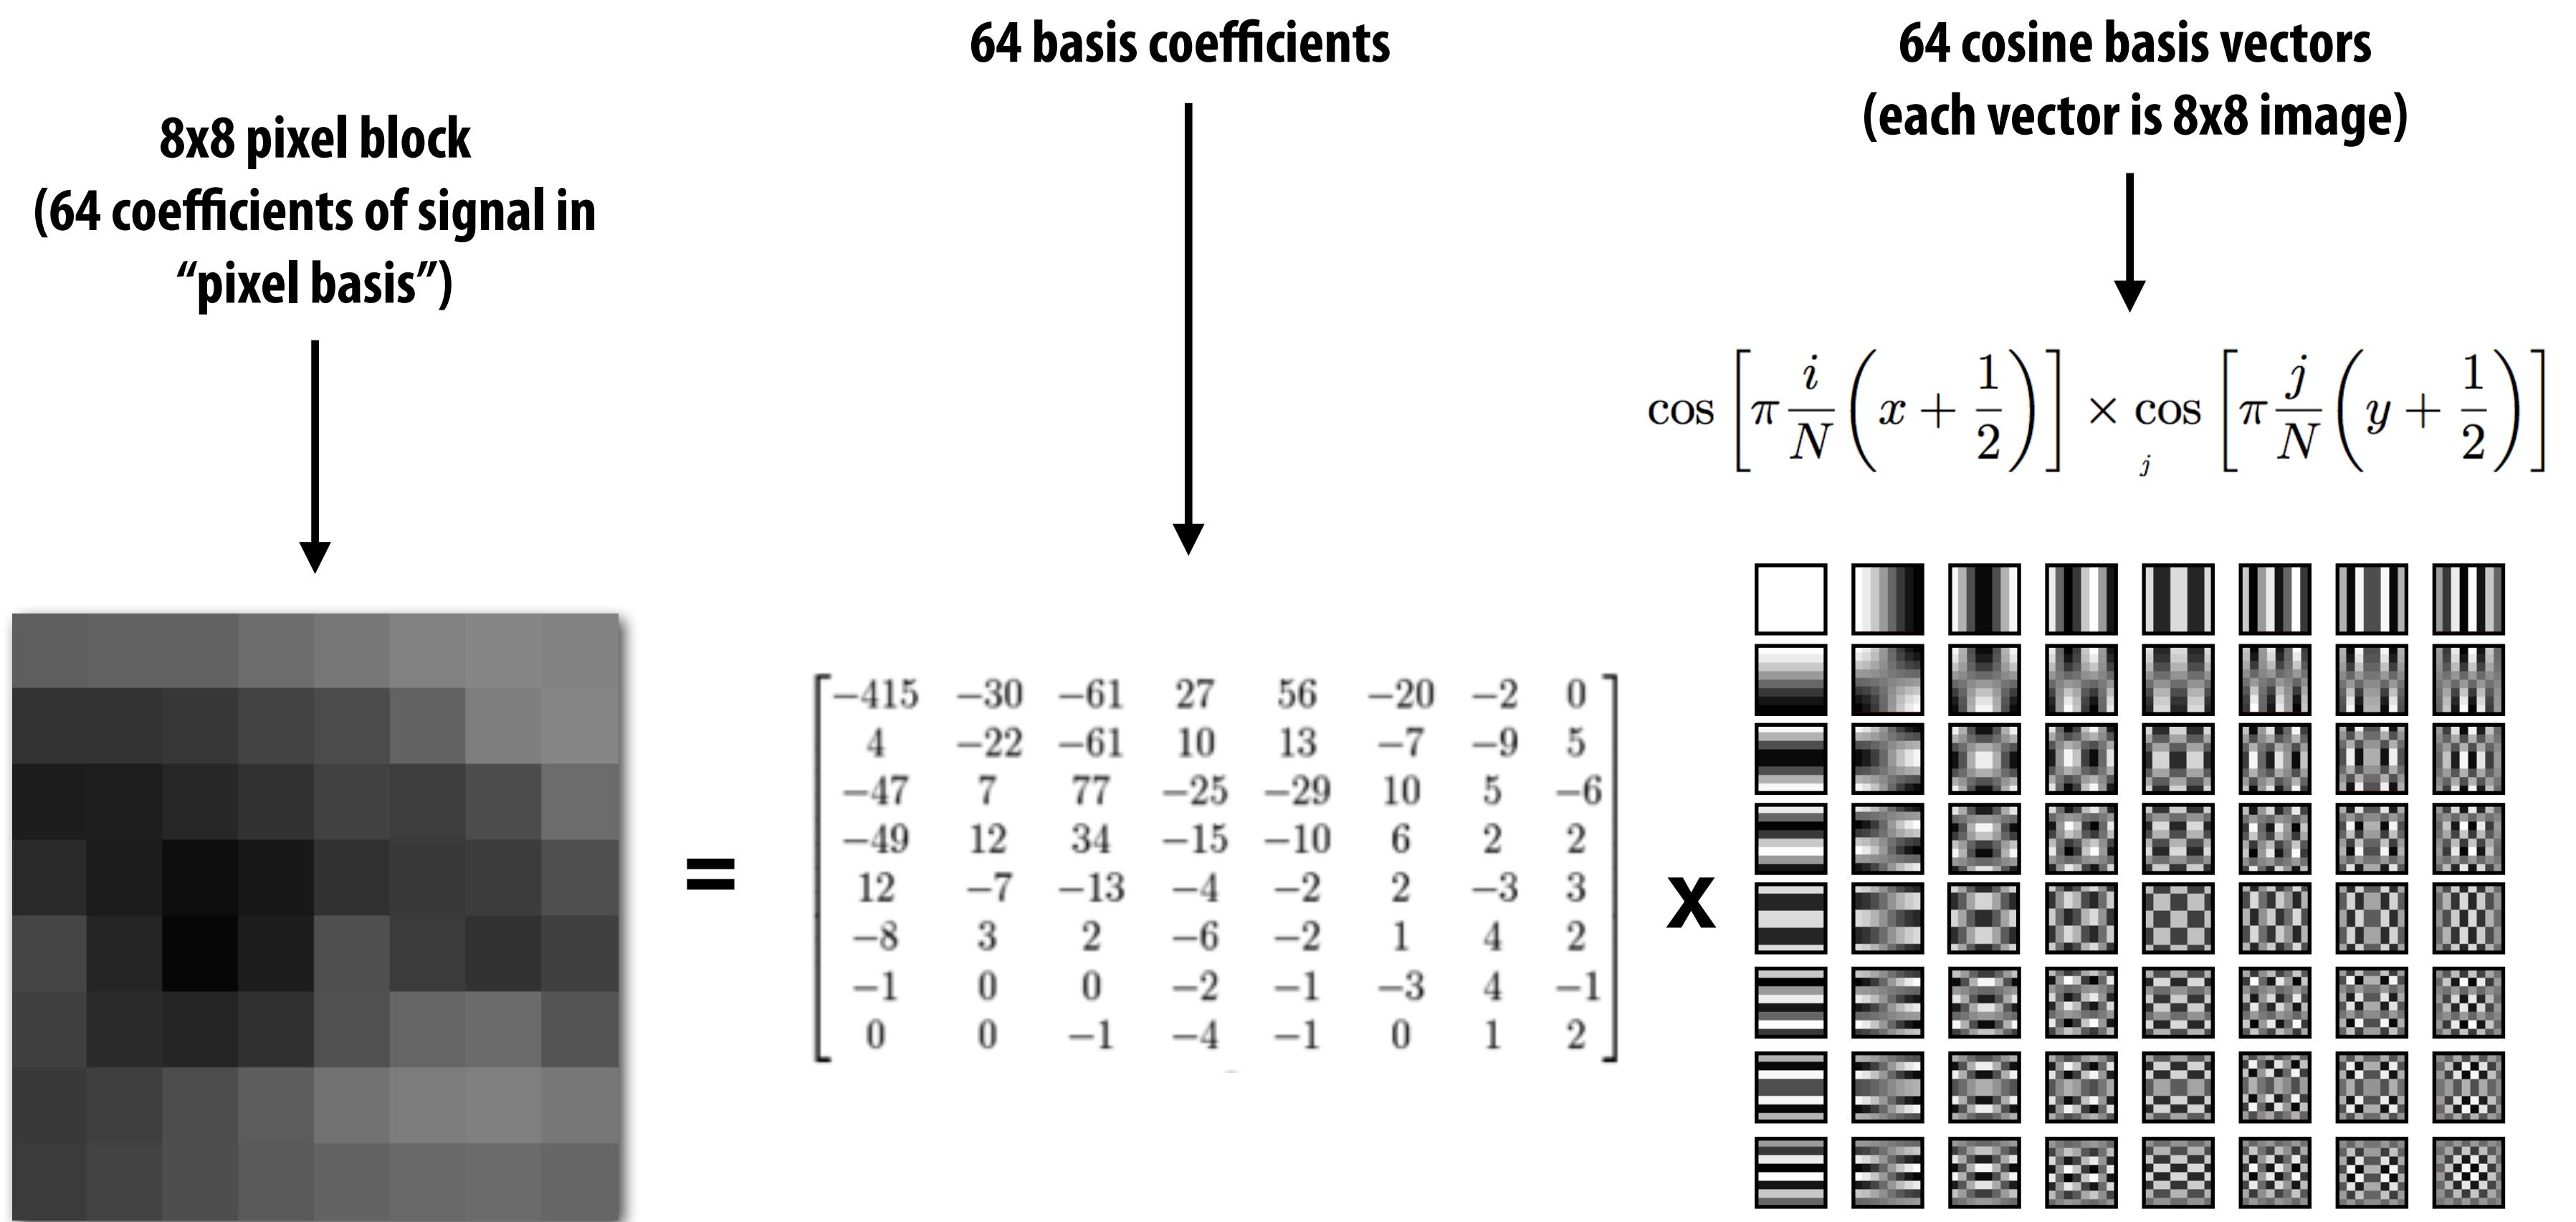
$$\text{basis}[i, j] = \cos \left[\pi \frac{i}{N} \left(x + \frac{1}{2} \right) \right] \times \cos \left[\pi \frac{j}{N} \left(y + \frac{1}{2} \right) \right]$$



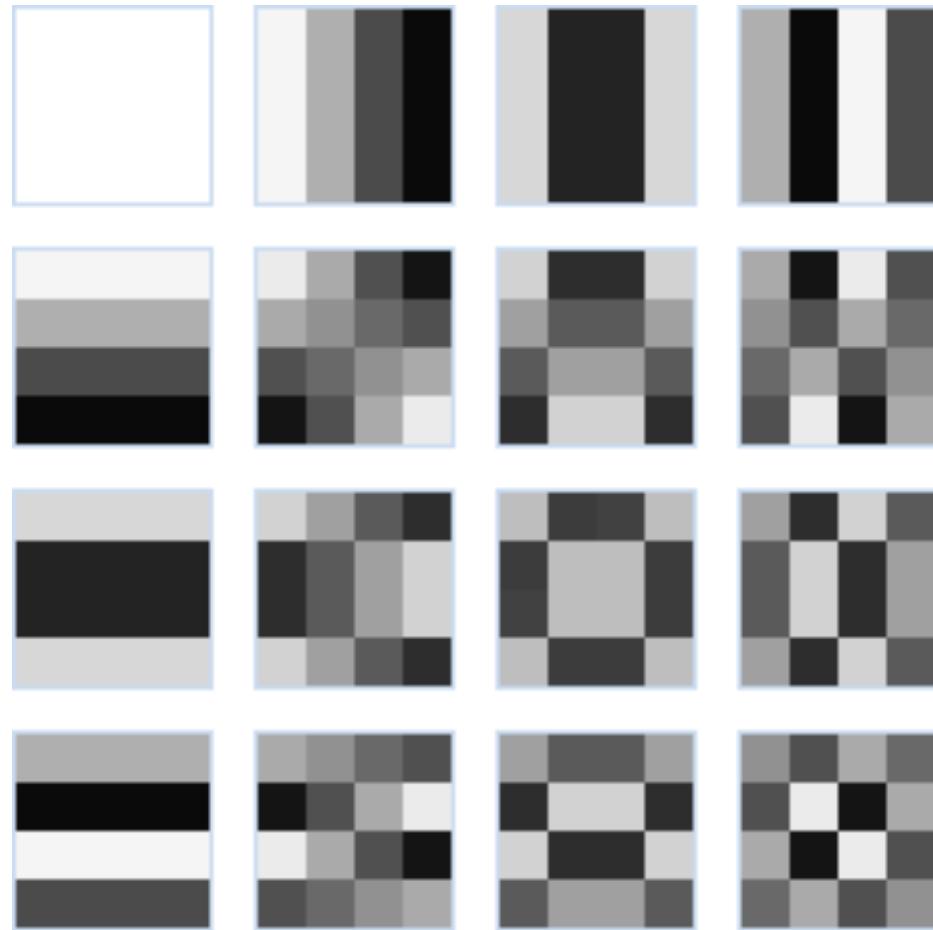
DCT computes projection of image onto 64 basis functions: $\text{basis}[i, j]$

DCT applied to 8x8 pixel blocks of Y' channel, 16x16 pixel blocks of Cb, Cr (assuming 4:2:0)

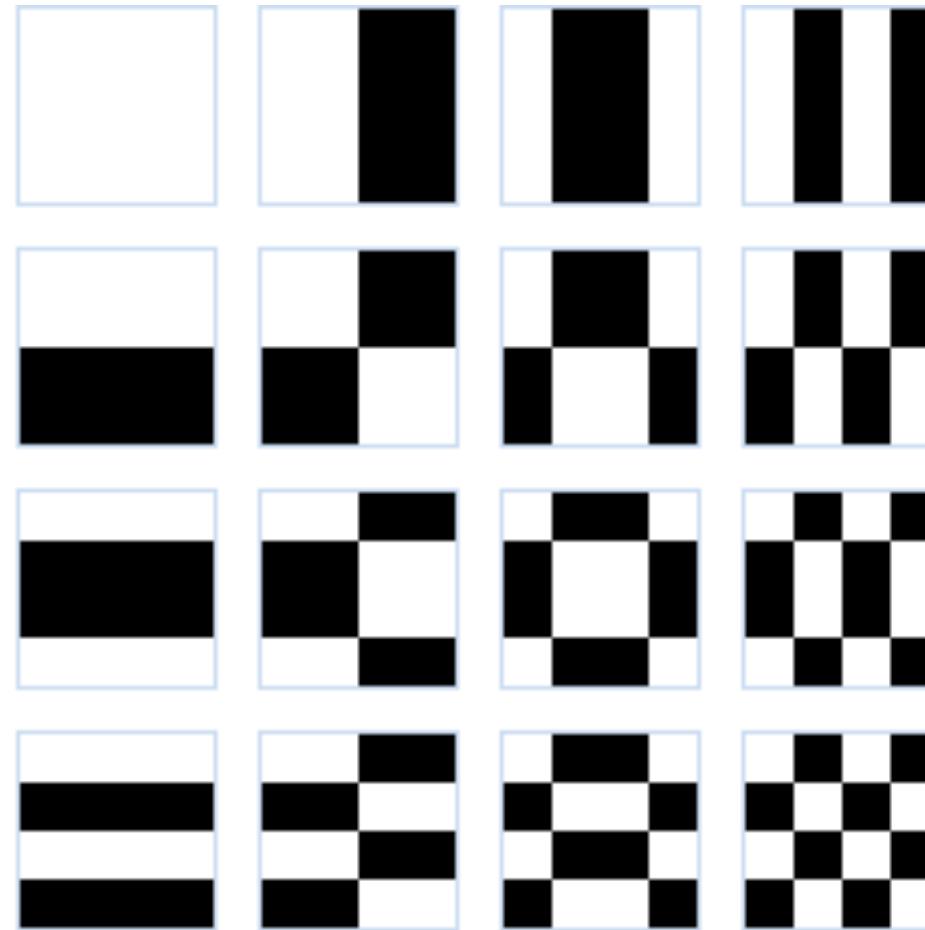
Image transform coding via DCT



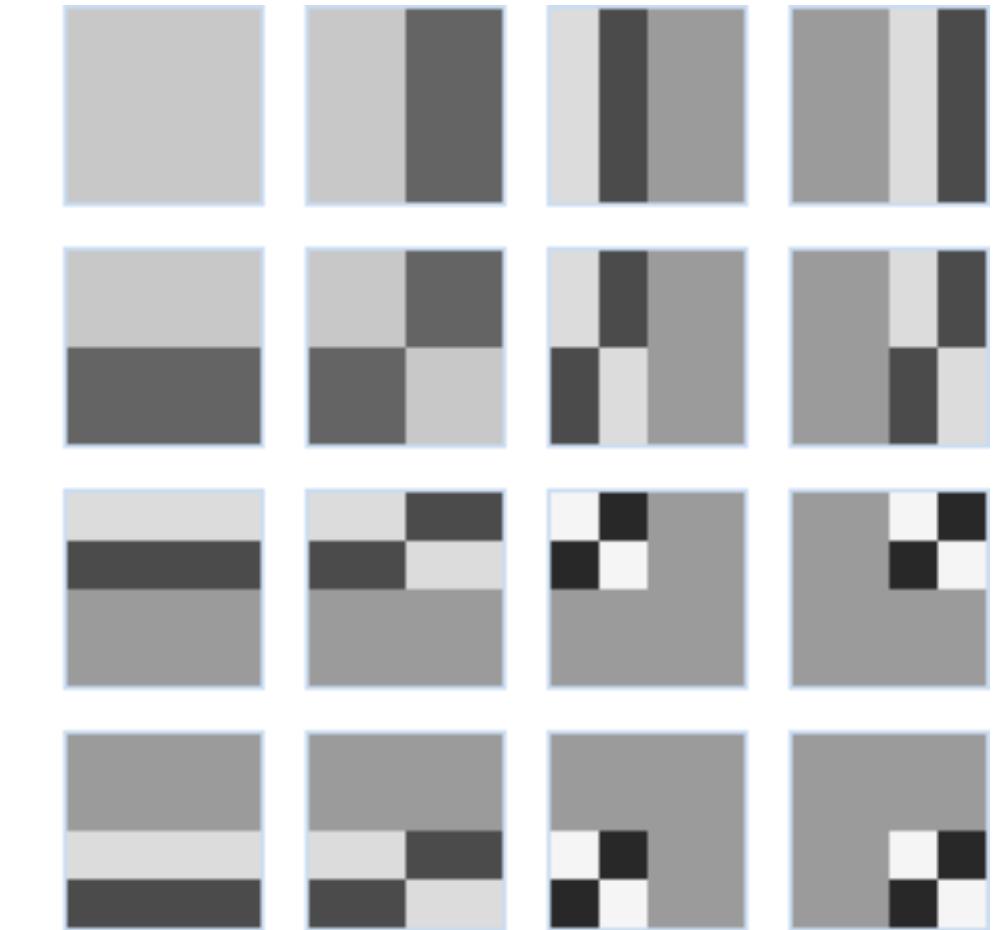
Other bases



DCT



Walsh-Hadamard



Haar Wavelet

This slide illustrates basis images for 4x4 image block

Quantization

$$\begin{bmatrix} -415 & -30 & -61 & 27 & 56 & -20 & -2 & 0 \\ 4 & -22 & -61 & 10 & 13 & -7 & -9 & 5 \\ -47 & 7 & 77 & -25 & -29 & 10 & 5 & -6 \\ -49 & 12 & 34 & -15 & -10 & 6 & 2 & 2 \\ 12 & -7 & -13 & -4 & -2 & 2 & -3 & 3 \\ -8 & 3 & 2 & -6 & -2 & 1 & 4 & 2 \\ -1 & 0 & 0 & -2 & -1 & -3 & 4 & -1 \\ 0 & 0 & -1 & -4 & -1 & 0 & 1 & 2 \end{bmatrix}$$

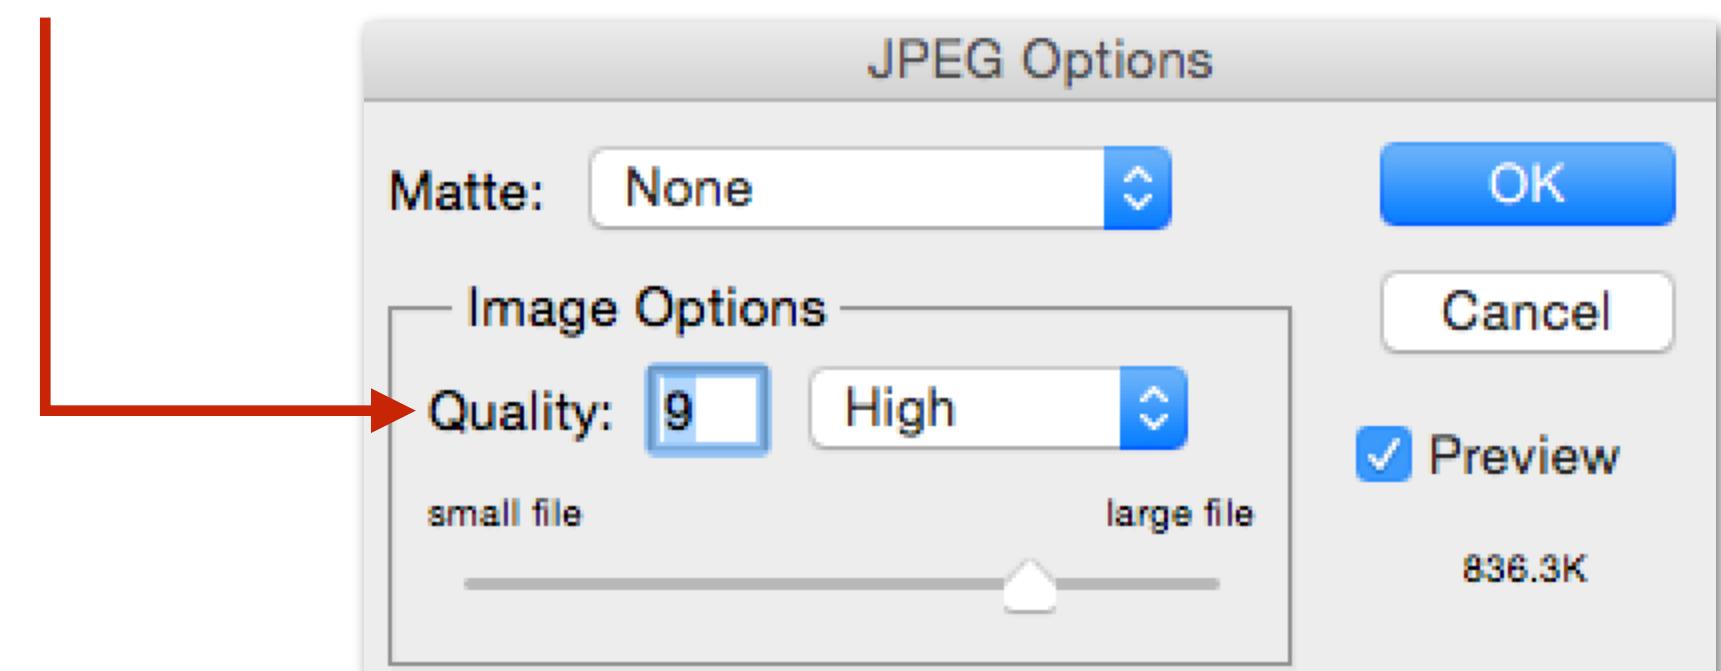
Result of DCT
(representation of image in cosine basis)

$$\begin{bmatrix} 16 & 11 & 10 & 16 & 24 & 40 & 51 & 61 \\ 12 & 12 & 14 & 19 & 26 & 58 & 60 & 55 \\ 14 & 13 & 16 & 24 & 40 & 57 & 69 & 56 \\ 14 & 17 & 22 & 29 & 51 & 87 & 80 & 62 \\ 18 & 22 & 37 & 56 & 68 & 109 & 103 & 77 \\ 24 & 35 & 55 & 64 & 81 & 104 & 113 & 92 \\ 49 & 64 & 78 & 87 & 103 & 121 & 120 & 101 \\ 72 & 92 & 95 & 98 & 112 & 100 & 103 & 99 \end{bmatrix}$$

Quantization Matrix

$$= \begin{bmatrix} -26 & -3 & -6 & 2 & 2 & -1 & 0 & 0 \\ 0 & -2 & -4 & 1 & 1 & 0 & 0 & 0 \\ -3 & 1 & 5 & -1 & -1 & 0 & 0 & 0 \\ -4 & 1 & 2 & -1 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

Changing JPEG quality setting in your favorite photo app modifies this matrix ("lower quality" = higher values for elements in quantization matrix)



Quantization produces small values for coefficients (only few bits needed per coefficient)
Quantization zeros out many coefficients

Slide credit: Wikipedia, Pat Hanrahan

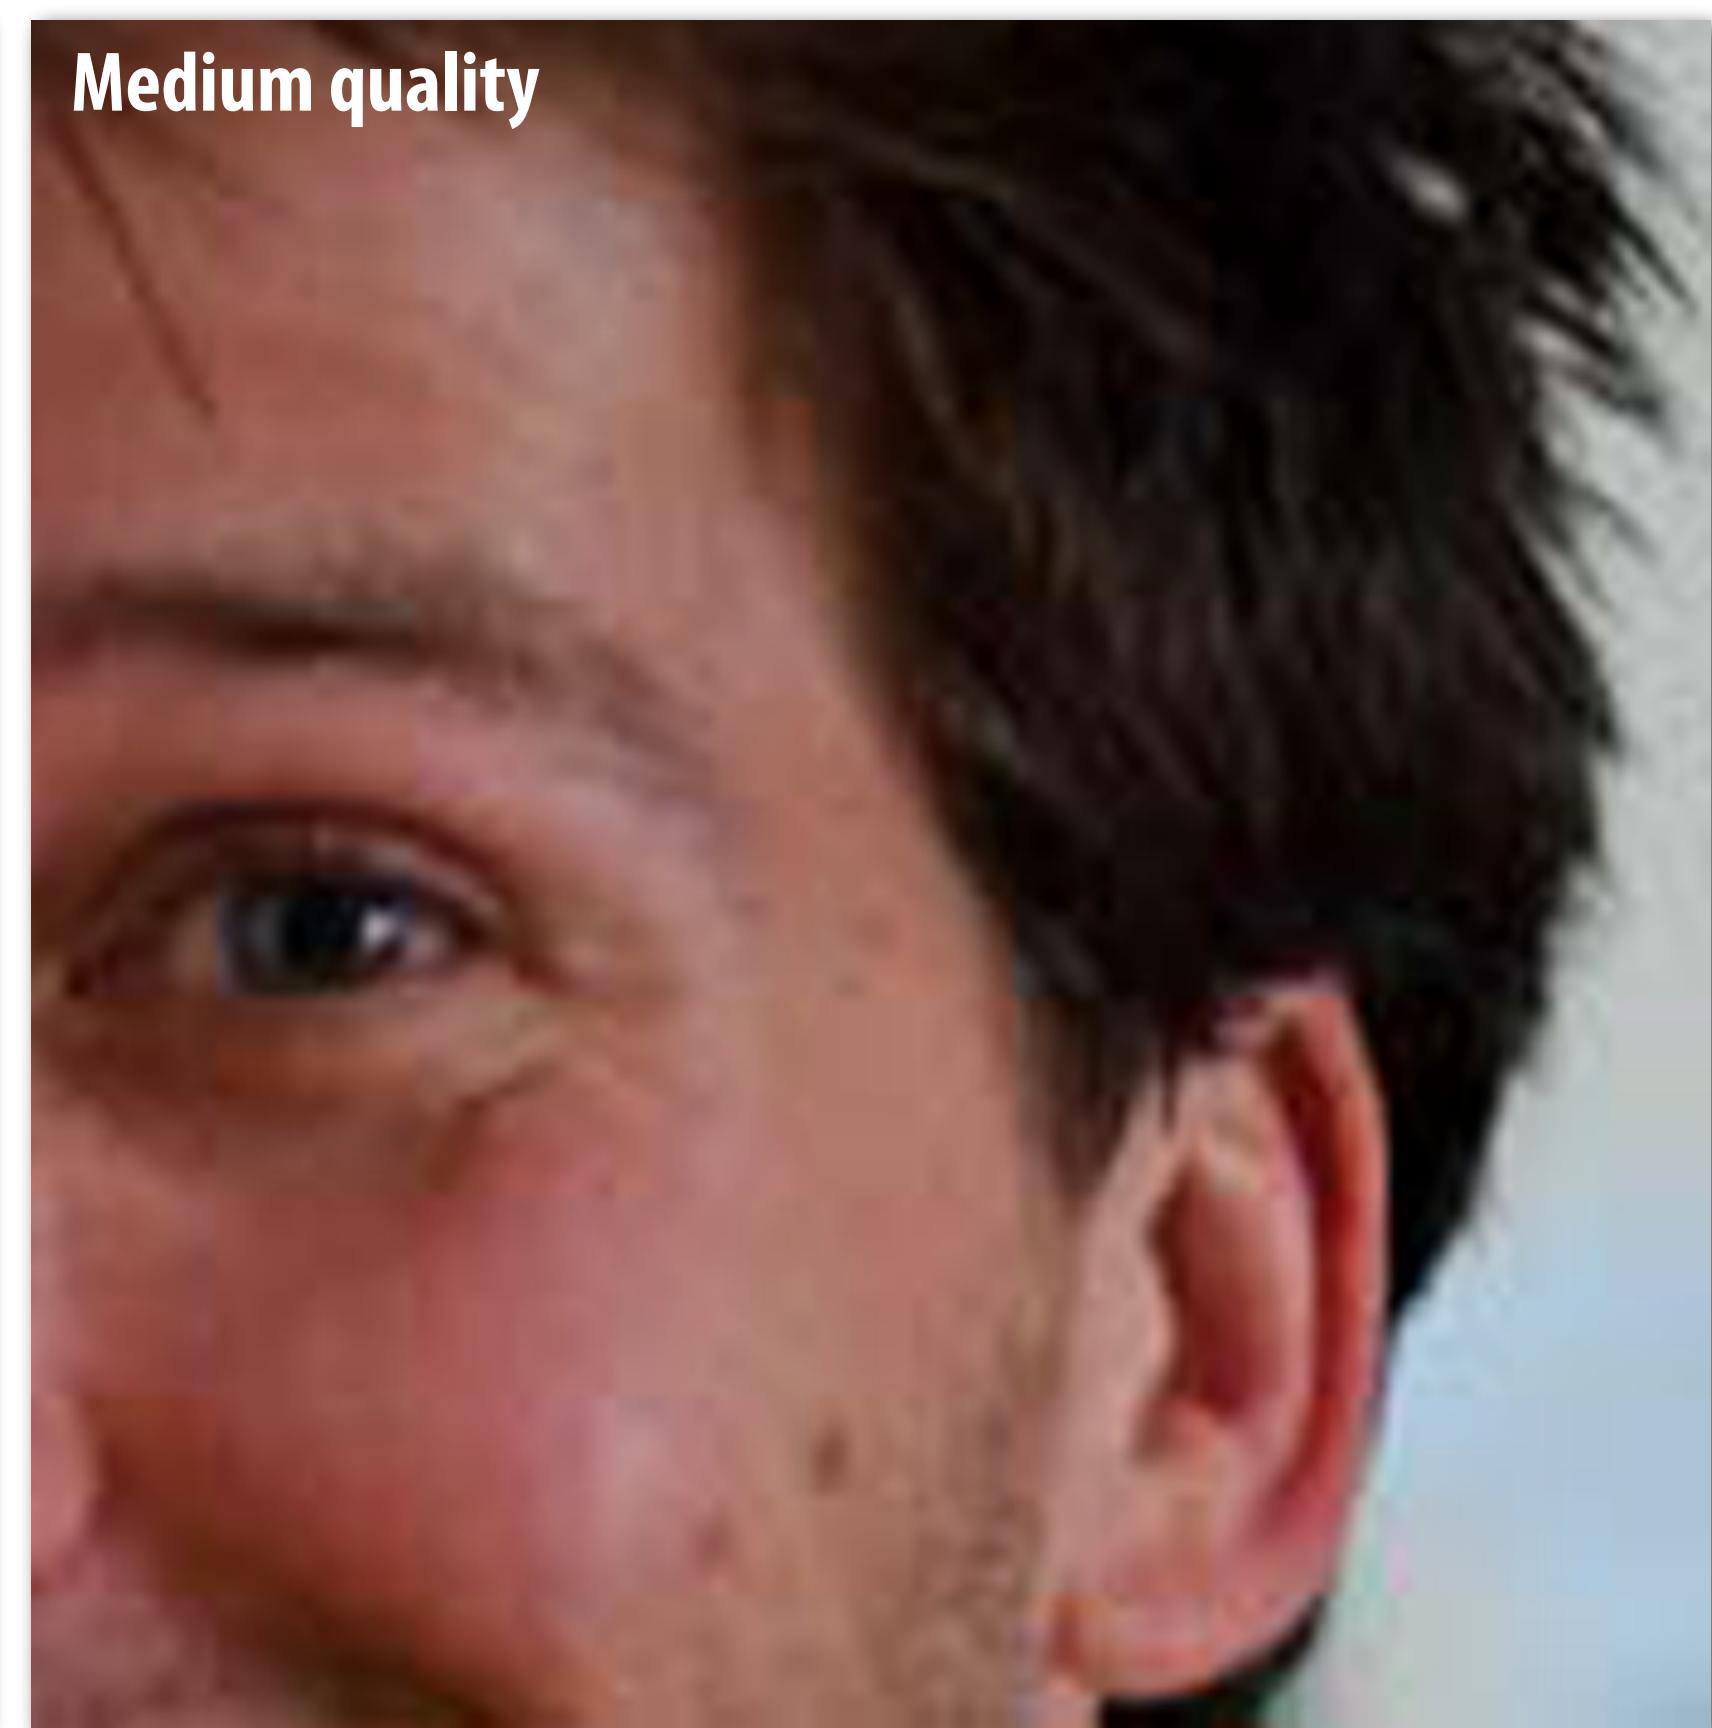
Stanford CS348V, Winter 2018

JPEG compression artifacts

Noticeable 8x8 pixel block boundaries

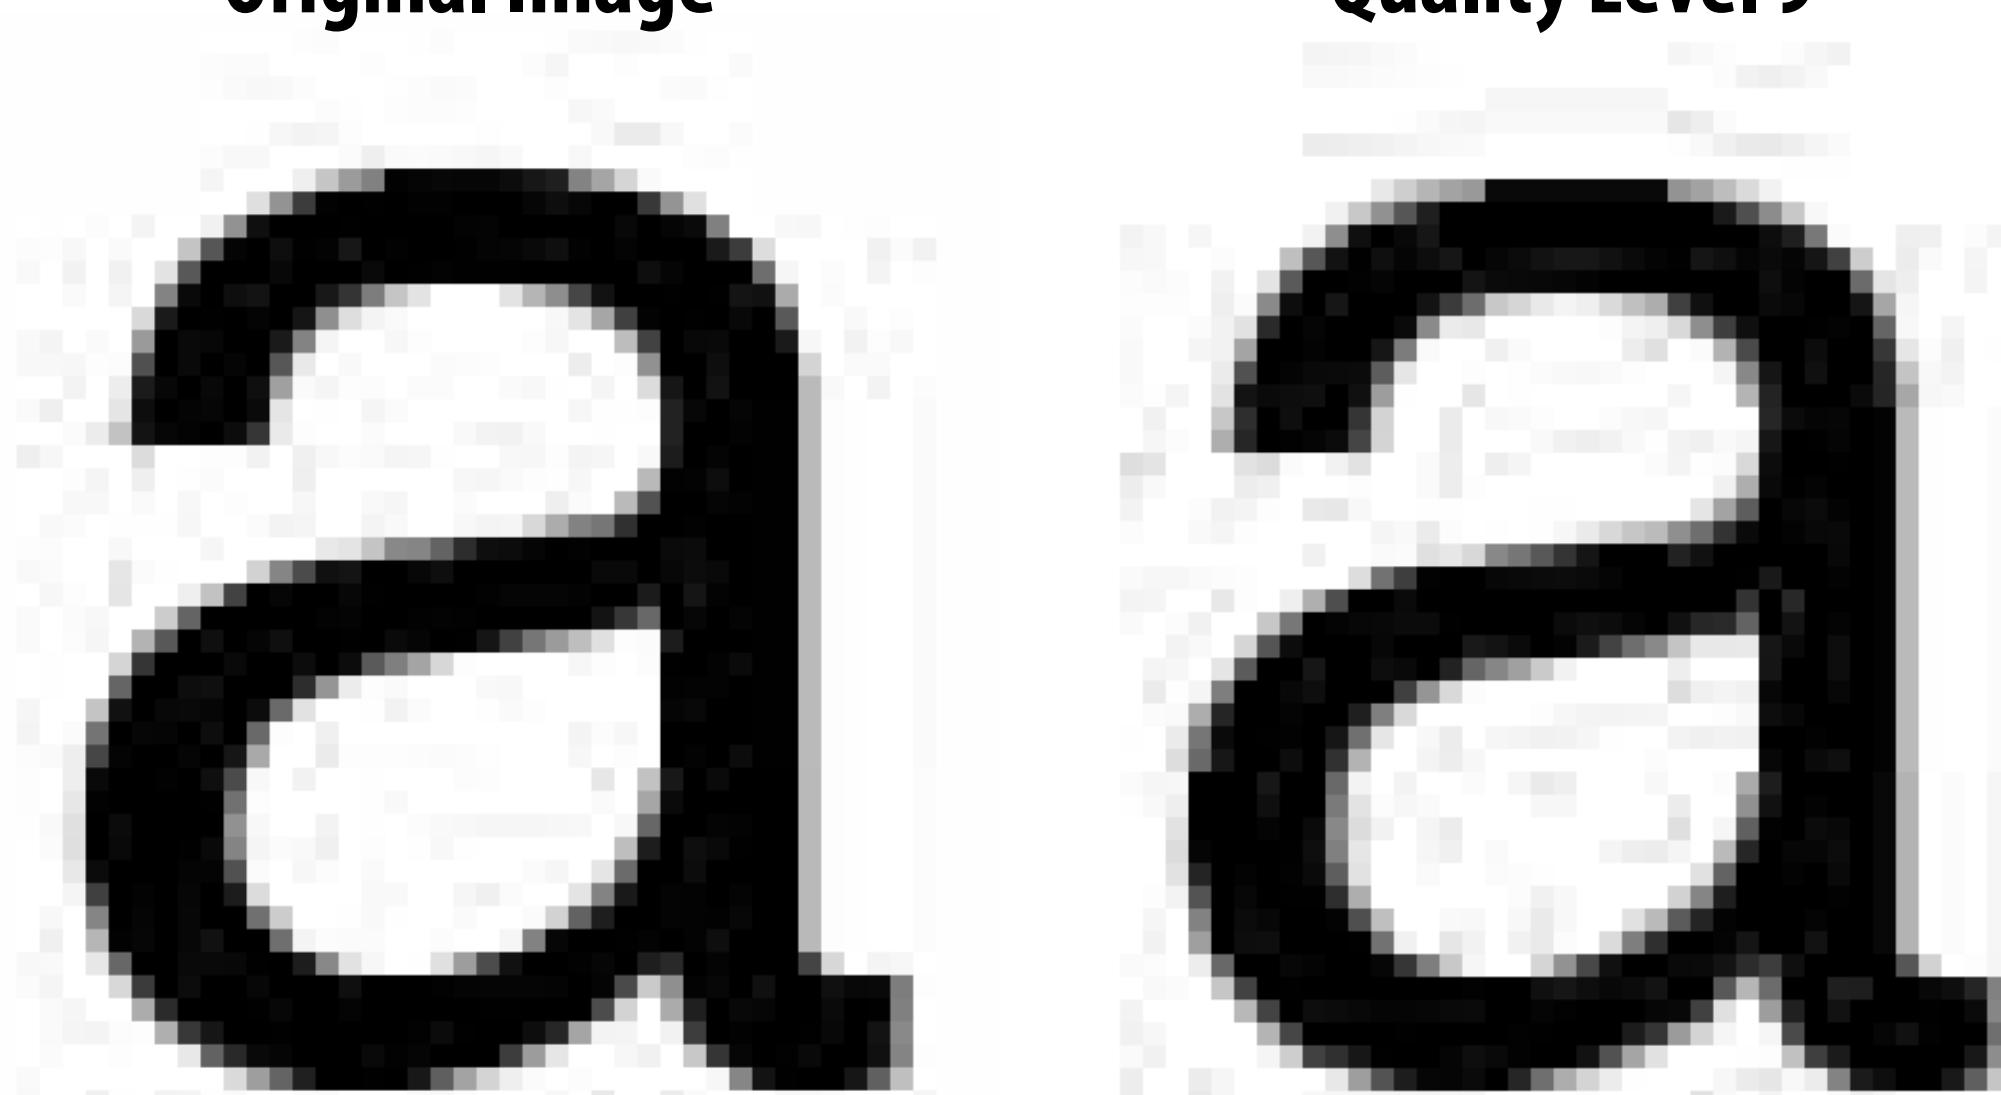


Noticeable error near large color gradients



Low-frequency regions of image represented accurately even under high compression

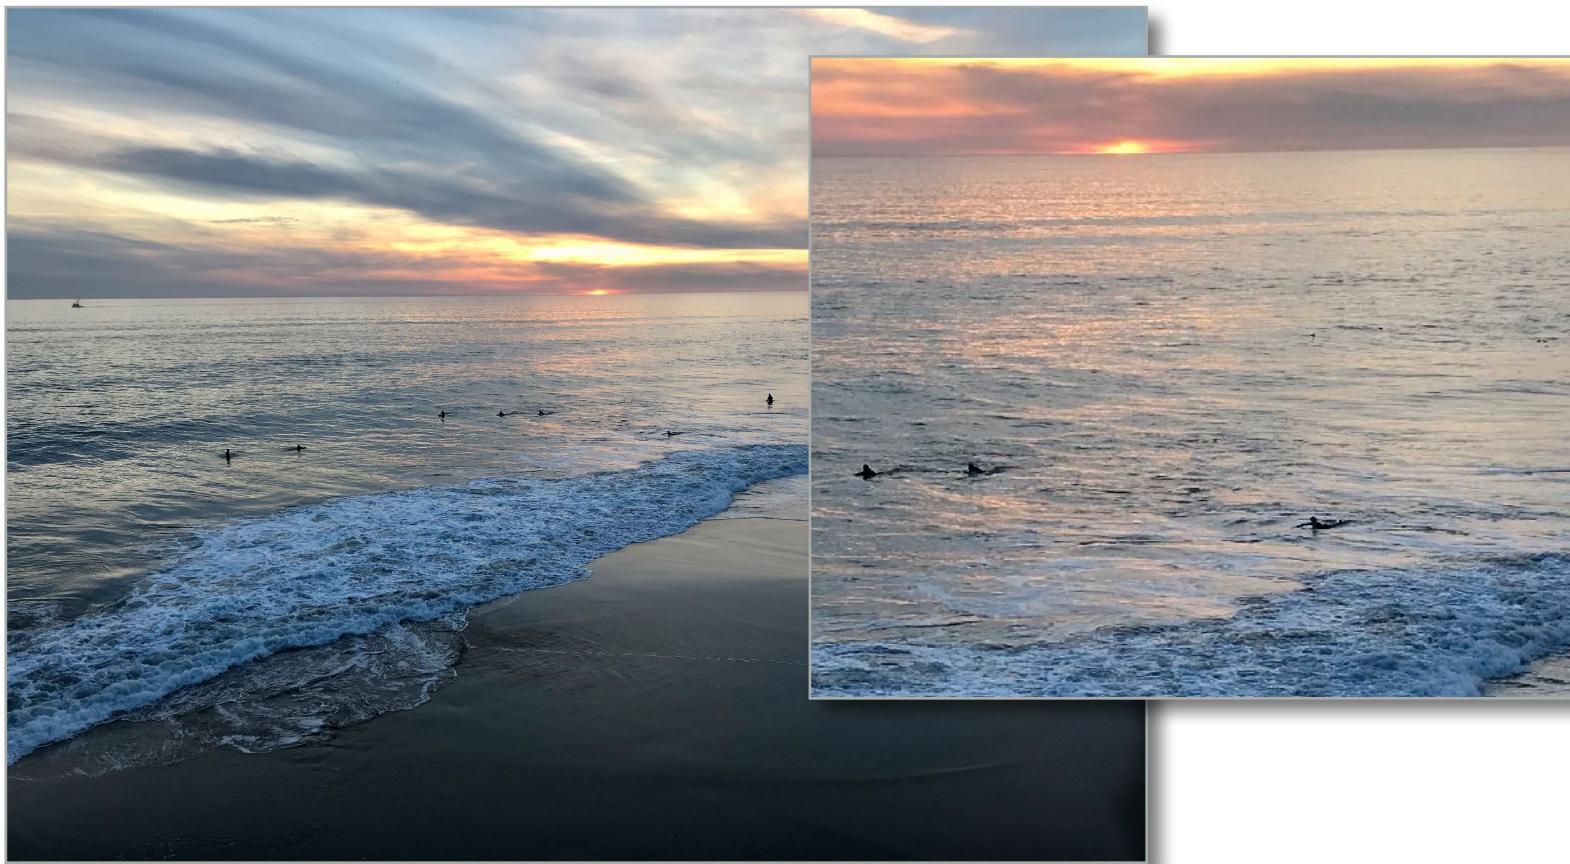
JPEG compression artifacts



Quality Level 3

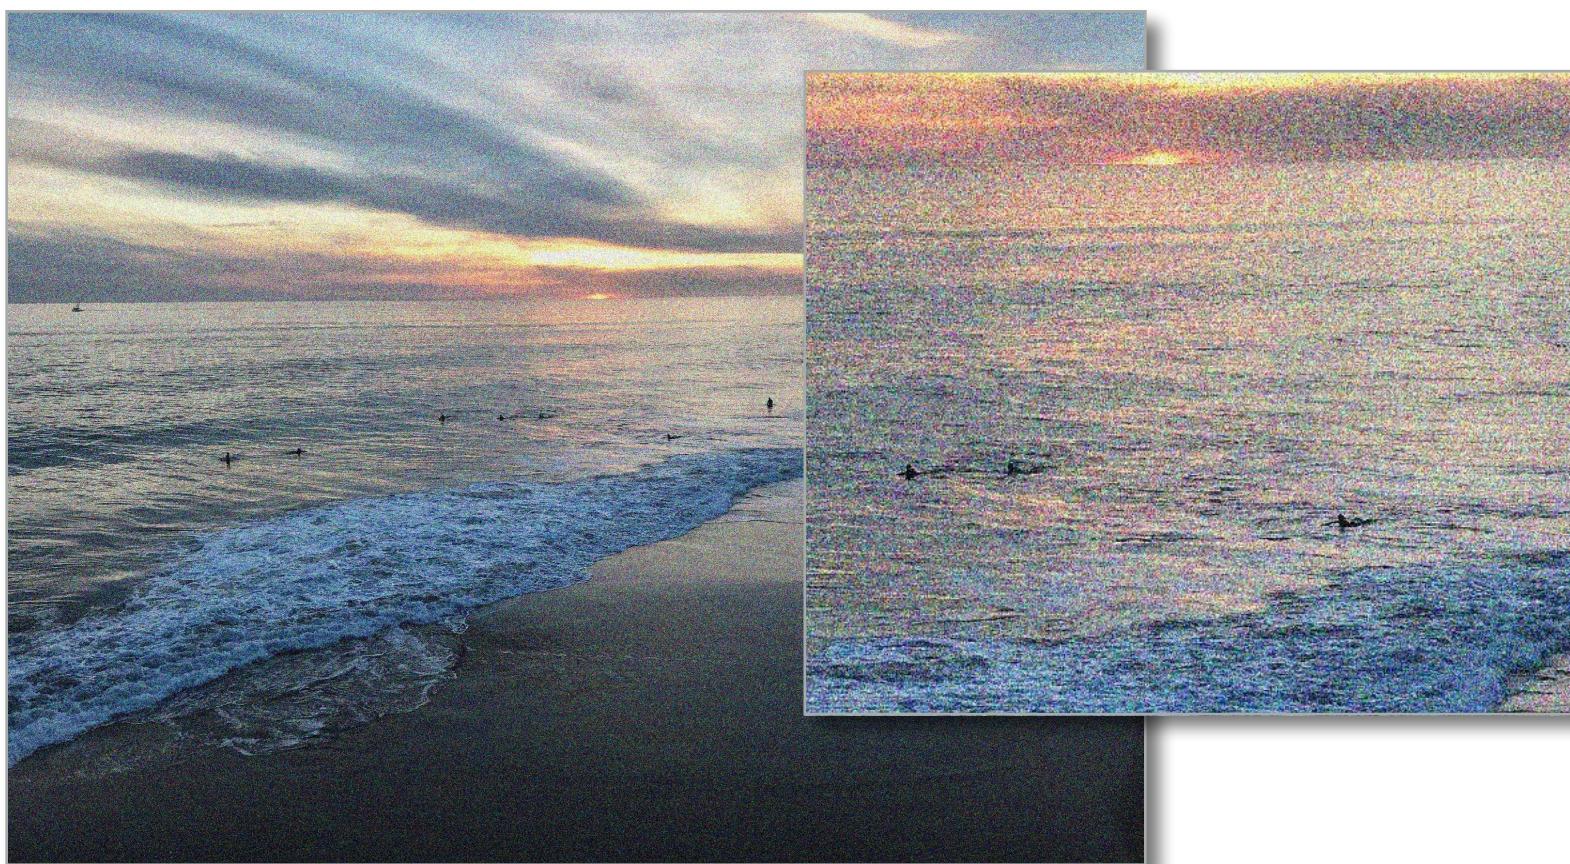
Quality Level 1

Why might JPEG compression not be a good compression scheme for illustrations and rasterized text?

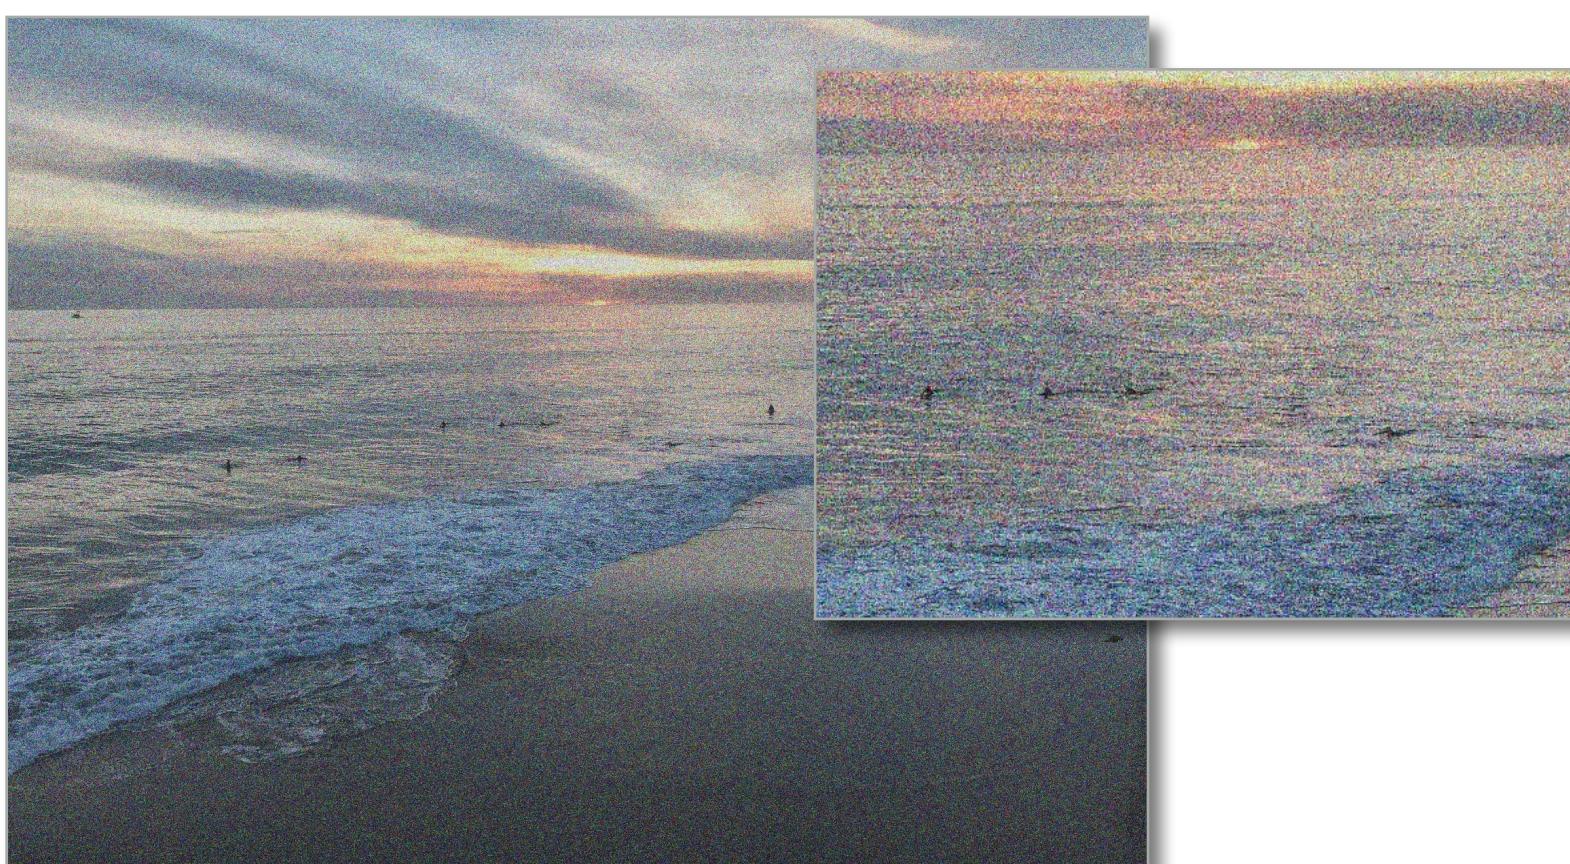


Images with high frequency content do not exhibit high compression ratios. Why?

Original image: 2.6MB JPG



Medium noise: 20.5MB JPG



High noise: 26.5MB JPG

**Photoshop JPG Compression level = 10
for all compressed images**

**Uncompressed image:
 $4032 \times 3024 \times 24 \text{ bpp} = 36.6 \text{ MB}$**

Lossless compression of quantized DCT values

-26	-3	-6	2	2	-1	0	0
0	-2	-4	1	1	0	0	0
-3	1	5	-1	-1	0	0	0
-4	1	2	-1	0	0	0	0
1	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0

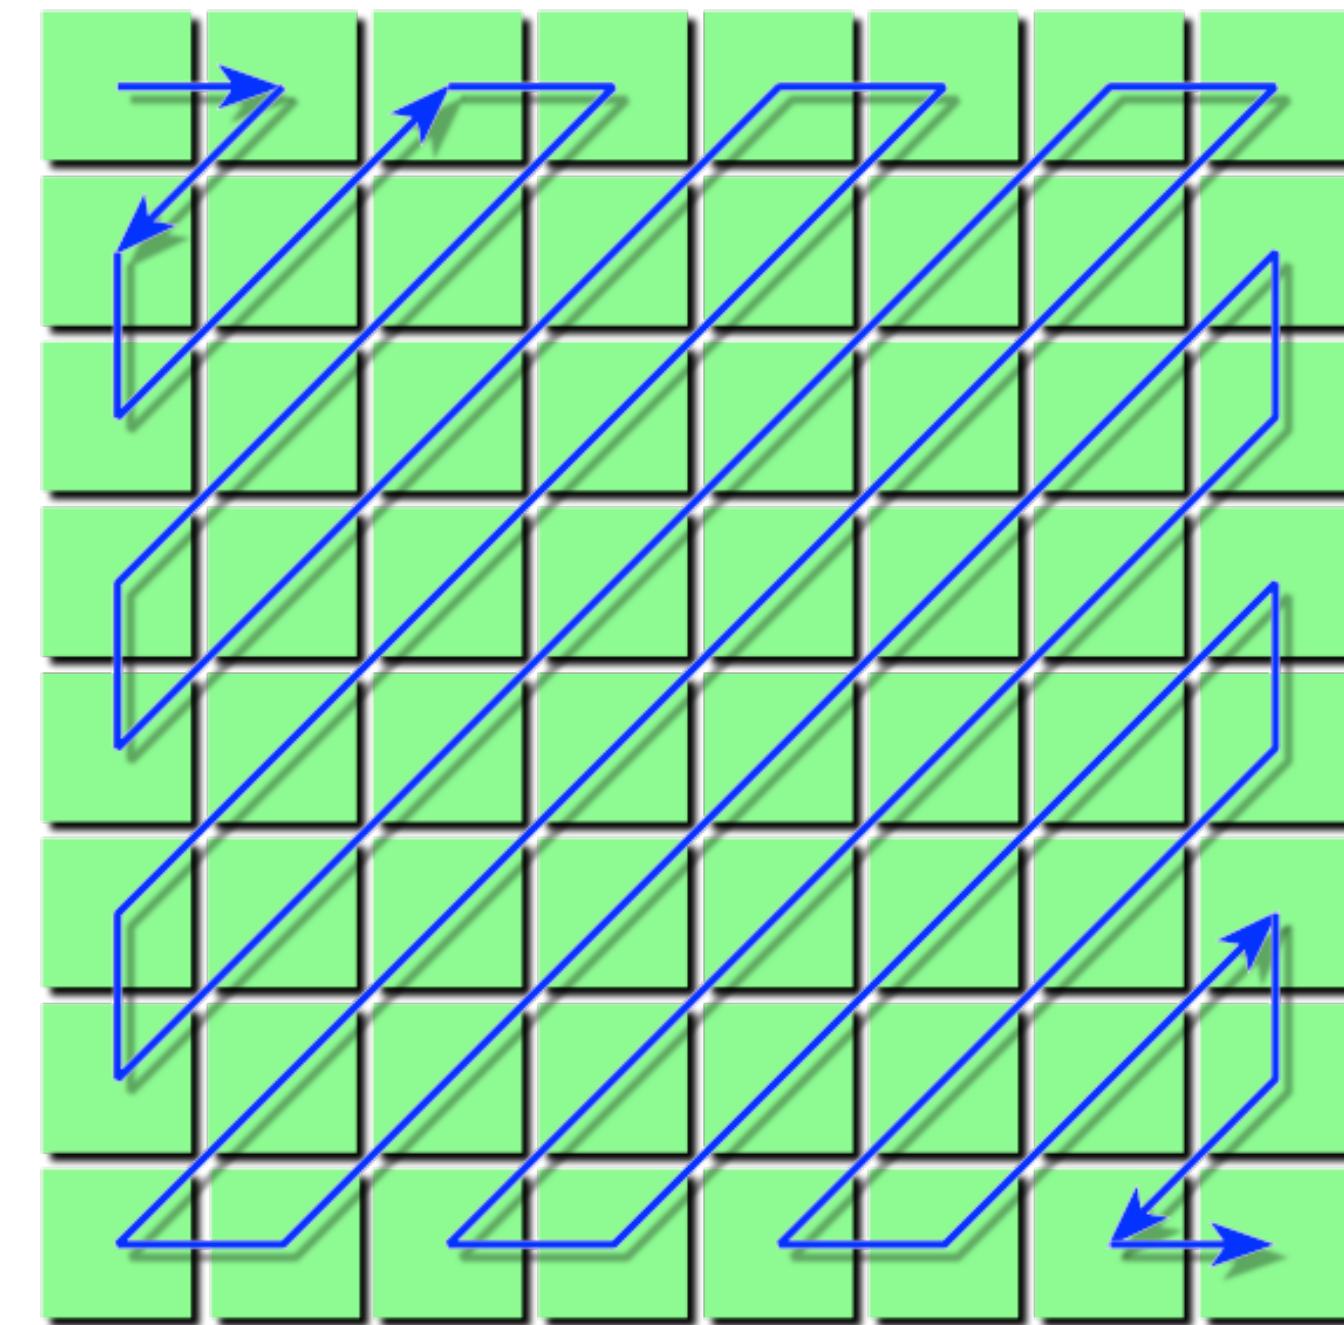
Quantized DCT Values

Entropy encoding: (lossless)

Reorder values

Run-length encode (RLE) 0's

Huffman encode non-zero values



Reordering

JPEG compression summary

$$\begin{bmatrix} -415 & -30 & -61 & 27 & 56 & -20 & -2 & 0 \\ 4 & -22 & -61 & 10 & 13 & -7 & -9 & 5 \\ -47 & 7 & 77 & -25 & -29 & 10 & 5 & -6 \\ -49 & 12 & 34 & -15 & -10 & 6 & 2 & 2 \\ 12 & -7 & -13 & -4 & -2 & 2 & -3 & 3 \\ -8 & 3 & 2 & -6 & -2 & 1 & 4 & 2 \\ -1 & 0 & 0 & -2 & -1 & -3 & 4 & -1 \\ 0 & 0 & -1 & -4 & -1 & 0 & 1 & 2 \end{bmatrix}$$

DCT

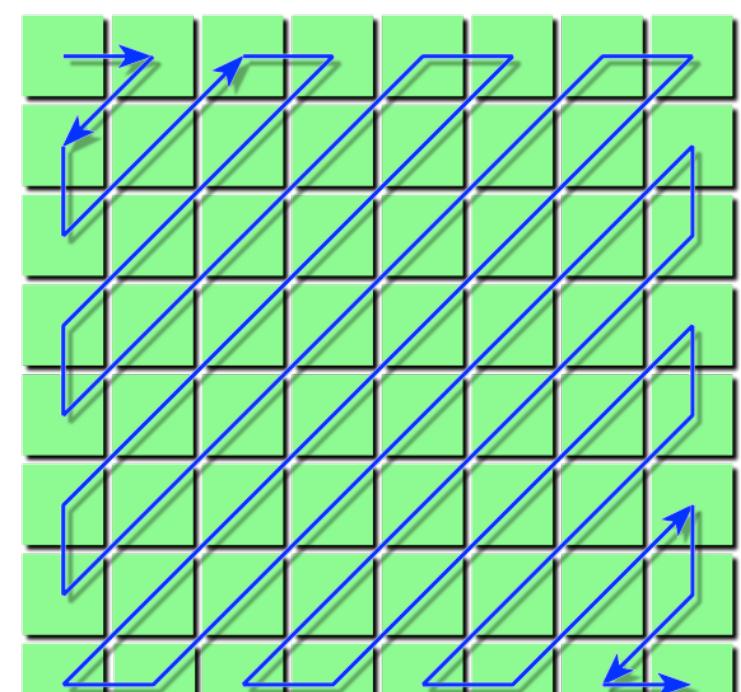
$$\begin{bmatrix} 16 & 11 & 10 & 16 & 24 & 40 & 51 & 61 \\ 12 & 12 & 14 & 19 & 26 & 58 & 60 & 55 \\ 14 & 13 & 16 & 24 & 40 & 57 & 69 & 56 \\ 14 & 17 & 22 & 29 & 51 & 87 & 80 & 62 \\ 18 & 22 & 37 & 56 & 68 & 109 & 103 & 77 \\ 24 & 35 & 55 & 64 & 81 & 104 & 113 & 92 \\ 49 & 64 & 78 & 87 & 103 & 121 & 120 & 101 \\ 72 & 92 & 95 & 98 & 112 & 100 & 103 & 99 \end{bmatrix}$$

Quantization Matrix

$$= \begin{bmatrix} -26 & -3 & -6 & 2 & 2 & -1 & 0 & 0 \\ 0 & -2 & -4 & 1 & 1 & 0 & 0 & 0 \\ -3 & 1 & 5 & -1 & -1 & 0 & 0 & 0 \\ -4 & 1 & 2 & -1 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

Quantized DCT

Quantization loses information
(lossy compression!)



Coefficient reordering

RLE compression of zeros
Entropy compression of
non-zeros

Lossless compression!

Compressed bits

JPEG compression summary

Convert image to Y'CbCr

Downsample CbCr (to 4:2:2 or 4:2:0) **(information loss occurs here)**

For each color channel (Y', Cb, Cr):

For each 8x8 block of values

Compute DCT

Quantize results

(information loss occurs here)

Reorder values

Run-length encode 0-spans

Huffman encode non-zero values

Key idea: exploit characteristics of human perception to build efficient image storage and image processing systems

- Separation of luminance from chrominance in color representation (Y'CrCb) allows reduced resolution in chrominance channels (4:2:0)
- Encode pixel values linearly in lightness (perceived brightness), not in luminance (distribute representable values uniformly in perceptual space)
- JPEG compression significantly reduces file size at cost of quantization error in high spatial frequencies
 - Human brain is more tolerant of errors in high frequency image components than in low frequency ones
 - Images of the real world are dominated by low-frequency components

H.264 Video Compression

Example video



Go Swallows!

30 second video: 1920 x 1080, @ 30fps

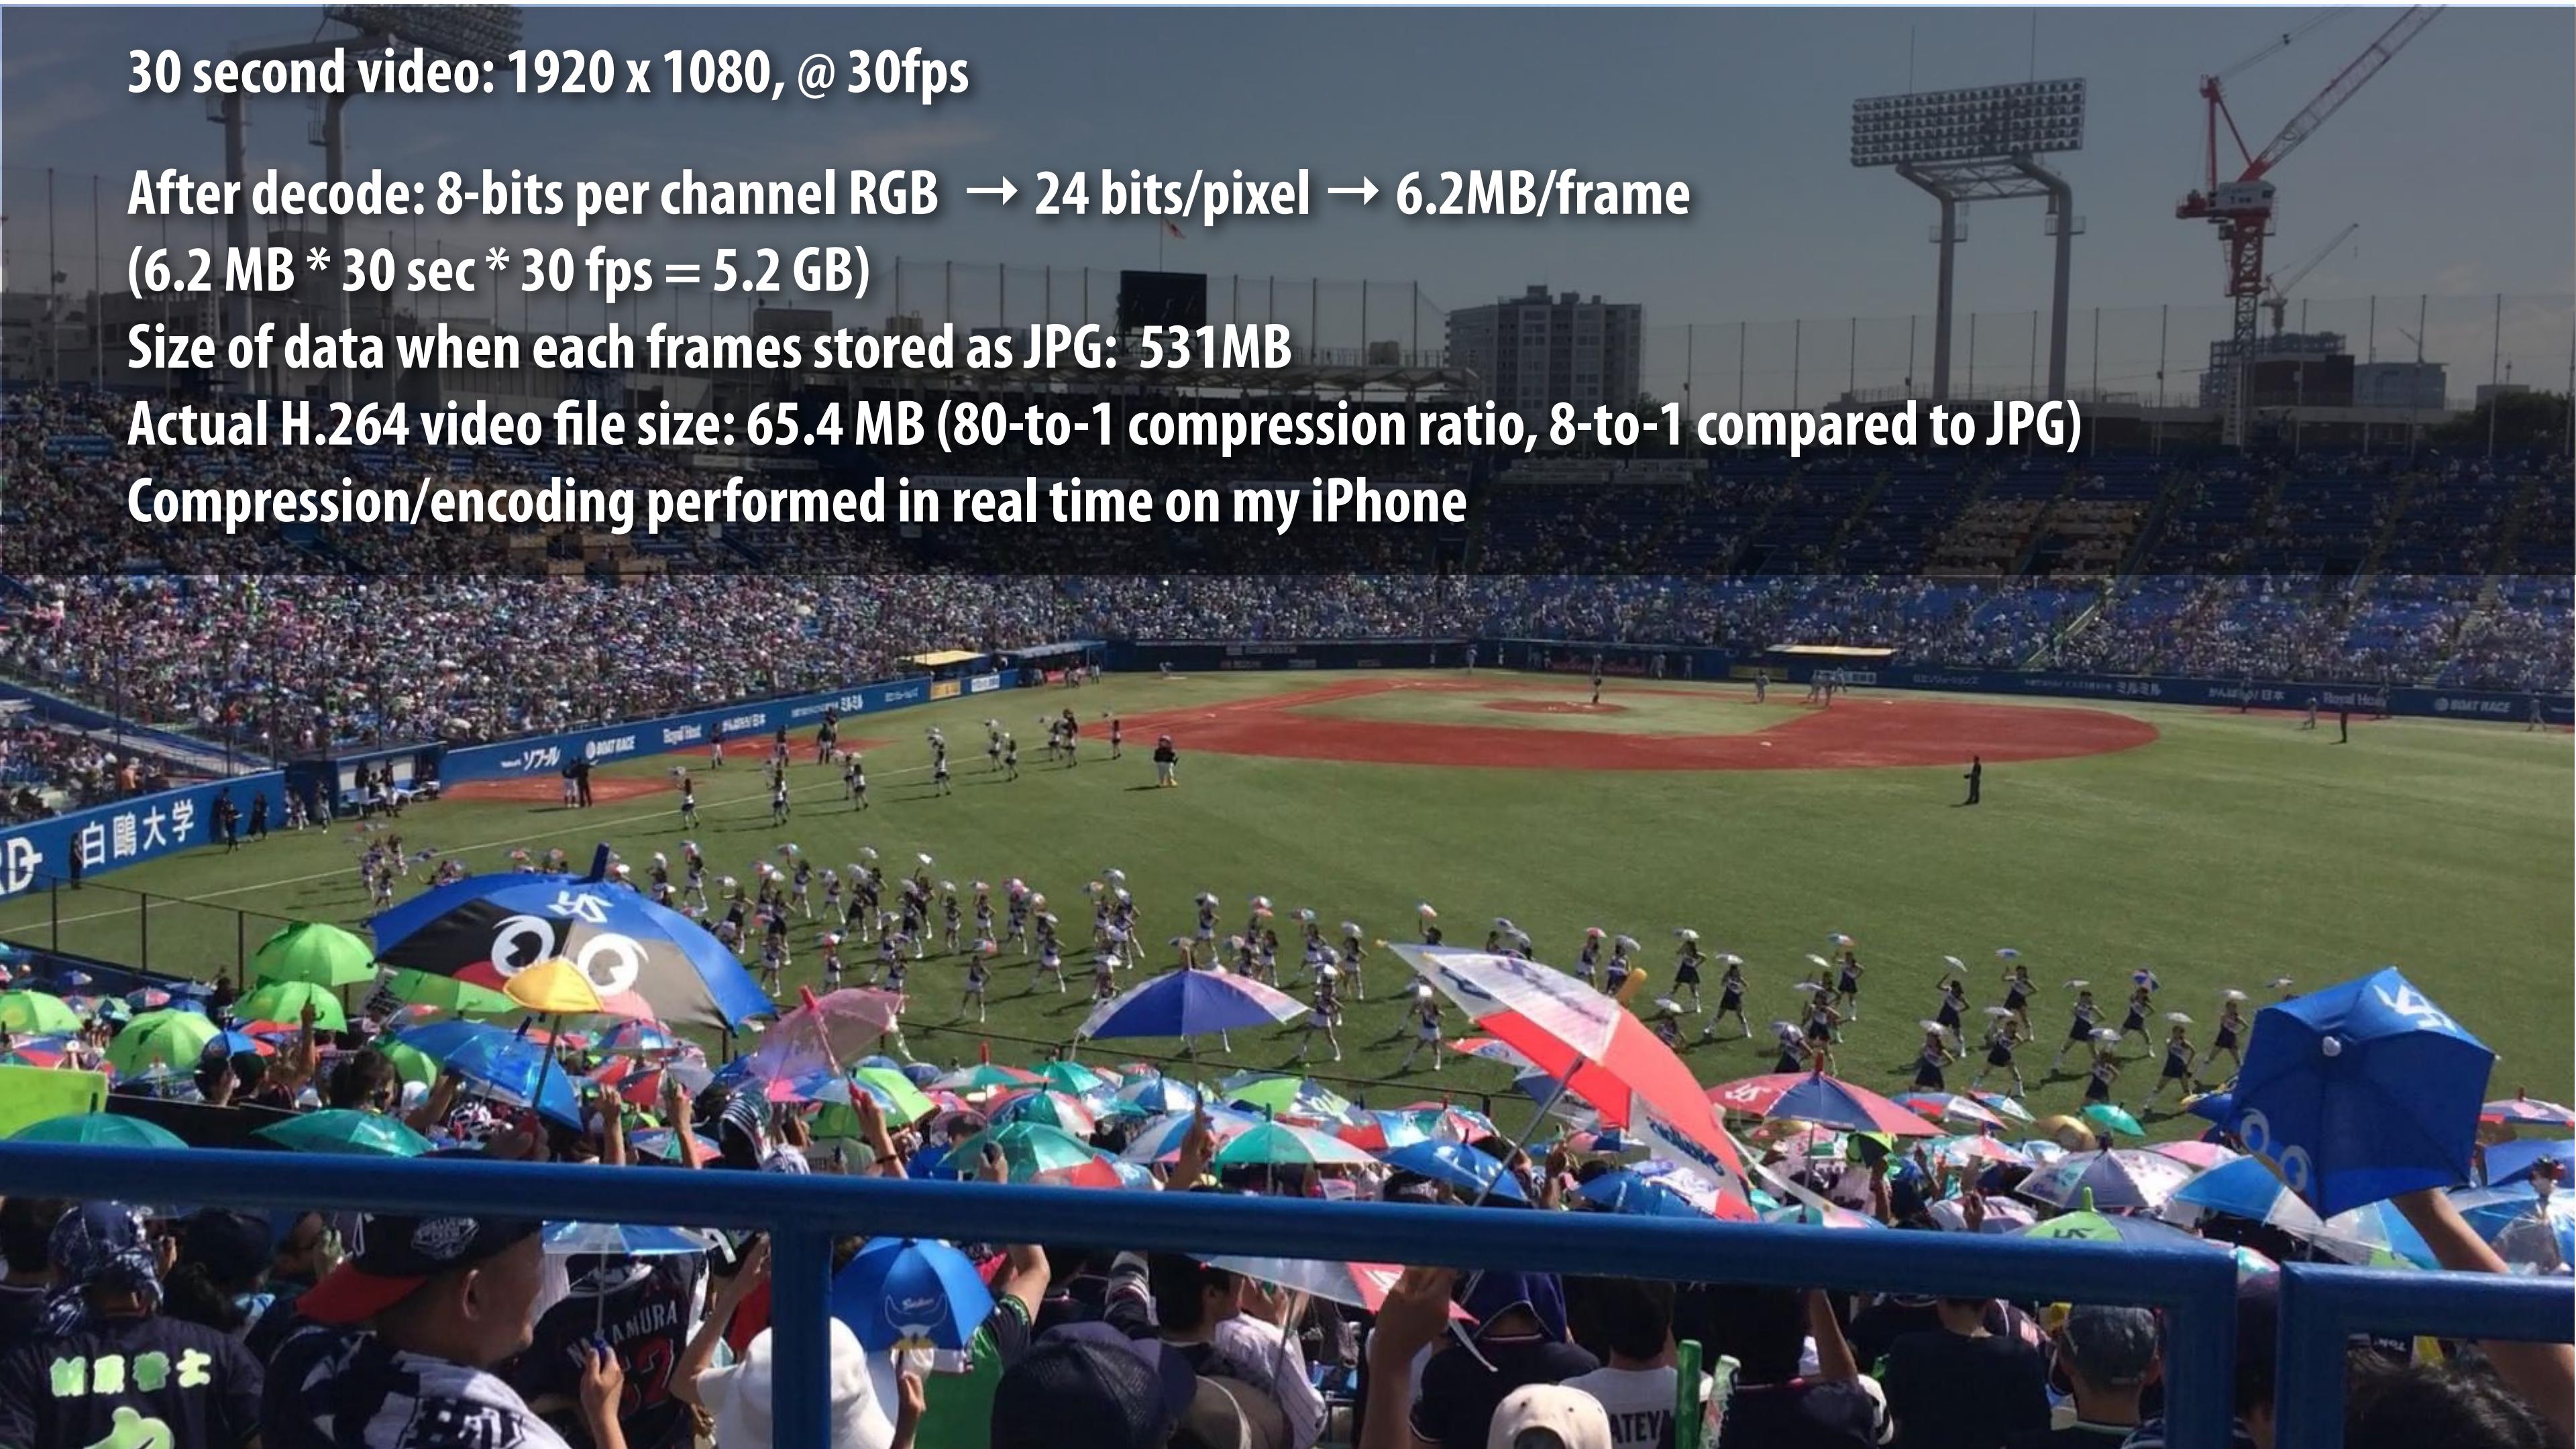
After decode: 8-bits per channel RGB → 24 bits/pixel → 6.2MB/frame

(6.2 MB * 30 sec * 30 fps = 5.2 GB)

Size of data when each frames stored as JPG: 531MB

Actual H.264 video file size: 65.4 MB (80-to-1 compression ratio, 8-to-1 compared to JPG)

Compression/encoding performed in real time on my iPhone



H.264/AVC video compression

- **AVC = advanced video coding**
- **Also called MPEG4 Part 10**
- **Common format in many modern HD video applications:**
 - **Blue Ray**
 - **HD streaming video on internet (Youtube, Vimeo, iTunes store, etc.)**
 - **HD video recorded by your smart phone**
 - **European broadcast HDTV (U.S. broadcast HDTV uses MPEG 2)**
 - **Some satellite TV broadcasts (e.g., DirecTV)**
- **Benefit: much higher compression ratios than MPEG2 or MPEG4**
 - **Alternatively, higher quality video for fixed bit rate**
- **Costs: higher decoding complexity, substantially higher encoding cost**
 - **Idea: trades off more compute for requiring less bandwidth/storage**

Hardware implementations

- Support for H.264 video encode/decode is provided by fixed-function hardware on many modern processors (not just mobile devices)
- Hardware encoding/decoding support existed in modern Intel CPUs since Sandy Bridge (Intel “Quick Sync”)
- Modern operating systems expose hardware encode decode support through hardware-accelerated APIs
 - e.g., DirectShow/DirectX (Windows), AVFoundation (iOS)

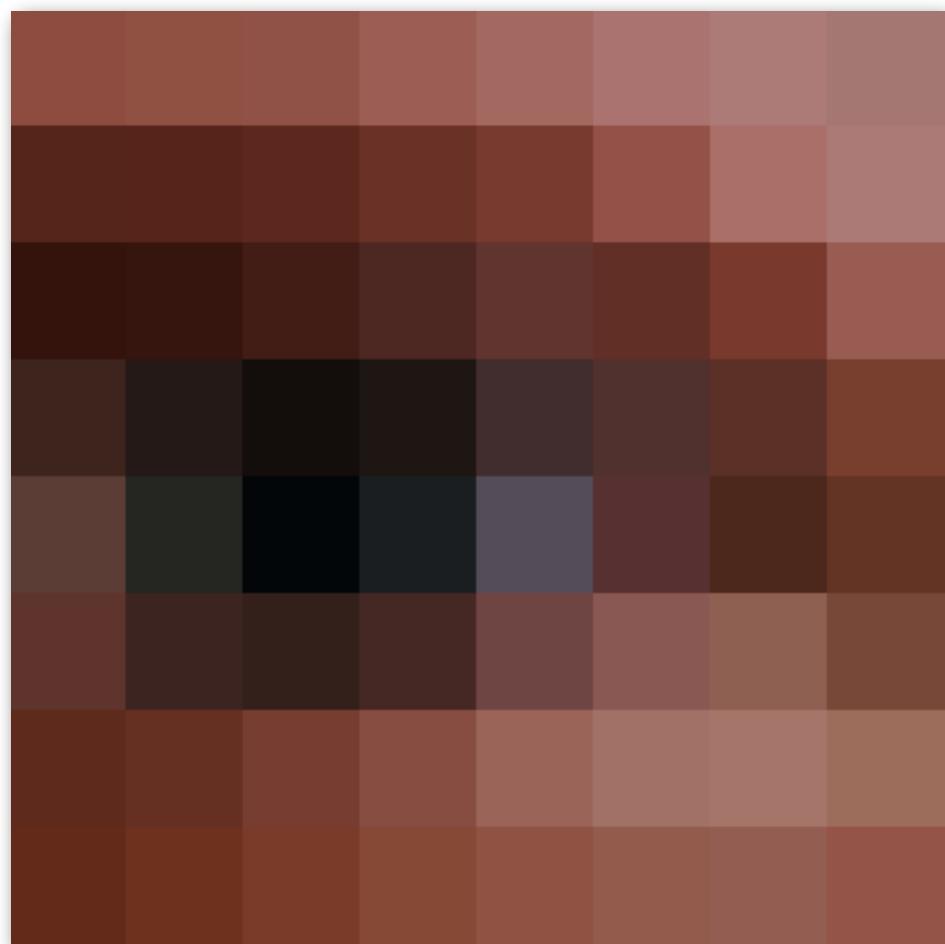
Video container format versus video codec

- **Video container (MOV, AVI) bundles media assets**
- **Video codec: H.264/AVC (MPEG 4 Part 10)**
 - **H.264 standard defines how to represent and decode video**
 - **H.264 does not define how to encode video (this is left up to implementations)**
 - **H.264 has many profiles**
 - **High Profile (HiP): supported by HDV and Blue Ray**

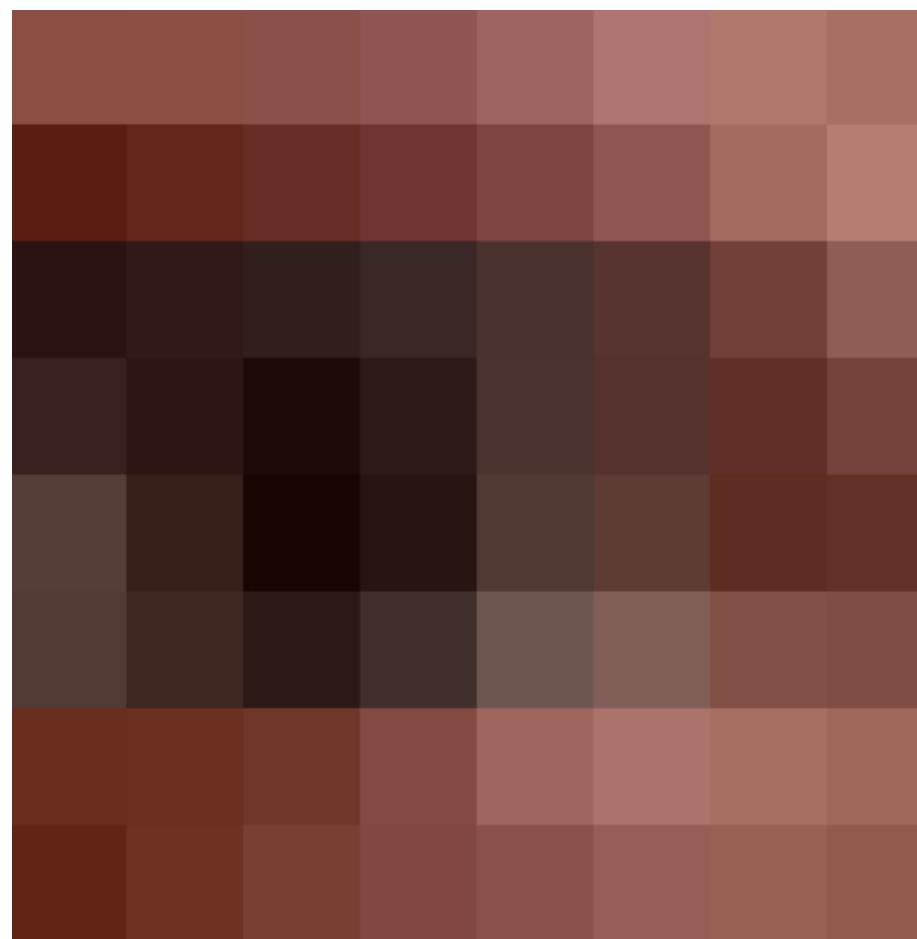
Video compression: main ideas

- **Compression is about exploiting redundancy in a signal**
 - **Intra-frame redundancy:** value of pixels in neighboring regions of a frame are good predictor of values for other pixels in the frame (**spatial redundancy**)
 - **Inter-frame redundancy:** pixels from nearby frames in time are a good predictor for the current frame's pixels (**temporal redundancy**)

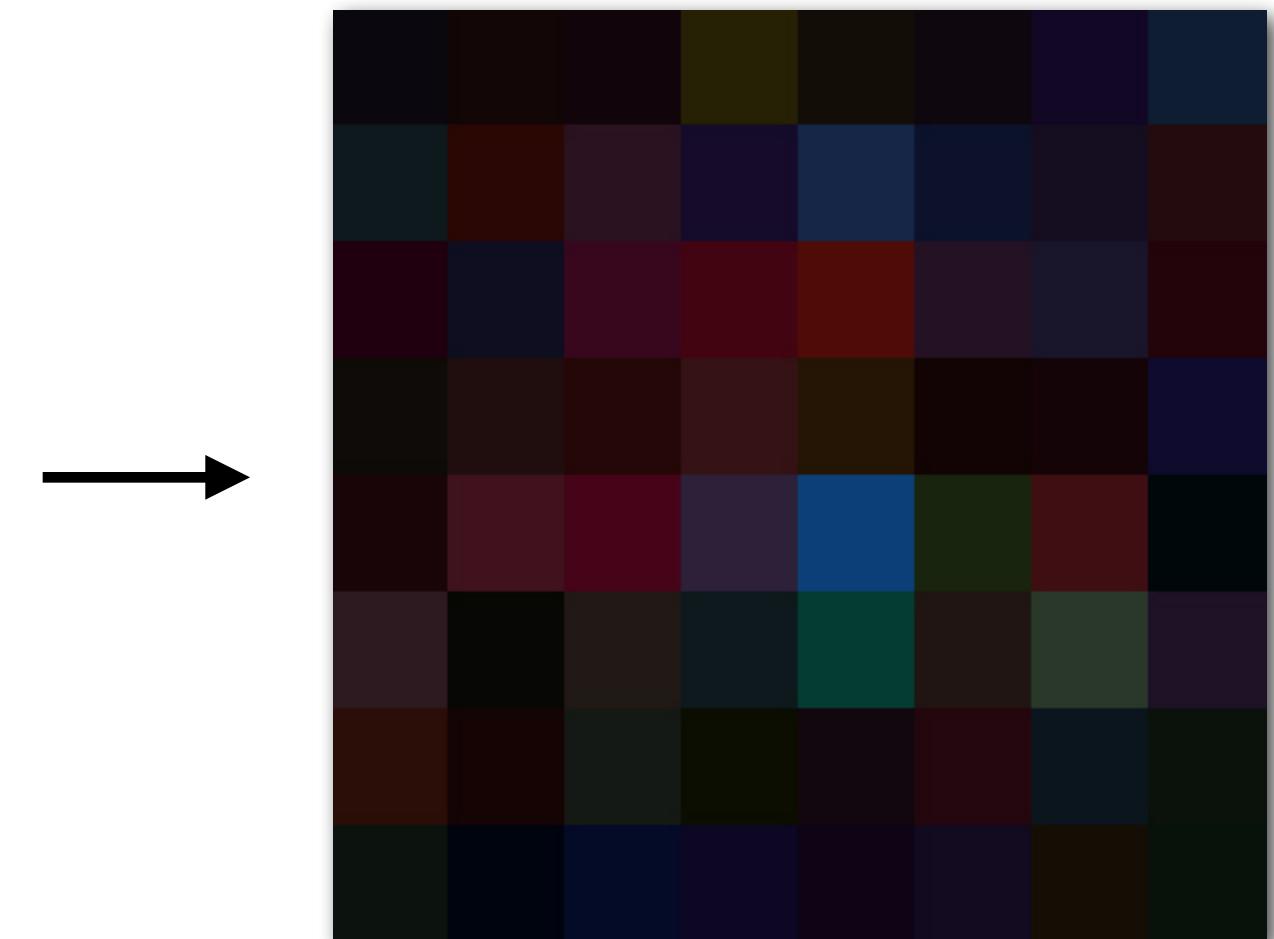
Residual: difference between compressed image and original image



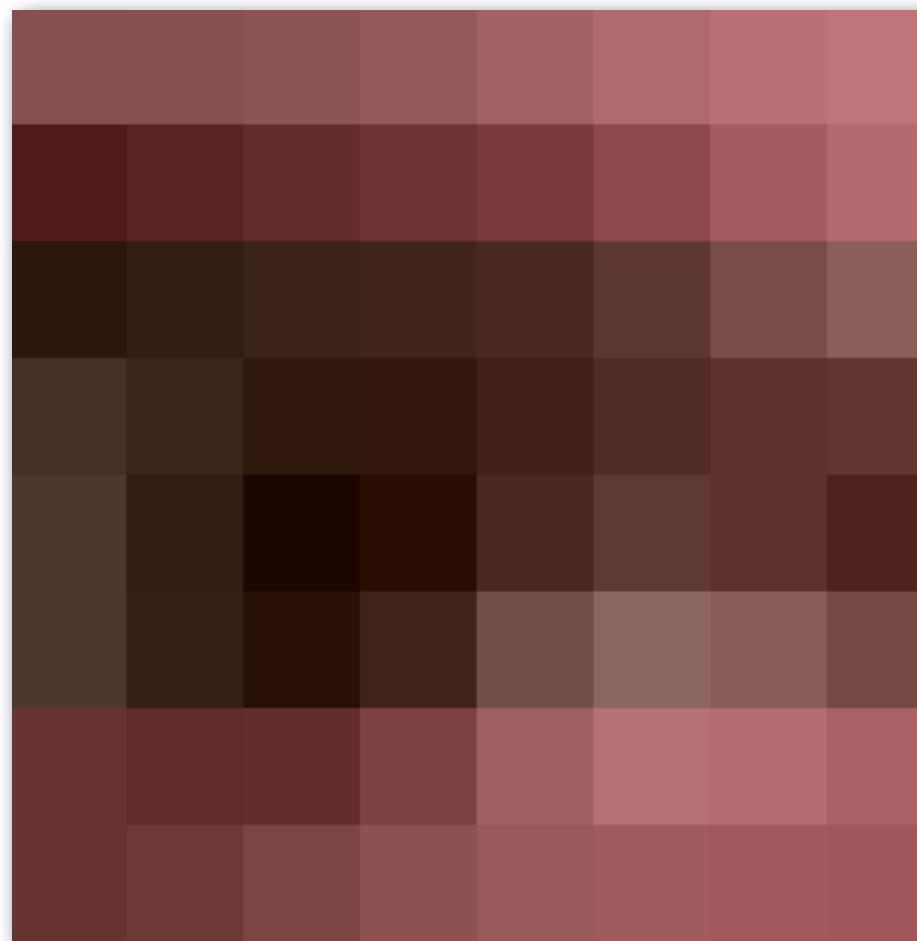
Original pixels



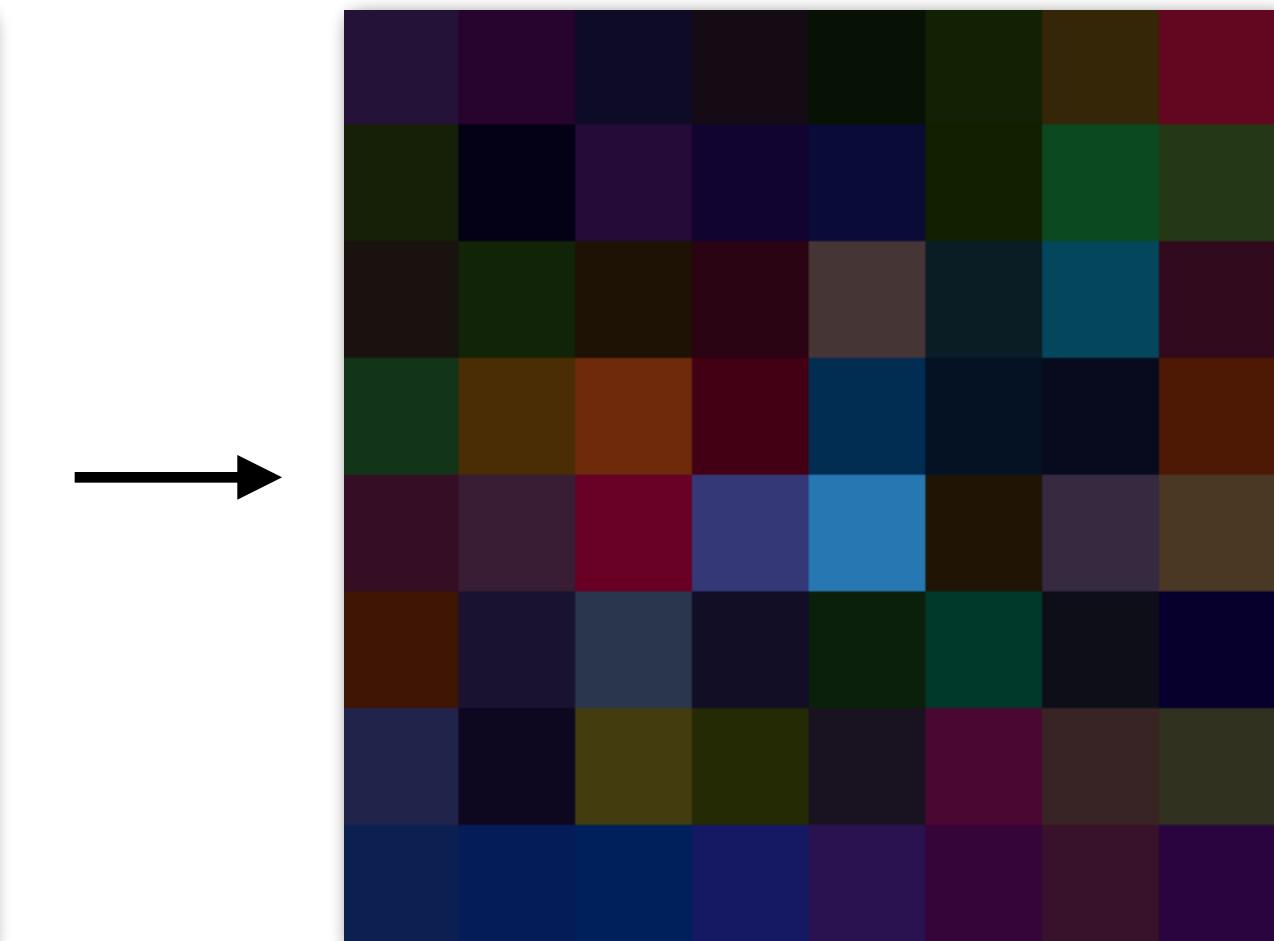
Compressed pixels
(JPEG quality level 6)



Residual
(amplified for visualization)

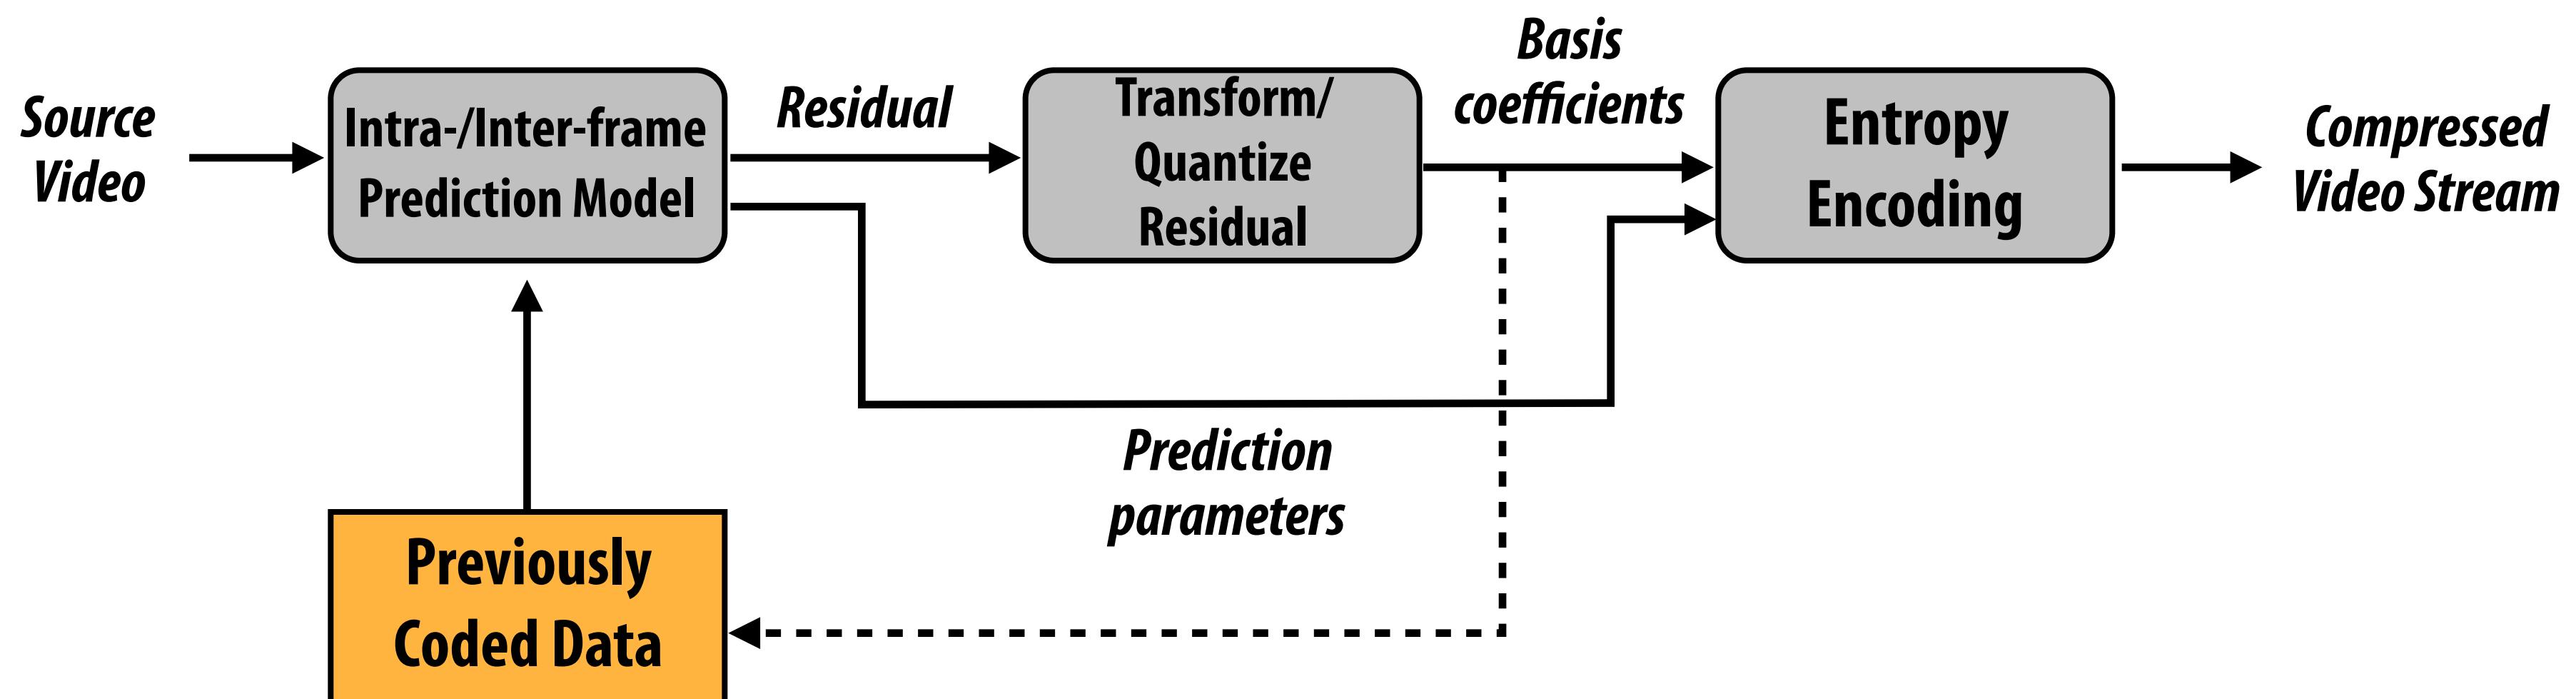


Compressed pixels
(JPEG quality level 2)



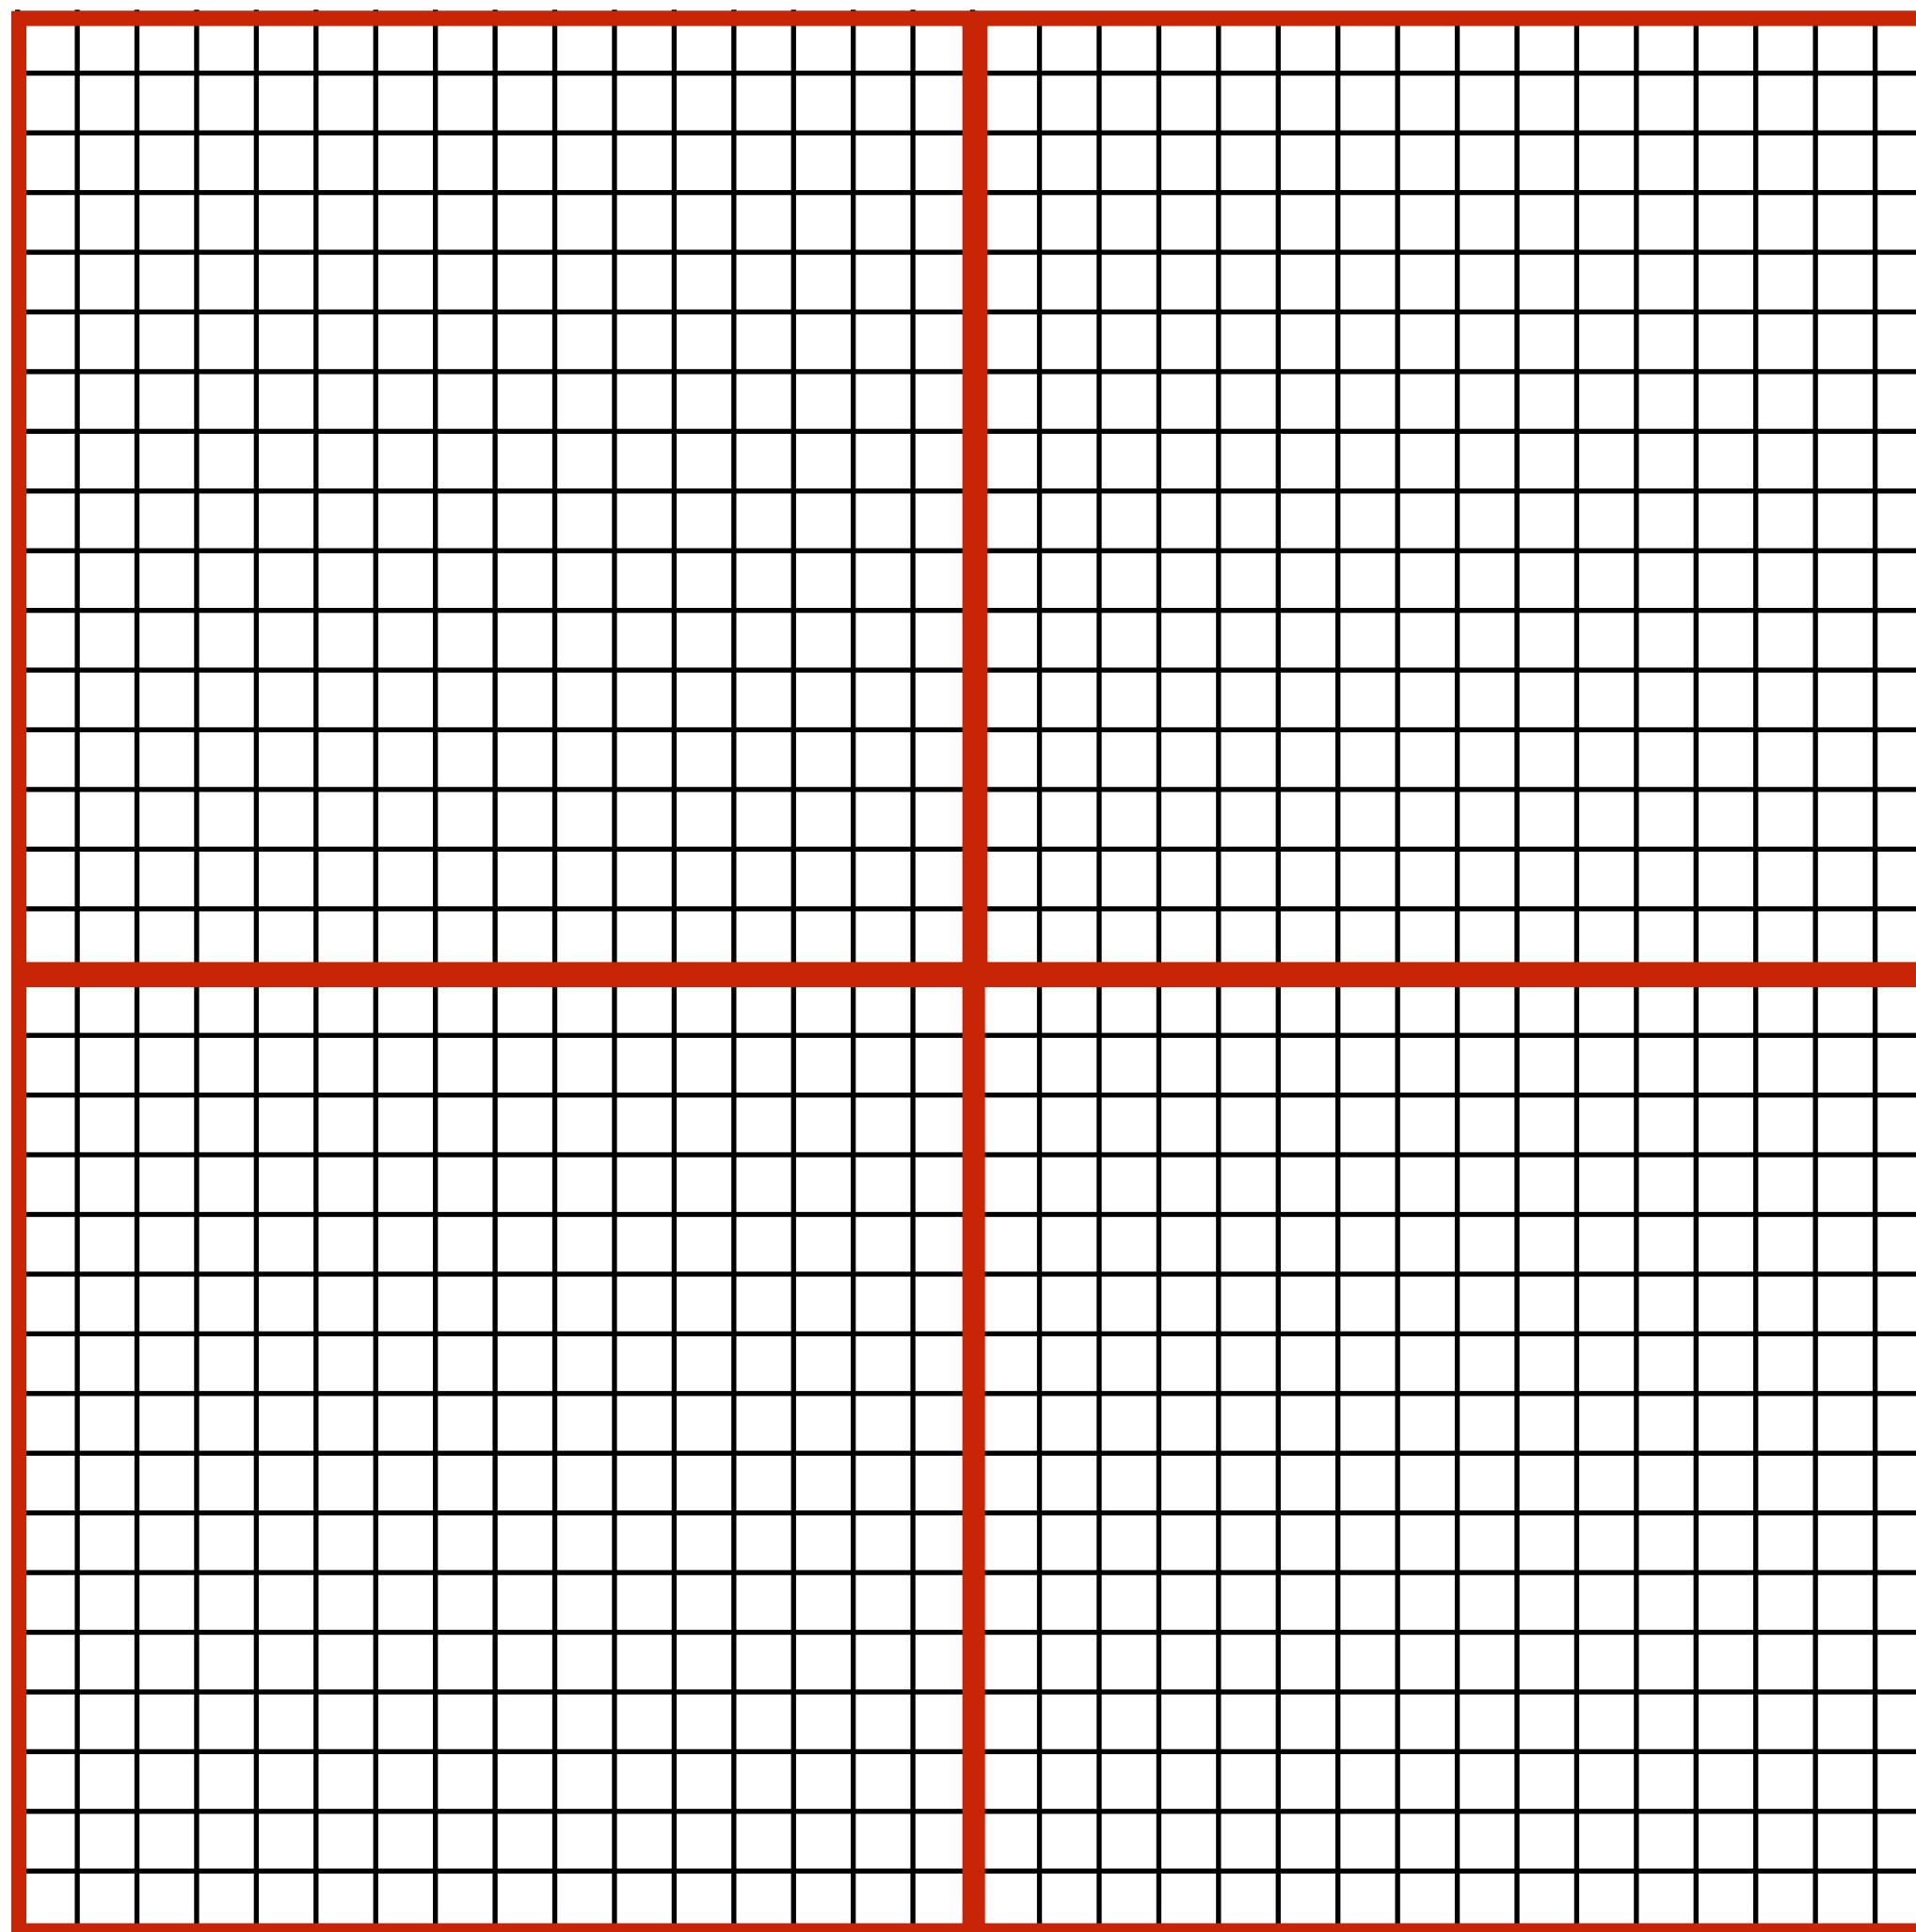
Residual
(amplified for visualization)

H.264/AVC video compression overview



Residual: difference between predicted pixel values and input video pixel values

16 x 16 macroblocks



Video frame is partitioned into 16 x 16 pixel macroblocks

**Due to 4:2:0 chroma subsampling,
macroblocks correspond to 16 x 16 luma
samples and 8 x 8 chroma samples**

Macroblocks in an image are organized into slices

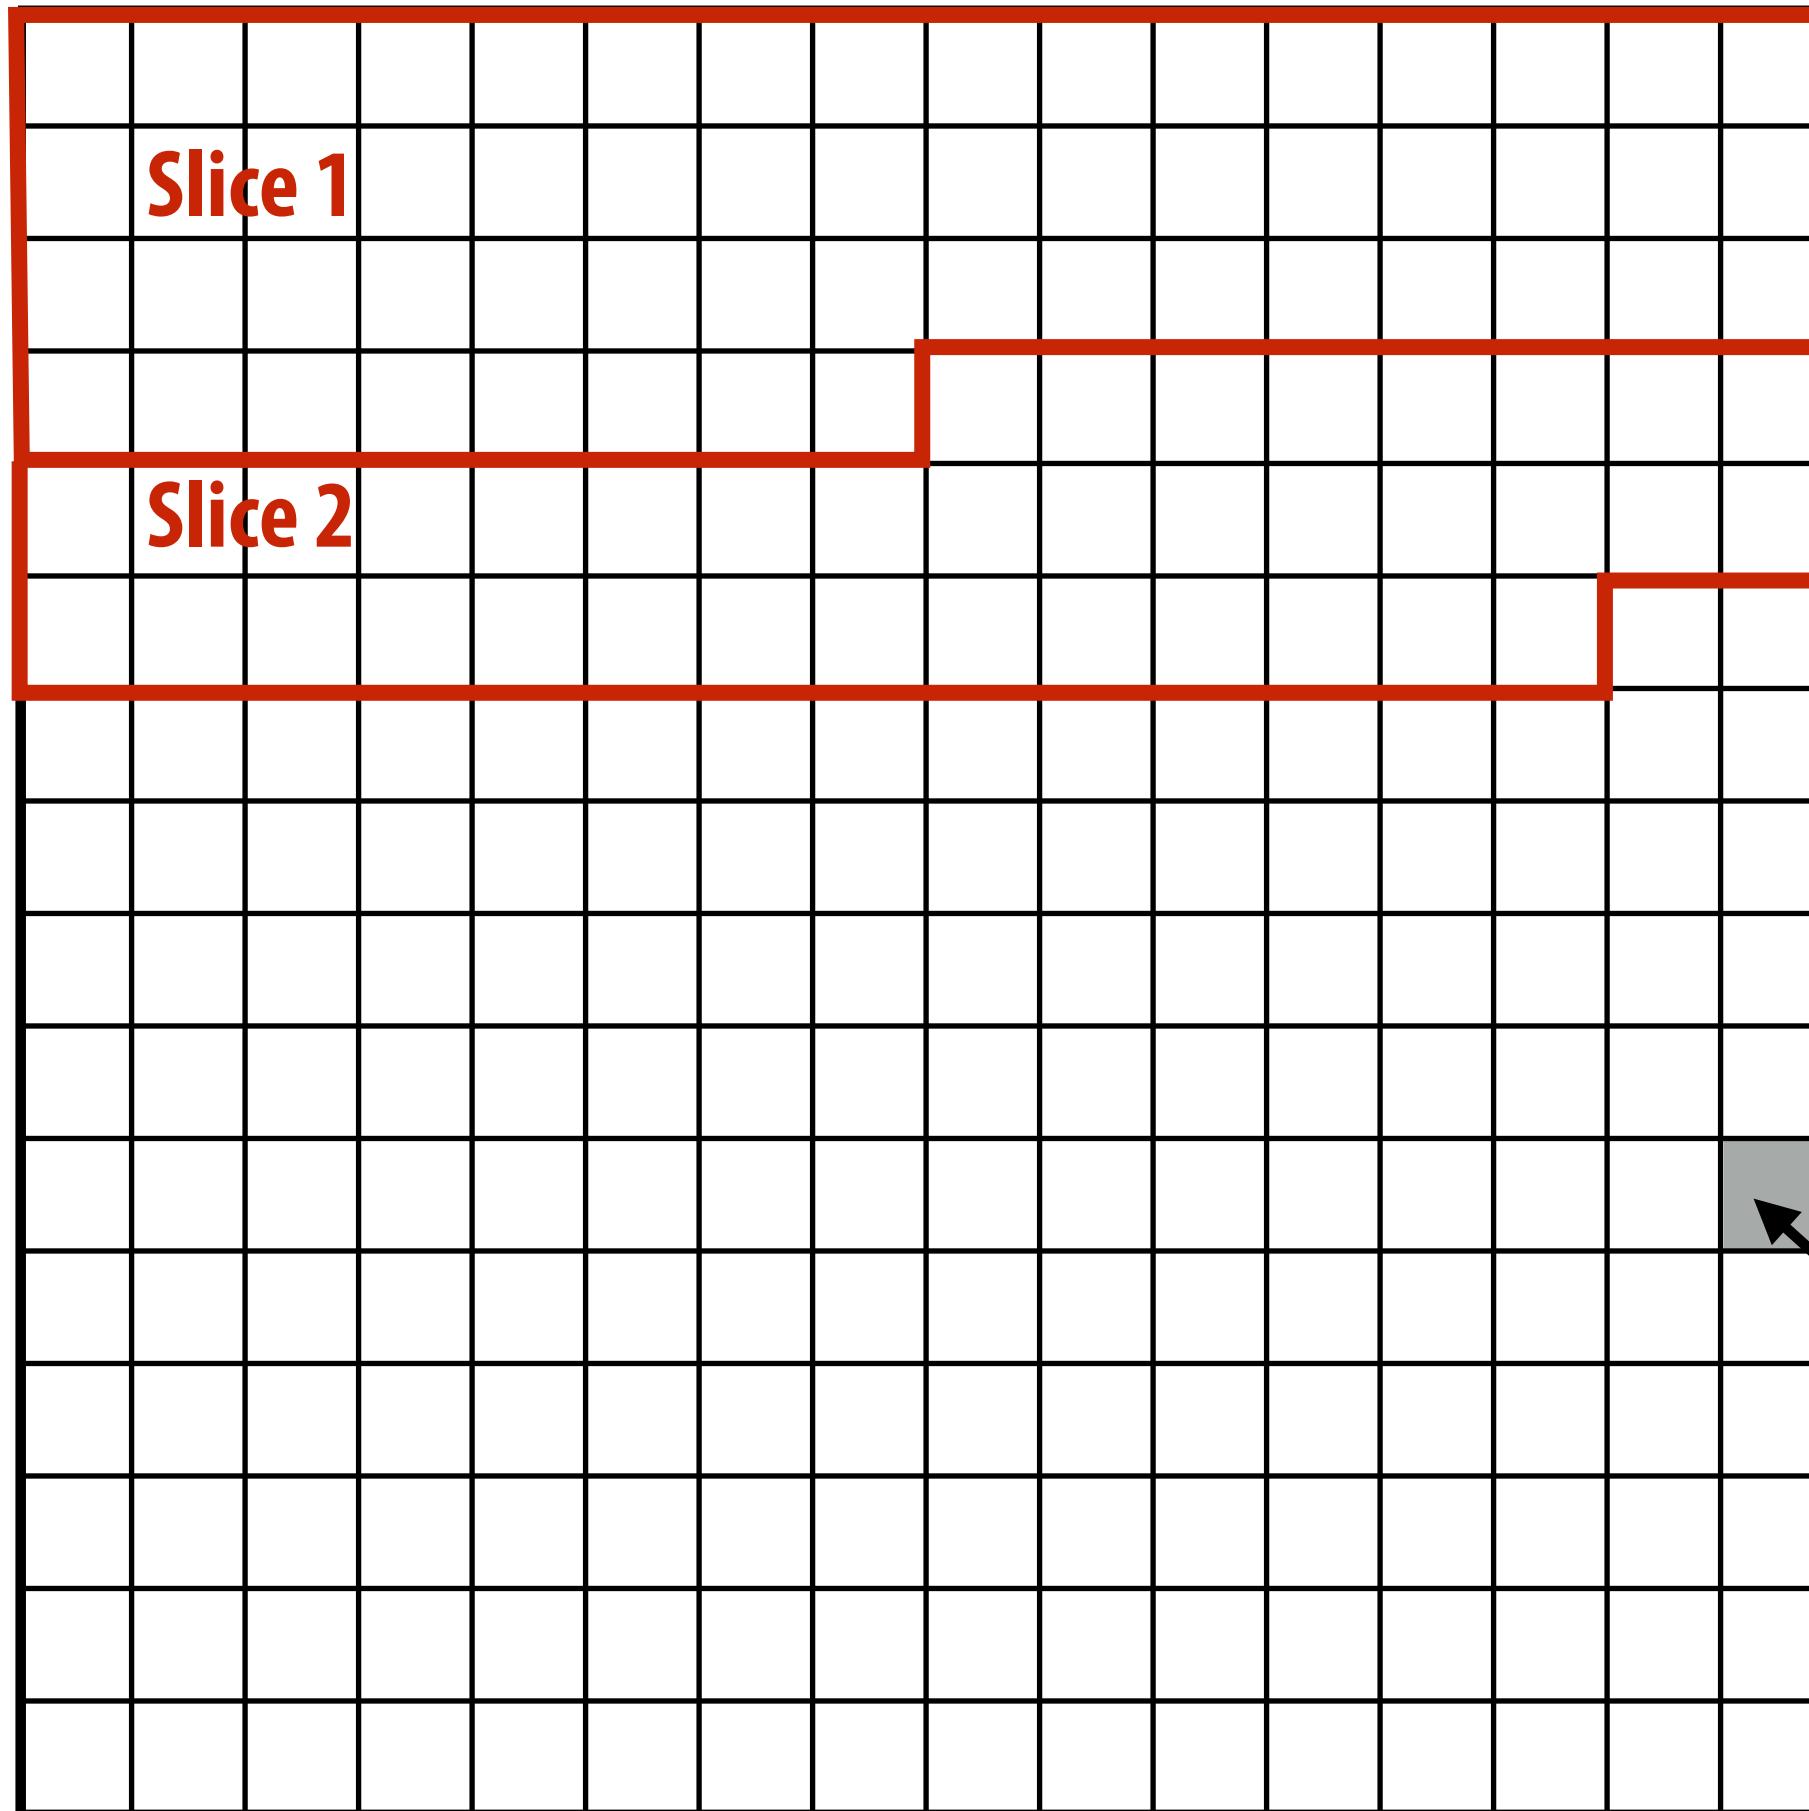


Figure to left shows the macro blocks in a frame

Macroblocks are grouped into slices

Can think of a slice as a sequence of macroblocks in raster scan order *

Slices can be decoded independently **

One 16x16 macroblock

* H.264 also has non-raster-scan order modes (FMO), will not discuss today.

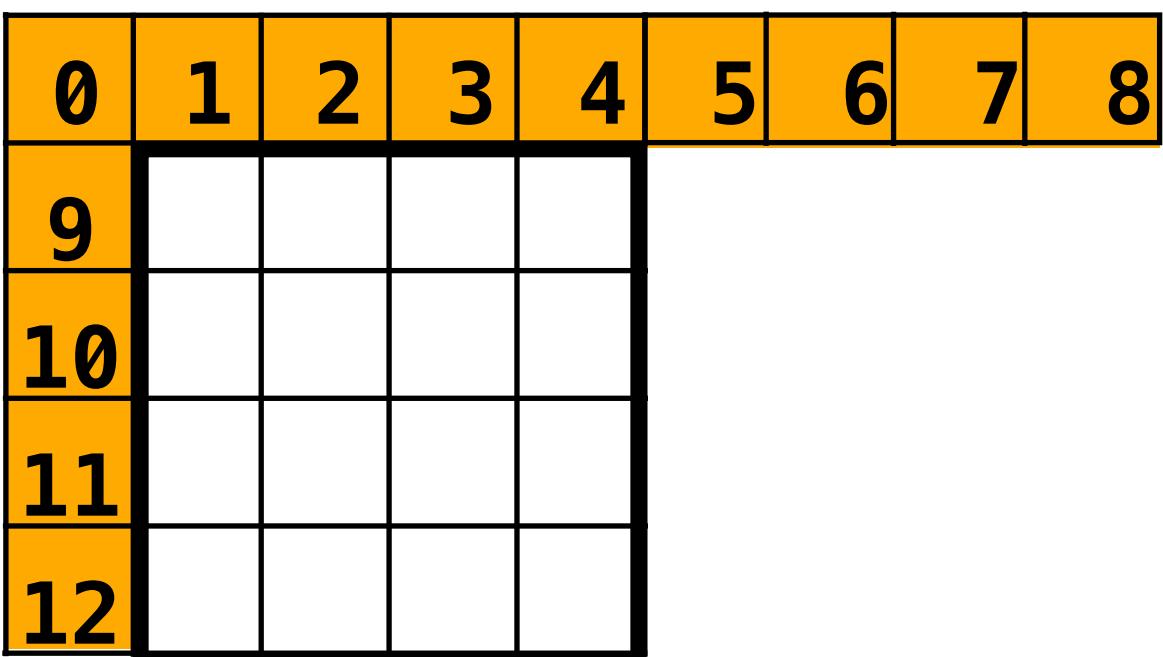
** Final “deblocking” pass is often applied to post-decode pixel data, so technically slices are not fully independent.

Decoding via prediction + correction

- During decode, samples in a macroblock are generated by:
 1. Making a prediction based on already decoded samples in macroblocks from the same frame (intra-frame prediction) or from other frames (inter-frame prediction)
 2. Correcting the prediction with a residual stored in the video stream
- Three forms of prediction:
 - **I-macroblock:** macroblock samples predicted from samples in previous macroblocks in the same slice of the current frame
 - **P-macroblock:** macroblock samples can be predicted from samples from one other frame (one prediction per macroblock)
 - **B-macroblock:** macroblock samples can be predicted by a weighted combination of multiple predictions from samples from other frames

Intra-frame prediction (I-macroblock)

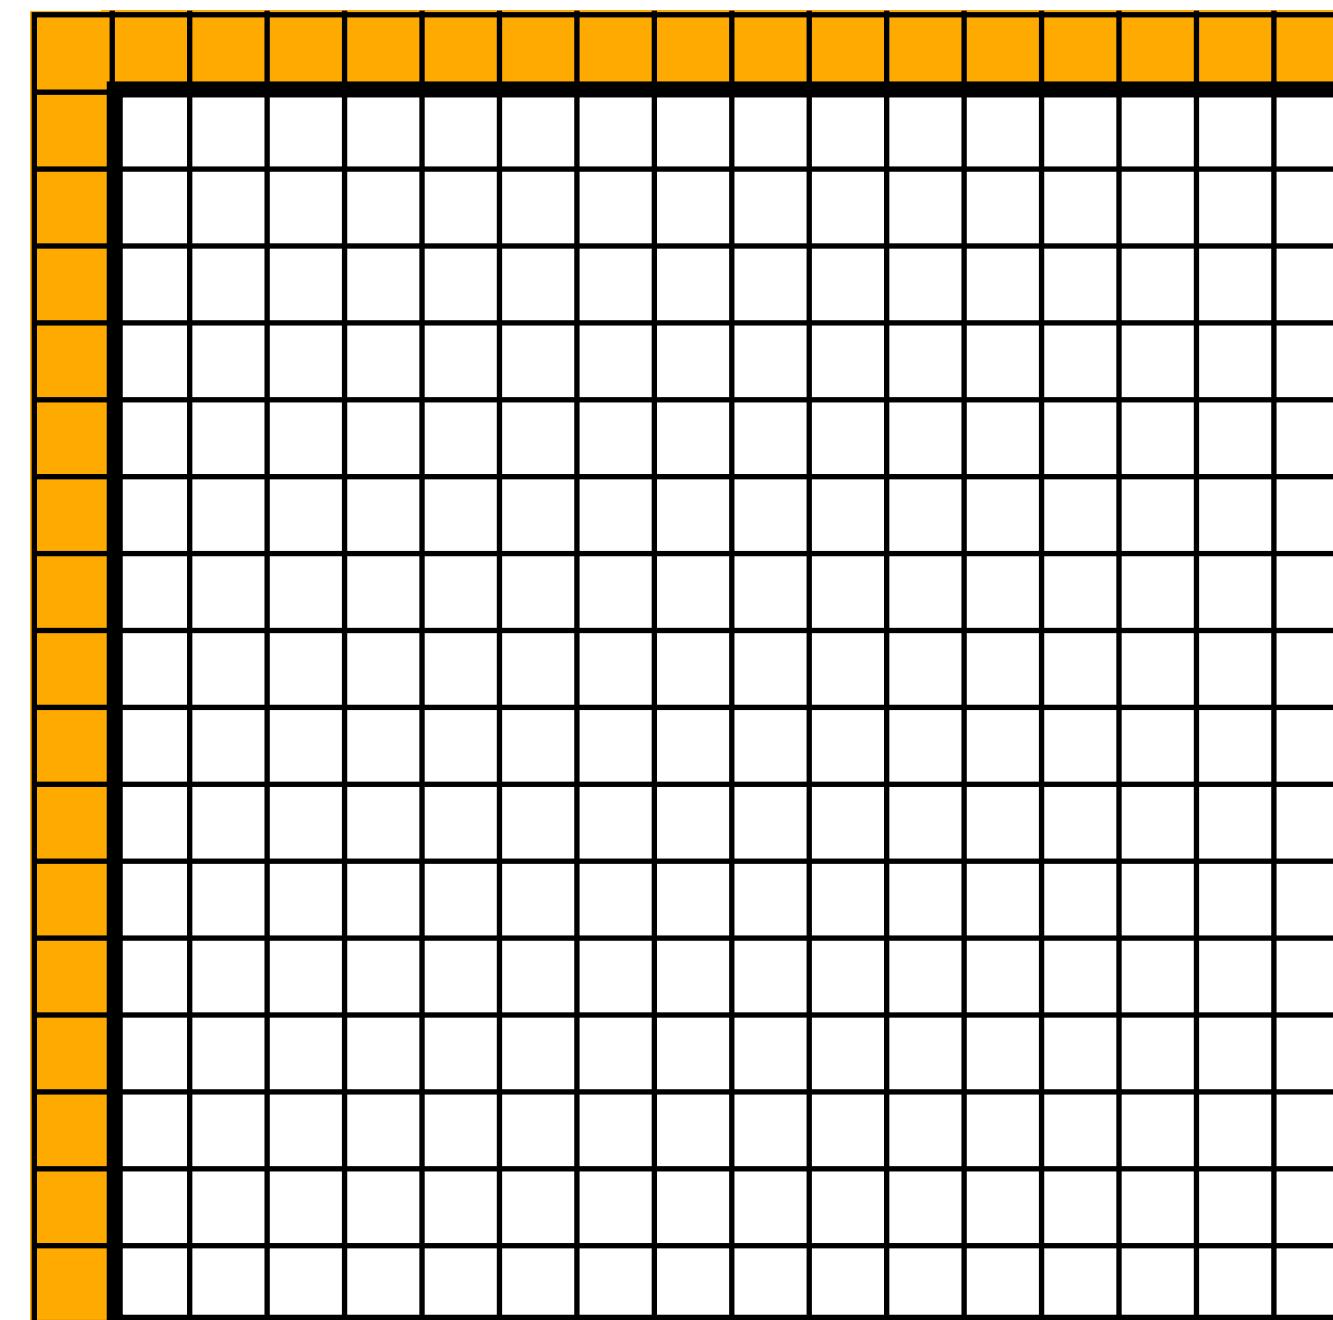
- Prediction of sample values is performed in spatial domain, not transform domain
 - Predicting pixel values, not basis coefficients
- Modes for predicting the 16x16 luma (Y) values: *
 - Intra_4x4 mode: predict 4x4 block of samples from adjacent row/col of pixels
 - Intra_16x16 mode: predict entire 16x16 block of pixels from adjacent row/col
 - I_PCM: actual sample values provided



Intra_4X4

Yellow pixels: already reconstructed (values known)

White pixels: 4x4 block to be reconstructed

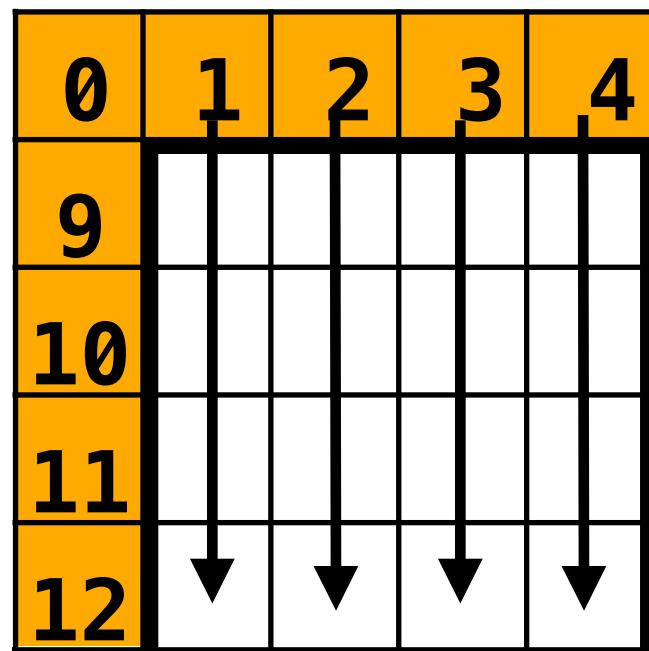


Intra_16x16

* An additional 8x8 mode exists in the H.264 High Profile

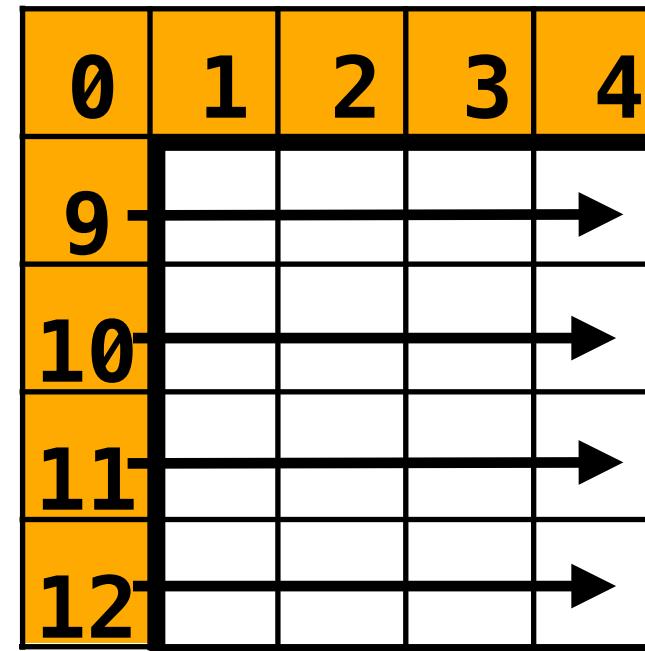
Intra_4x4 prediction modes

- Nine prediction modes (6 shown below)
 - Other modes: horiz-down, vertical-left, horiz-up



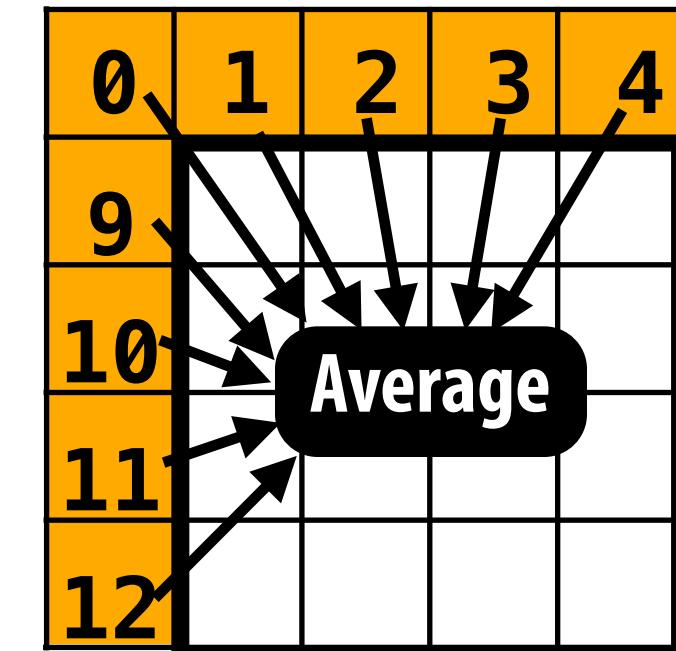
Mode 0: vertical

(4x4 block is copy of
above row of pixels)



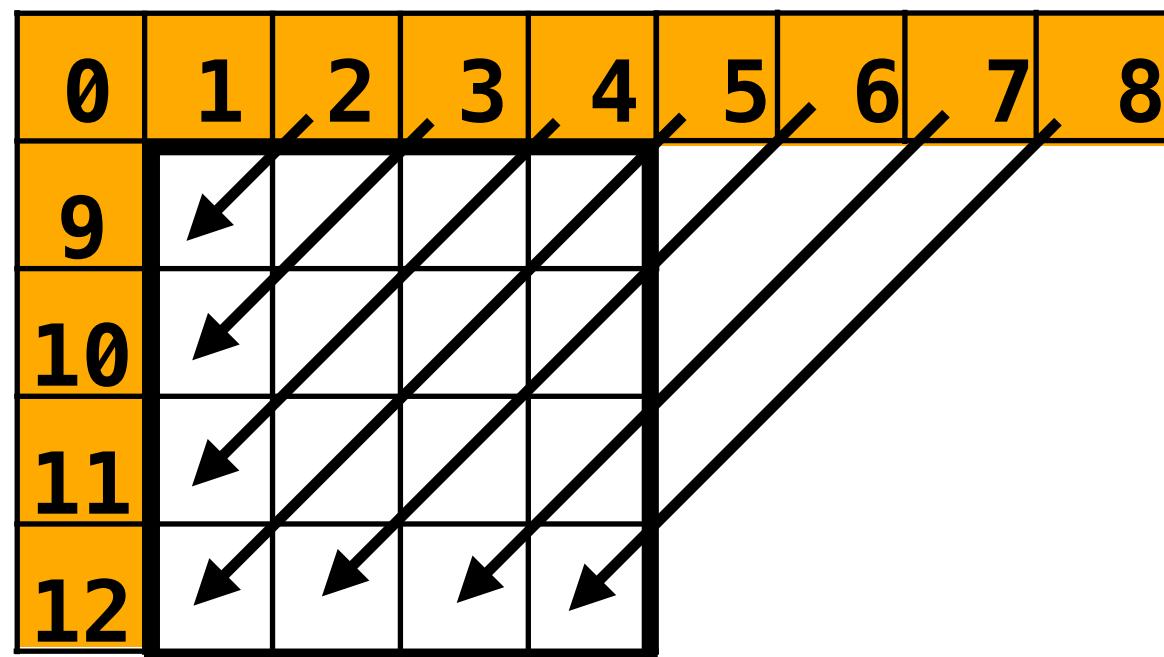
Mode 1: horizontal

(4x4 block is copy of left
col of pixels)

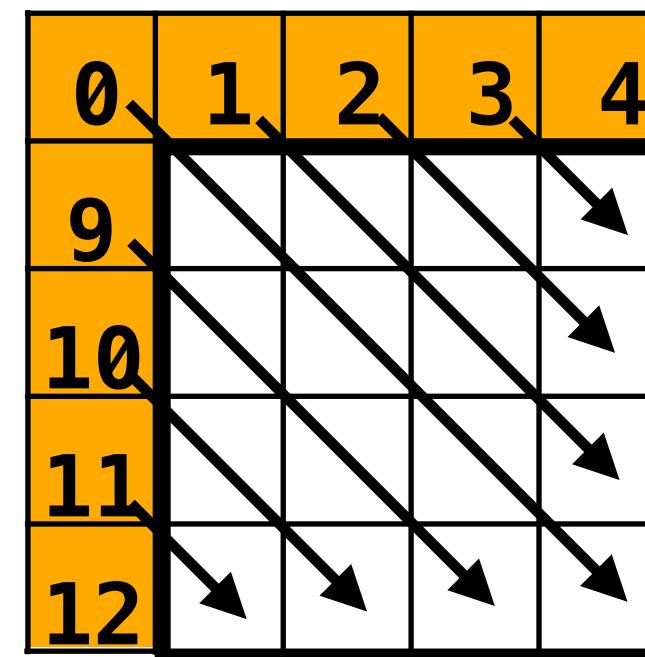


Mode 2: DC

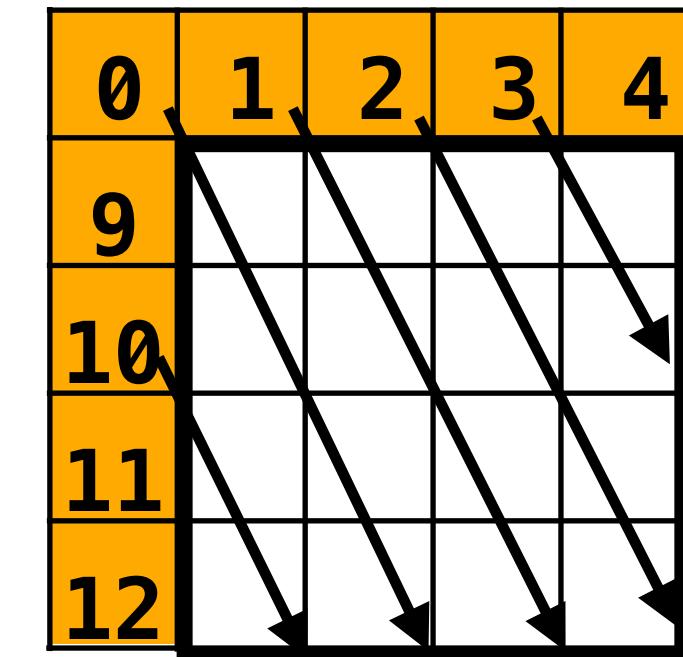
(4x4 block is average of above
row and left col of pixels)



Mode 3: diagonal down-left (45°)



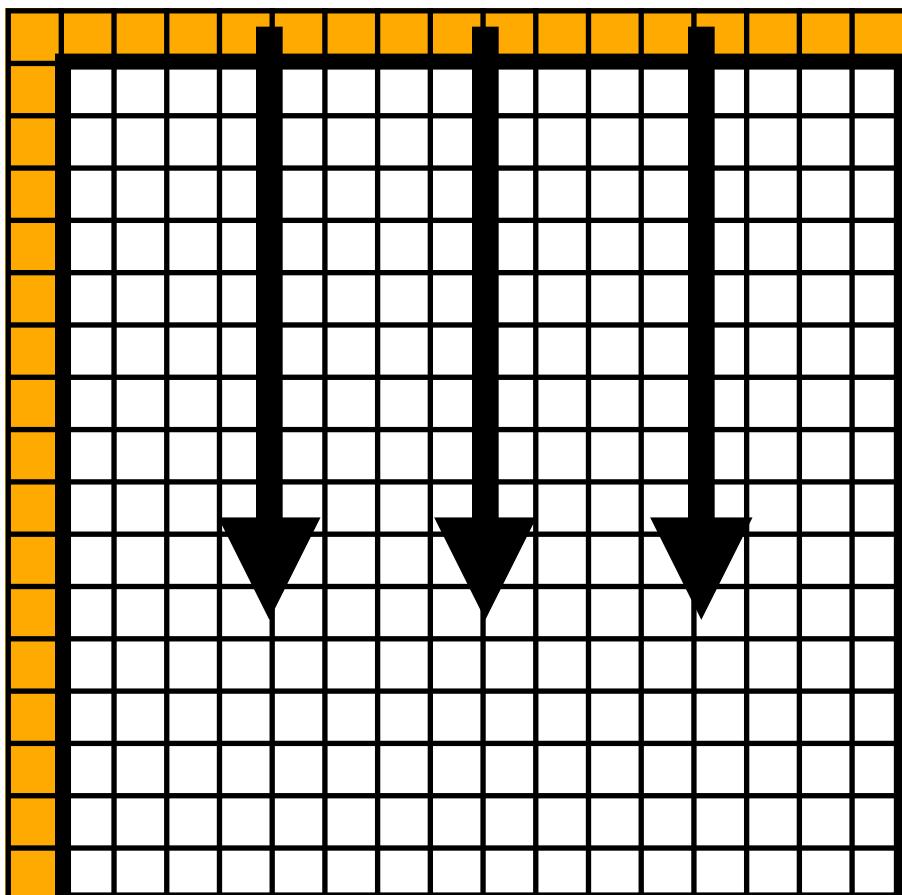
Mode 4: diagonal down-right (45°)



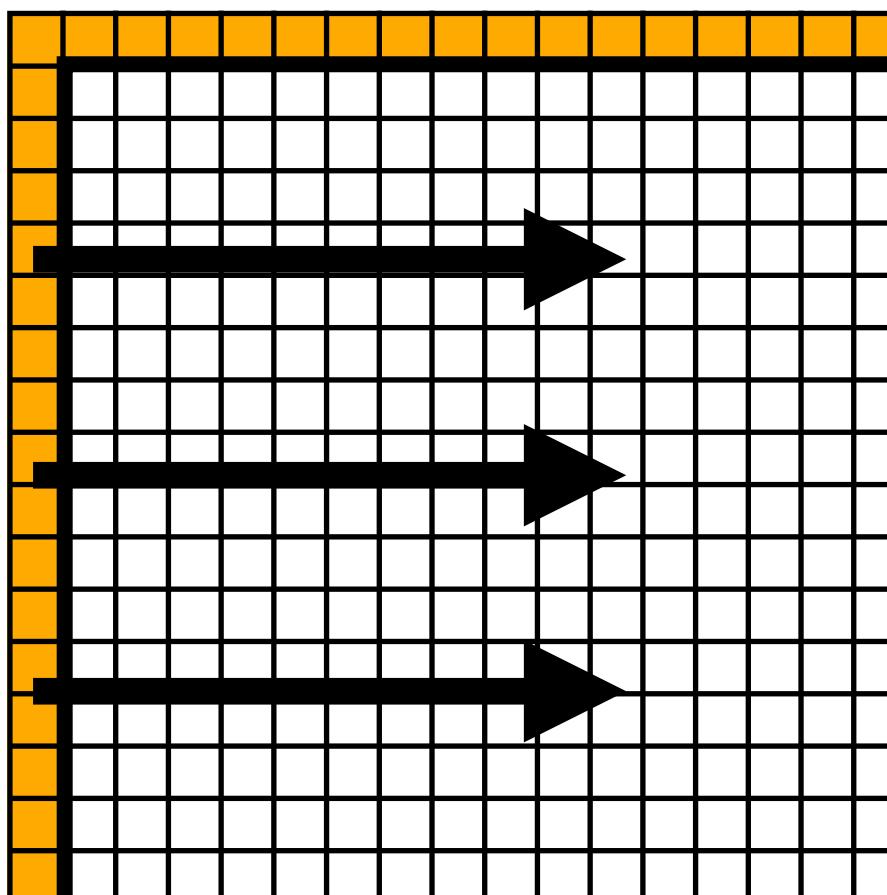
Mode 5: vertical-right (26.6°)

Intra_16x16 prediction modes

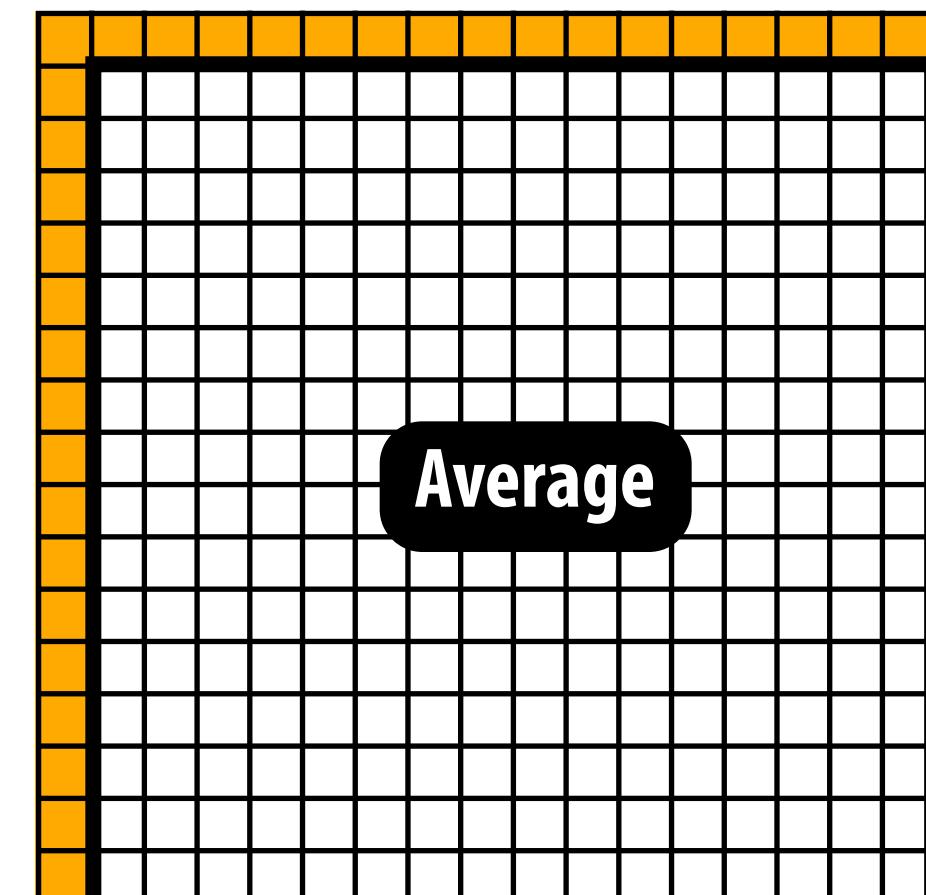
- 4 prediction modes: vertical, horizontal, DC, plane



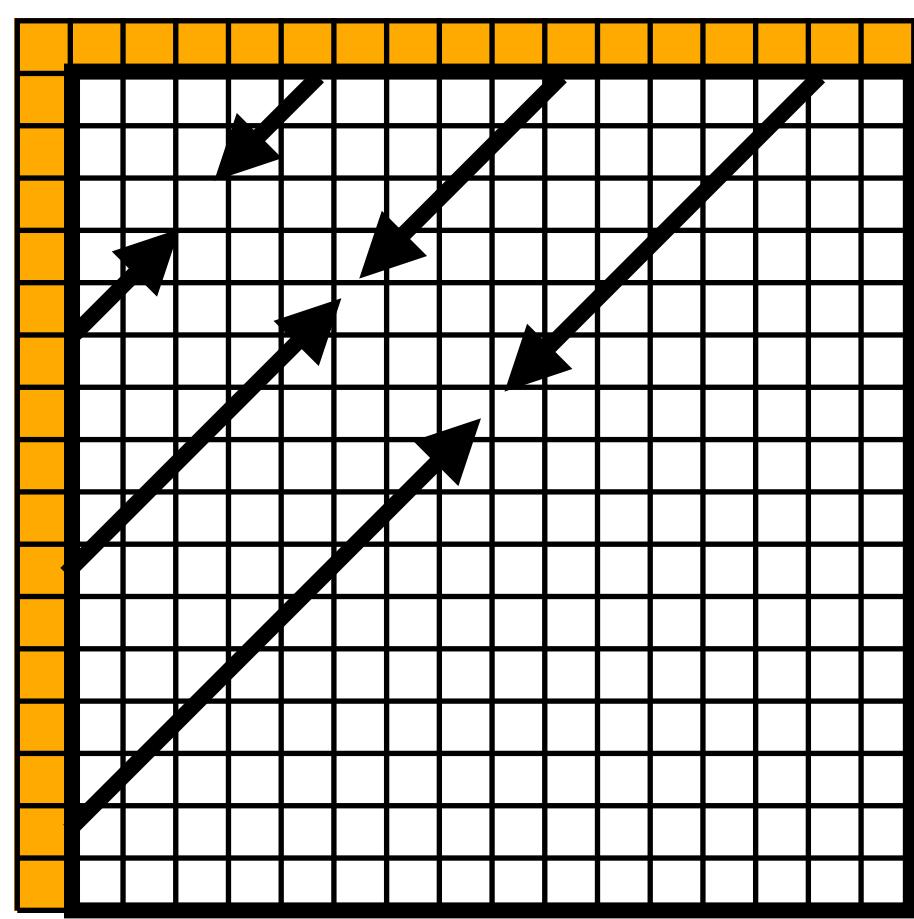
Mode 0: vertical



Mode 1: horizontal



Mode 2: DC



Mode 4: plane

$$P[i,j] = A_i * B_j + C$$

A derived from top row, B derived from left col, C from both

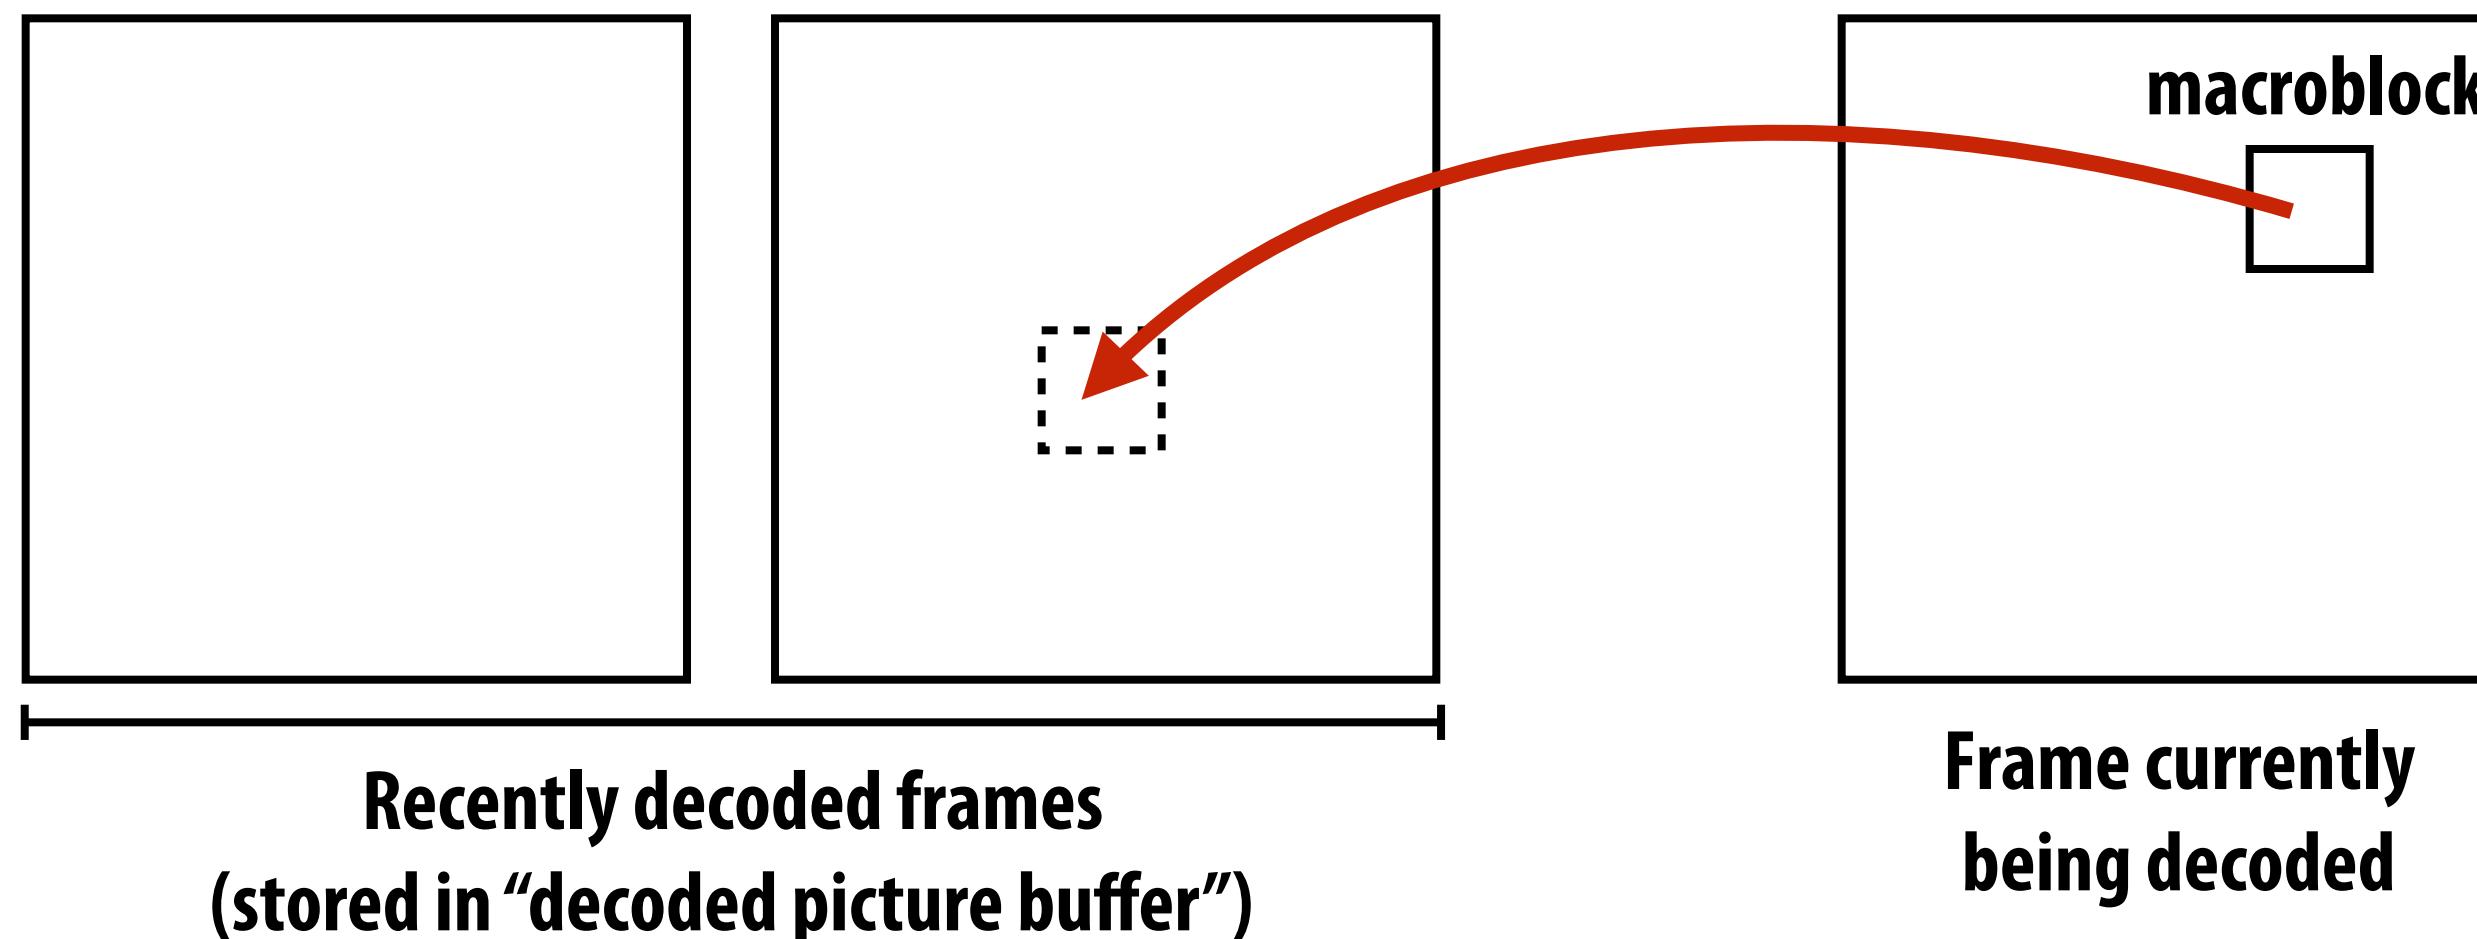
Further details

- Intra-prediction of chroma (8x8 block) is performed using four modes similar to those of intra_16x16 (except reordered as: DC, vertical, horizontal, plane)
- Intra-prediction scheme for each 4x4 block within macroblock encoded as follows:
 - One bit per 4x4 block:
 - if 1, use most probable mode
 - Most probable = lower of modes used for 4x4 block to left or above current block
 - if 0, use additional 3-bit value `rem_intra4x4_pred_mode` to encode one of nine modes
 - if `rem_intra4x4_pred_mode` is smaller than most probable mode, use mode given by `rem_intra4x4_pred_mode`
 - else, mode is `rem_intra4x4_pred_mode+1`

	mode=8
mode=2	mode=??

Inter-frame prediction (P-macroblock)

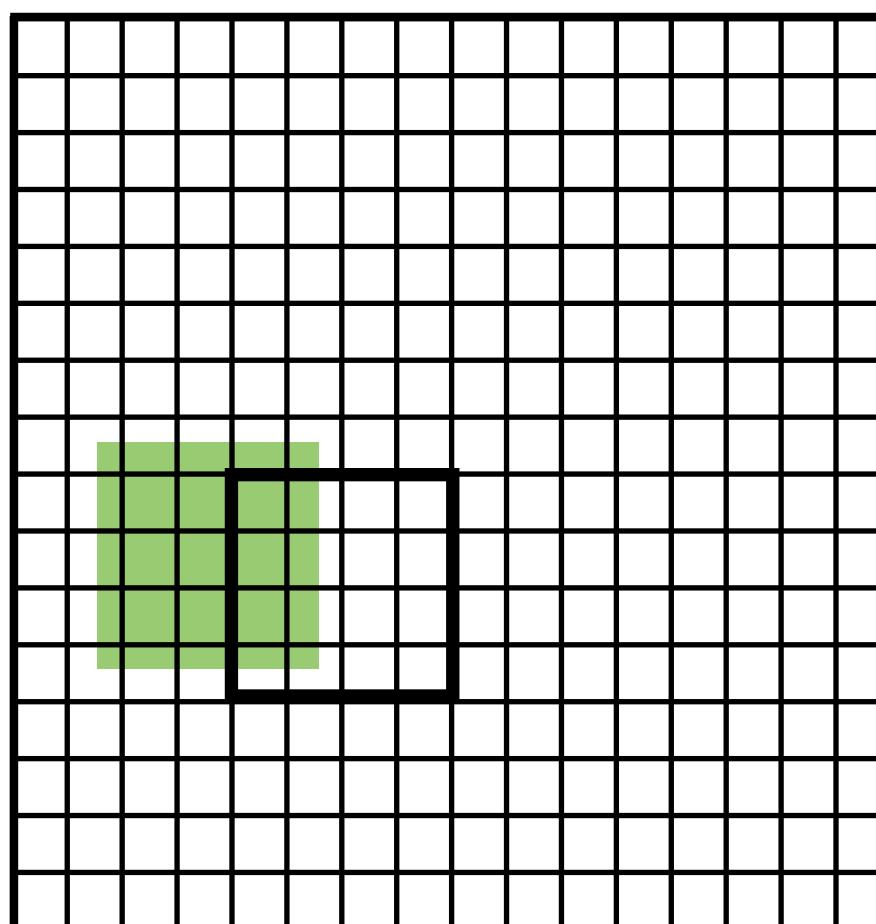
- Predict sample values using values from a block of a previously decoded frame *
- Basic idea: current frame formed by translation of pixels from temporally nearby frames (e.g., object moved slightly on screen between frames)
 - “Motion compensation”: use of spatial displacement to make prediction about pixel values



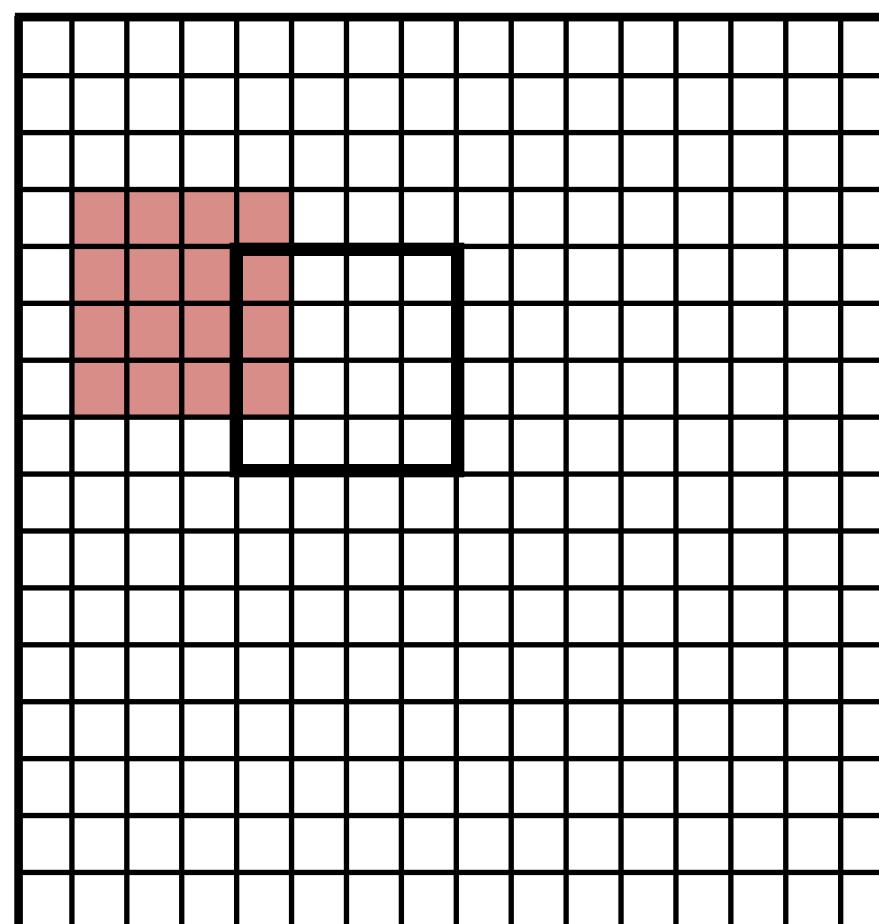
* Note: “previously decoded” does not imply source frame must come before current frame in the video sequence.
(H.264 supports decoding out of order.)

P-macroblock prediction

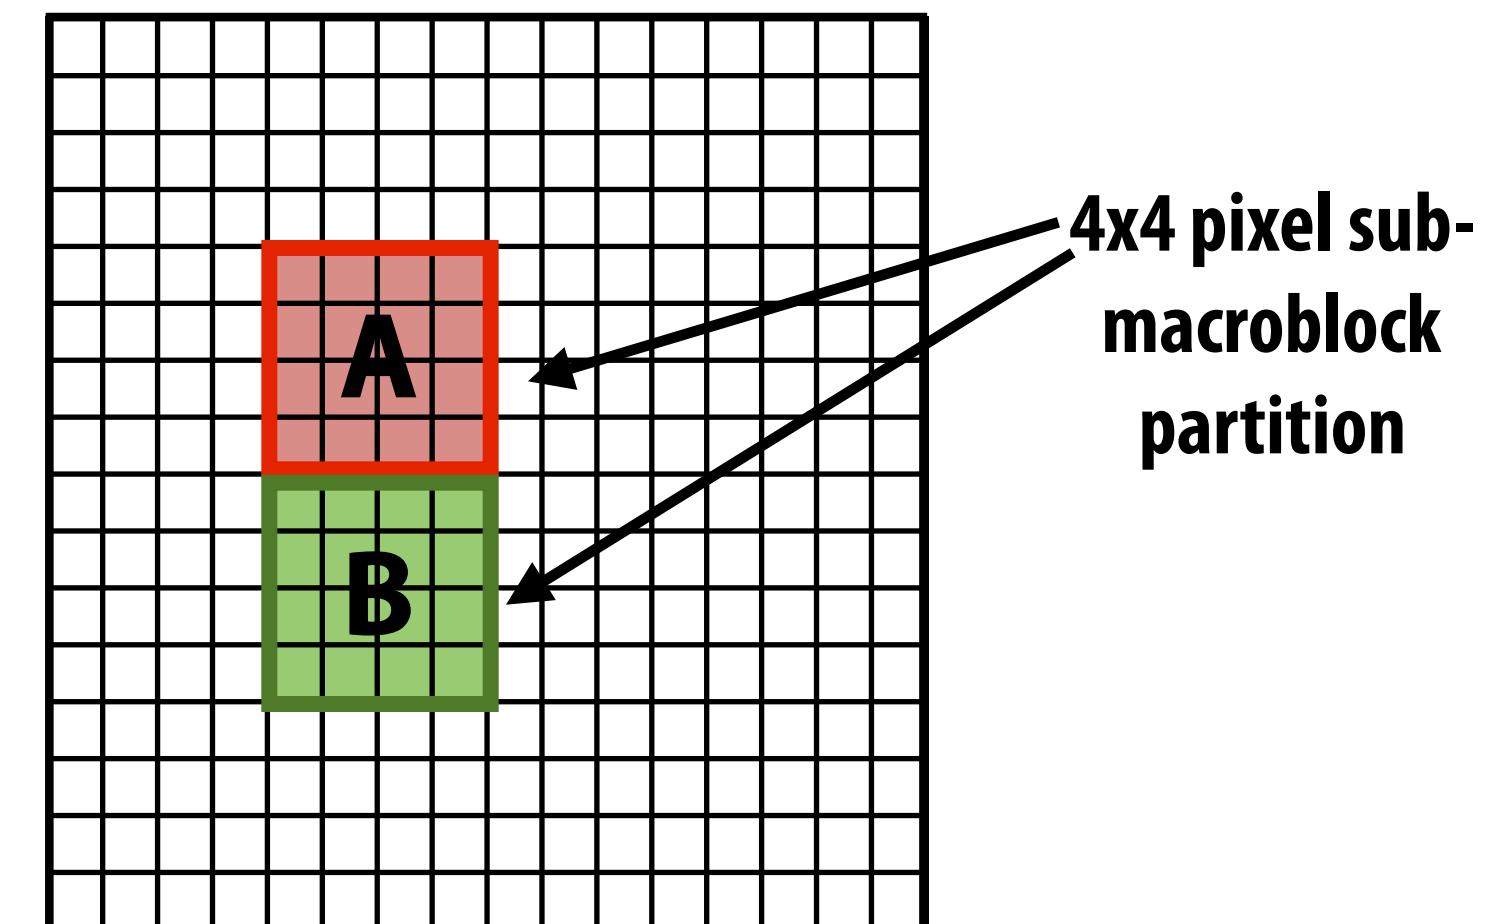
- Prediction can be performed at macroblock or sub-macroblock granularity
 - Macroblock can be divided into 16x16, 8x16, 16x8, 8x8 “partitions”
 - 8x8 partitions can be further subdivided into 4x8, 8x4, 4x4 sub-macroblock partitions
- Each partition predicted by sample values defined by:
(reference frame id, motion vector)



Decoded picture
buffer: frame 1



Decoded picture
buffer: frame 0



Current frame

Block A: predicted from (frame 0, motion-vector = [-3, -1])

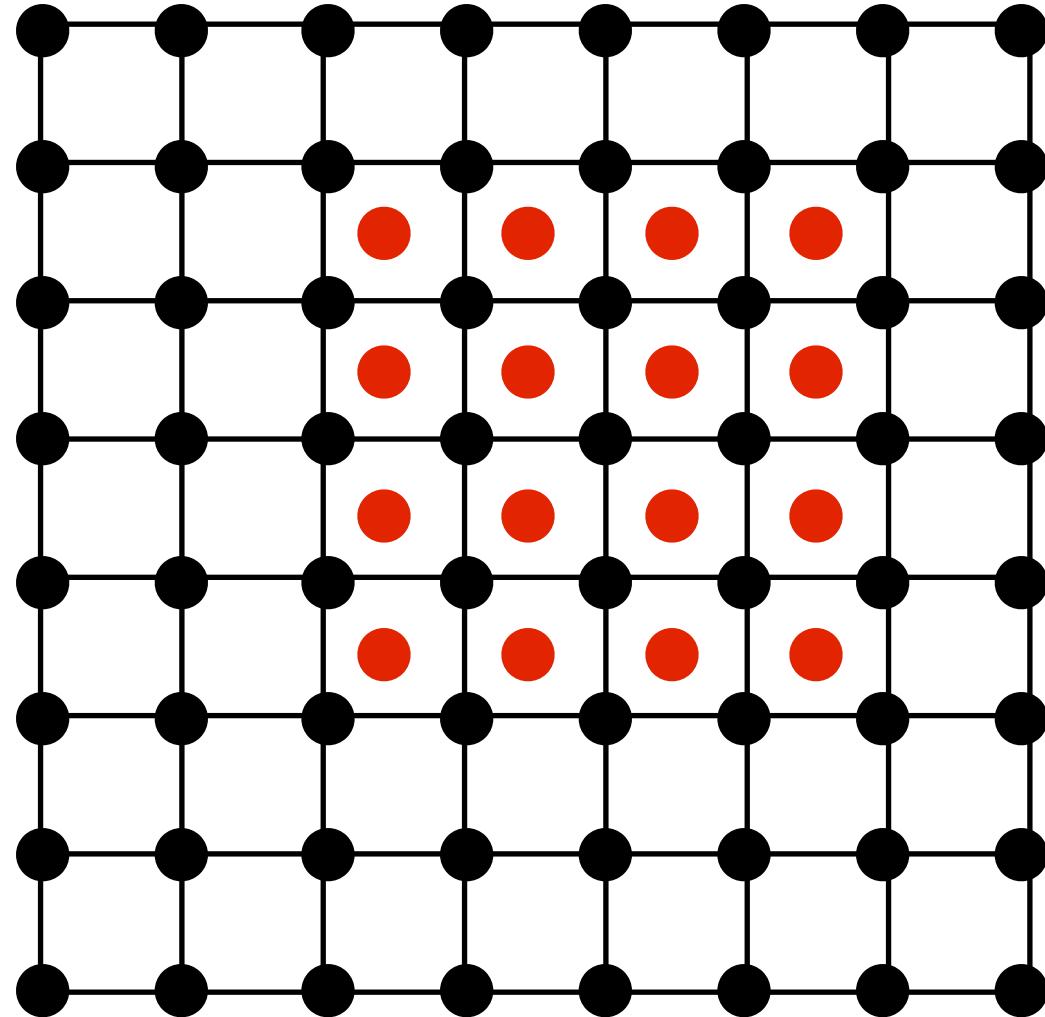
Block B: predicted from (frame 1, motion-vector = [-2.5, -0.5])

Note: non-integer motion vector

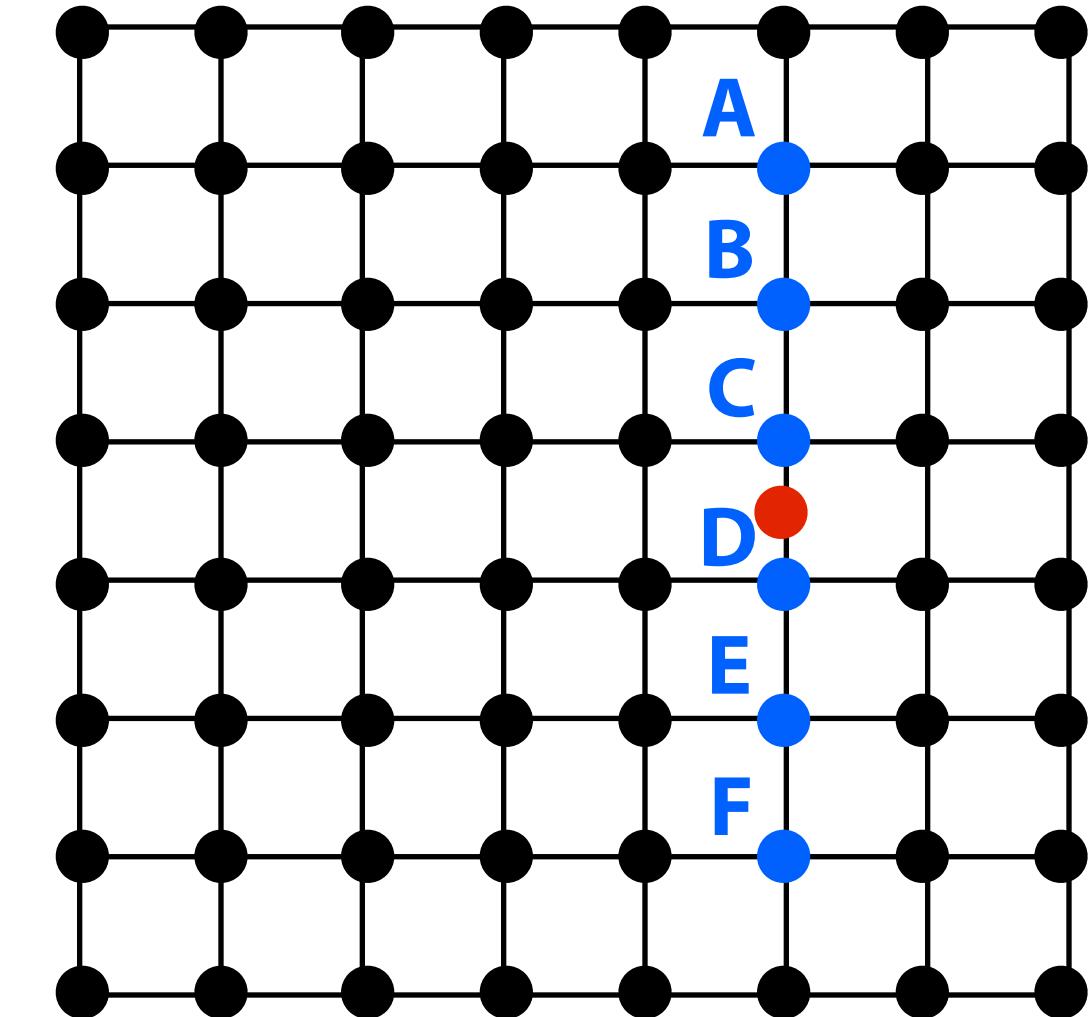
Motion vector visualization



Non-integer motion vectors require resampling



Example: motion vector with 1/2 pixel values.
Must resample reference block at positions given by red dots.



Interpolation to 1/2 pixel sample points via 6-tap filter:

`half_integer_value = clamp((A - 5B + 20C + 20D - 5E + F) / 32)`

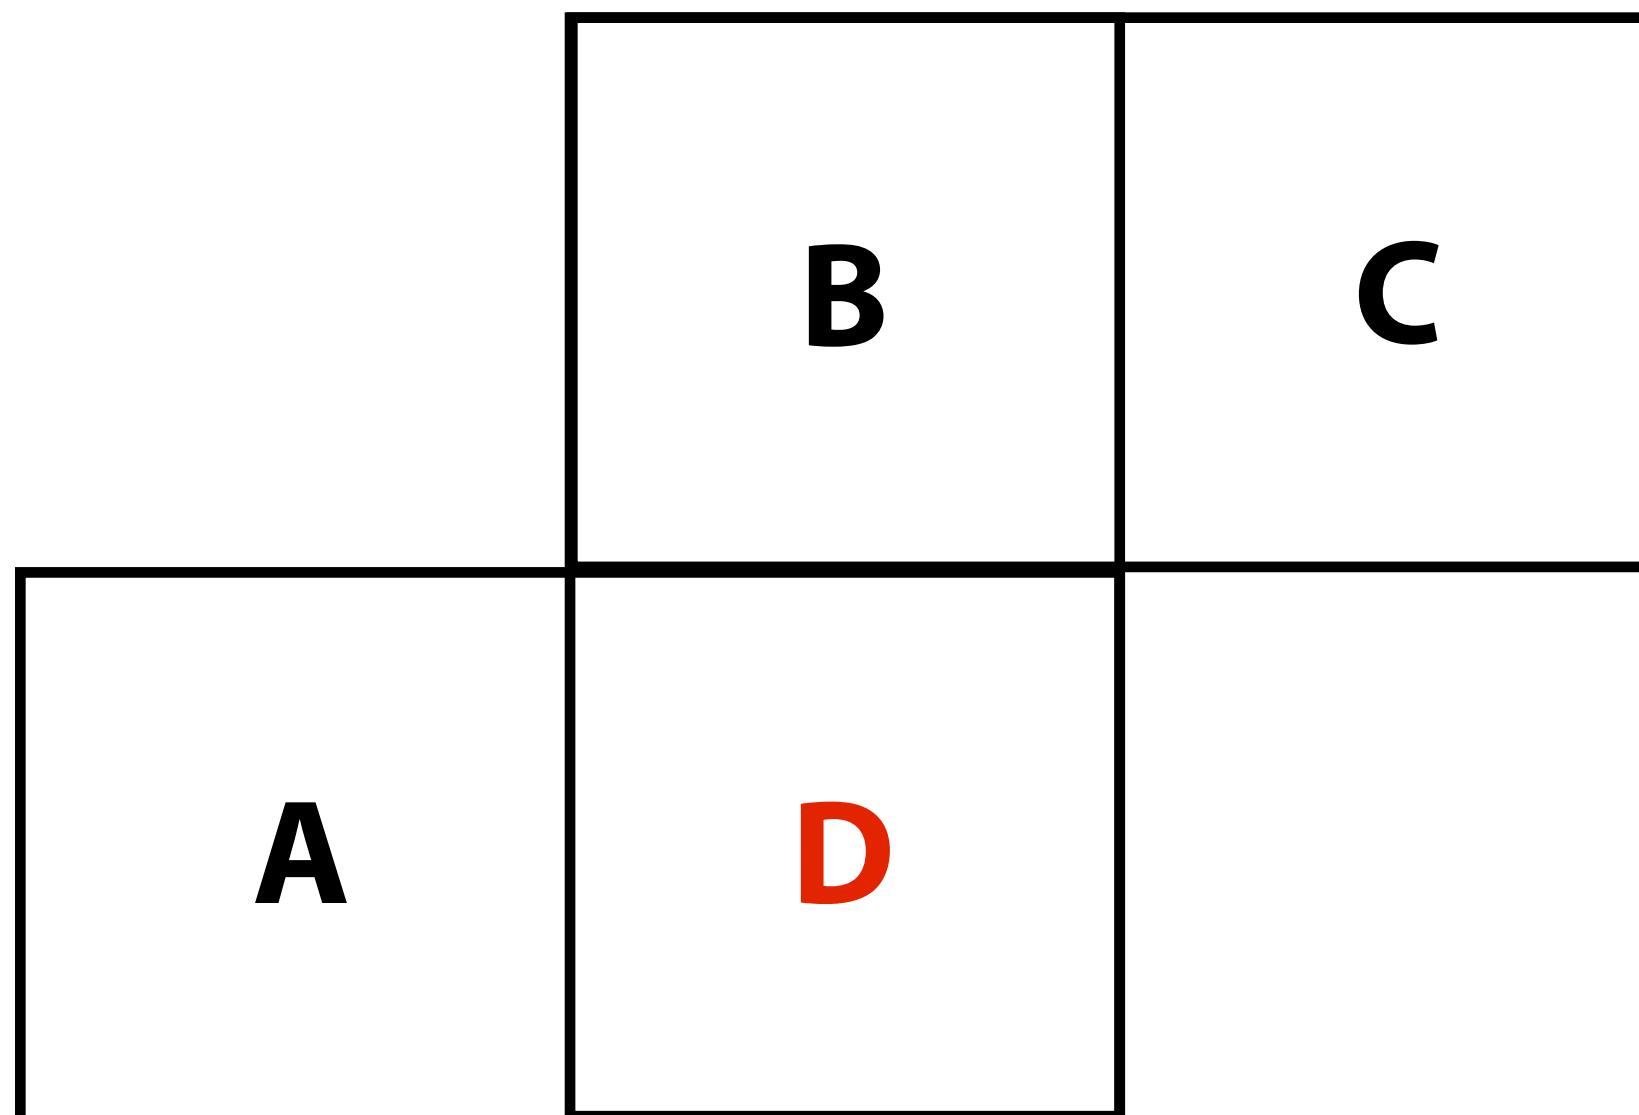
H.264 supports both 1/2 pixel and 1/4 pixel resolution motion vectors

1/4 resolution resampling performed by bilinear interpolation of 1/2 pixel samples

1/8 resolution (chroma only) by bilinear interpolation of 1/4 pixel samples

Motion vector prediction

- **Problem: per-partition motion vectors require significant amount of storage**
- **Solution: predict motion vectors from neighboring partitions and encode residual in compressed video stream**
 - Example below: predict D's motion vector as average of motion vectors of A, B, C
 - Prediction logic becomes more complex when partitions of neighboring blocks are of different size

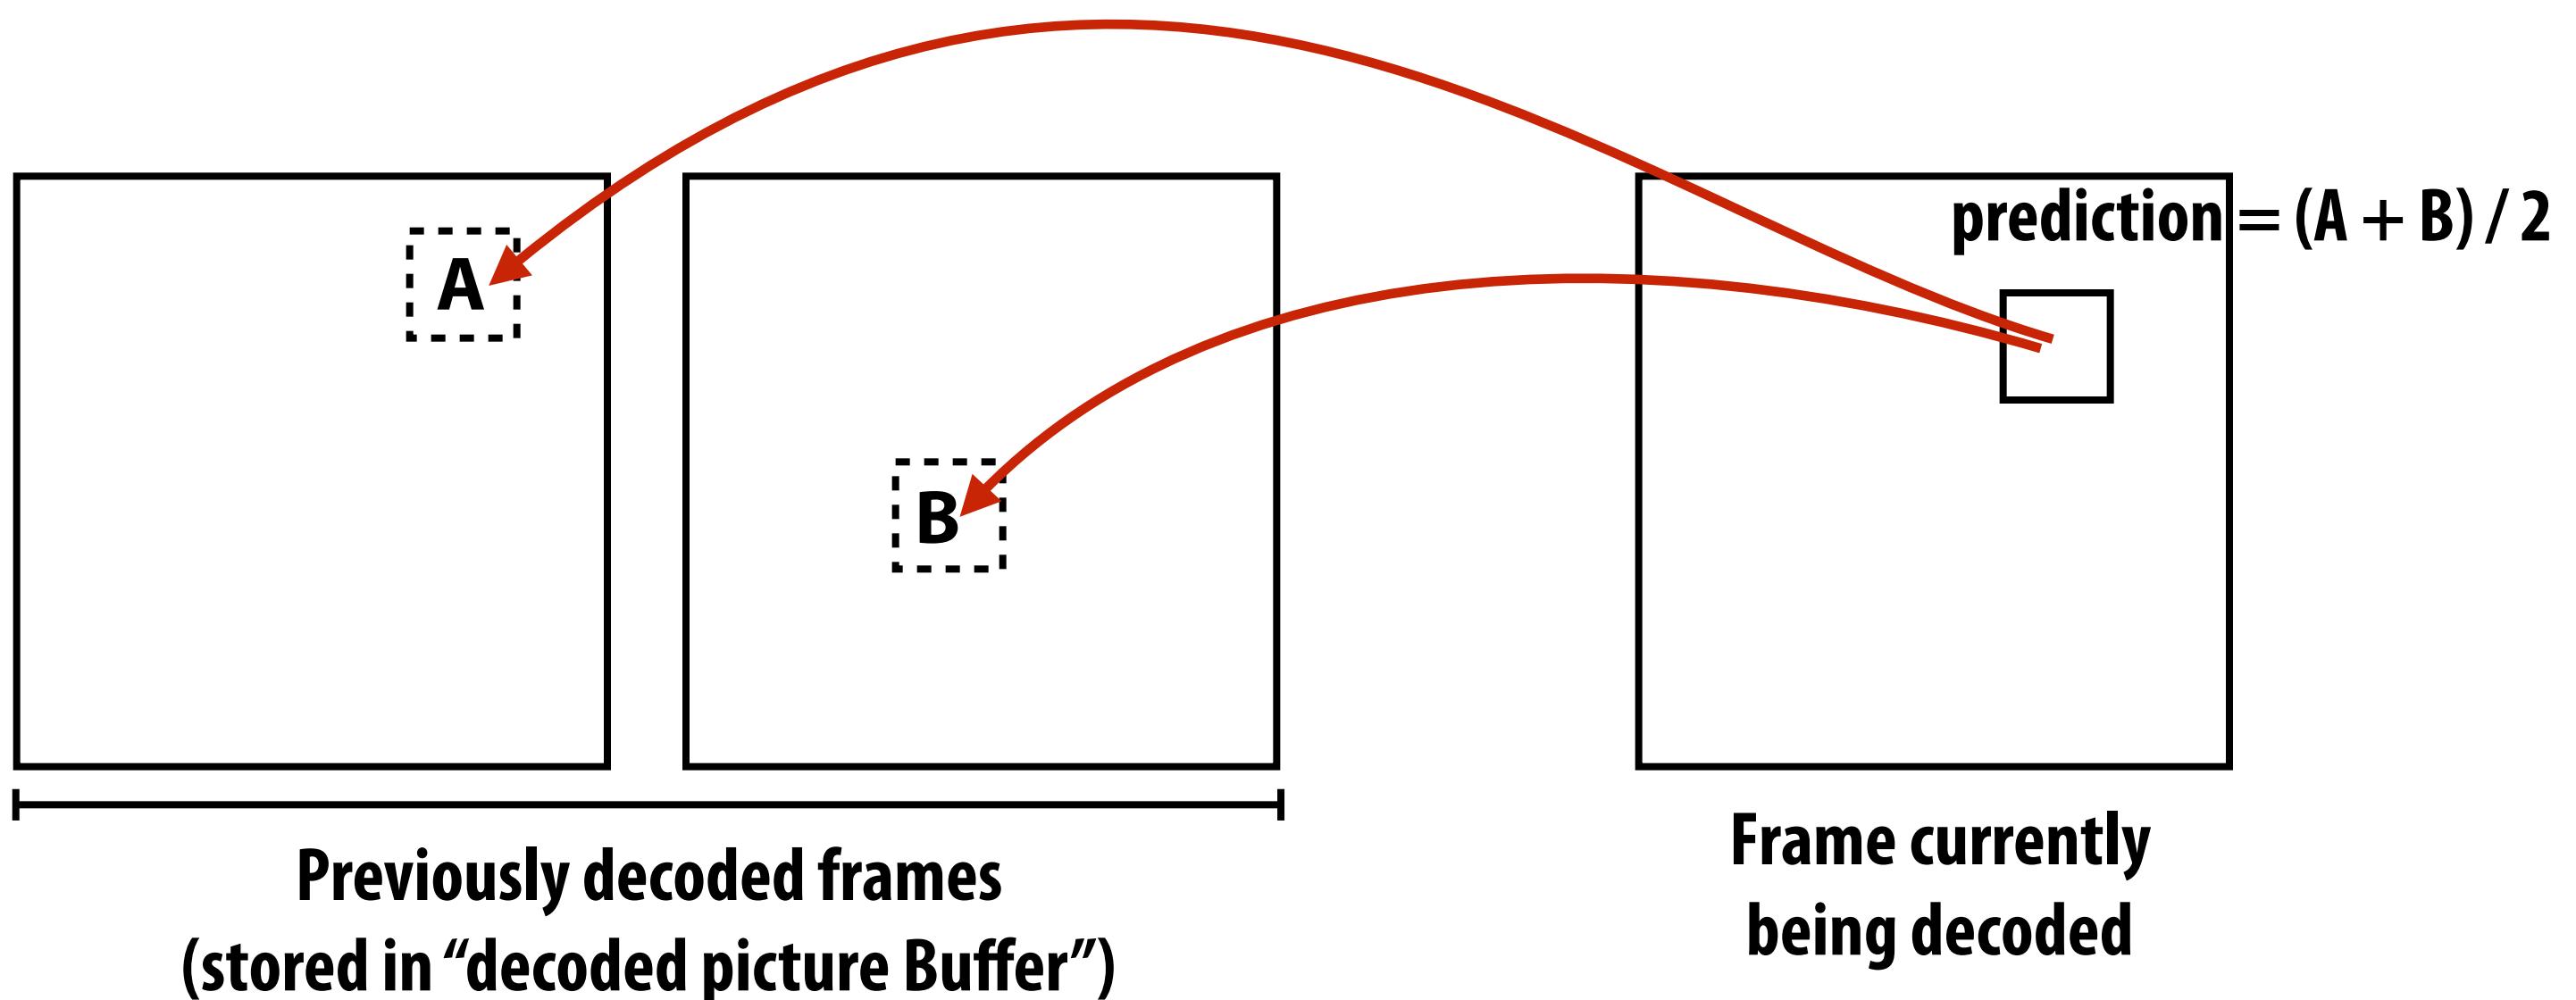


Question: what partition size is best?

- **Smaller partitions likely yield more accurate prediction**
 - Fewer bits needed for residuals
- **Smaller partitions require more bits to store partition information (diminish benefits of prediction)**
 - Reference picture id
 - Motion vectors (note: motion vectors are more coherent with finer sampling, so they likely compress well)

Inter-frame prediction (B-macroblock)

- Each partition predicted by up to two source blocks
 - Prediction is the average of the two reference blocks
 - Each B-macroblock partition stores two frame references and two motion vectors (recall P-macroblock partitions only stored one)



Additional prediction details

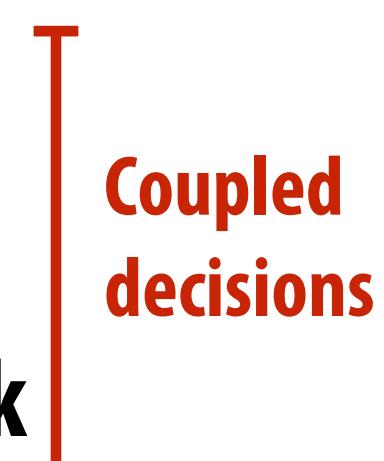
■ Optional weighting to prediction:

- Per-slice explicit weighting (reference samples multiplied by weight)
- Per-B-slice implicit weights (reference samples weights by temporal distance of reference frame from current frame in video)
 - Idea: weight samples from reference frames nearby in time more

■ Deblocking

- Blocking artifacts may result as a result of macroblock granularity encoding
- After macroblock decoding is complete, optionally perform smoothing filter across block edges.

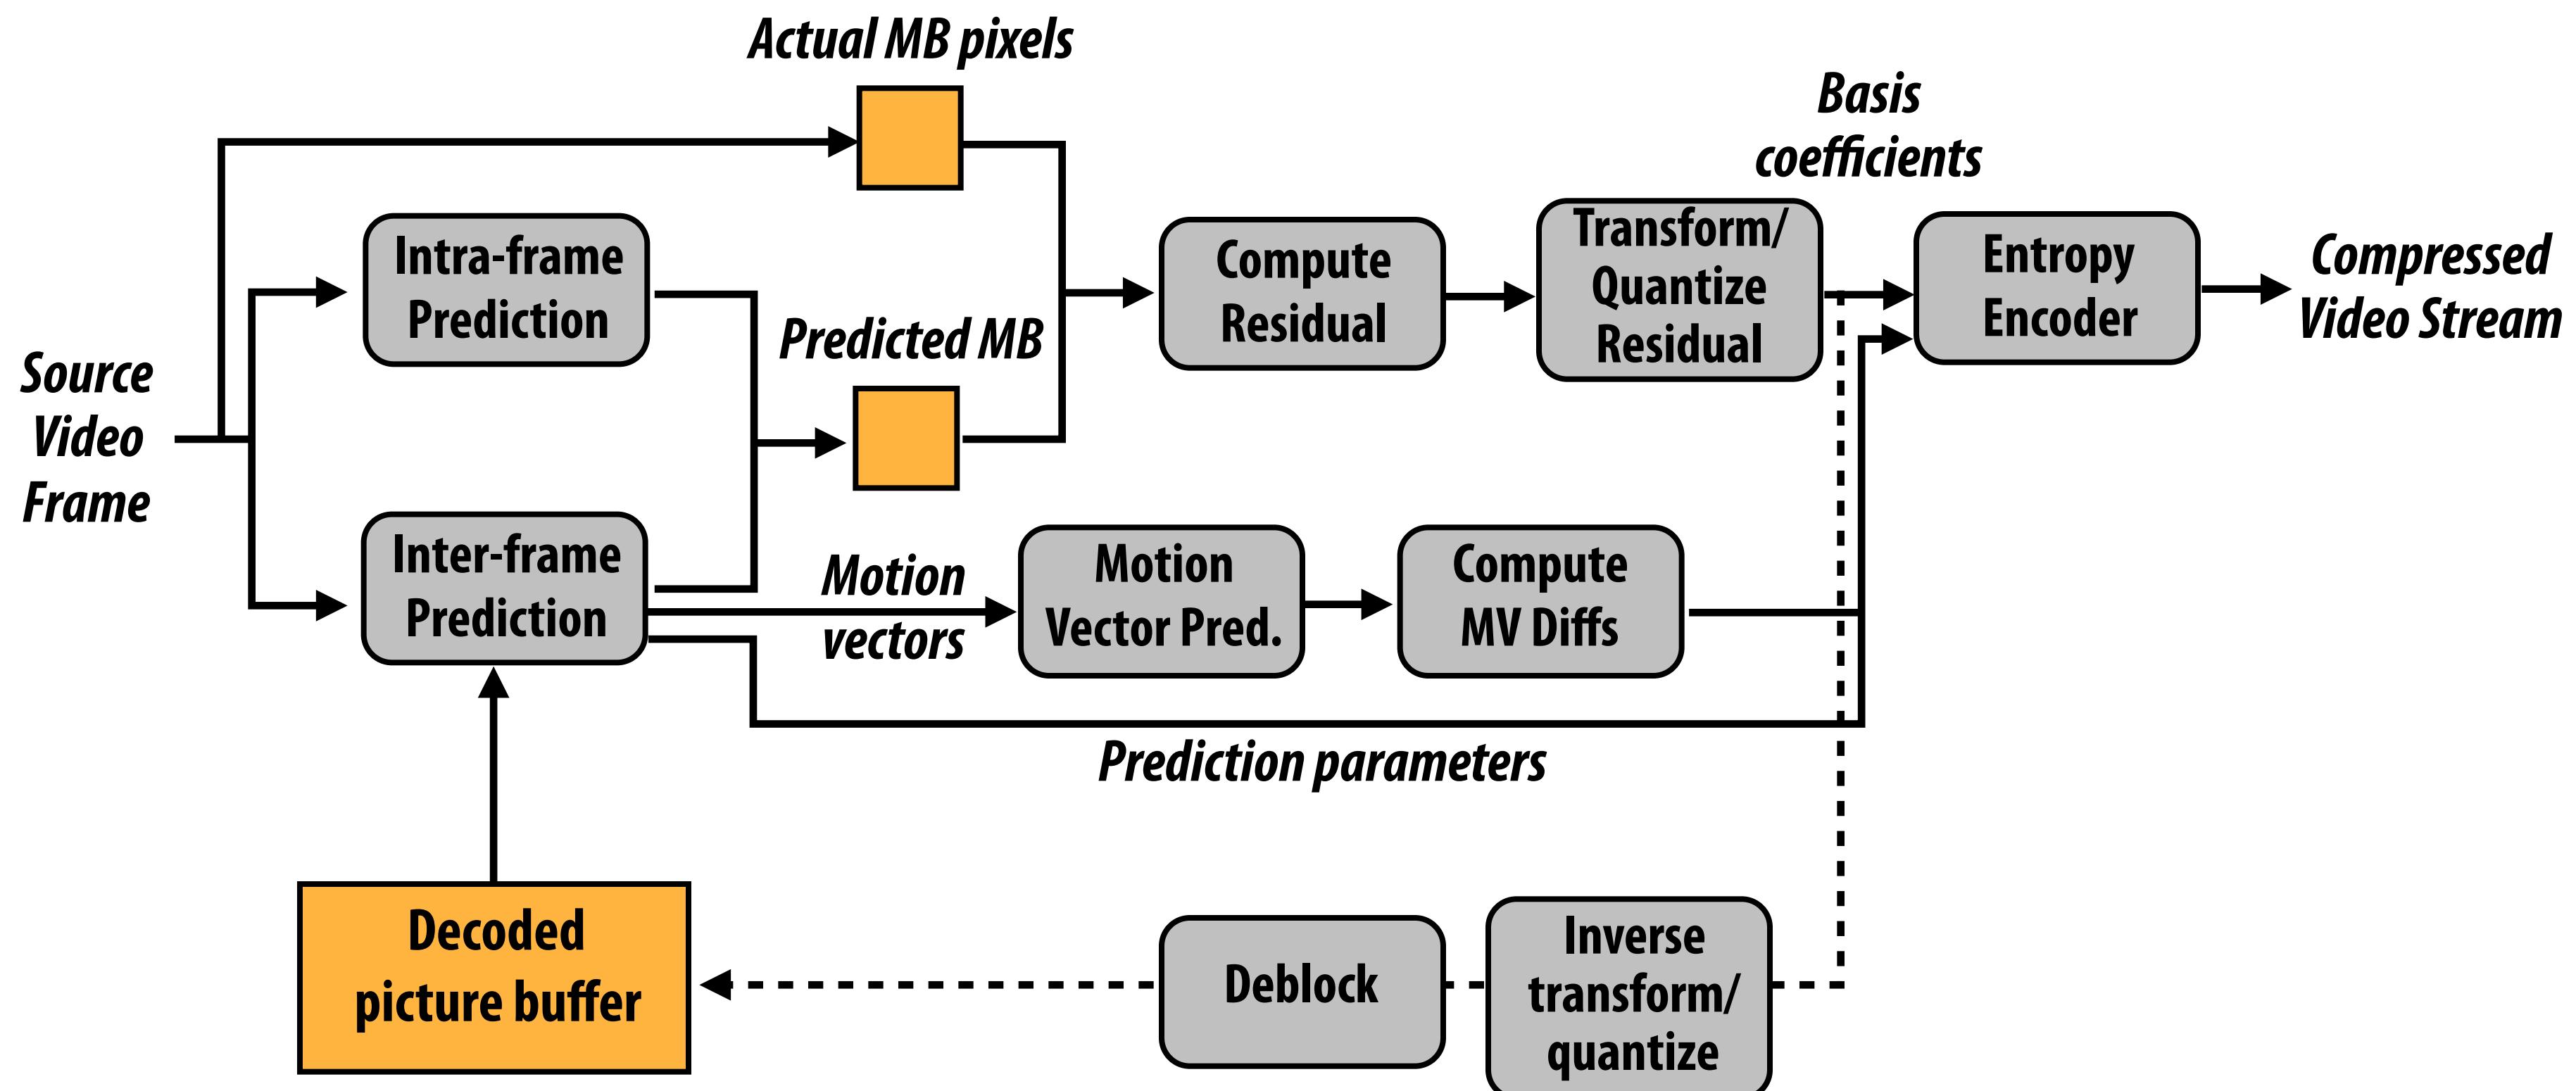
Putting it all together: encoding an inter-predicted macroblock

- **Inputs:**
 - Current state of decoded picture buffer (state of the decoder)
 - 16x16 block of input video to encode
 - **General steps: (need not be performed in this order)**
 - Resample images in decoded picture buffer to obtain 1/2, and 1/4, 1/8 pixel resampling
 - Choose prediction type (P-type or B-type)
 - Choose reference pictures for prediction
 - Choose motion vectors for each partition (or sub-partition) of macroblock
 - Predict motion vectors and compute motion vector difference
 - Encode choice of prediction type, reference pictures, and motion vector differences
 - Encode residual for macroblock prediction
 - Store reconstructed macroblock (post deblocking) in decoded picture buffer to use as reference picture for future macroblocks
- 
- Coupled decisions

H.264/AVC video encoding

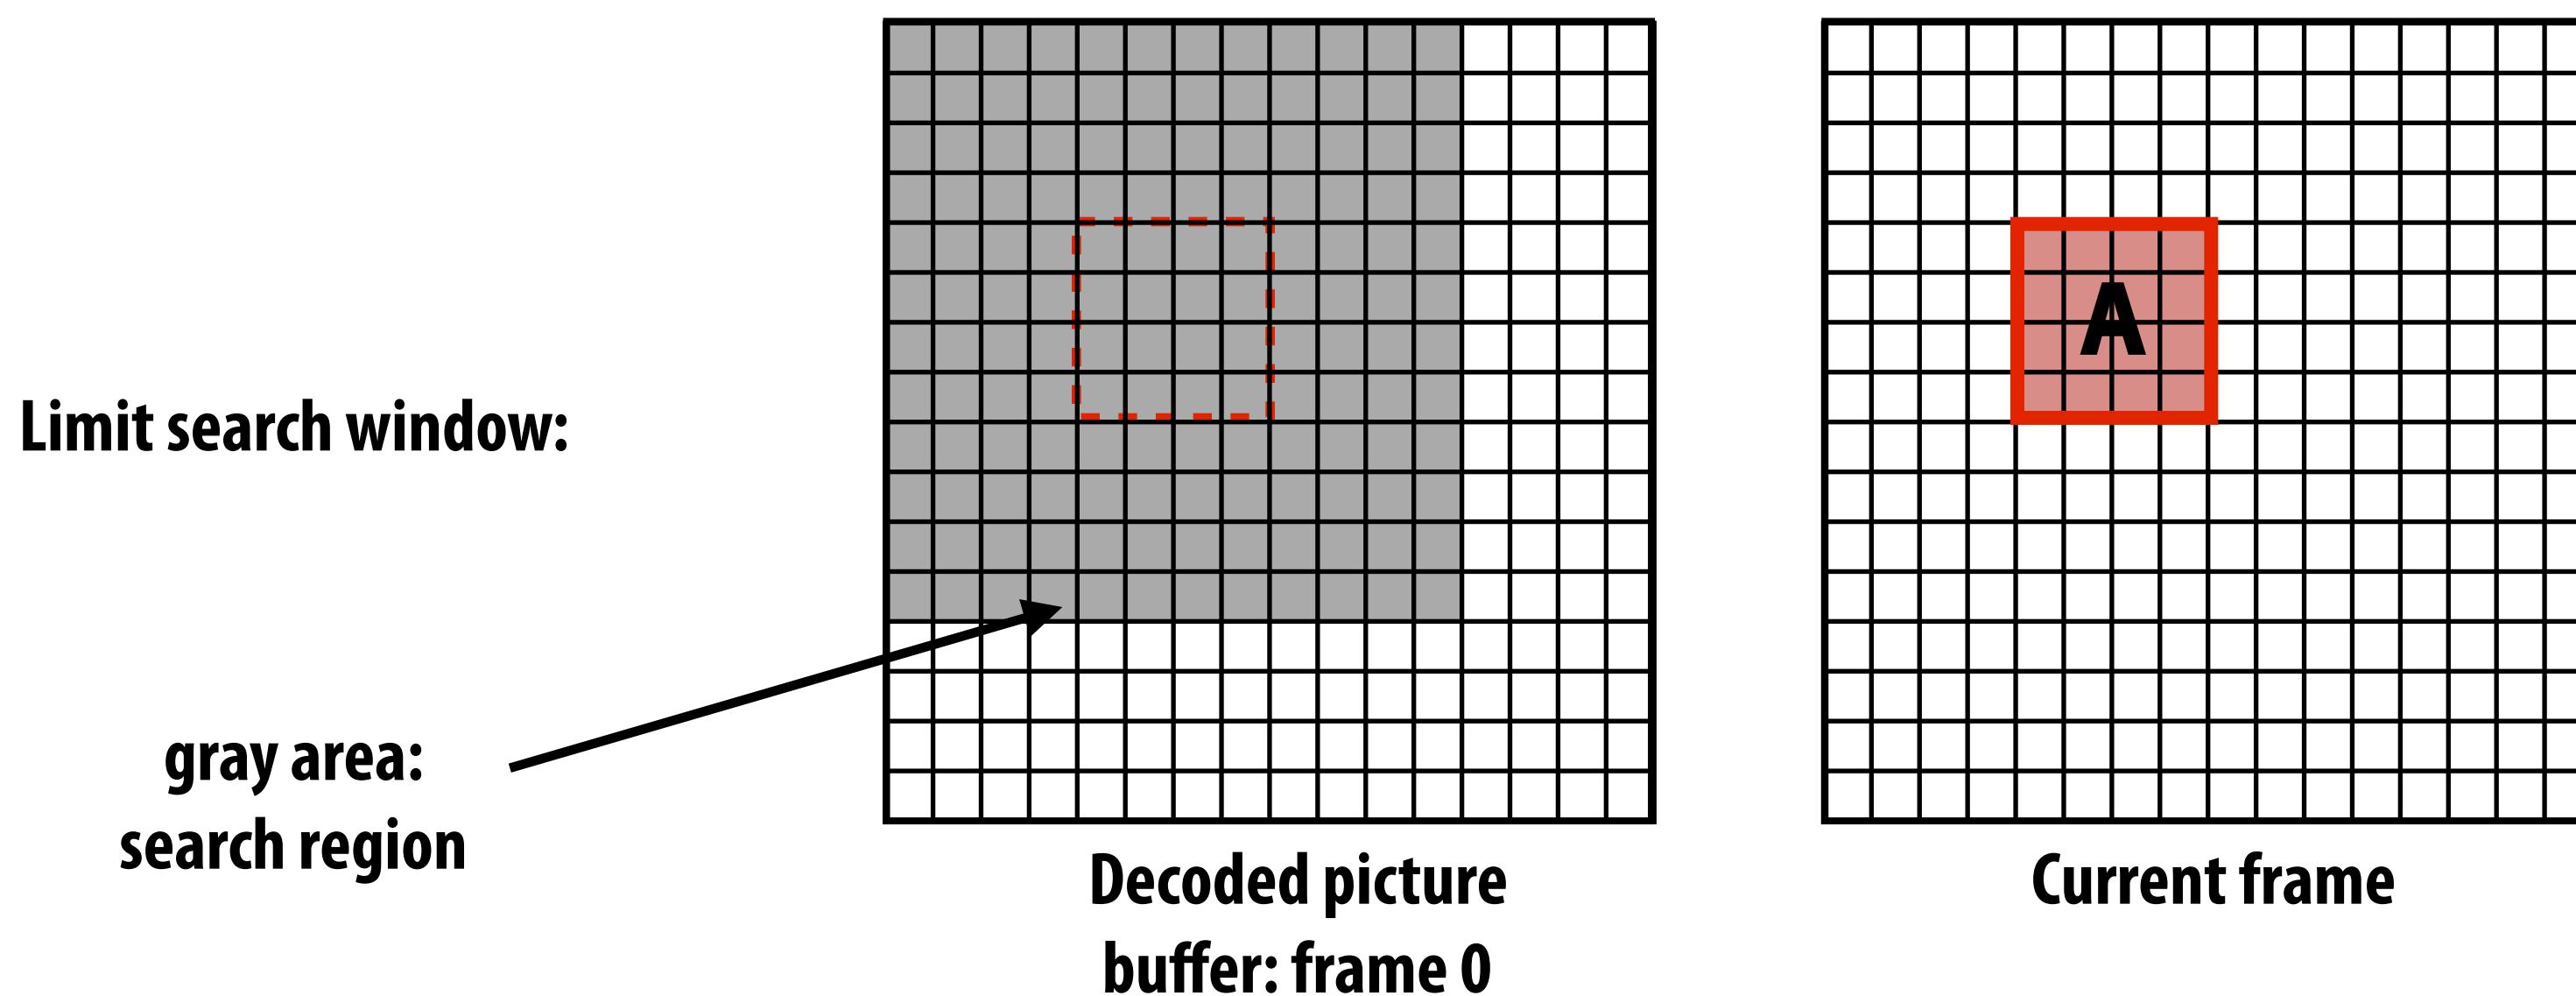
MB = macroblock

MV = motion vector



Motion estimation

- Encoder must find reference block that predicts current frame's pixels well.
 - Can search over multiple pictures in decoded picture buffer + motion vectors can be non-integer (huge search space)
 - Must also choose block size (macroblock partition size)
 - And whether to predict using combination of two blocks
 - Literature is full of heuristics to accelerate this process
 - Remember, must execute motion estimation in real-time for HD video (1920x1080), on a low-power smartphone



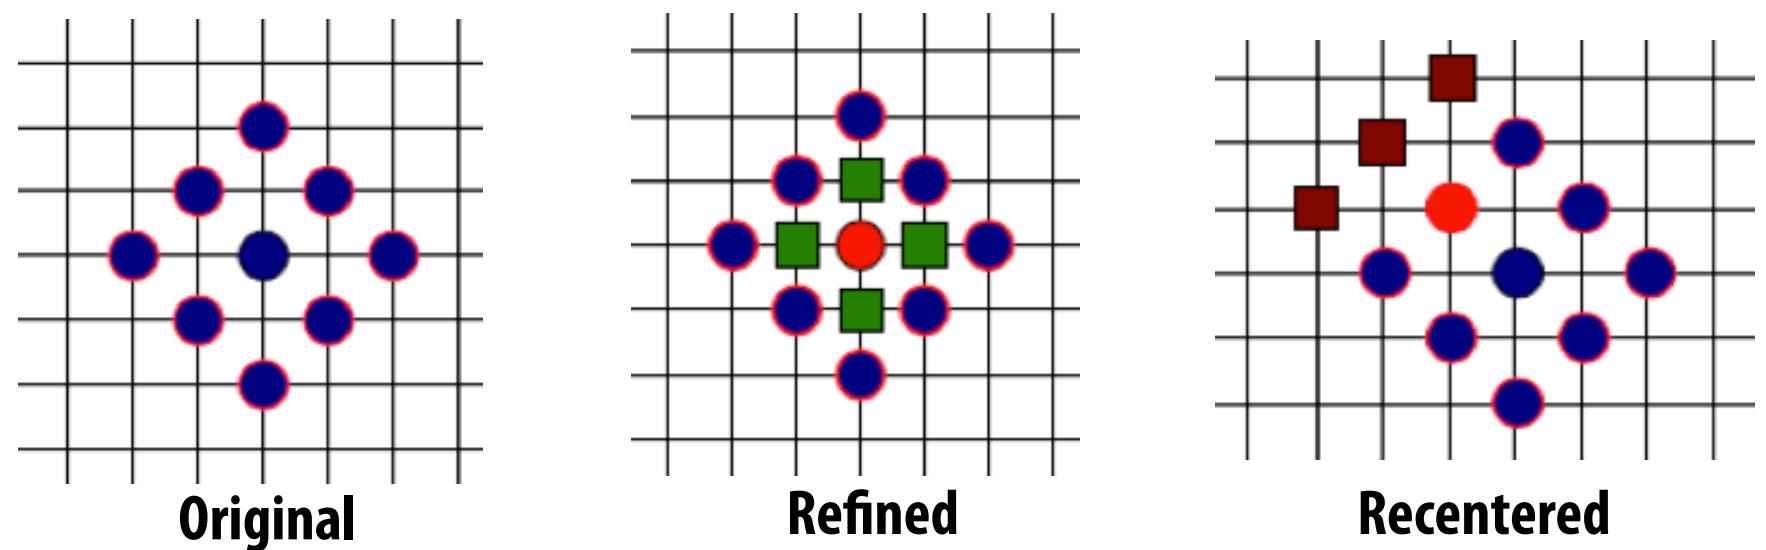
Motion estimation optimizations

■ Coarser search:

- Limit search window to small region
- First compute block differences at coarse scale (save partial sums from previous searches)

■ Smarter search:

- Guess motion vectors similar to motion vectors used for neighboring blocks
- Diamond search: start by test large diamond pattern centered around block
 - If best match is interior, refine to finer scale
 - Else, recenter around best match



■ Early termination: don't find optimal reference patch, just find one that's "good enough": e.g., compressed representation is lower than threshold

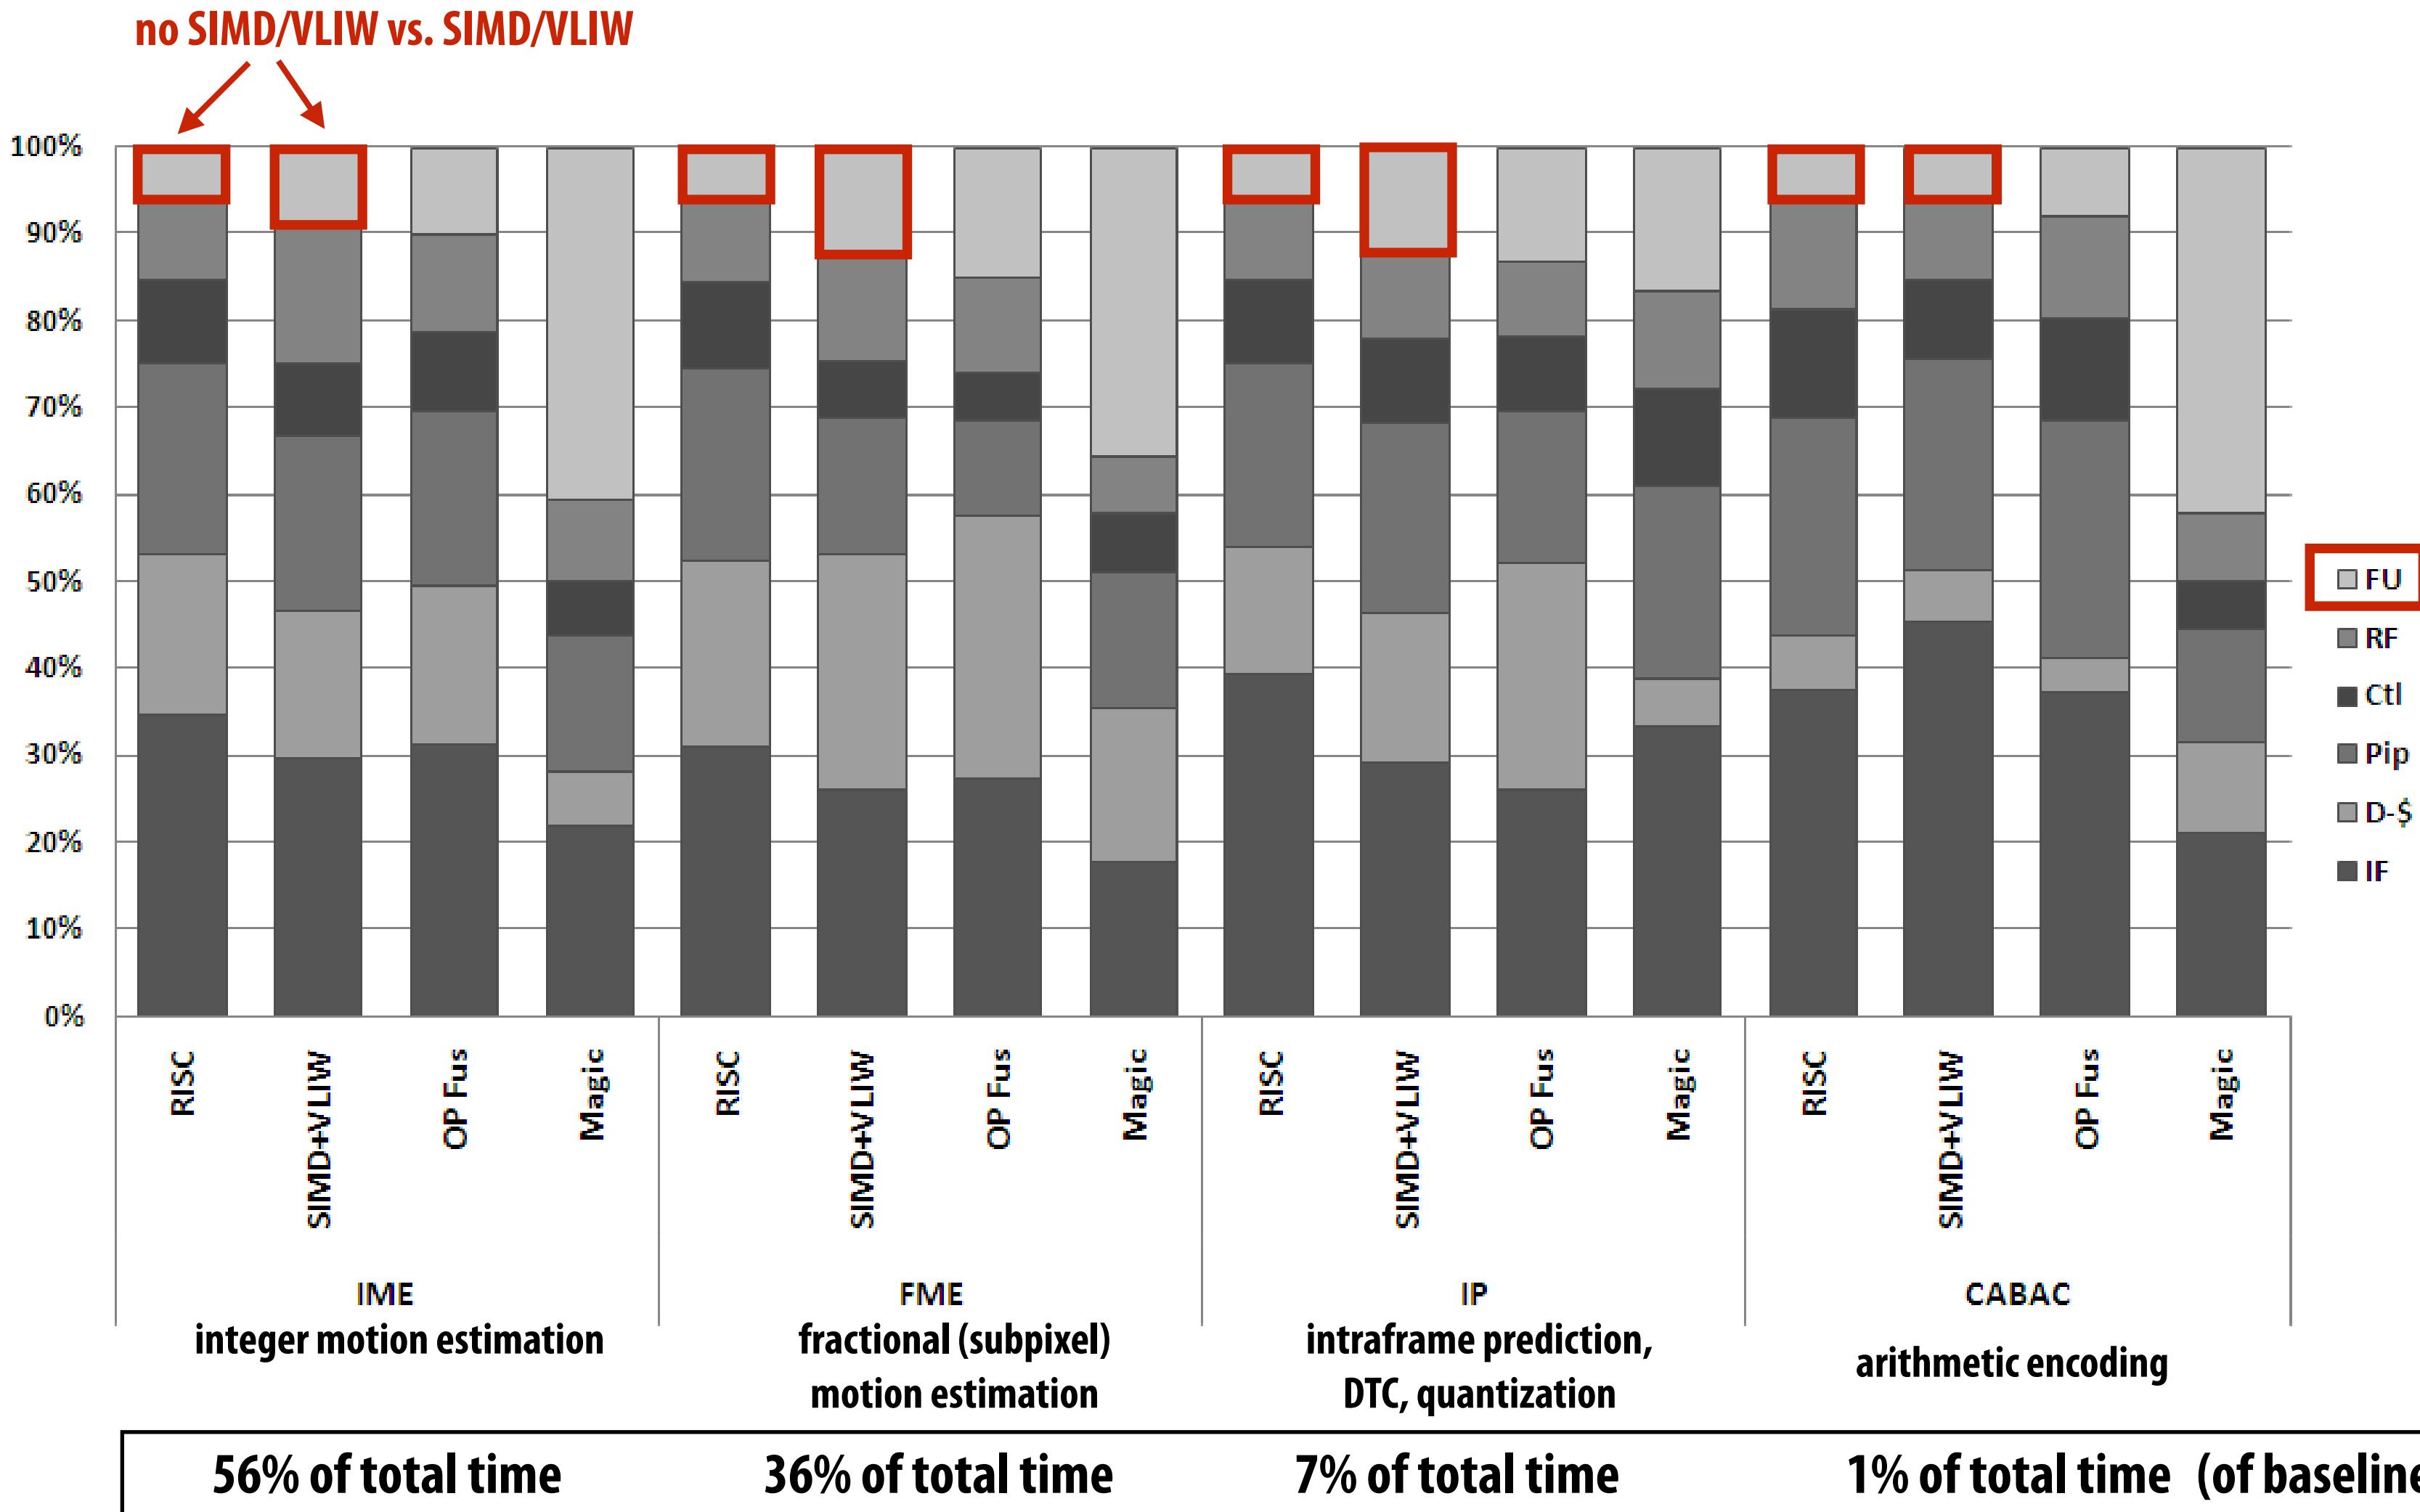
- Test zero-motion vector first (optimize for non-moving background)

■ Optimizations for subpixel motion vectors:

- Refinement: find best reference block given only pixel offsets, then try 1/2, 1/4-subpixel offsets around this match

Fraction of energy consumed by different parts of instruction pipeline (H.264 video encoding)

[Hameed et al. ISCA 2010]



FU = functional units

RF = register fetch

Ctrl = misc pipeline control

Pip = pipeline registers (interstage)

D-\$ = data cache

IF = instruction fetch + instruction cache

H.265 (HEVC)

- Standard ratified in 2013
- Goal: ~2X better compression than H.264
- Main ideas:
 - Macroblock sizes up to 64x64
 - Prediction block size and residual block sizes can be different
 - 35 intra-frame prediction modes (recall H.264 had 9)
 - ...

Summary

- **JPG image compression and H.264 video compression are “lossy” compression techniques that discard information is that less likely to be noticed by the human eye**
 - **Key principle: “Lossy, but still looks good enough to humans!”**
- **But moving forward most videos in the world will be analyzed by computers, not viewed by humans**
 - **What is appropriate form of compression for a particular set of computer vision tasks?**
 - **New principle: “Lossy, but image analysis tasks still work!”**
 - **Can we “learn” domain-specific compressors for particular scenes, types of tasks, etc?**