

Lecture 9:

Parallel Deep Network Training

**Visual Computing Systems
Stanford CS348V, Winter 2018**

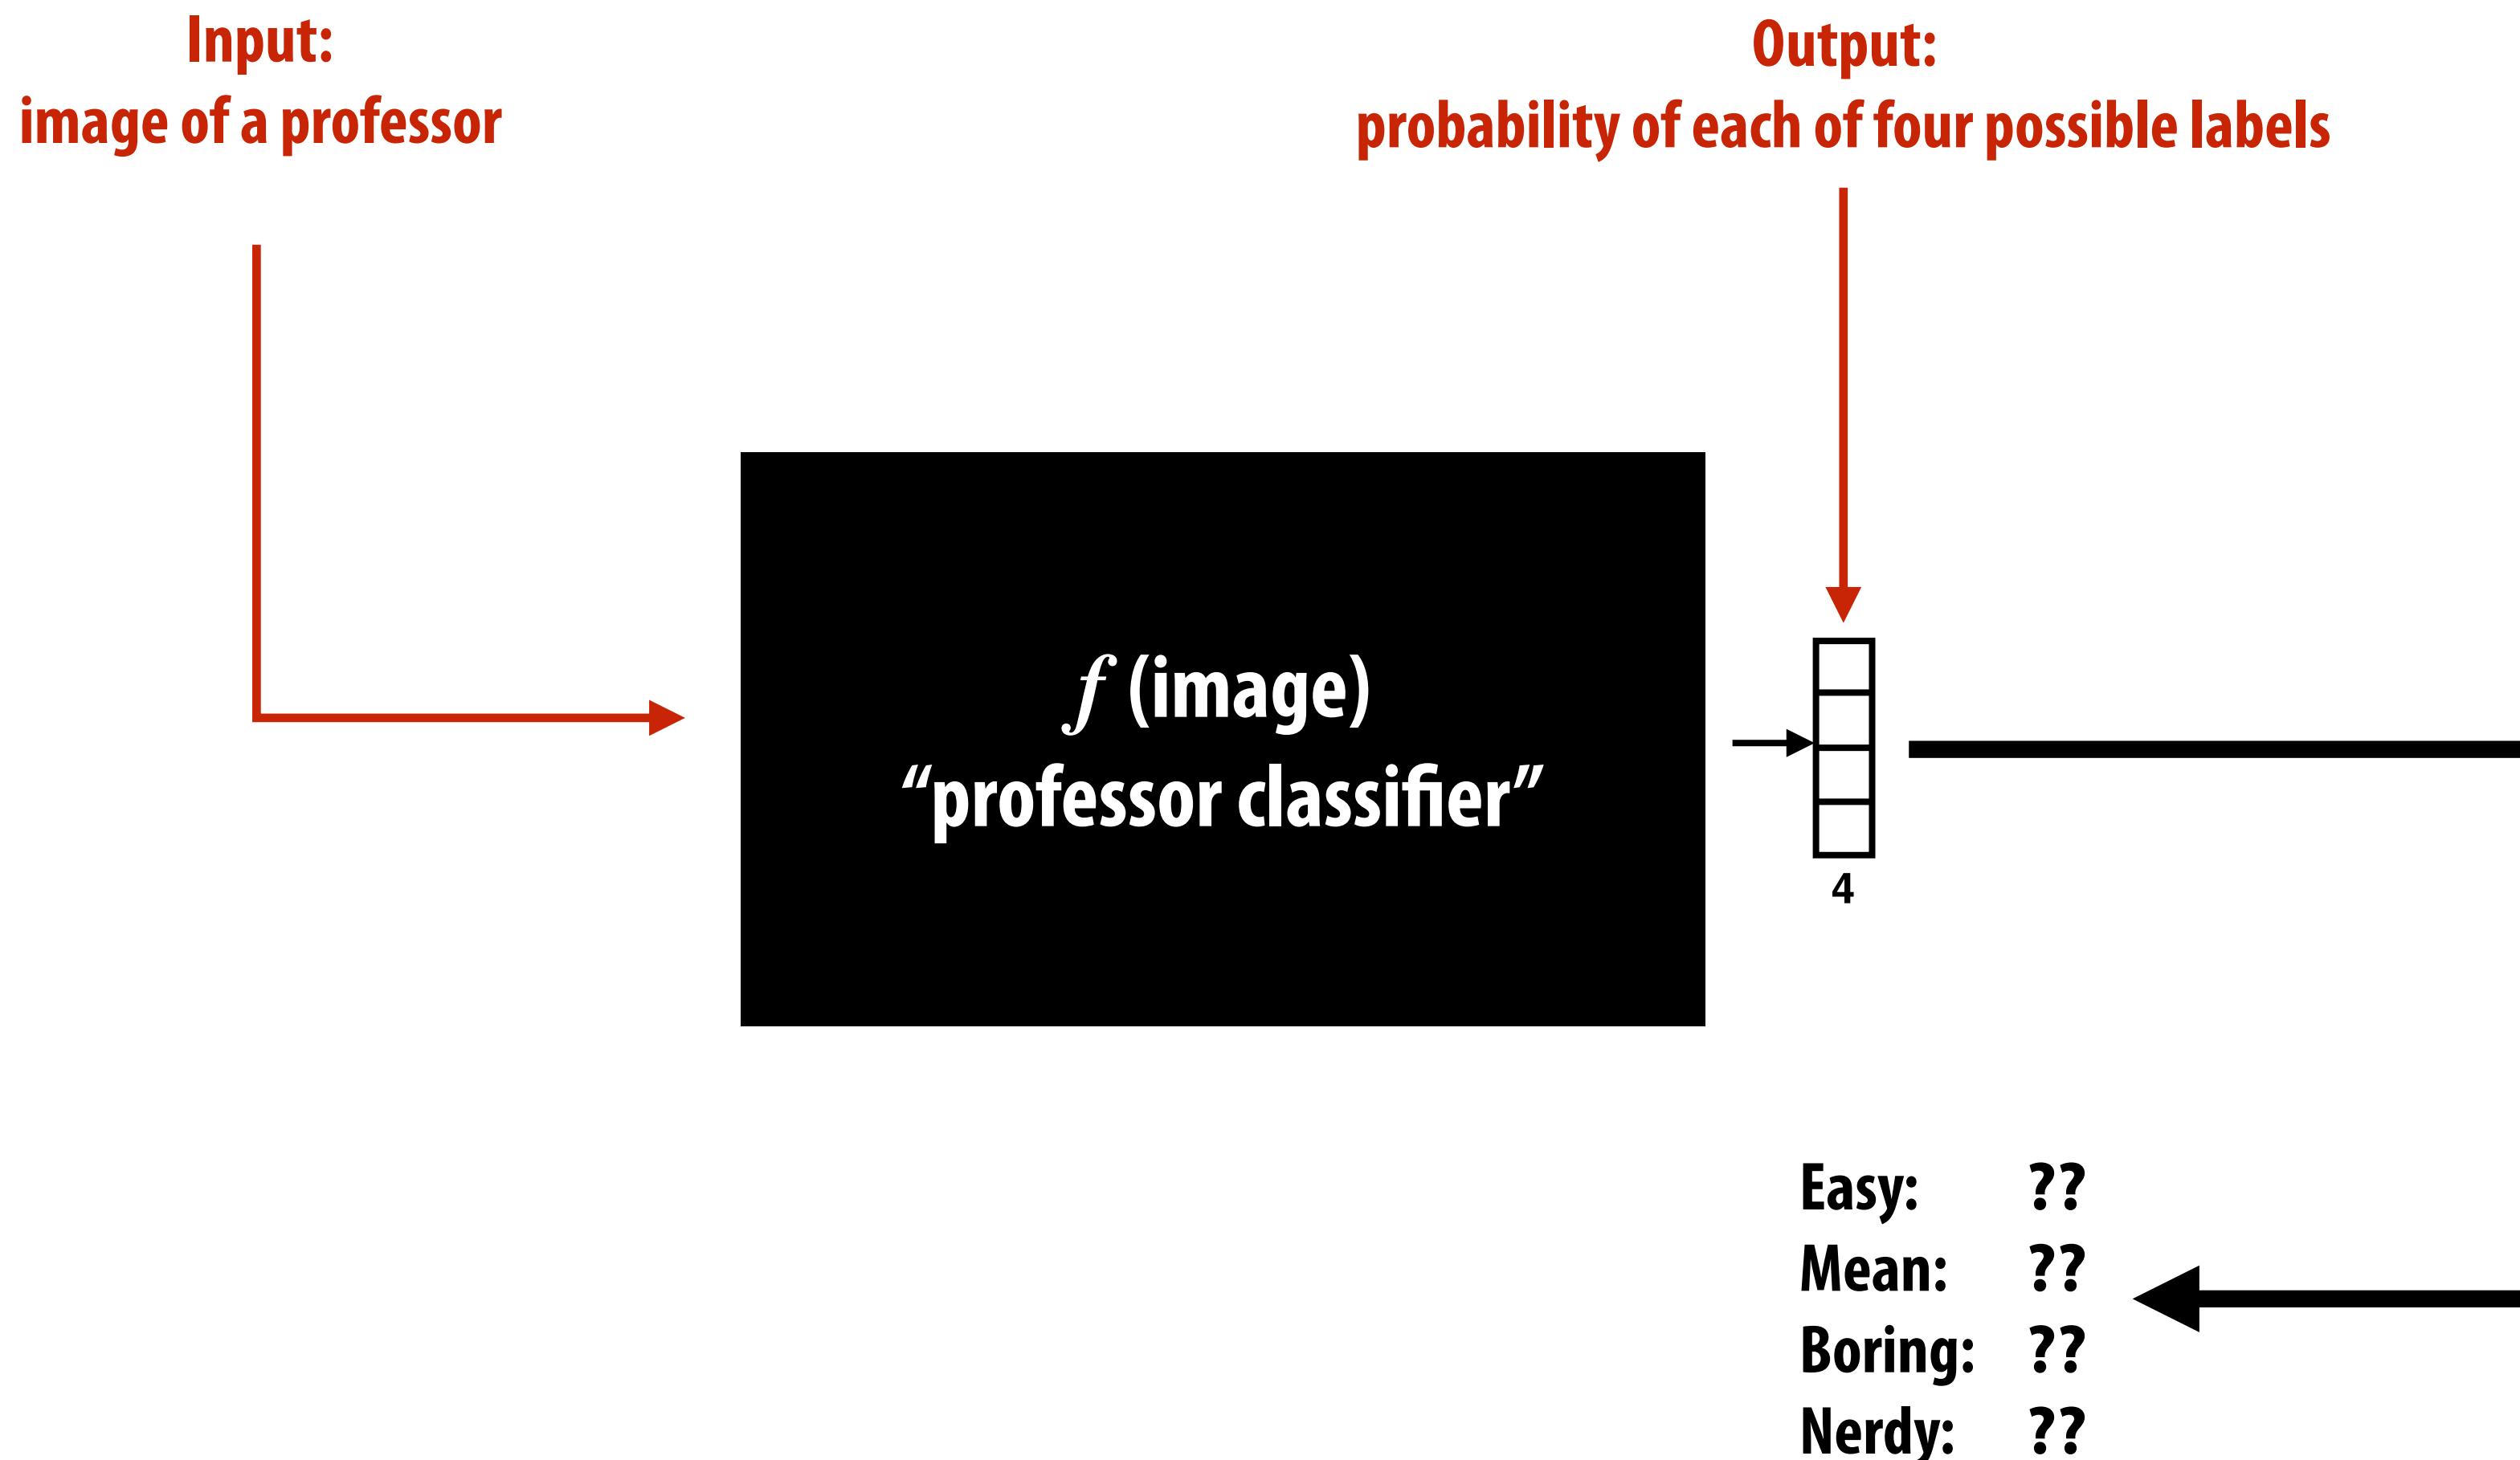
How would you describe this professor?



Easy?
Mean?
Boring?
Nerdy?

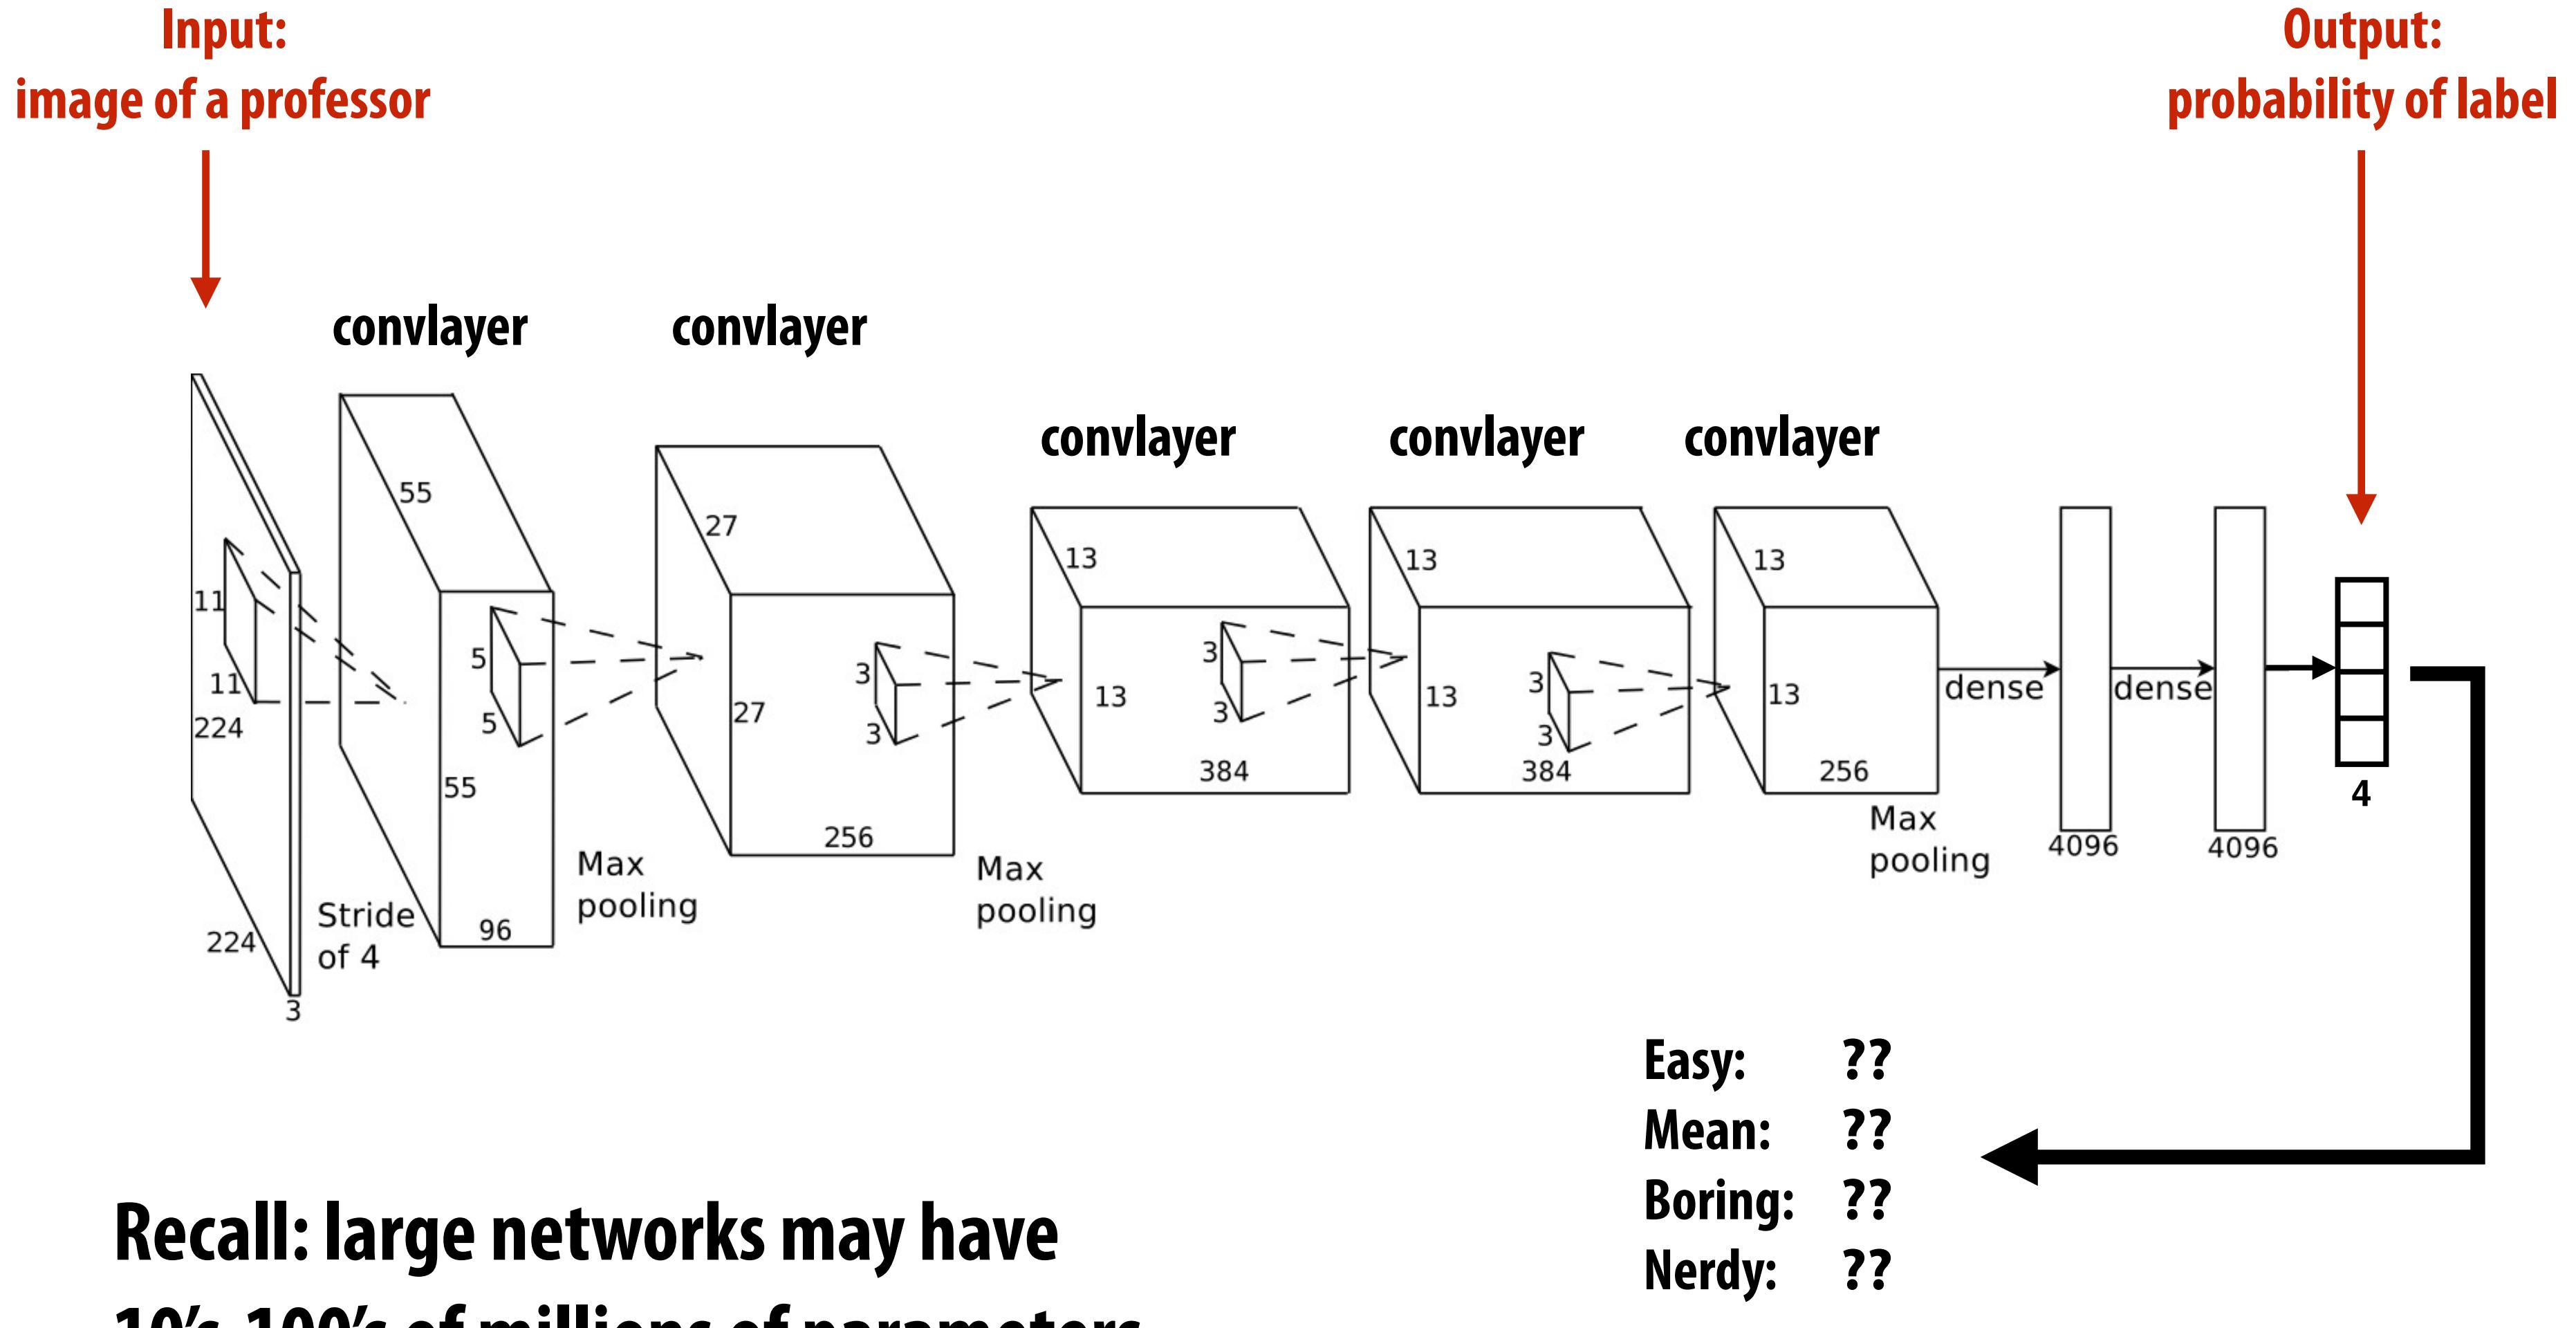
Professor classification task

Classifies professors as easy, mean, boring, or nerdy based on their appearance.

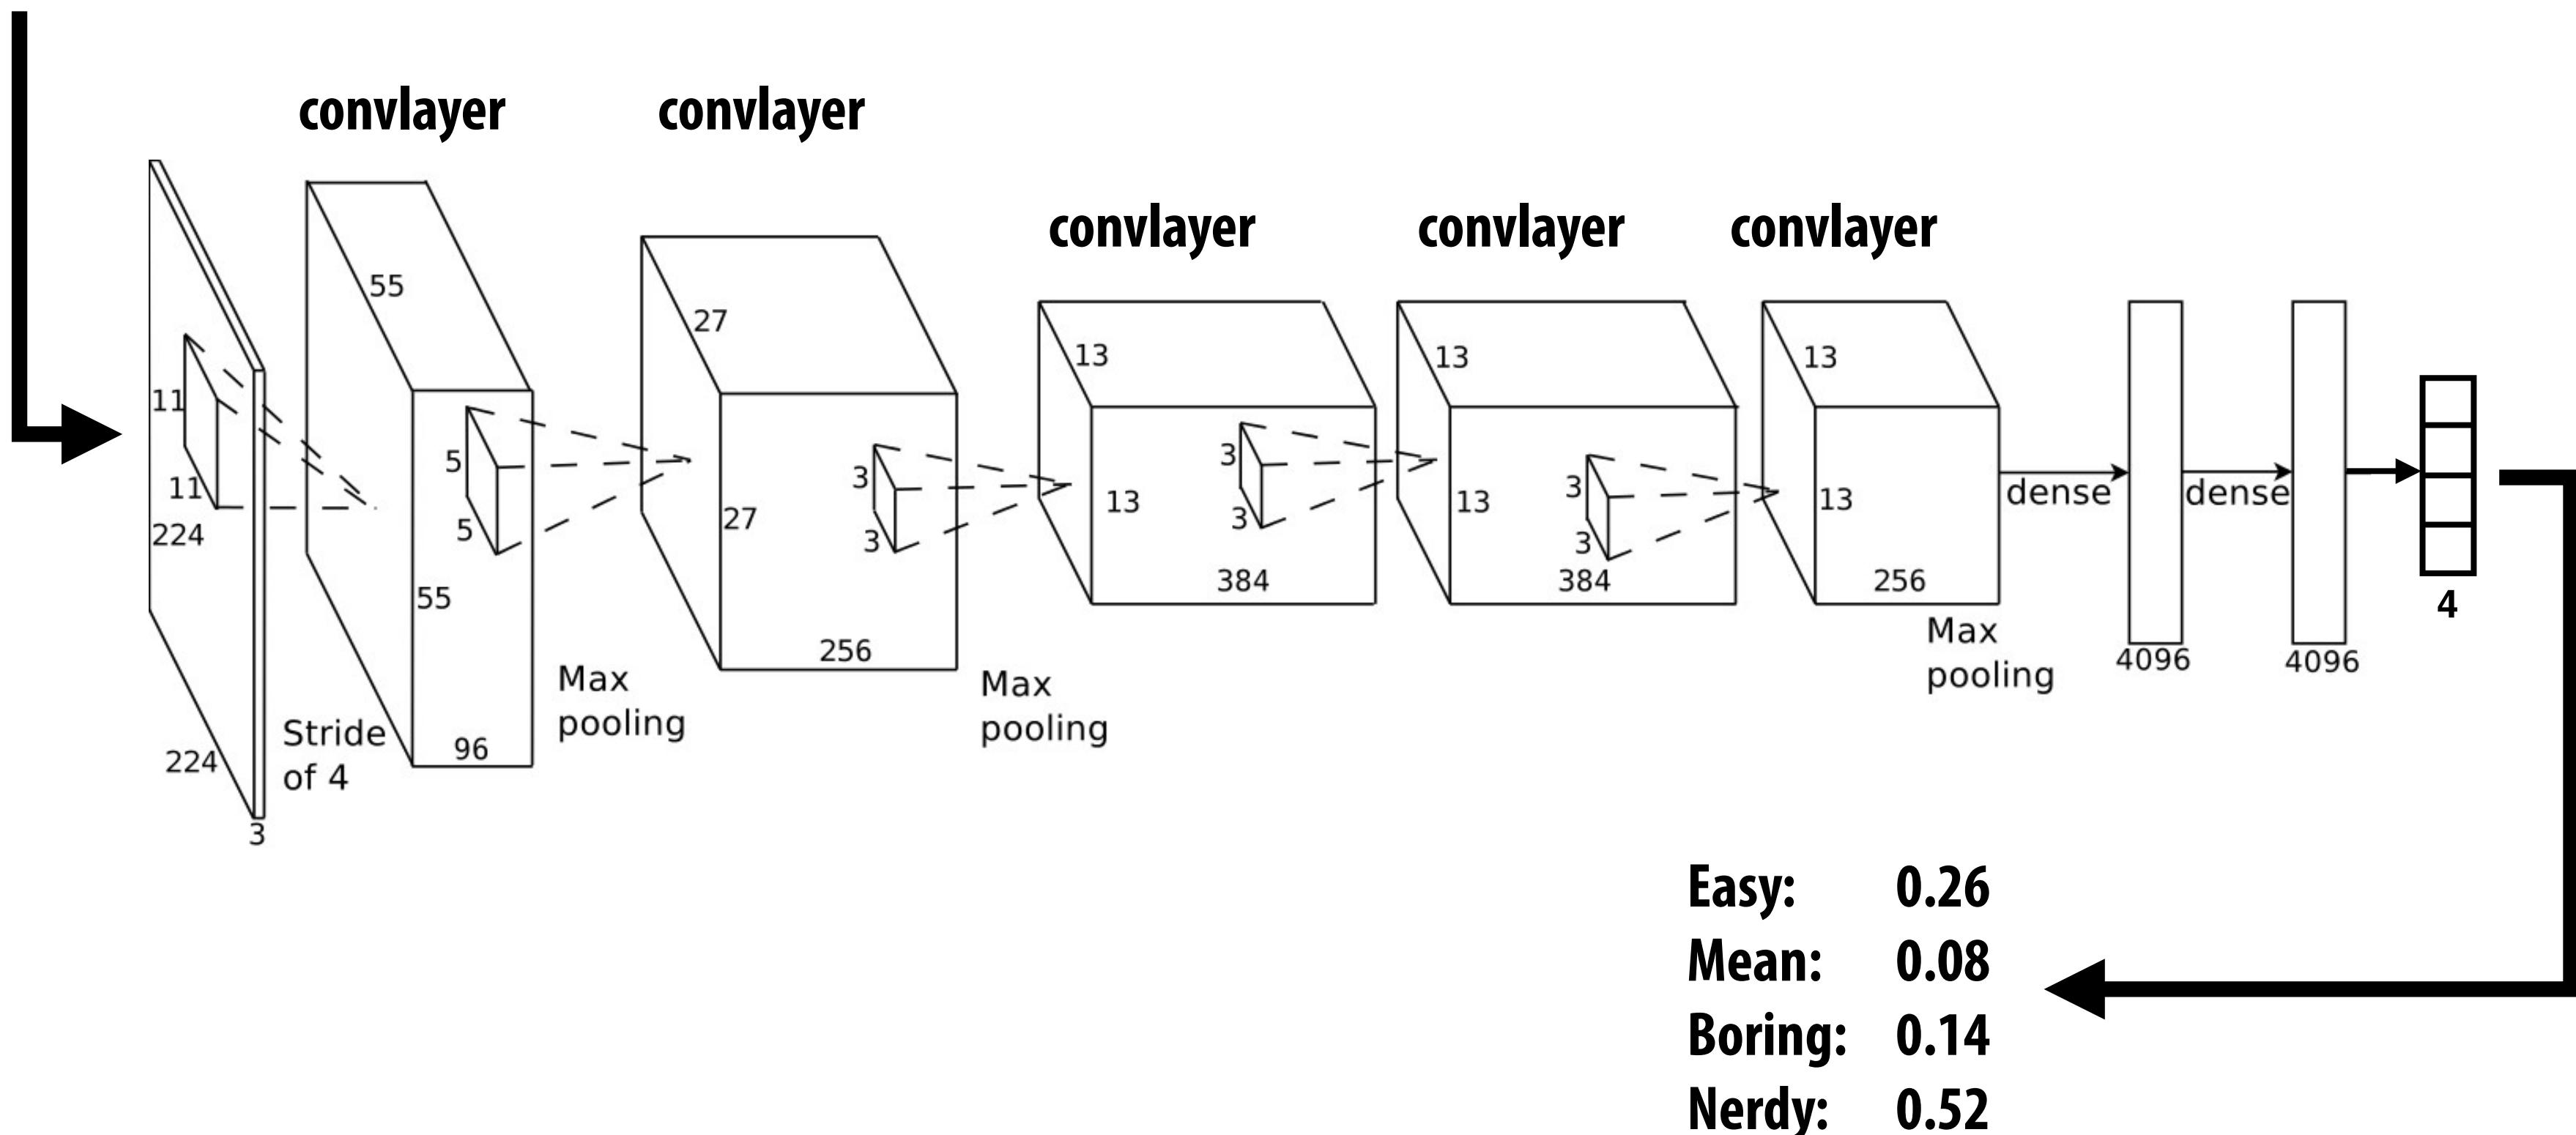


Professor classification network

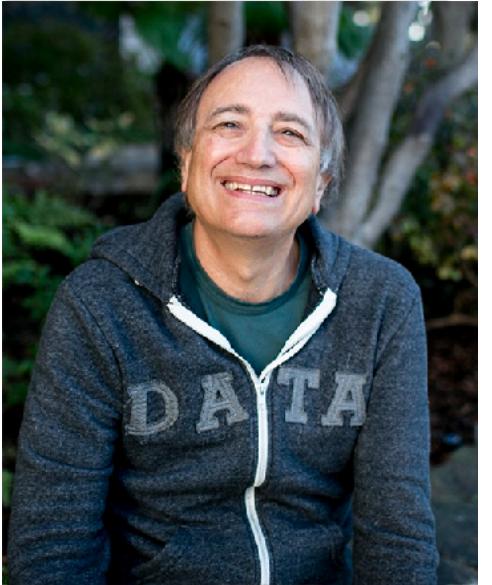
Classifies professors as easy, mean, boring, or nerdy based on their appearance.



Professor classification network



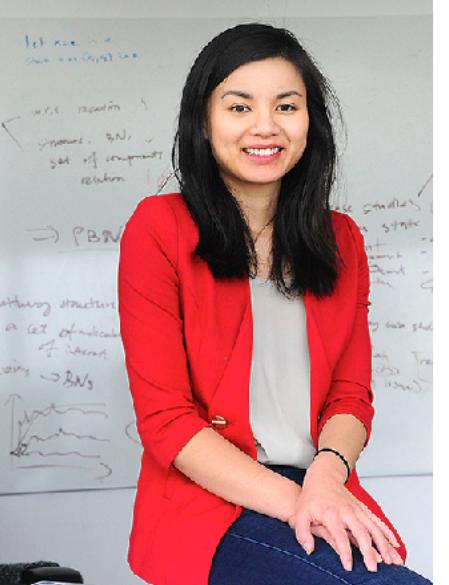
Training data (ground truth answers)



[label omitted]



[label omitted]



[label omitted]



Nerdy



[label omitted]



[label omitted]



[label omitted]



[label omitted]



[label omitted]



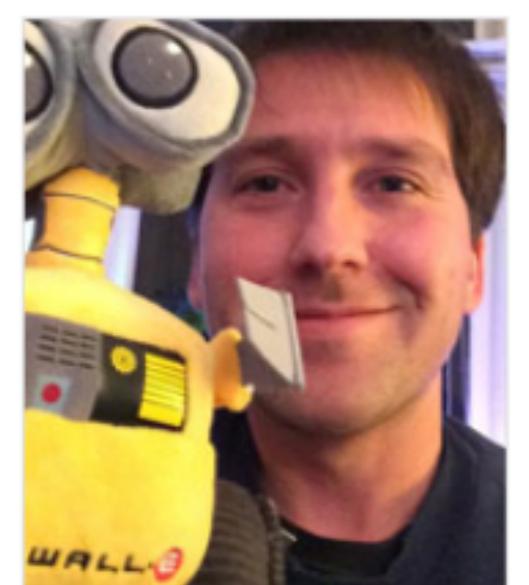
Nerdy



[label omitted]



[label omitted]



Nerdy



[label omitted]



[label omitted]



Nerdy



Nerdy



[label omitted]



[label omitted]



[label omitted]



Nerdy

Professor classification network



New image of Kayvon (not in training set)

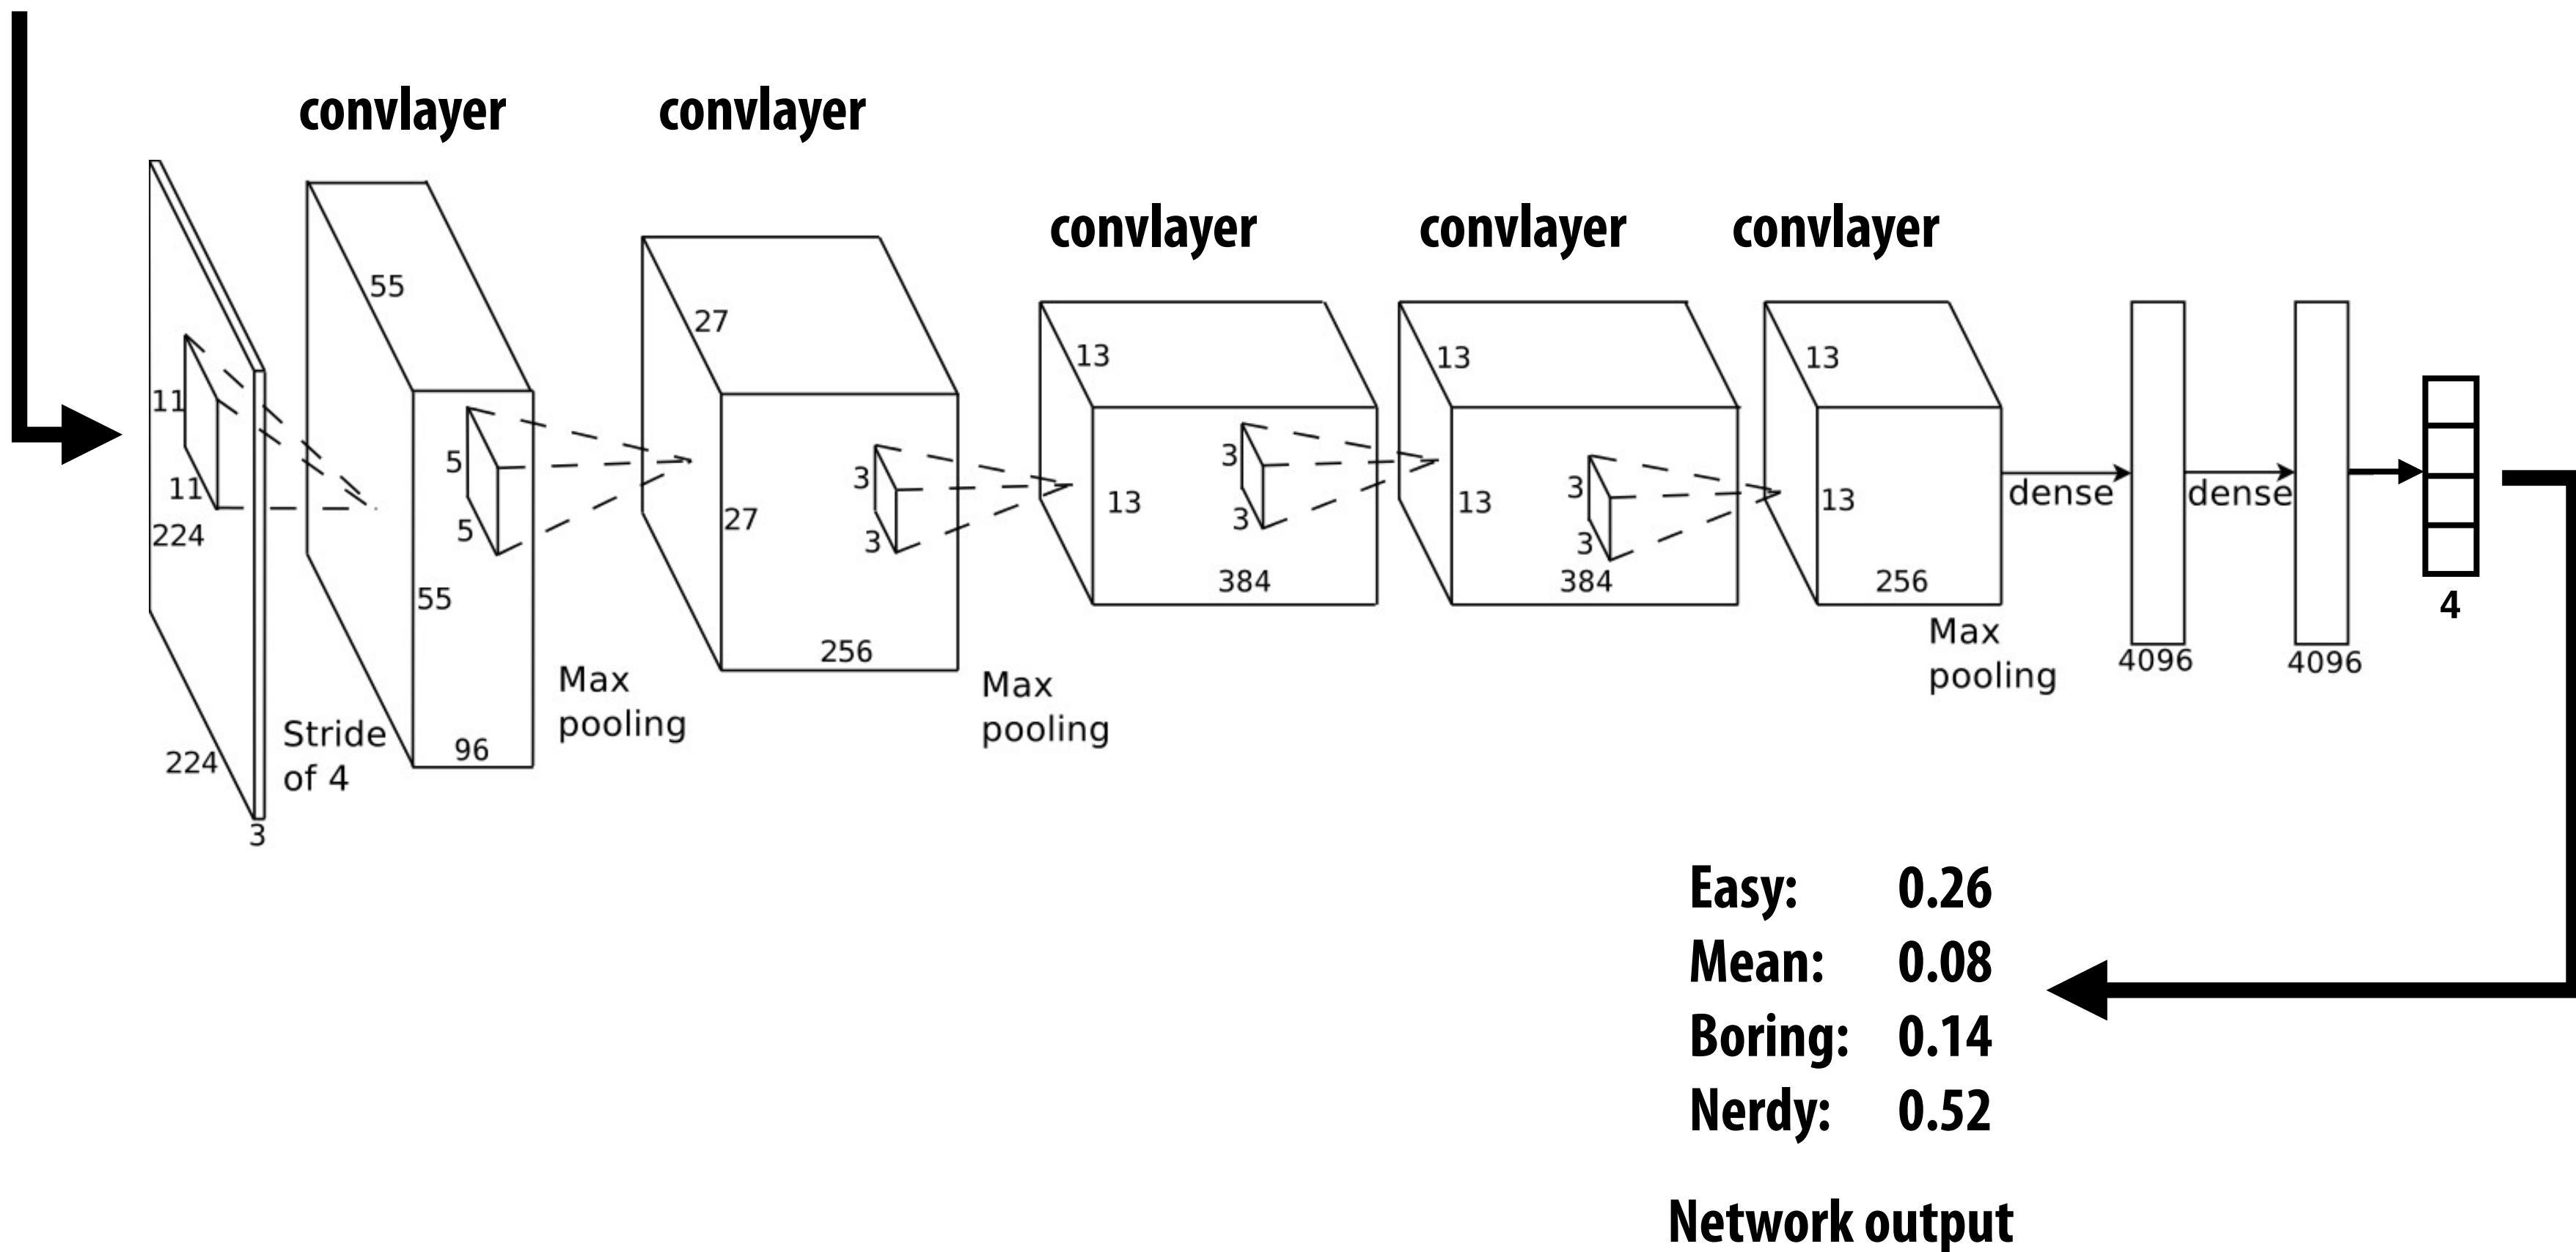
Ground truth (what the answer should be)

Easy: 0.0

Mean: 0.0

Boring: 0.0

Nerdy: 1.0



Error (loss)

Ground truth:
(what the answer should be)

Easy: 0.0

Mean: 0.0

Boring: 0.0

Nerdy: 1.0

Network output: *

Easy: 0.26

Mean: 0.08

Boring: 0.14

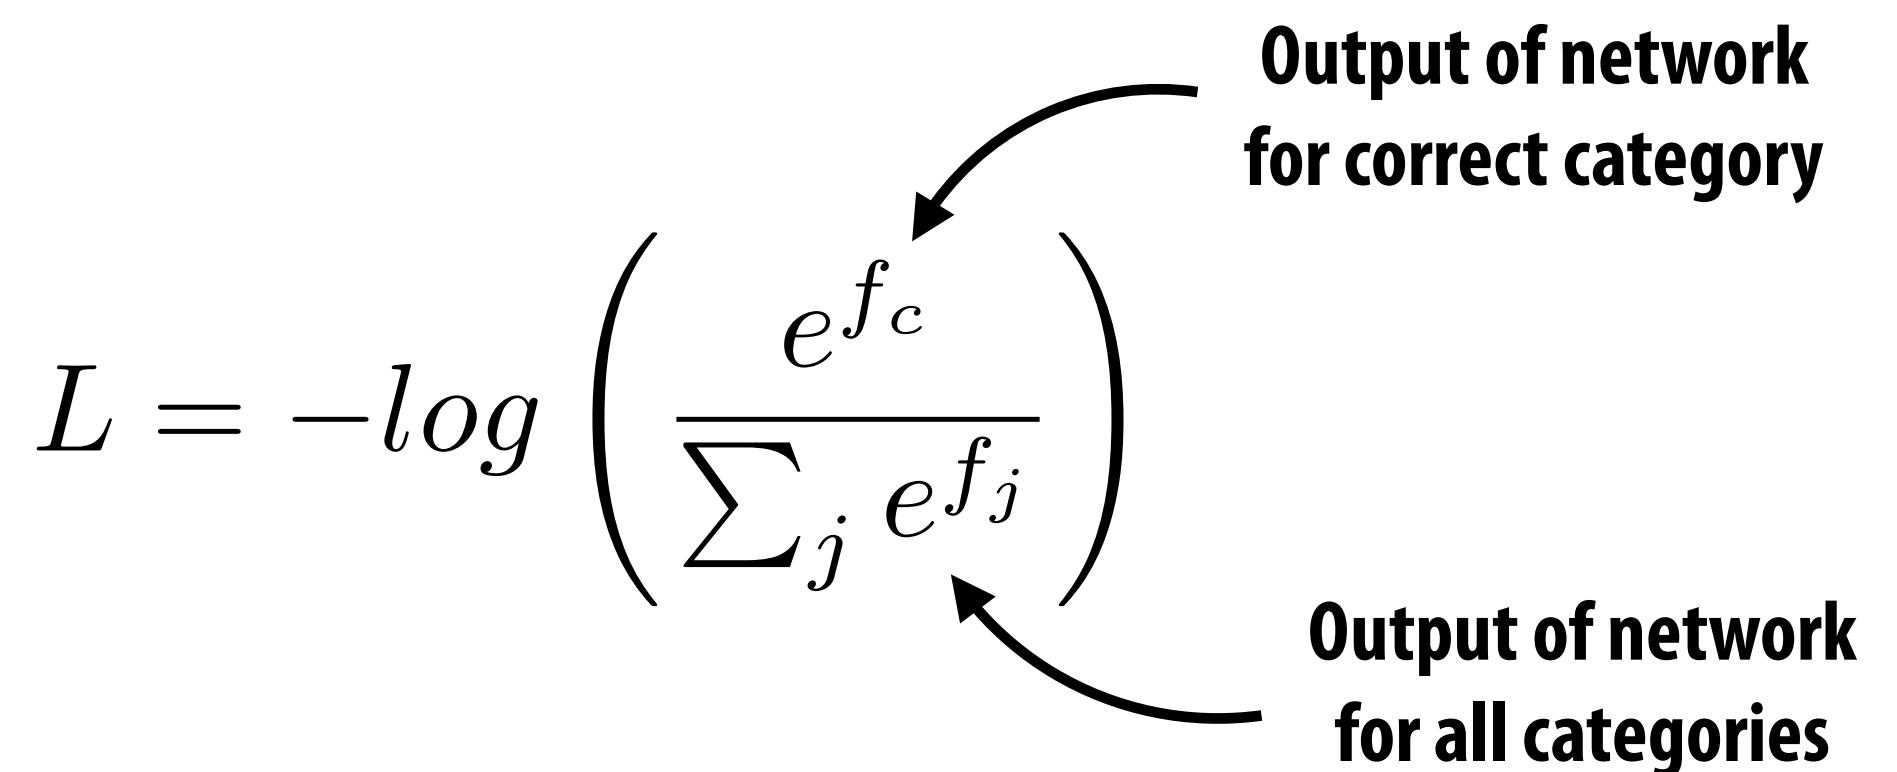
Nerdy: 0.52

Common example: softmax loss:

$$L = -\log \left(\frac{e^{f_c}}{\sum_j e^{f_j}} \right)$$

Output of network
for correct category

Output of network
for all categories

A diagram illustrating the softmax loss formula. The formula is $L = -\log \left(\frac{e^{f_c}}{\sum_j e^{f_j}} \right)$. Two arrows point from the labels 'correct category' and 'all categories' to the respective terms in the equation. One arrow points from 'correct category' to the term e^{f_c} , and another arrow points from 'all categories' to the term $\sum_j e^{f_j}$.

* In practice a network using a softmax classifier outputs unnormalized, log probabilities (f_j),
but I'm showing a probability distribution above for clarity

Training

Goal of training: learning good values of network parameters so that the network outputs the correct classification result for any input image

Idea: minimize loss for all the training examples (for which the correct answer is known)

$$L = \sum_i L_i \quad (\text{total loss for entire training set is sum of losses } L_i \text{ for each training example } x_i)$$

Intuition: if the network gets the answer correct for a wide range of training examples, then hopefully it has learned parameter values that yield the correct answer for future images as well.

Intuition: gradient descent

Say you had a function f that contained hidden parameters p_1 and p_2 : $f(x_i)$

And for some input x_i , your training data says the function should output 0.

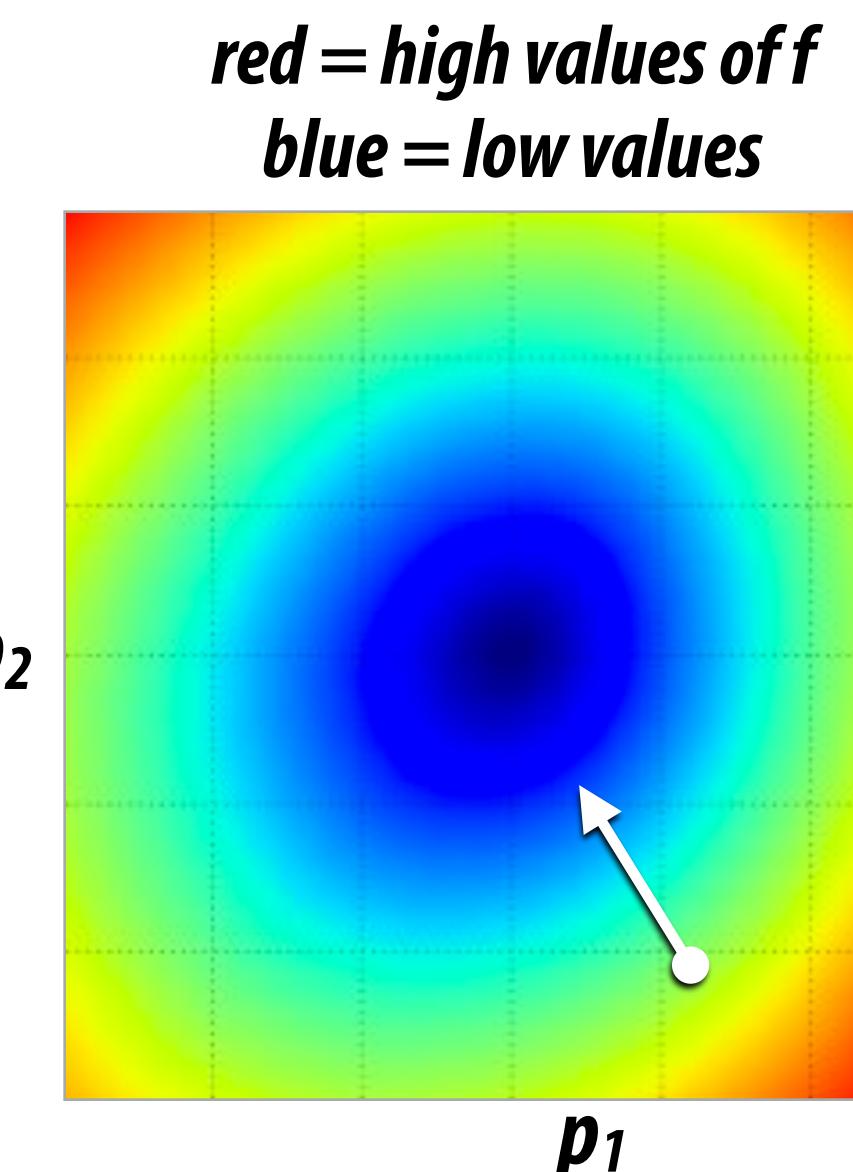
But for the current values of p_1 and p_2 , it currently outputs 10.

$$f(x_i, p_1, p_2) = 10$$

And say I also gave you expressions for the derivative of f with respect to p_1 and p_2 so you could compute their value at x_i .

$$\frac{df}{dp_1} = 2 \quad \frac{df}{dp_2} = -5$$

$$\nabla f = [2, -5]$$



How might you adjust the values p_1 and p_2 to reduce the error for this training example?

Basic gradient descent

```
while (loss too high):  
    for each epoch: // a pass through the training dataset  
        for each item  $x_i$  in training set:  
            grad = evaluate_loss_gradient( $f$ , params, loss_func,  $x_i$ )  
            params += -grad * learning_rate;
```

Mini-batch stochastic gradient descent (mini-batch SGD):

choose a random (small) subset of the training examples to use to compute the gradient in each iteration of the while loop

```
while (loss too high):  
    for each epoch: // a pass through the training dataset  
        for all mini batches in training set:  
            grad = 0;  
            for each item  $x_i$  in minibatch:  
                grad += evaluate_loss_gradient( $f$ , params, loss_func,  $x_i$ )  
            params += -grad * learning_rate;
```

How do we compute $d\text{Loss}/dp$ for a deep neural network with millions of parameters?

Quick review of back-propagation

Derivatives using the chain rule

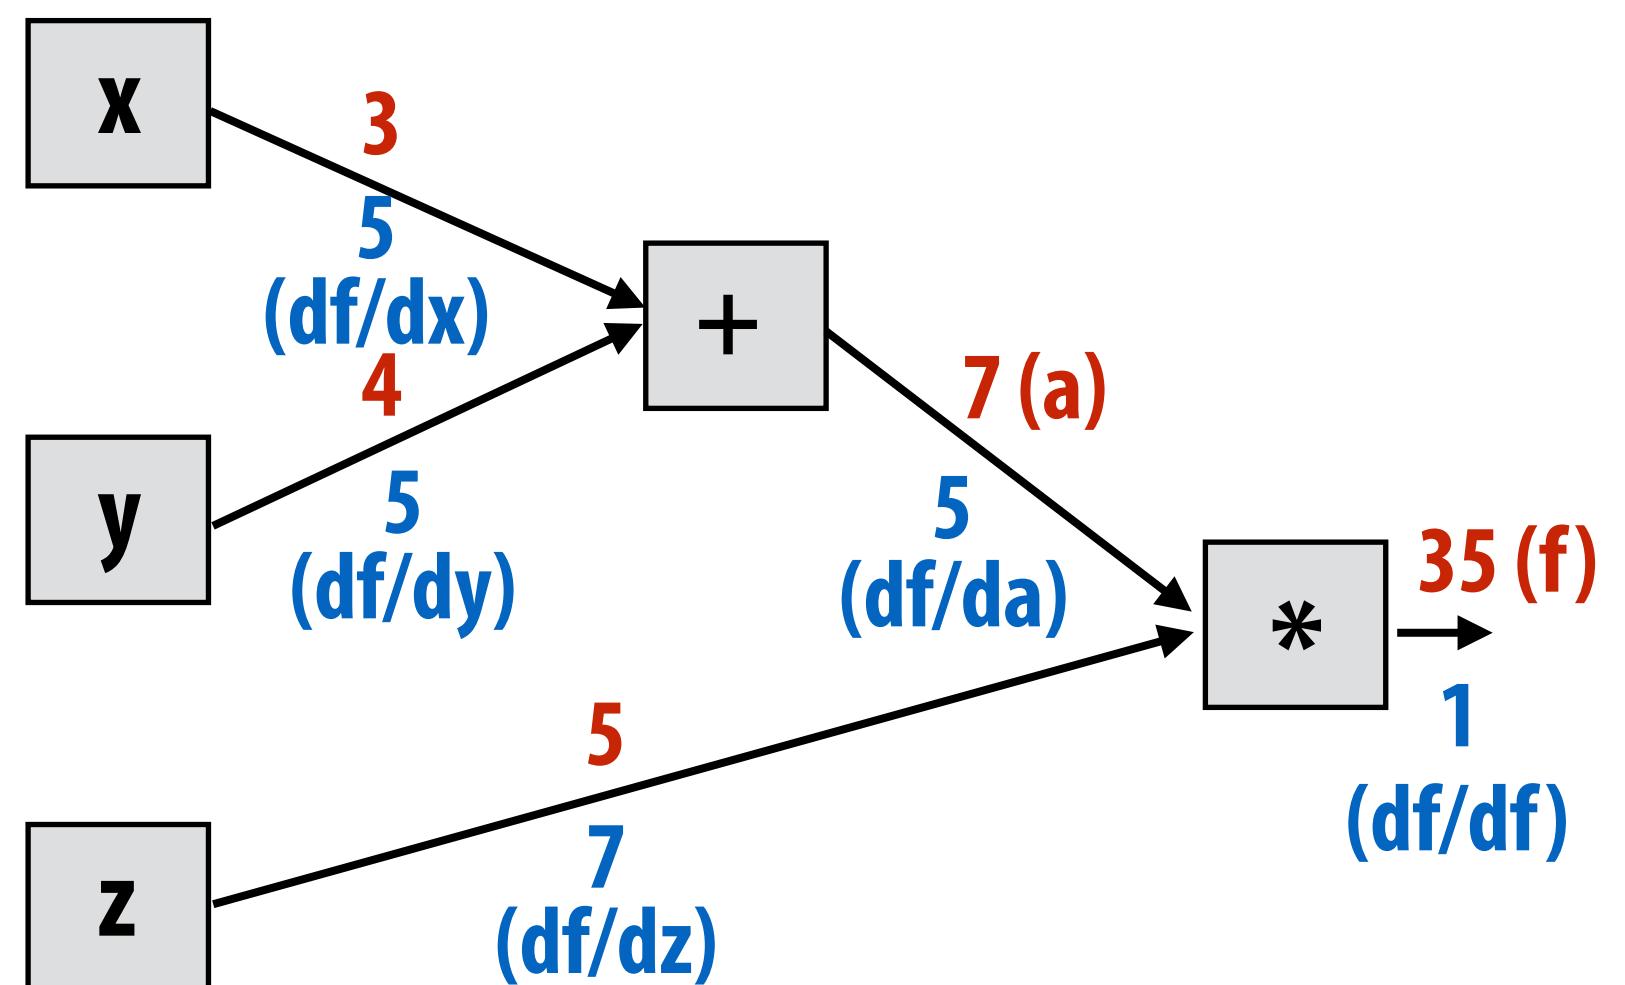
$$f(x, y, z) = (x + y)z = az$$

Where: $a = x + y$

$$\frac{df}{da} = z \quad \frac{da}{dx} = 1 \quad \frac{da}{dy} = 1$$

So, by the derivative chain rule:

$$\frac{df}{dx} = \frac{df}{da} \frac{da}{dx} = z$$



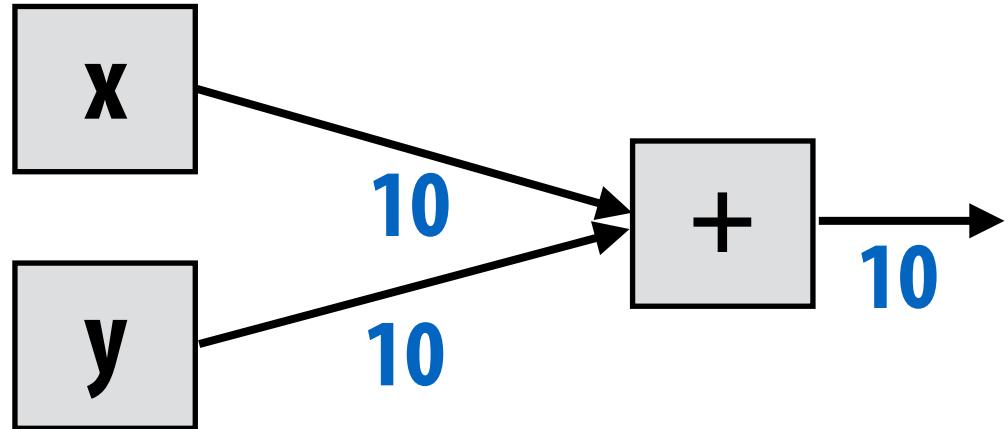
Red = output of node
Blue = $df/dnnode$

Backpropagation

Red = output of node

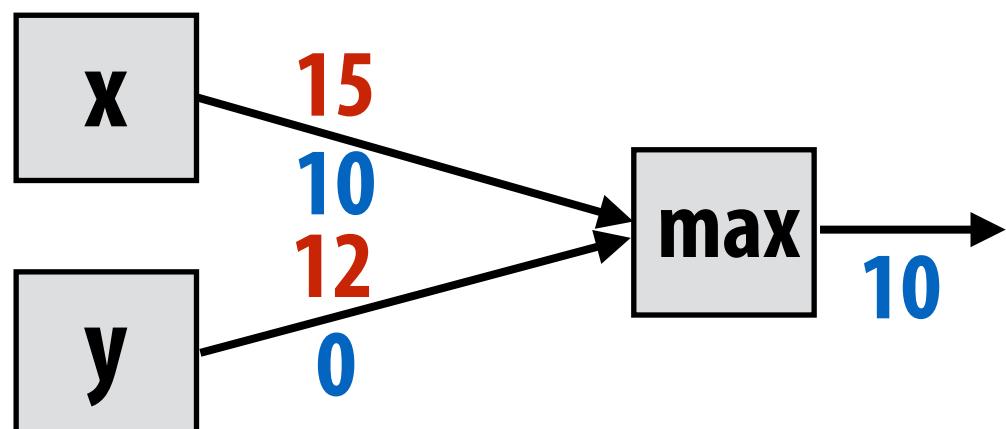
Blue = df/dnode

Recall: $\frac{df}{dx} = \frac{df}{dg} \frac{dg}{dx}$



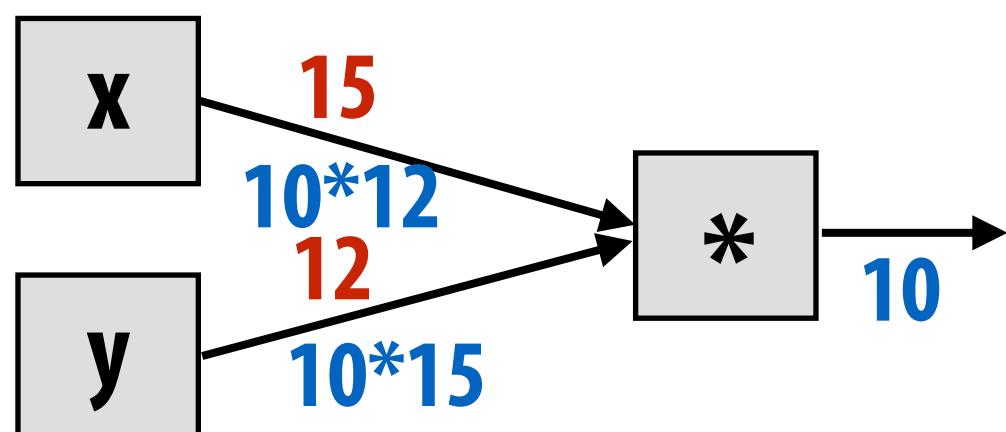
$$g(x, y) = x + y$$

$$\frac{dg}{dx} = 1, \frac{dg}{dy} = 1$$



$$g(x, y) = \max(x, y)$$

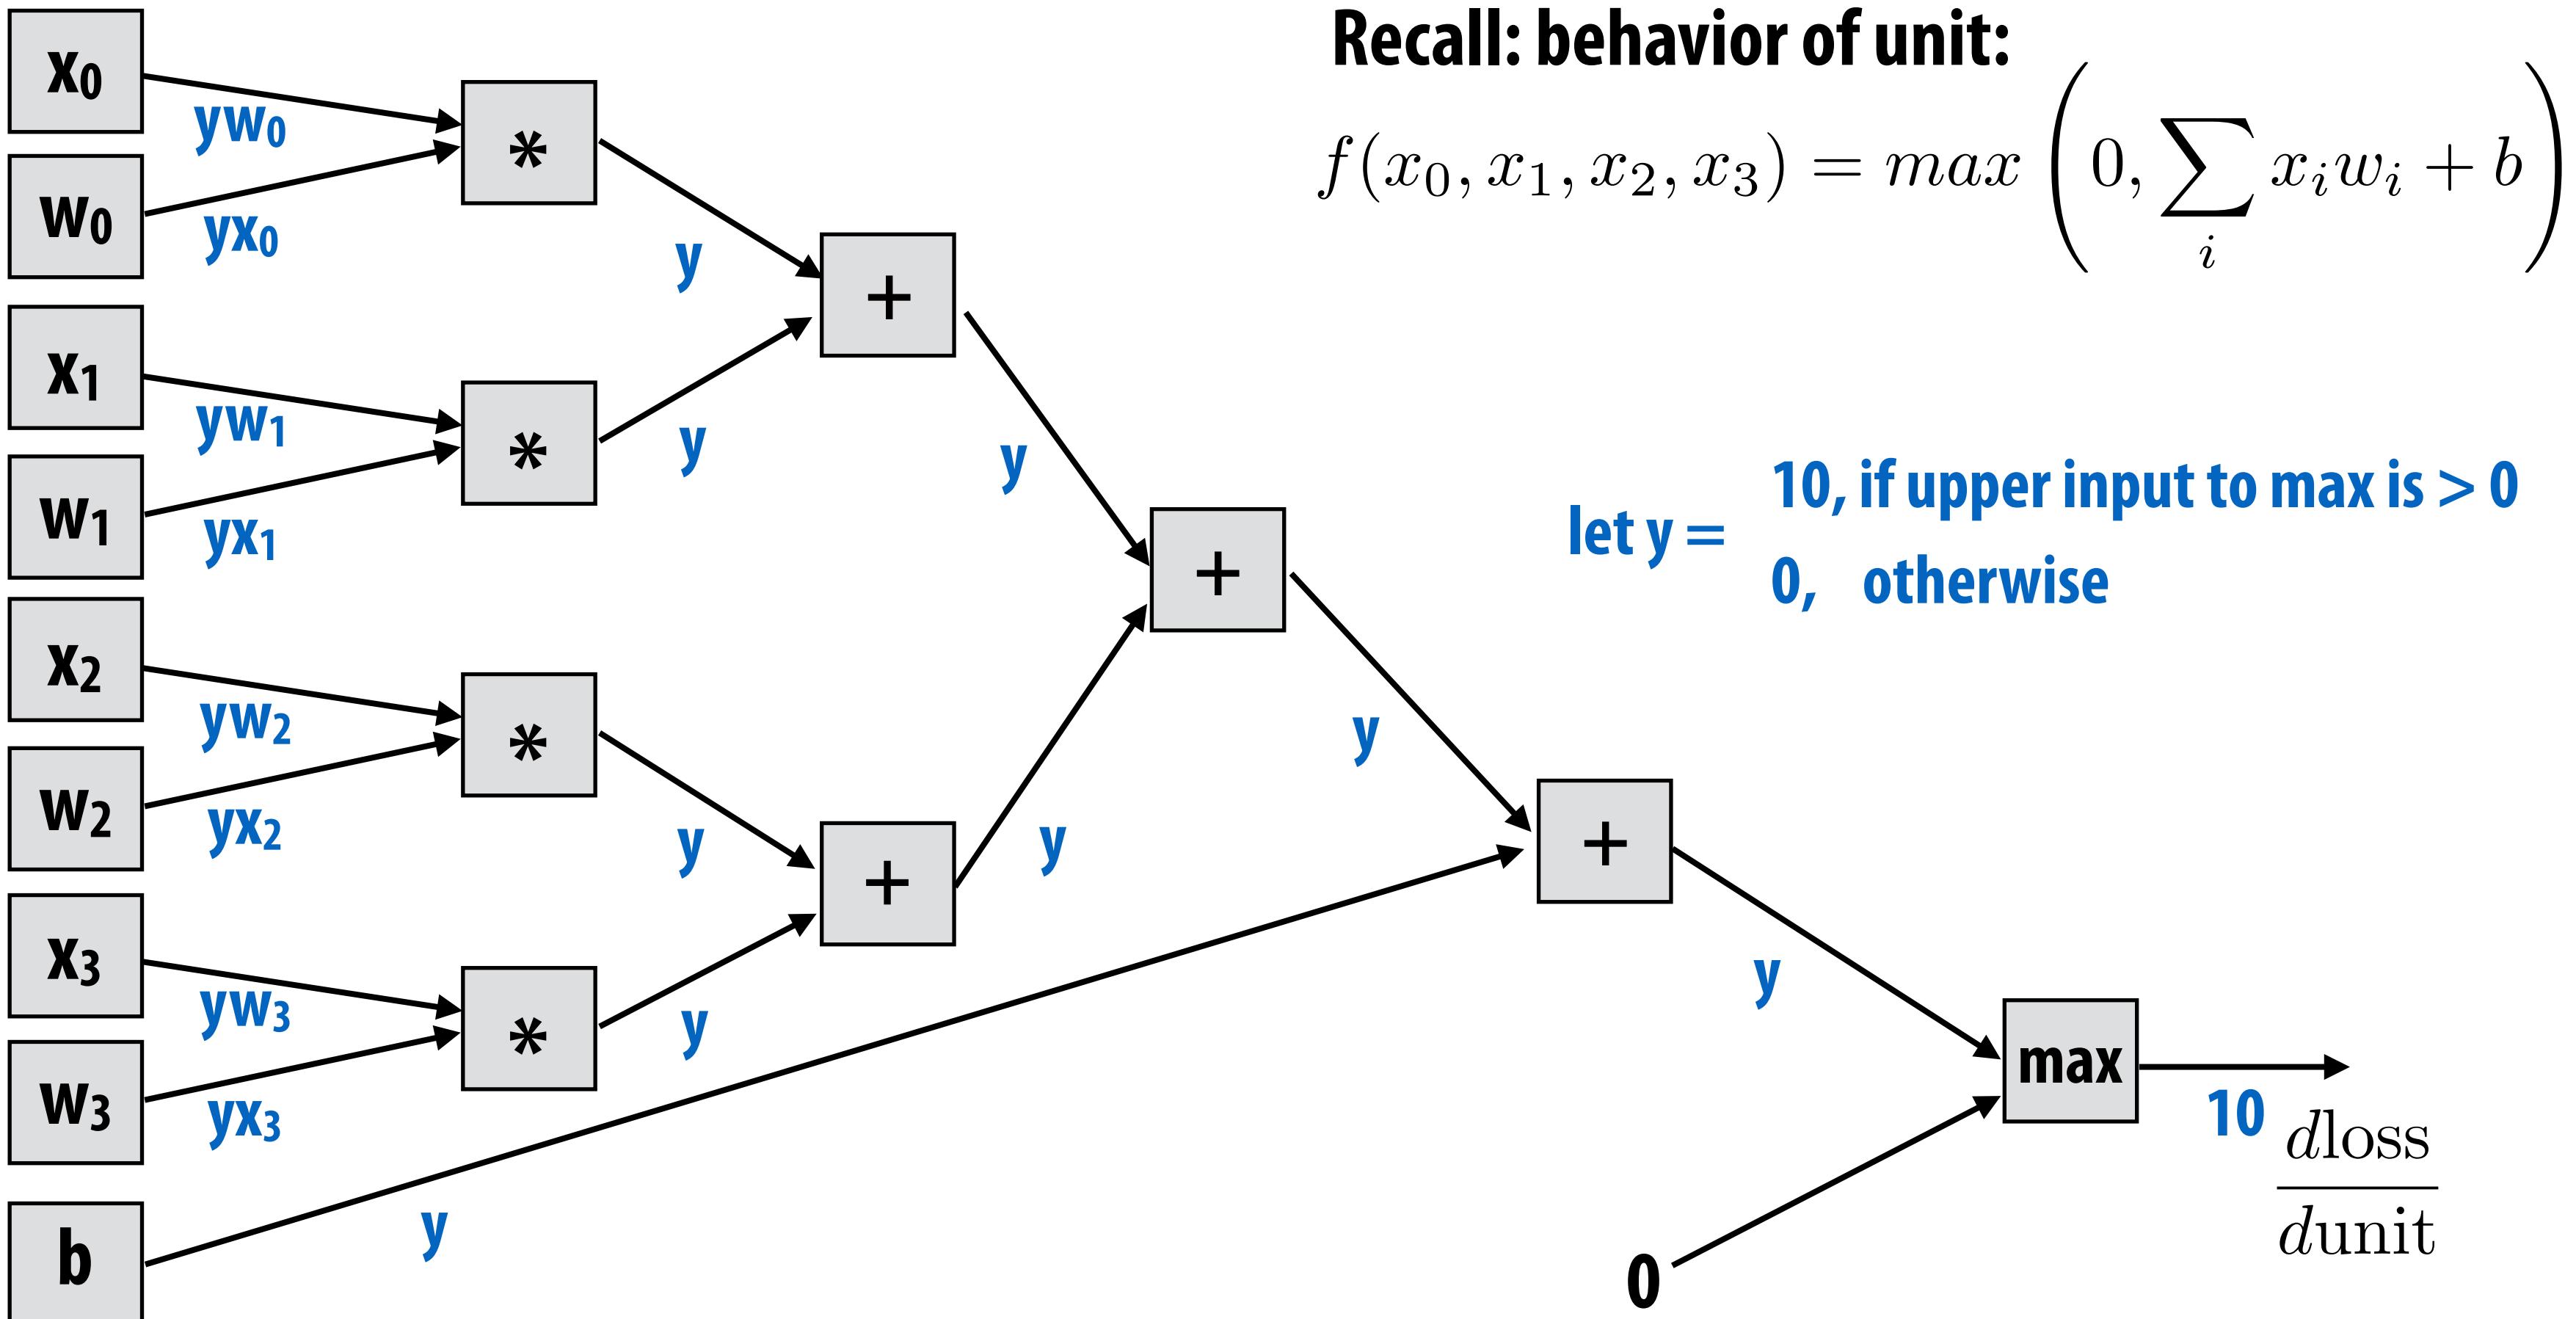
$$\frac{dg}{dx} = \begin{cases} 1, & \text{if } x > y \\ 0, & \text{otherwise} \end{cases}$$



$$g(x, y) = xy$$

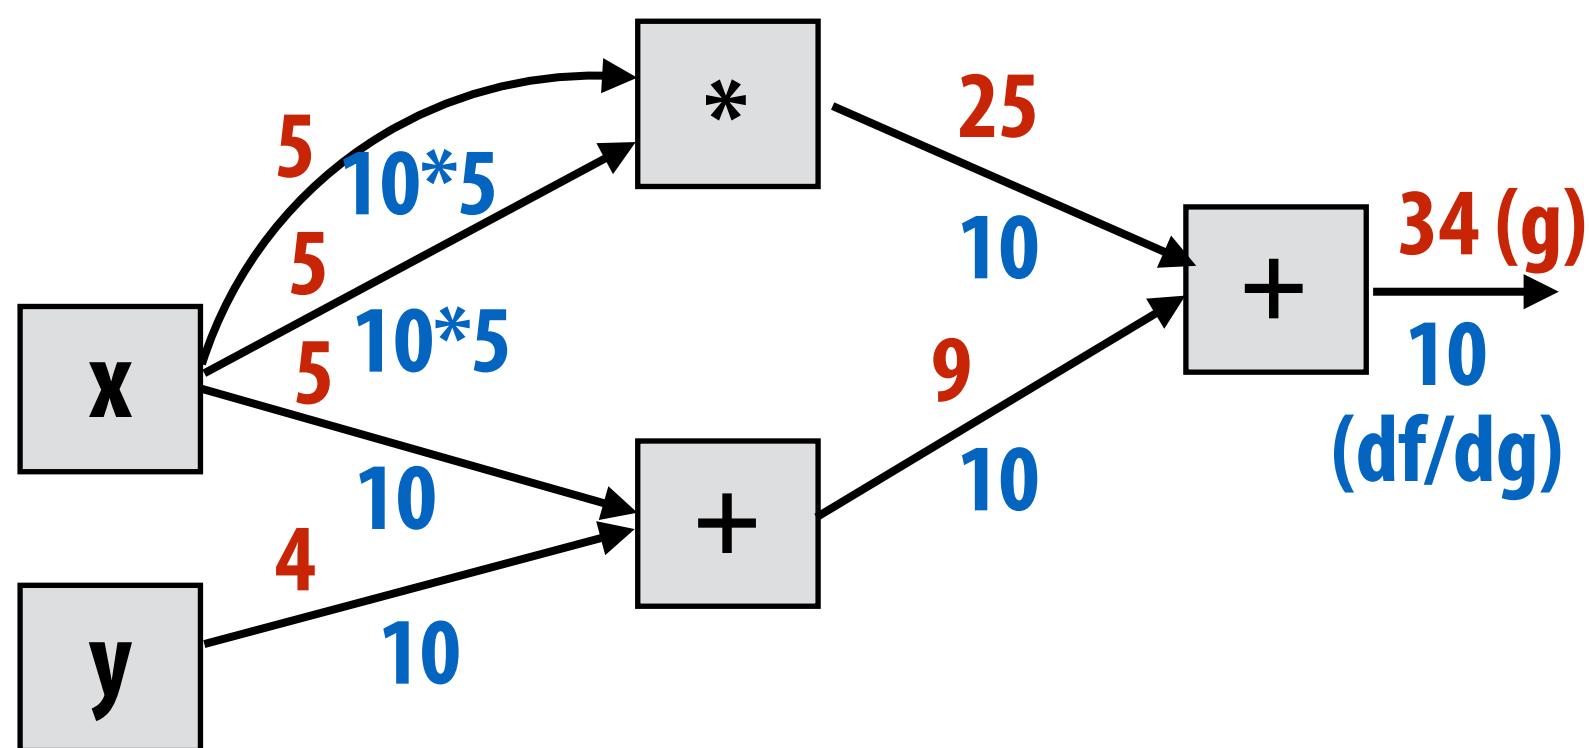
$$\frac{dg}{dx} = y, \frac{dg}{dy} = x$$

Back-propagating through single unit



Observe: output of prior layer must be retained in order to compute weight gradients for this unit during backprop.

Multiple uses of an input variable



$$g(x, y) = (x + y) + x * x = a + b$$

$$\frac{da}{dx} = 1, \quad \frac{db}{dx} = 2x$$

$$\frac{dg}{dx} = \frac{dg}{da} \frac{da}{dx} + \frac{dg}{db} \frac{db}{dx} = 2x + 1$$

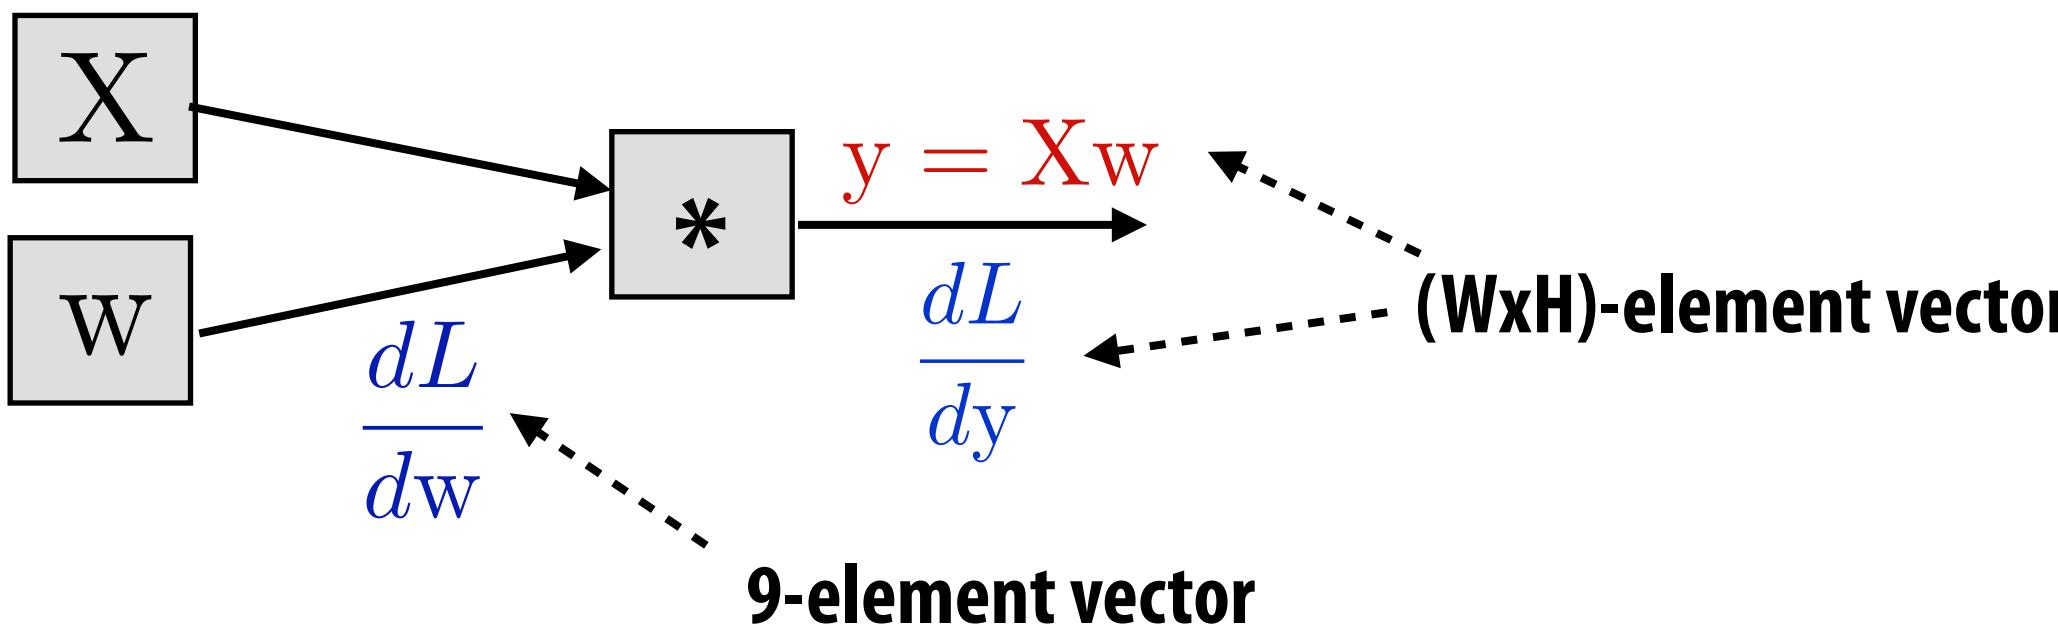
Implication: backpropagation through all units in a convolutional layer adds gradients computed from each unit to the overall gradient for the shared weights

Sum gradients from each use of variable:

Here:

$$\begin{aligned}\frac{df}{dx} &= \frac{df}{dg} \frac{dg}{dx} \\ &= 10 \frac{dg}{dx} \\ &= 10(2x + 1) \\ &= 10(10 + 1) = 110\end{aligned}$$

Back-propagation: matrix form

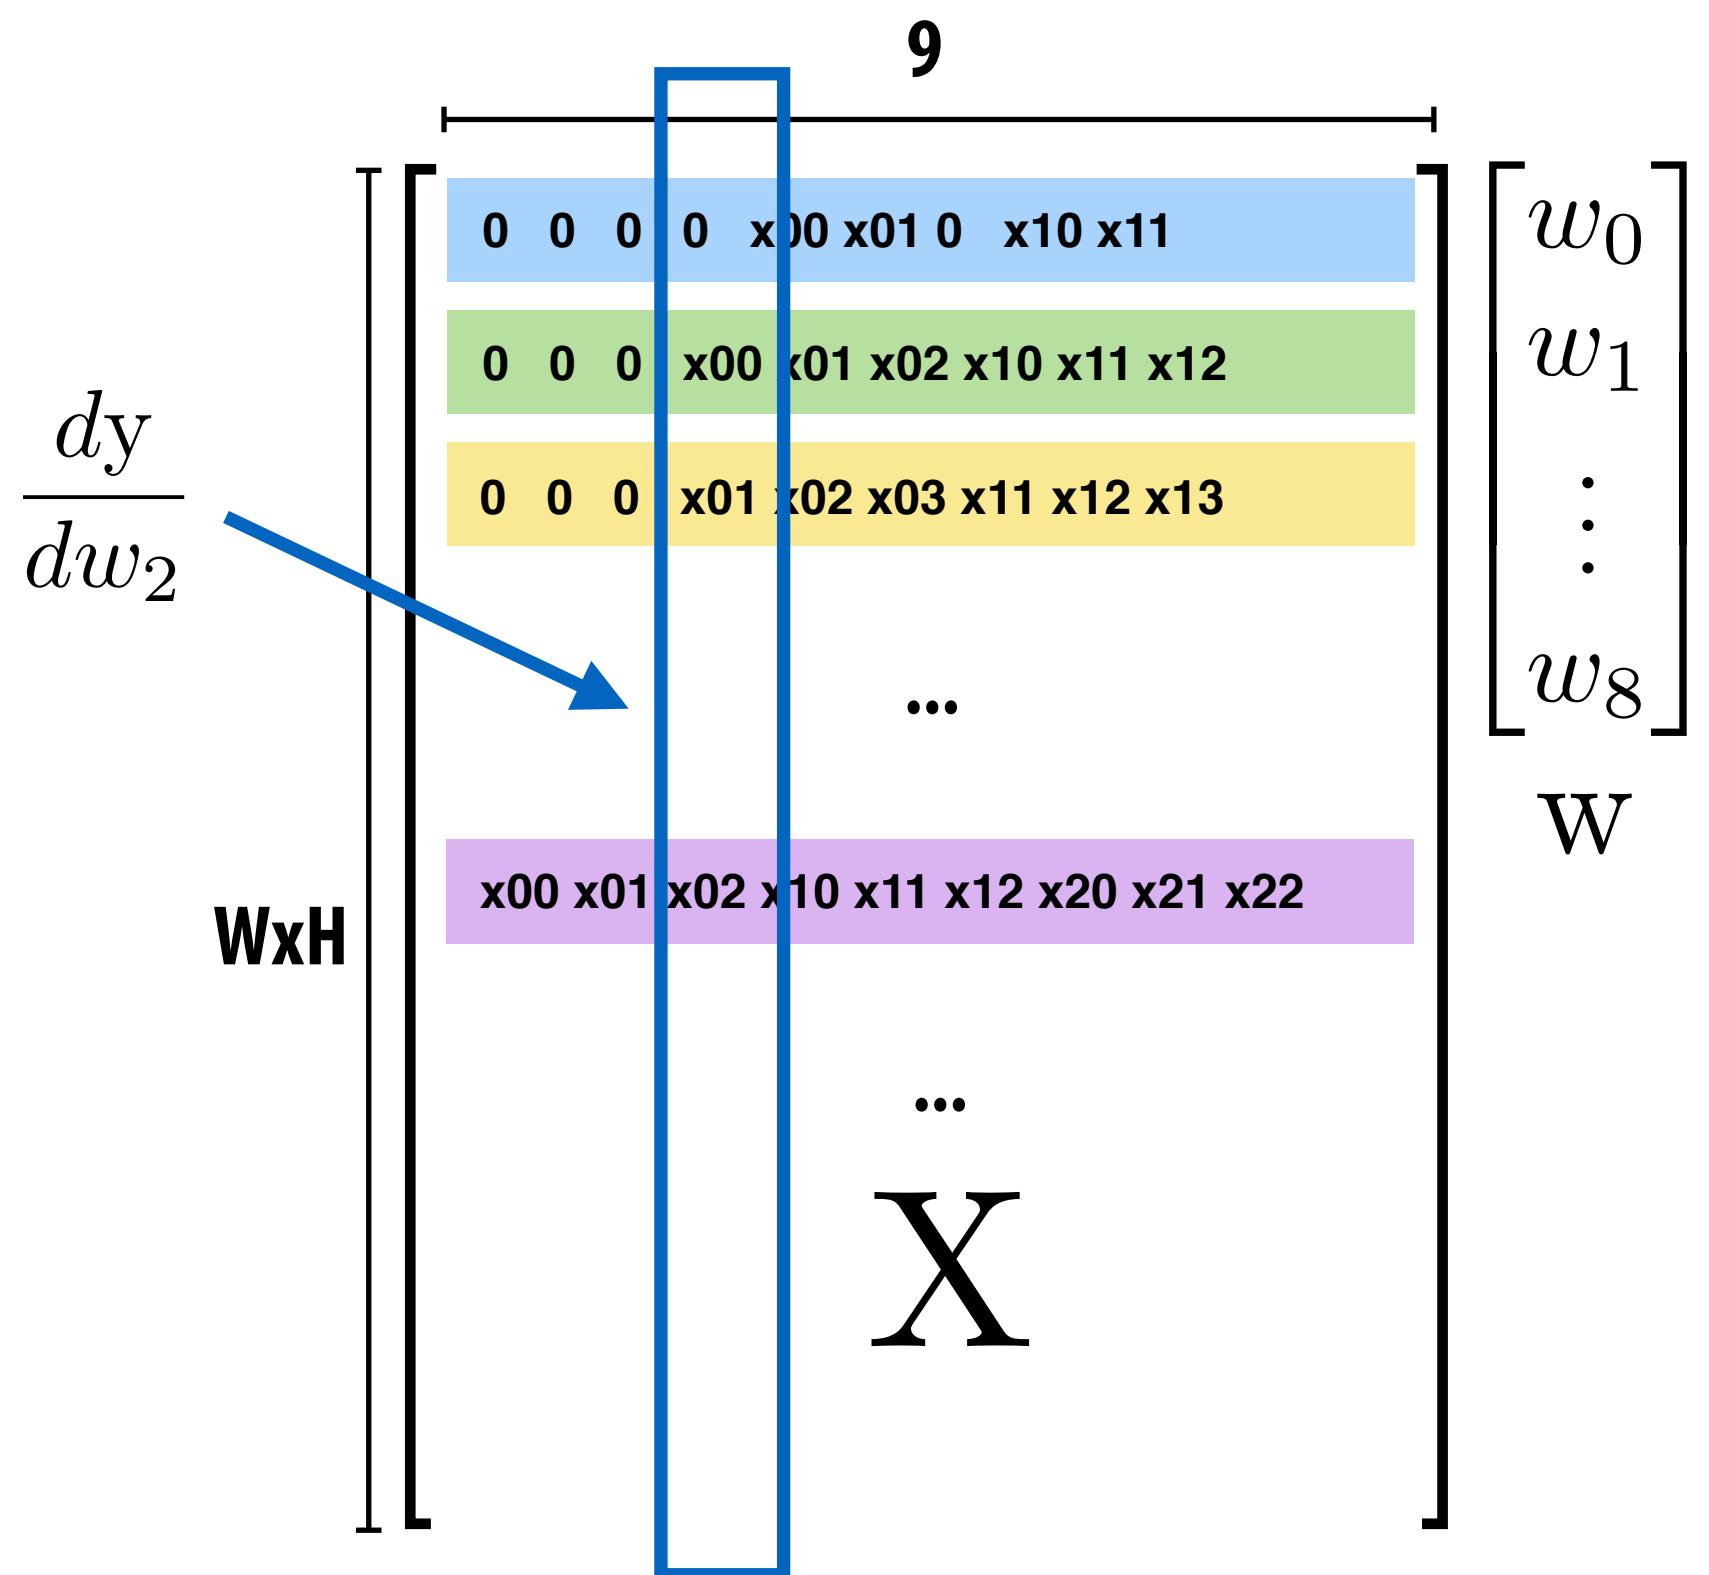


$$\frac{dy_j}{dw_i} = X_{ji}$$

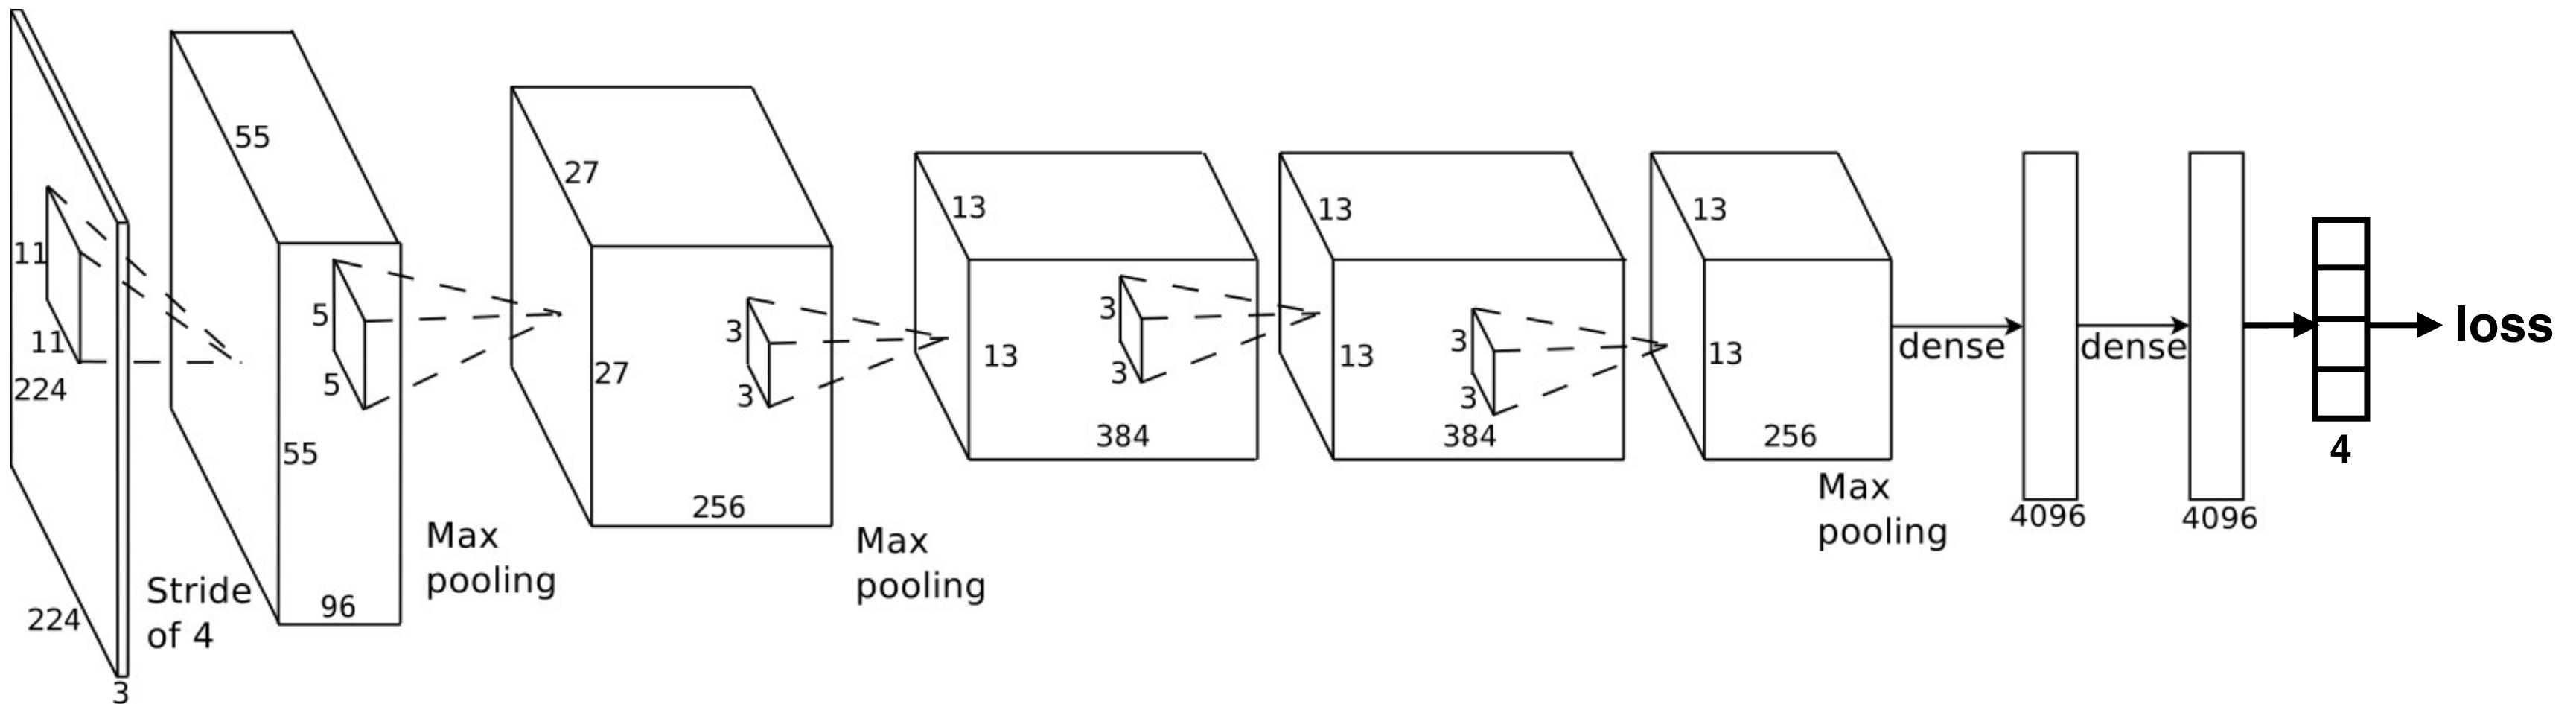
$$\begin{aligned}\frac{dL}{dw_i} &= \sum_j \frac{dL}{dy_j} \frac{dy_j}{dw_i} \\ &= \sum_j \frac{dL}{dy_j} X_{ji}\end{aligned}$$

Therefore:

$$\frac{dL}{dw} = X^T \frac{dL}{dy}$$



Backpropagation through the entire professor classification network



For each training example x_i in mini-batch:

- Perform forward evaluation to compute loss for x_i
 - Compute gradient of loss w.r.t. final layer's outputs
 - Backpropagate gradient to compute gradient of loss w.r.t. all network parameters
 - Accumulate gradients (over all images in minibatch)
- Update all parameter values: $w_{\text{new}} = w_{\text{old}} - \text{learning_rate} * \text{grad}$

Recall from last class: VGG memory footprint

Calculations assume 32-bit values (image batch size = 1)

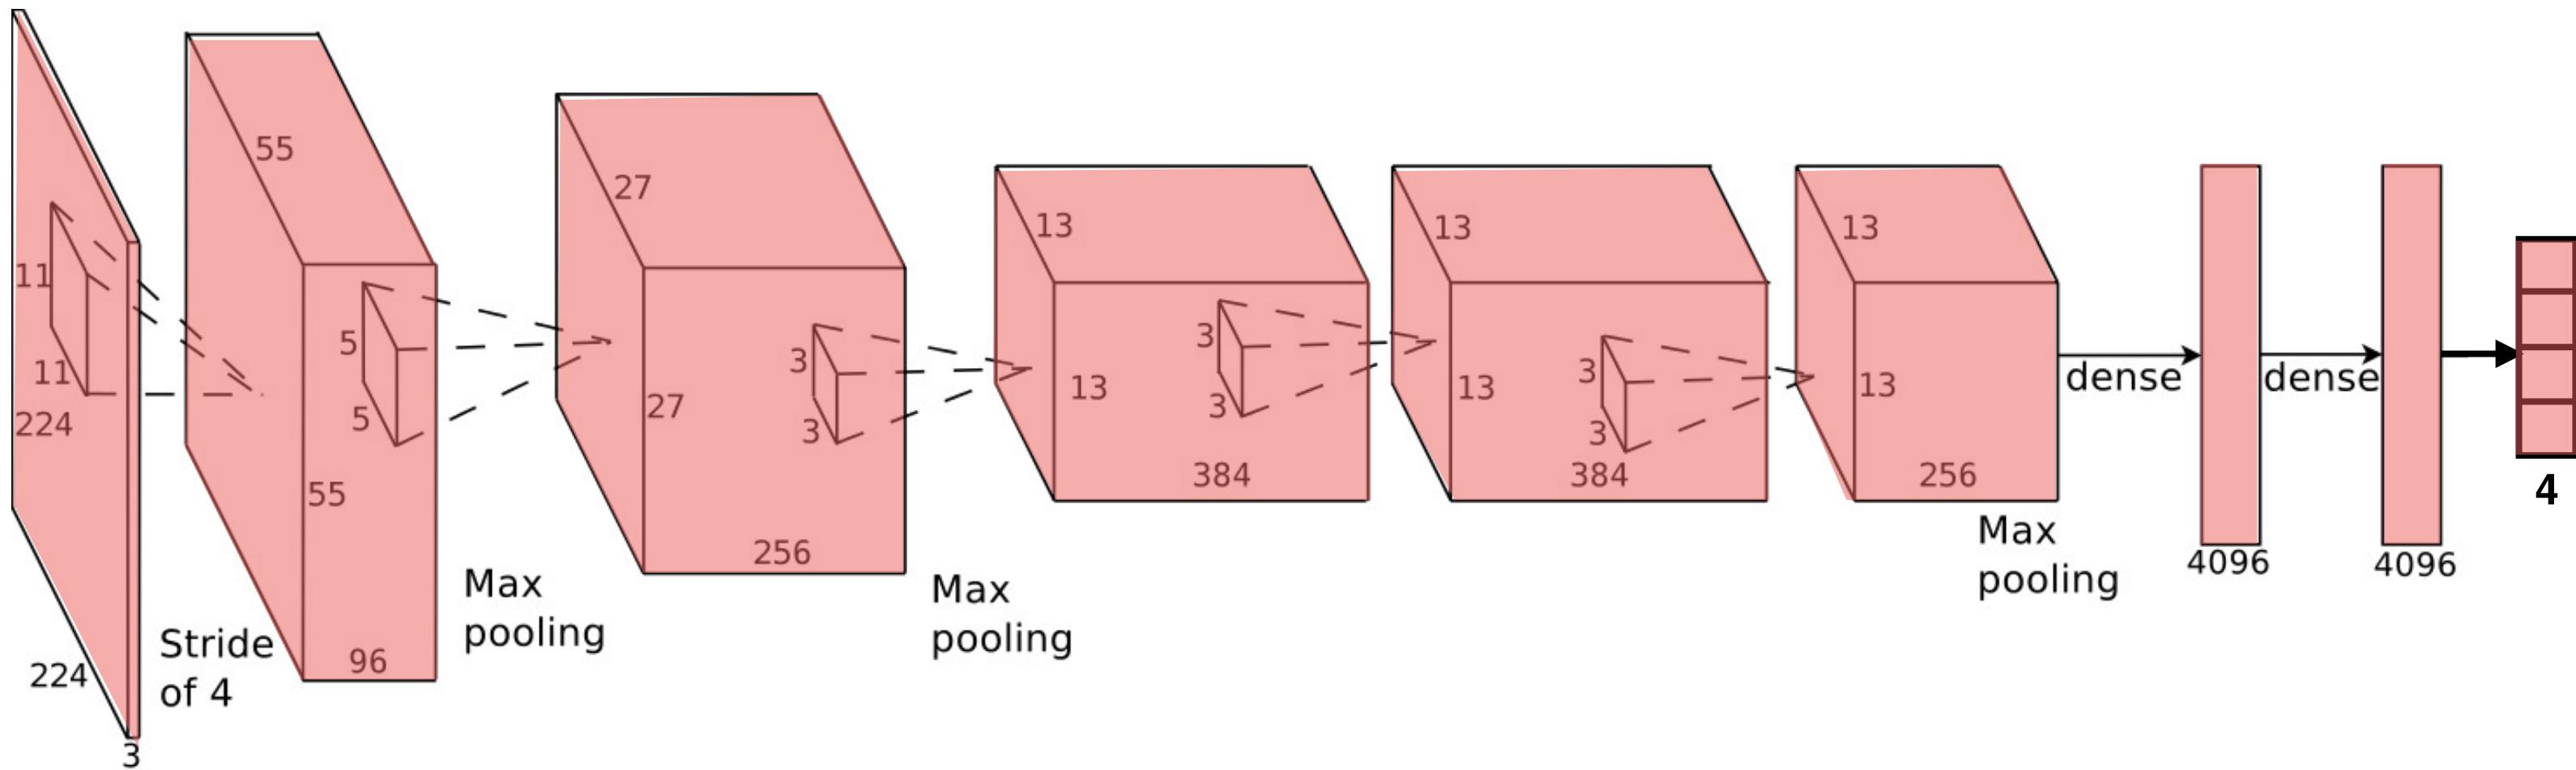
	weights mem:	output size (per image)	(mem)
input: 224 x 224 RGB image	—	224x224x3	150K
conv: (3x3x3) x 64	6.5 KB	224x224x64	12.3 MB
conv: (3x3x64) x 64	144 KB	224x224x64	12.3 MB
maxpool	—	112x112x64	3.1 MB
conv: (3x3x64) x 128	228 KB	112x112x128	6.2 MB
conv: (3x3x128) x 128	576 KB	112x112x128	6.2 MB
maxpool	—	56x56x128	1.5 MB
conv: (3x3x128) x 256	1.1 MB	56x56x256	3.1 MB
conv: (3x3x256) x 256	2.3 MB	56x56x256	3.1 MB
conv: (3x3x256) x 256	2.3 MB	56x56x256	3.1 MB
maxpool	—	28x28x256	766 KB
conv: (3x3x256) x 512	4.5 MB	28x28x512	1.5 MB
conv: (3x3x512) x 512	9 MB	28x28x512	1.5 MB
conv: (3x3x512) x 512	9 MB	28x28x512	1.5 MB
maxpool	—	14x14x512	383 KB
conv: (3x3x512) x 512	9 MB	14x14x512	383 KB
conv: (3x3x512) x 512	9 MB	14x14x512	383 KB
conv: (3x3x512) x 512	9 MB	14x14x512	383 KB
maxpool	—	7x7x512	98 KB
fully-connected 4096	392 MB	4096	16 KB
fully-connected 4096	64 MB	4096	16 KB
fully-connected 1000	15.6 MB	1000	4 KB
soft-max		1000	4 KB

Many weights in fully-connected players

Storing convolution layer outputs (unit “activations”) can get big in early layers with large input size and many filters

Note: multiply these numbers by N for batch size of N images

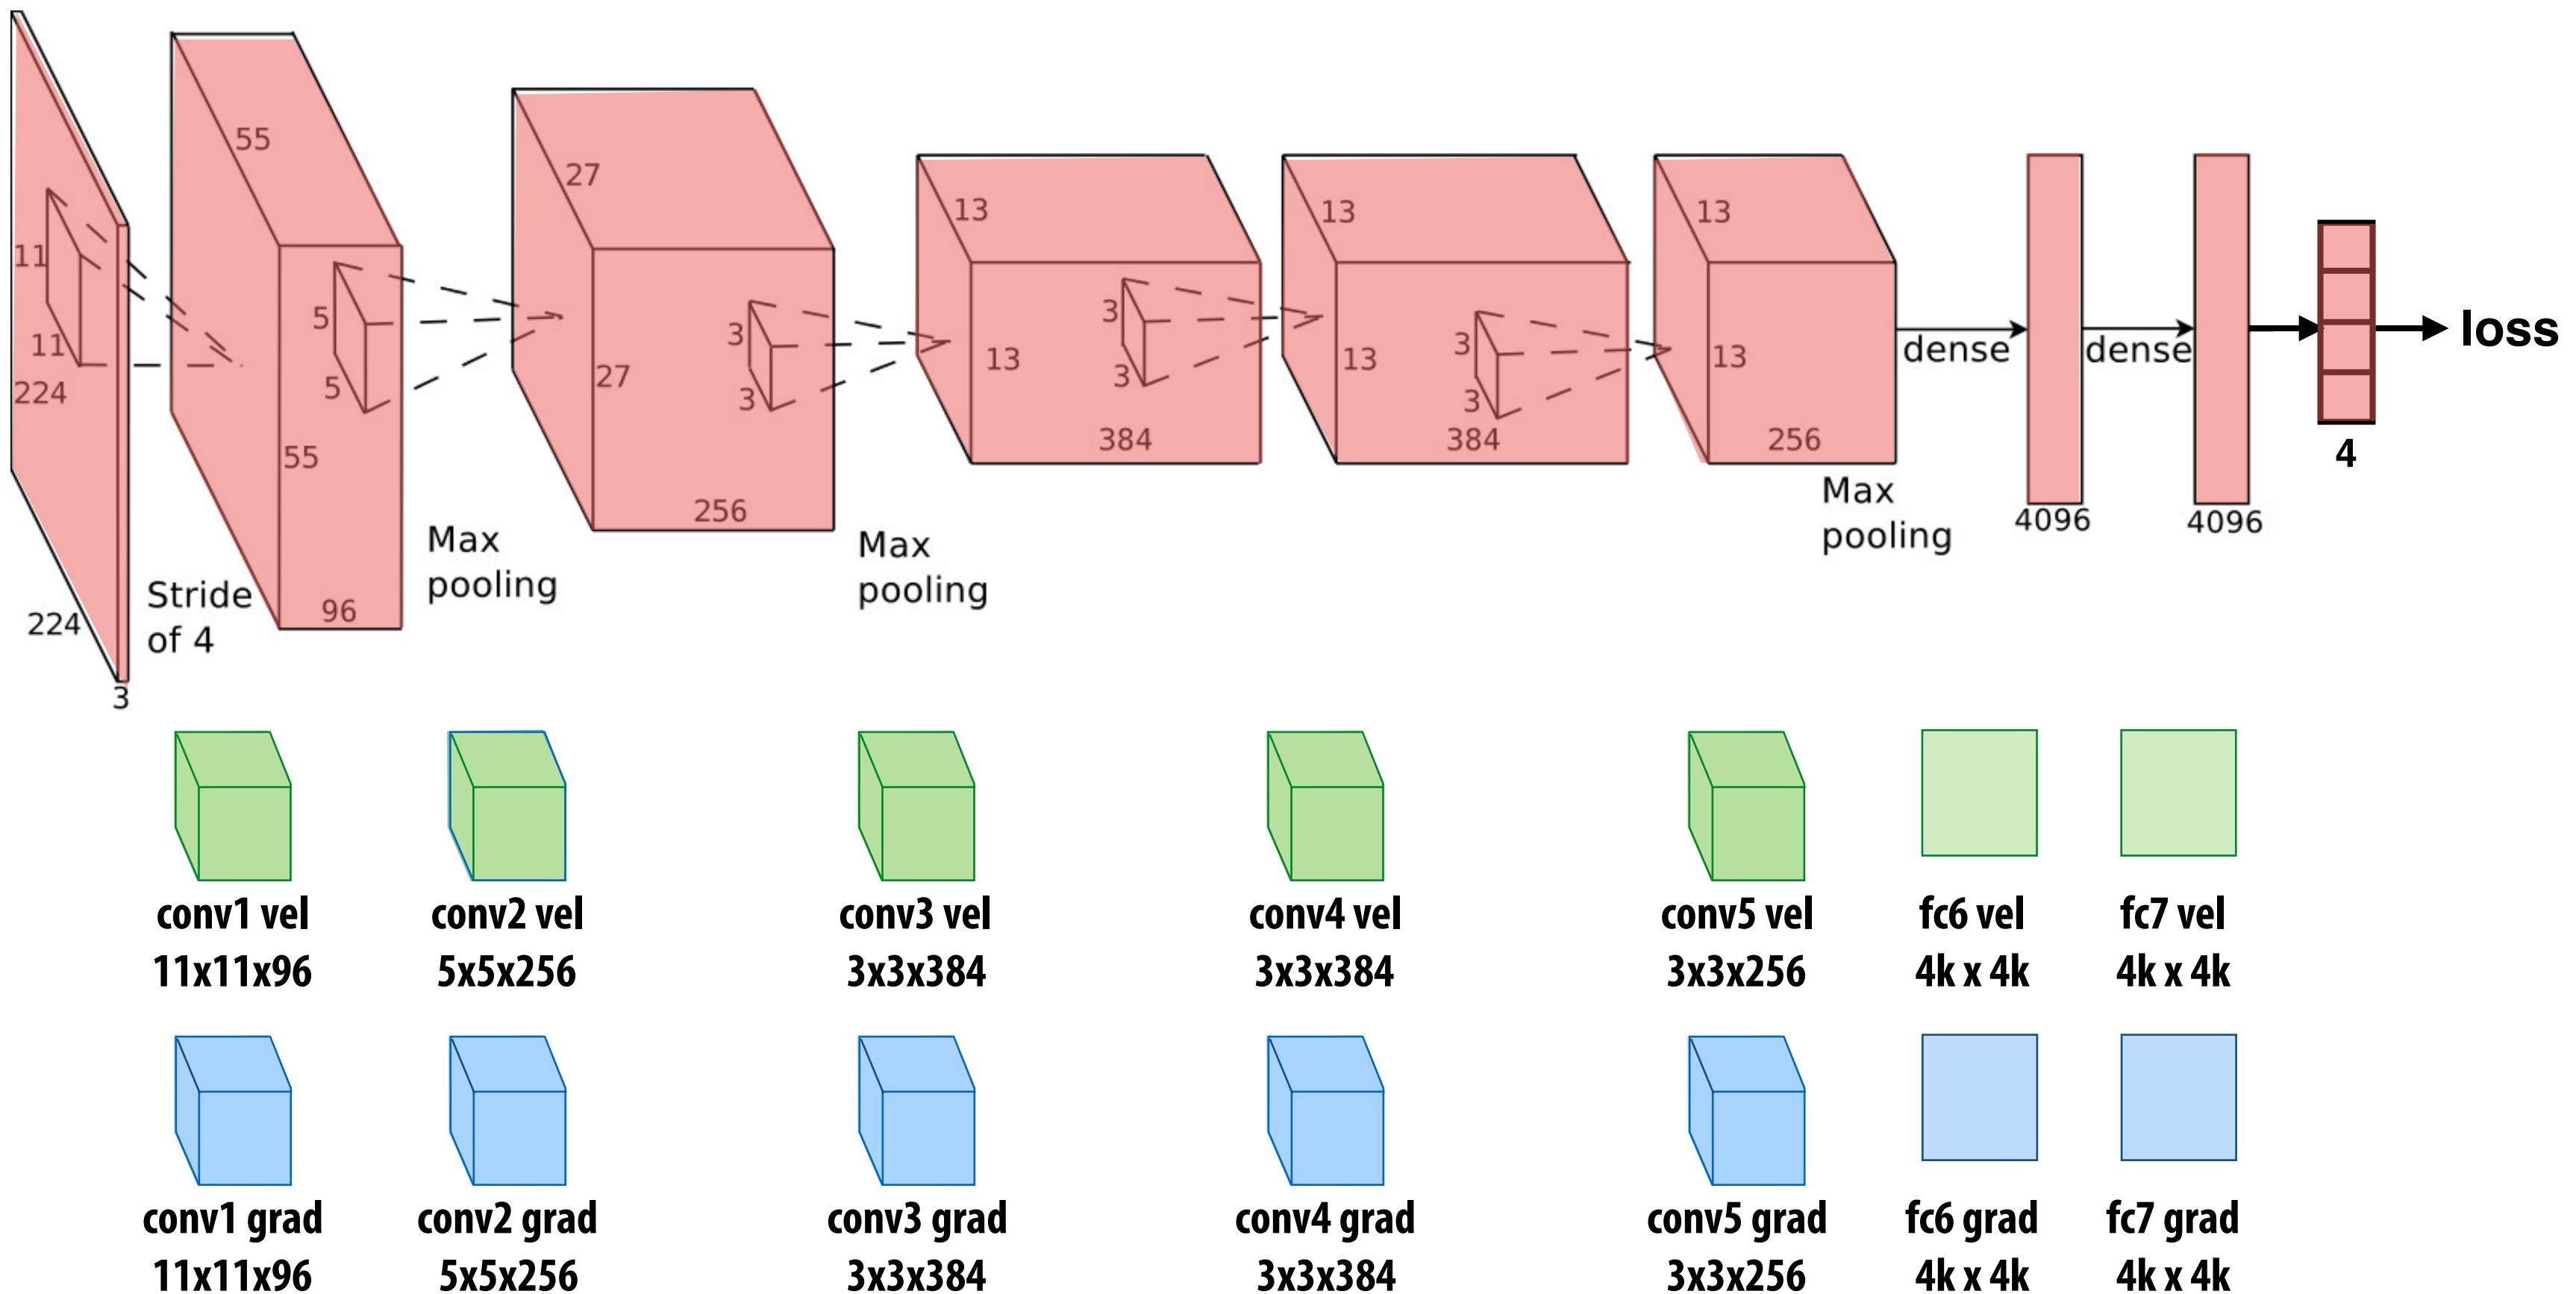
Data lifetimes during network evaluation



Weights (read-only) reside in memory

After evaluating layer i , can free outputs from layer $i-1$

Data lifetimes during training



- Must retain outputs for all layers because they are needed to compute gradients during back-prop
- Parallel back-prop will require storage for per-weight gradients (more about this in a second)
- In practice: may also store per-weight gradient velocity (if using SGD with "momentum") or step size cache in adaptive step size schemes like Adagrad

$$\text{vel_new} = \mu * \text{vel_old} - \text{step_size} * \text{grad}$$

$$\text{w_new} = \text{w_old} + \text{vel_new}$$

VGG memory footprint

Calculations assume 32-bit values (image batch size = 1)

input: 224 x 224 RGB image

conv: (3x3x3) x 64

conv: (3x3x64) x 64

maxpool

conv: (3x3x64) x 128

conv: (3x3x128) x 128

maxpool

conv: (3x3x128) x 256

conv: (3x3x256) x 256

conv: (3x3x256) x 256

maxpool

conv: (3x3x256) x 512

conv: (3x3x512) x 512

conv: (3x3x512) x 512

maxpool

conv: (3x3x512) x 512

conv: (3x3x512) x 512

conv: (3x3x512) x 512

maxpool

fully-connected 4096

fully-connected 4096

fully-connected 1000

soft-max

weights mem:

—

6.5 KB

144 KB

—

228 KB

576 KB

—

1.1 MB

2.3 MB

2.3 MB

—

4.5 MB

9 MB

9 MB

—

9 MB

9 MB

9 MB

—

392 MB

64 MB

15.6 MB

inputs/outputs get multiplied by mini-batch size

output size
(per image)

224x224x3

224x224x64

224x224x64

112x112x64

112x112x128

112x112x128

56x56x128

56x56x256

56x56x256

56x56x256

28x28x256

28x28x512

28x28x512

28x28x512

14x14x512

14x14x512

14x14x512

14x14x512

7x7x512

4096

4096

1000

1000

Unlike forward evaluation:
cannot immediately free
outputs once consumed by
next level of network

Must also store per-weight gradients

Many implementations
also store gradient
“momentum” as well
(multiply by 3)

(mem)

150K

12.3 MB

12.3 MB

3.1 MB

6.2 MB

6.2 MB

1.5 MB

3.1 MB

3.1 MB

3.1 MB

766 KB

1.5 MB

1.5 MB

1.5 MB

383 KB

383 KB

383 KB

383 KB

98 KB

16 KB

16 KB

4 KB

4 KB

SGD workload

```
while (loss too high): ←
```

At first glance, this loop is sequential (each step of “walking downhill” depends on previous)

```
for each item x_i in mini-batch: ← Parallel across images  
    grad += evaluate_loss_gradient(f, loss_func, params, x_i)
```

↑
sum reduction

large computation with its own parallelism
(but working set may not fit on single machine)

```
params += -grad * step_size;
```

← trivial data-parallel over parameters

DNN training workload

■ Large computational expense

- Must evaluate the network (forward and backward) for millions of training images
- Must iterate for many iterations of gradient descent (100's of thousands)
- Training modern networks on big datasets takes days

■ Large memory footprint

- Must maintain network layer outputs from forward pass
- Additional memory to store gradients/gradient velocity for each parameter
- Recall parameters for popular VGG-16 network require ~500 MB of memory (training requires GBs of memory for academic networks)
- Scaling to larger networks requires partitioning DNN across nodes to keep DNN + intermediates in memory

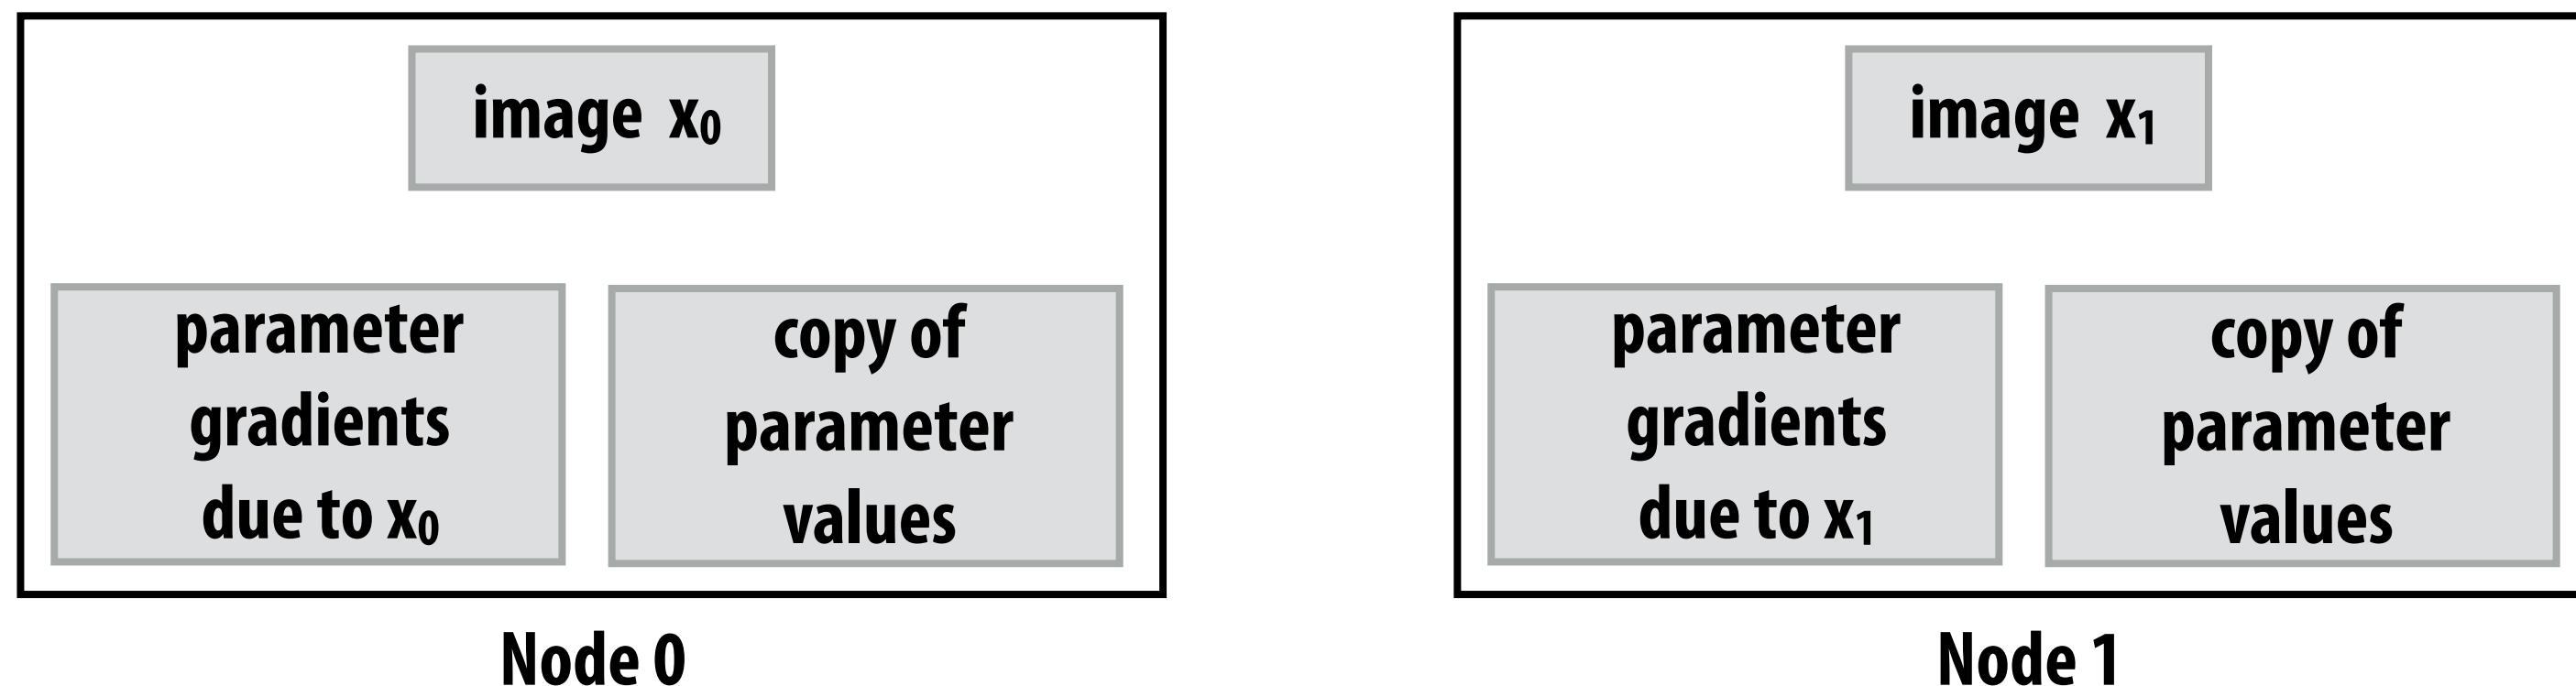
■ Dependencies /synchronization (not embarrassingly parallel)

- Each parameter update step depends on previous
- Many units contribute to same parameter gradients (fine-scale reduction)
- Different images in mini batch contribute to same parameter gradients

Synchronous data-parallel training (across images)

```
for each item x_i in mini-batch:  
    grad += evaluate_loss_gradient(f, loss_func, params, x_i)  
params += -grad * learning_rate;
```

Consider parallelization of the outer for loop across machines in a cluster



```
partition dataset across nodes  
for each item x_i in mini-batch assigned to local node:  
    // just like single node training  
    grad += evaluate_loss_gradient(f, loss_func, params, x_i)  
barrier();  
sum reduce gradients, communicate results to all nodes  
barrier();  
update copy of parameter values
```

Synchronous training

- All nodes cooperate to compute gradients for a mini-batch *
- Gradients are summed (across the entire machine)
 - All-to-all communication (e.g., MPI_Allreduce)
 - Good implementations will sum gradients for layer i when computing backprop for $i+1$ (overlap communication and computation).
- Update model parameters
 - Typically done without wide parallelism (e.g. each machine computes its own update)
- All nodes proceed to work on next mini-batch given new model parameters

* If curious about batch norm in a parallel training setting. In practice each of k nodes works on a set of n images, with batch norm statistics computed independently for each set of n (mini-batch size is kn).

Challenges of scaling out (many nodes)

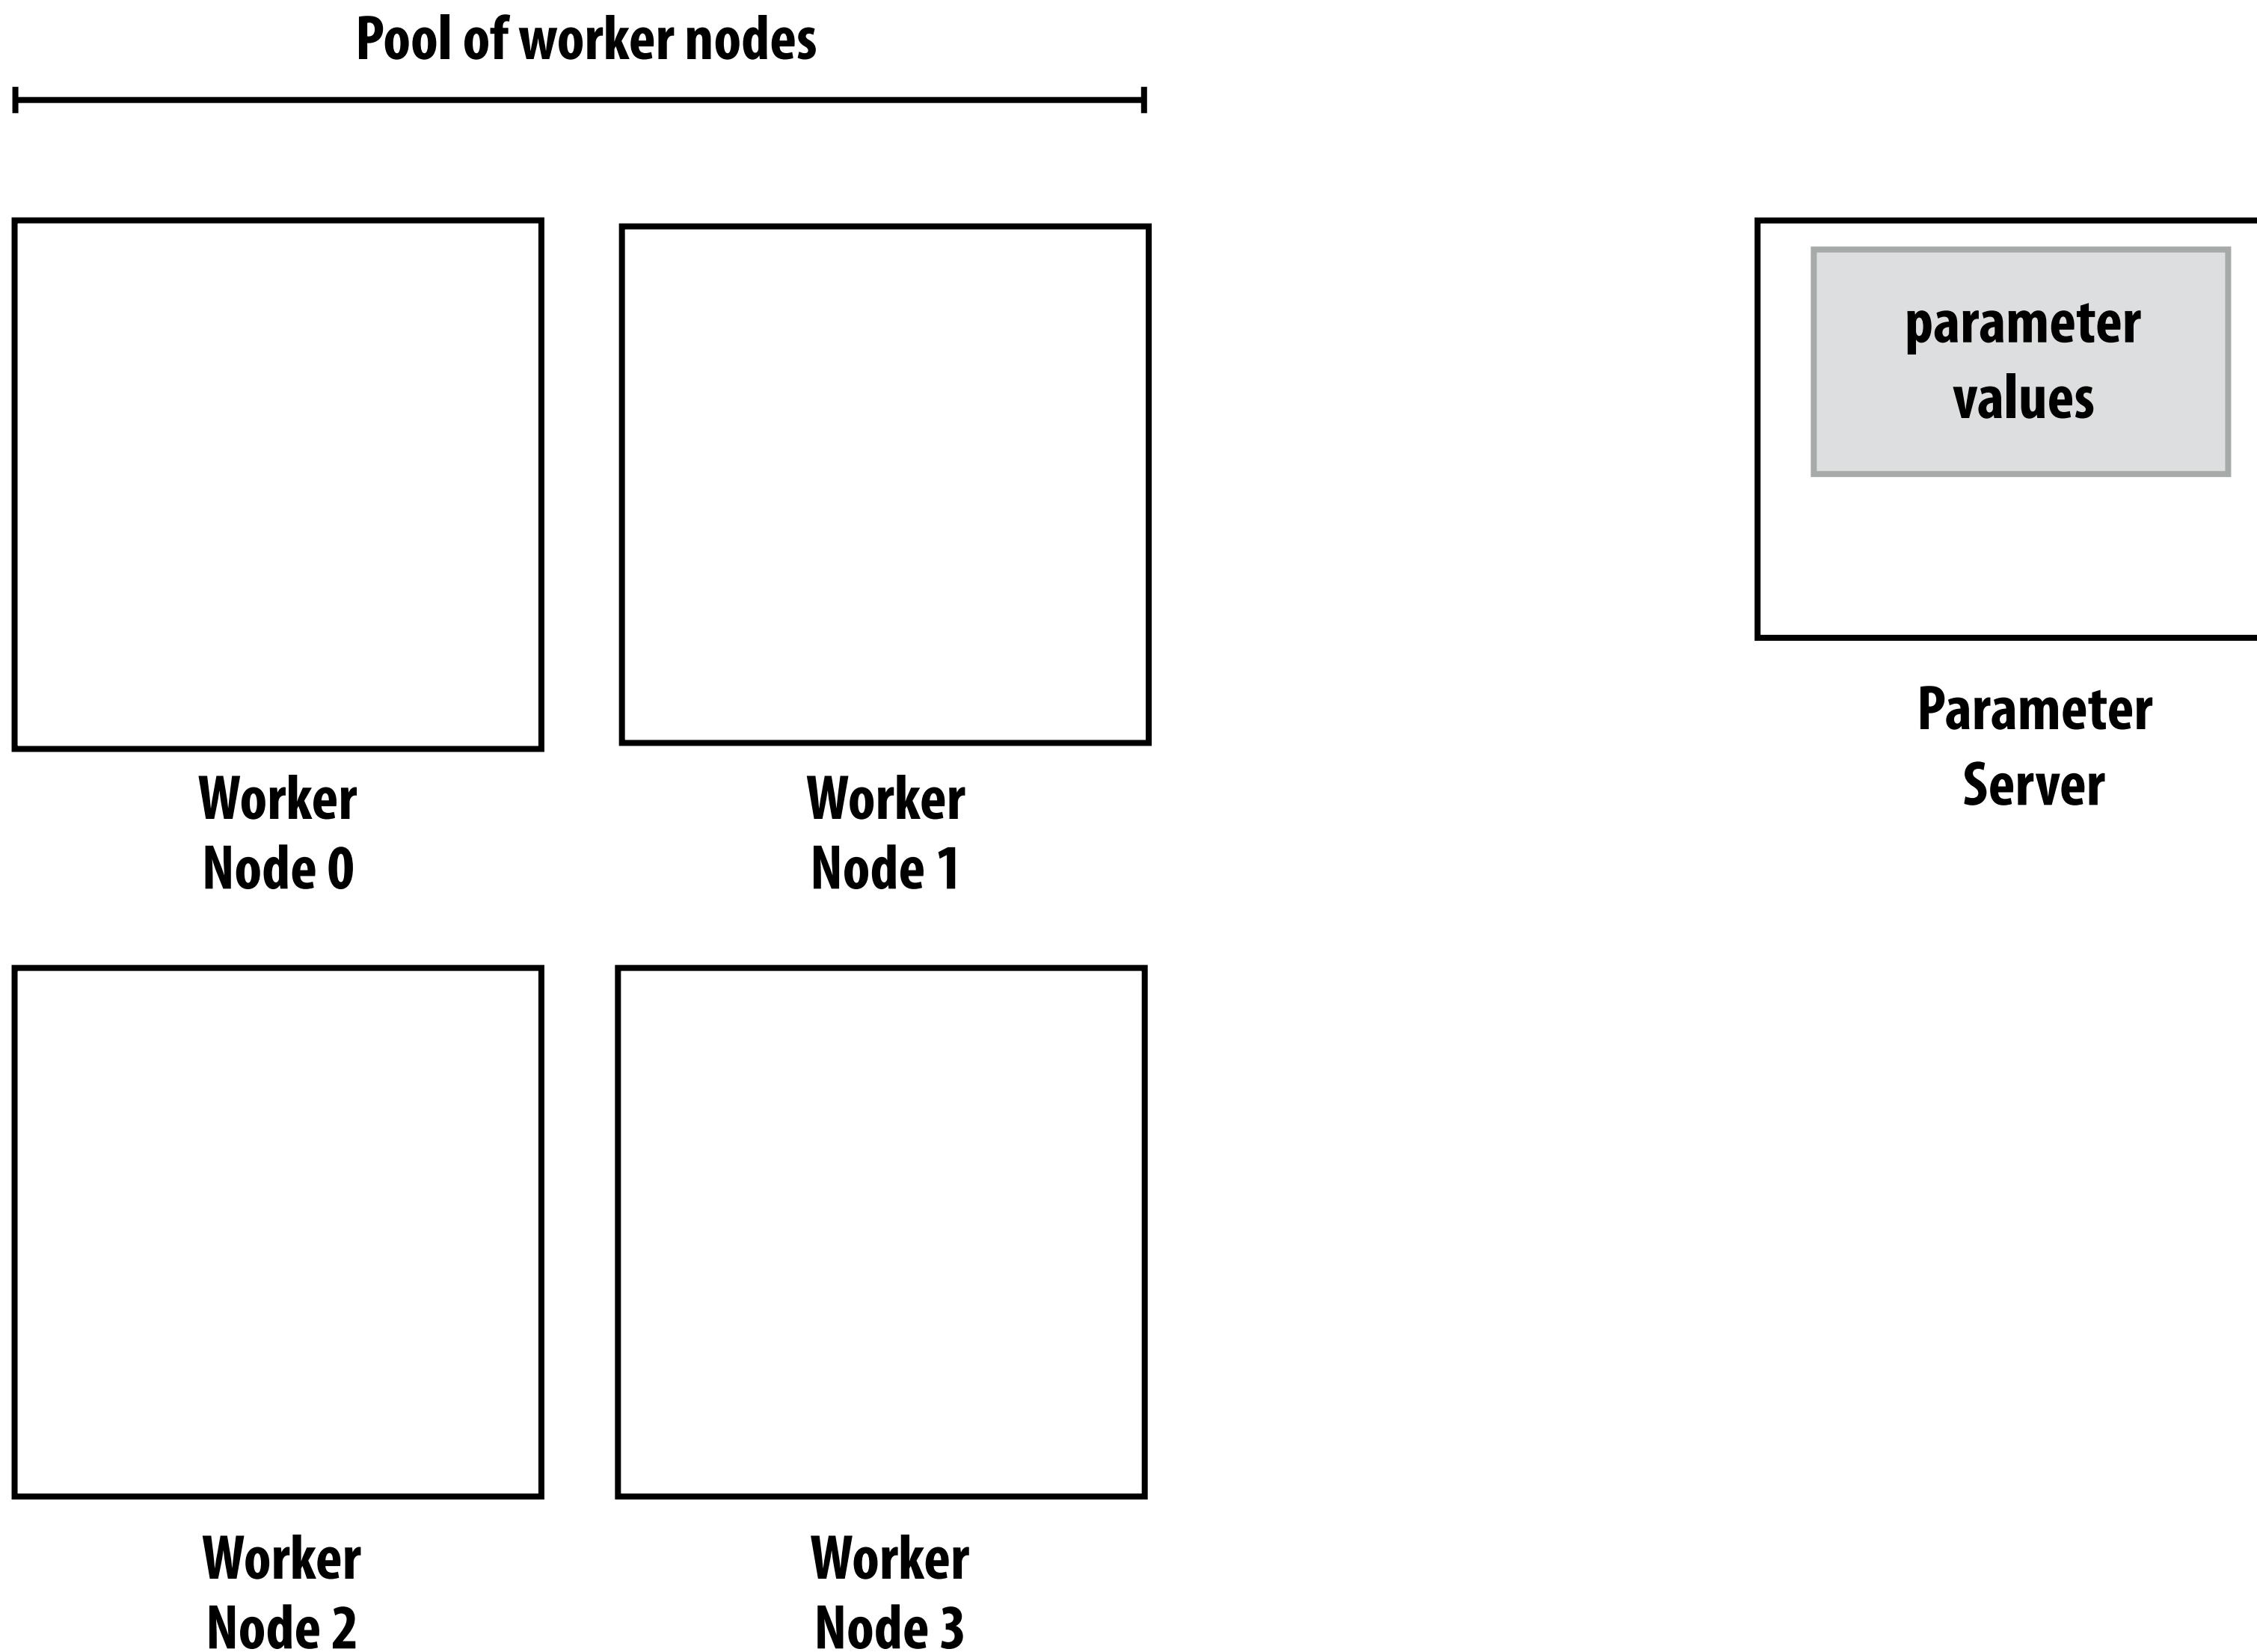
- Slow communication between nodes
 - Commodity clusters do not feature high-performance interconnects (e.g., infiniband) typical of supercomputers
 - Synchronous SGD involves all to all communication after each minibatch

- Nodes with different performance (even if machines are the same)
 - Workload imbalance at barriers (sync points between nodes)

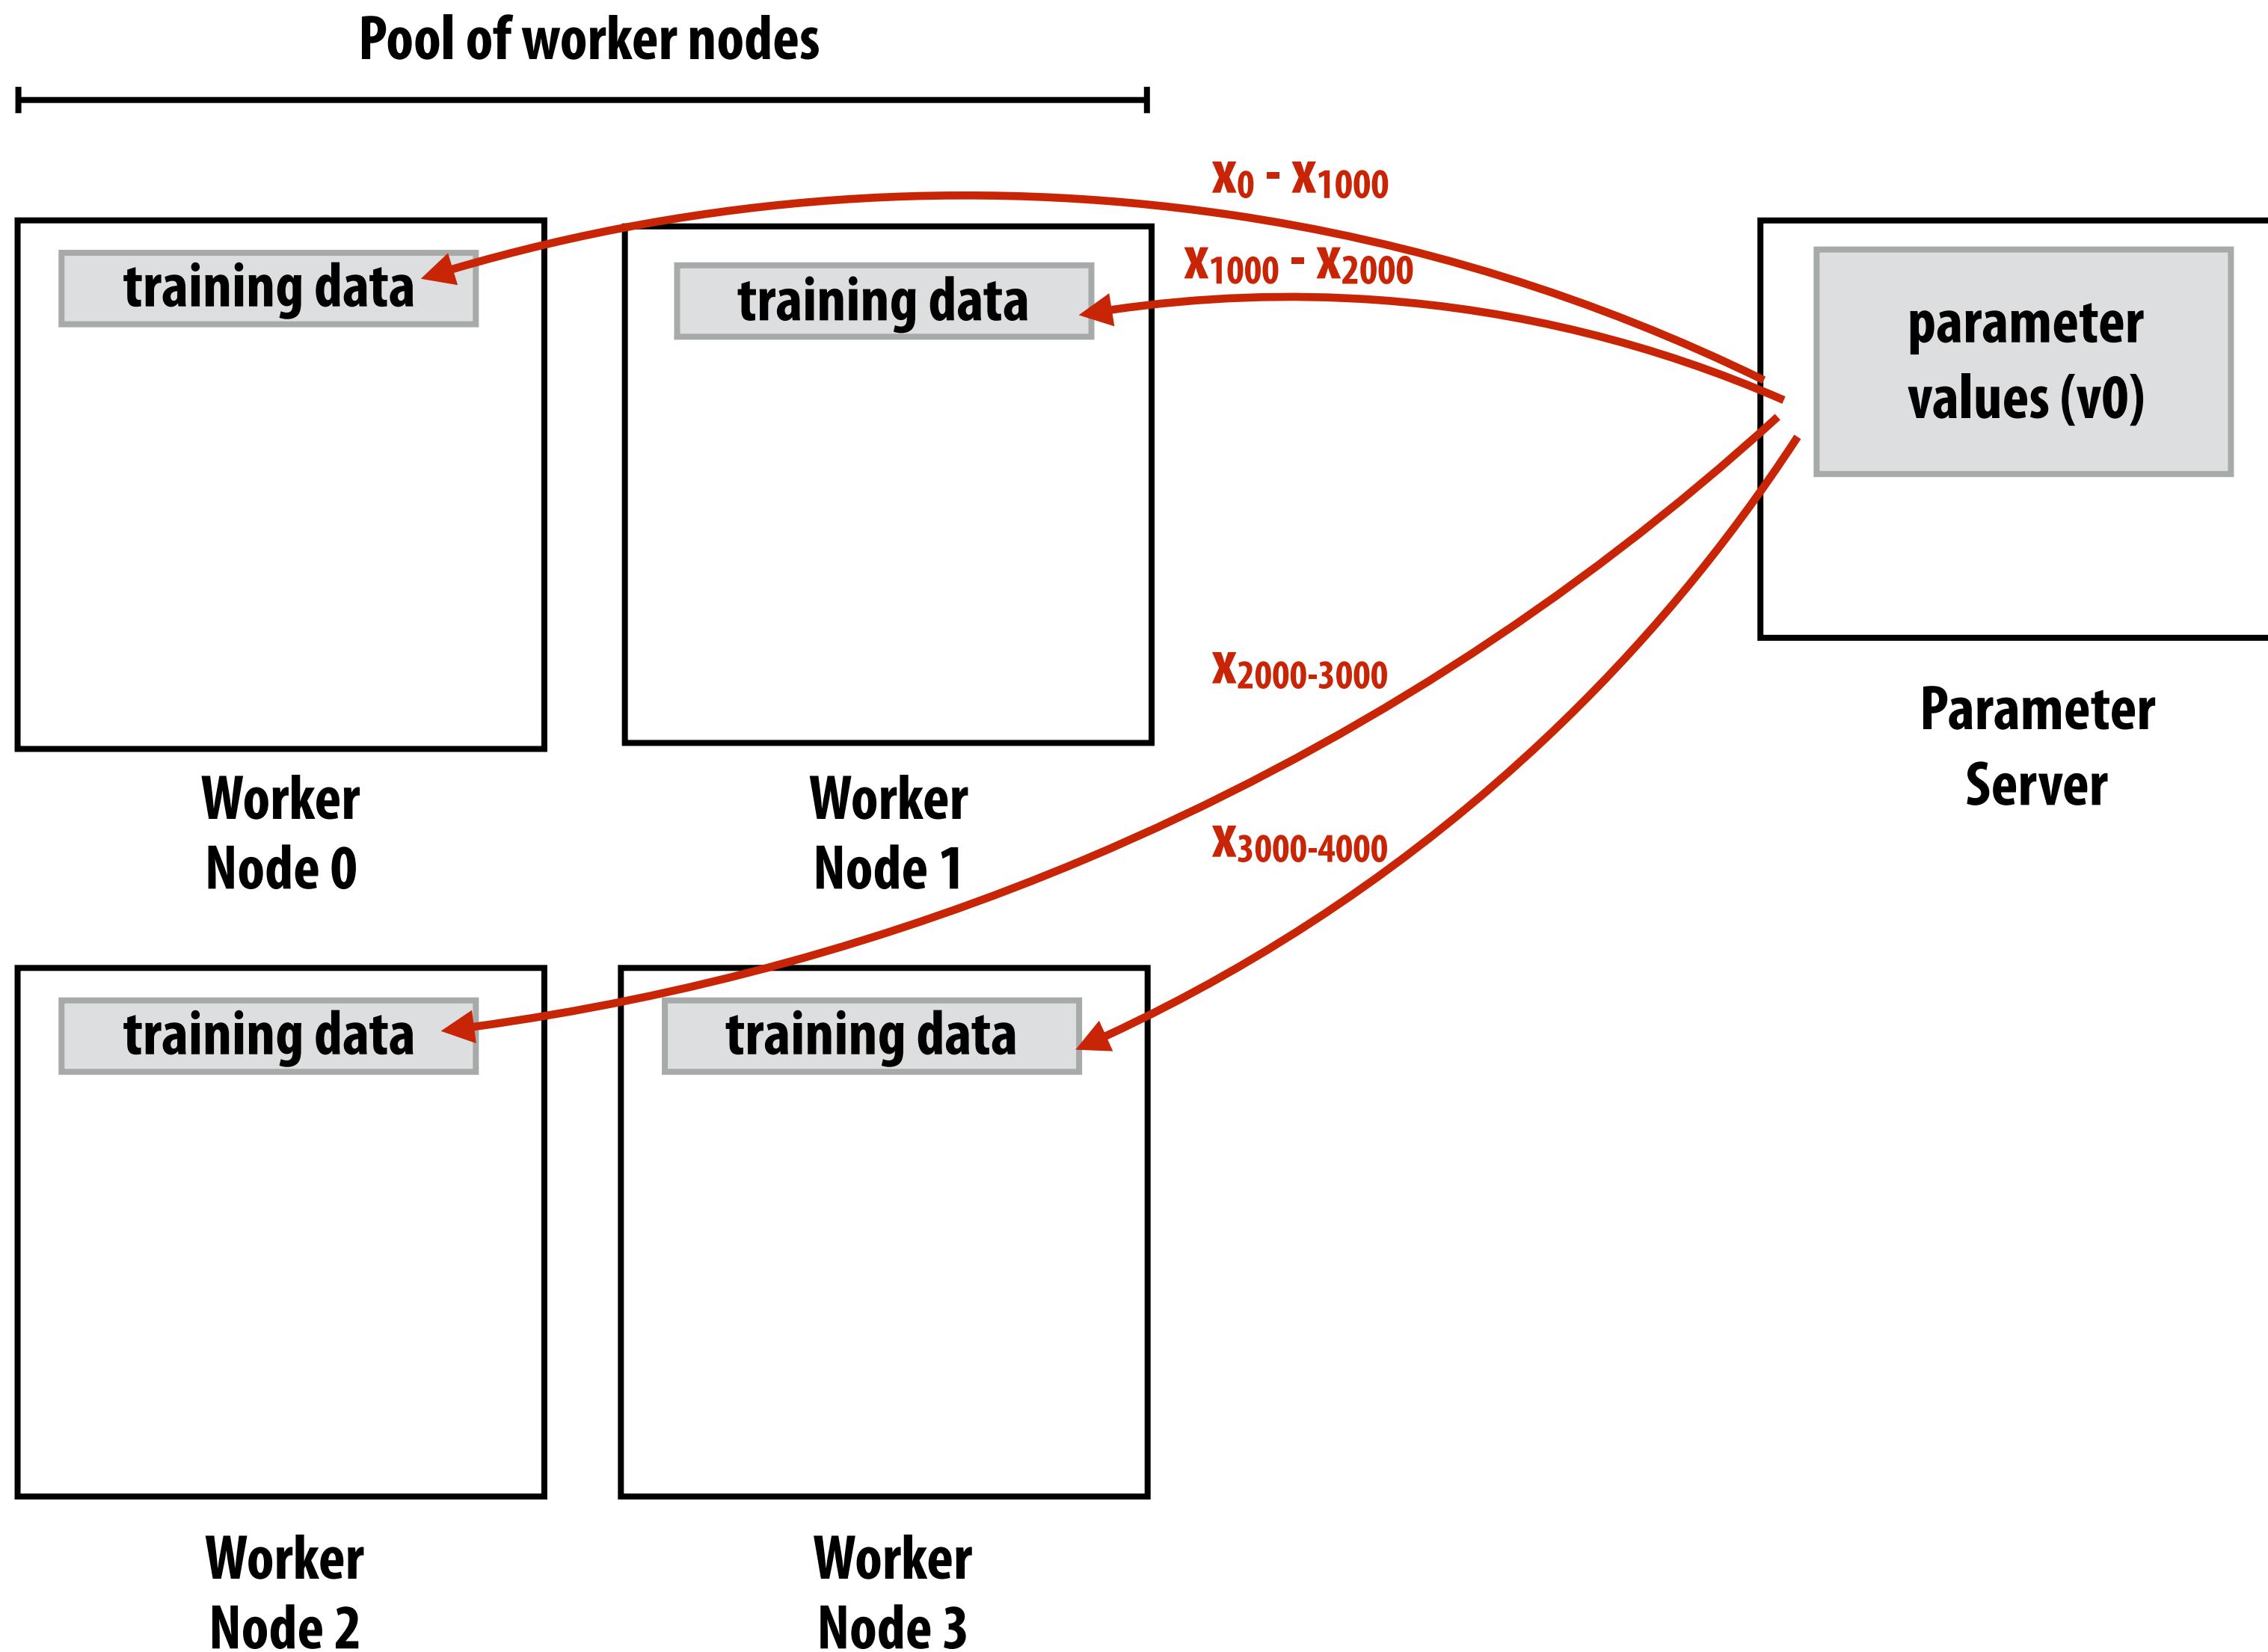
Alternative solution: exploit properties of SGD by using asynchronous execution

Parameter server design

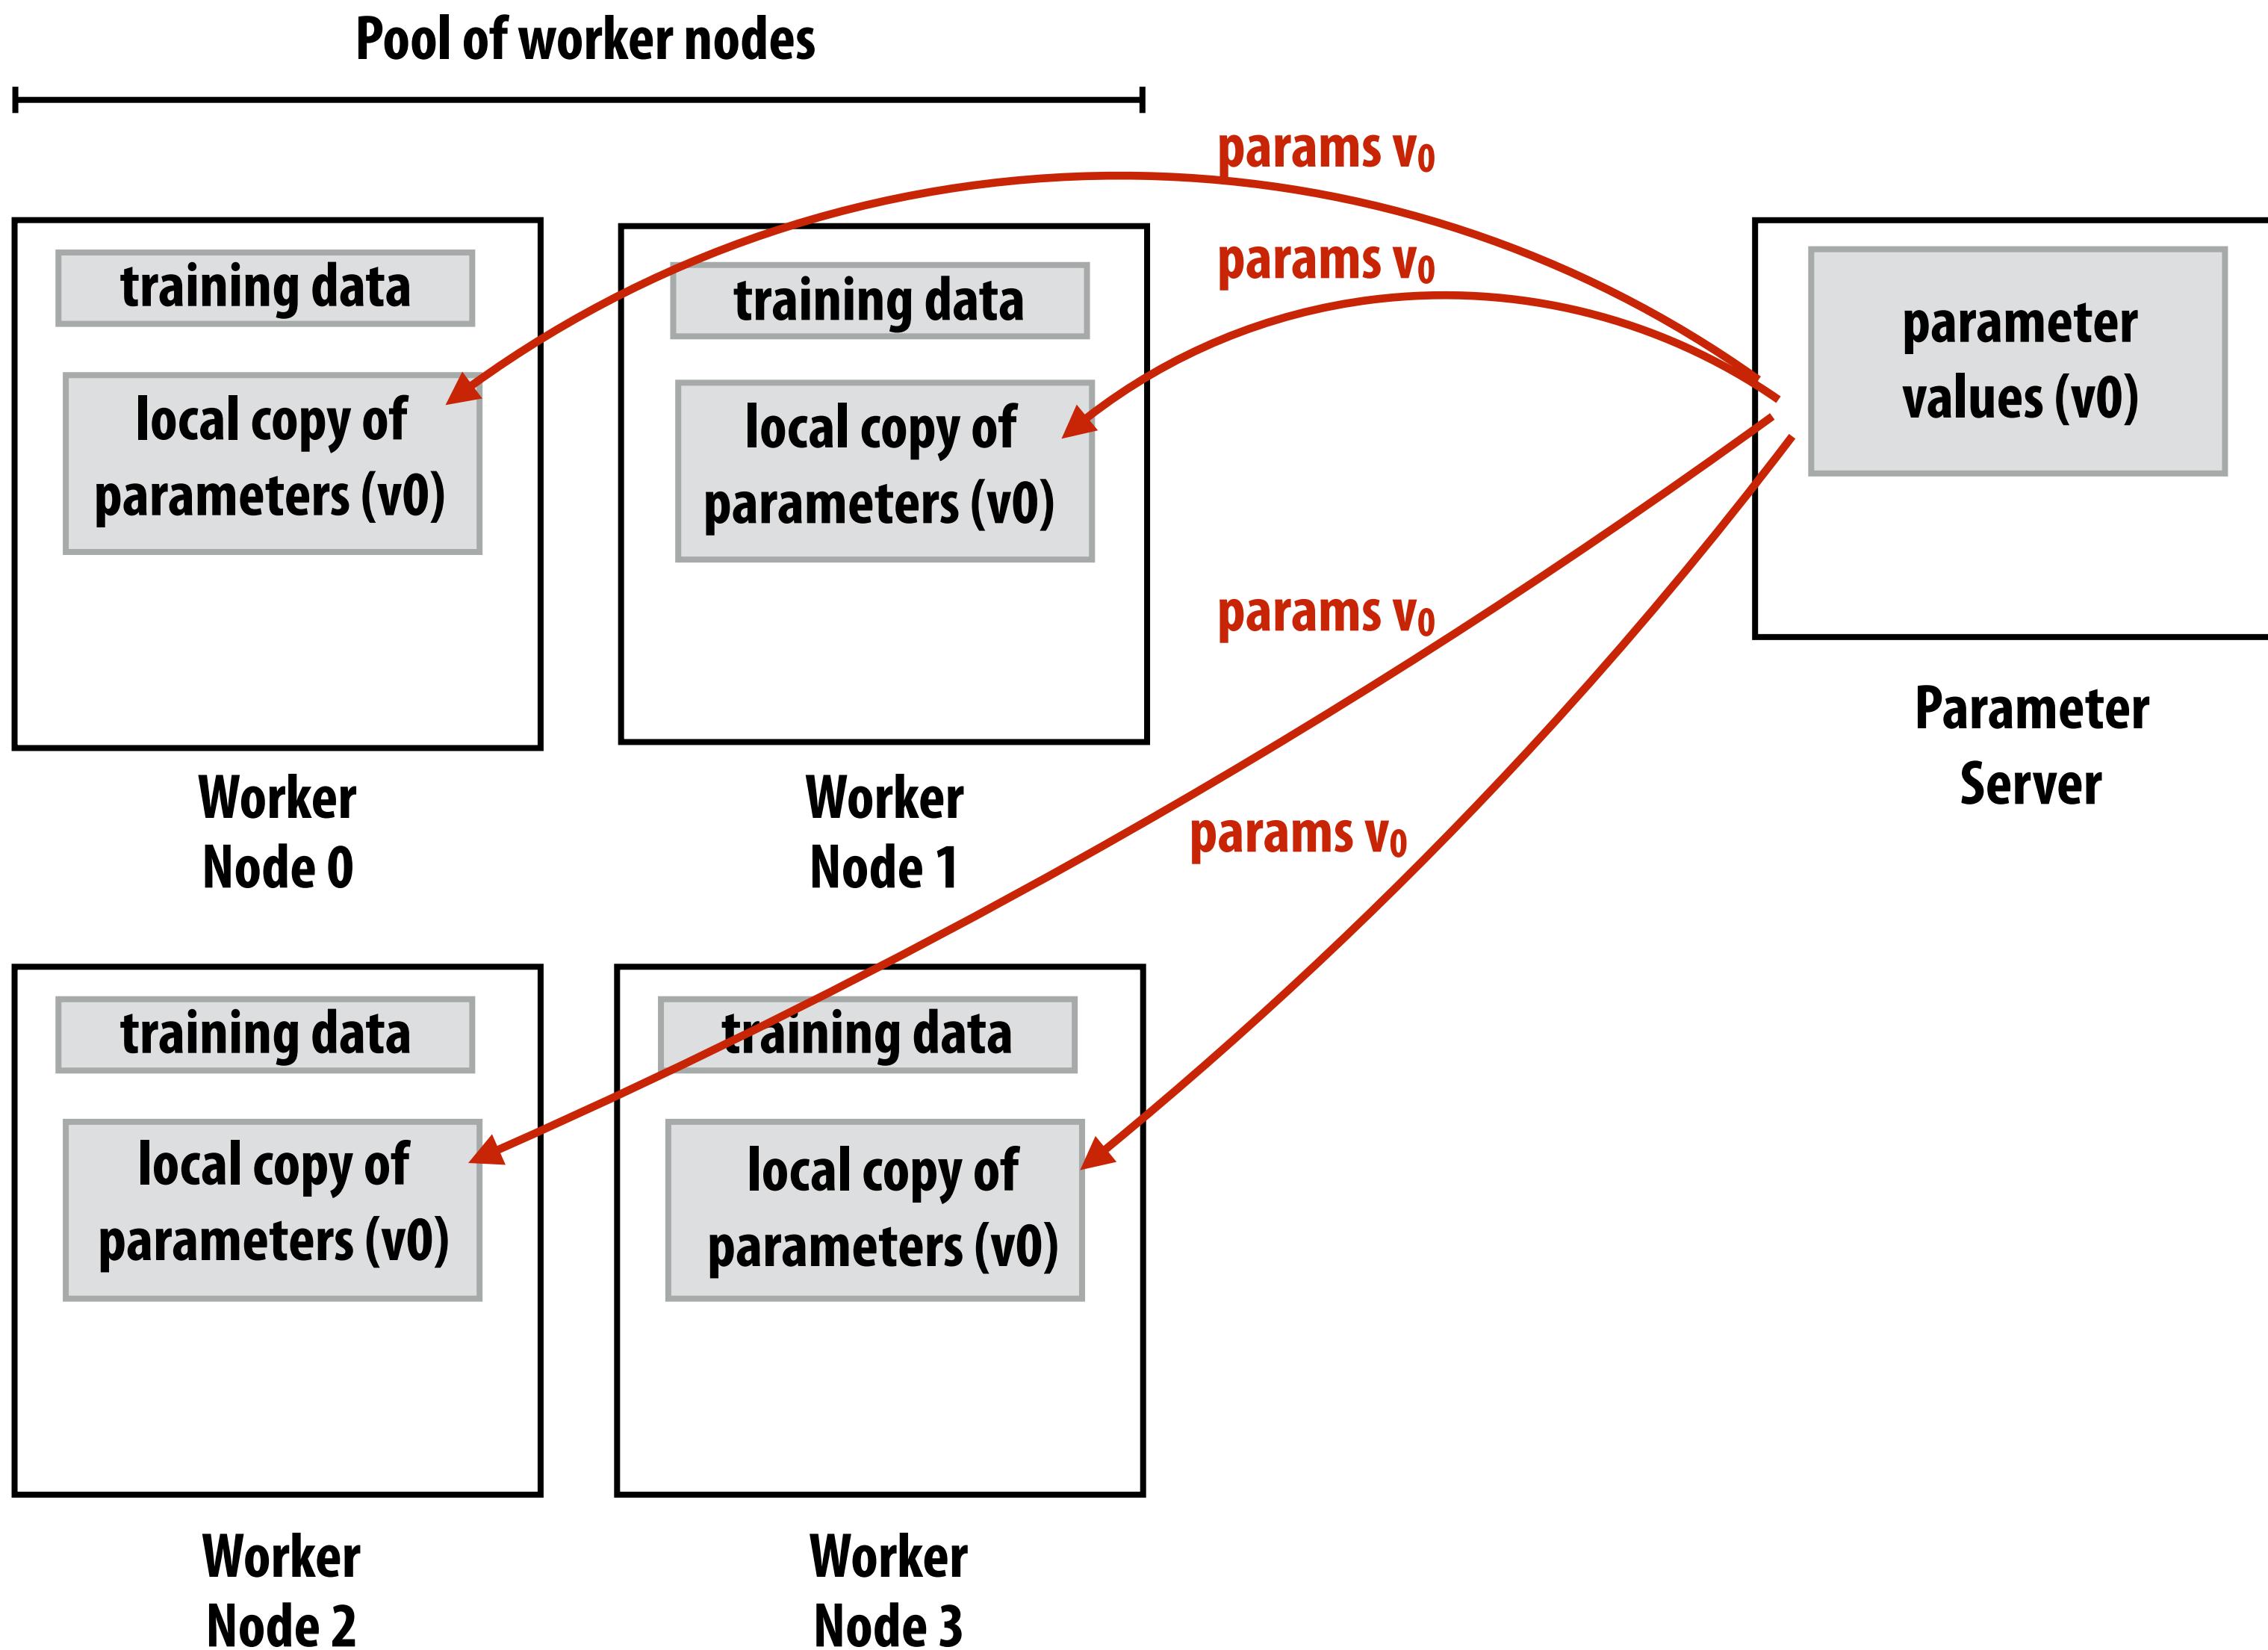
Parameter Server [Li OSDI14]
Google's DistBelief [Dean NIPS12]
Microsoft's Project Adam [Chilimbi OSDI14]



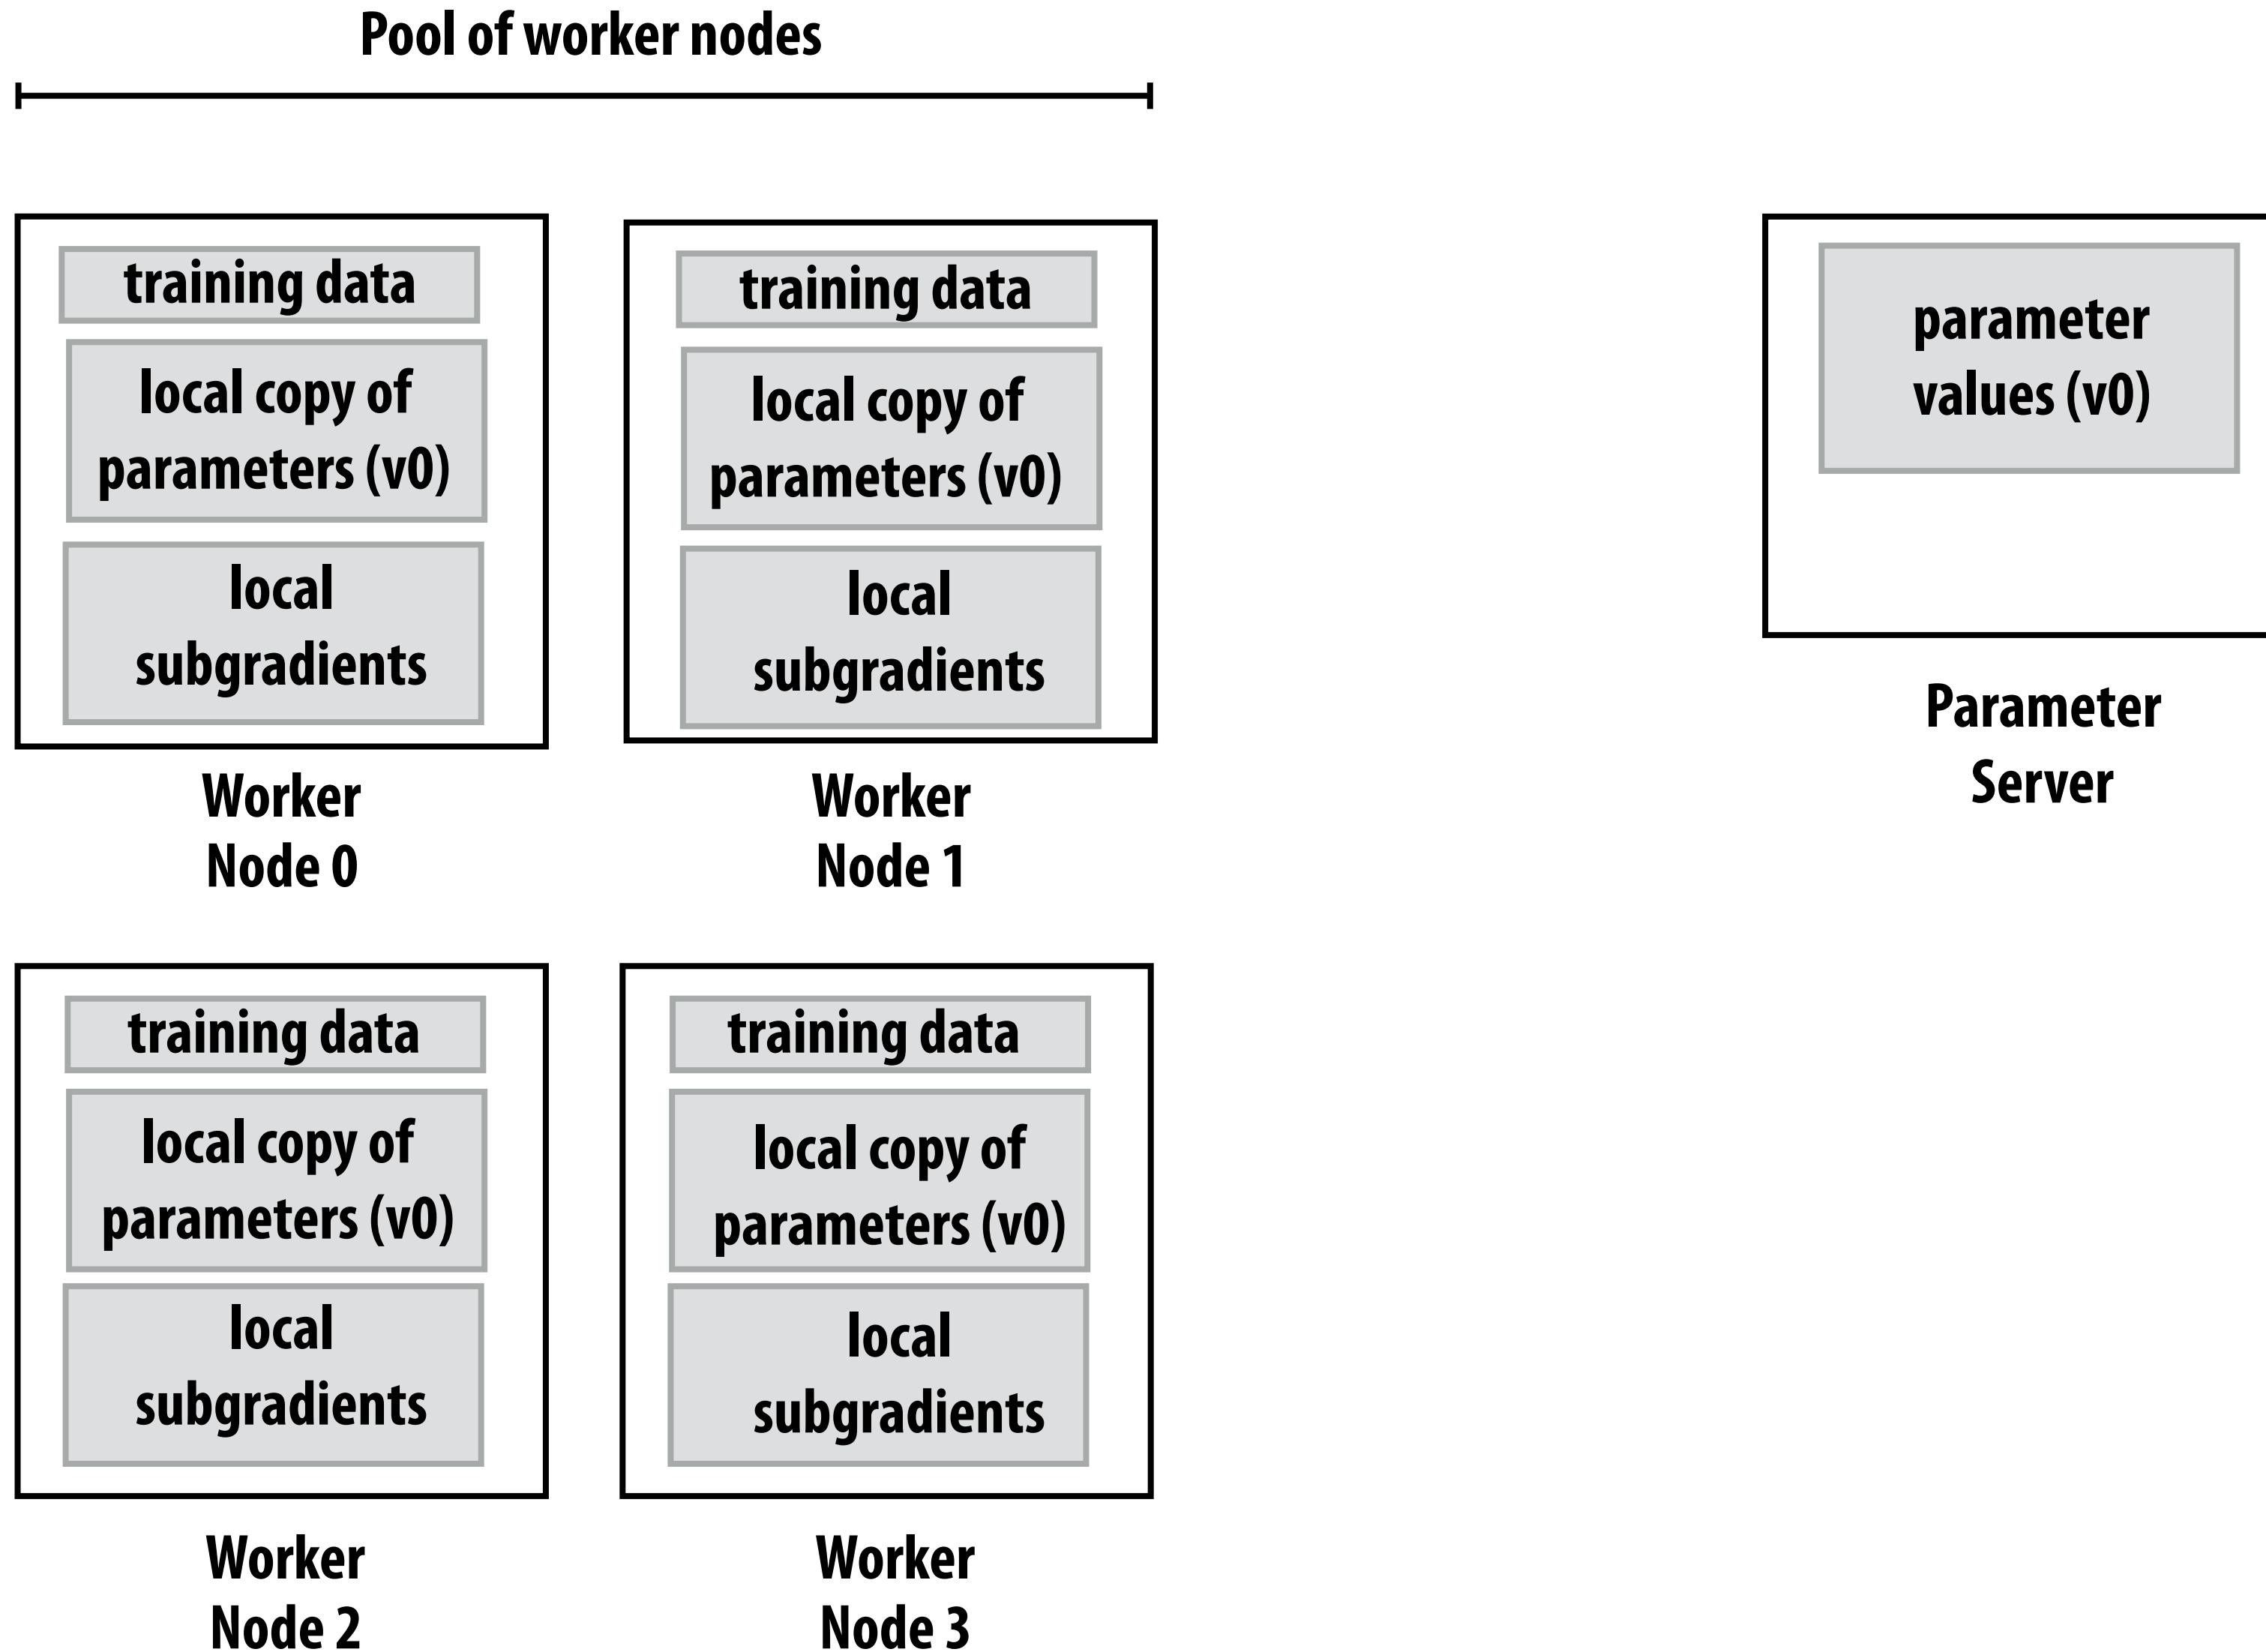
Training data partitioned among workers



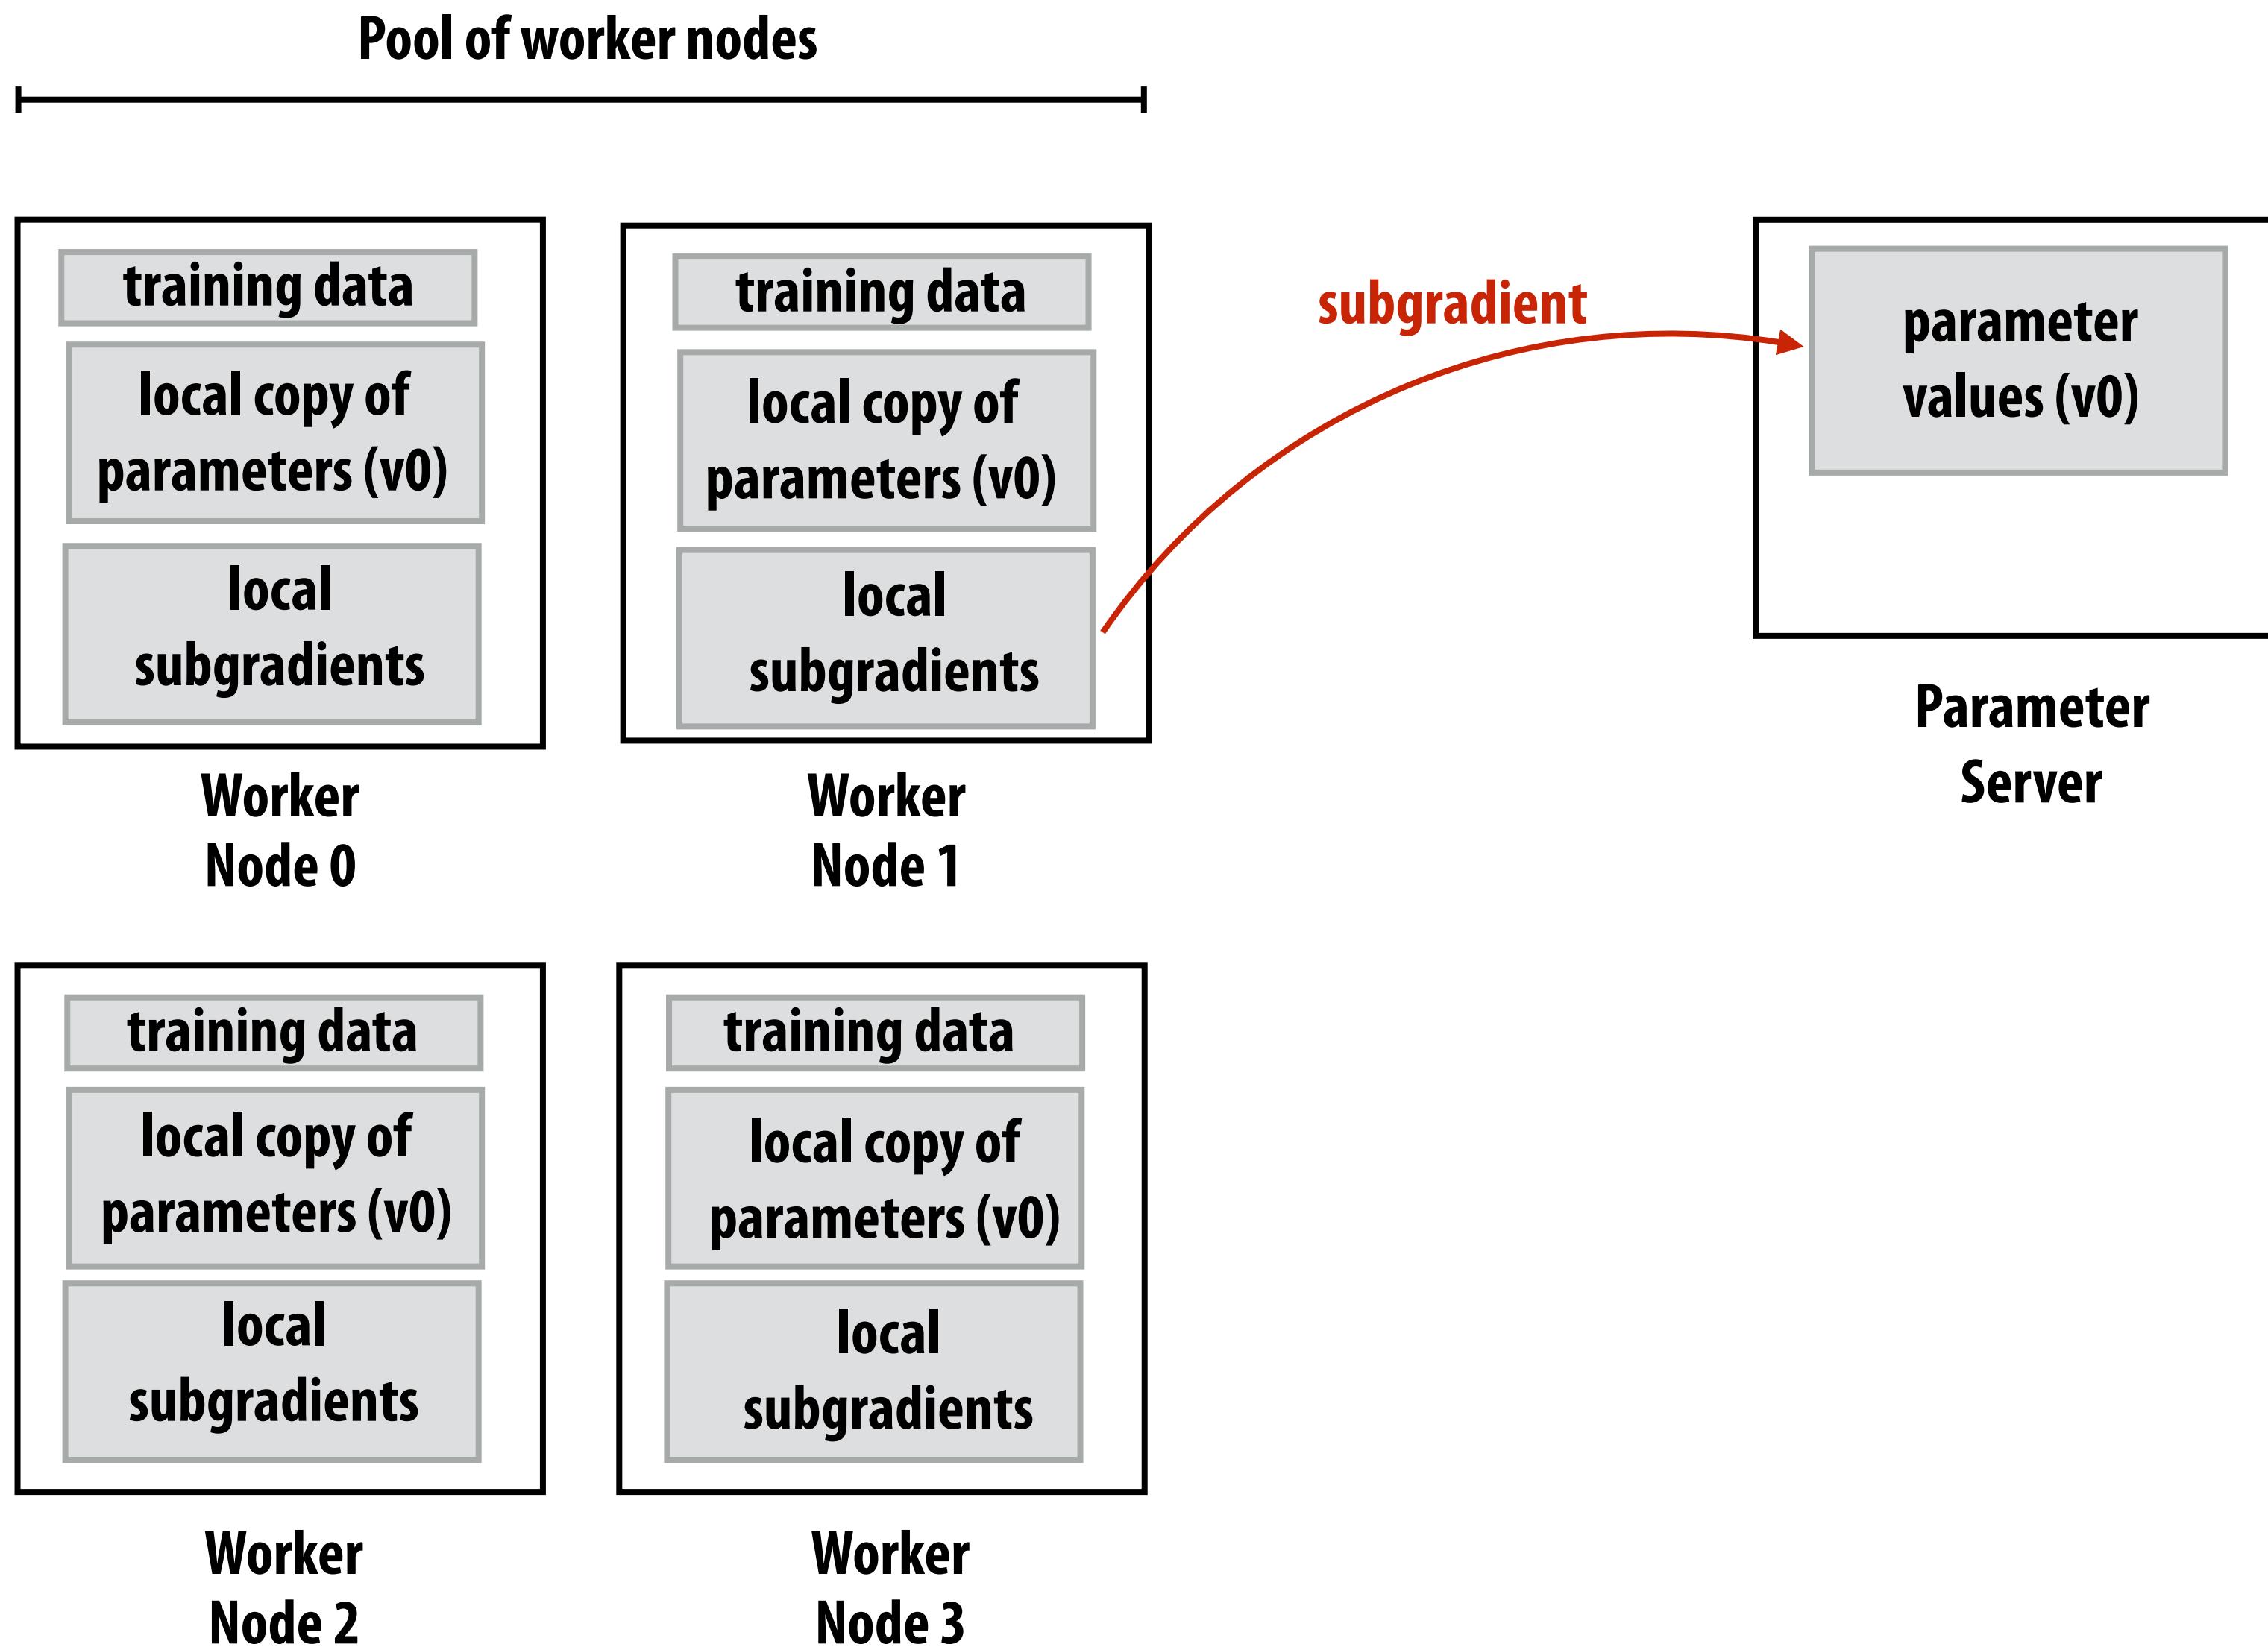
Copy of parameters sent to workers



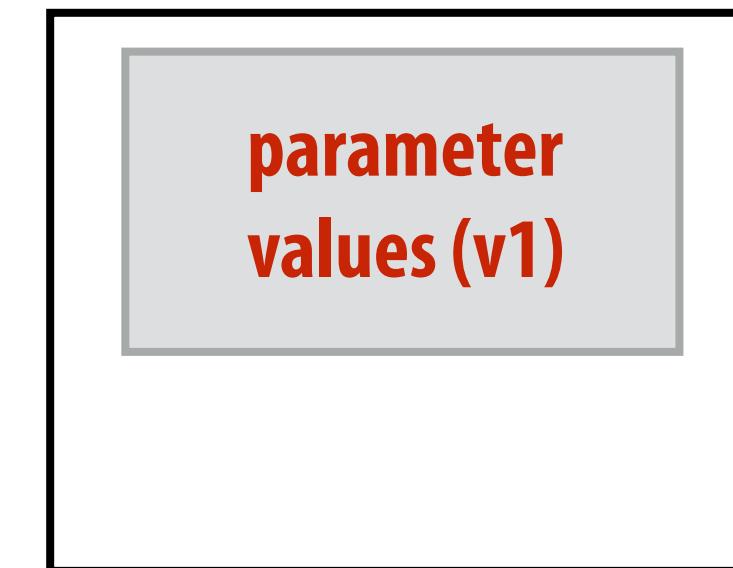
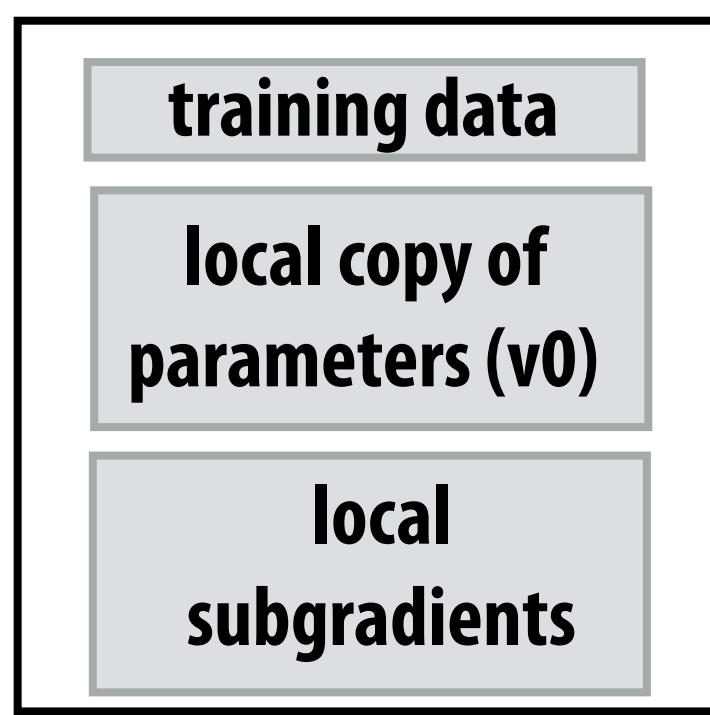
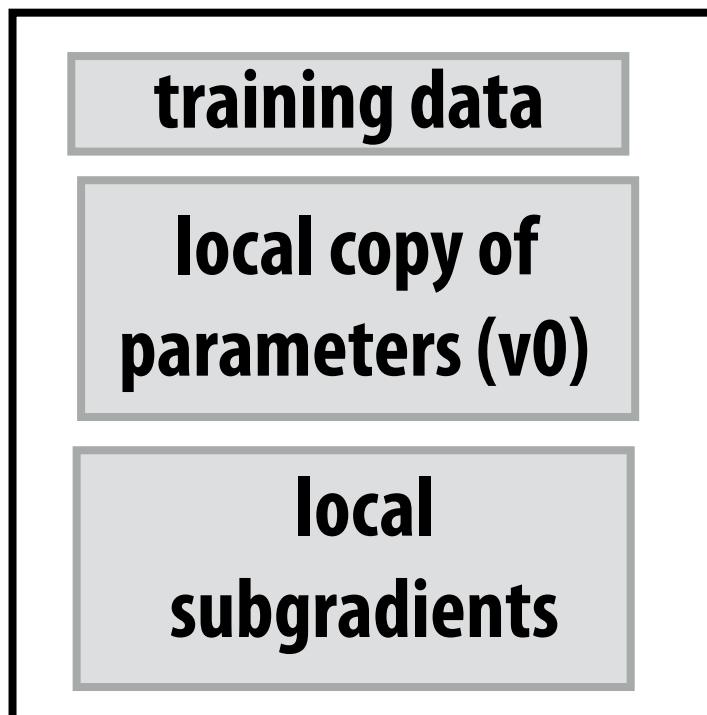
Workers independently compute local “subgradients”



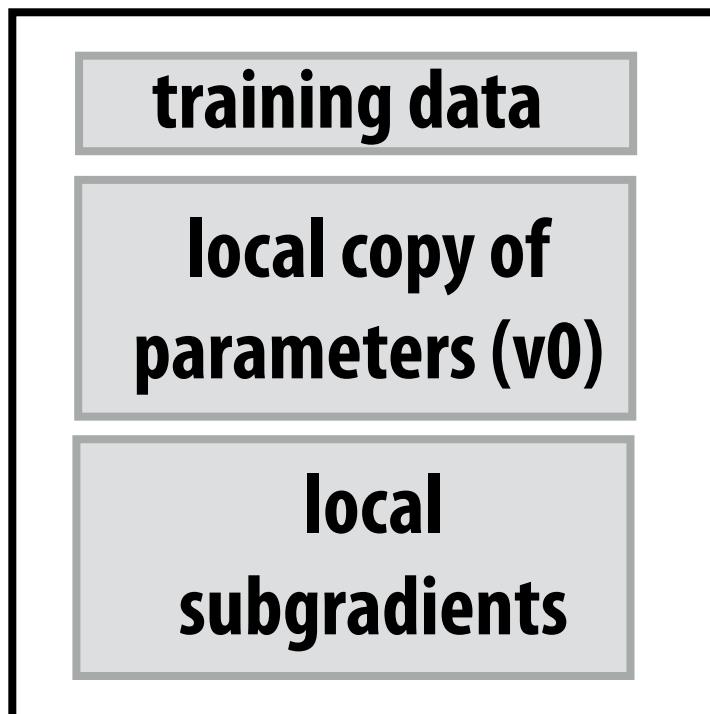
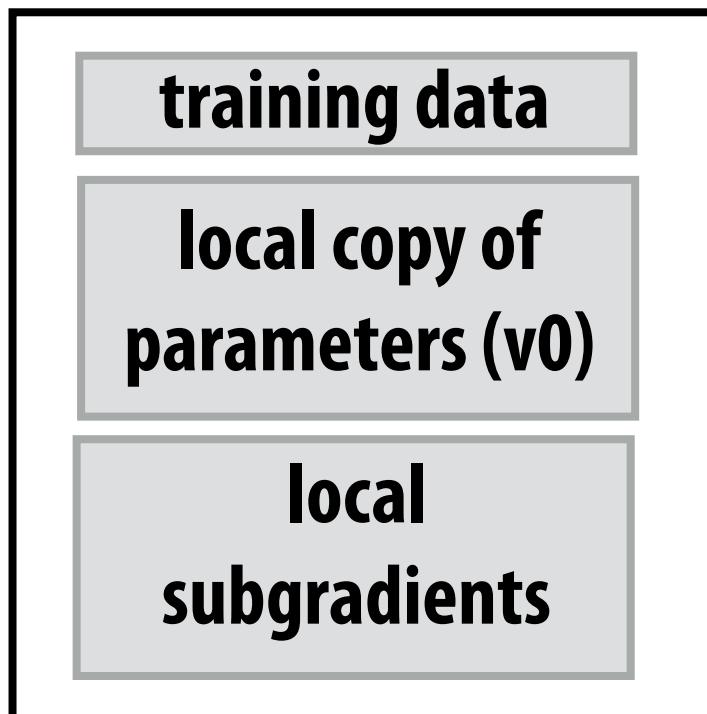
Worker sends subgradient to parameter server



Server updates global parameter values based on subgradient

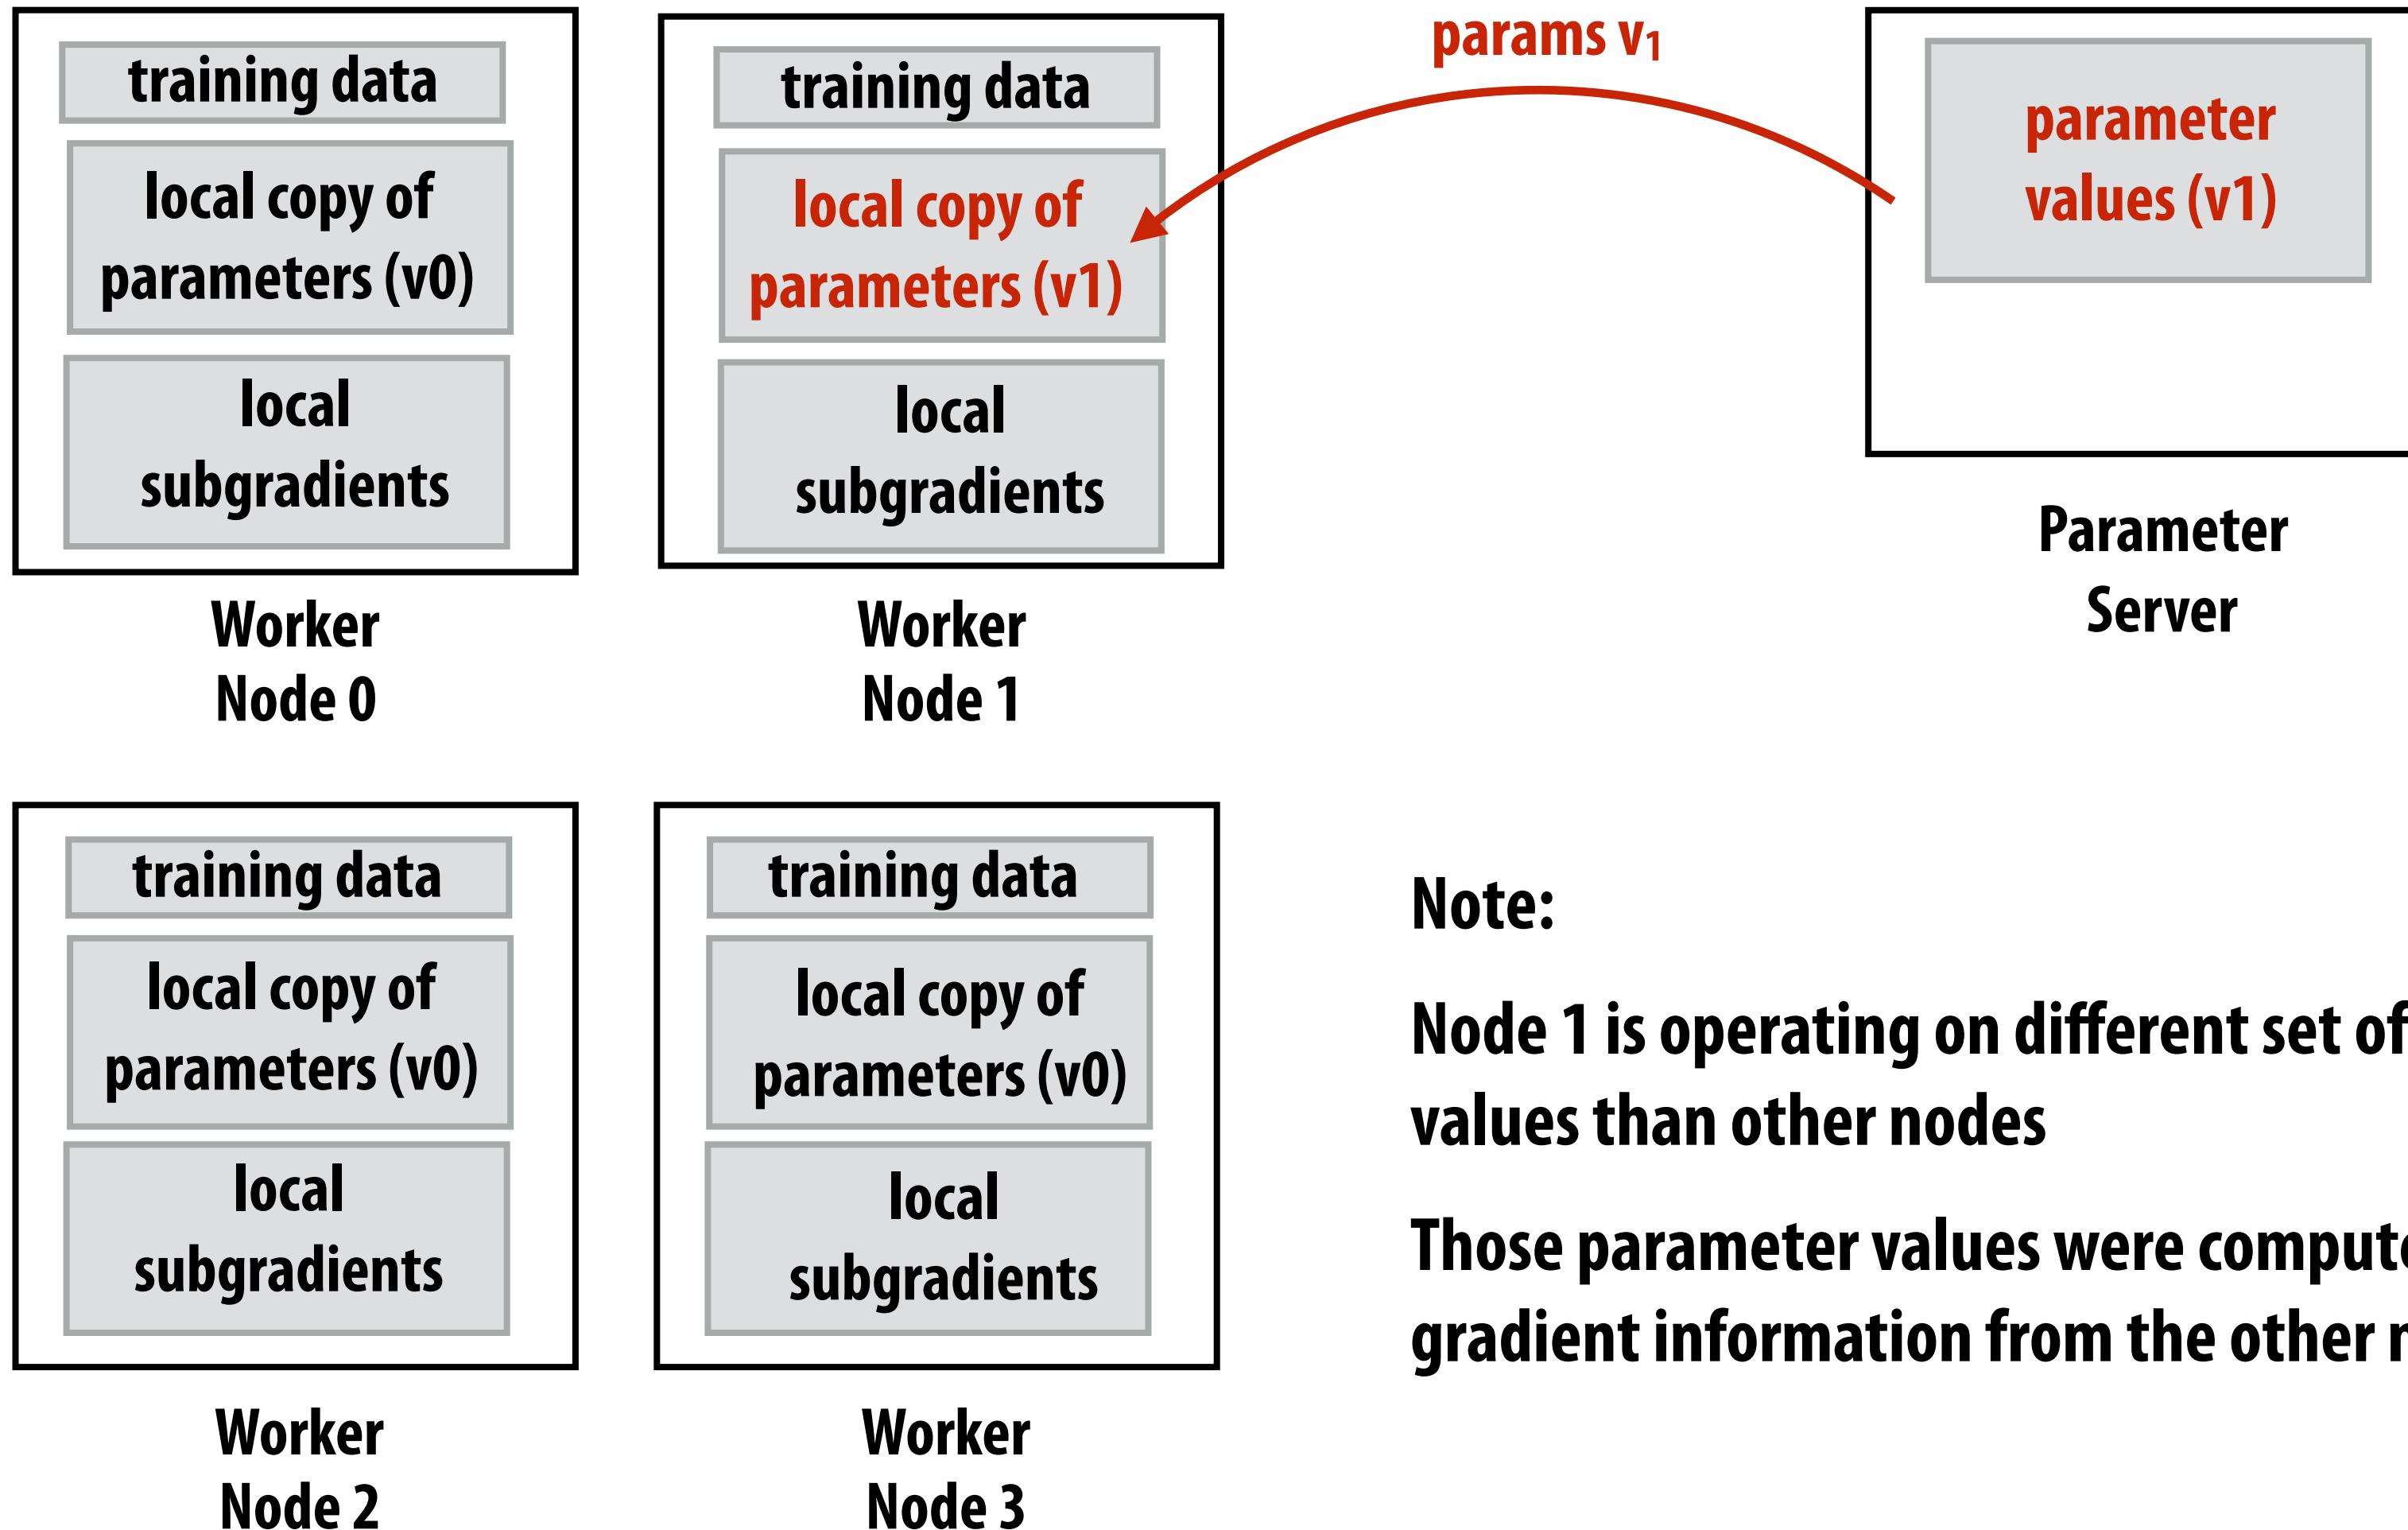


```
params += -subgrad * step_size;
```



Updated parameters sent to worker

Worker proceeds with another gradient computation step

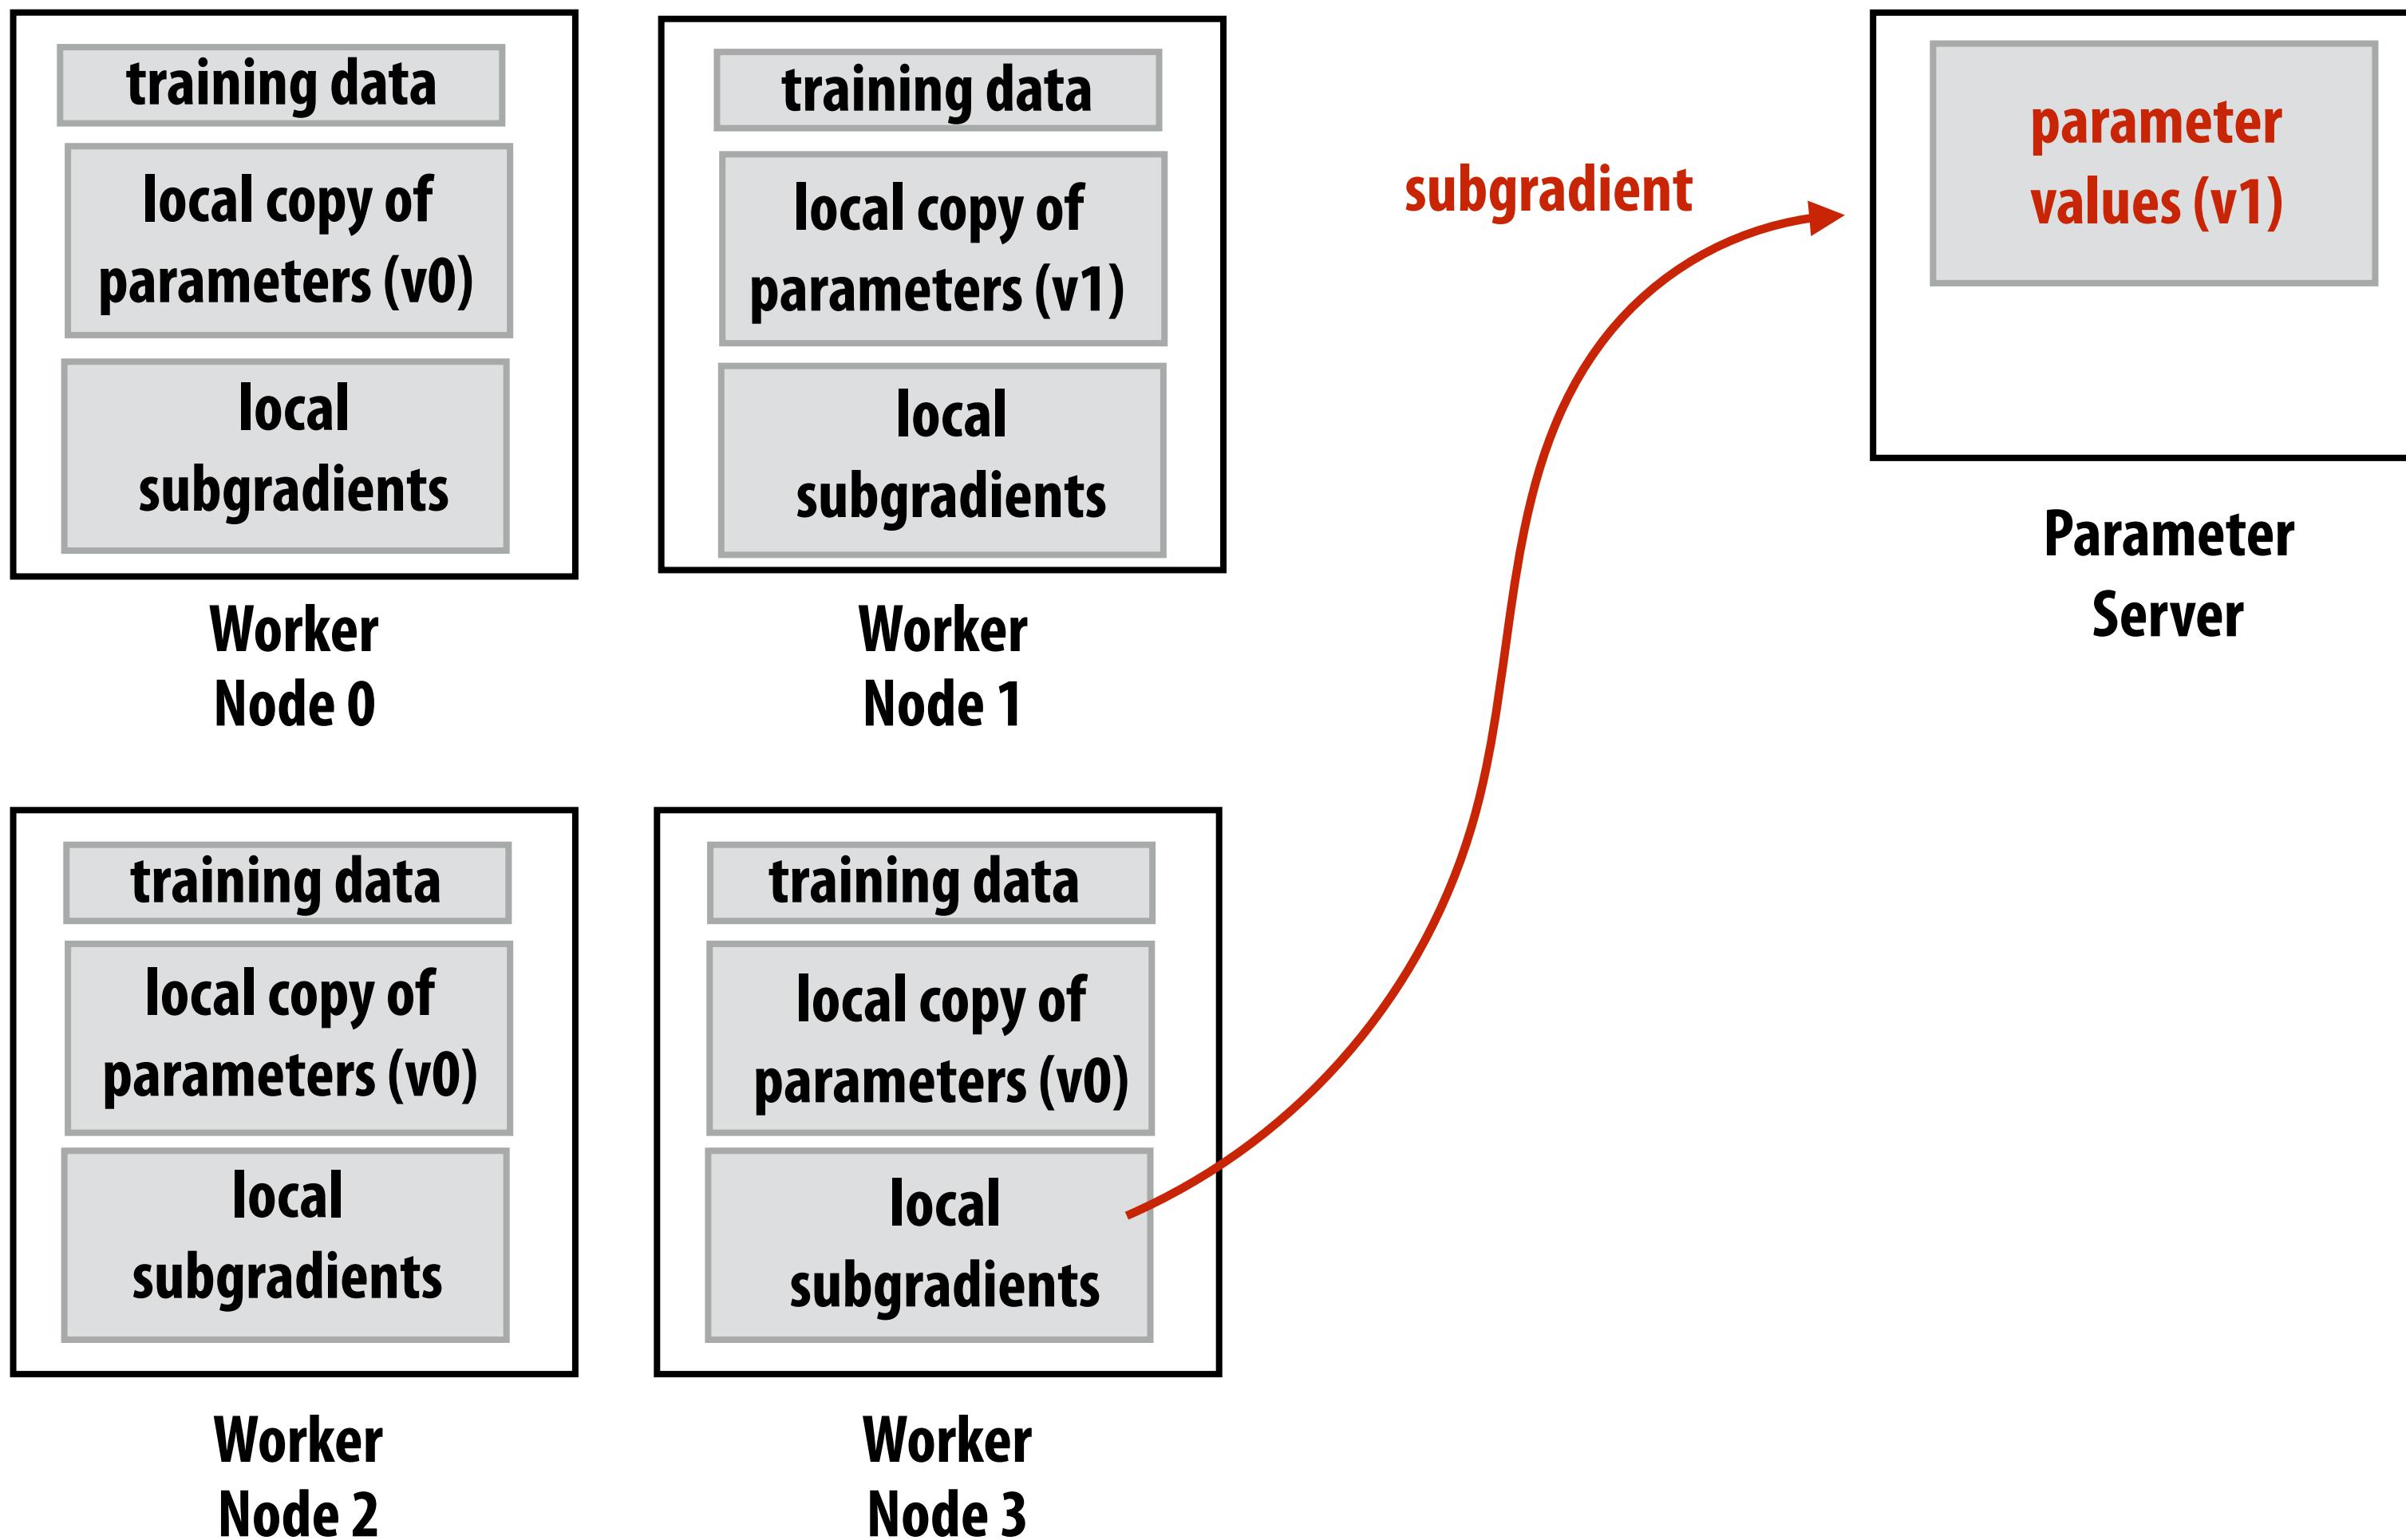


Note:

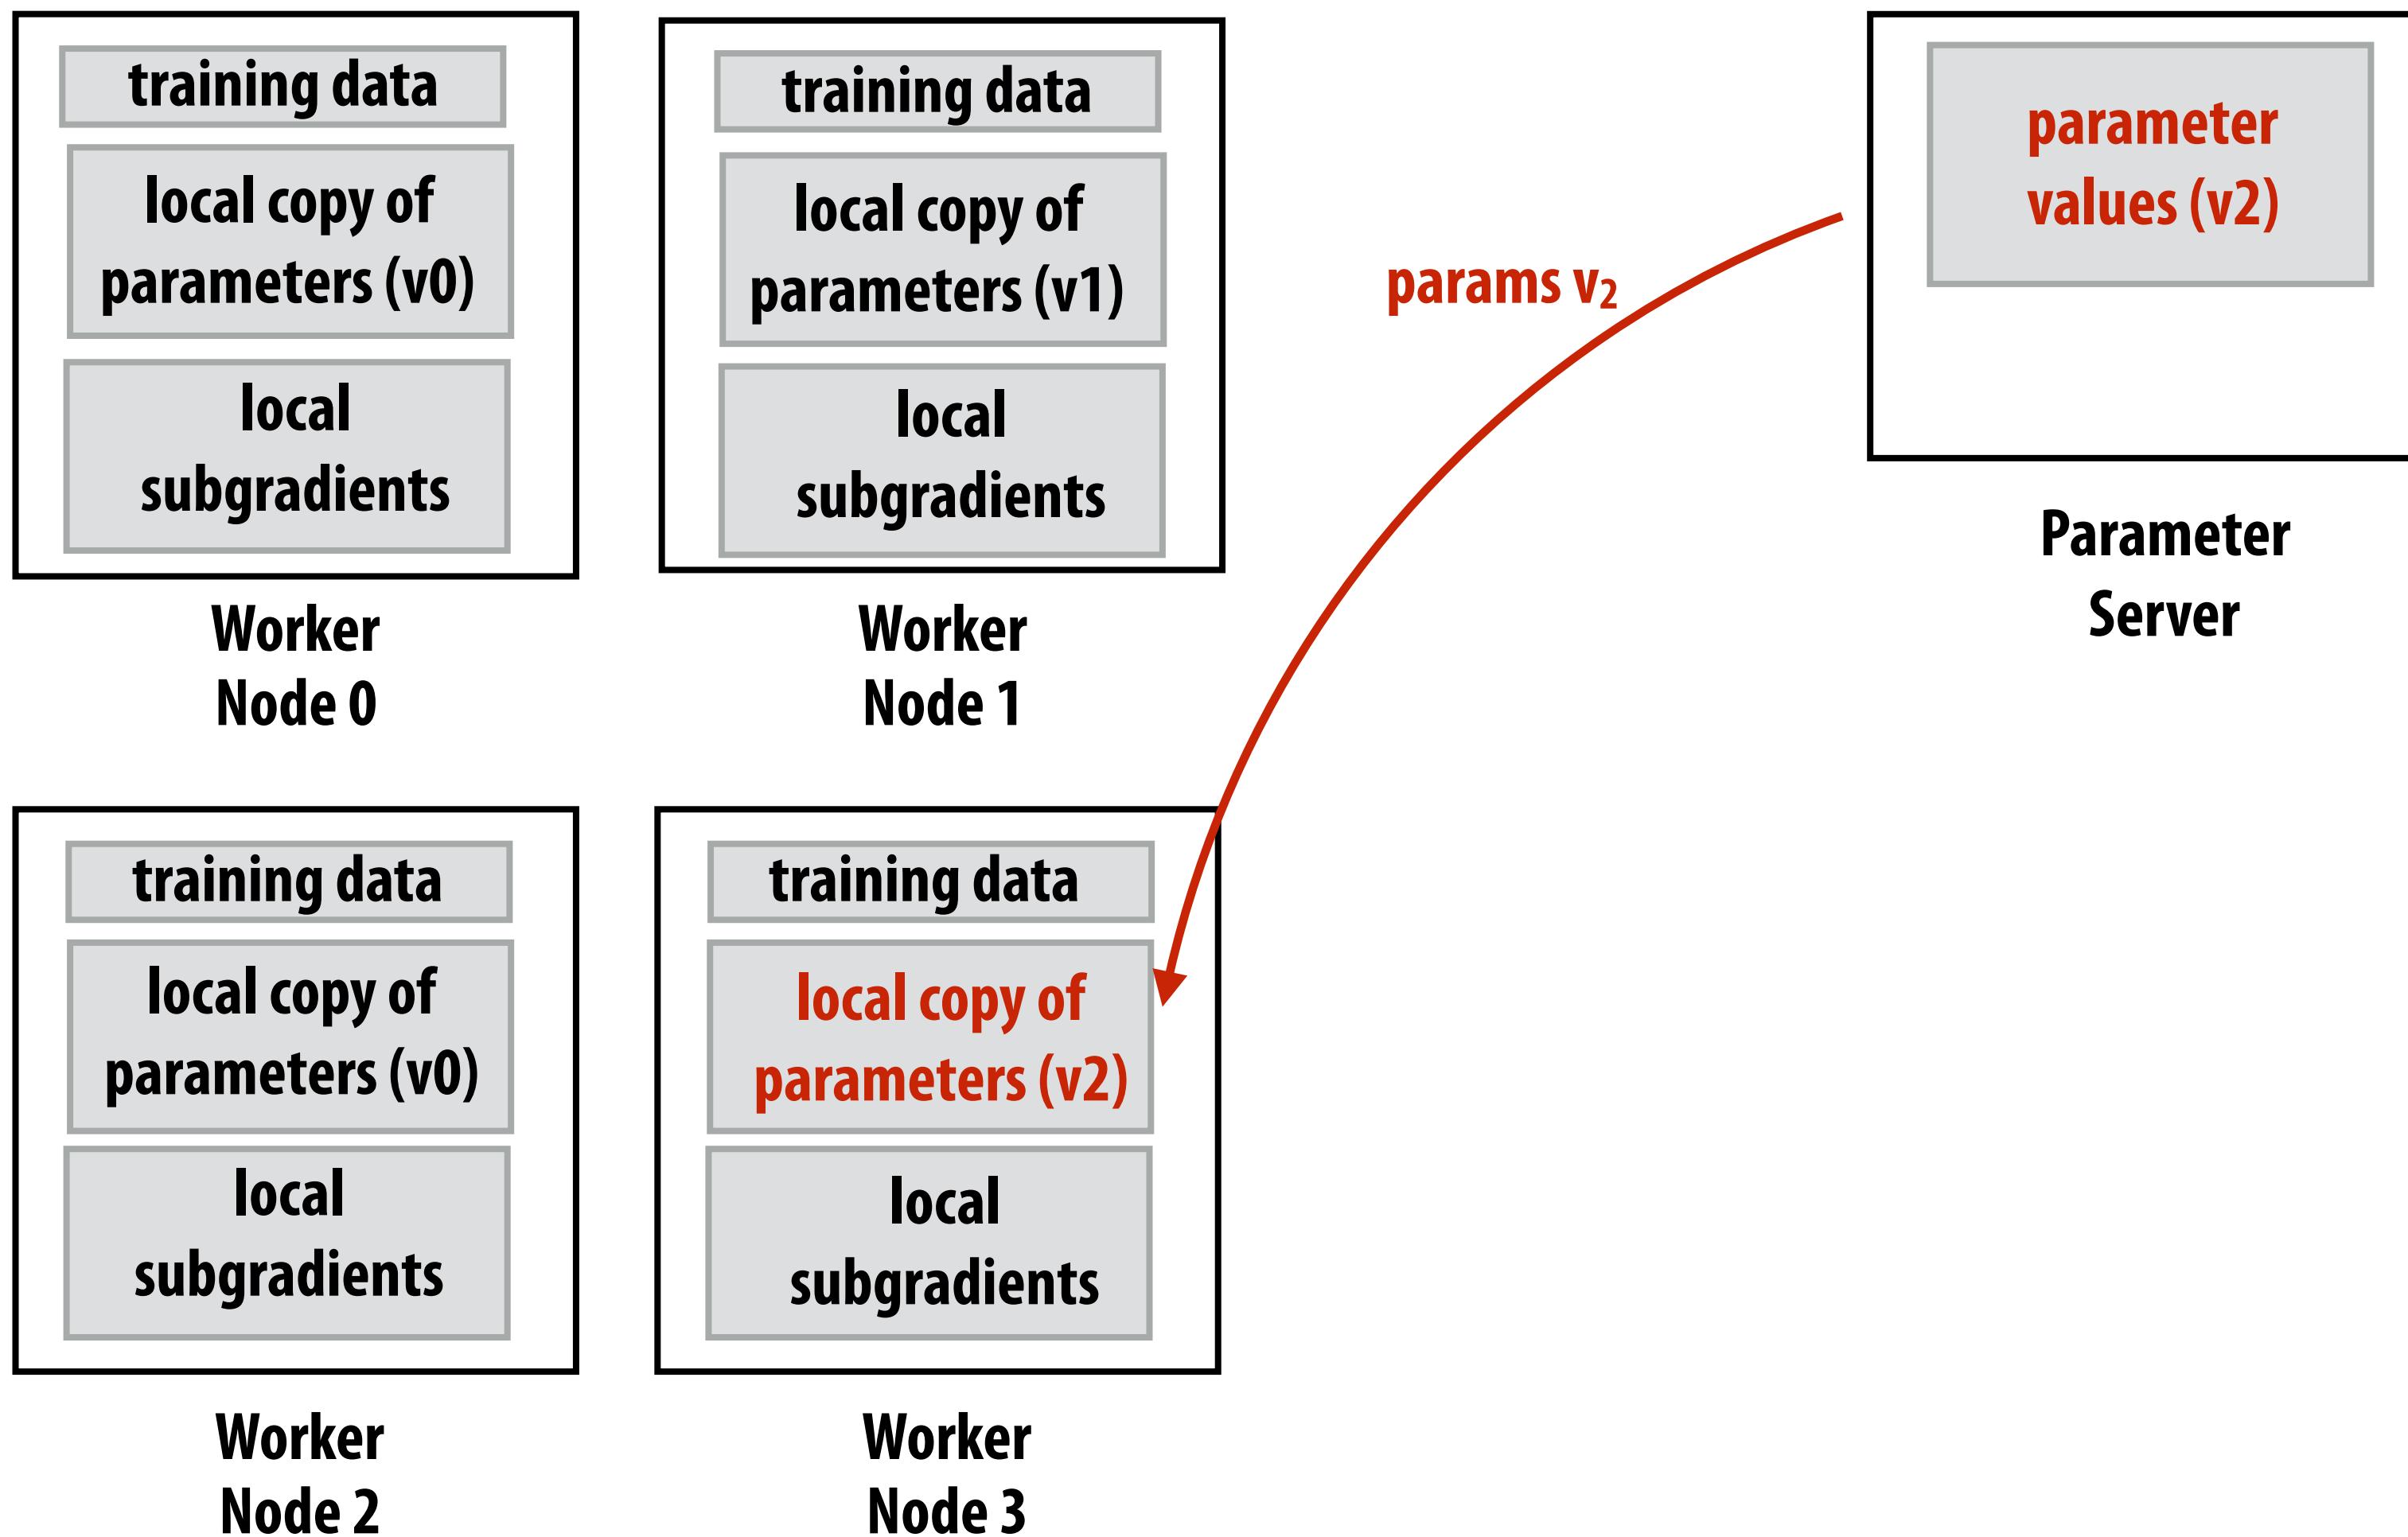
Node 1 is operating on different set of parameter values than other nodes

Those parameter values were computed without gradient information from the other nodes

Updated parameters sent to worker (again)



Worker continues with updated parameters

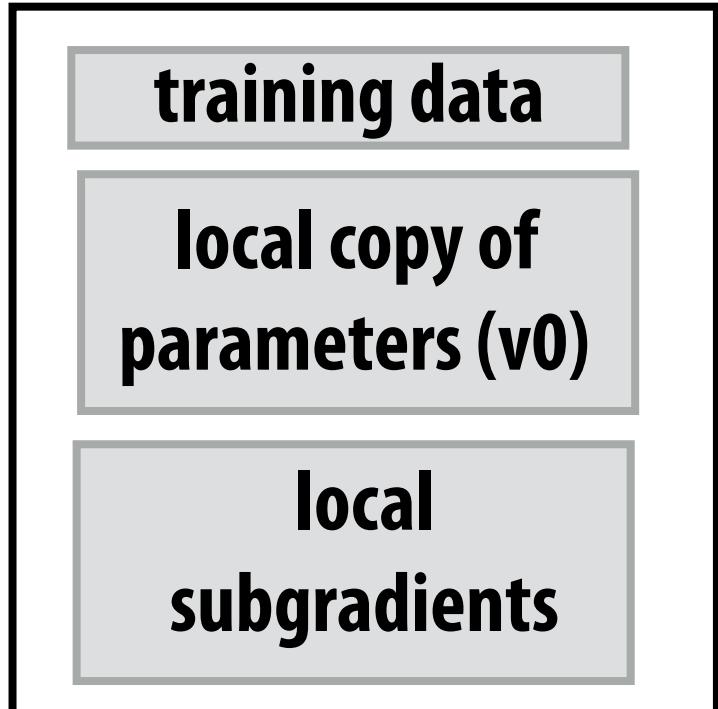


Summary: asynchronous parameter update

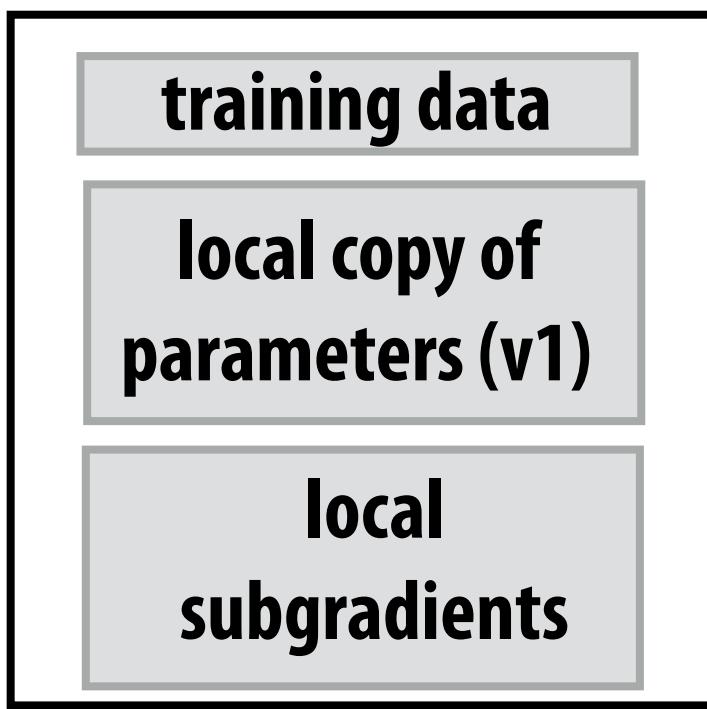
- Idea: avoid global synchronization on all parameter updates between each SGD iteration
 - Design reflects realities of cluster computing:
 - Slow interconnects
 - Unpredictable machine performance
- Solution: asynchronous (and partial) subgradient updates
- Will impact convergence of SGD
 - Node N working on iteration i may not have parameter values that result the results of the $i-1$ prior SGD iterations

Bottleneck?

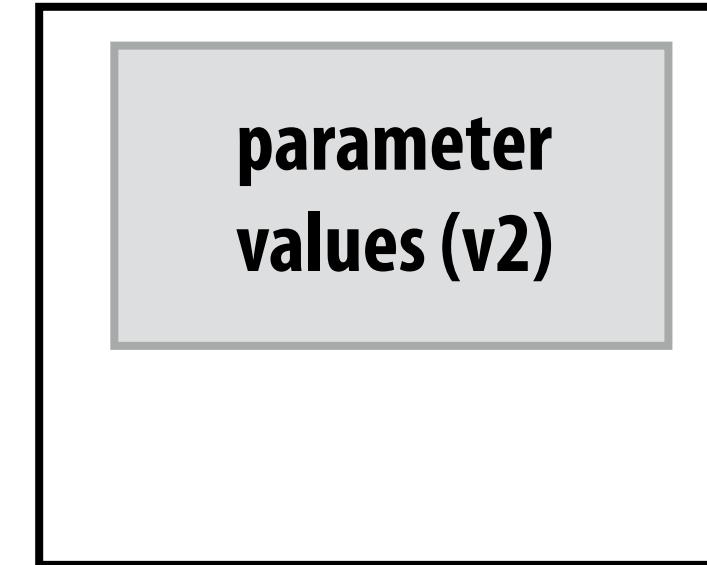
What if there is heavy contention for parameter server?



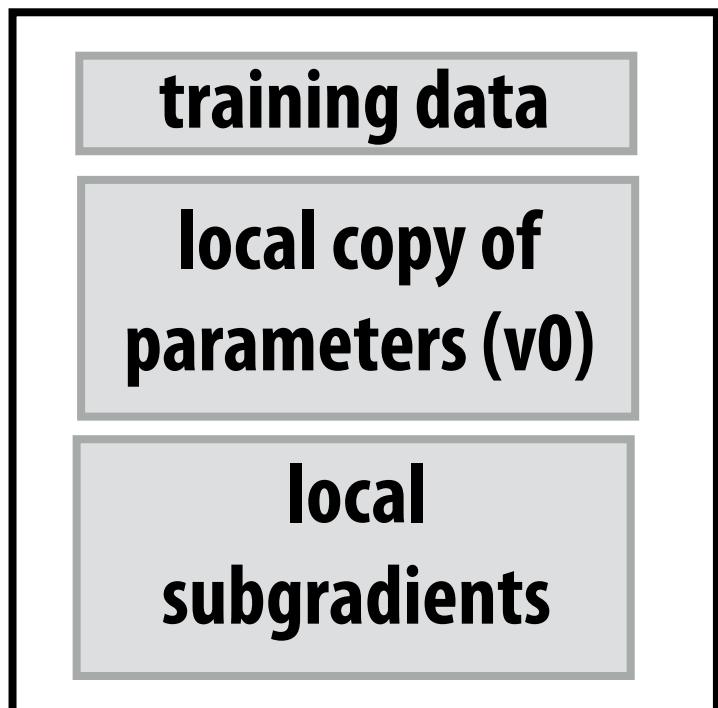
Worker
Node 0



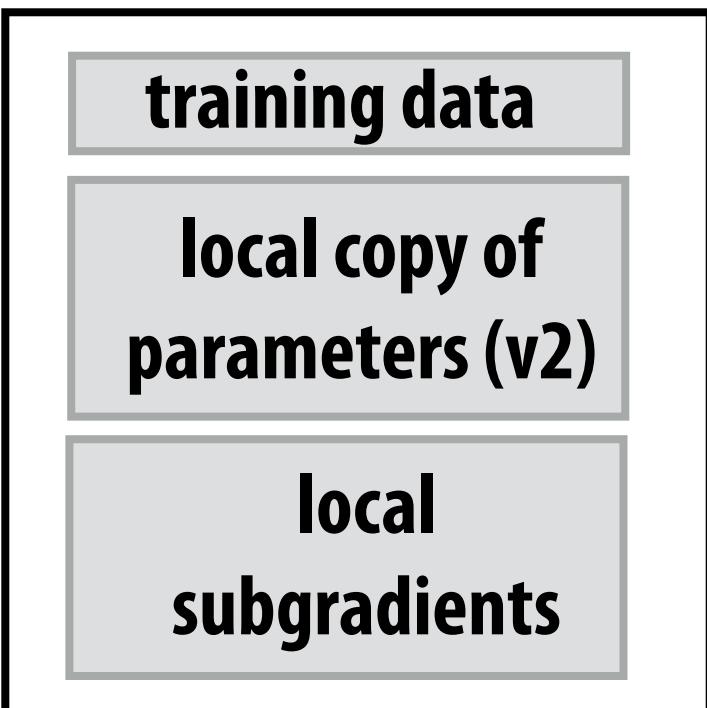
Worker
Node 1



Parameter
Server



Worker
Node 2

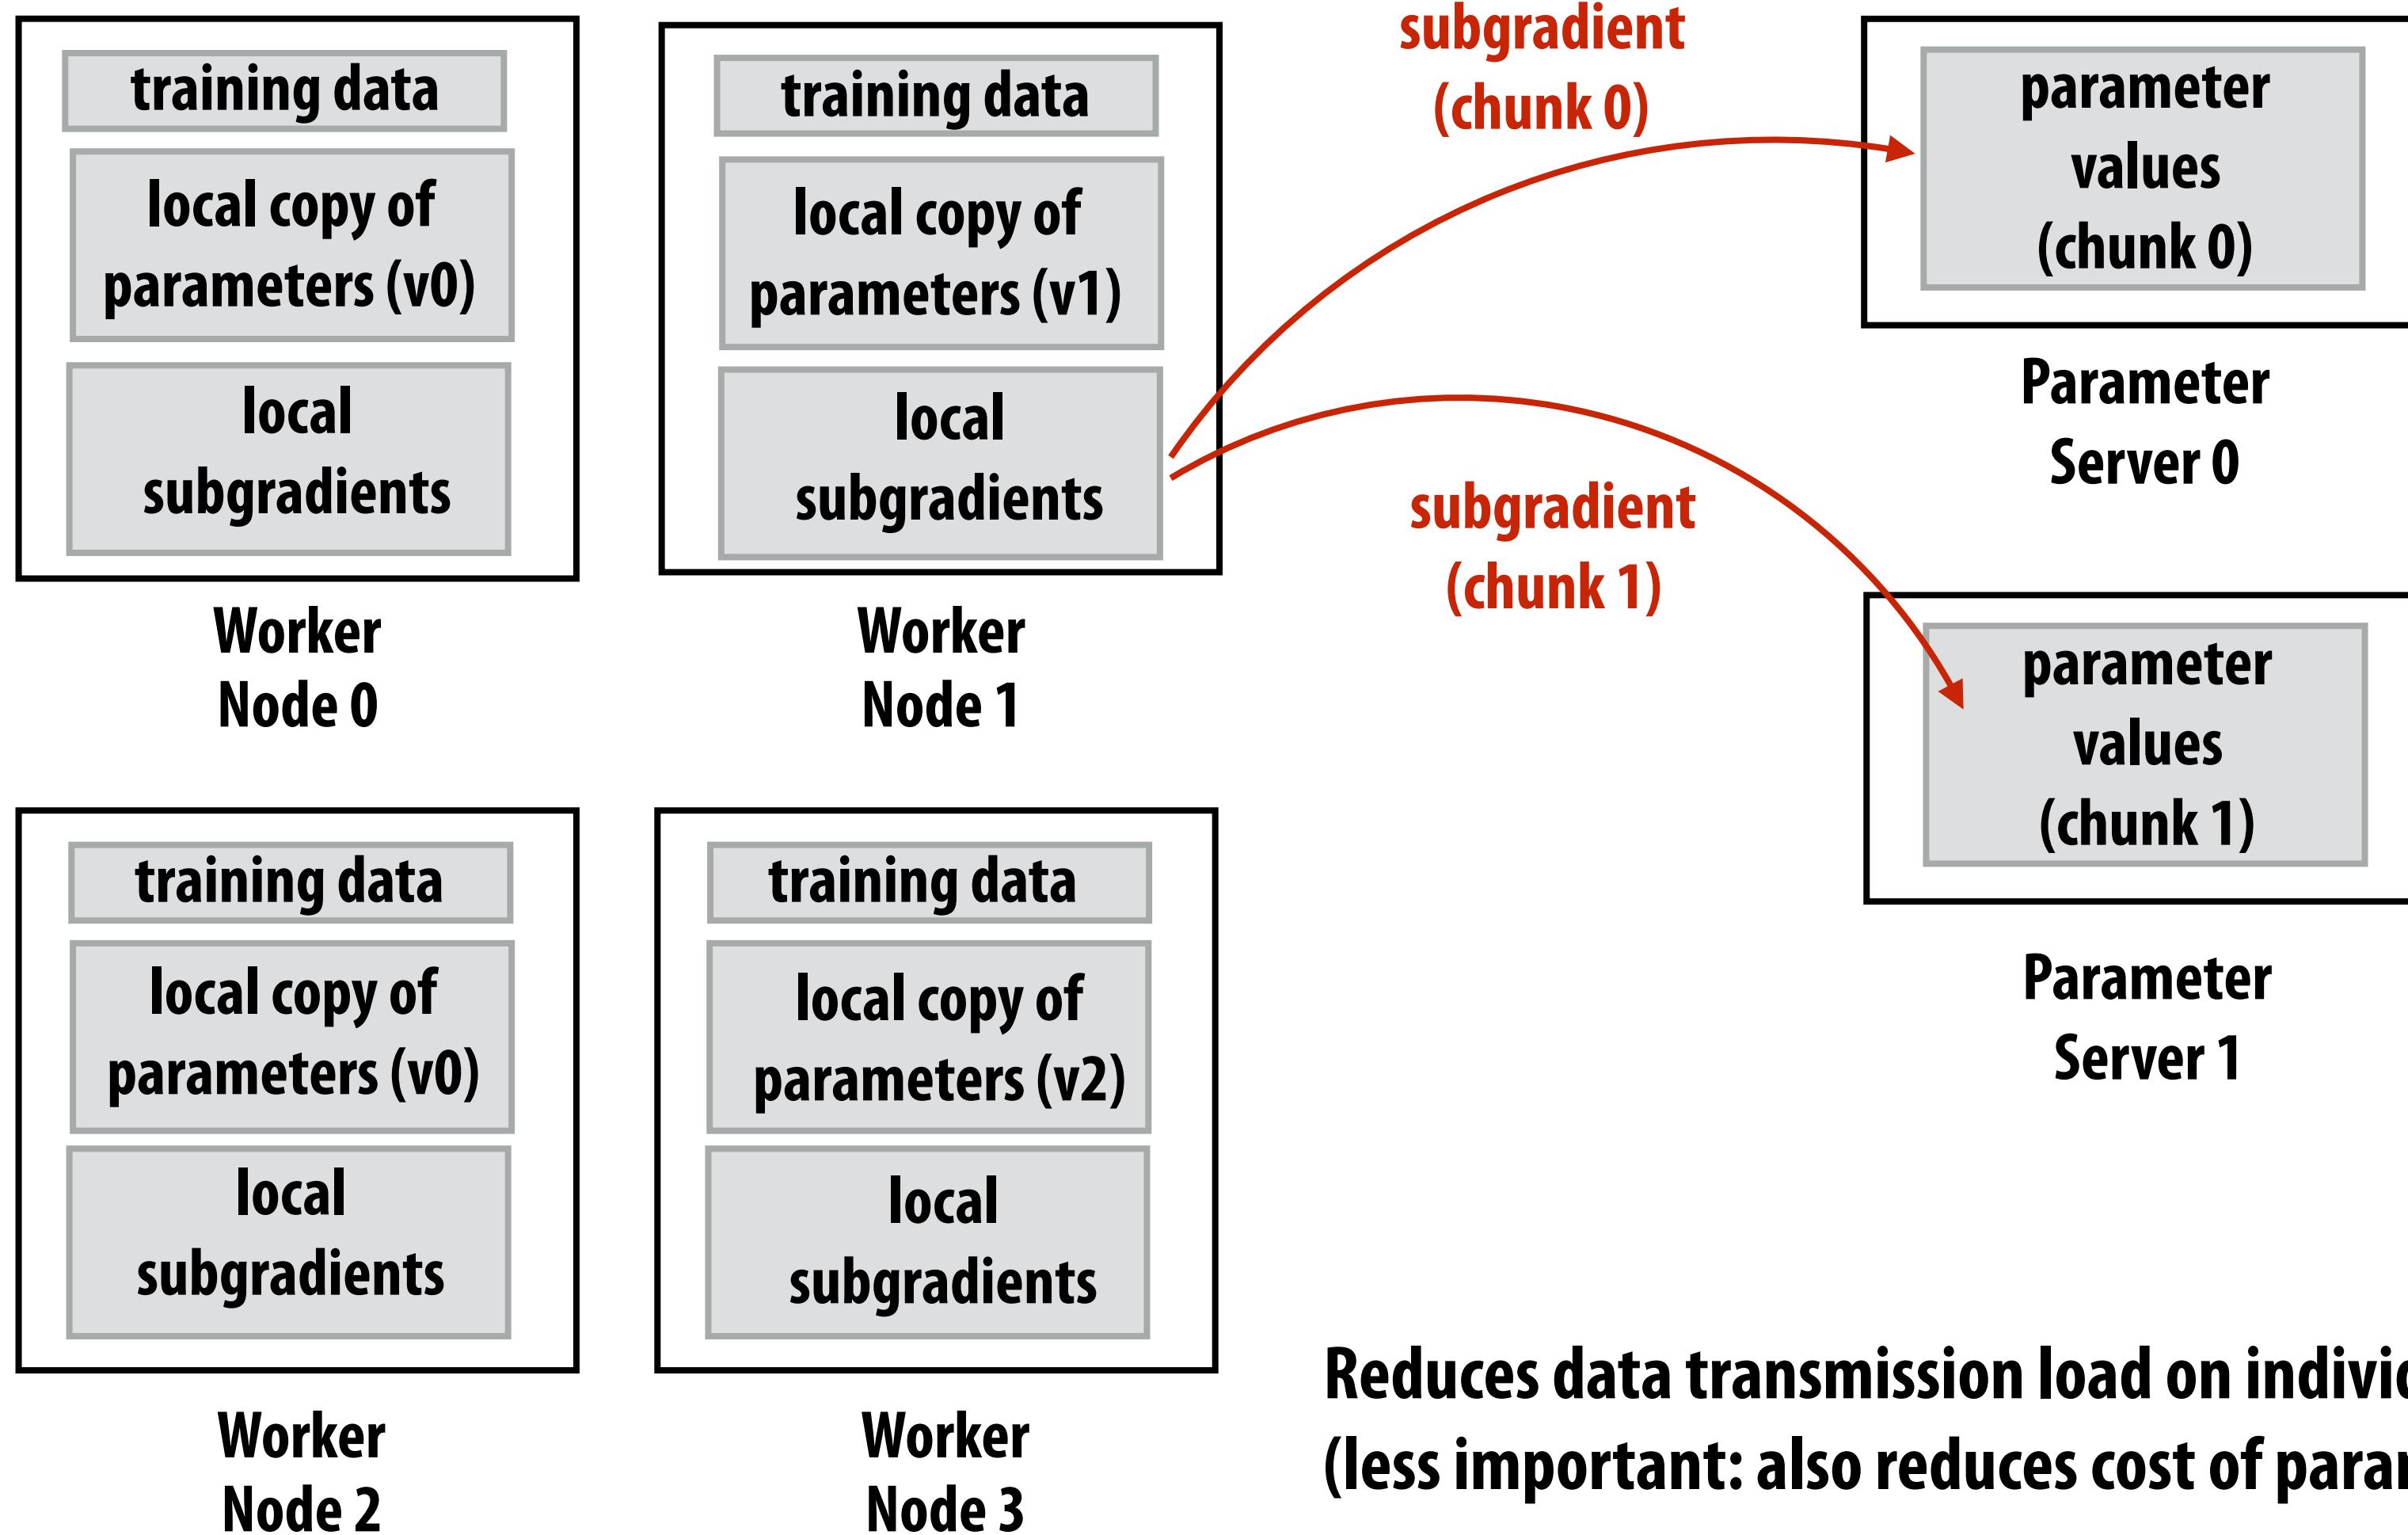


Worker
Node 3

Shard the parameter server

Partition parameters across servers

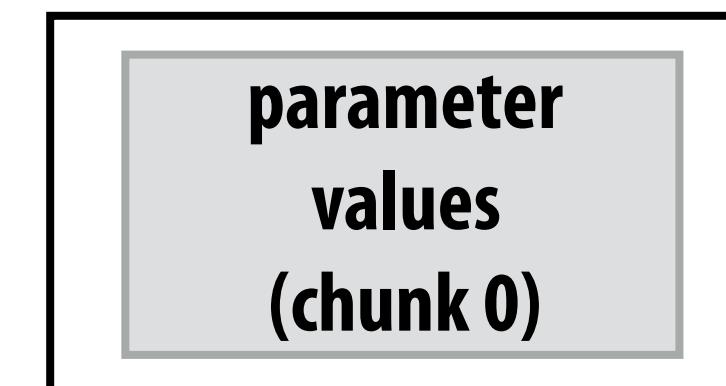
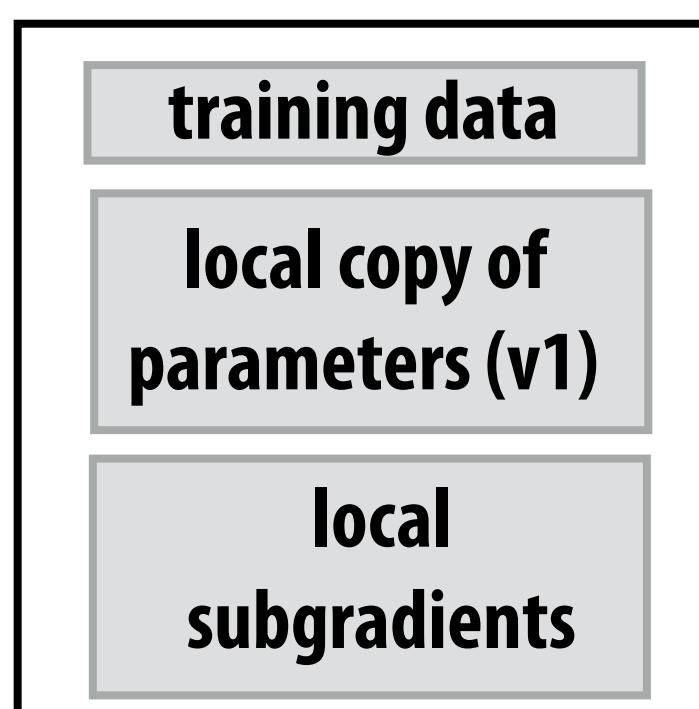
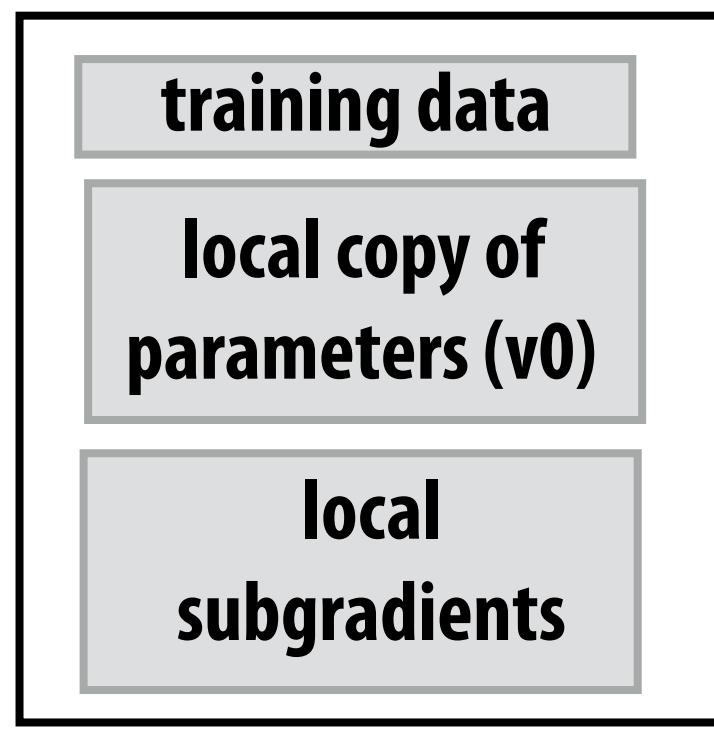
Worker sends chunk of subgradients to owning parameter server



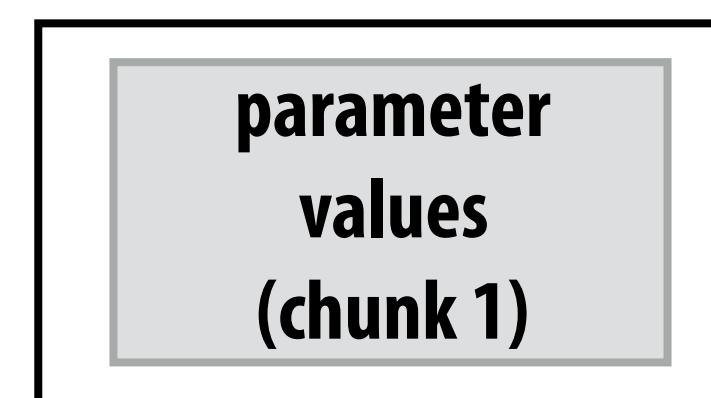
Reduces data transmission load on individual servers
(less important: also reduces cost of parameter update)

What if model parameters do not fit on one worker?

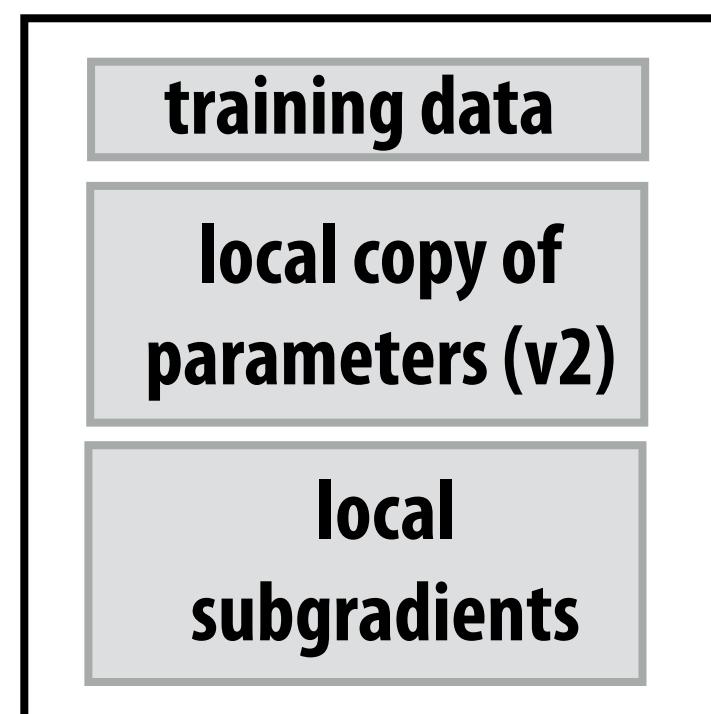
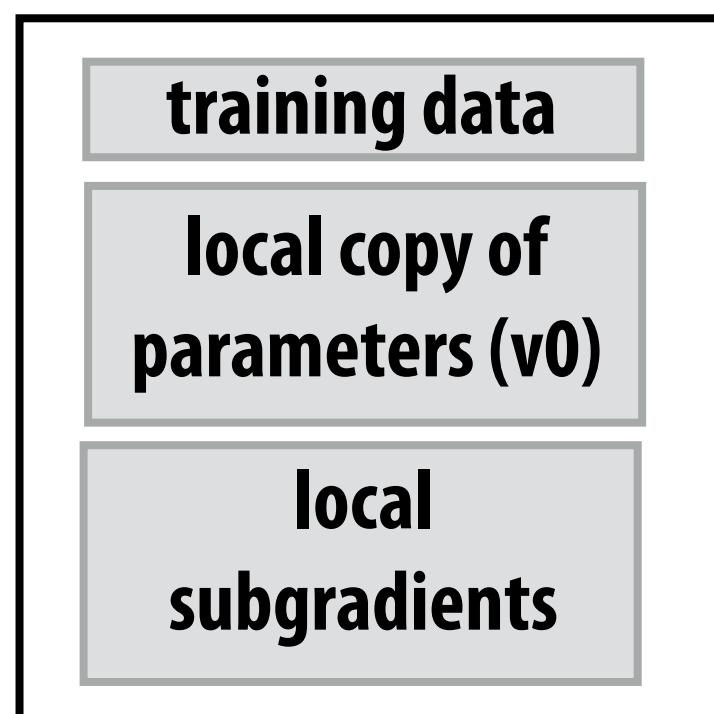
Recall high footprint of training large networks
(particularly with large mini-batch sizes)



Parameter Server 0



Parameter Server 1

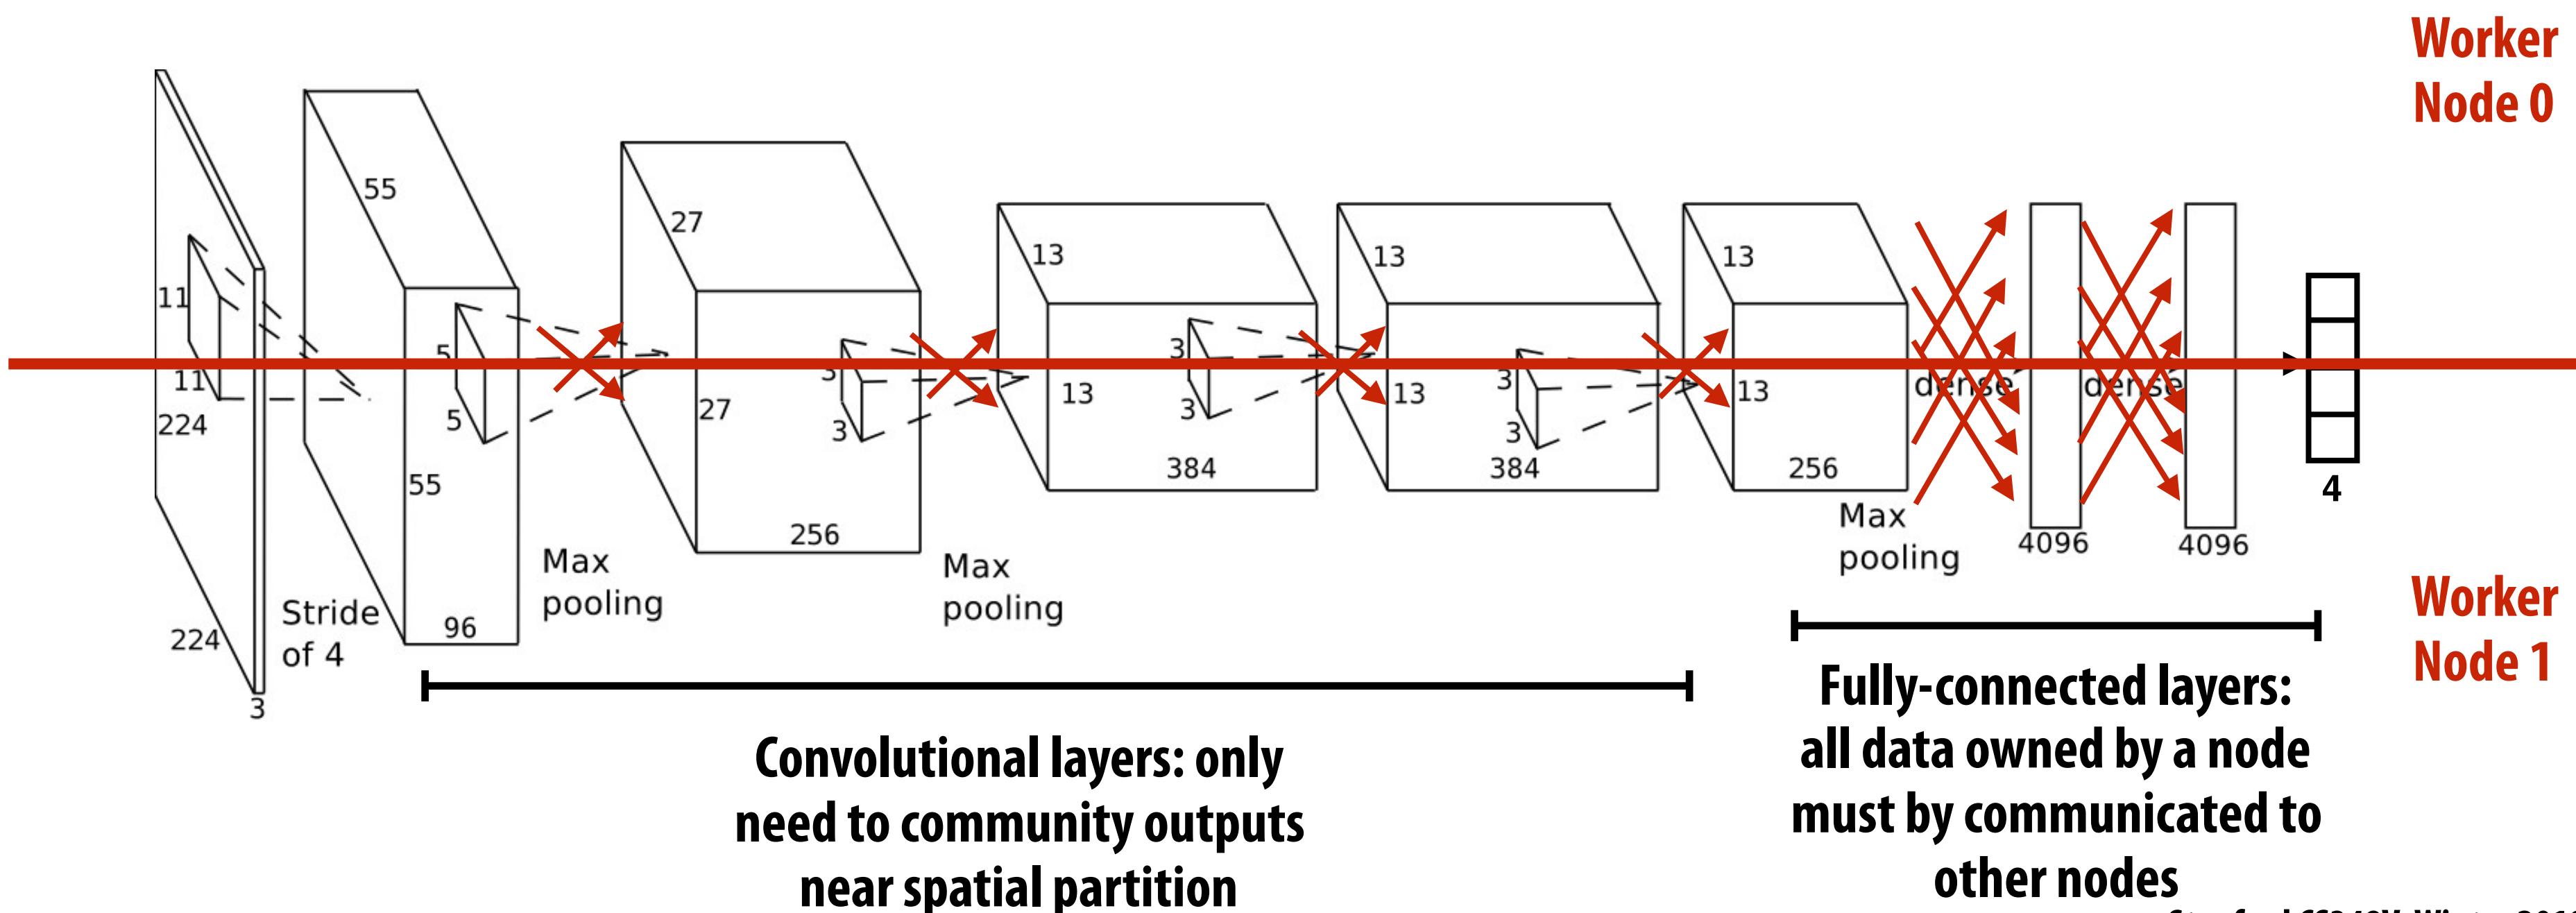


Model parallelism

Partition network parameters across nodes
(spatial partitioning to reduce communication)

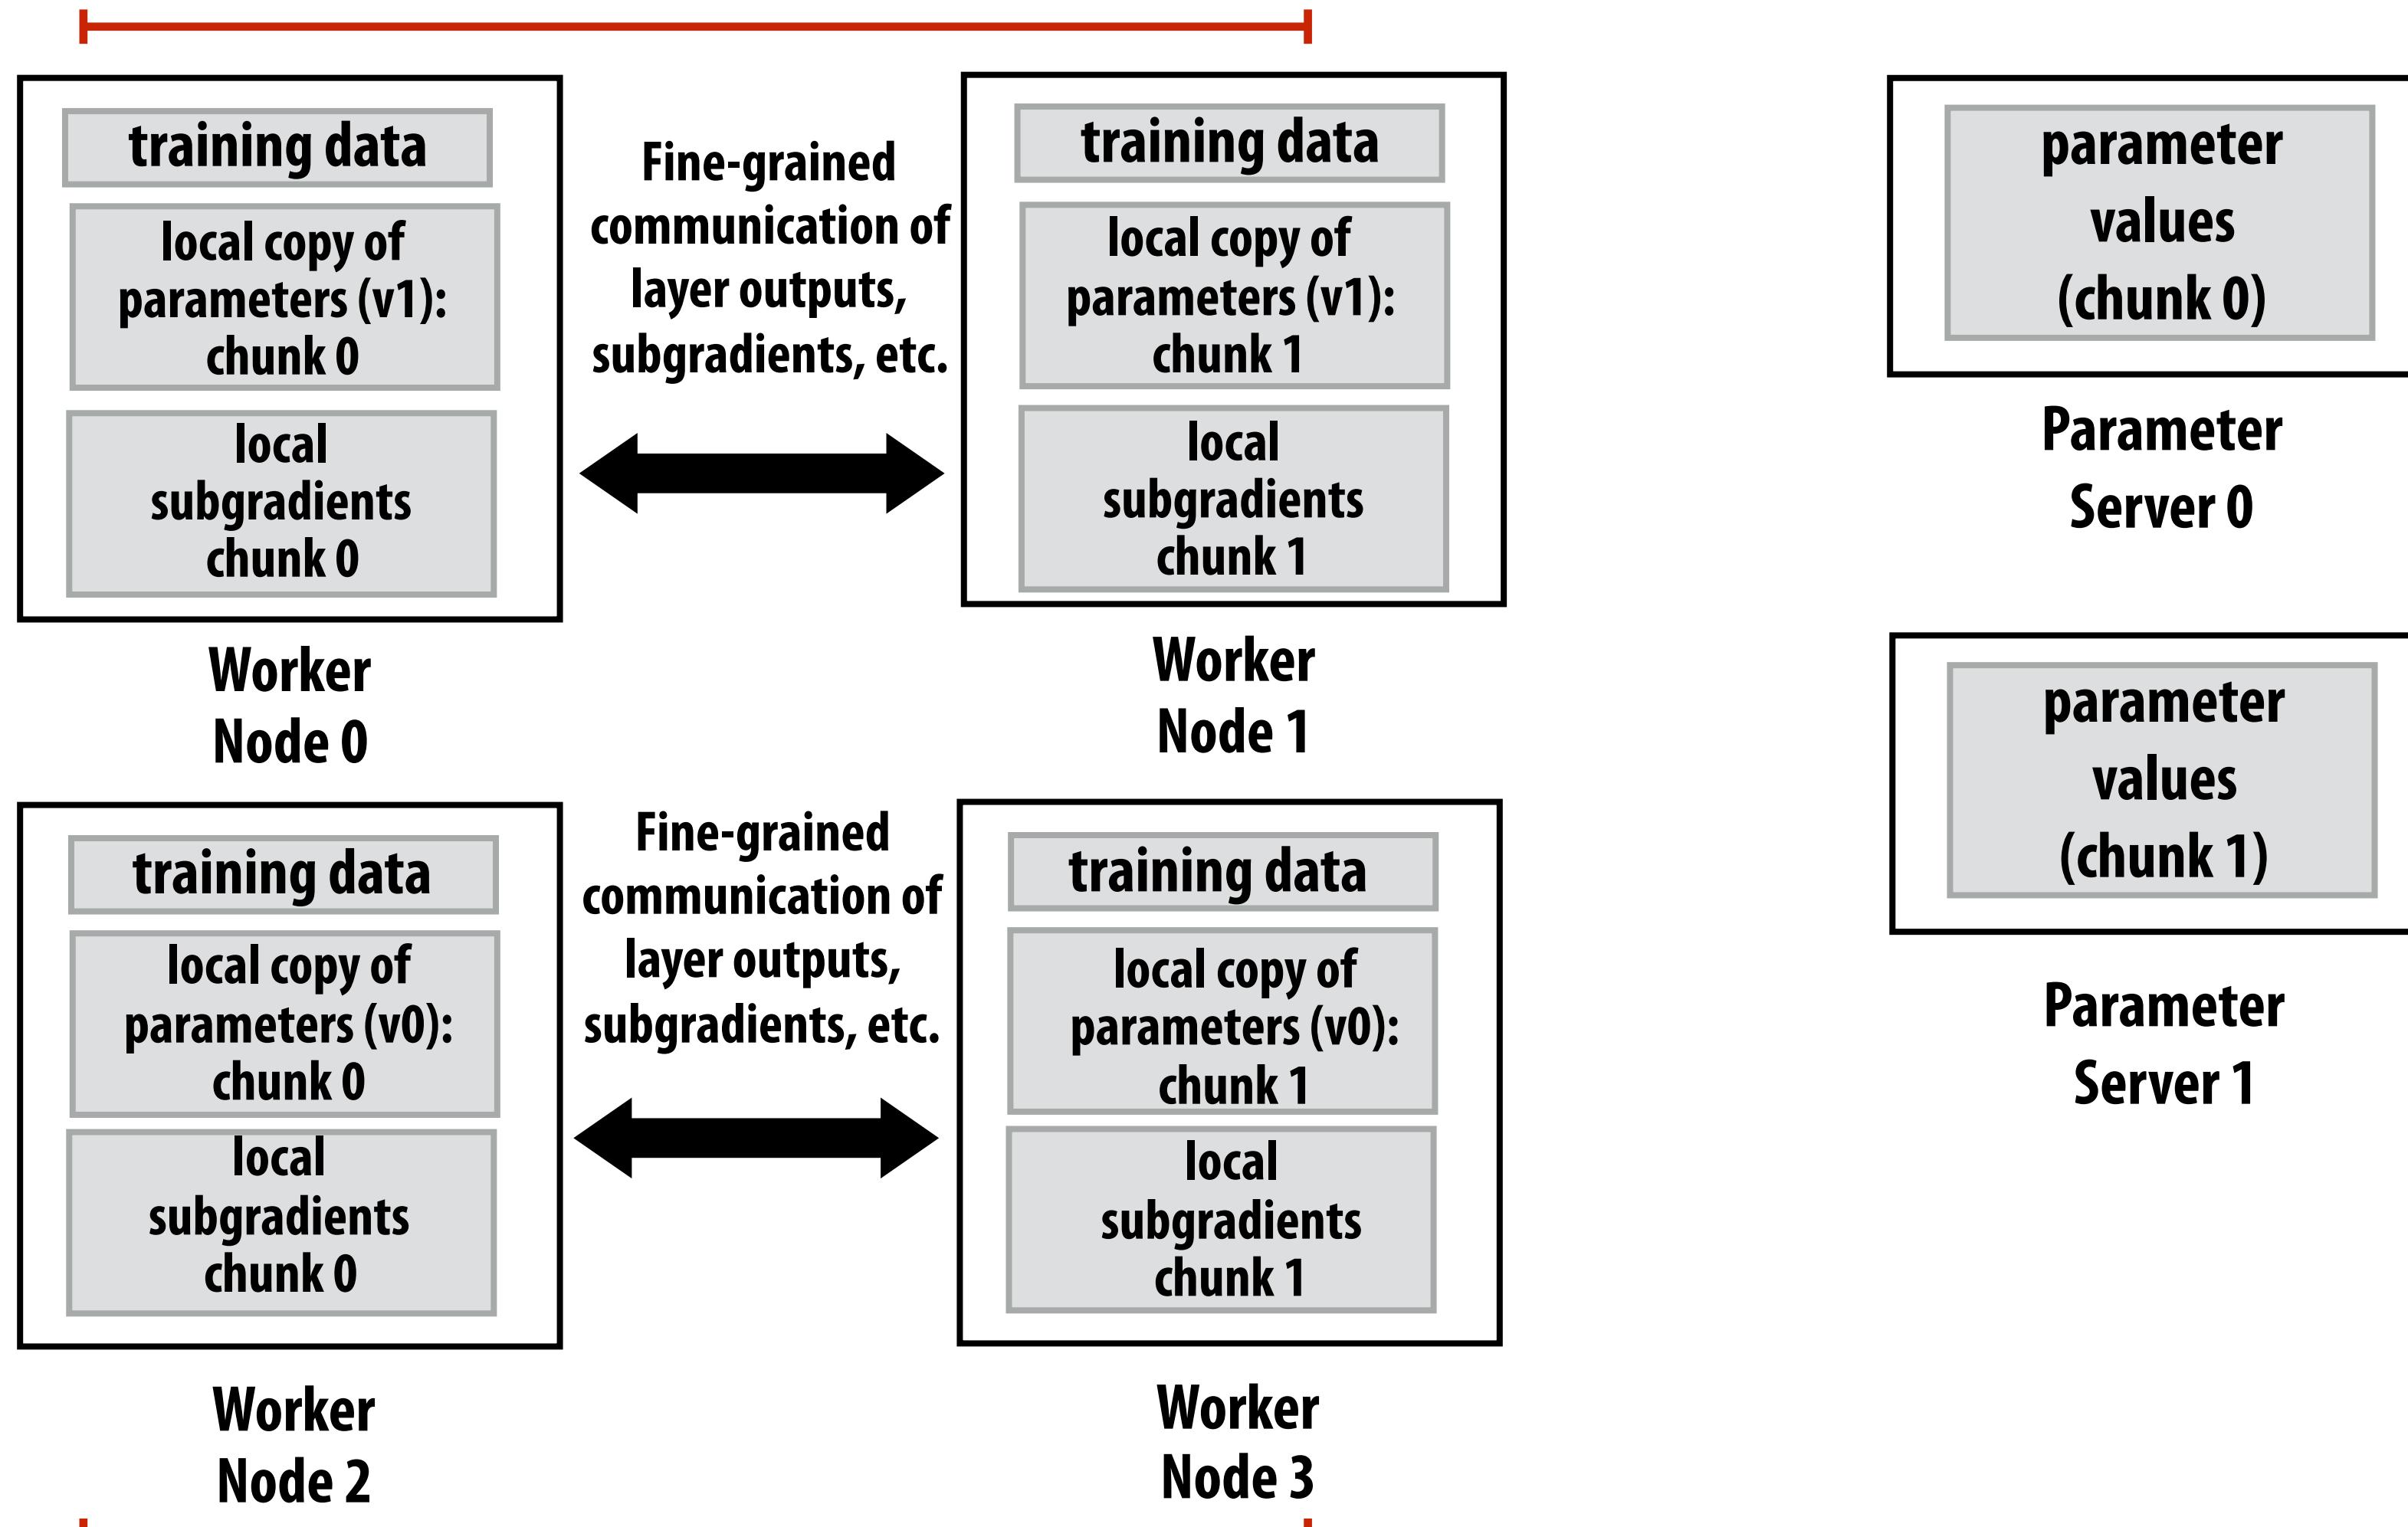
Reduce internode communication through network design:

- Use small spatial convolutions (1x1 convolutions)
- Reduce/shrink fully-connected layers



Training data-parallel and model-parallel execution

Working on subgradient computation
for a single copy of the model



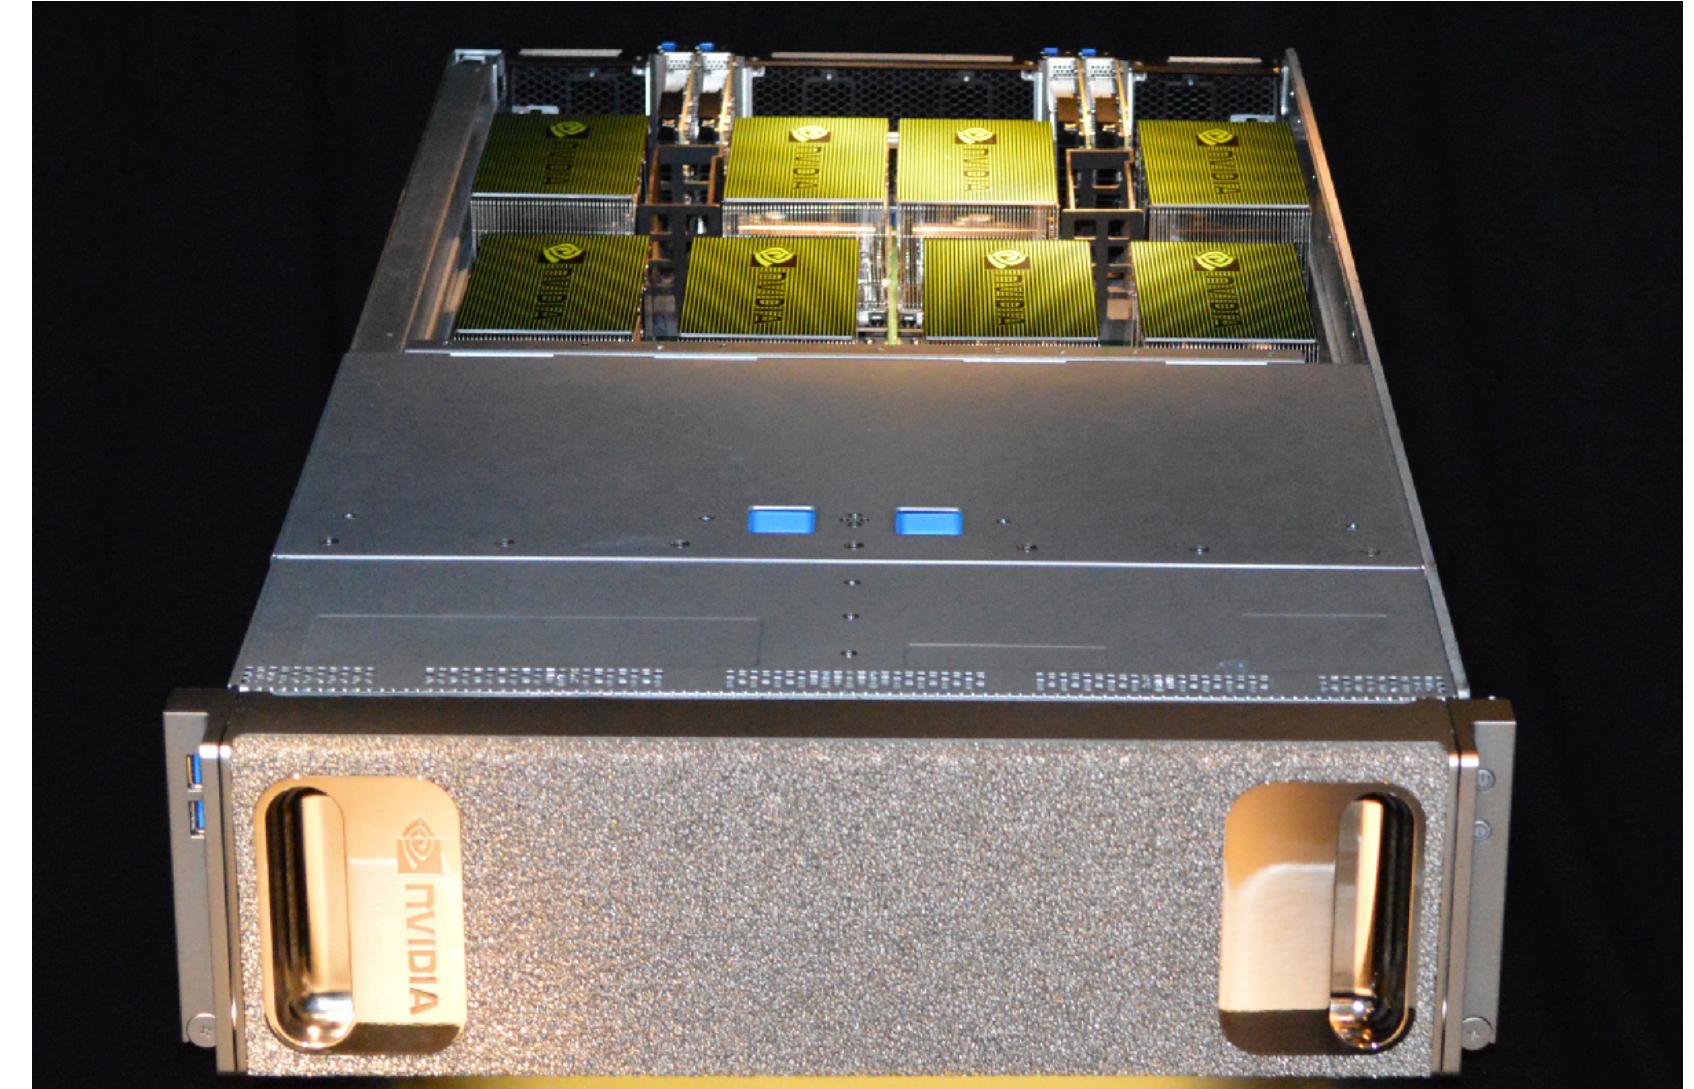
Working on subgradient computation
for a single copy of the model

Better hardware? using supercomputers for training?

- Fast interconnects critical for model-parallel training
 - Fine-grained communication of outputs and gradients
- Fast interconnects diminish need for async training algorithms
 - Avoid randomness in training due to schedule of computation (yes, there remains randomness due to SGD algorithm)



**OakRidge Titan Supercomputer
(Cray low-latency interconnect)**



**NVIDIA DGX-1: 8 Pascal GPUs connected
via high speed NV-Link interconnect**

Better algorithmic techniques (again): improving scalability of synchronous training...

- Larger mini-batches increase compute to communication ratio:
communicate gradients summed over B training inputs

```
for each item x in mini-batch on this node:  
    grad += evaluate_loss_gradient(f, loss_func, params, x)  
barrier();  
sum reduce gradients across all nodes, communicate results to all nodes  
barrier();  
update copy of local parameter values
```

- But large mini-batches (if used naively) reduce accuracy of
model trained

Linear scaling rule

Recall: minibatch SGD parameter update

$$w_{t+1} = w_t - \eta \frac{1}{n} \sum_{x \in \mathcal{B}} \nabla l(x, w_t)$$

size of mini batch = n
SGD learning rate = η

Consider processing of k minibatches (k steps of gradient descent)

$$w_{t+k} = w_t - \eta \frac{1}{n} \sum_{j < k} \sum_{x \in \mathcal{B}_j} \nabla l(x, w_{t+j})$$

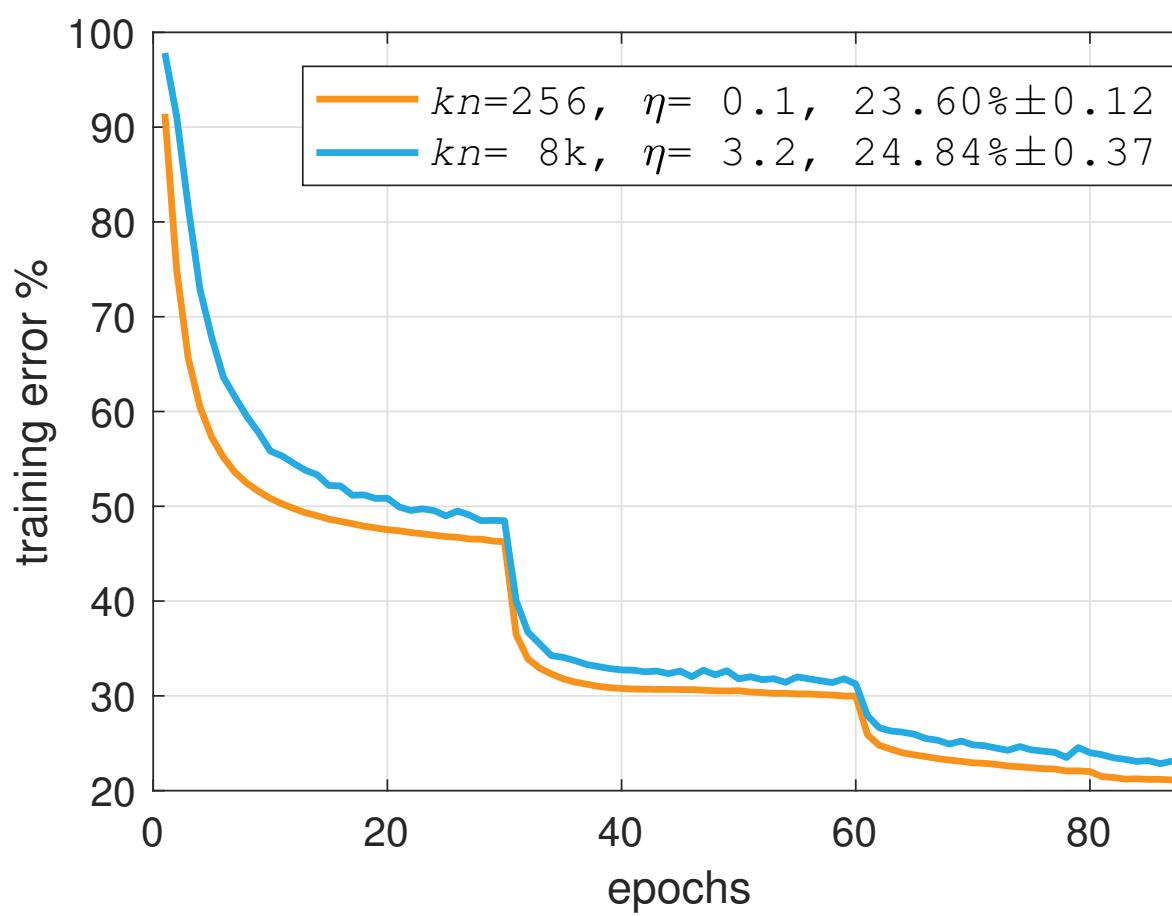
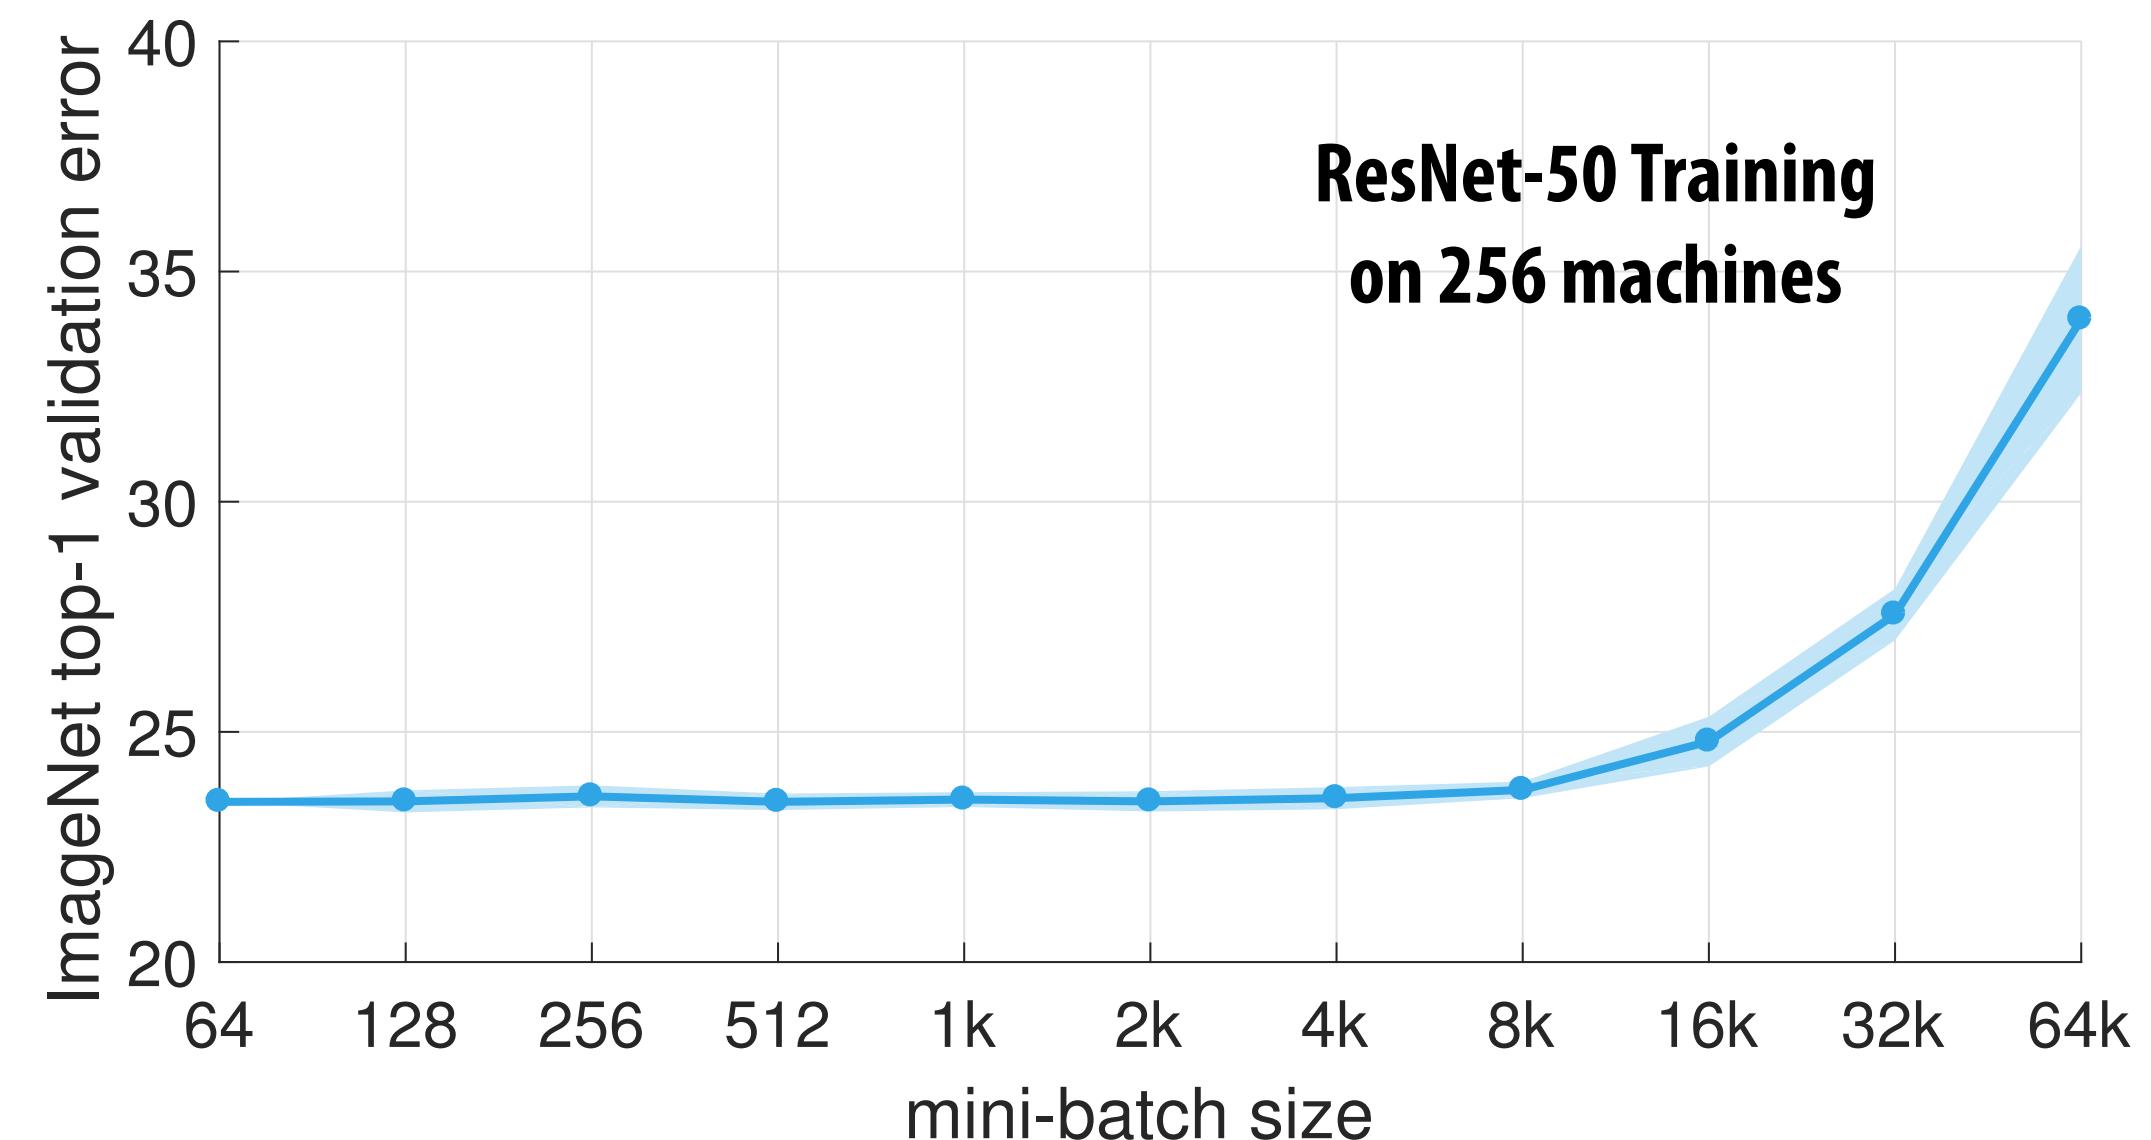
Consider processing one minibatch that is of size kn (one step of gradient descent)

$$\hat{w}_{t+1} = w_t - \hat{\eta} \frac{1}{kn} \sum_{j < k} \sum_{x \in \mathcal{B}_j} \nabla l(x, w_t)$$

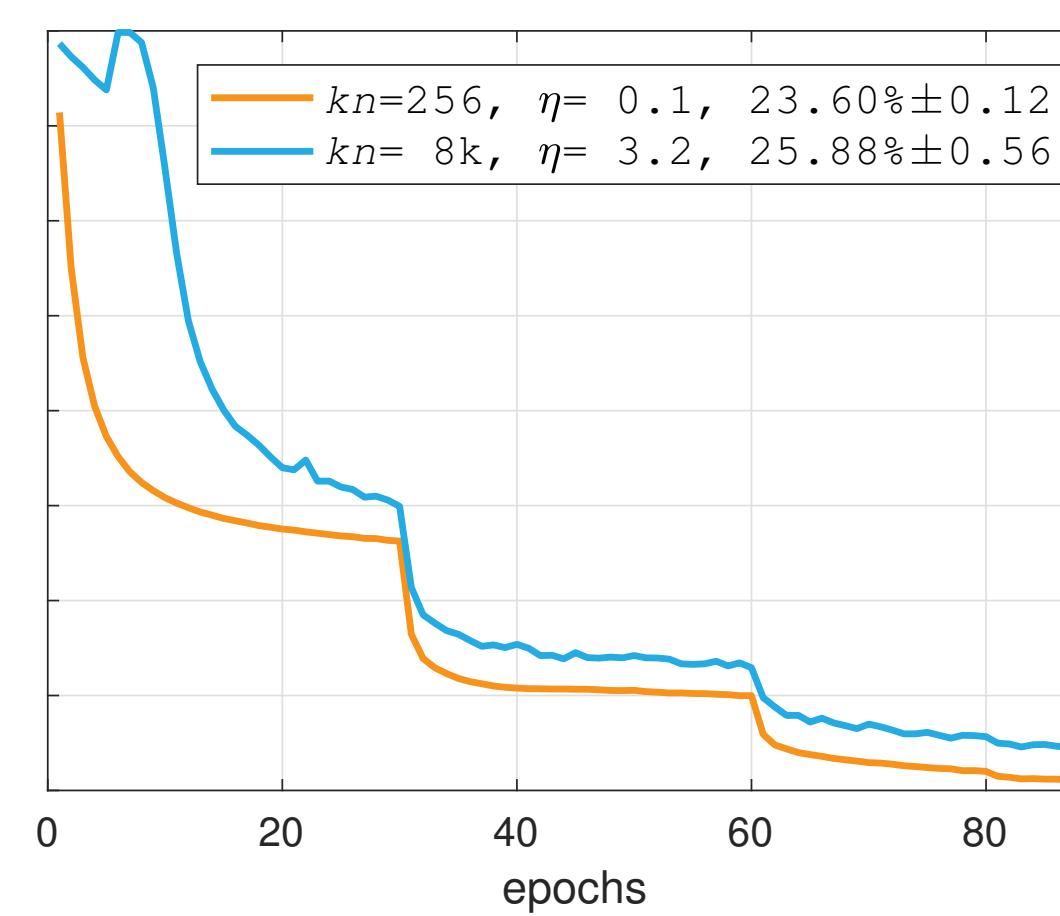
Suggests that if $\nabla l(x, w_t) \approx \nabla l(x, w_{t+j})$ for $j < k$ then minibatch SGD with size n and learning rate η can be approximated by large mini batch SGD with size kn if the learning rate is also scaled to $k\eta$

When does $\nabla l(x, w_t) \approx \nabla l(x, w_{t+j})$ not hold?

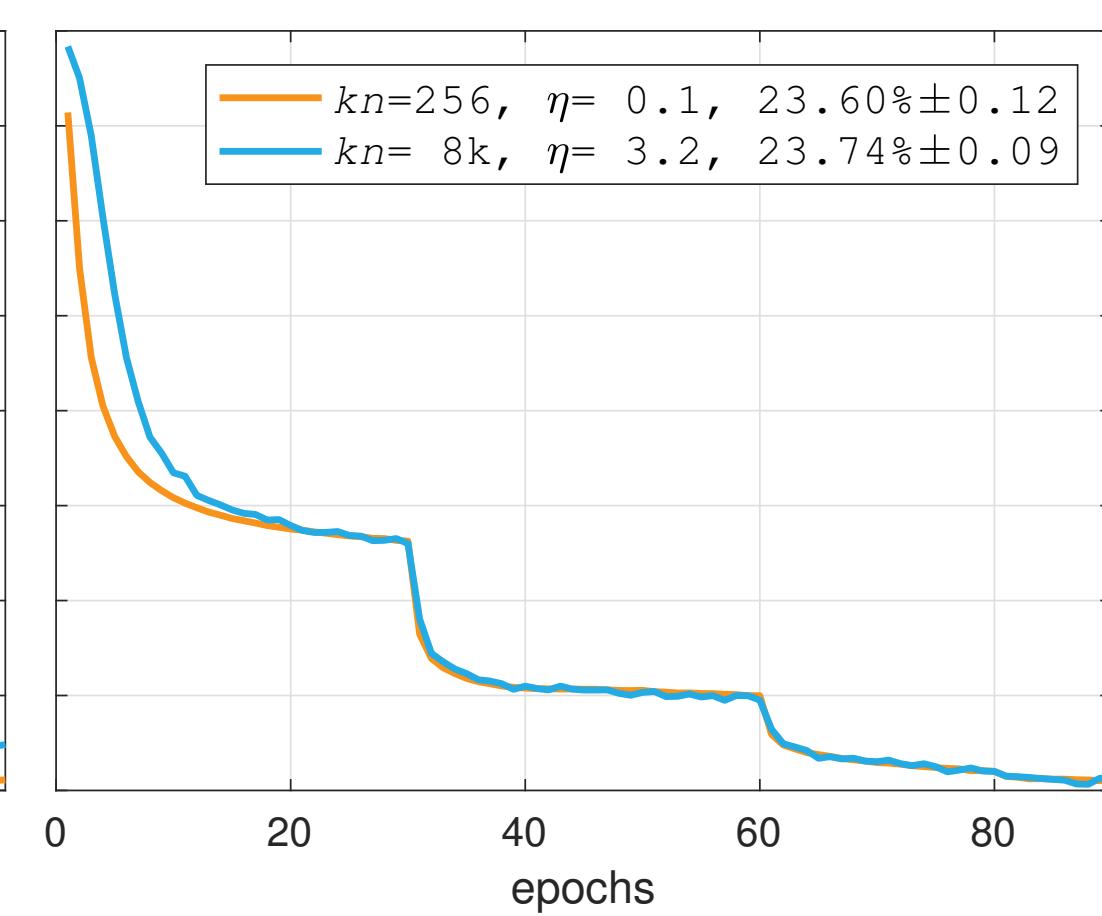
- At beginning of training
 - Suggests starting training with smaller learning rate (learning rate “warmup”)
- When minibatch size begins to get too large (there is a limit to scaling minibatch size)



(a) no warmup



(b) constant warmup



(c) gradual warmup

Minibatch size 256 (orange) vs. 8192 (blue)

Summary: training large networks in parallel

- Many cluster systems rely on data-parallel training with asynchronous update to efficiently use clusters of commodity machines
 - Modification of SGD algorithm to meet constraints of modern parallel systems
 - Effects on convergence are problem dependent and not particularly well understood
 - Efficient use of fast interconnects may provide alternative to these methods (facilitate tightly orchestrated solutions much like supercomputing applications)
- Modern DNN designs (with fewer weights), large minibatch sizes, careful learning rate schedules enable scalability without asynchronous execution on commodity clusters
- High-performance training of deep networks is an interesting example of constant iteration of algorithm design and parallelization strategy (a key theme of this course!)