



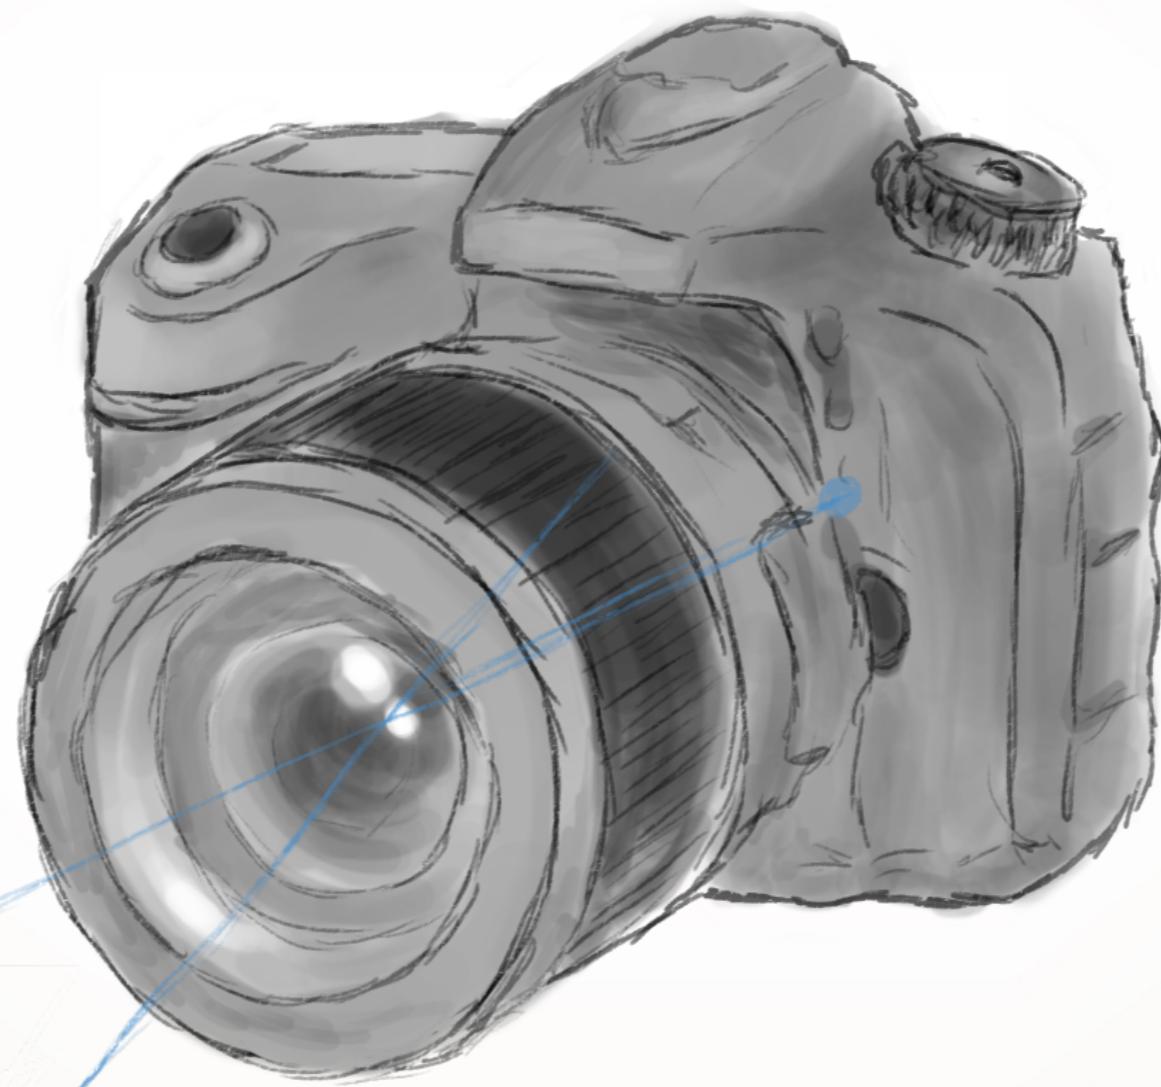
Deep Joint Demosaicking and Denoising

Michaël Gharbi¹ gharbi@mit.edu
&

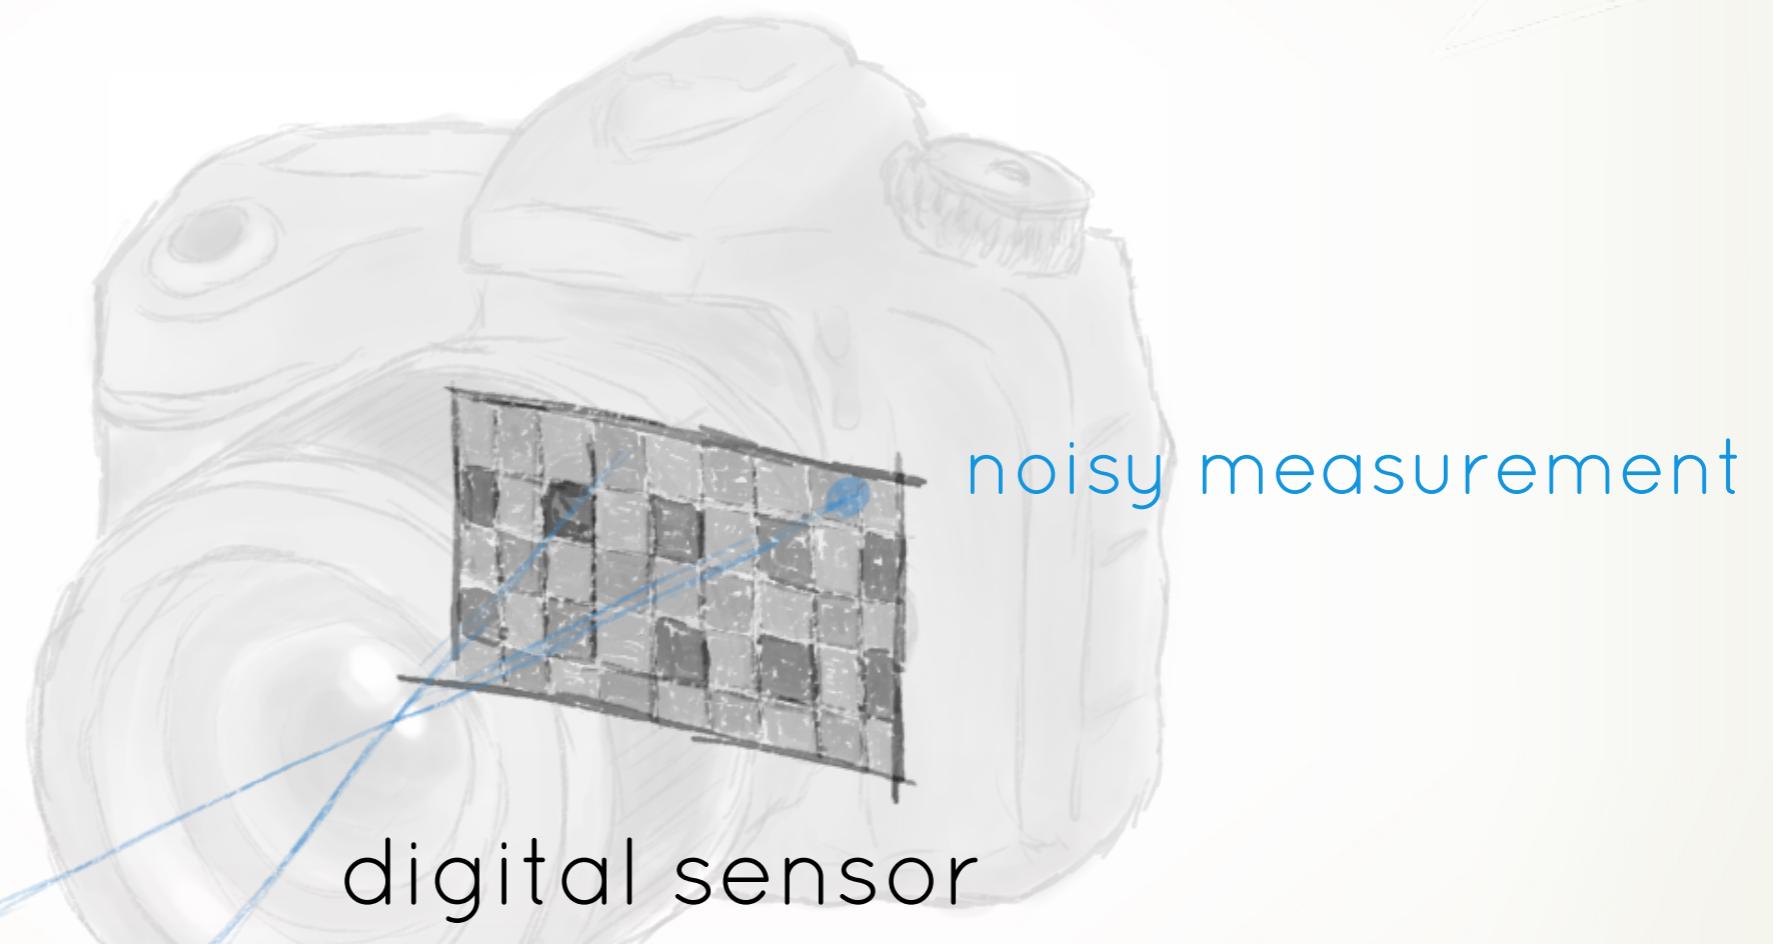
G. Chaurasia¹ S. Paris² F. Durand¹



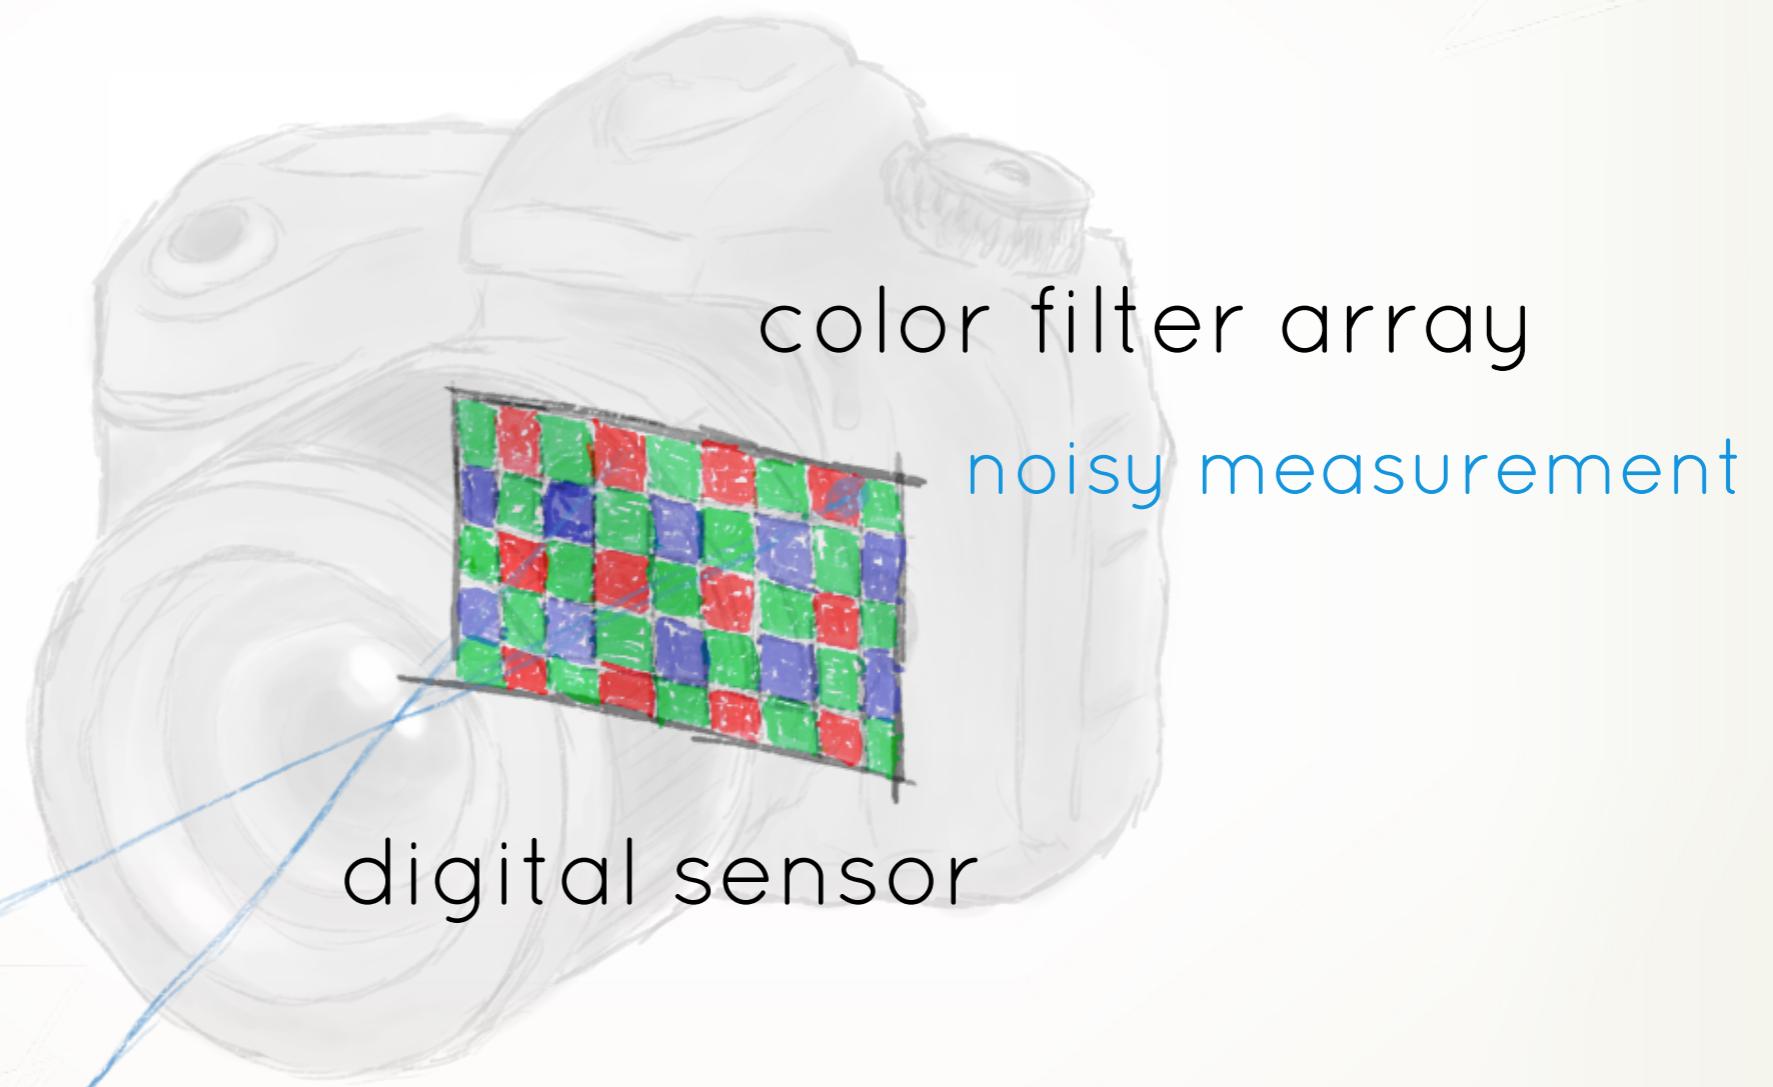
Digital imaging pipeline



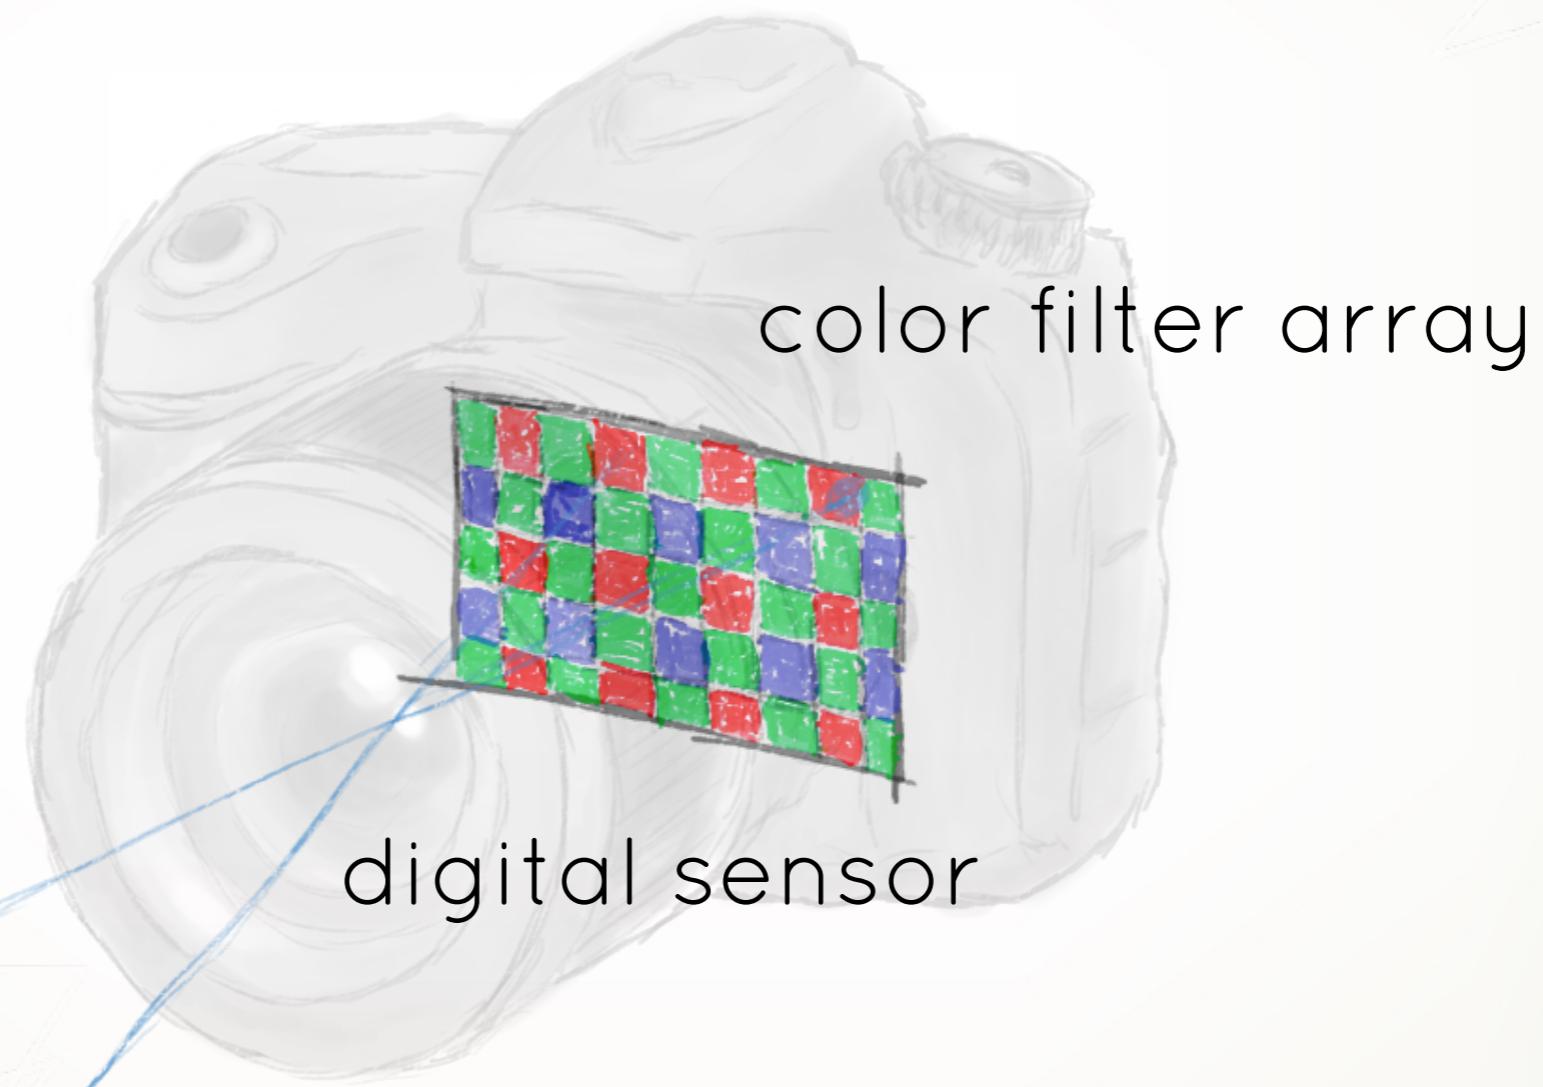
Digital imaging pipeline



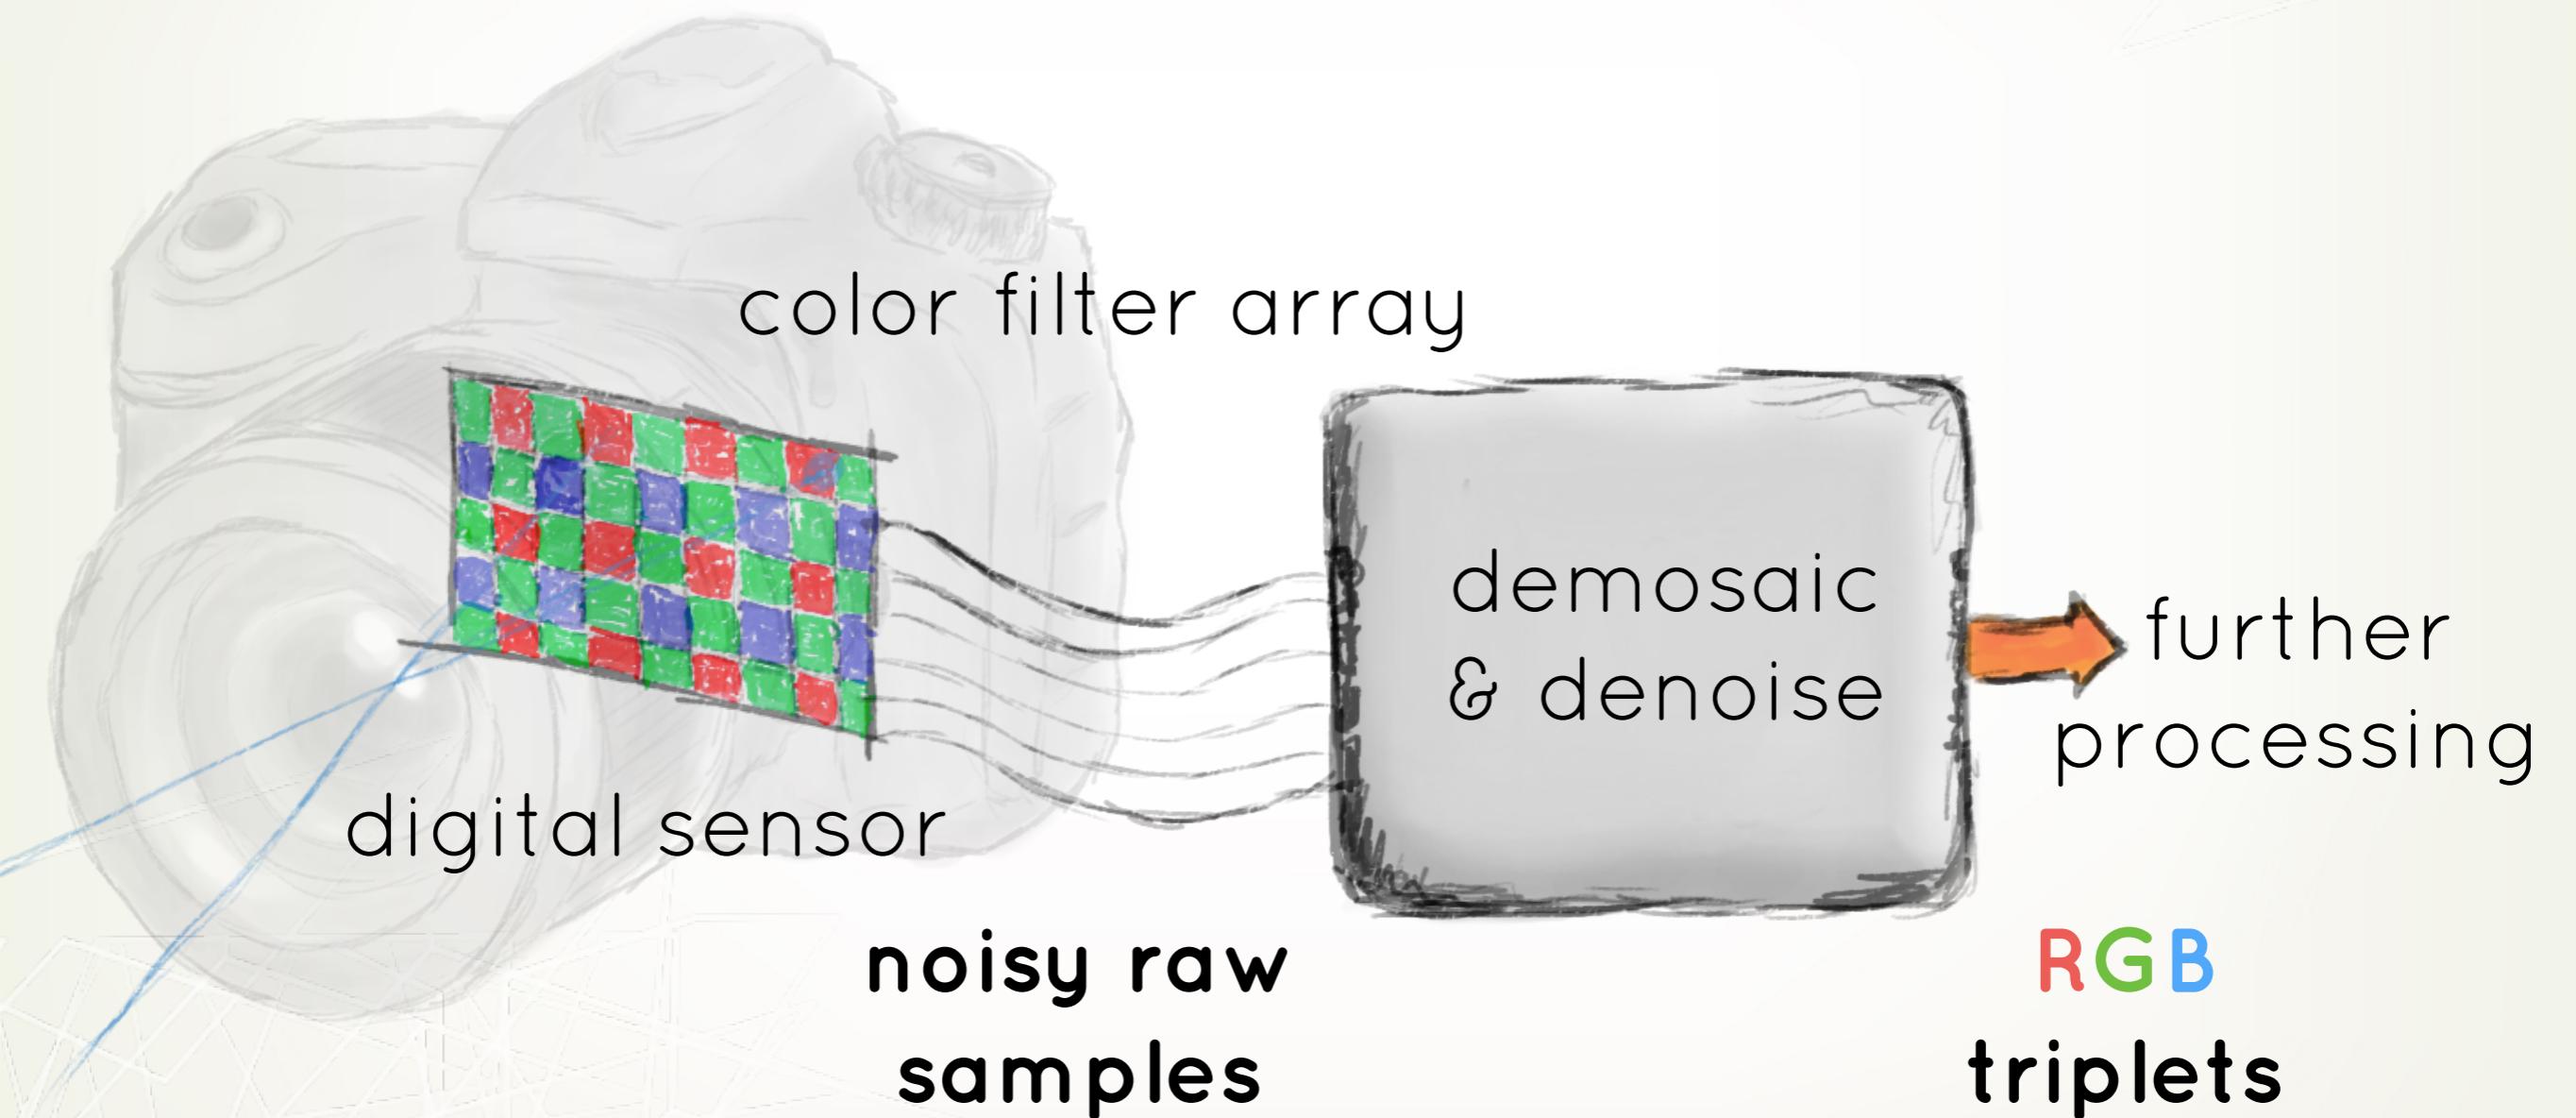
Digital imaging pipeline



Digital imaging pipeline



Digital imaging pipeline

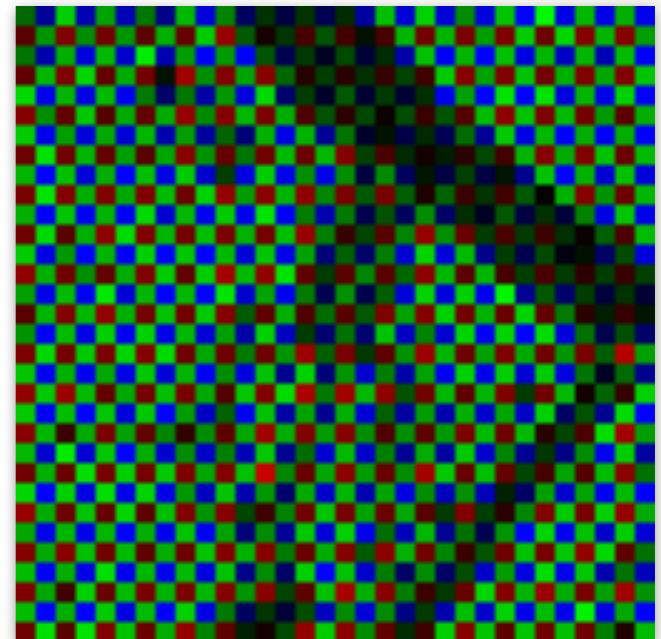




Demosaicking-denoising:
critical first stage!

Ill-posed problem

- **incomplete** color information: mosaic
 - 3 unknowns per input sample
- **noisy** sensor measurements
 - photon, thermal noise
- **mis-aligned** samples
 - 3 interdependent interpolations
 - spatial multiplexing (e.g. Bayer)



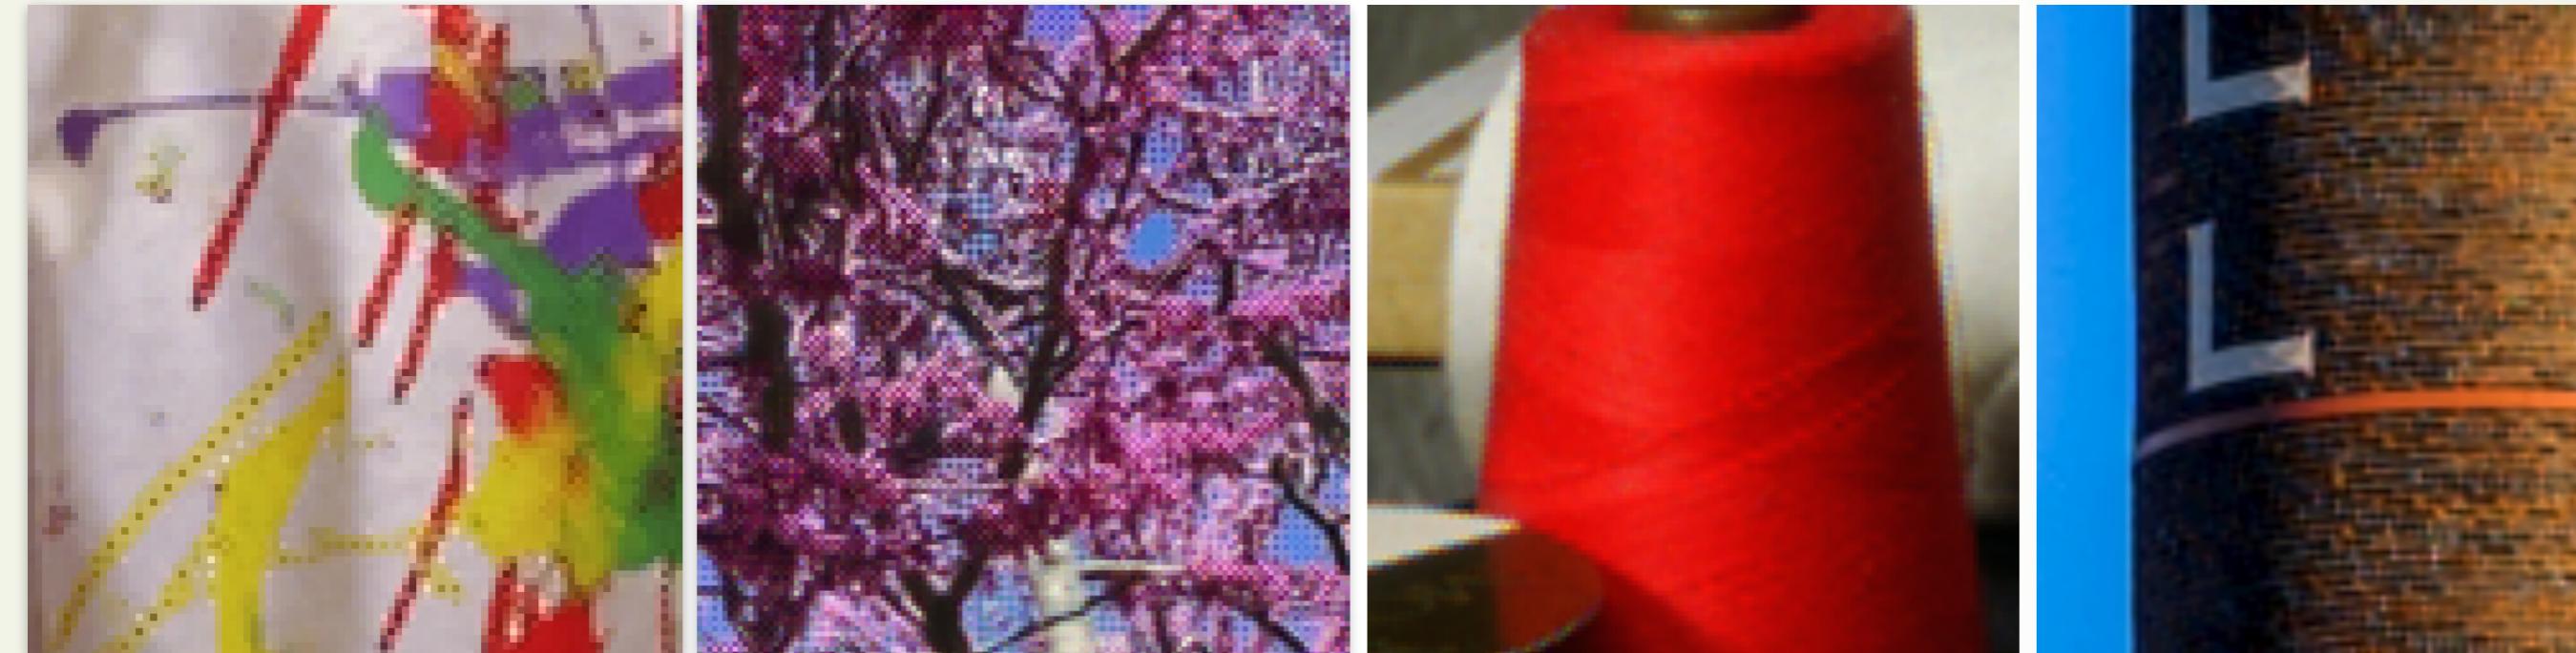
Previous work

- **faster**, traditional methods
 - sequential [Park 2009, Akiyama 2015]
 - filter design [Laroche 1994, Li 2008]
 - ad-hoc post-processing [Hirakawa 2005]
- **more accurate**, modern approaches:
 - joint demosaicking-denoising [Condat 2012]
 - non-local priors [Buades 2009, Zhang 2011]
 - global optimization [Heide 2014]
 - machine-learning [Kashabi 2014, Klatzer 2016]

It mostly works

Artifacts on challenging images

previous methods struggle with them!



zippering
[Buañes 2009]

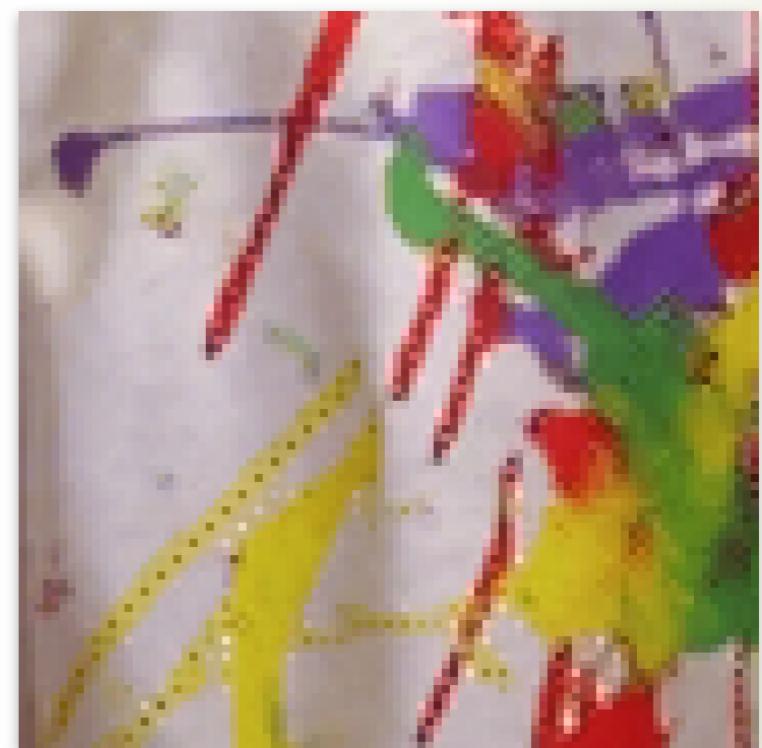
discoloration
[Heide 2014]

blur
[Getreuer 2011]

moiré
(Photoshop)

Rare but catastrophic failures

- **salient** artifacts
 - aliasing, zippering, blur...
 - all cameras prone to them
- **less than 1%** of the pixels
- scarcity **impedes progress**
 - good training data is hard to get



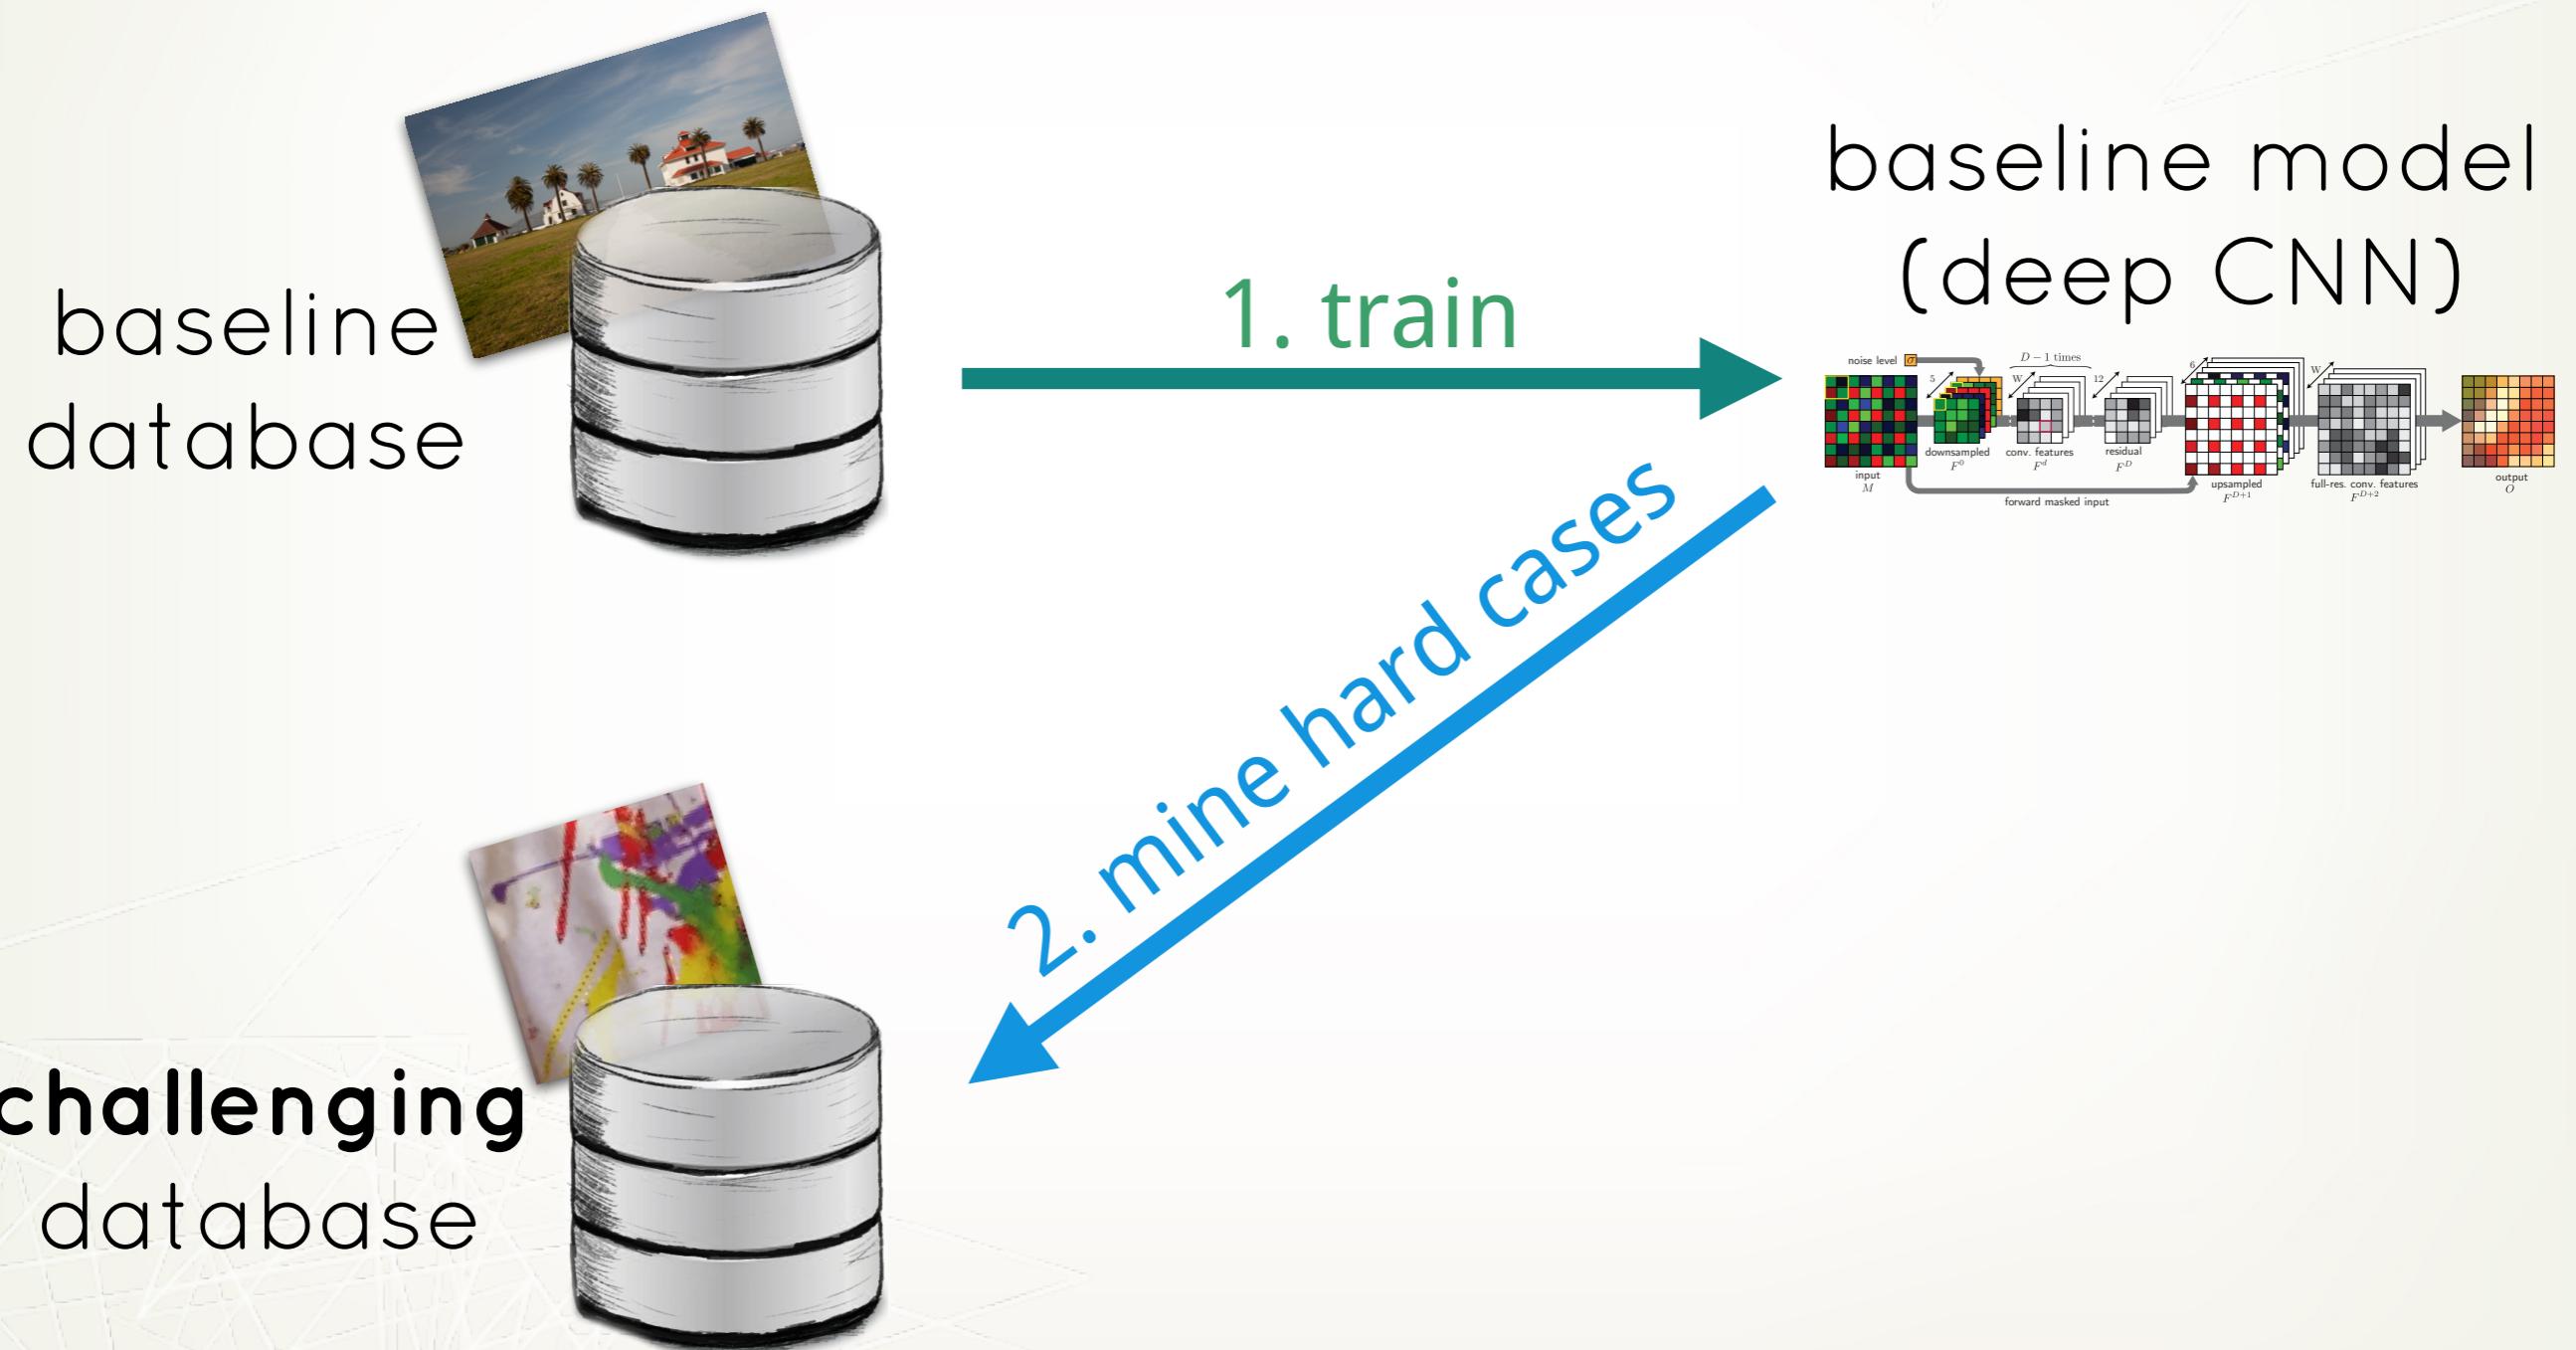


Our approach

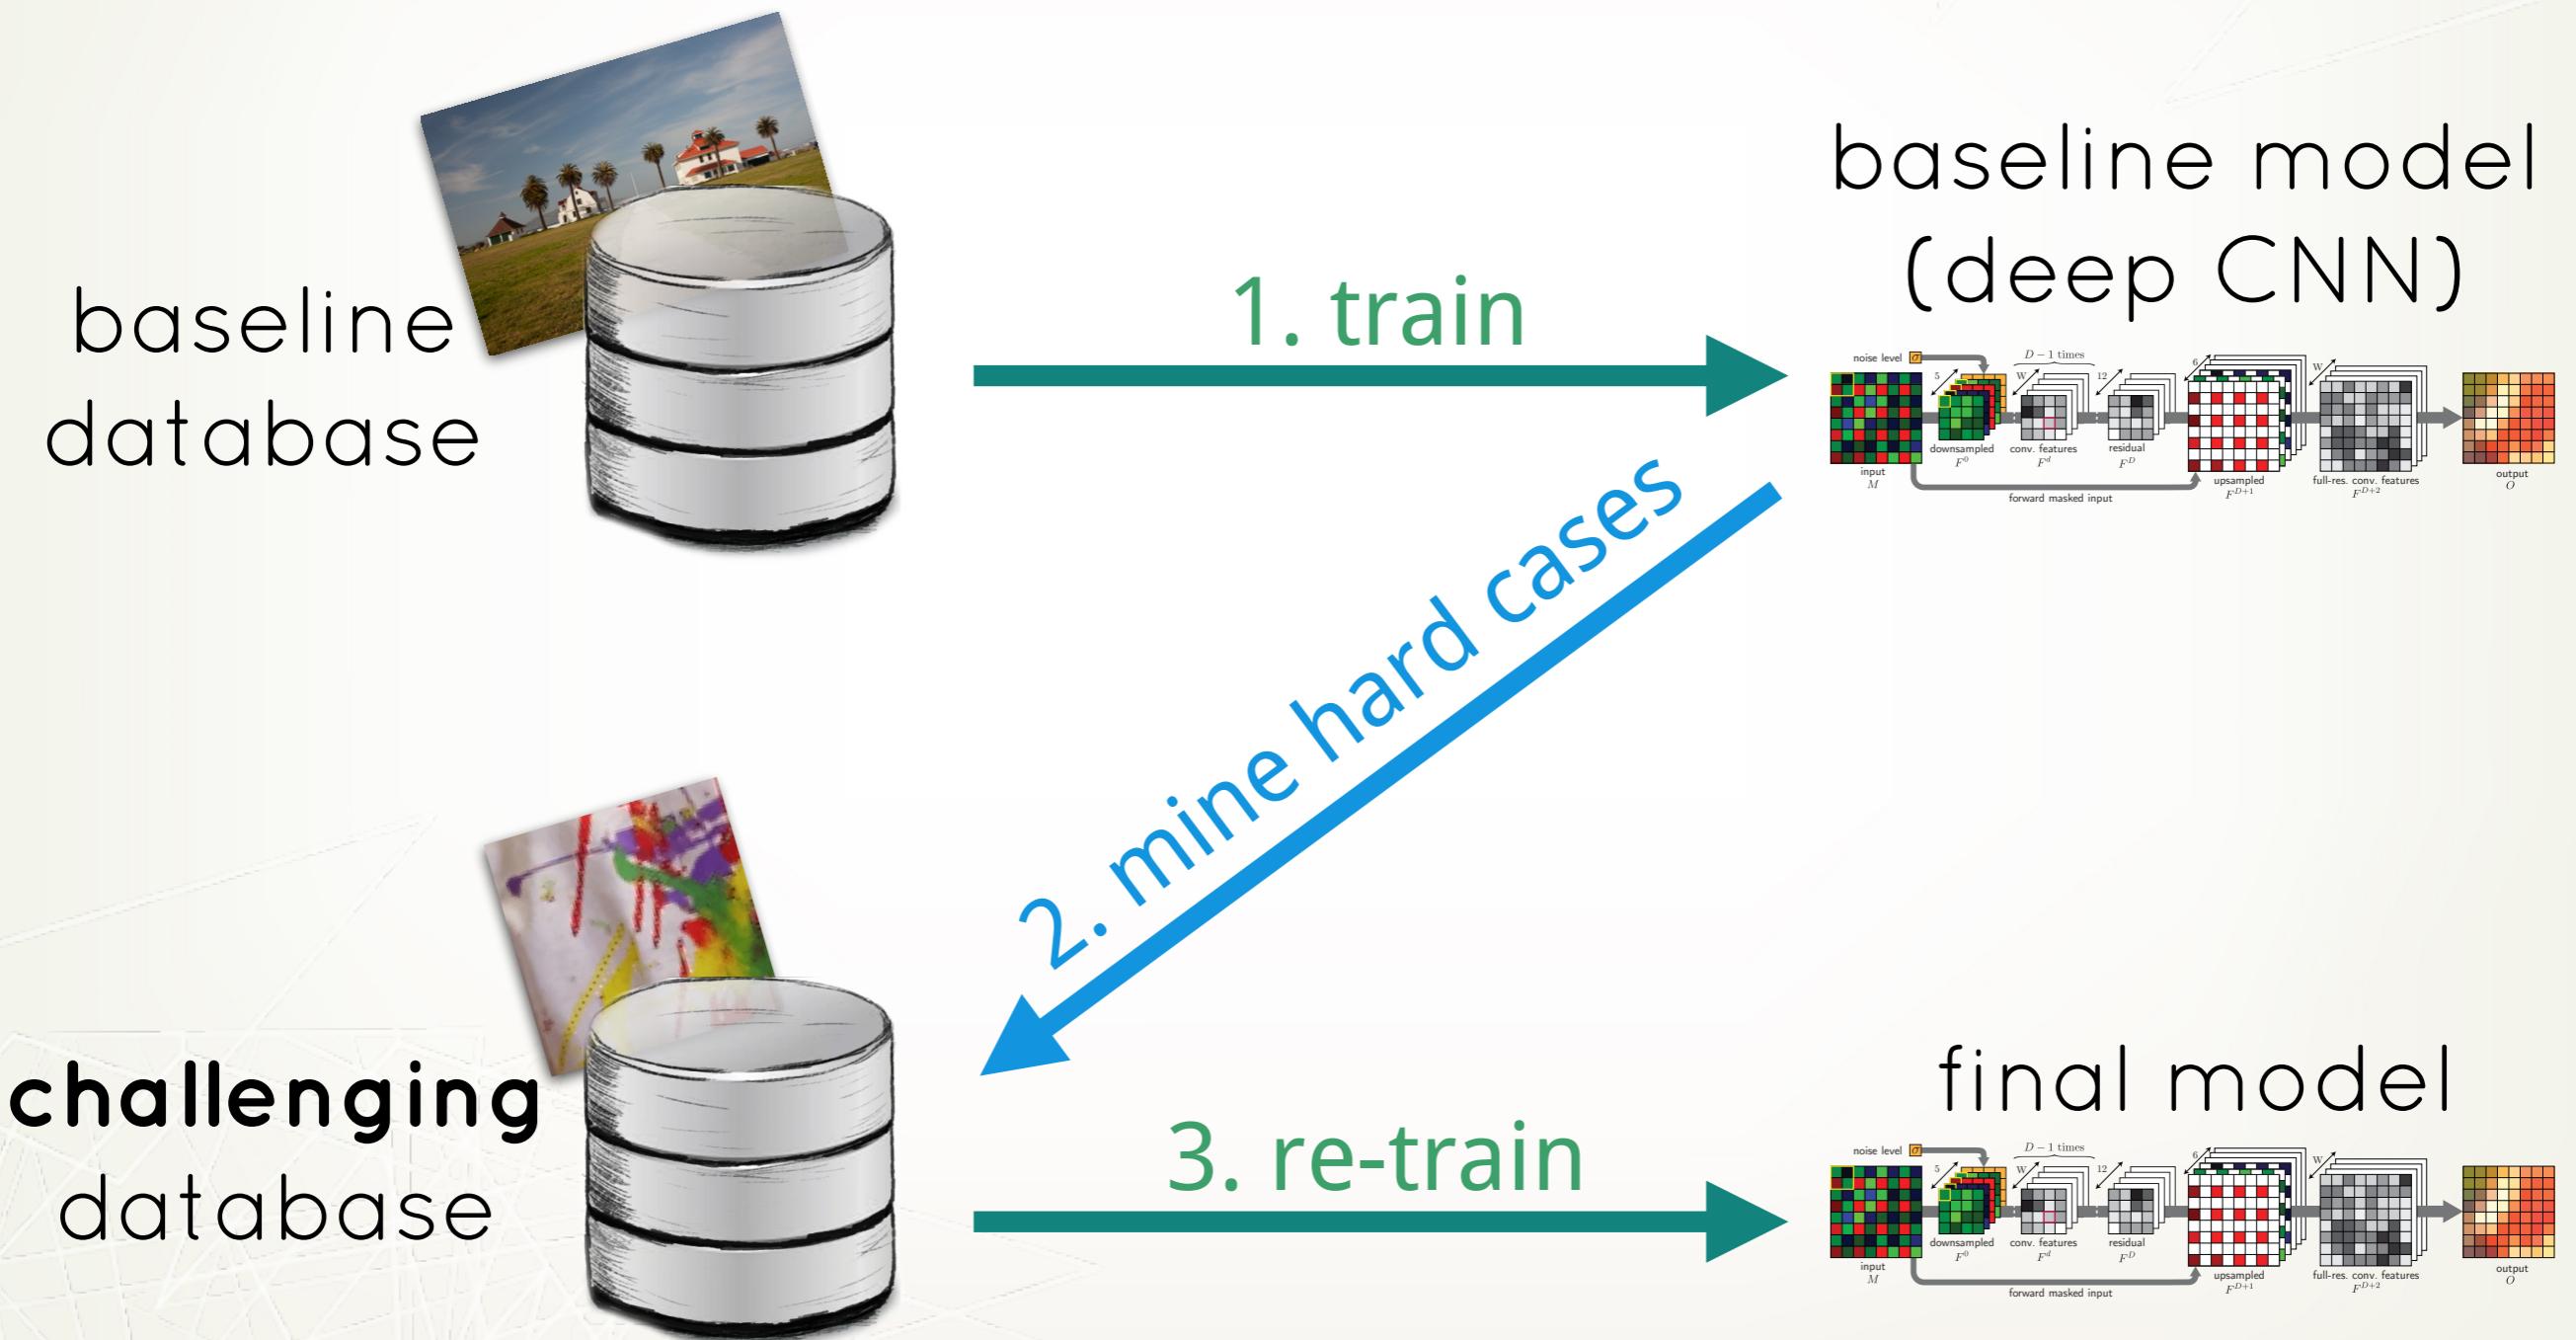
Three-step learning approach to joint demosaicking-denoising



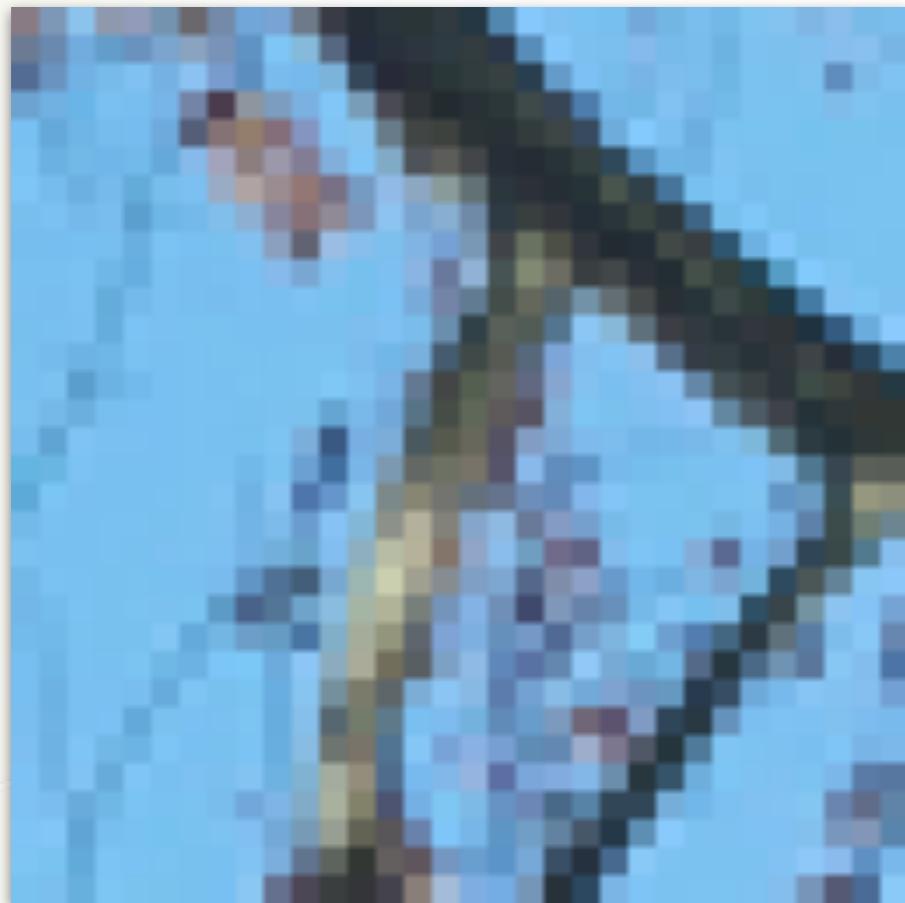
Three-step learning approach to joint demosaicking-denoising



Three-step learning approach to joint demosaicking-denoising



Pseudo ground truth



sRGB image from
the web

Pseudo ground truth

becomes 1 pixel



4x downsample

sRGB image from
the web

Pseudo ground truth

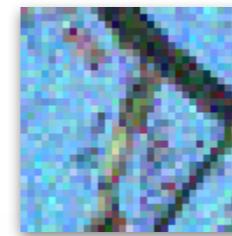
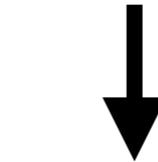
becomes 1 pixel



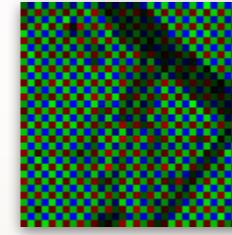
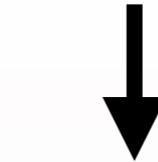
sRGB image from
the web



4x downsample
artifacts and noise
greatly reduced



Gaussian noise
[Jeon 2013]



mosaick



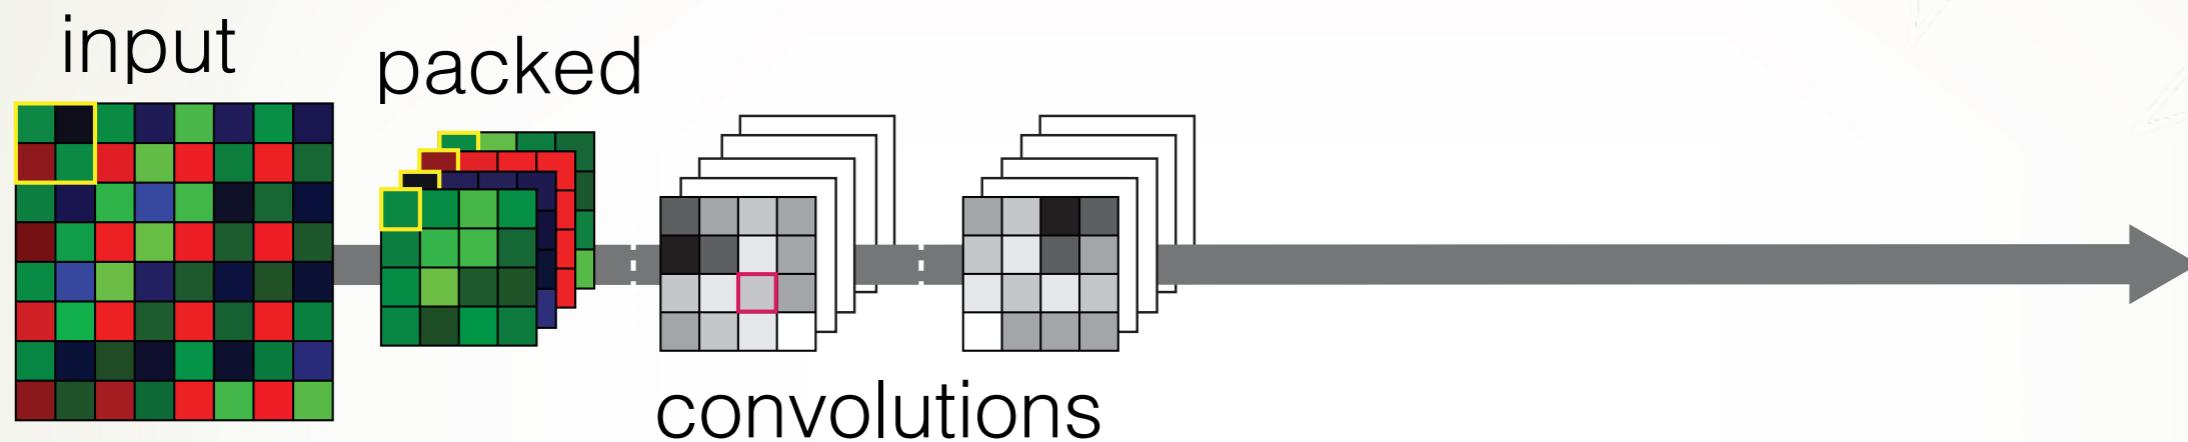
Model architecture

Deep CNN architecture



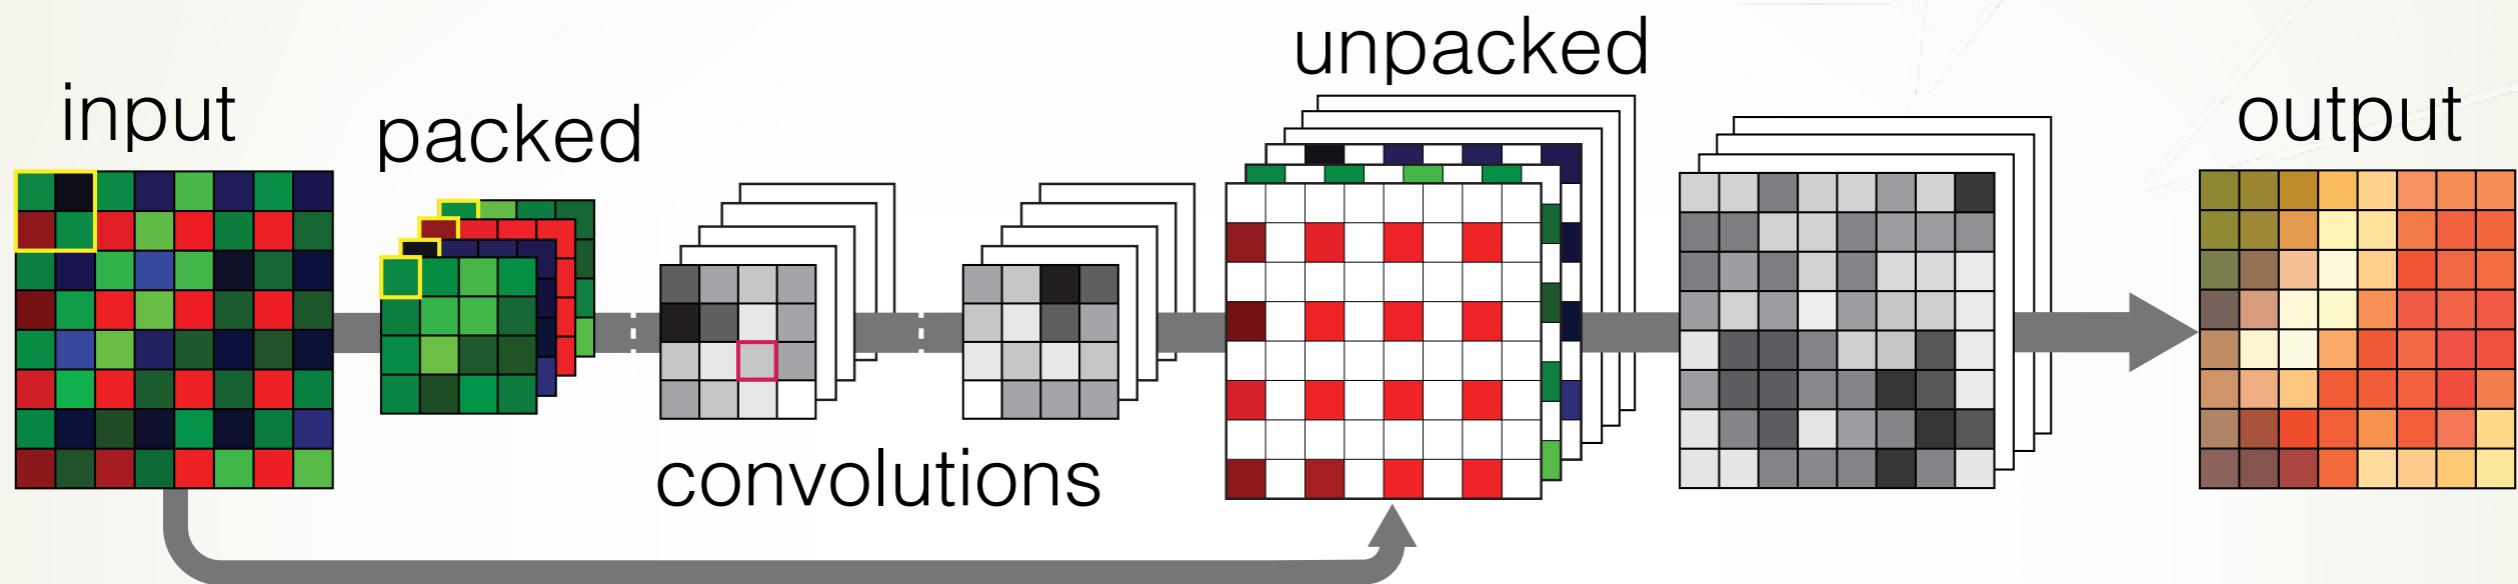
- translation invariance

Deep CNN architecture



- translation invariance
- trainable stack of convolutions
 - large footprint non-linear filter

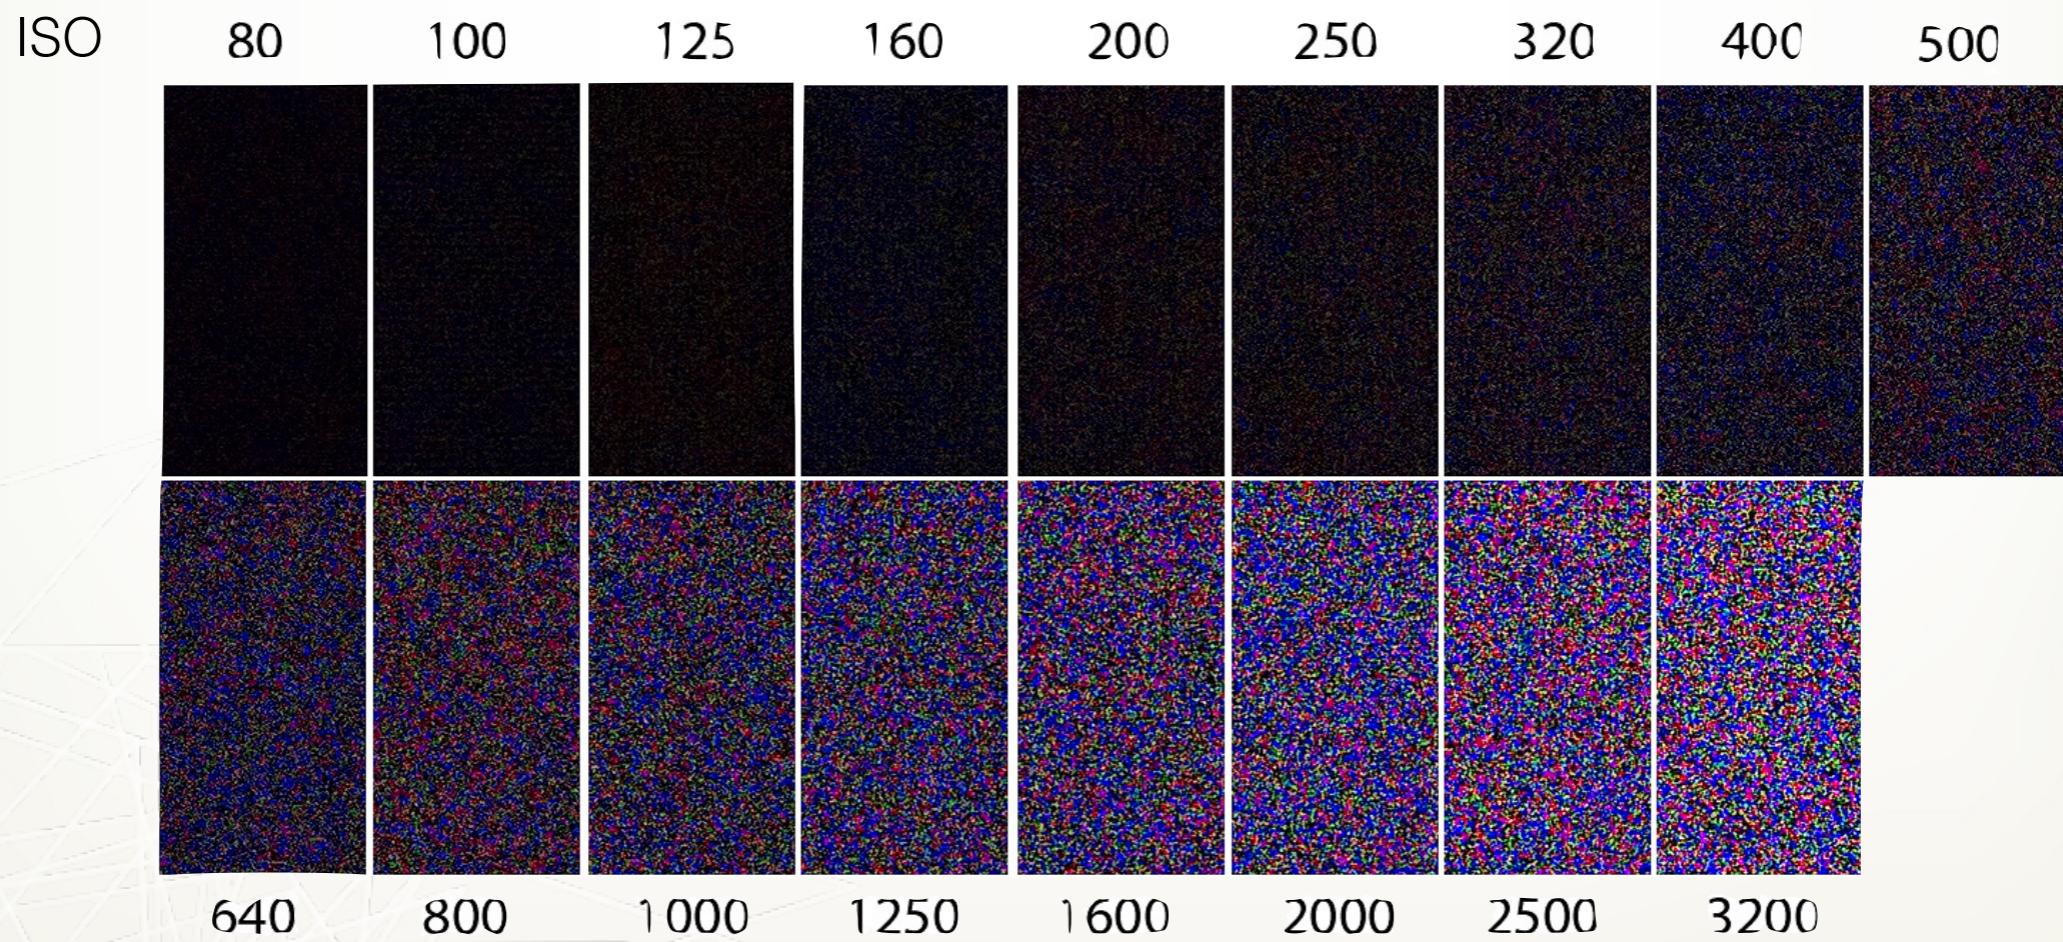
Deep CNN architecture



- translation invariance
- trainable stack of convolutions
 - large footprint non-linear filter
- predict difference from input: residual
 - easier than synthesizing output from scratch
 - similar to ResNet [He 2016]

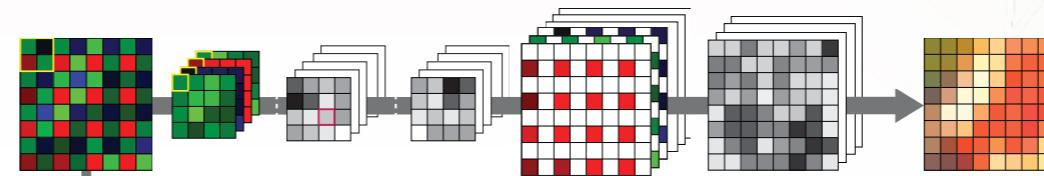
Noise varies with ISO

- need to handle multiple noise levels
- noise characteristics **known** in advance

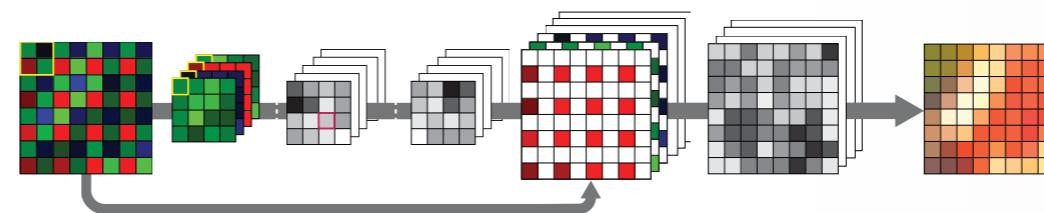


Naive: one model per noise level

noise level $\sigma = 0$

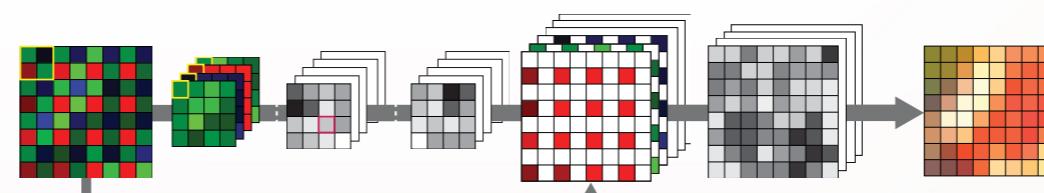


$\sigma = 4$

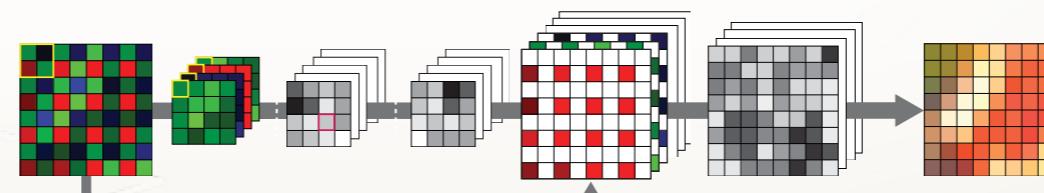


•
•
•

$\sigma = 16$



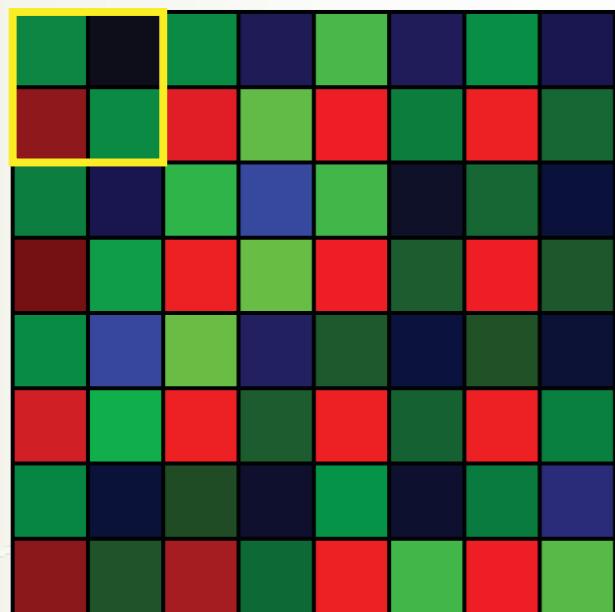
$\sigma = 20$



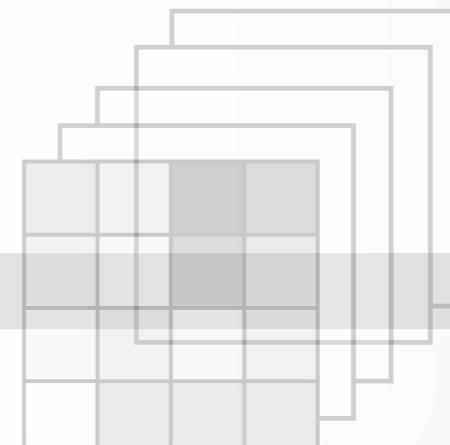
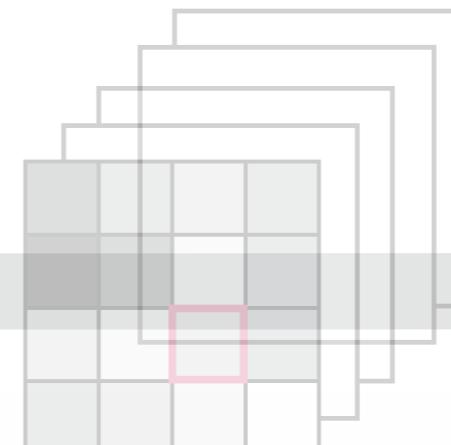
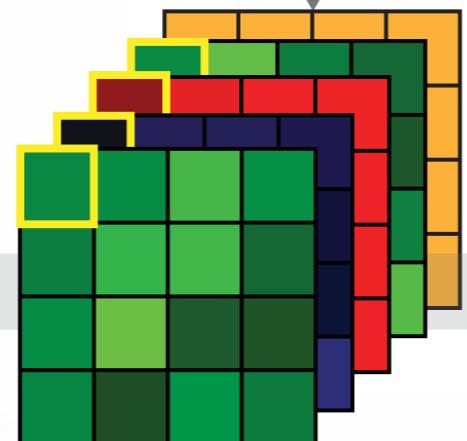
Parametrize by noise level

noise estimate

$$\sigma$$

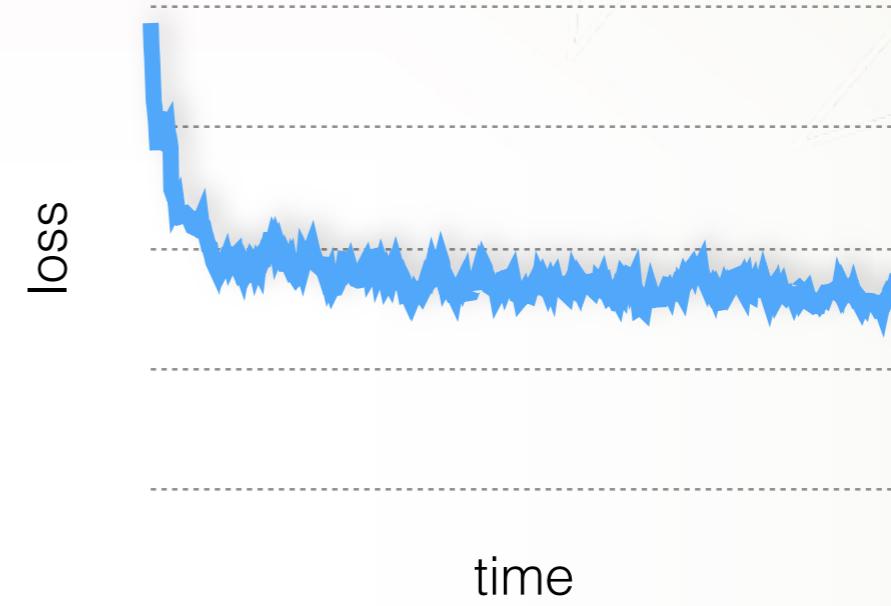


noisy input



Training procedure

- minimize L2 loss
 - Adam [Kingma 2014]
 - 1 week of training
- range of noise levels
 - train jointly on all noise levels
 - random noise variance per training image
 - fixed range $\sigma \in [0, 20]$





Baseline trained on Imagenet

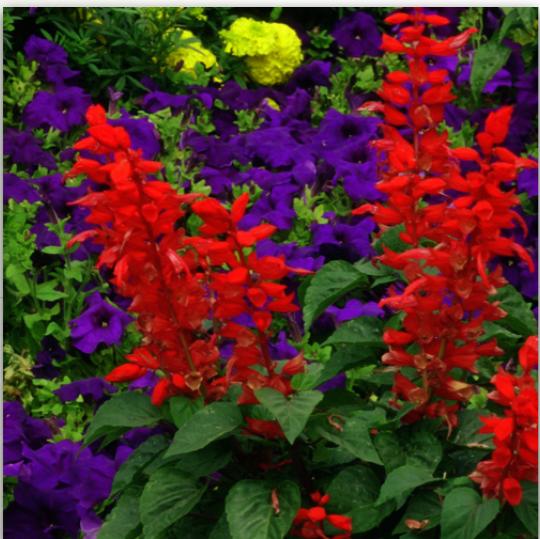
1.5 million images



Imagenet baseline

Standard benchmarks look good...

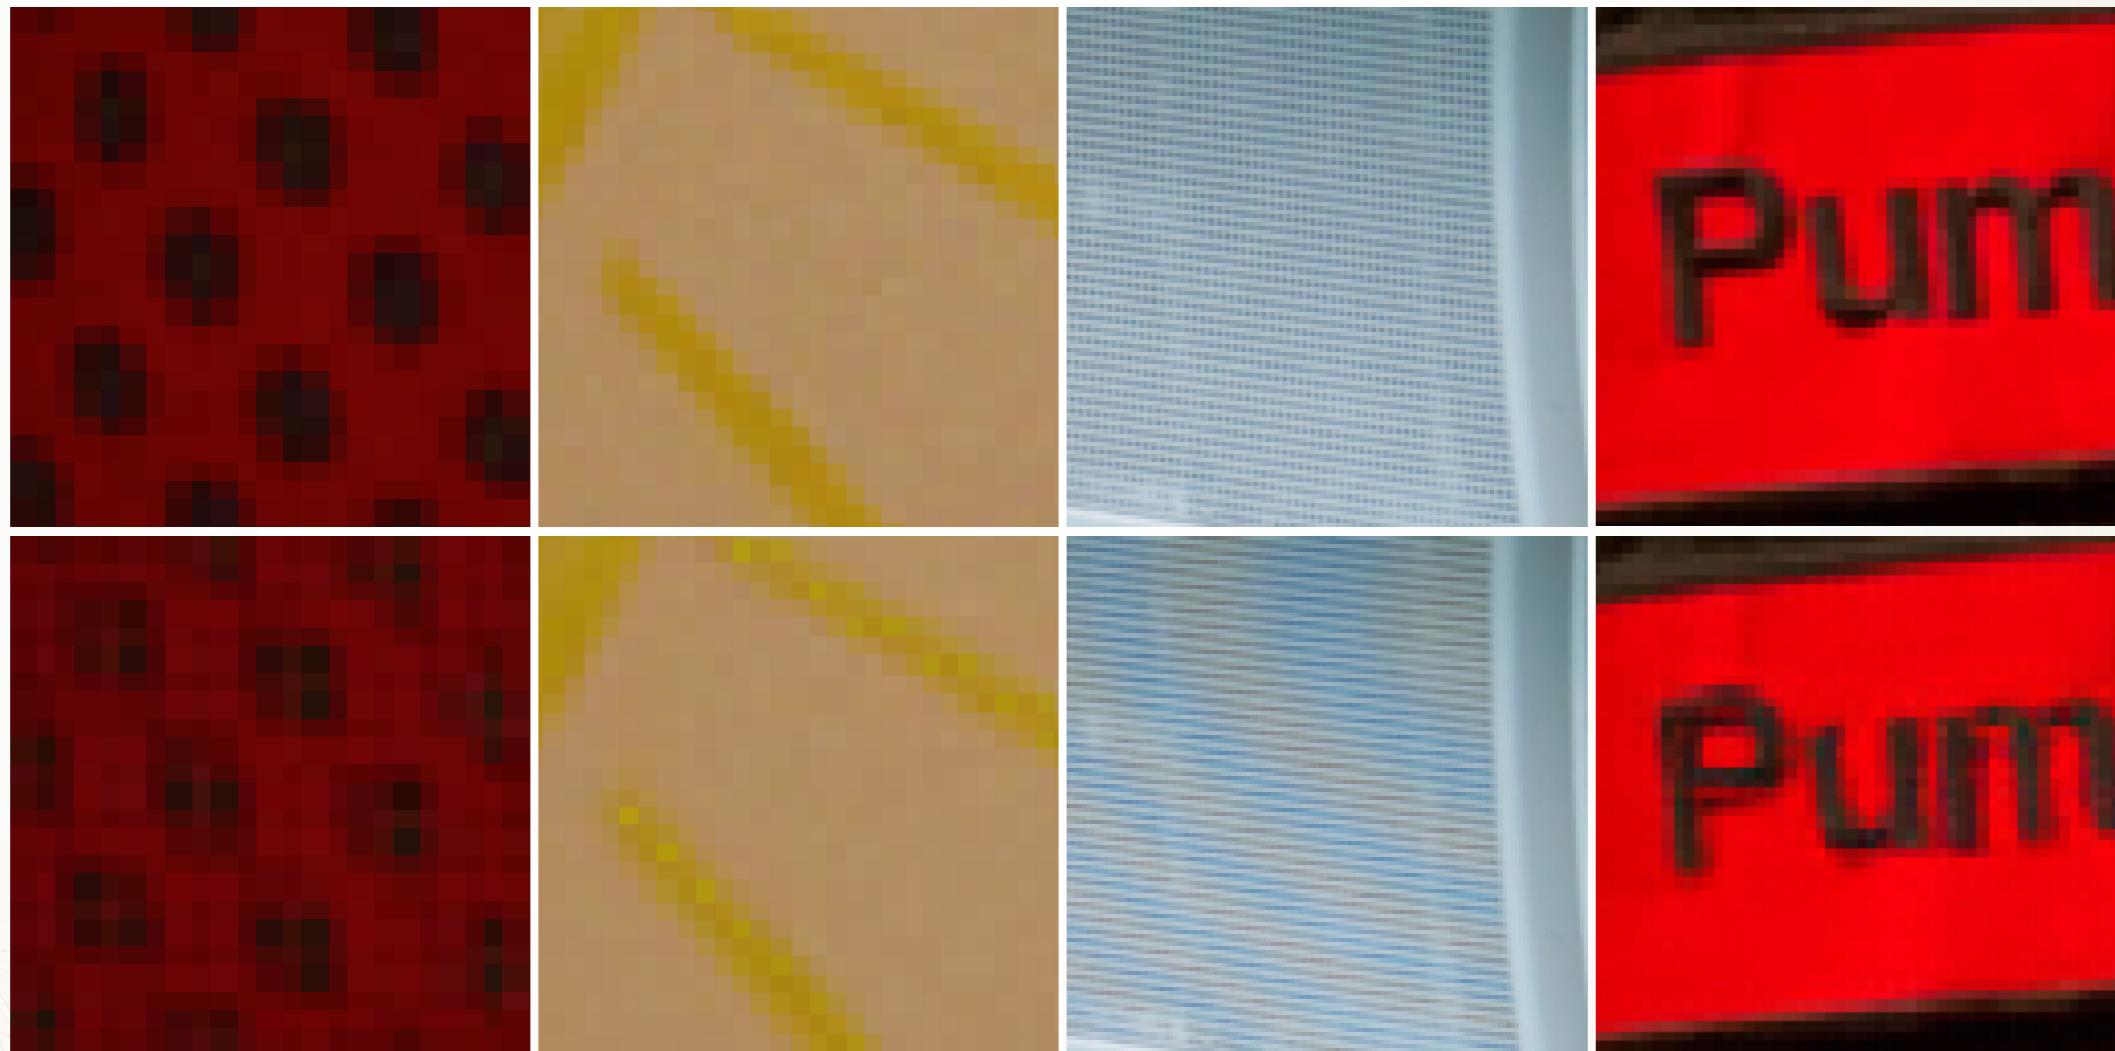
(numbers too)



...but major artifacts remain

Imagenet baseline

ground truth



trained on
imagenet

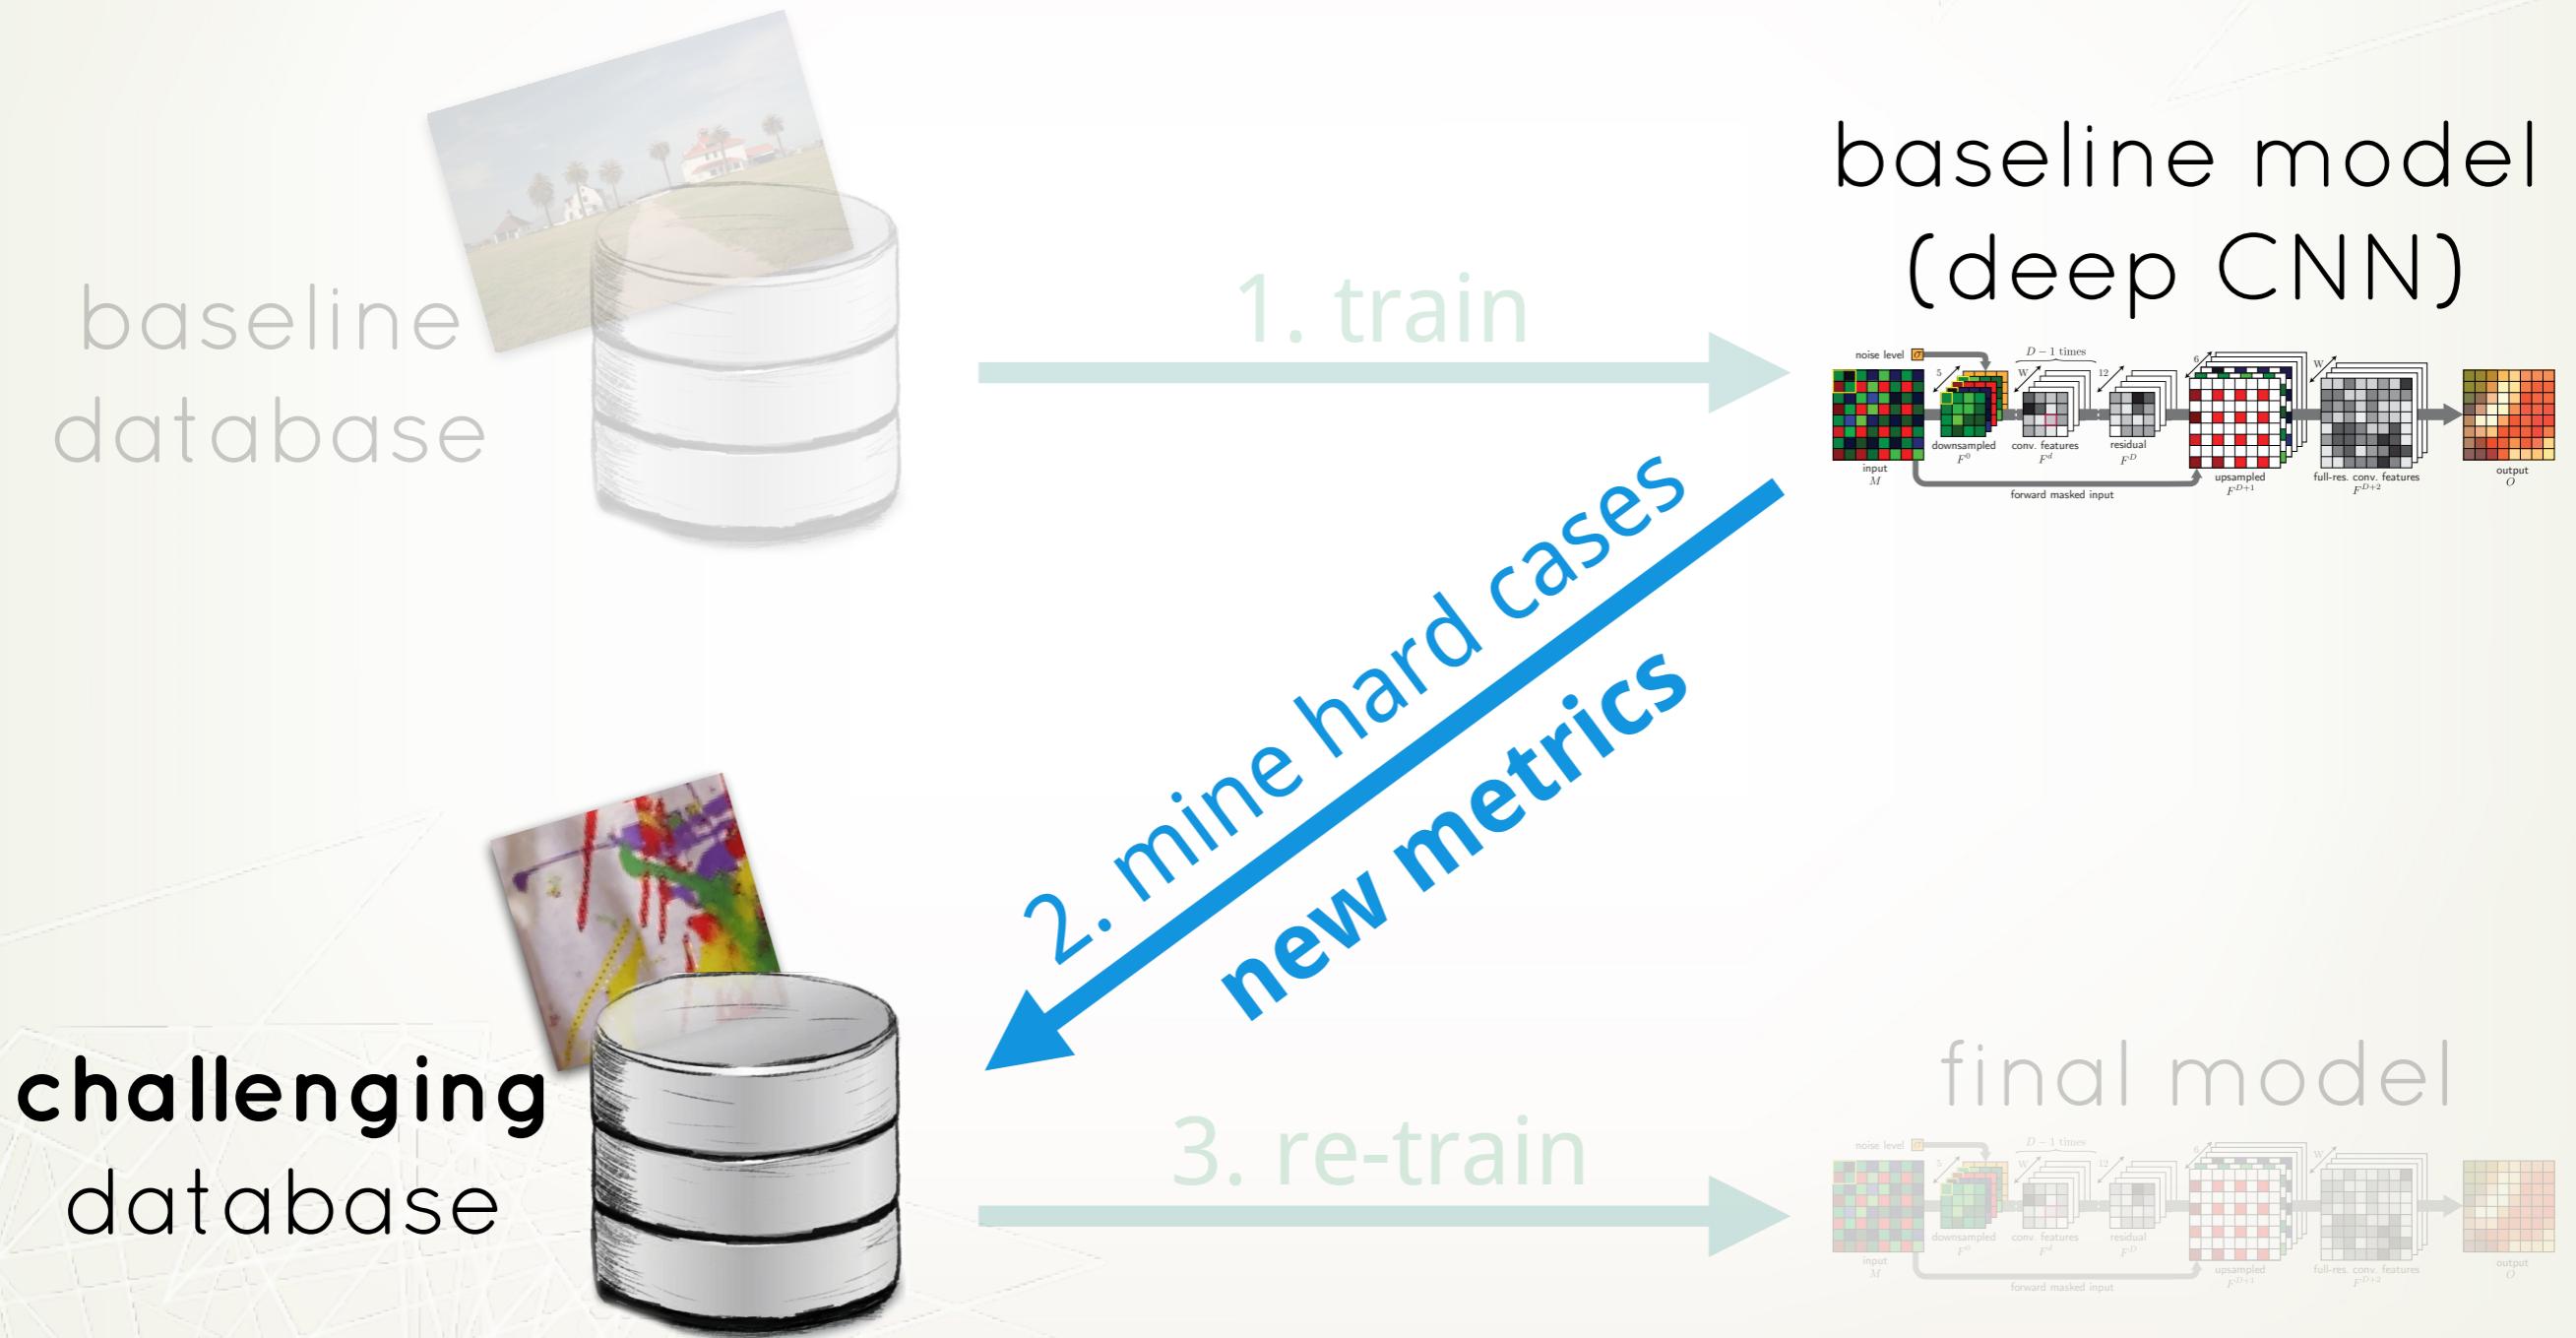
Why does this happen?

- too few challenging images
 - most patches are smooth [Levin 2012]
 - only 1 in 2,000 has artifacts

Why does this happen?

- too few challenging images
 - most patches are smooth [Levin 2012]
 - only 1 in 2,000 has artifacts
- metrics cannot detect artifacts [Sergej 2014]
 - SSIM, MSE: low correlation with human perception

Three-step learning approach



Detecting artifacts

- **analyze** millions of photographs
 - 1 month of scraping, 100 computers

Detecting artifacts

- **analyze** millions of photographs
 - 1 month of scraping, 100 computers
- **process** them with our baseline model
 - we now have before/after ground-truth

Detecting artifacts

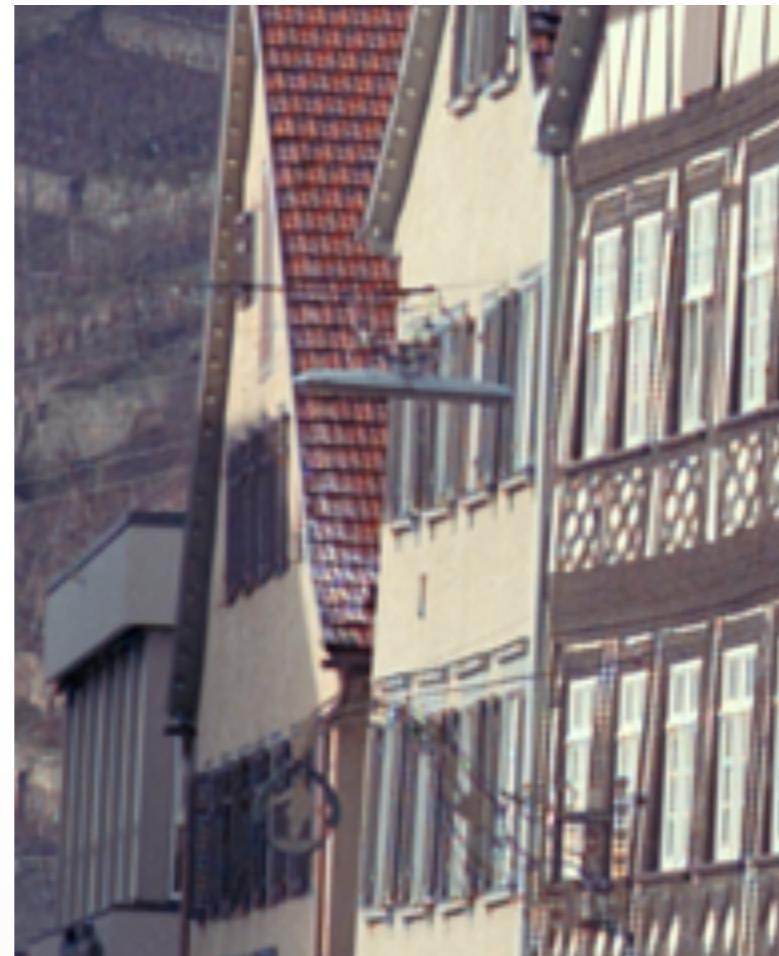
- **analyze** millions of photographs
 - 1 month of scraping, 100 computers
- **process** them with our baseline model
 - we now have before/after ground-truth
- **rejection-sampling**: keep hard cases
 - **2 criteria**: luminance errors, color moiré

Finding luminance errors

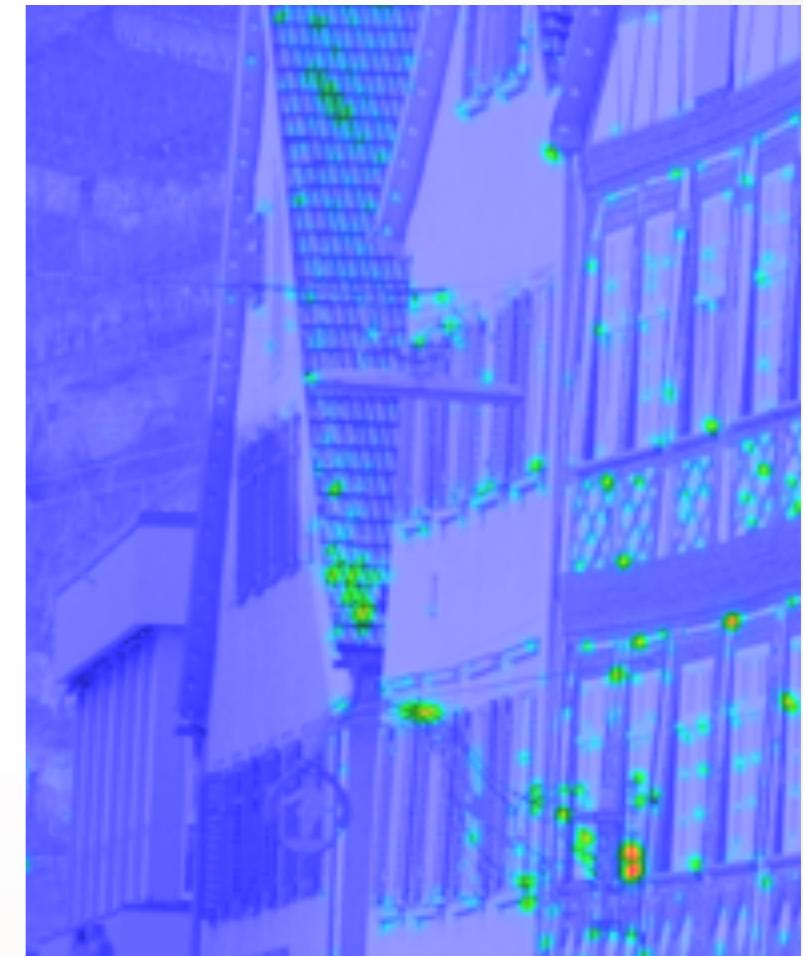
use HDR-VDP, perceptual model [Mantiuk 2012]



ref.



corrupted



error map

Exposing color moiré

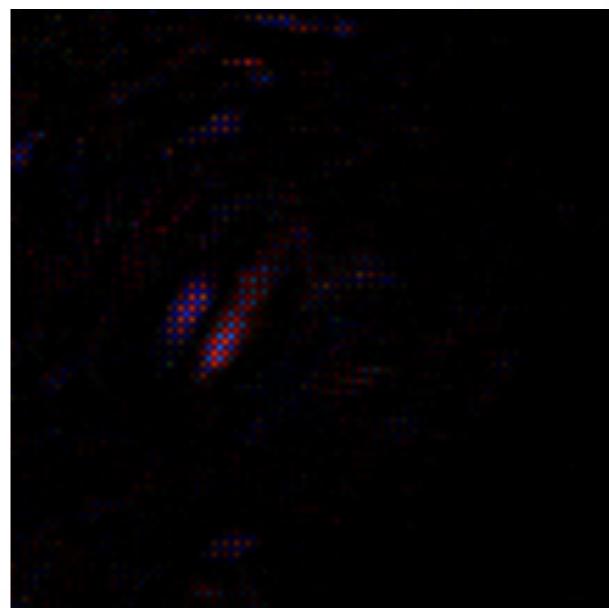
detect added low-frequency chroma



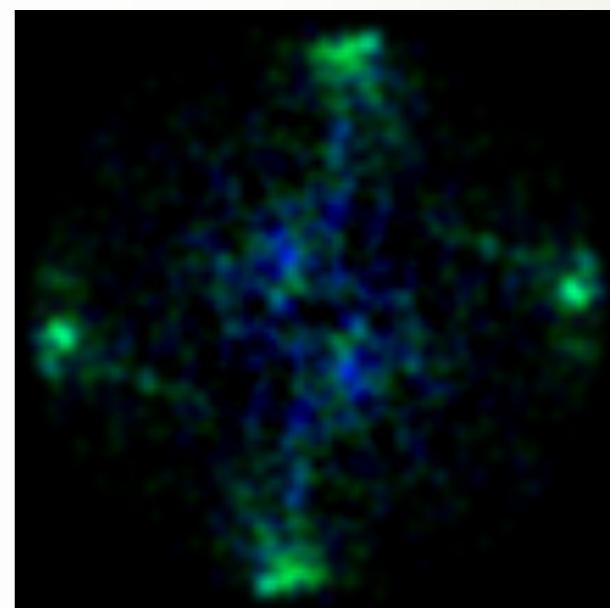
ref.



corrupted

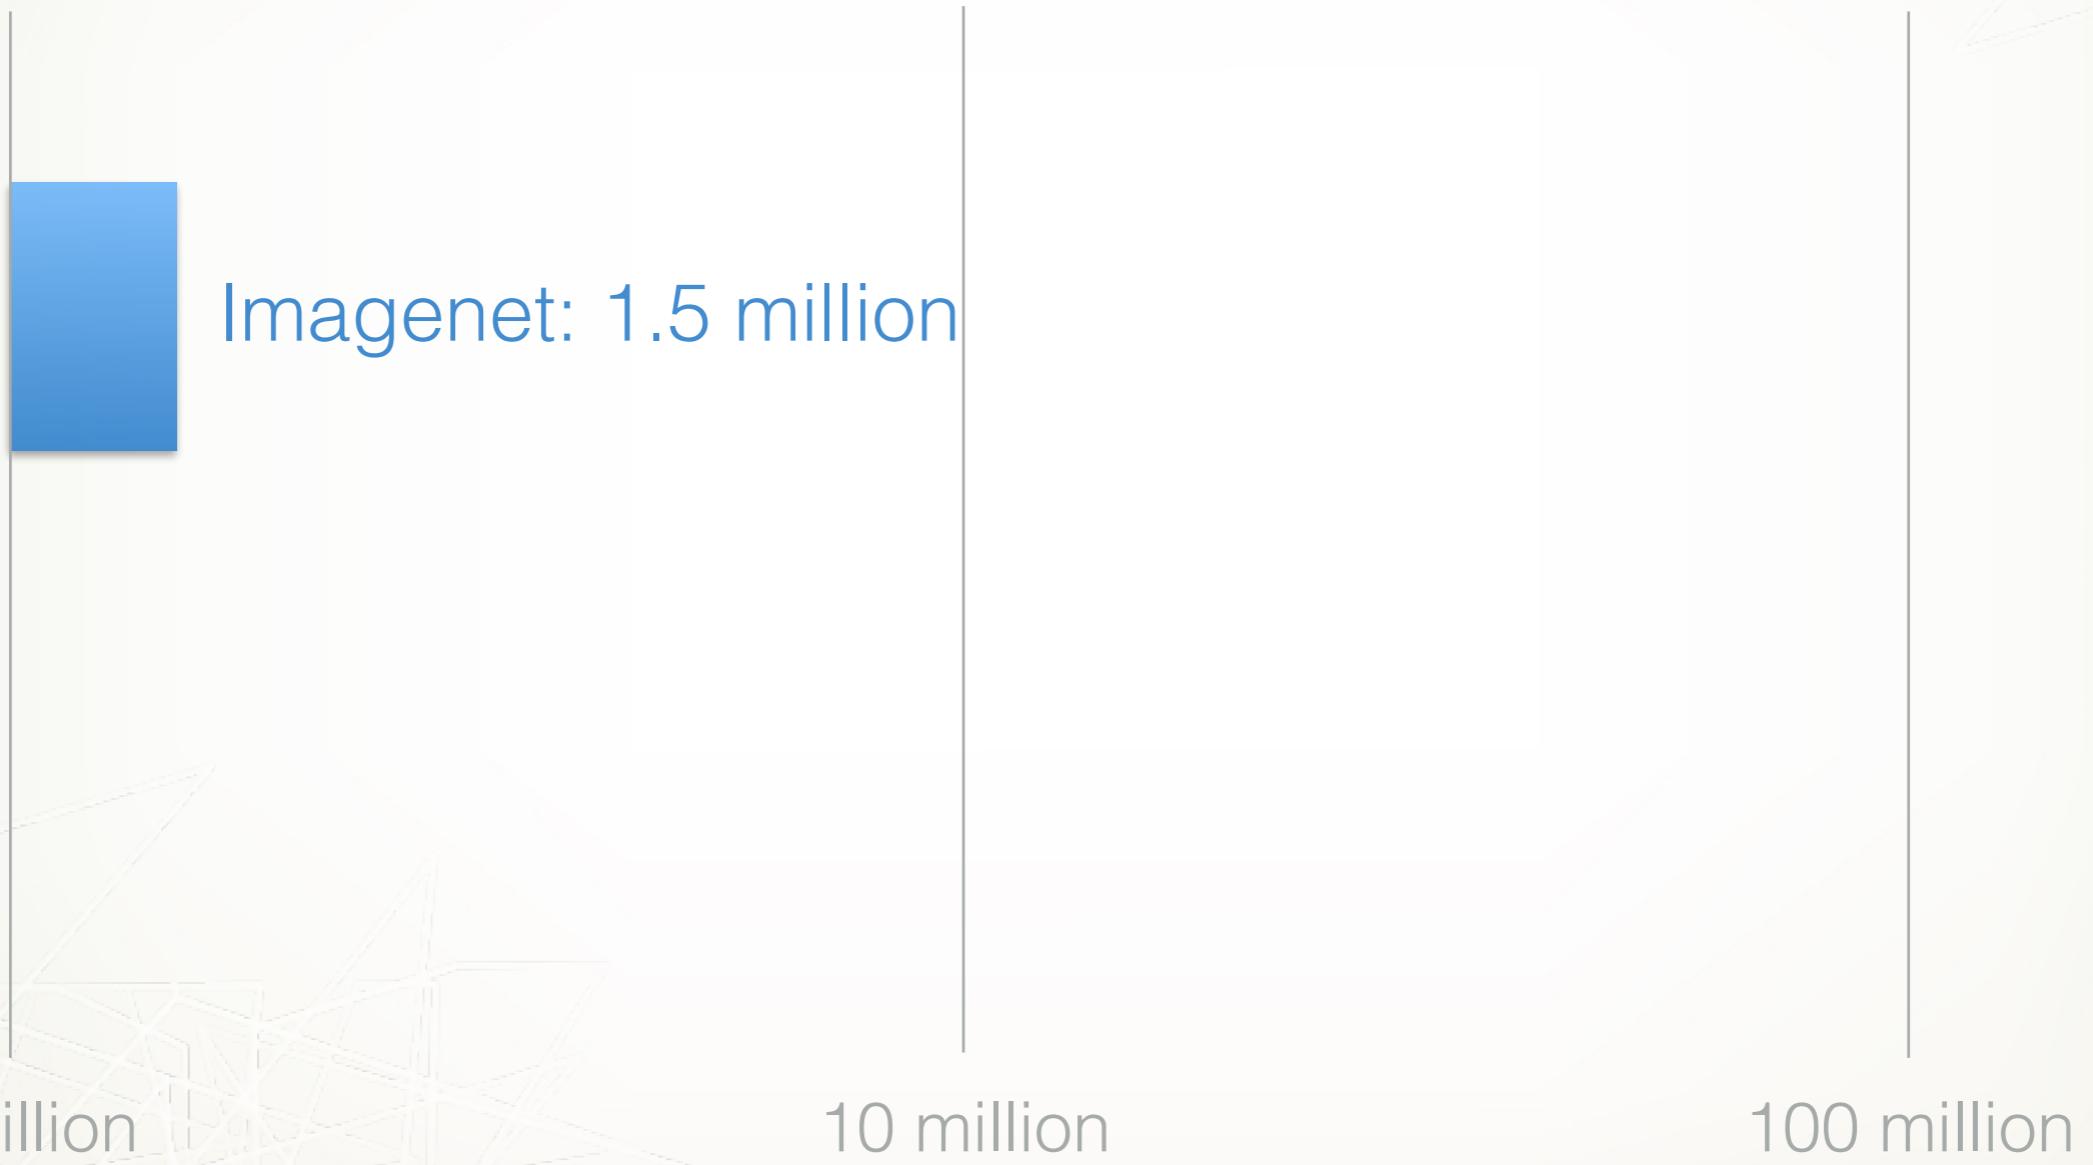


difference

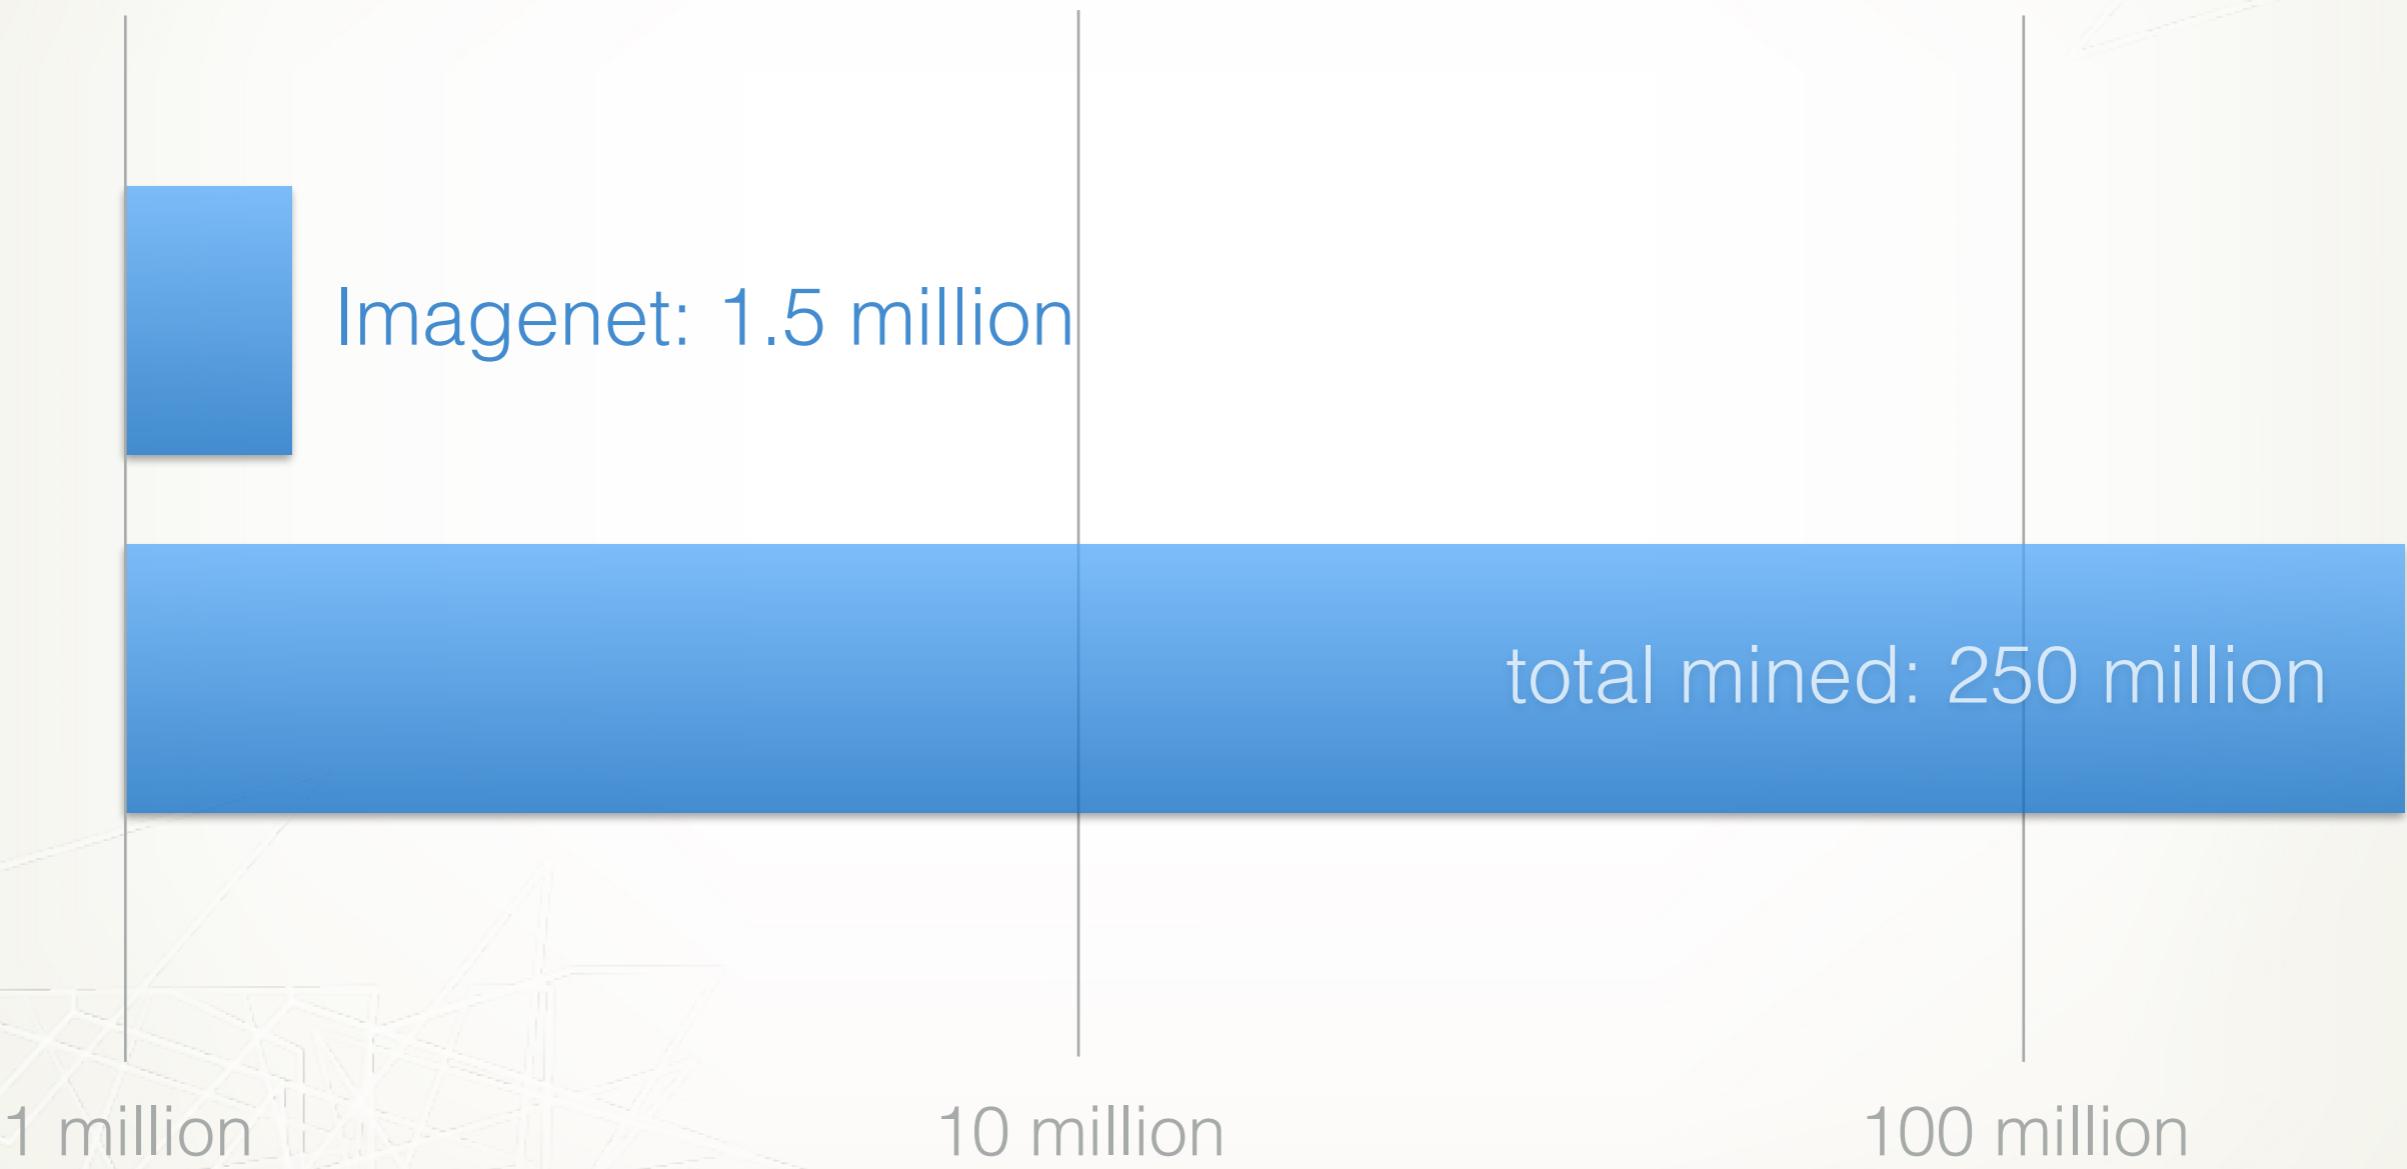


amplitude gain
in Fourier
domain

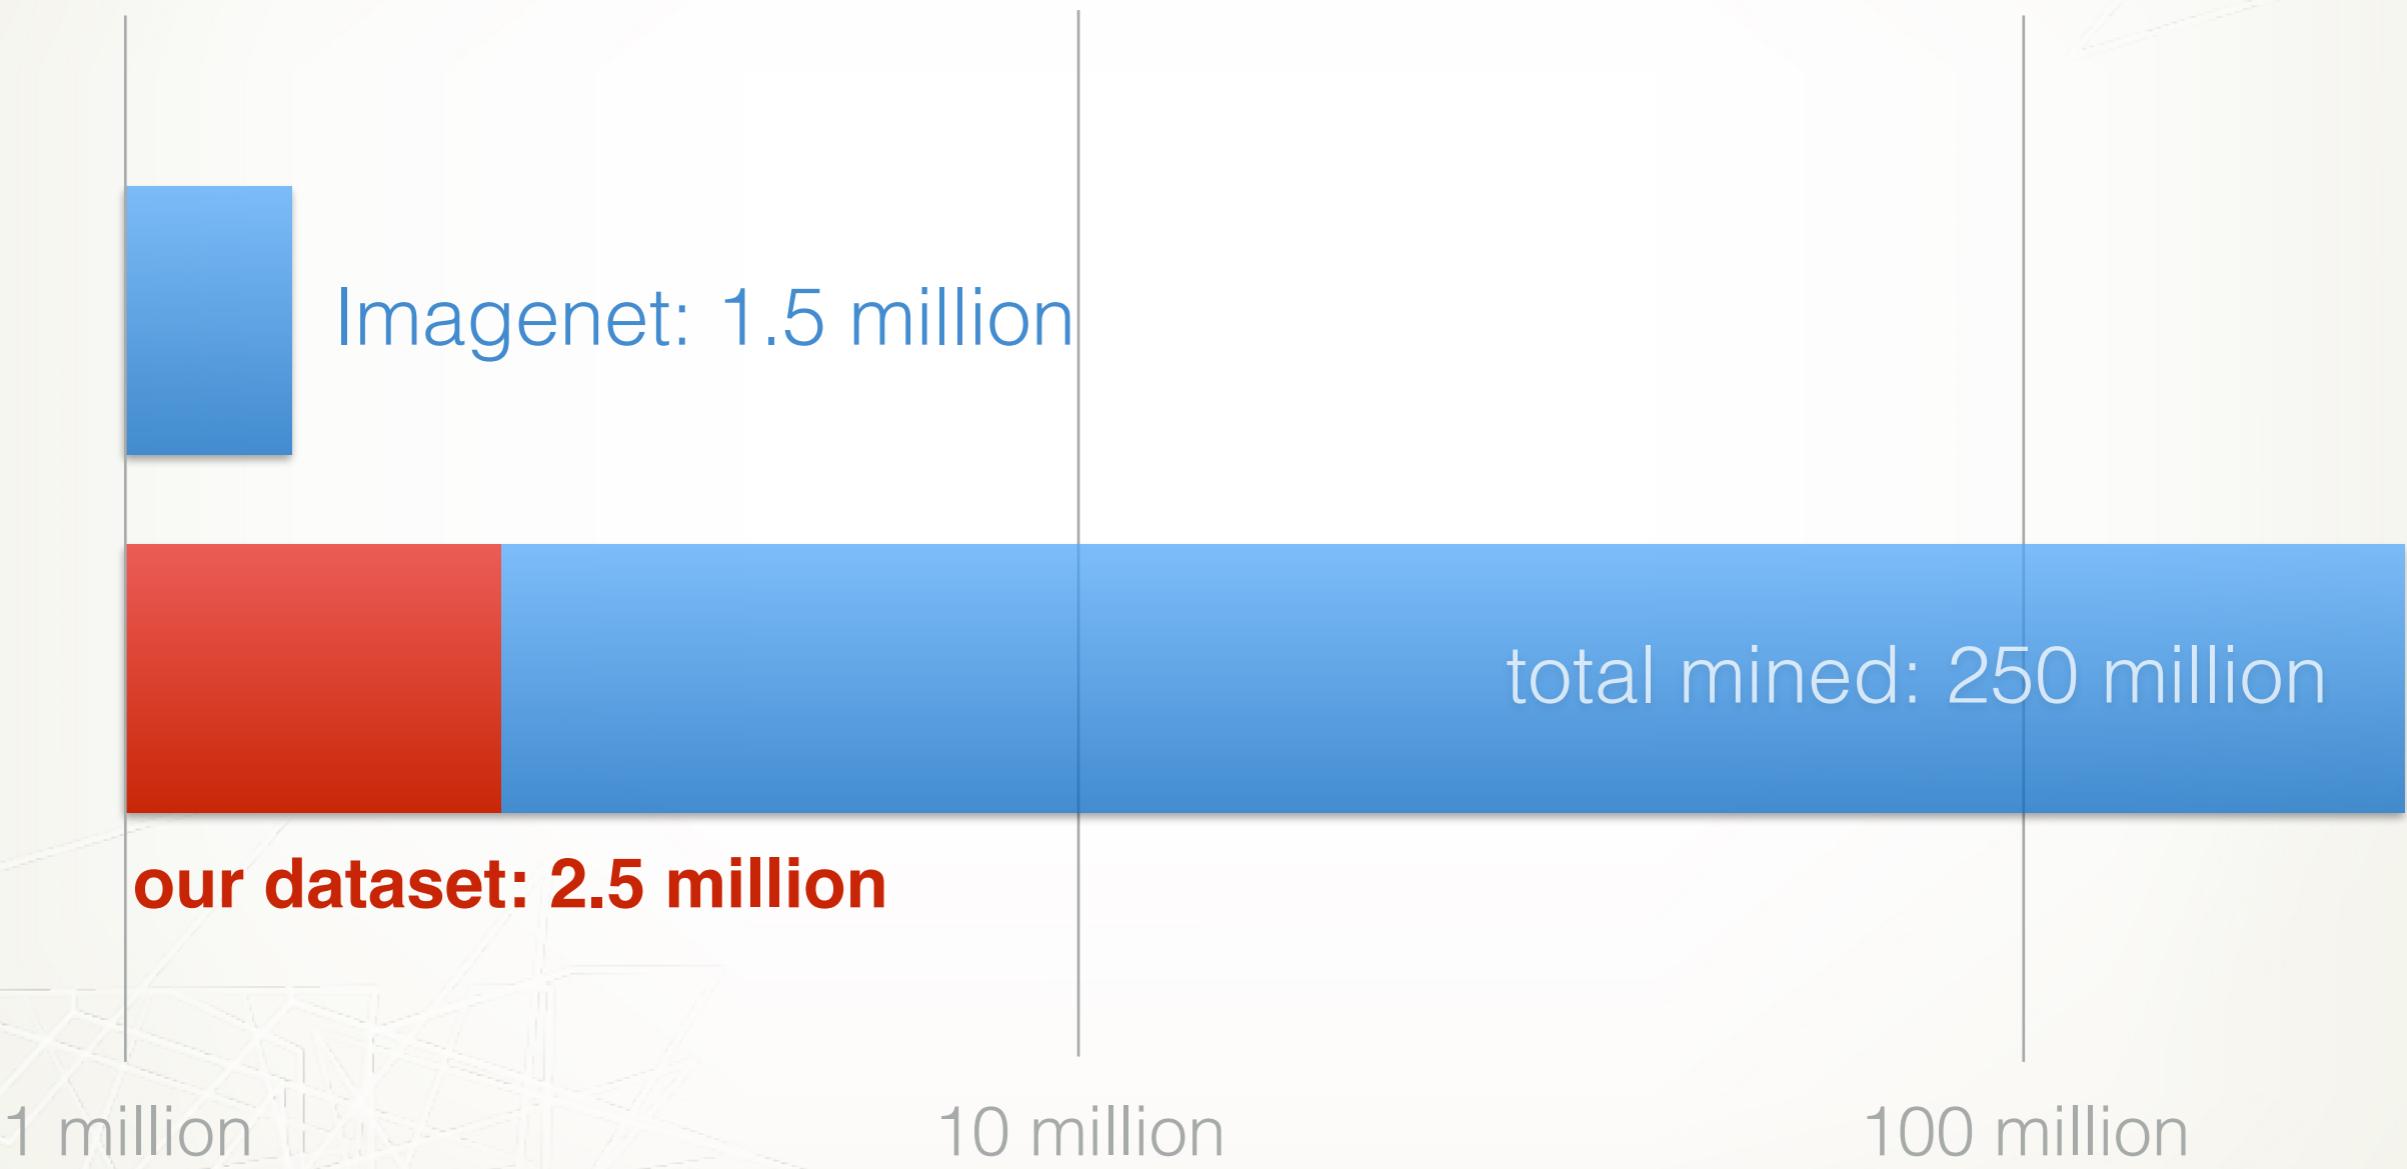
Challenging dataset



Challenging dataset



Challenging dataset





Results

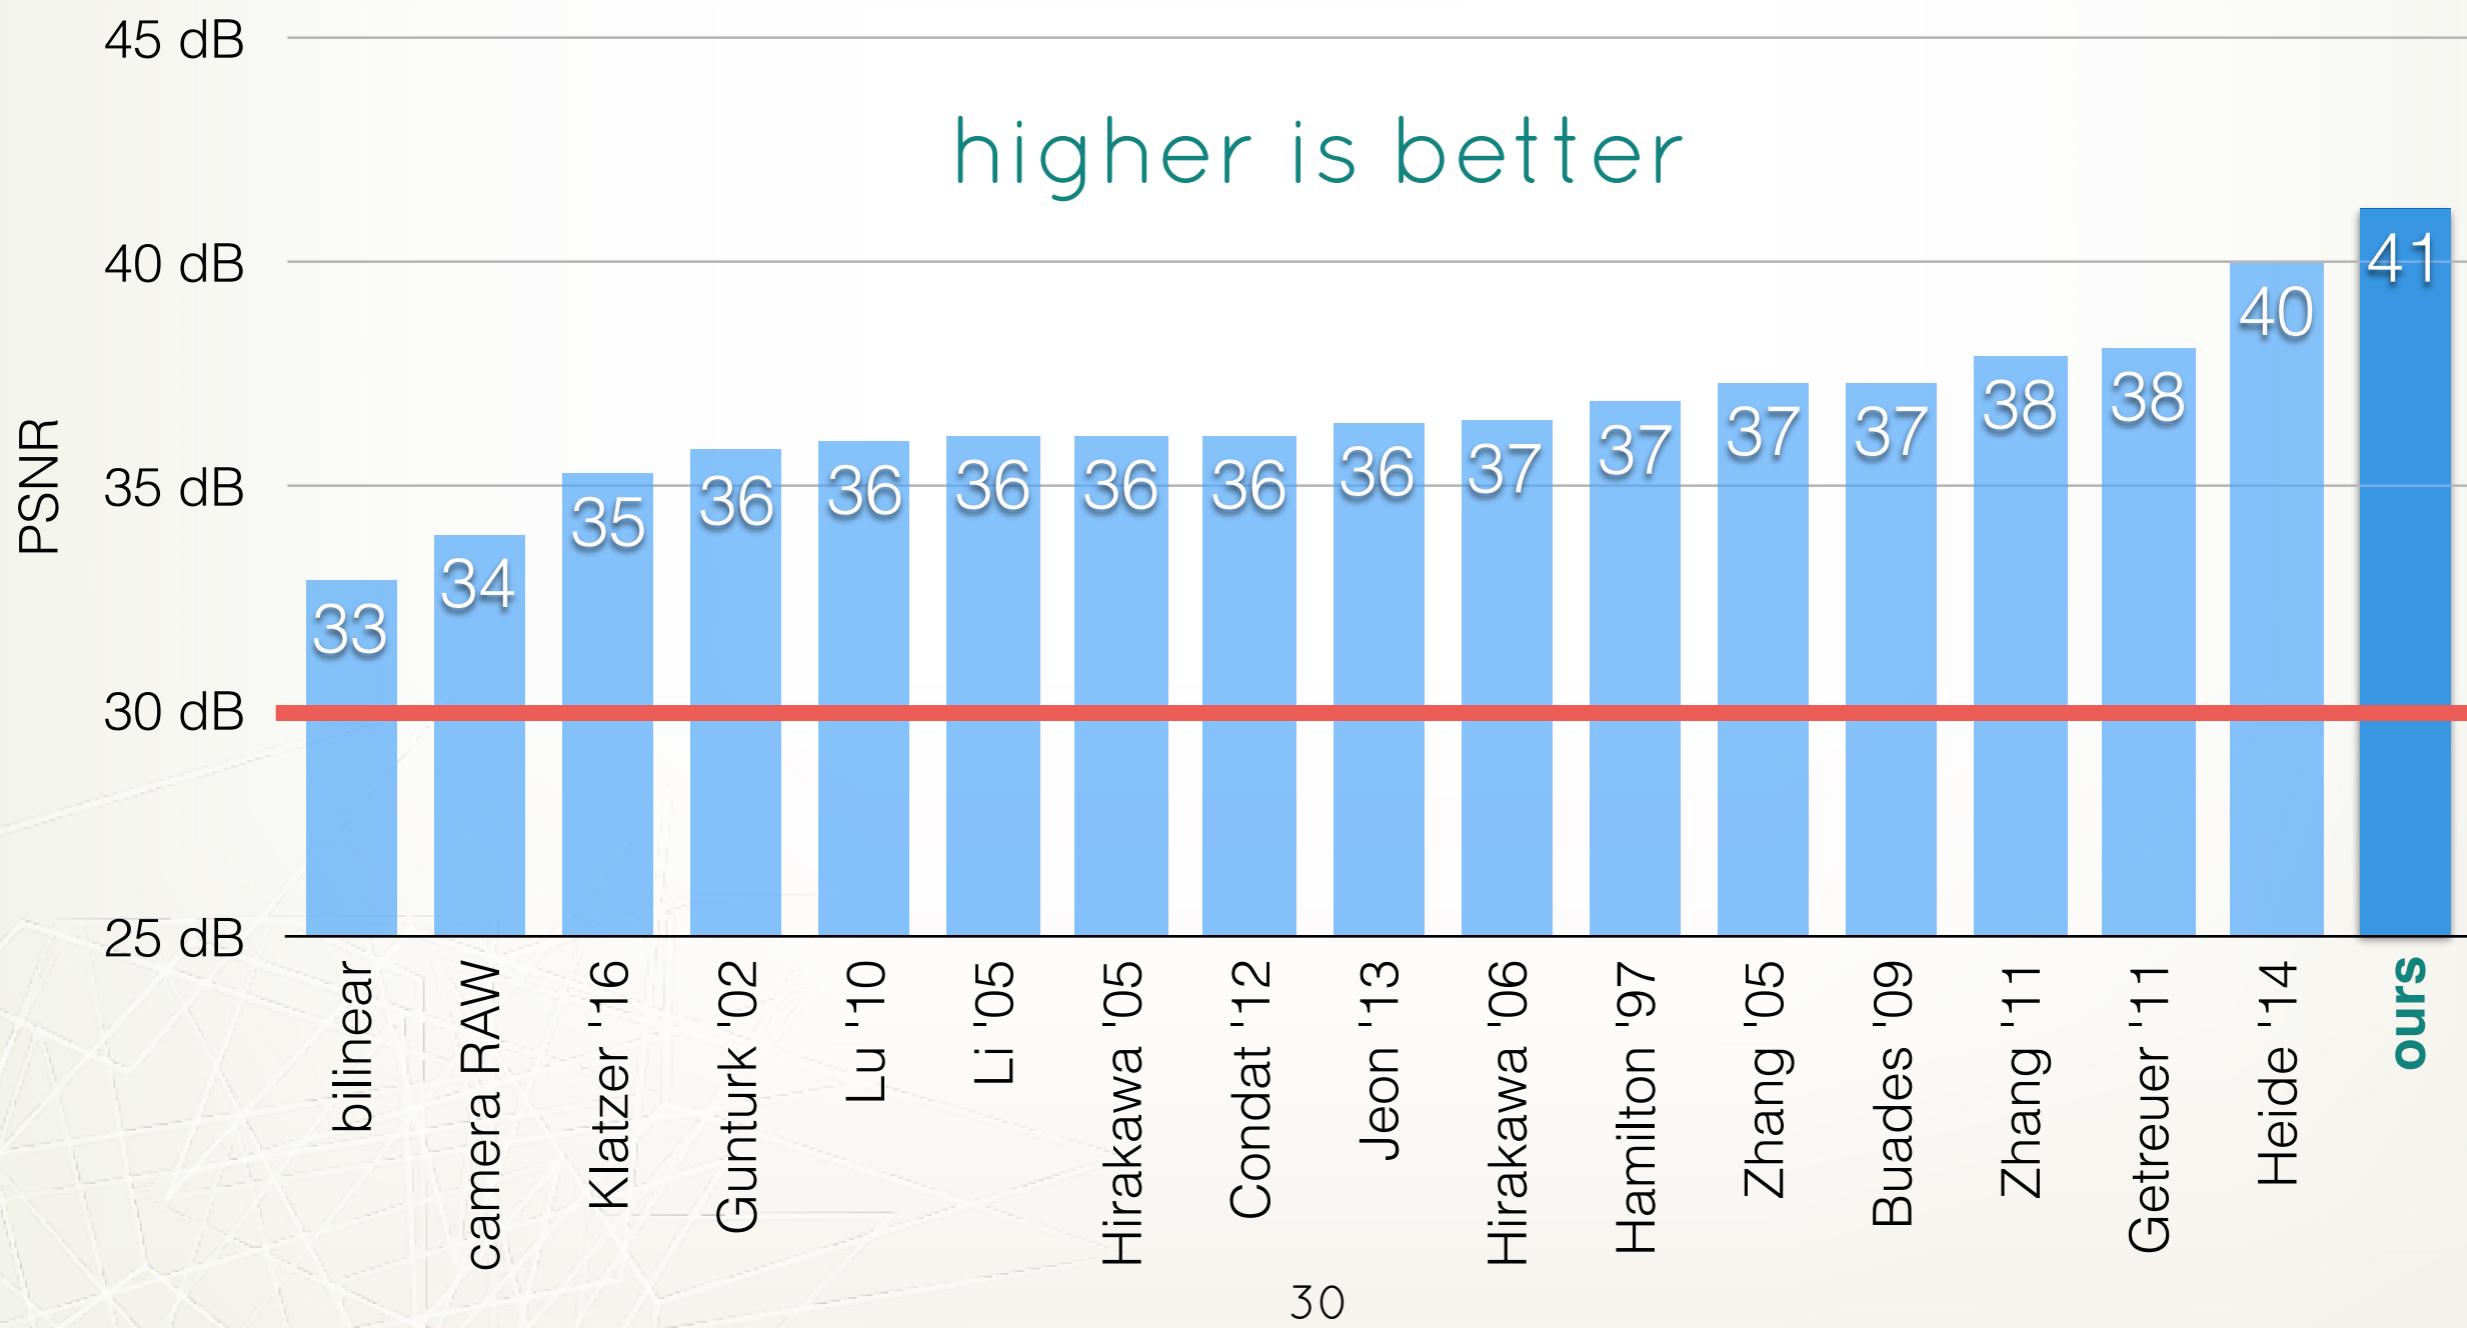
training our model on this new dataset

Evaluation

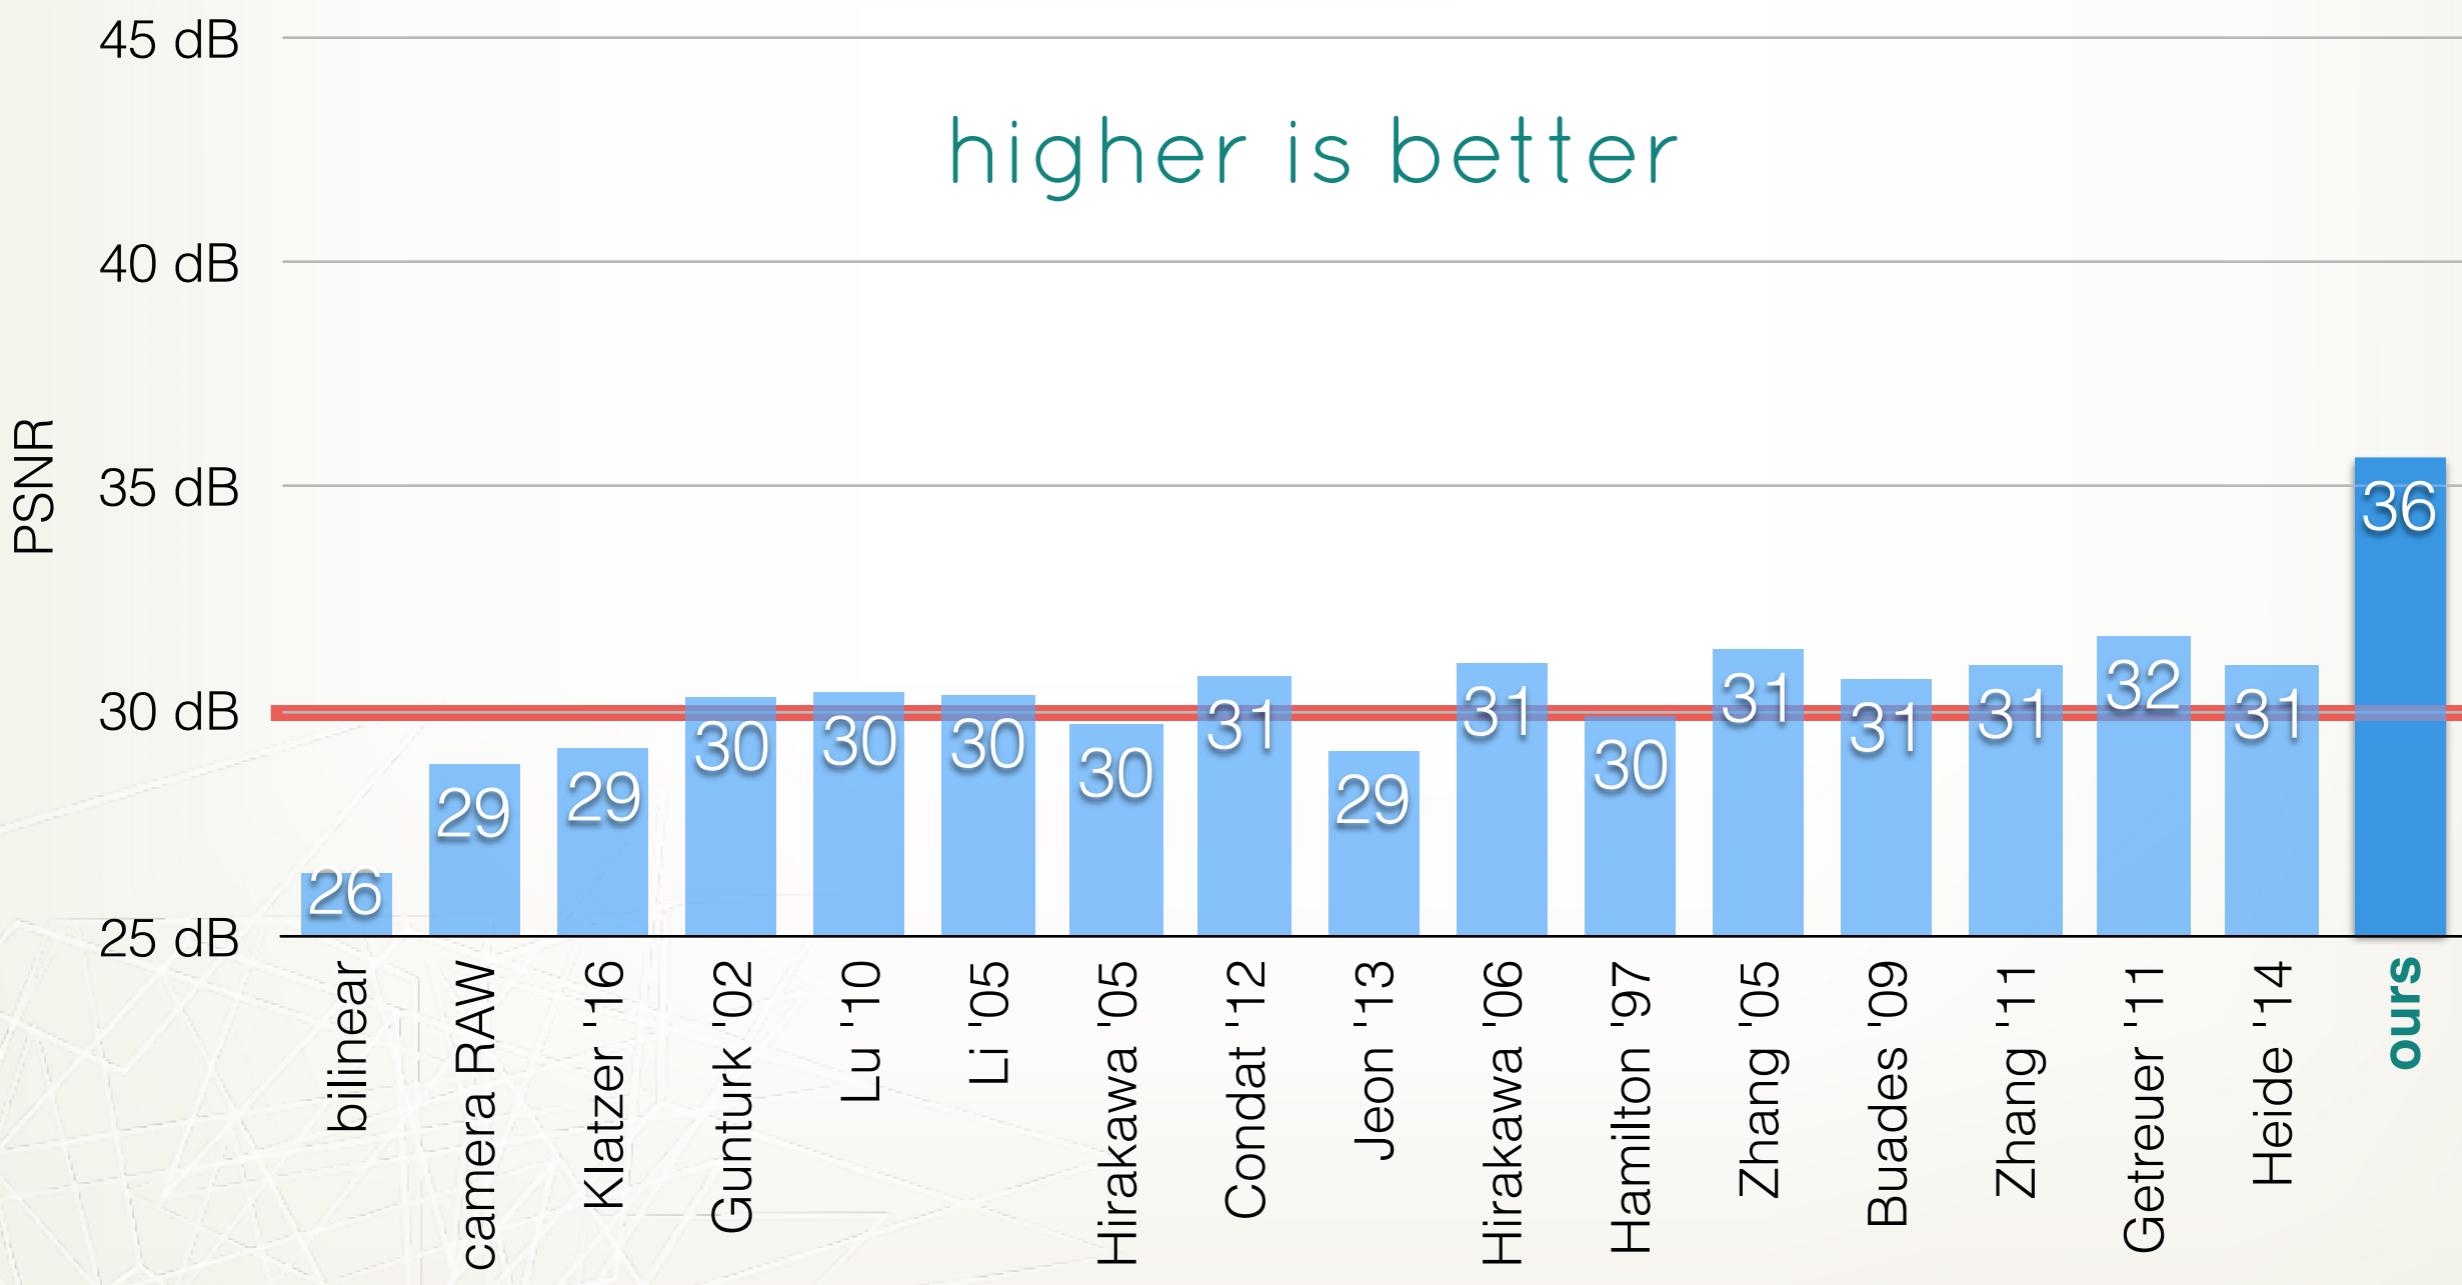
- new, separate test and validation sets
 - 2,000 images each
 - standard datasets are too easy
- **more accurate** than state-of-the-art
- **faster** than best previous work
 - though not real time yet

Demosaicking only

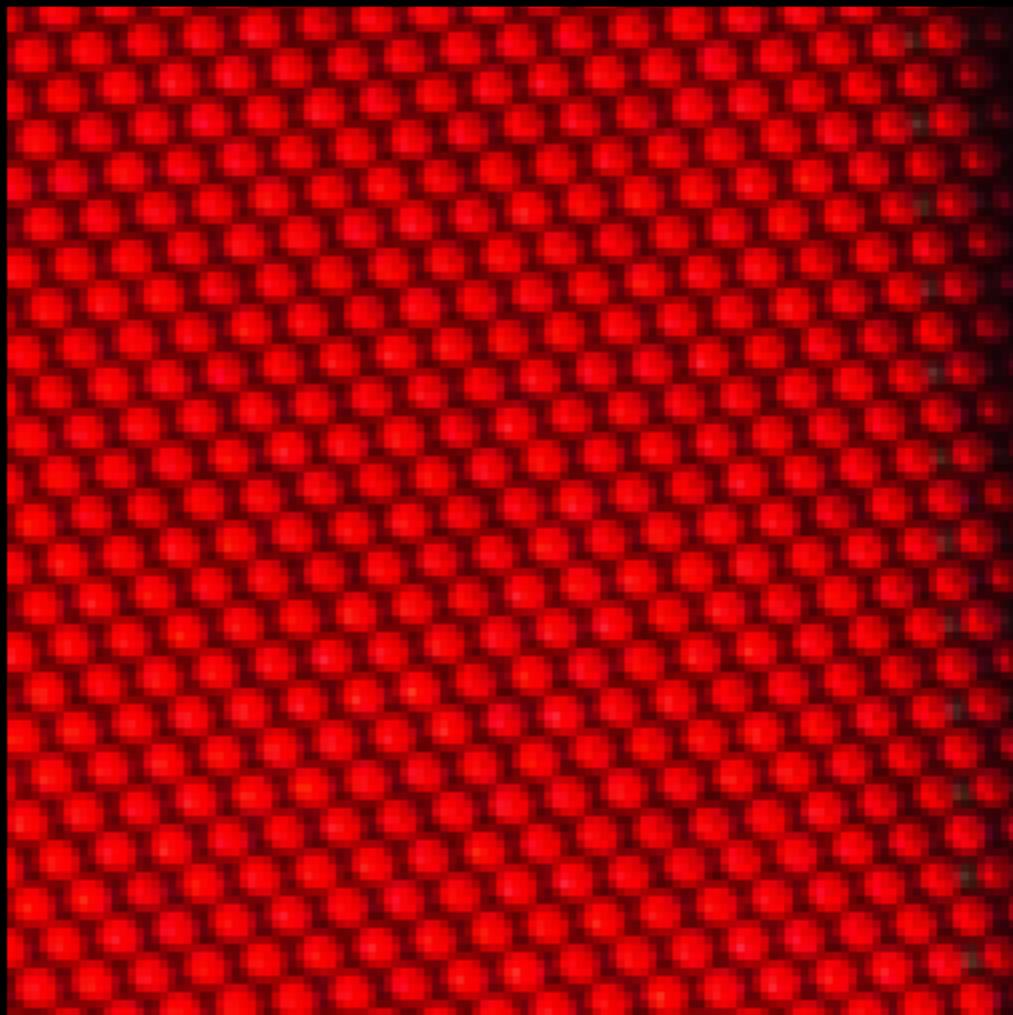
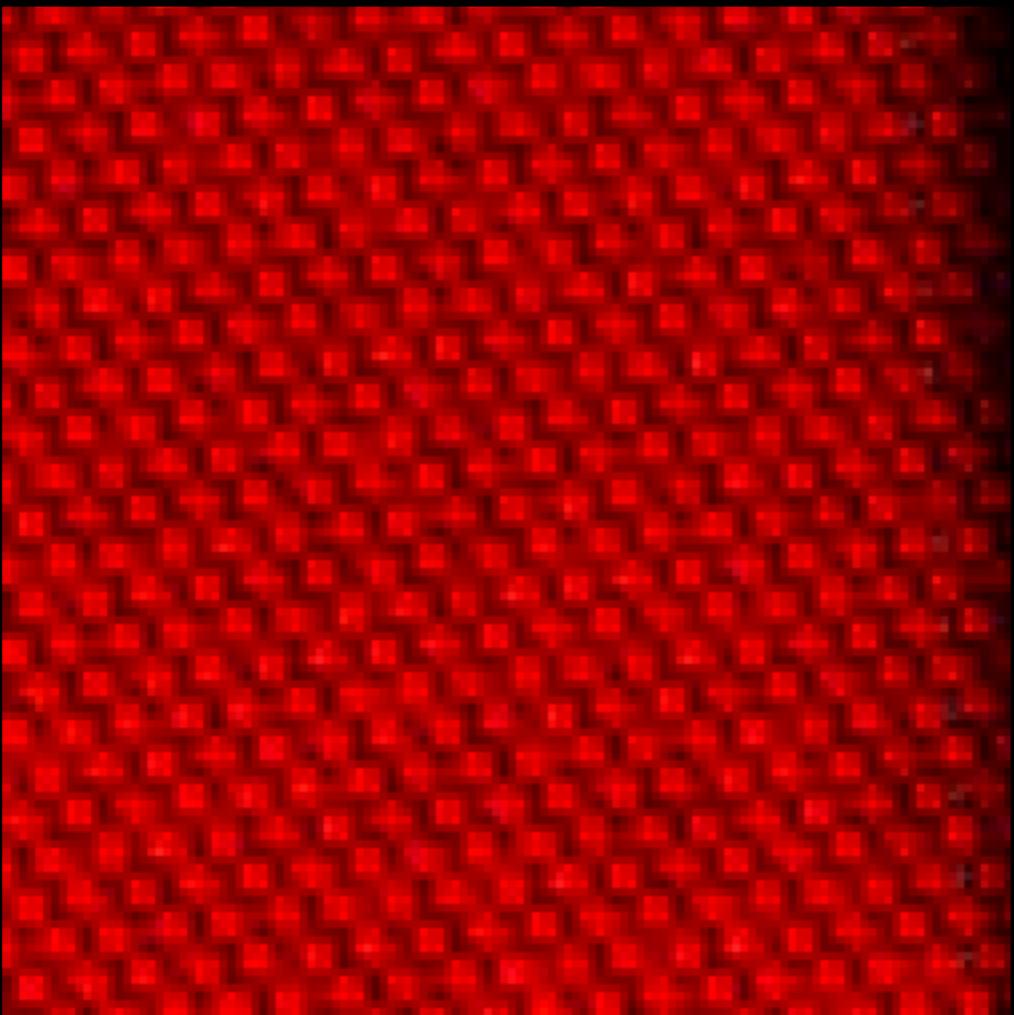
Kodak and McMaster datasets (noise free)



Demosaicking only our test dataset

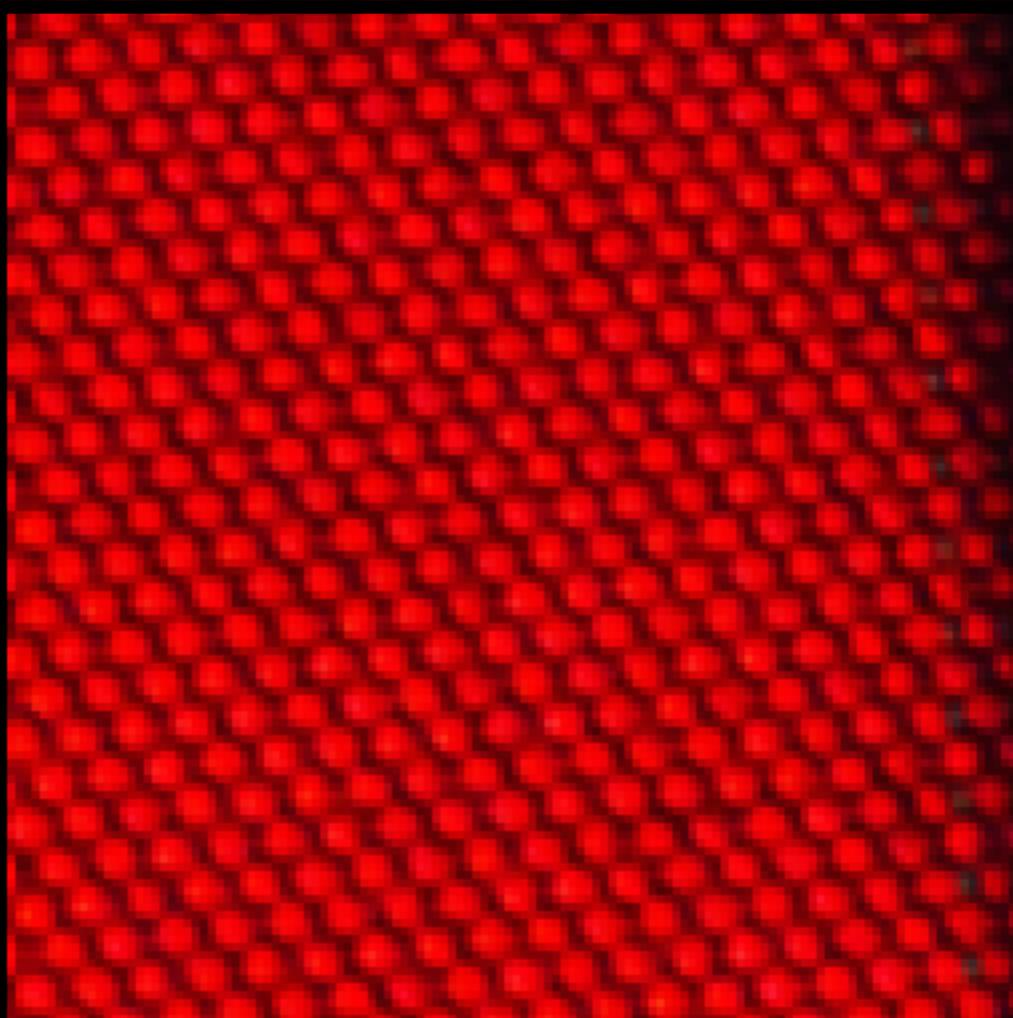
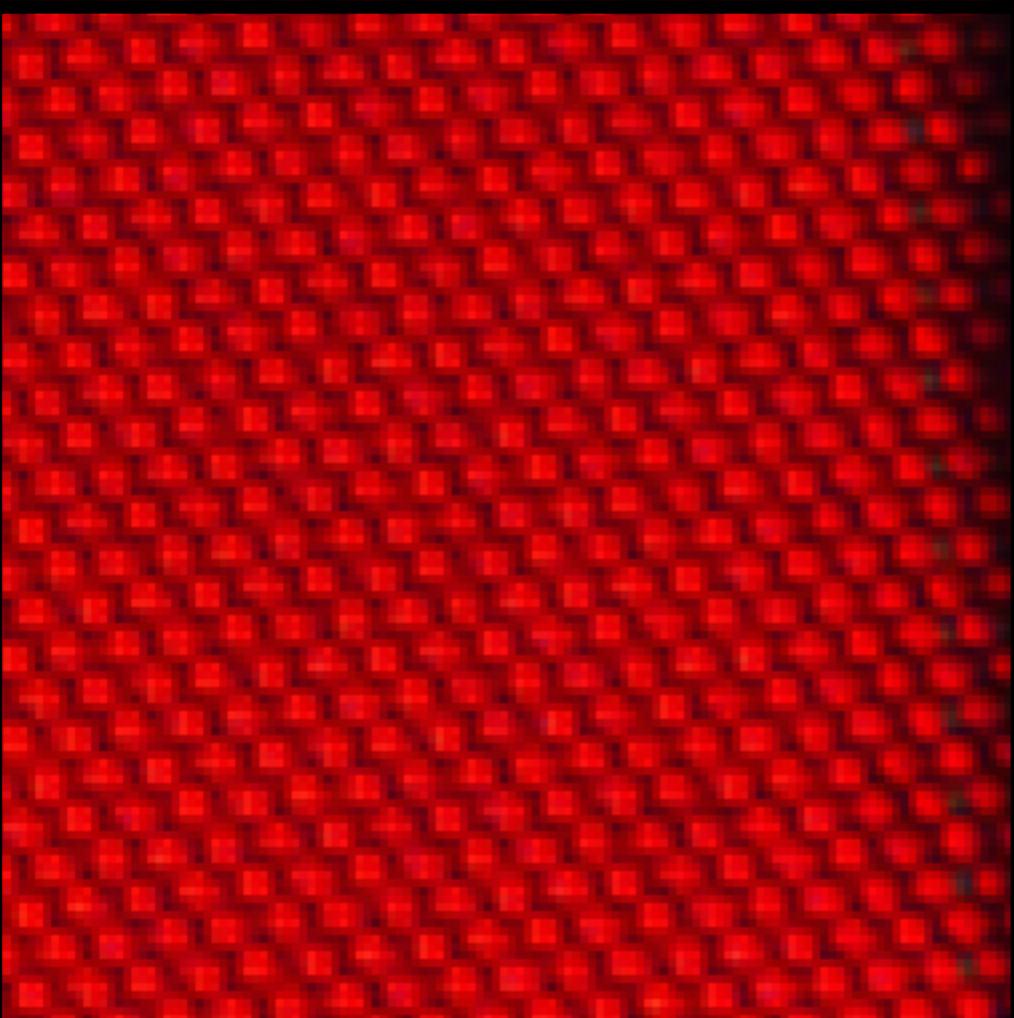


Photoshop
25 dB



ref

FlexISP
27 dB



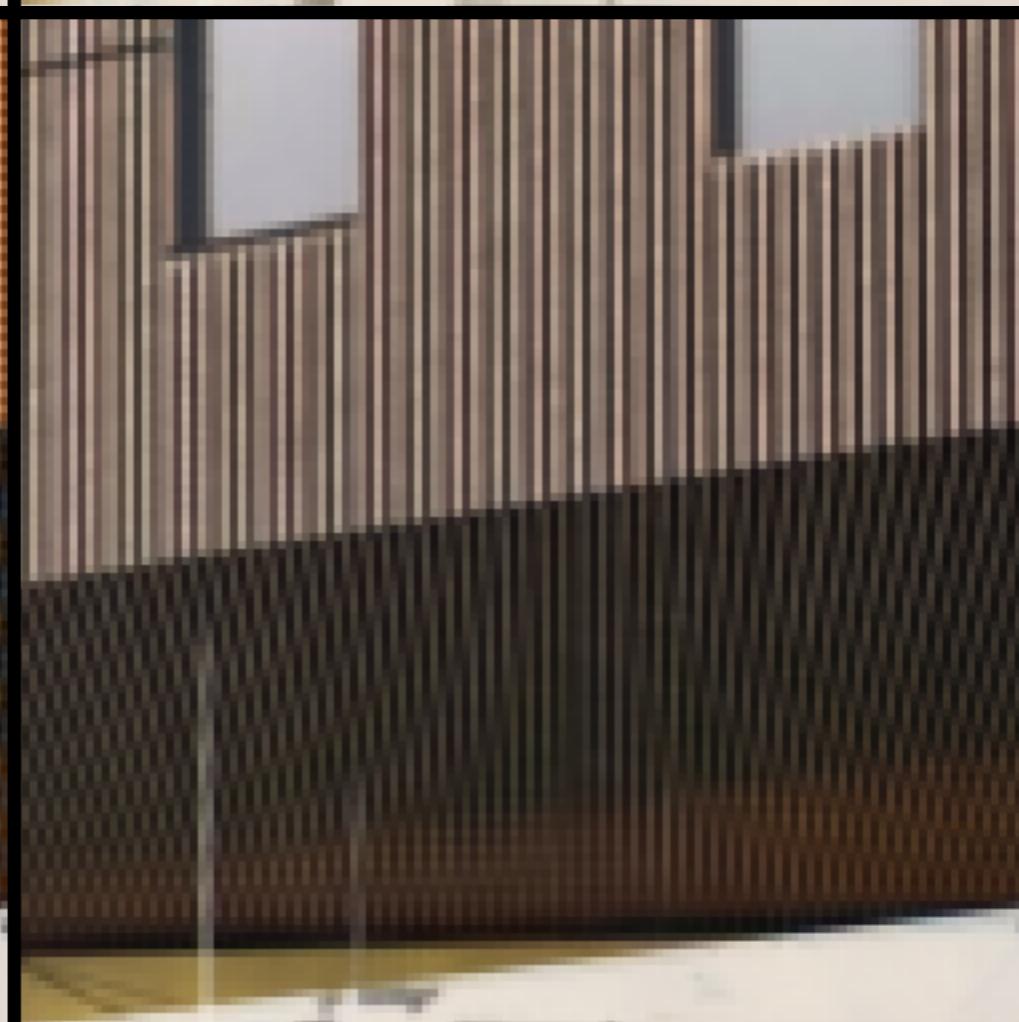
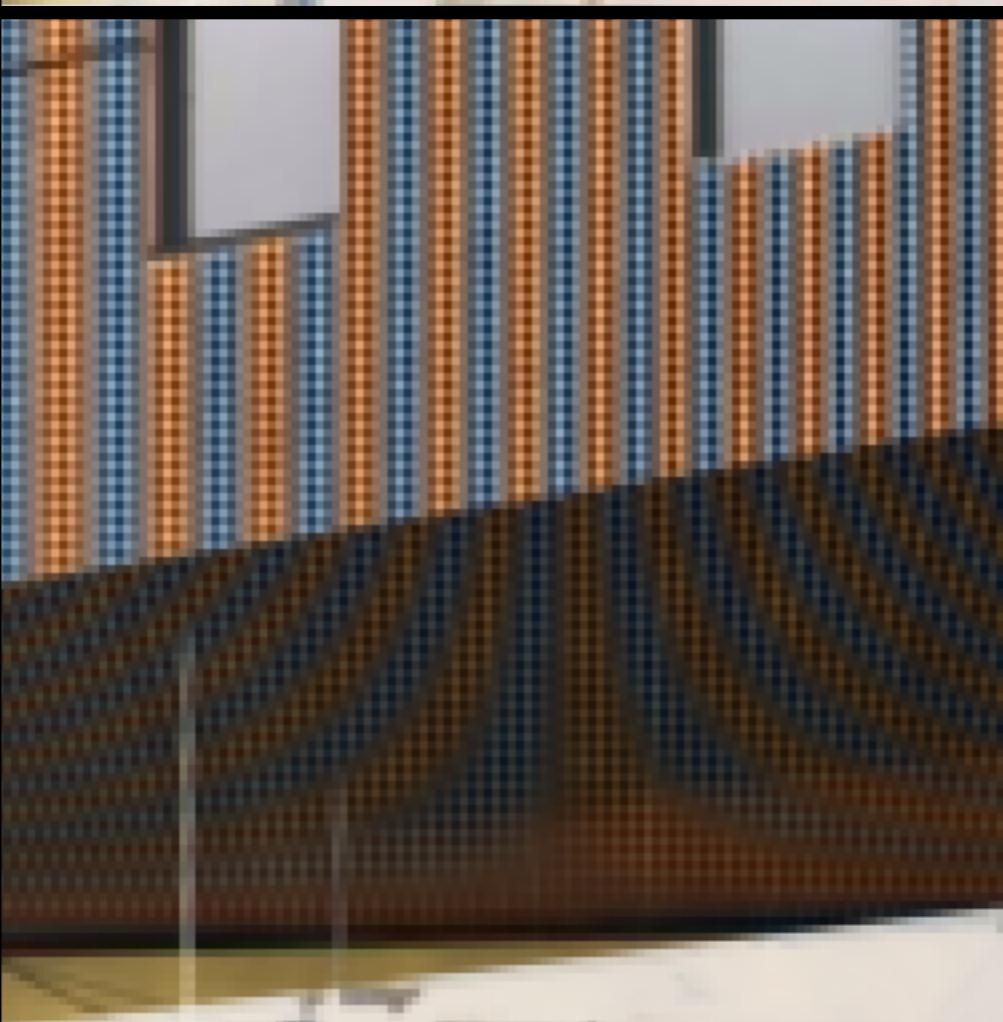
ours
31 dB

Photoshop
30 dB



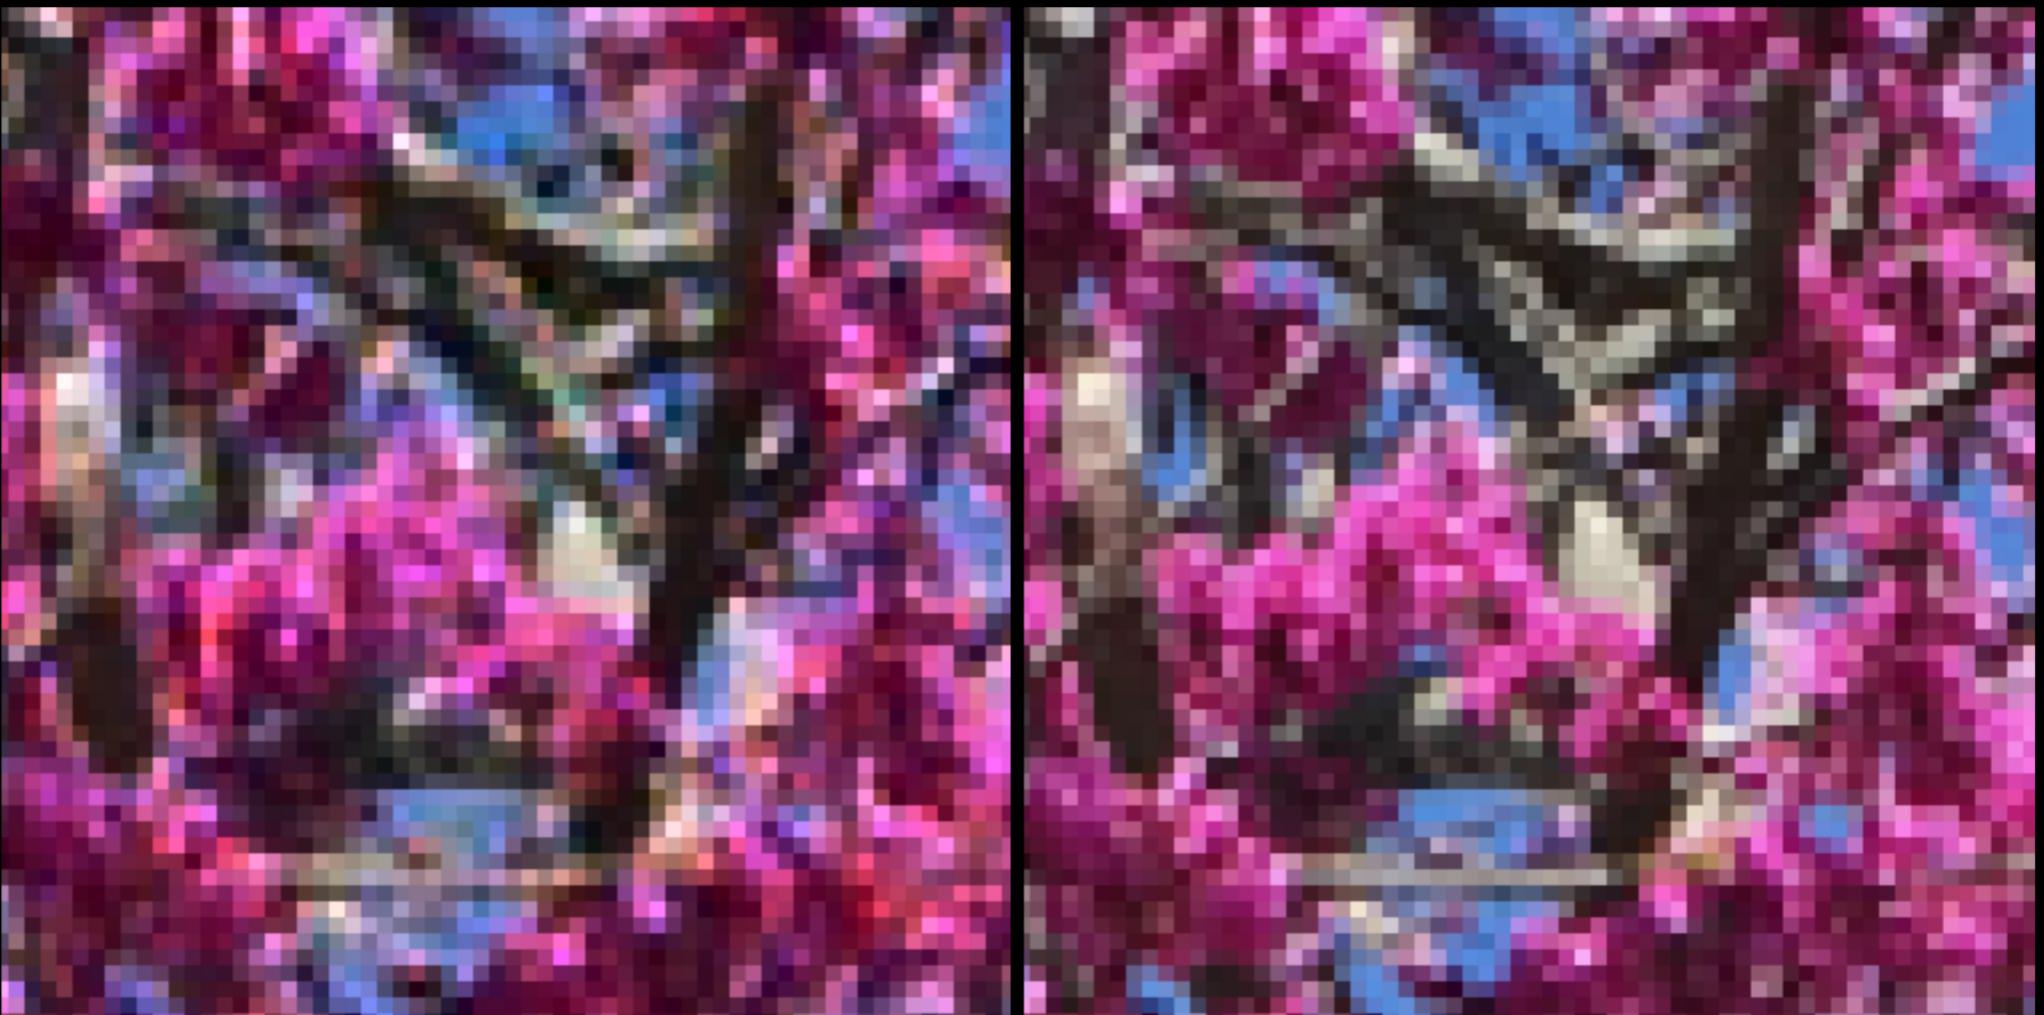
ref

FlexISP
22 dB

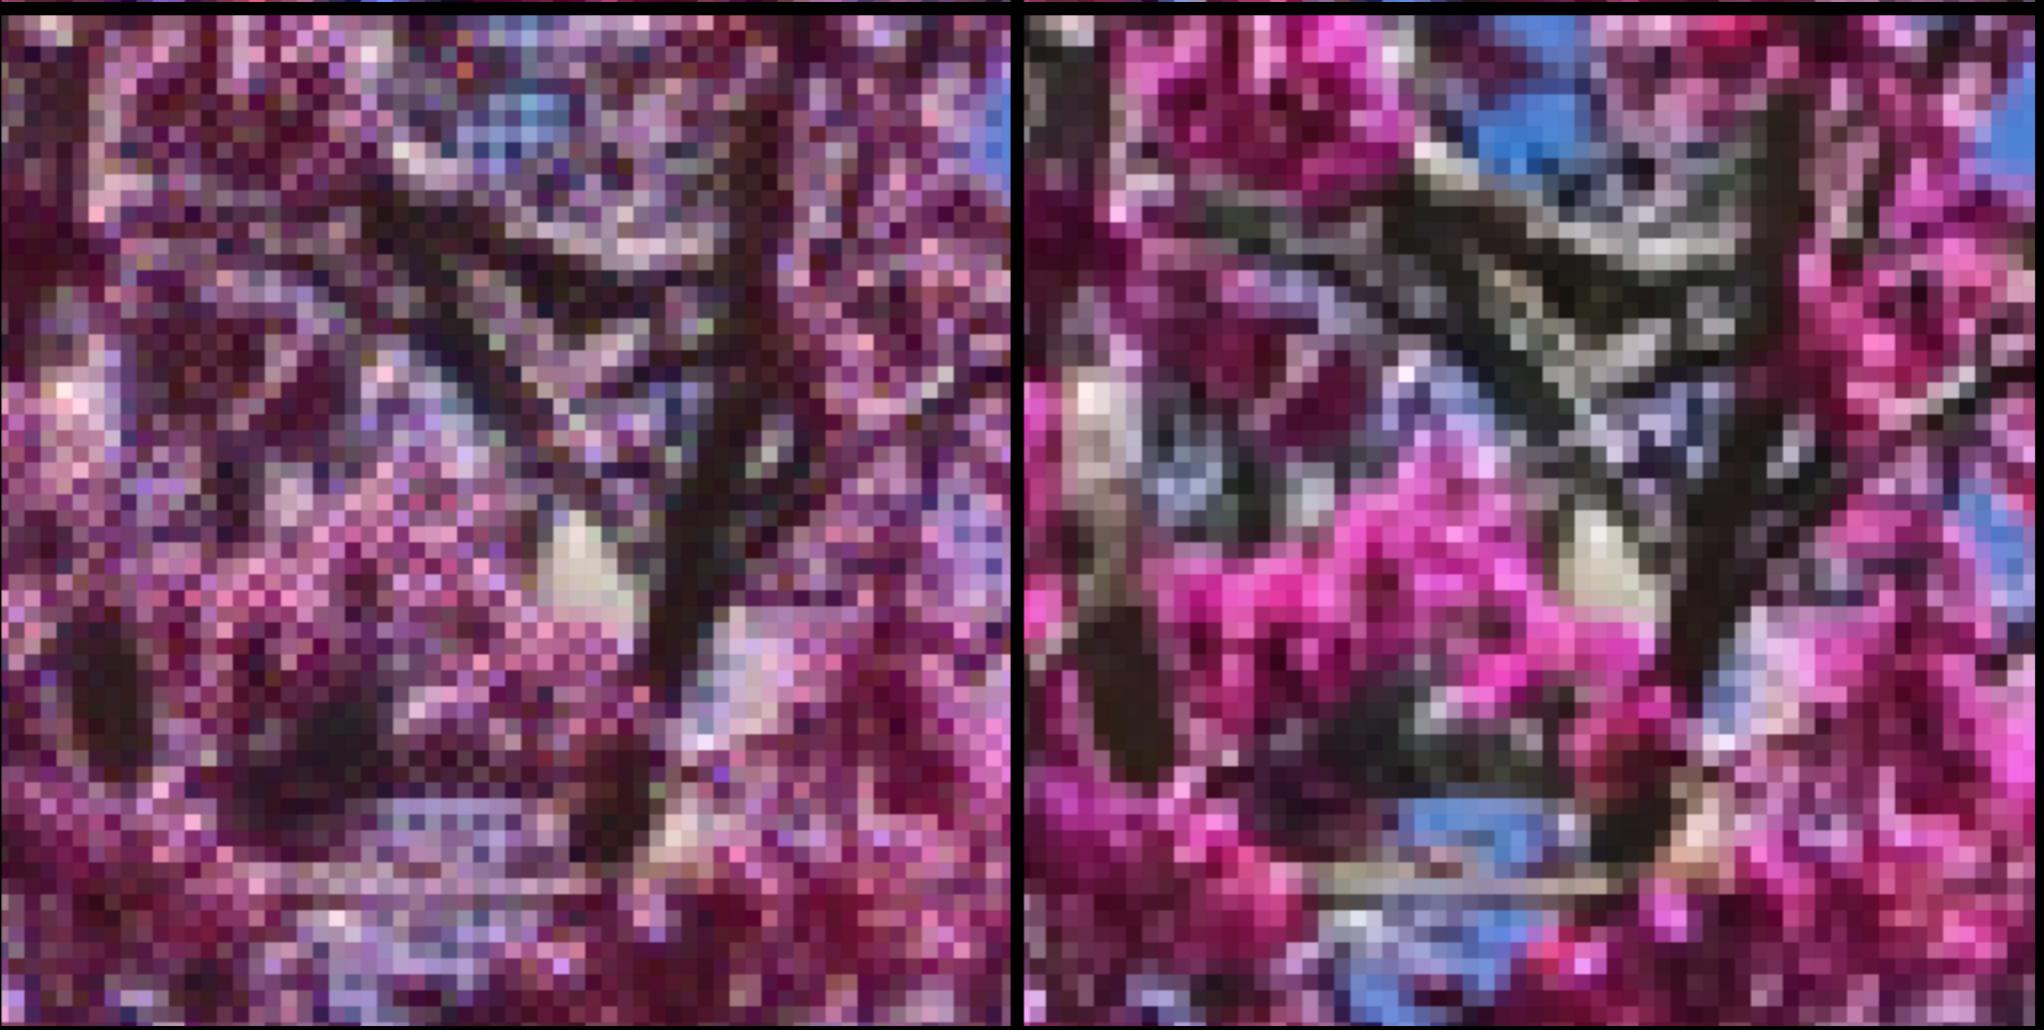


ours
40 dB

Photoshop
24 dB



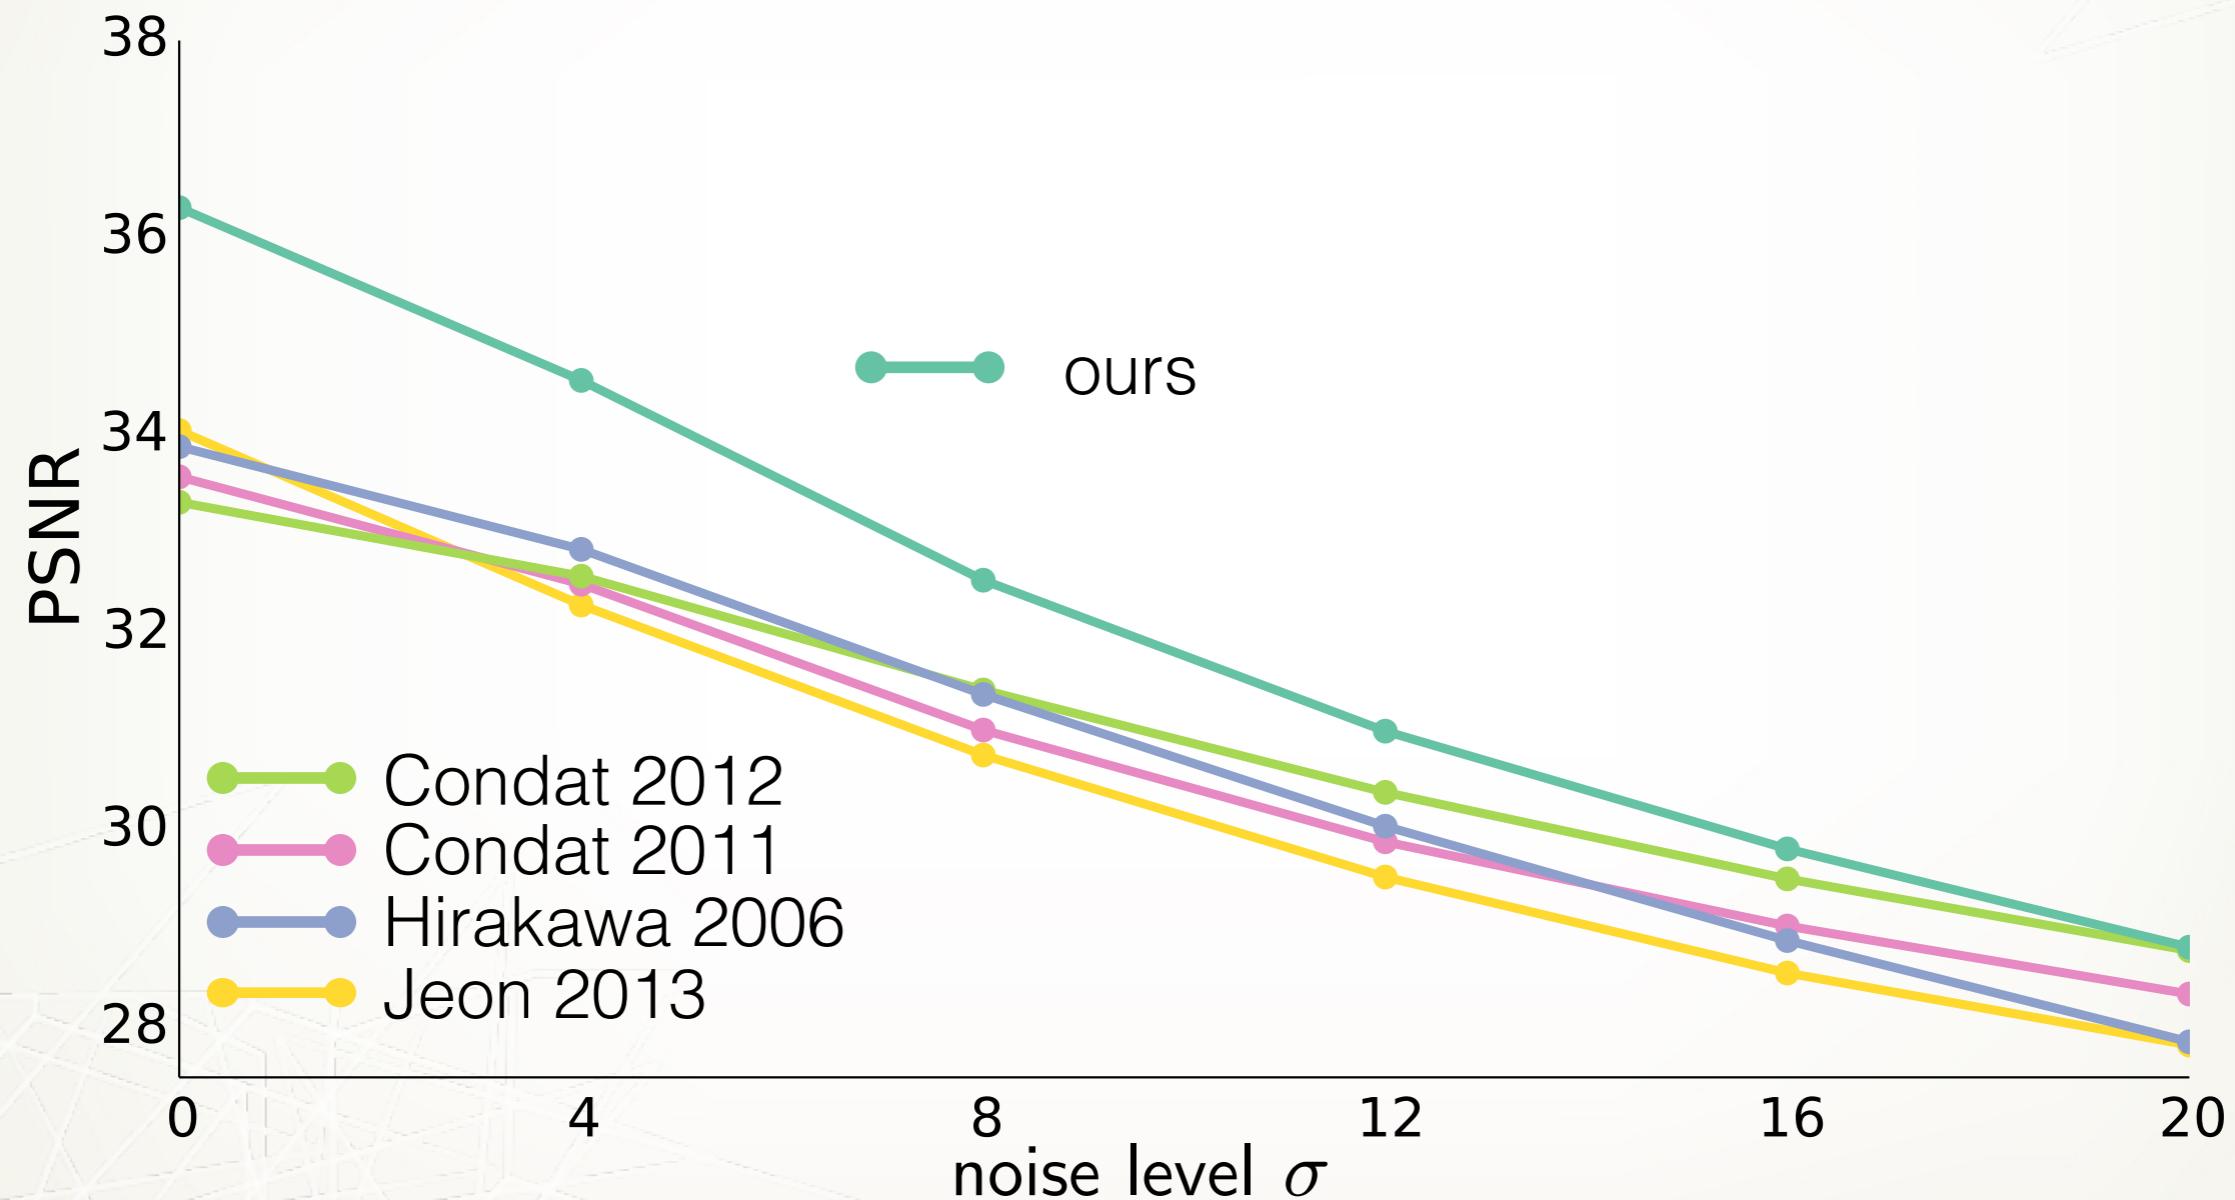
FlexISP
22 dB



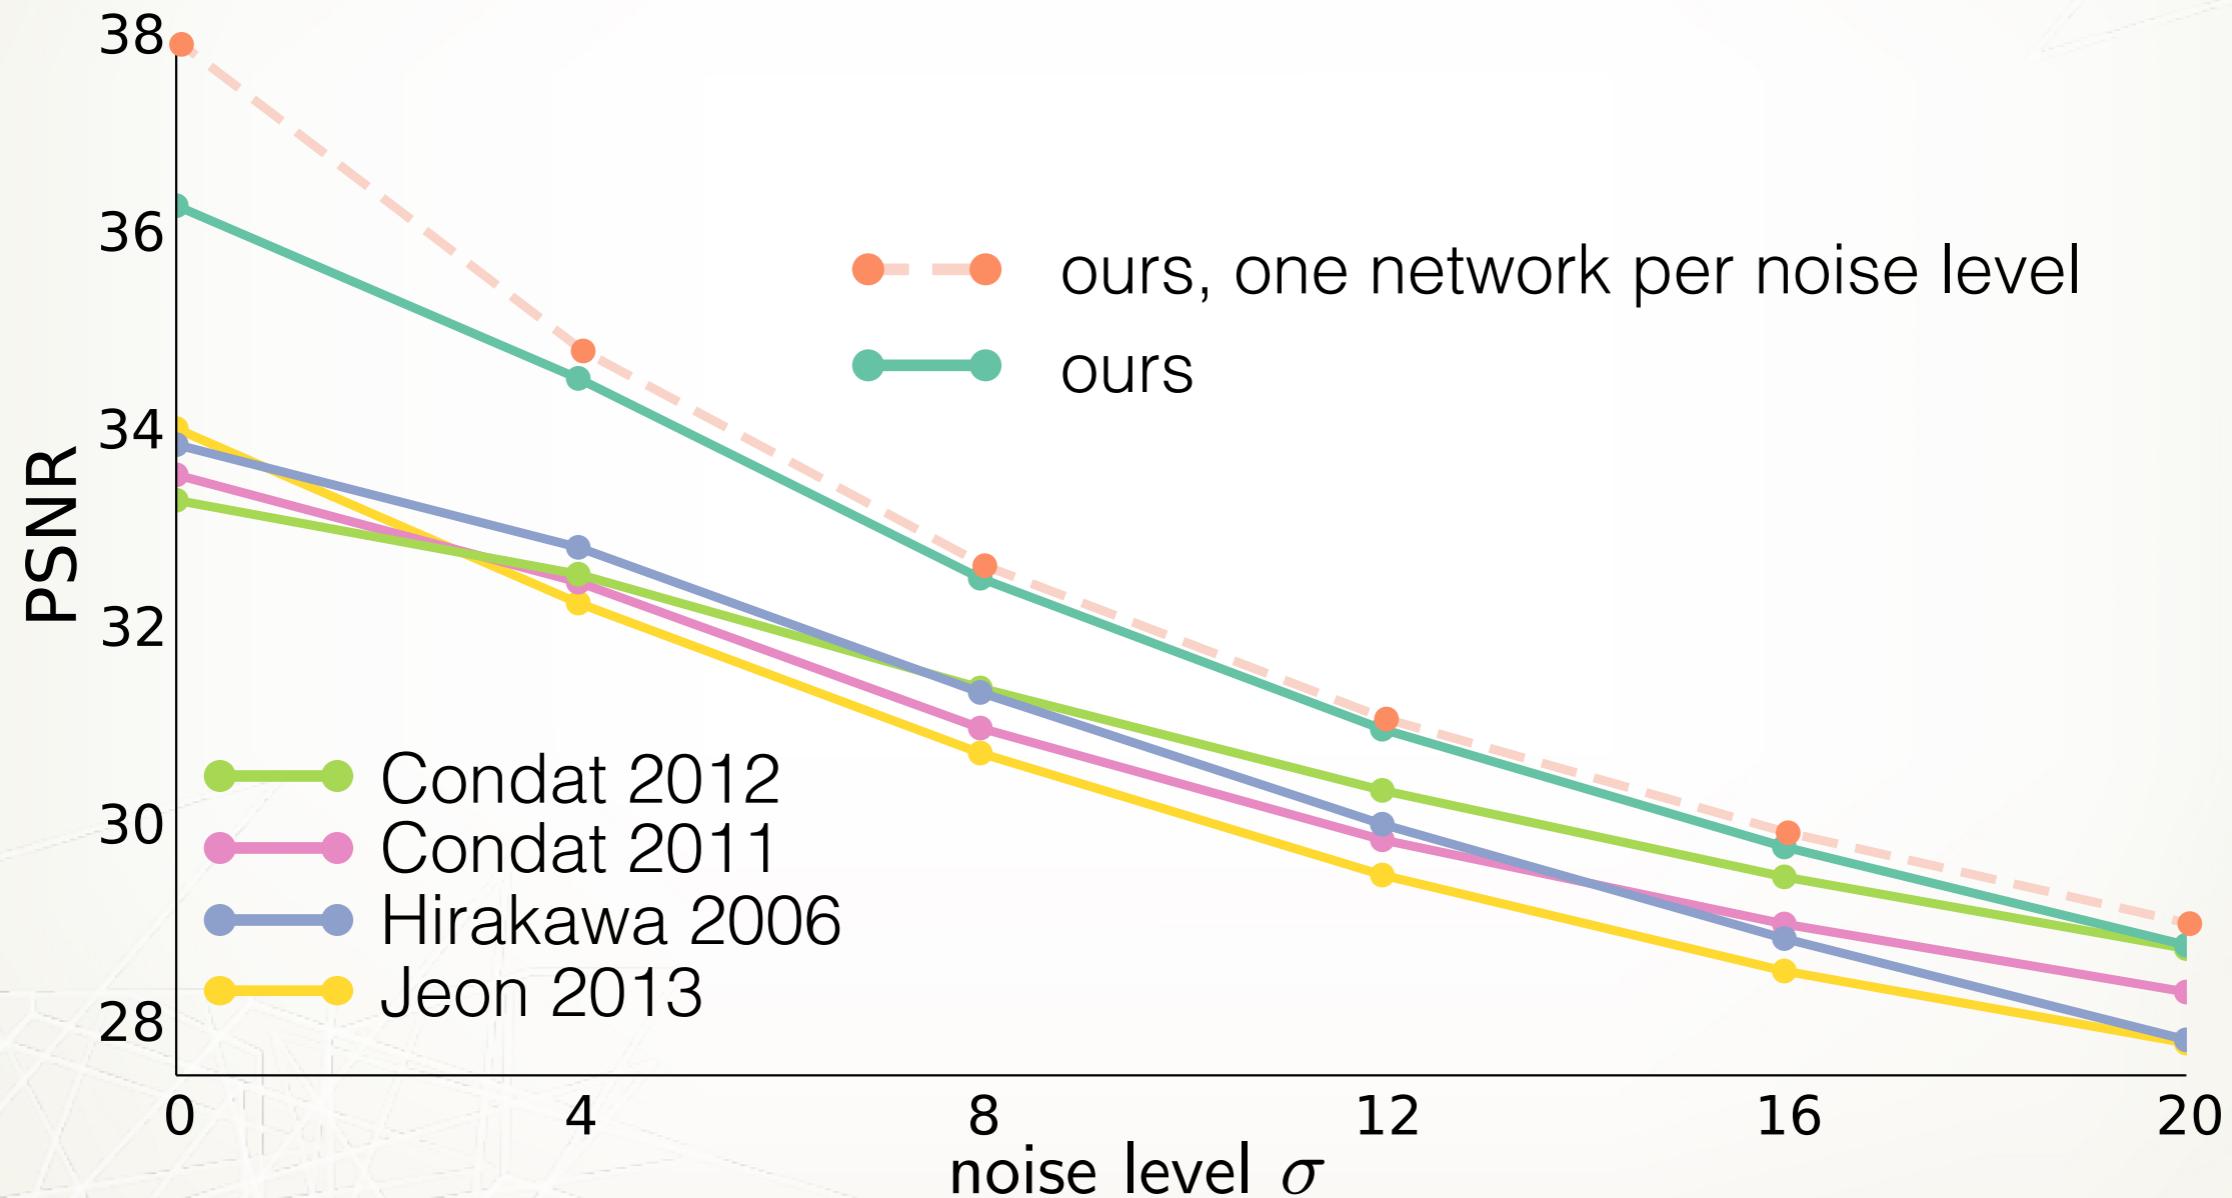
ref

ours
27 dB

Joint demosaicking/denoising



Joint demosaicking/denoising



noisy
 $\sigma=4$



Condat 2012
27 dB



ref

ours
31 dB

noisy
 $\sigma=12$



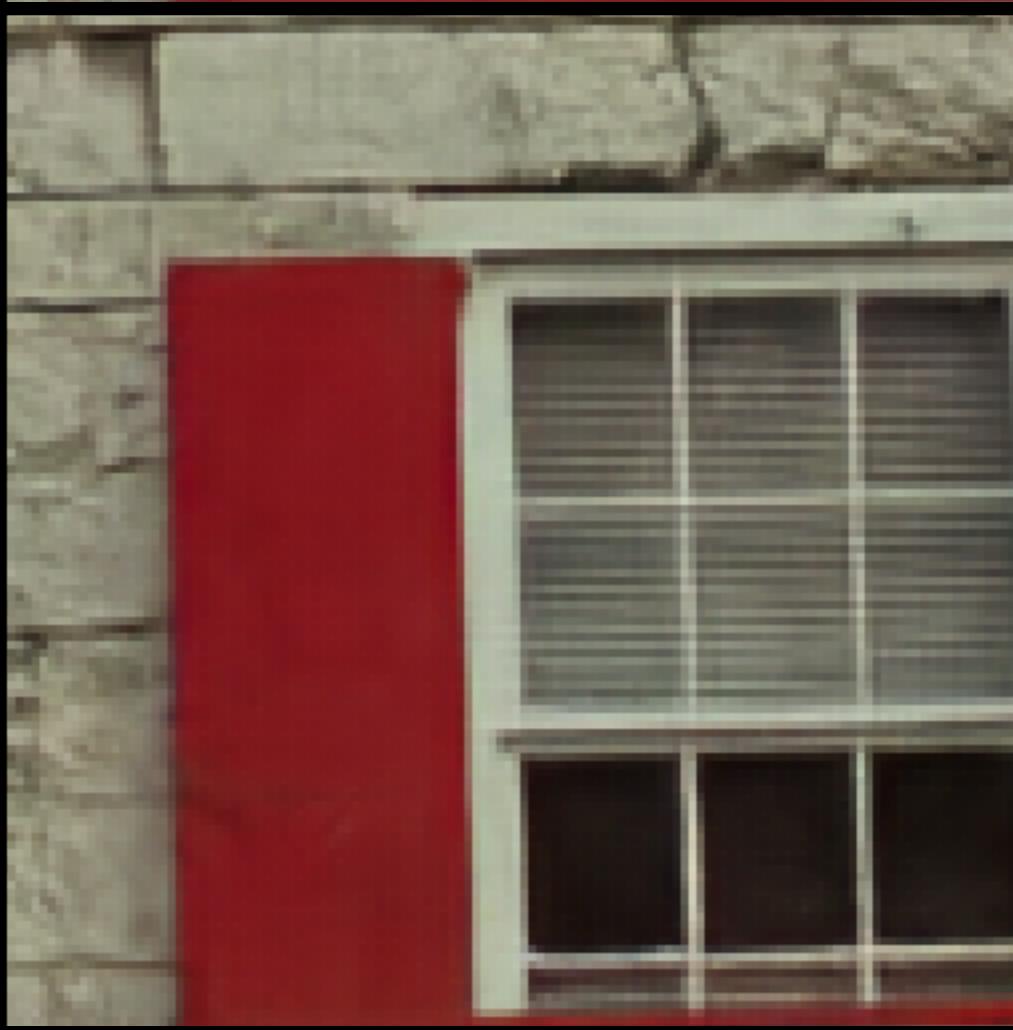
ref



Condat 2012
27 dB



ours
31 dB

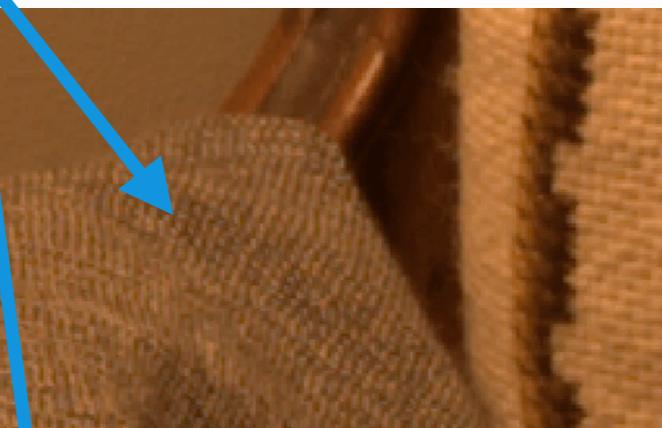


Real RAW data

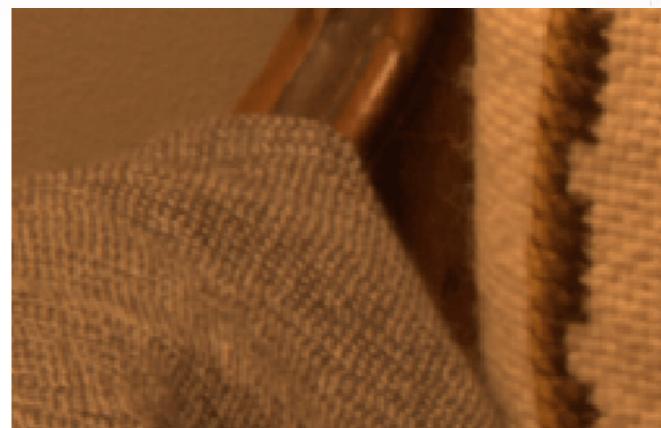
false colors



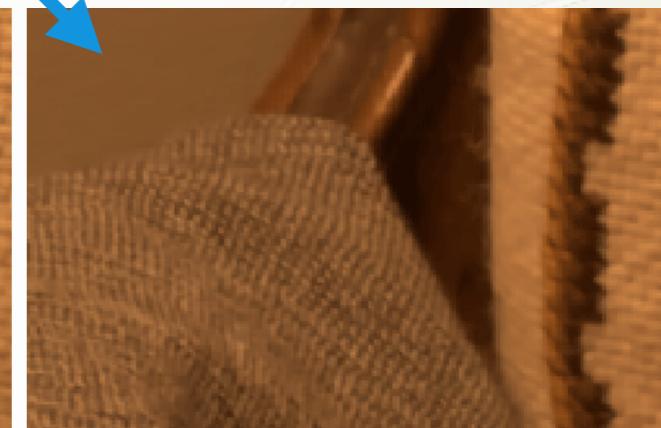
dcraw



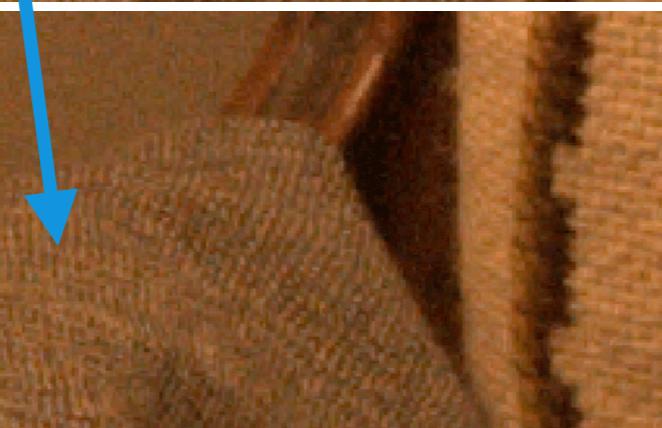
over-smoothing
ours



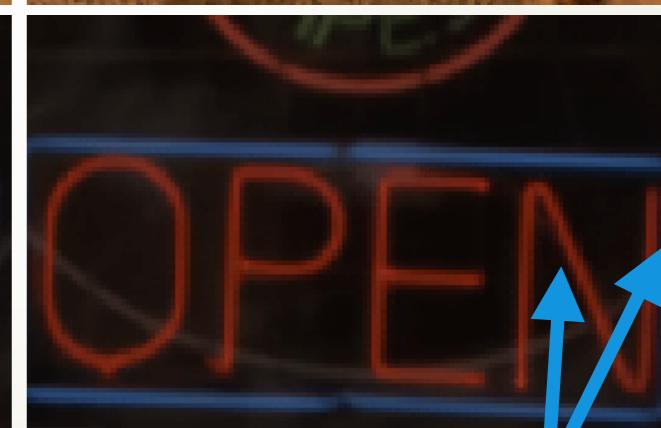
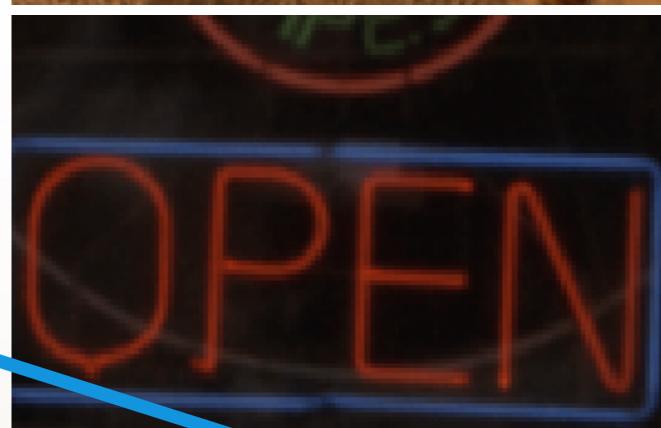
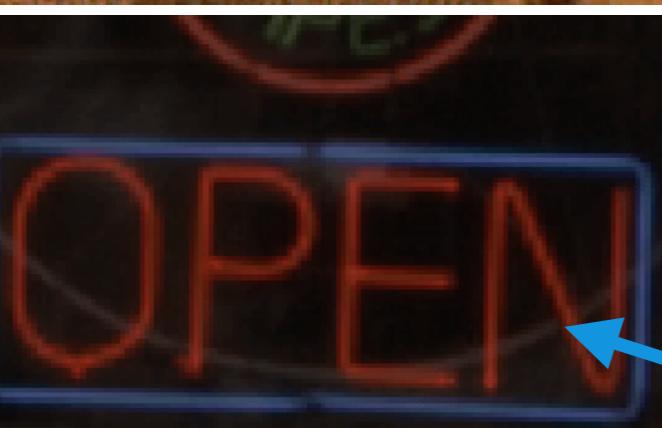
Klatzer 2016



ISO6400



ISO100



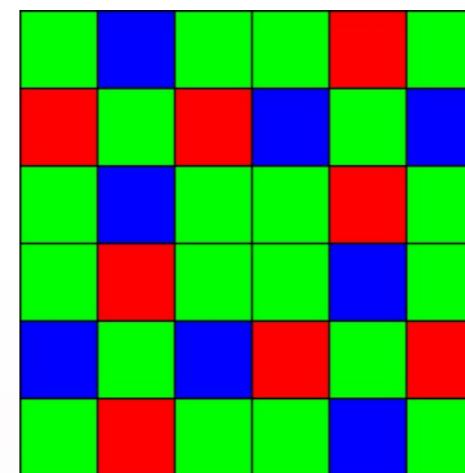
captured with a Canon 5D mark II

edge artifacts



Non-Bayer mosaic

RAW Fuji X-Trans

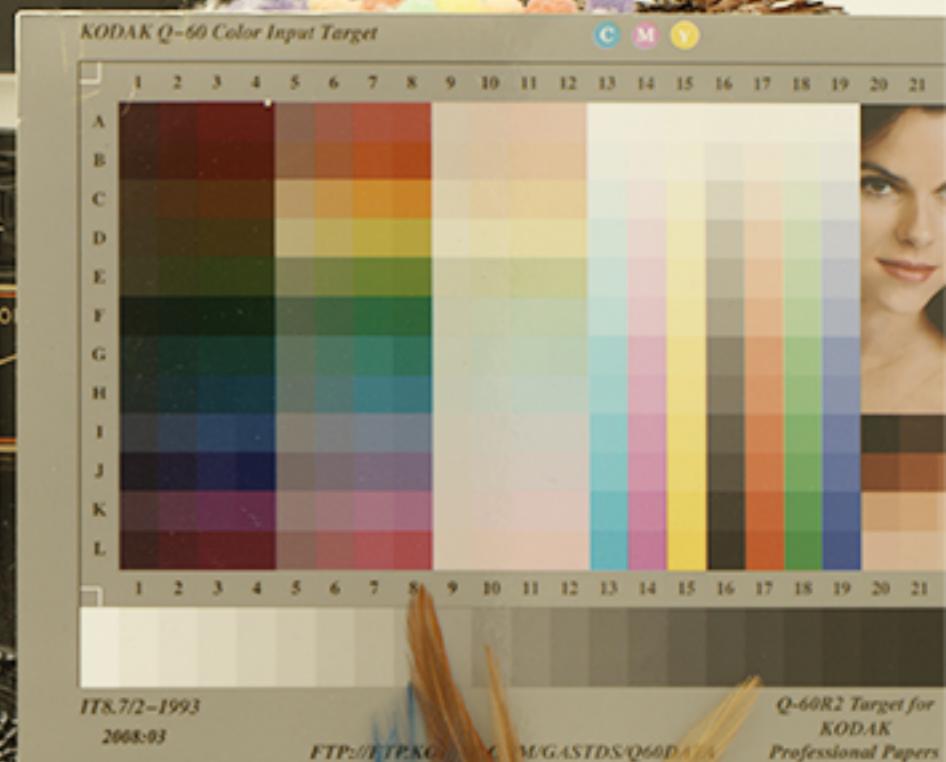


KODAK Gray Scale

A row of three colored circles representing the primary colors used in printing: Cyan (C), Magenta (M), and Yellow (Y).

A small Kodak logo is located in the bottom right corner of the slide.

A 1 2 3 4 5 M 8 9 10 11 12 13 14 15 B 17 18 19

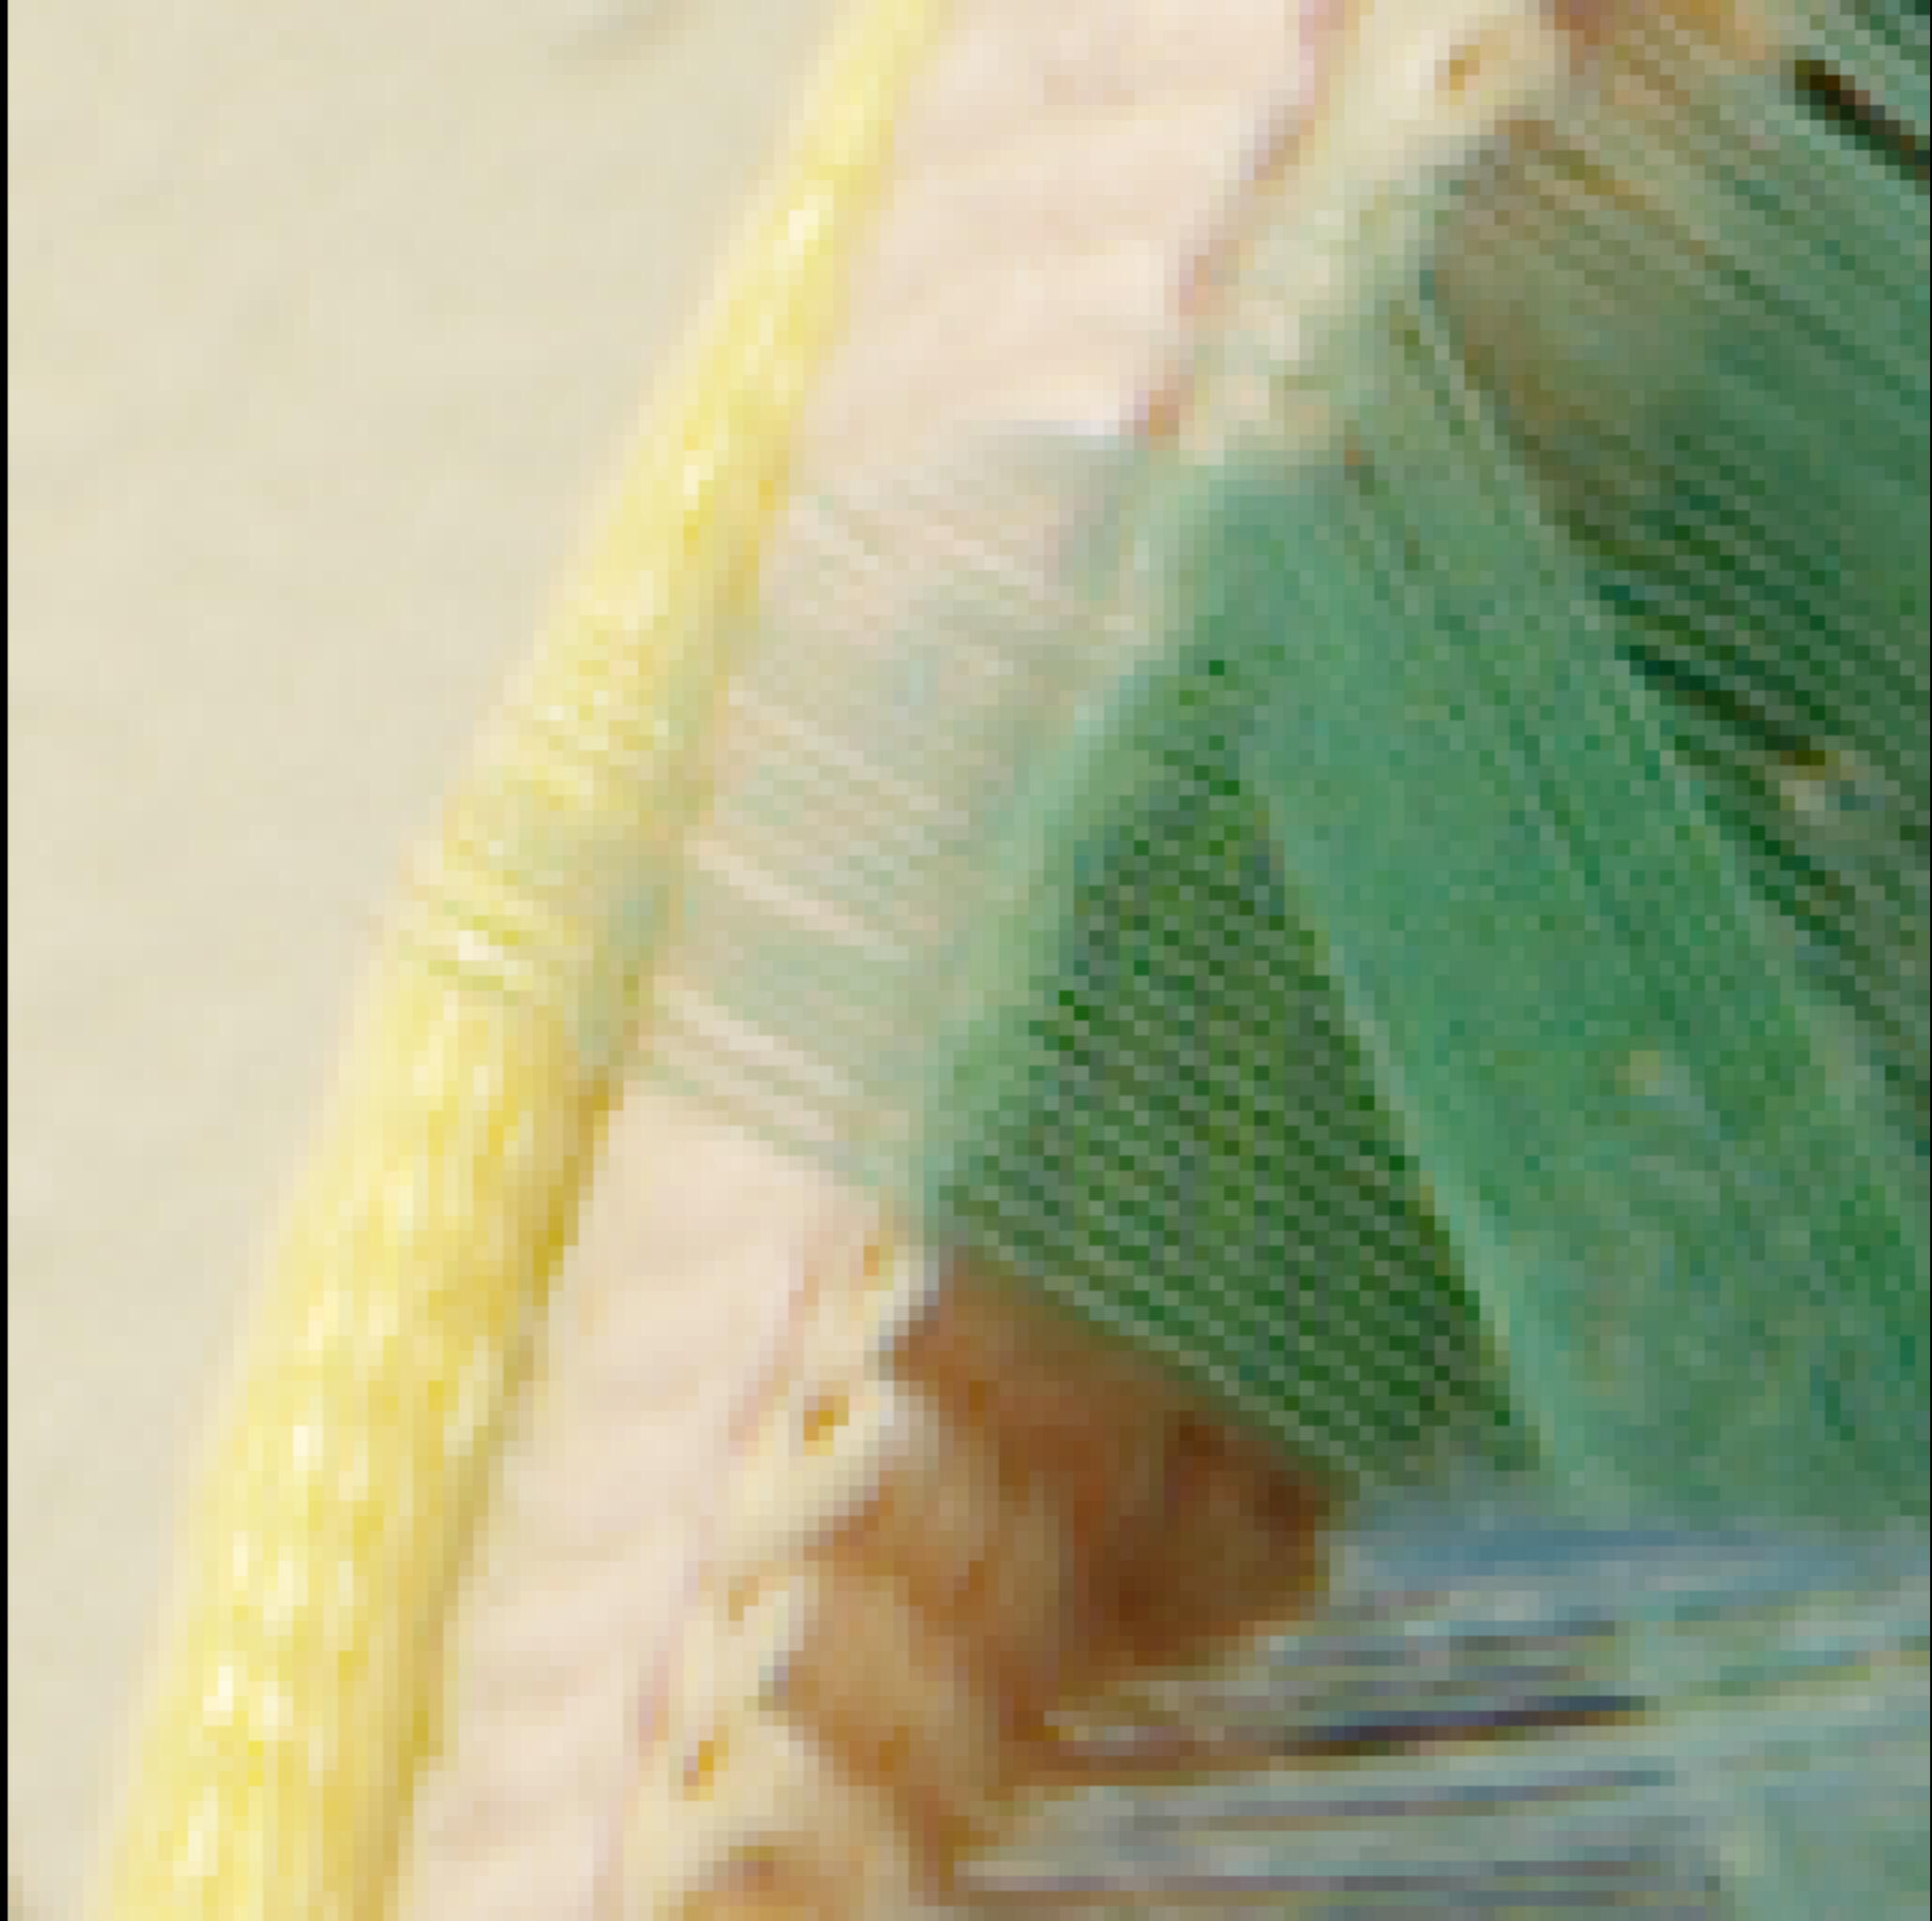


*Q-60R2 Target for
KODAK
Professional Papers*

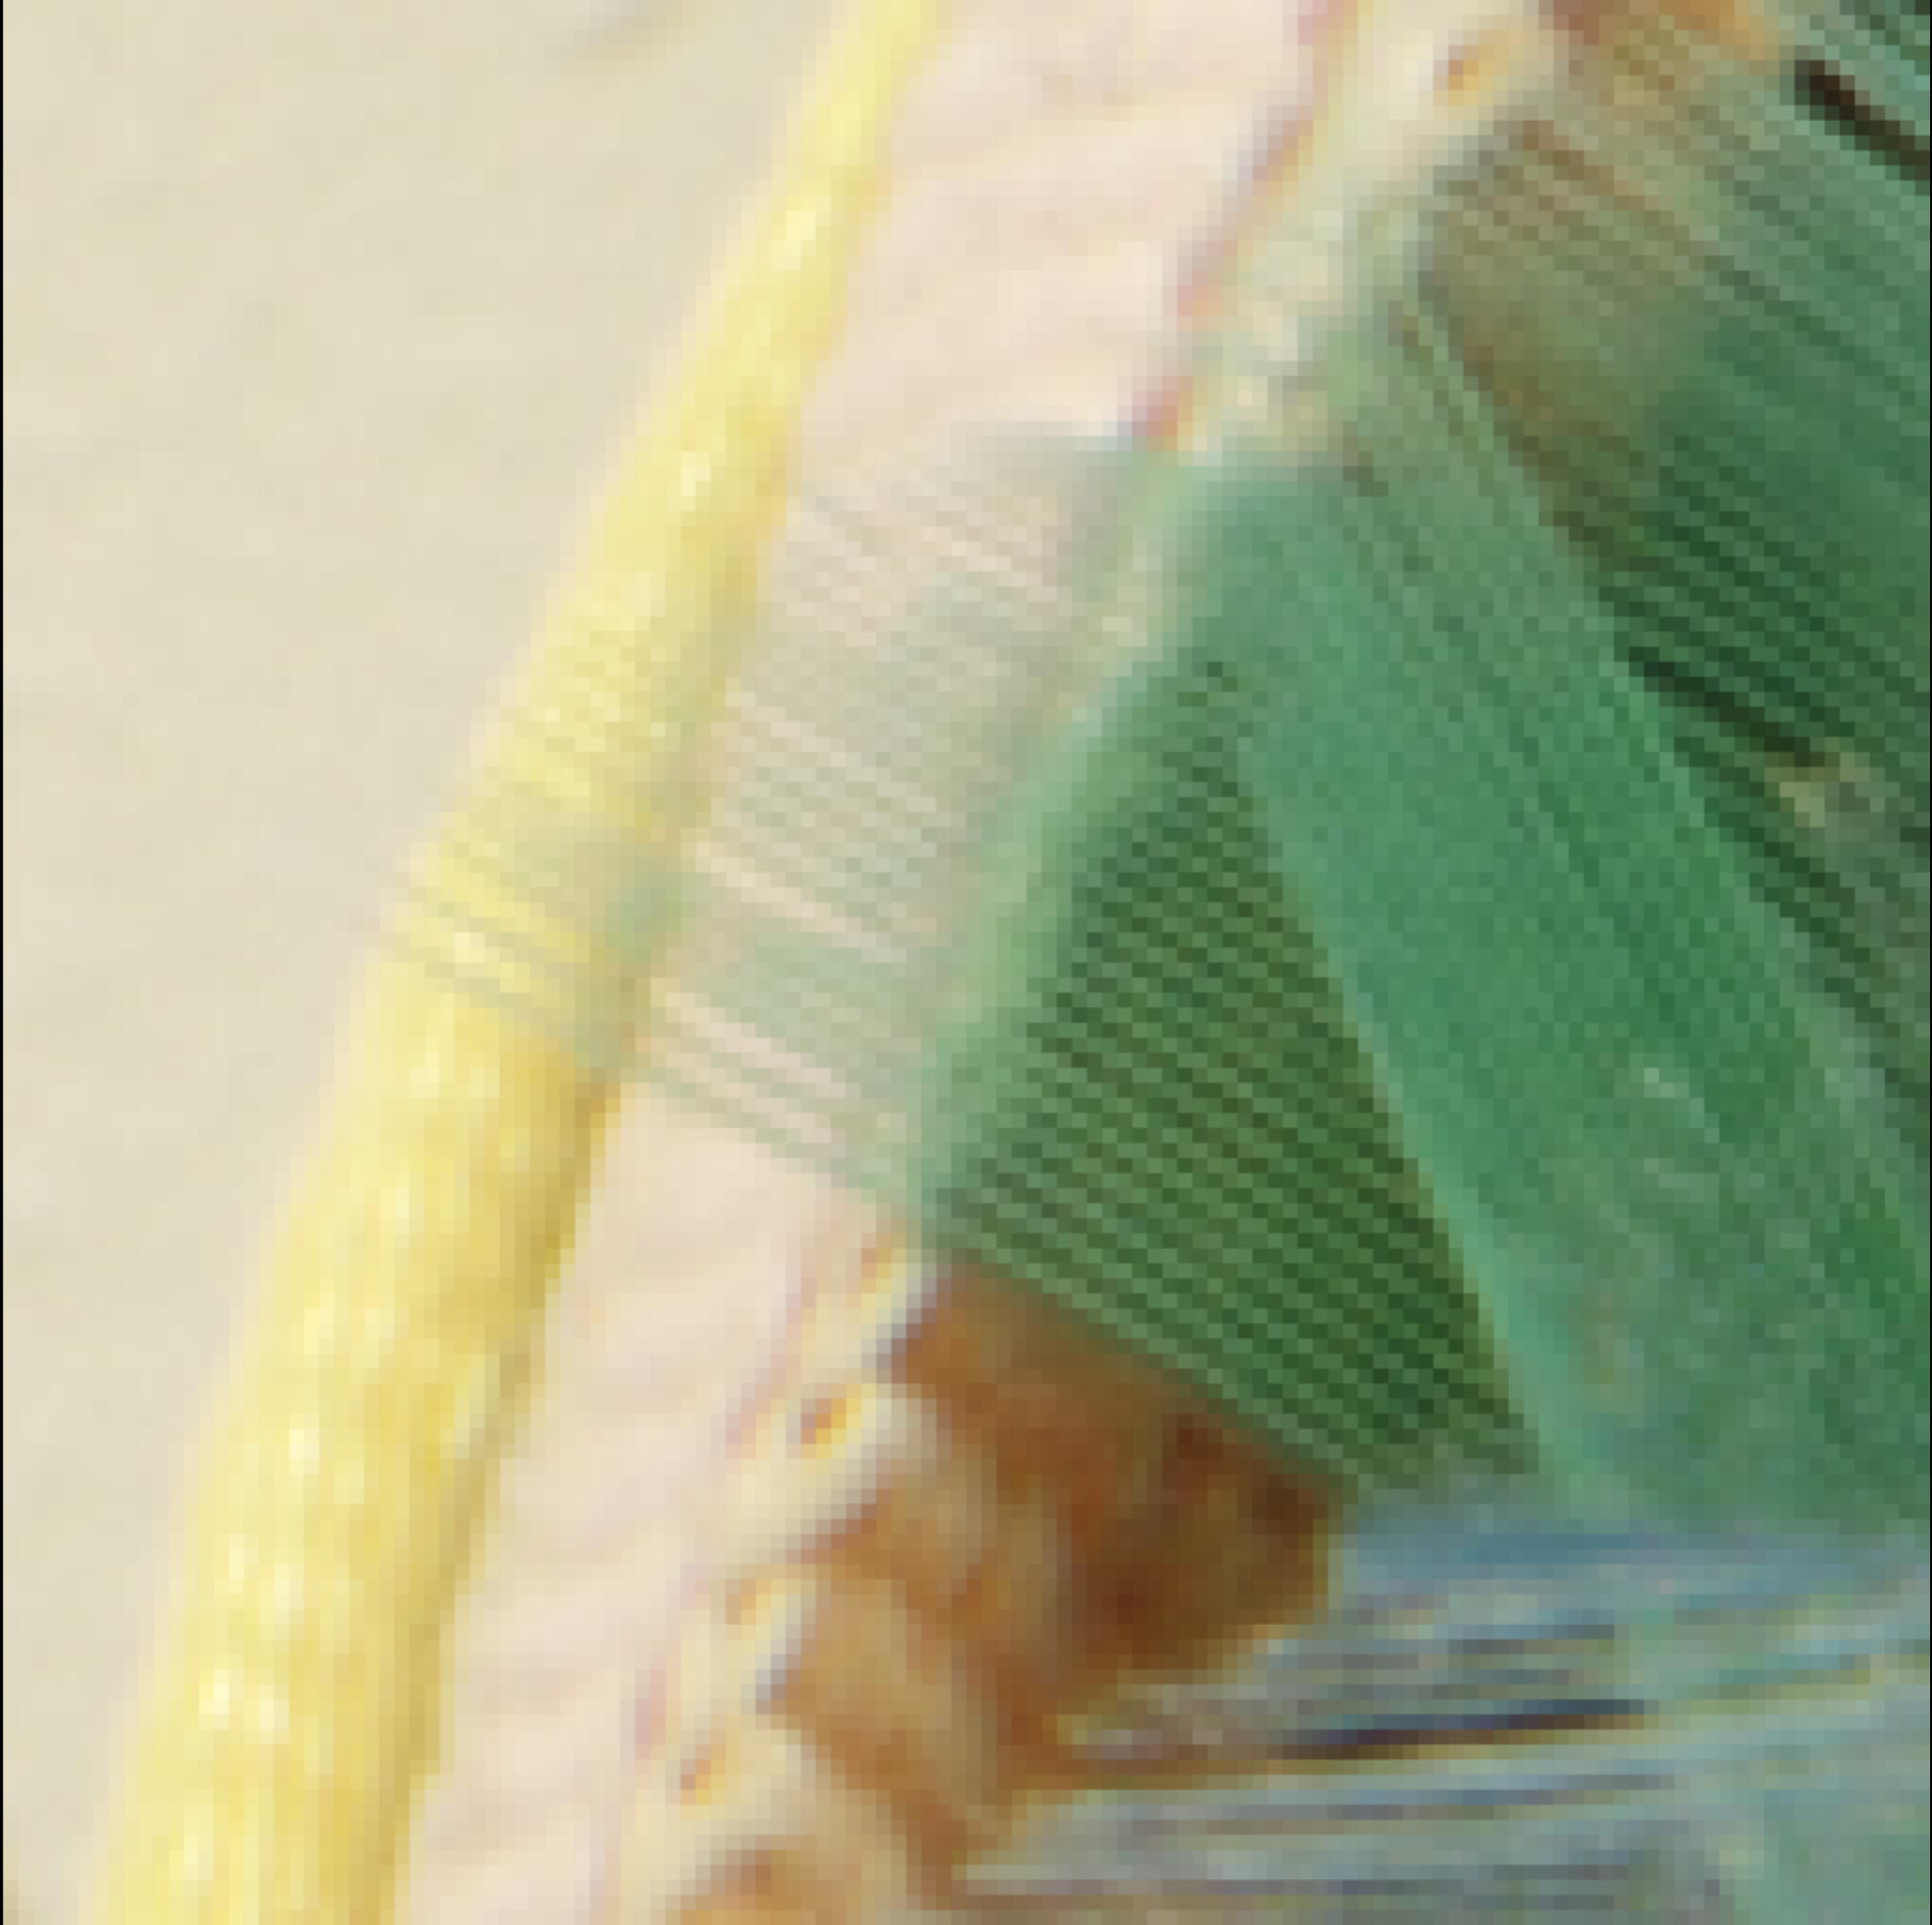
image from www.imaging-resource.com

Fuji X-Trans

drawing

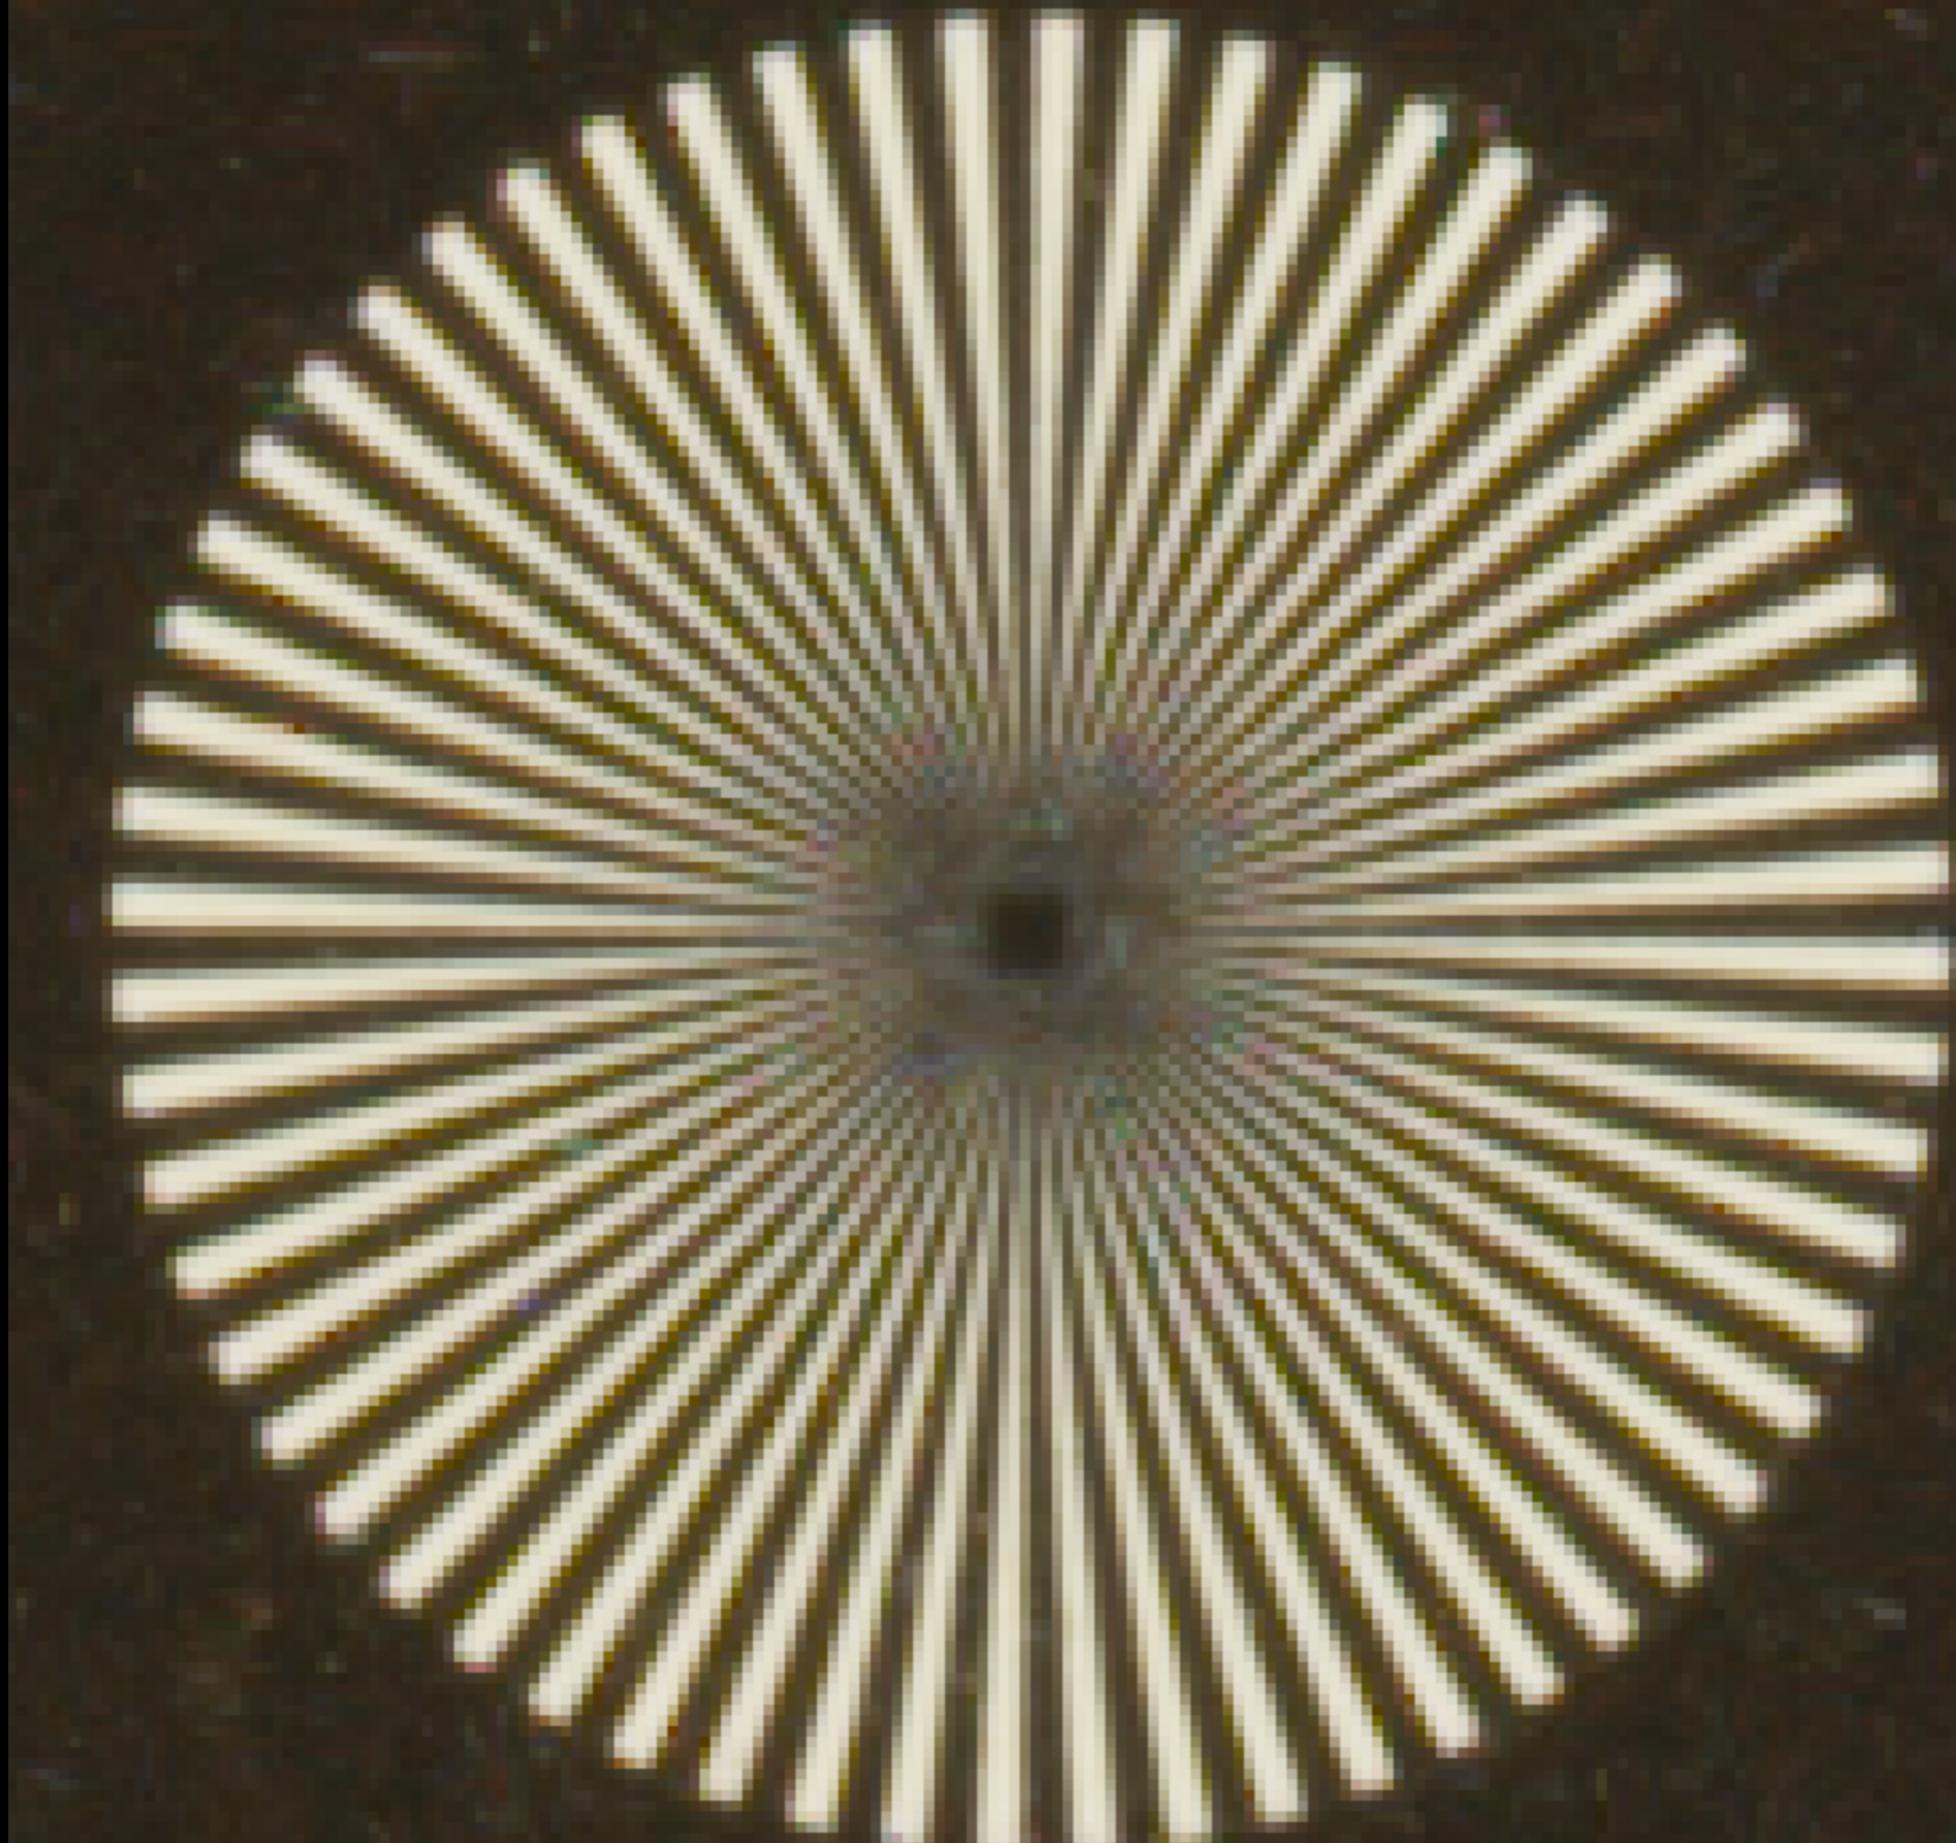


Fuji X-Trans
ours

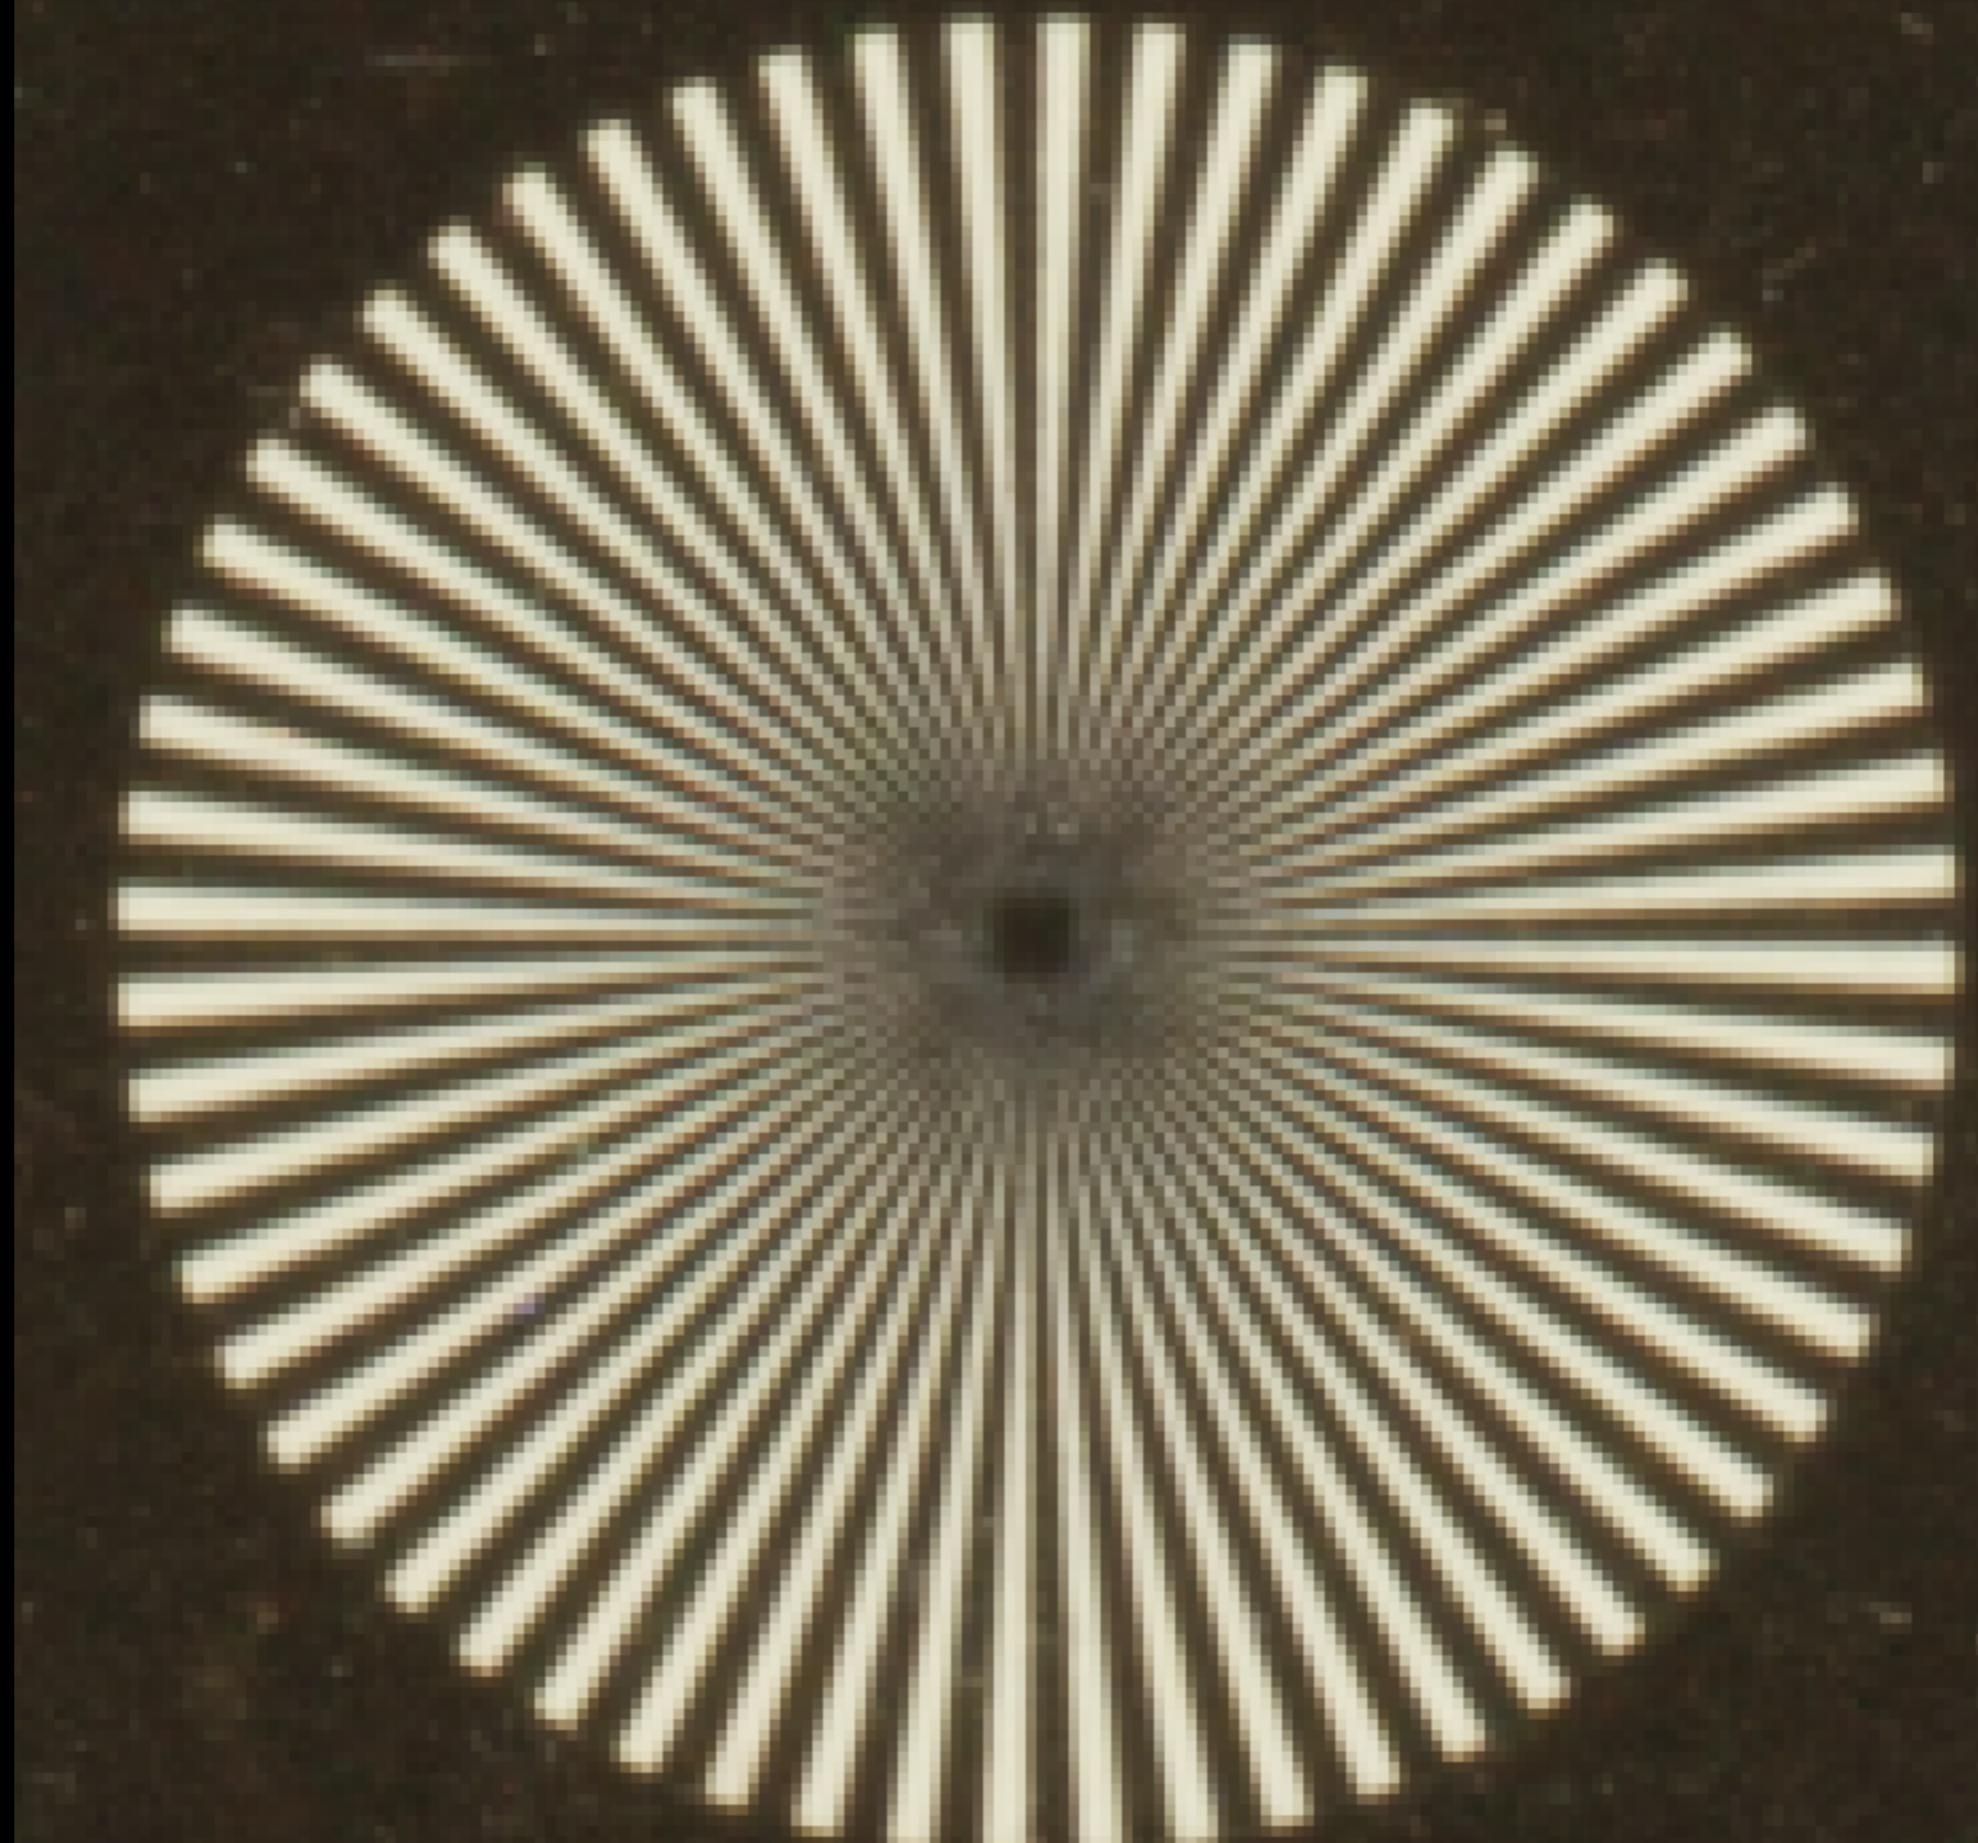


Fuji X-Trans

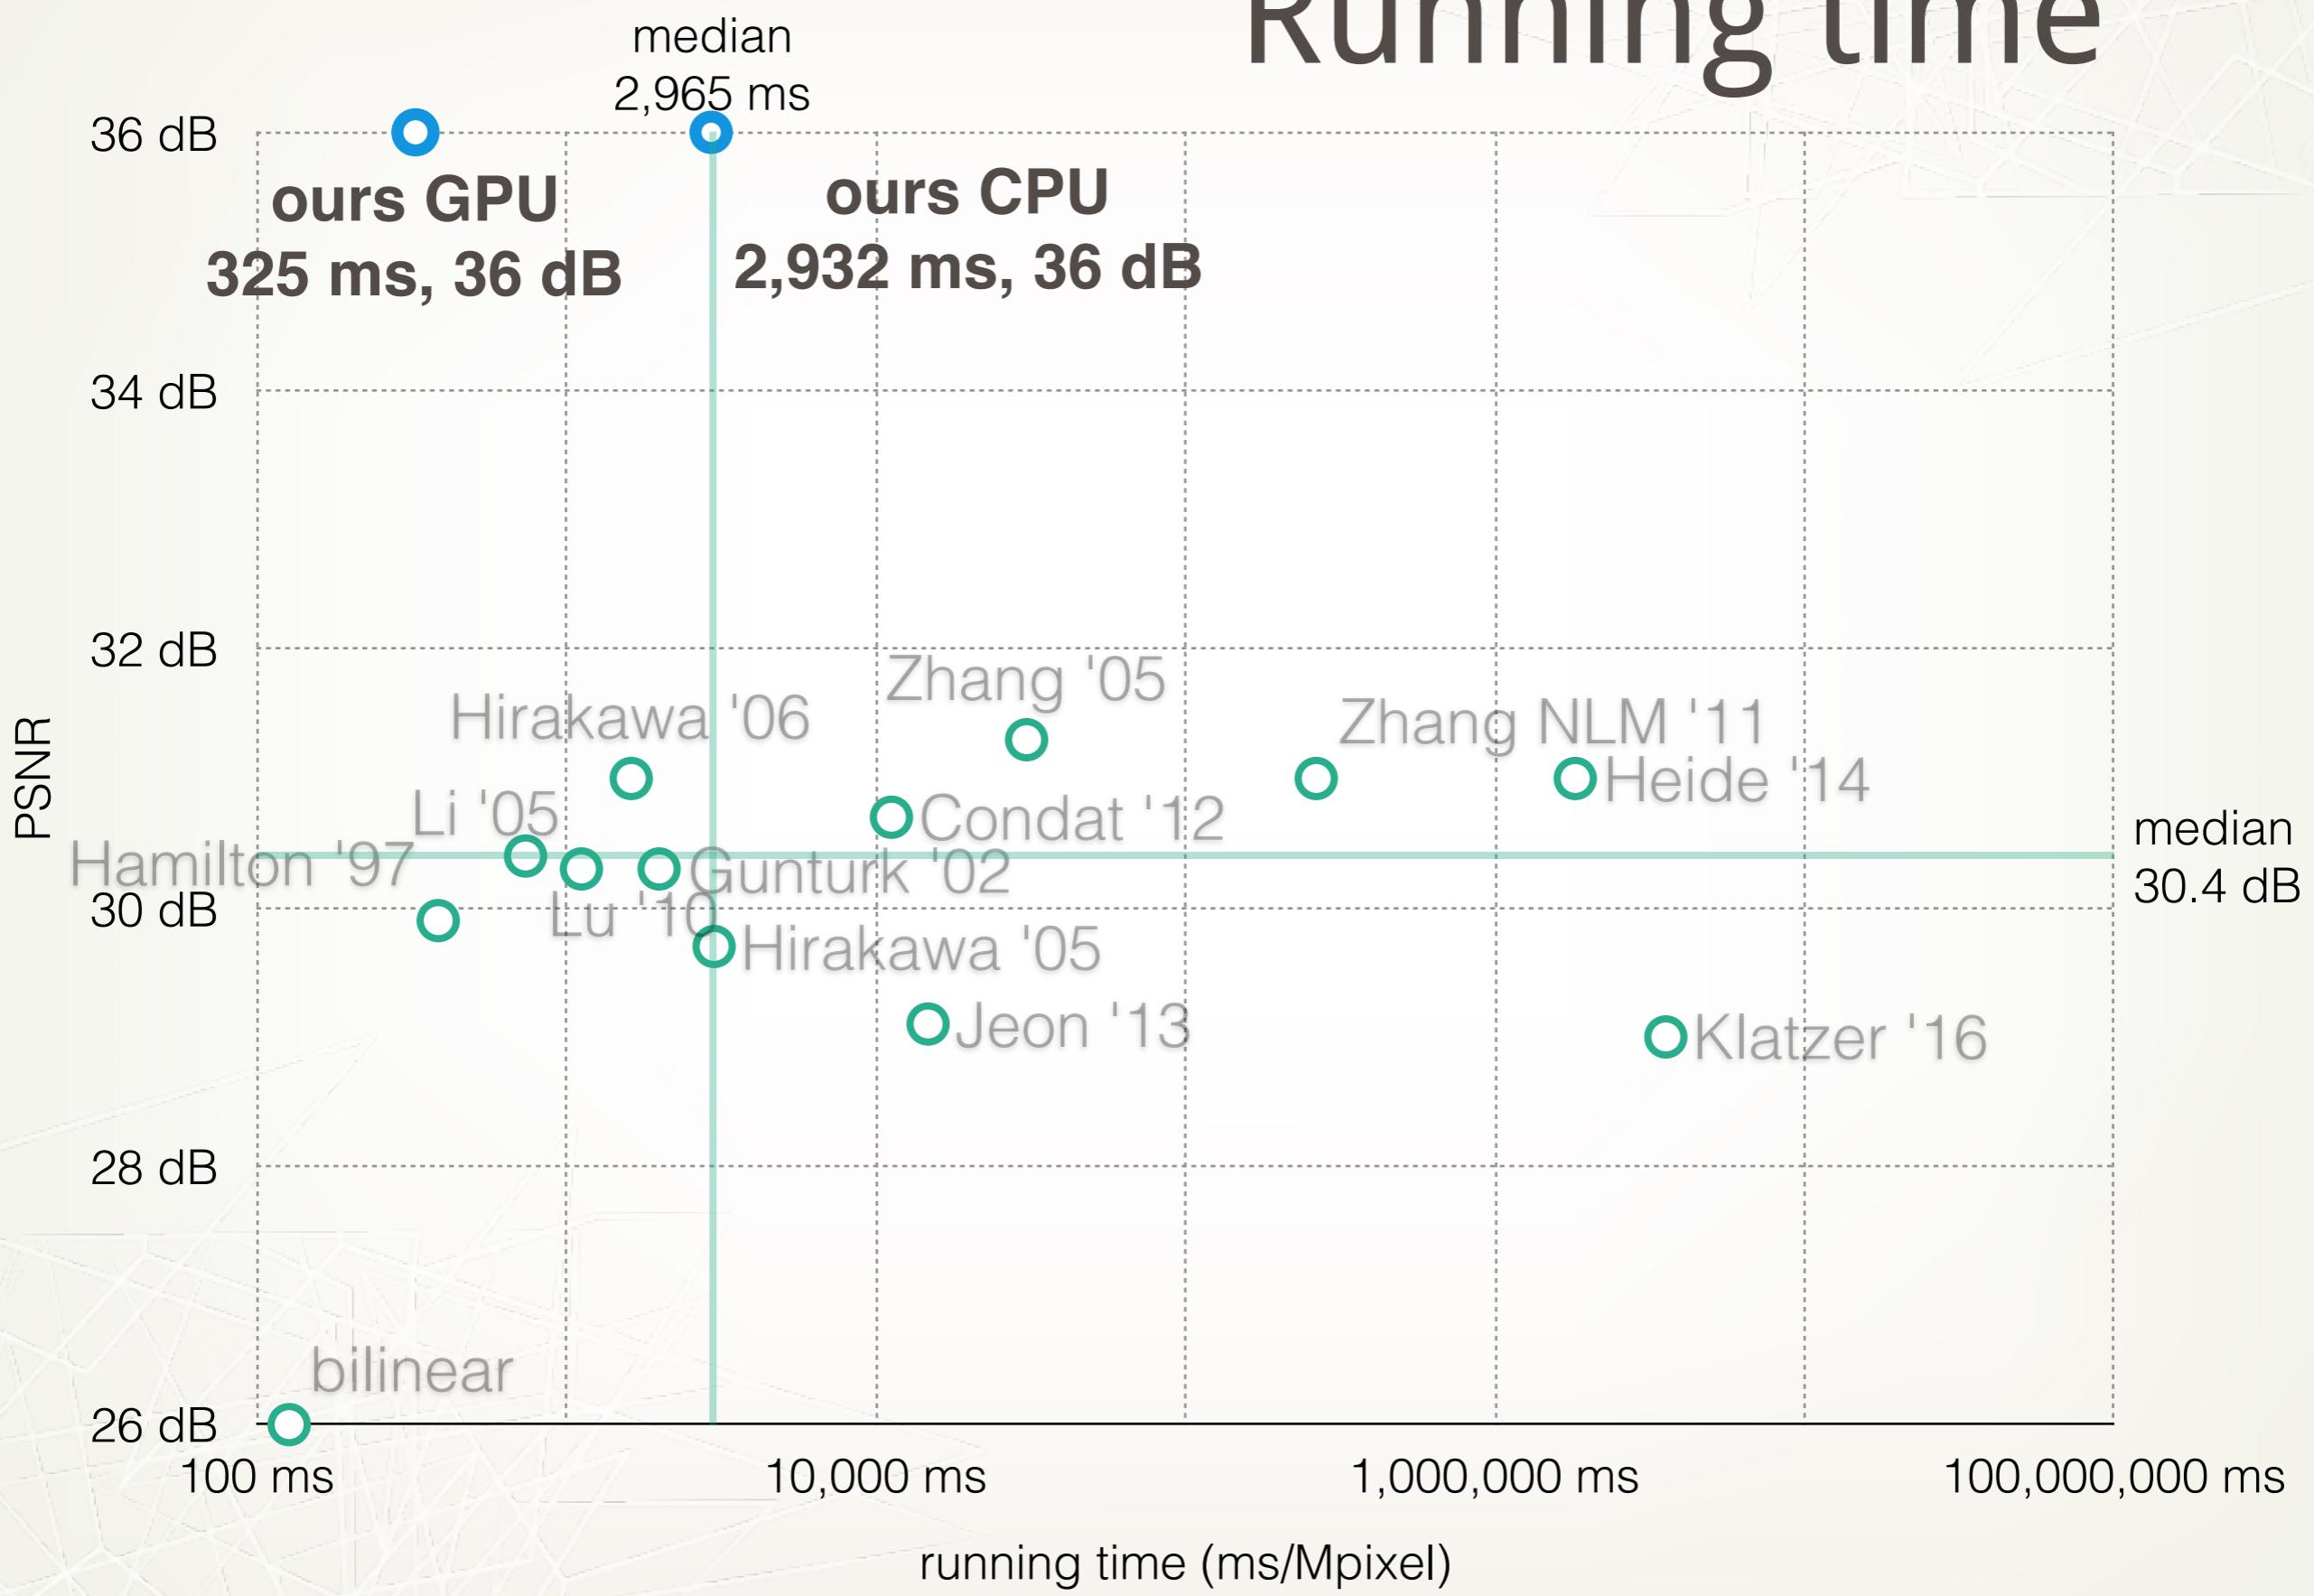
drawing



Fuji X-Trans
ours

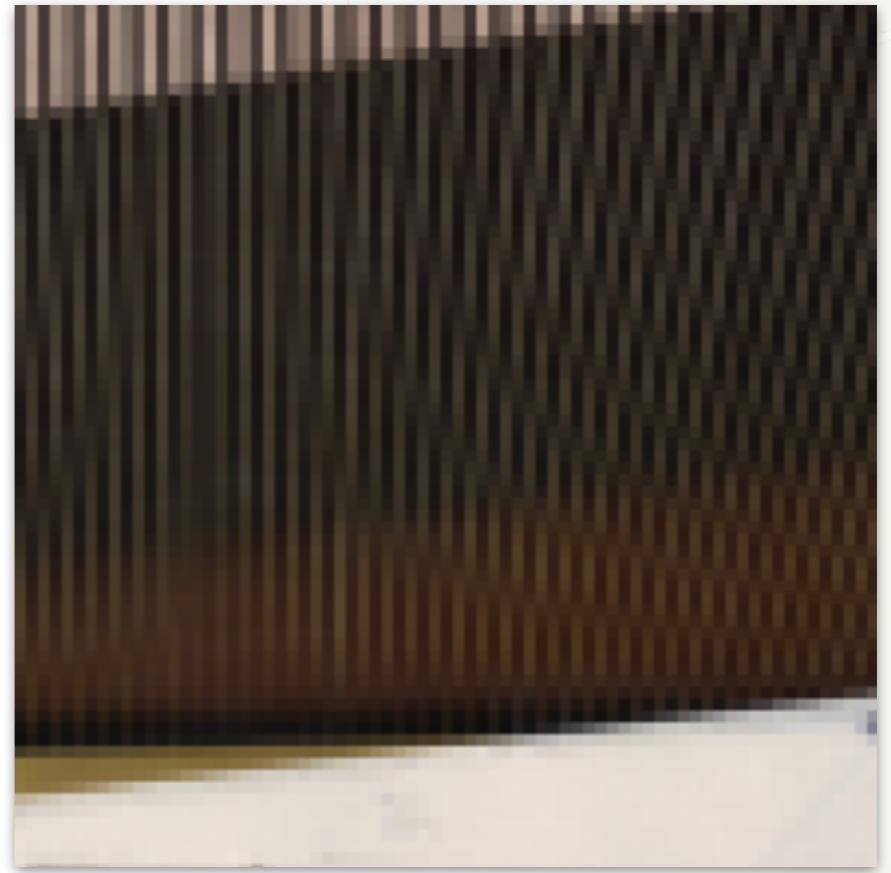


Running time



Limitations and future work

- better ground truth
 - e.g. moiré in the reference
- better metrics
 - HDR-VDP does not capture all the luminance errors



Conclusion

- fast, state-of-the-art demosaicking/denoising
- noise-parametrized network
- three-step process to mine challenging data
- code and data **available online!**

Michaël GHARBI gharbi@mit.edu
code & data: www.mgharbi.com